

Just Delegate It: Comprehensive Scenario Assessment

50 Applied Multiple-Choice Questions on Autonomous Delegation, Brief Engineering & Verification

Assessment Directives: Each scenario assesses your grasp of the Delegation Loop (Define, Delegate, Observe, Intervene, Verify, Accept), the 6-line Delegation Brief, verification rigor (ITCD), and authority boundaries. Evaluate each situation based on course principles, not superficial assistant fluency. Answers and rationales are indexed at the end of the document.

PART I: TASK FRAMING & THE DELEGATION SHIFT

QUESTIONS 1 – 3

1. For the third week running, Omar pastes the same two paragraphs of background before his weekly summary job. What does the course say this is a sign of?

- A. The context belongs in a Project or a saved note he pastes first, so the job needs one message instead of three.
- B. The job is too hard for AI, because a task that needs two paragraphs of setup every single time is not a delegation.
- C. He should switch assistants, because one that needed the background every week has no memory worth using.
- D. He should shorten the background, because two paragraphs is more context than a weekly summary can use.

2. The assistant cited a page for a claim. Ali opened the page and found it says something slightly different. Which decision fits, and why?

- A. Accept, because there is a citation and a page that is slightly different still counts as support for the claim.
- B. Reject the whole result, because one bad citation means none of the others can be trusted.
- C. Ask for another citation and accept that one, because a second page is more likely to say what the claim says.
- D. Correct or Investigate, because a citation is not verification and the sentence did not support the claim.

3. Sana types “What are the best free tools for making a CV?” and gets a good list. She then opens each tool, compares them, and picks one. Which step did she keep that she could have delegated?

- A. Verifying that the tools exist, which the assistant could have done by opening each page and confirming it is still live today.
- B. Deciding which tool to use, since a recommendation from the assistant would have settled the choice for her.
- C. The comparison itself, because the question returned advice and a job asking for a table with links would have moved it.
- D. Choosing free tools, because the assistant could have filtered by price before she had to look at anything.

4. Maryam ran the same brief in two assistants on the same afternoon. One finished in two messages, the other in four. She writes “Assistant A is better.” What does the course say?

- A. She has an observation about two runs on one day, not a ranking, and it belongs in her Delegation Record as such.
- B. She is right, and should use Assistant A from now on, because finishing in fewer messages is the measure that matters.
- C. She should run it ten more times in each assistant, because a ranking needs a larger sample before it can be trusted.
- D. She should change the brief, because a brief that takes four messages in one assistant was not written clearly.

5. A result is wrong in one figure. The reader replies “Redo the whole thing.” Which move would the course prefer, and why?

- A. Stop, because a wrong figure means the job itself was wrong and the outcome needs restating in one line.
- B. Escalate to a stronger model, because a figure that came out wrong means the level was too low for the whole of the job.
- C. Correct: name the error and what should replace it, because the smallest message that fixes it is the intervention.
- D. Take the job back, because a result with an error in it has shown that this job belongs with the reader.

6. A brief has outcome, context, constraints, and deliverable, and the result recommends a paid course without saying it is paid. Which missing line caused this?

- A. Verification, because a line saying what would be checked would have made the assistant flag the price itself.
- B. A longer context line, because the assistant did not know enough about the reader to rule out paid courses.
- C. A stricter deliverable line, because a column for price in the table would have forced the fact into the open for the reader.
- D. Authority, because “say clearly if a course is paid” is a limit on it, and without it AI decided for the reader silently.

7. Usman pastes a notice and asks three questions about it. The assistant answers all three fluently, and one answer describes a deadline that is not in the notice. What did the brief lack?

- A. A shorter notice, because a long paste makes the assistant lose track of which sentence answers which question.
- B. A request for a table, because answers laid out in rows are checked more carefully than answers in paragraphs.
- C. “Use only the text below” and “quote the sentence you used”, so that “not in the source” becomes the right answer.
- D. Web search, because the deadline was probably published somewhere online and the assistant should have gone and found it there first.

8. Aisha's result lists three courses with confident descriptions and no links or dates. What does this tell her?

- A. The assistant probably answered from what it remembered rather than from a search, and she should not assume it looked.
 - B. The courses probably do not exist, because a real course would have produced a link the assistant could show.
 - C. The assistant is broken and she should start a new chat, because search normally attaches a link and a date to every fact.
 - D. Nothing in particular, because confident descriptions are enough to act on when the courses are well known.
-

9. Zara's result has fourteen facts. She has ten minutes. What does the Identify step tell her to do?

- A. Check all fourteen quickly, because a fast pass over everything catches more than a slow pass.
 - B. Check the ones that matter most, such as prices, dates, and the recommendation she will act on.
 - C. Check the first three, because errors cluster near the start of a result where the assistant was still settling in.
 - D. Ask the assistant which ones to check, because it knows which facts it was least sure about when it wrote them.
-

10. Hira asks for a plan first, approves it, and then finds the result used a source that was not in the plan. Which loop step is this concept training, and what is the right response?

- A. Accept, because the result is fine and the extra source only means the assistant was being thorough.
 - B. Define, because a plan that changed means the outcome was never clear enough and should be rewritten from the start.
 - C. Verify, because the right move is to trace the new source, confirm it is good, and move on with the result.
 - D. Observe and Intervene, because every difference between the approved plan and what happened is a question to ask.
-

11. Bilal writes: "Go to two job sites, open the ten newest listings for data entry, copy the salary from each, and tell me the average." The assistant returns a neat table with an average. What is the most likely problem?

- A. The assistant cannot do arithmetic reliably, so an average computed from ten listings should never be trusted.
- B. He gave a recipe, so the result is only as good as the two sites he chose, and nothing in the result tells him so.
- C. He should have asked for a chart instead of a table, because a picture of the salaries would have exposed the spread between the listings.
- D. He forgot to say which currency to use, so the average may mix figures that cannot be added together.

12. A chat answer carried two “unverified” marks. The one-page document made from it carries none. What happened, and what should the reader do?

- A. The facts were verified during formatting, so there is nothing to do and the document can be sent as it is.
- B. Deliverables lose caveats, so find the two marks and put them back before the document leaves the chat.
- C. The document is wrong and should be rejected, because a deliverable that changed cannot be trusted.
- D. Ask for a longer document, because the marks were dropped to fit the page and more room will bring them back.

13. Tariq types: “Can you suggest methods to reduce cloud infrastructure costs?” The model outputs six plausible best practices. Why does this interaction fail to qualify as a delegation?

- A. Because cost-cutting advice is inherently subjective and cannot be measured mathematically.
- B. Because Tariq did not specify which cloud vendor he uses, leading to generic hallucinated answers.
- C. Because the question yielded advice; all research, tool evaluations, and actual execution remain entirely on Tariq.
- D. Because cloud infrastructure requires access via an API connector rather than a conversational chat window.

14. Zoya provides an AI assistant with five strict search keywords and instructs it to inspect only three specific blogs for industry benchmark data. How does the course evaluate this approach?

- A. Excellent prompt engineering, because narrowing the web perimeter prevents model drift and saves token budget.
- B. Flawed delegation, because Zoya dictated the clicks (a recipe), constraining the quality of the findings to her own plan.
- C. Sub-optimal, because blogs are never primary sources and should be globally banned under Constraints.
- D. Valid, provided she enabled deep research mode before supplying the three specific URLs.

15. Hamza constructs a Delegation Brief containing an Outcome, Context, Constraints, and a Deliverable. The output recommends an expensive enterprise software tier. Which omission triggered this failure?

- A. Hamza lacked an Authority line establishing boundaries on autonomous decisions like budget limits.
- B. The Context line was under 200 words, depriving the assistant of his enterprise organization size.
- C. Hamza failed to include an external API connector to his procurement department software.
- D. The model defaulted to commercial bias because Hamza was operating on an enterprise paid subscription.

16. Nida needs to summarize an internal quarterly dispute between two departments. What must she do before pasting the emails into a browser assistant?

- A. Convert the email thread into an encrypted PDF attachment so the model cannot retain it.
- B. Strip all employee names, contact info, and sensitive data, utilizing a temporary/incognito session.
- C. Instruct the model in Constraints: "Delete this chat from your servers immediately after processing."
- D. Upgrade to a paid plan, because consumer free plans retain data whereas commercial tiers guarantee zero retention.

17. Farhan asks an assistant to evaluate a pasted corporate policy. The response provides a verbatim quote supporting a claim, but searching the source yields no match. What occurred?

- A. The assistant hallucinated a fabricated quote because it lacked strict grounding and quoting constraints.
- B. The browser clipboard truncated the source text during Farhan's copy-paste action.
- C. The model paraphrased the text into higher-level business prose to improve clarity.
- D. Farhan experienced a context-window overflow caused by pasting more than two pages of text.

18. Fatima's analysis contains a single erroneous VAT percentage in an otherwise flawless 20-row table. She types: "Discard and run the entire brief again." What is the primary critique?

- A. Re-running the brief consumes excessive API tokens, causing rate-limit throttling.
- B. She violated the intervention rule: apply the smallest message that fixes it (Correct: name the error and replacement).
- C. She should have accepted the table, since a 5% margin of mathematical variance is expected in LLMs.
- D. She should have escalated immediately to a reasoning model rather than repeating the baseline run.

19. Dawood adds: "Before executing, list your proposed steps in order and wait for my go-ahead." What technique is this, and why does it matter?

- A. Few-shot prompting; it provides the model with in-context examples of procedural reasoning.
- B. Plan-first; it leverages the cheapest moment in a job to catch trajectory errors before execution begins.
- C. Guardrailing; it locks the assistant's weights to prevent non-deterministic sampling drift.
- D. Sandboxing; it forces the assistant to execute in a temporary cache without generating permanent files.

20. Mehwish updates her Project Knowledge files, but notices the assistant keeps applying a deprecated refund rule from last quarter. What maintenance step was omitted?

- A. Purging outdated standing instructions and reviewing assistant memory to remove superseded facts.
- B. Deleting and recreating the entire Project container from scratch each week.
- C. Forcing web search mode on, so that live web documents override Project knowledge files.
- D. Downgrading to a temporary incognito chat window to disable background token caching.

21. Kamran receives an AI report citing an official whitepaper. When he opens the link, the whitepaper actually claims the reverse of what the report stated. What course maxim applies?

- A. "Generative AI cannot parse technical whitepapers without a vector database."
- B. "A citation is not verification. Checking the citation is verification."
- C. "When a secondary source disagrees with a primary source, always trust the AI synthesis."
- D. "Models never invent URLs; broken links indicate web server downtime rather than model error."

22. Rabia writes a prompt containing specific vendor UI instructions: "Click the blue Artifacts button and open the Canvas slider." Why is this considered an anti-pattern?

- A. The prompt describes transient vendor buttons rather than the job; a true brief is model-agnostic.
- B. Assistant controls should only be manipulated through keyboard shortcuts rather than written instructions.
- C. Describing interface controls triggers safety filters designed to prevent robotic process automation.
- D. Browser assistants are completely blind to their own interface controls unless vision is enabled.

23. Asad delegates an analysis of competitive pricing. He immediately forwards the raw output to his Chief Operating Officer without opening a single link. Who is legally and professionally accountable?

- A. The AI model provider, under enterprise terms of service and commercial output indemnification.
- B. Asad alone; delegating work to AI never transfers accountability to the machine.
- C. The Chief Operating Officer, who bears ultimate managerial responsibility for departmental deliverables.
- D. Accountability is split 50/50 between Asad and the AI under the course's 10-80-10 operational framework.

24. Sara has been collaborating in a single chat session for over 90 minutes. The assistant begins contradicting its earlier recommendations. What is the recommended operational recovery?

- A. Continuously argue with the assistant in the same thread until it recognizes its logical contradiction.
- B. Instruct the assistant to summarize decisions made so far, copy the summary, and start a fresh chat with it.
- C. Refresh the browser tab and clear the local storage cache while keeping the thread active.
- D. Upgrade the model context tier mid-session to force retroactive re-indexing of all tokens.

25. Imran requests current exchange rates and tuition fees. The output provides exact numbers with zero citations, timestamps, or URLs. What should Imran conclude?

- A. The assistant retrieved live data via internal browser caching without needing external URLs.
- B. The assistant answered from stale training memory; he cannot assume it conducted a search.
- C. The assistant executed a web search but omitted URLs to comply with markdown brevity rules.
- D. The figures are verified because modern foundation models only output numbers when confident.

26. Zainab configures an assistant to execute a recurring weekly competitive brief via a scheduled task. When the Monday morning output arrives, how should she treat it?

- A. Skim the headline and file it immediately, because scheduled tasks have pre-compiled verified pipelines.
- B. Apply full ITCD verification; a job that ran autonomously is not the same as a job that worked.
- C. Verify only if the automated run took more than 5 minutes longer than standard scheduled execution.
- D. Forward it automatically to stakeholders, scheduling a manual verification audit once per quarter.

27. Waleed drafts the Verification line for his brief. Which of the following entries adheres strictly to the course definition?

- A. "Verification: Read through the final output very carefully and make sure it sounds smart."
- B. "Verification: I will check the enrollment page to confirm zero tuition and mark unconfirmed items as unverified."
- C. "Verification: Assistant must self-verify its internal logic against standard Wikipedia citations."
- D. "Verification: Ensure the resulting document looks attractive and fits onto exactly one printed page."

28. Ayesha's assistant begins drafting an essay on enterprise leadership when she specifically briefed it for a spreadsheet comparison of three CRM tools. Which intervention move fits?

- A. Redirect: provide an intermediate paragraph guiding the essay back toward software themes.
- B. Stop: restate the requested outcome in one line, terminating the runaway text generation.
- C. Escalate: switch immediately to a specialized coding model with higher reasoning parameters.
- D. Take the job back: assume generative models cannot distinguish spreadsheets from essays.

29. Daniyal provides an assistant access to his customer database with the rule: "Inspect for duplicate accounts, but do not merge or delete." Why is "Inspect, but do not delete" a foundational boundary?

- A. Database deletion queries require more compute than standard chat tokens accommodate.
- B. Deleting destroys the primary evidence needed to audit, verify, and reverse mistakes.
- C. Consumer LLMs are prohibited from executing SQL queries under international software standards.
- D. Assistants experience emotional degradation when instructed to execute destructive operations.

30. Lubna needs to locate the physical headquarters address of a competitor down the street. She activates an autonomous "deep research mode" that browses for 12 minutes. Why is this a procedural error?

- A. Deep research modes are restricted to academic peer-reviewed scientific discovery.
- B. She violated "need first, then tool"; a single address is settled by a simple search, not a deep report.
- C. Physical addresses cannot be indexed by web search engines without an enterprise Google Maps API.
- D. Deep research modes automatically redact physical addresses to prevent corporate stalking.

31. Rehan runs the same brief on Monday morning and again on Thursday afternoon. The recommendations change slightly. What does the course identify as the cause?

- A. Model weight degradation caused by weekly fine-tuning updates rolled out by the provider.
- B. Either the external world changed (prices/links moved), or the model navigated a different search path.
- C. The browser assistant detected a duplicate session and intentionally diversified the output.
- D. Thursday web server latencies trigger lower-parameter fallback models.

32. Kiran's brief states: "Search for current information rather than answering from what you already know." What role does this instruction play?

- A. It functions as an explicit Constraint forcing the assistant to ground claims in live web search over memory.
- B. It serves as a Context line establishing Kiran's personal philosophical bias toward freshness.
- C. It operates as an Authority grant allowing the model to bypass intellectual property firewalls.
- D. It acts as a Deliverable specification ensuring footnotes are rendered in APA format.

33. Junaid tests an AI by intentionally omitting Constraints and Authority. It produces an articulate summary recommending an advanced, paid Python course to a complete beginner. What did this illustrate?

- A. Without explicit boundaries, the assistant makes silent, unreviewed decisions on your behalf.
- B. Generative AI has an inherent architectural defect making it incapable of teaching programming.
- C. Beginners should never attempt to delegate learning tasks to artificial intelligence.
- D. Paid courses are universally superior to free courses across all educational benchmarks.

34. Mahnoor audits an AI deliverable comparing two enterprise ERP vendors. Vendor A's pricing sheet is current, but Vendor B's cited page is a dead 404 link. What is her correct ITCD decision?

- A. Accept: assume Vendor B's pricing remains equivalent to Vendor A's published rates.
- B. Investigate: trace Vendor B again or find an active primary source before making a decision.
- C. Reject: dismiss both vendors immediately and terminate the entire enterprise software procurement.
- D. Correct: invent a temporary price estimate based on general industry benchmarks.

35. Saad directs an assistant connected to an email tool: "Draft a response to the tenant, and if the terms look fair, go ahead and send it." Why does the course forbid this pattern?

- A. Email plugins introduce latency that often results in corrupted multipart message headers.
- B. "Draft, but do not send" is a crucial boundary; sending creates external legal promises you cannot retract.
- C. Real estate contracts are legally invalid if processed through large language models.
- D. Assistants lack sufficient vocabulary to negotiate lease terms respectfully.

36. Anum needs to extract key obligations from a service agreement. Which two prompt instructions turn pasted text into a rigorous, groundable source?

- A. "Answer as an expert corporate attorney" and "Provide a comprehensive analysis."
- B. "Using only the text below" and "Quote the sentence you used; if not present, state 'not in the source'."
- C. "Synthesize this document into professional Markdown" and "Do not make any mistakes."
- D. "Check your knowledge base" and "Compare this against standard commercial contracts."

37. Haris runs a brief in Model X and Model Y. Model X finishes in 1 message; Model Y takes 3 messages but asks a clarifying question that prevents an omission. What does Haris log in his Delegation Record?

- A. He logs that Model X is objectively superior because minimal turn count is the primary metric of AI IQ.
- B. He logs an observation about this specific run on this day, noting how each model navigated the brief.
- C. He logs a formal benchmark ranking declaring Model Y defective due to redundant message overhead.
- D. He deletes Model X's output because any single-turn response is proof of hallucination.

38. Bushra wants an AI to research project management tools but wants to ensure it never recommends anything requiring credit card setup. Where in the brief does this requirement live?

- A. Deliverable.
- B. Context.
- C. Constraints.
- D. Outcome.

39. Fahad has exactly six minutes before entering a executive steering meeting. His AI brief returned a dense 25-fact briefing. Under the Identify step of ITCD, what does he check?

- A. The grammatical syntax of every sentence to prevent embarrassing typographic errors.
- B. The high-stakes facts: budget totals, critical delivery dates, and the core recommendation he will endorse.
- C. The first three bullet points only, assuming errors always appear at the top of an output.
- D. None; he trusts the assistant's output since it was generated by an enterprise reasoning model.

40. Sana sets up a Project workspace for recurring vendor reconciliation. What belongs in Project Knowledge versus Project Standing Instructions?

- A. Knowledge: standing rules; Instructions: static invoice PDFs.
- B. Knowledge: vendor rate sheets and source files; Instructions: permanent rules and formatting requirements.
- C. Knowledge: temporary session chats; Instructions: web browsing search query strings.
- D. Both Knowledge and Instructions must contain duplicate copies of every prompt sent to the assistant.

41. Adeel watches an assistant follow an approved 4-step research plan. At step 3, it encounters a broken link and silently substitutes a random personal forum post. What loop phase is this, and what is the action?

- A. Observe and Intervene: catch the unauthorized deviation and redirect the assistant to an official portal.
- B. Accept: recognize that autonomous agents must exhibit flexibility when facing 404 errors.
- C. Define: delete the brief and rewrite the initial outcome from scratch.
- D. Verify: wait until the entire 4-step job finishes before checking if the forum post was accurate.

42. Maria asks: "Give me an update on our logistics project." The assistant outputs a vague three-paragraph essay. Maria complains the AI is incompetent. What did Maria fail to do?

- A. She asked a question yielding conversational advice rather than defining an outcome and a deliverable.
- B. She failed to activate the assistant's voice mode to hear the inflection of the logistics update.
- C. She should have attached a spreadsheet of logistics formulas before typing the question.
- D. She forgot to state "Act as a Chief Logistics Officer" at the start of the chat.

43. Why does the course prioritize "The job first, then the machine, then the collaboration, then the techniques"?

- A. Because learning prompt engineering tricks without clear outcome ownership creates neat answers to useless tasks.
- B. Because modern AI architectures change monthly, rendering all technical prompting skills obsolete.
- C. Because enterprise software licenses require outcome documentation prior to user onboarding.
- D. Because collaborative teaming frameworks must be legally registered before accessing browser assistants.

44. In the 6-step Delegation Loop, what are the exact ownership boundaries between human supervisor and artificial intelligence?

- A. Human owns Define and Accept; AI owns Delegate, Observe, Intervene, and Verify.
- B. Human owns both ends (Define, Delegate at start; Verify, Accept at end); AI owns the execution in the middle.
- C. Human owns only the budget payment; AI owns all six stages from Define through Accept.
- D. Human owns Intervene and Verify; AI owns Define, Delegate, Observe, and Accept.

45. Waqas audits an AI claim: “Product A has 98% satisfaction (TechReview Blog, 2025).” He traces the link and finds TechReview was citing a 2021 marketing pamphlet funded by Product A. What ITCD step was this?

- A. Identify.
- B. Trace and Challenge.
- C. Accept.
- D. Define.

46. Huma drafts a brief to evaluate two potential office rental leases. Her Authority line reads: “Authority: None.” What does this command instruct the assistant to do?

- A. It forbids the assistant from using any compute tokens, causing the request to error out.
- B. It instructs the assistant to present factual comparisons only, reserving every recommendation and choice for Huma.
- C. It indicates that Huma lacks managerial clearance to authorize the AI workflow in her company.
- D. It commands the assistant to browse only public government databases without login credentials.

47. Salman delegates a multi-step financial depreciation model. The default fast model produces incoherent calculations across multiple rows. Which intervention move fits?

- A. Redirect: ask the assistant to try being friendlier and more confident in its arithmetic.
- B. Escalate: switch to a stronger reasoning model or decompose the model into sequential sub-tasks.
- C. Stop: cancel the account subscription and file a vendor support ticket for a refund.
- D. Accept: assume financial depreciation is inherently approximate when computed digitally.

48. Nabeel observes that every time an assistant exports research into a polished HTML document, the “unverified” flags disappear. What cognitive/systemic bias explains this?

- A. Deliverables are architected to look authoritative and clean; formatting templates naturally shed caveats.
- B. Browser security sandboxes automatically strip brackets and quotation marks to prevent XSS attacks.
- C. The underlying foundation model internally verified all claims during the code rendering compilation.
- D. HTML tables lack the unicode font glyphs required to render caveat characters.

49. Saima's assistant stops executing before step 1 and asks four redundant background questions that were explicitly answered in the Context line. Which intervention move should she use?

- A. Stop and Redirect: state in one sentence that context is provided and direct it to begin step 1 immediately.
- B. Re-paste the entire 6-line brief five consecutive times in the chat to force attention weighting.
- C. Apologize to the assistant and spend twenty minutes answering all four redundant questions in depth.
- D. Delete the browser cookies and restart the operating system to clear memory registers.

50. What is the foundational shift separating a passive AI user from an autonomous delegator?

- A. Moving from asking for conversational advice that keeps the work on you, to handing over checkable work orders.
- B. Purchasing high-tier enterprise GPU subscriptions to bypass standard chat throttles.
- C. Memorizing 500-token prompt templates and automated shell scripts for every daily task.
- D. Allowing the artificial intelligence to communicate directly with clients without human oversight.

SCORING & COMPETENCY RUBRIC

45 – 50 Correct

Master Delegator: Full command of the 10-80-10 rule, brief boundaries, grounding discipline, and verification rigor.

38 – 44 Correct

Practitioner: Solid brief structuring, occasional vulnerability to deliverable caveat-dropping or recipe-style micromanagement.

Below 38 Correct

Needs Reinforcement: Tendency to confuse asking with delegating, or trusting fluent citations without verification.

#	Ans	Core Concept Tested	Operational Rationale & Course Citation
1	A	Project vs. Chat Context	Repetition asks for context storage. The 3rd time context repeats, it belongs in a Project or saved note (Concept 8).
2	D	Verification (ITCD)	A citation is not verification; checking the citation is. Slight discrepancies require Correct or Investigate (Concept 10).
3	C	Asking vs. Delegating	A question yields advice; asking for a comparative table with links delegates the synthesis work (Concept 1).
4	A	Model Agnosticism	Two runs on one day represent single observations for a Delegation Record, not a permanent vendor ranking (Concept 11).
5	C	Intervention Moves	The smallest message that fixes the drift is the intervention: "Correct" names the error and replacement (Concept 9).
6	D	Authority Line	Without an Authority boundary ("say if paid"), the AI makes silent decisions on your behalf (Concept 3 & 7).
7	C	Grounding Instructions	"Use only the text below" and "quote the sentence" force fidelity and make "not in the source" a valid answer (Concept 5).
8	A	Search vs. Memory	Confident assertions lacking dates/links indicate pre-trained memory recall rather than active web search (Concept 4).
9	B	Identify Step (ITCD)	Identify focuses checking where errors cost the most: prices, dates, metrics, and actionable recommendations (Concept 10).
10	D	Observe & Intervene	Any silent deviation from an approved plan represents an authority failure requiring observation and intervention (Concept 7).
11	B	Outcome vs. Recipe	Dictating specific sites limits the result to the quality of the user's plan without revealing better alternatives (Concept 2).
12	B	Deliverable Caveats	Deliverables naturally shed caveats to look clean; the supervisor must audit and restore unverified tags (Concept 6).
13	C	Work Hand-Off	Asking for general tips produces conversational advice, leaving all analytical legwork on the human (Concept 1).
14	B	Micromanaging Steps	Specifying exact clicks handcuffs the model's exploratory potential and confines findings to user biases (Concept 2).
15	A	Boundary Governance	Omitting Authority permits the assistant to silently select expensive software tiers without human sign-off (Concept 7).
16	B	Privacy & Data Hygiene	Strip PII/sensitive data and utilize temporary/incognito sessions; consumer browser tools are not filing cabinets (Concept 5).
17	A	Source Fidelity	LLMs fabricate quotes unless constrained by strict source grounding ("quote the sentence used") (Concept 5).
18	B	Intervention Sizing	Never redo the entire job for one wrong number; name the specific error and supply the fix (Concept 9).
19	B	Plan-First Method	Reviewing steps before execution is the cheapest moment to catch misunderstandings and redirect trajectory (Concept 7).
20	A	Project Maintenance	Stored context is a responsibility; audit Projects and memory monthly to delete obsolete guidelines (Concept 8).
21	B	Rigor over Citations	A citation is merely a pointer; human verification requires opening the link and validating the claim (Concept 4 & 10).
22	A	Model Agnosticism	A good brief describes the job, not vendor UI buttons which change unpredictably across software versions (Concept 11).
23	B	Accountability Invariant	Delegating work to AI never transfers accountability; the professional name on the deliverable remains yours (Concept 10).
24	B	Context Freshness	Long sessions degrade; asking for a summary and pasting it into a fresh chat restores fidelity (Concept 5).
25	B	Verifying Web Retrieval	Current claims appearing without traceable links or dates indicate recall from pre-trained weights (Concept 4).

#	Ans	Core Concept Tested	Operational Rationale & Course Citation
26	B	Scheduled Autonomy	Scheduled runs still demand human verification; a job that ran is not identical to a job that succeeded (Concept 8).
27	B	Actionable Verification	Effective verification lines state exactly what external source will be audited and demand “unverified” flags (Concept 3).
28	B	Intervention: Stop	When an assistant pursues a completely false outcome, immediately stop it and restate the objective (Concept 9).
29	B	Authority Guardrails	Deleting records destroys audit trails and source evidence necessary for human verification (Concept 7).
30	B	Need First, Then Tool	Deep research modes are for synthesising complex multi-source reports, not answering single facts (Concept 4).
31	B	Variance Dynamics	Outputs vary because the live web moves or non-deterministic search sampling explores different sources (Concept 8).
32	A	Constraints Mechanics	Instructing the model to search forces it to override static pre-training memory with live external retrieval (Concept 4).
33	A	Silent Decision Trap	Omitting authority boundaries results in AI silently resolving user trade-offs without notification (Concept 3 & 7).
34	B	ITCD: Investigate	Broken links or contradictory evidence require an “Investigate” decision to trace secondary sources (Concept 10).
35	B	Action Permissions	“Draft, but do not send” prevents autonomous agents from binding you to external commitments (Concept 7).
36	B	Source Grounding	“Using only text below” and “quote sentence” turn arbitrary copy-paste into an auditable source (Concept 5).
37	B	Comparative Testing	Cross-assistant tests document behavioral observations on a given day, not absolute vendor scorecards (Concept 11).
38	C	Constraints Allocation	Hard requirements, negative boundaries, and financial limits belong squarely in Constraints (Concept 3).
39	B	ITCD: Identify	Time-constrained verification must prioritize high-stakes claims: money, dates, and actionable counsel (Concept 10).
40	B	Project Architecture	Knowledge stores reference documentation; Standing Instructions house persistent behavioural rules (Concept 8).
41	A	Plan Drift Detection	Unapproved deviations from an agreed plan must be caught via Observation and redirected immediately (Concept 7).
42	A	Outcome Definition	Vague queries yield conversational chat; specifying deliverables produces structured, checkable artifacts (Concept 1 & 6).
43	A	Pedagogical Hierarchy	True competence begins with problem framing and outcome ownership, long before button mastery (Foundations).
44	B	10-80-10 Split	Humans own the initial 10% (Define/Delegate) and terminal 10% (Verify/Accept); AI runs the middle 80% (Thesis).
45	B	ITCD: Trace & Challenge	Tracing the citation revealed a secondary source repeating stale, biased commercial claims (Concept 10).
46	B	Restricting Authority	An Authority line of “None” requires the assistant to stop at analysis and leave decisions to the human (Concept 7).
47	B	Intervention: Escalate	Complex analytical failures call for model escalation (reasoning/thinking modes) or task decomposition (Concept 9).
48	A	Format vs. Caveats	Polished documents possess a systemic bias toward clean presentation, shedding messy “unverified” tags (Concept 6).
49	A	Intervention: Redirect	A single concise sentence pointing the assistant back to provided context stops unnecessary stalling (Concept 9).
50	A	The Delegation Shift	Delegation transforms AI from an advisory chatbot into a managed system producing checkable deliverables (Concept 1).