

AI Fluency (The 4Ds): Comprehensive Assessment

50 Scenario-Based Practice Problems · Foundations to Agent Factory

Examiner's Edition

1. Why does the course open by saying a new colleague and an AI assistant are alike in one way and different in another?

- A** Because AI is faster than a colleague, so the same instructions give results in seconds and there is no need to plan the work in advance.
- B** Because a colleague can refuse an unreasonable request, and knowing AI will attempt anything is what makes delegation the first competency.
- C** Because both need context you forgot to give, but AI keeps nothing between chats, so what a person hears once you must supply every time.
- D** Because a colleague improves with feedback over many months, while AI improves inside a single session and then stops improving for good.

2. The course says AI can automate work but cannot automate accountability. Which situation illustrates that most directly?

- A** A model produces different answers to the same question on different days, so a team cannot reproduce a result shown to a client last week.
- B** A coding assistant introduces a security flaw, an engineer ships it, and the engineer and the organization still own the consequences.
- C** An assistant refuses a request that falls outside its safety guidelines, so the work must be completed by a person instead.
- D** An agent escalates far more cases to humans than expected, which erases most of the time the team hoped automation would save.

3. A writer has spent forty minutes steering an assistant through a 600-word article she could have drafted in thirty. The final piece is good. What does the course say she should conclude?

- A** The process is failing even though the product is fine, and her moves are a new performance description, another tool, or taking it back.
- B** The session succeeded, because a good final product is what matters, and the time spent steering is a fair cost of learning any tool.
- C** She has a description problem, so the fix is to spend more time up front writing a much fuller brief before the AI drafts anything.
- D** She has a delegation problem that can only be solved by splitting the article into smaller pieces and assigning each one to the AI.

4. A clinic manager opens an assistant and types 'build me a system for following up with patients who miss appointments.' The output looks impressive but nobody can use it. Which competency failed, and why?

- A** Delegation, because nobody had decided who may be contacted, after how long, in what tone, and which decisions a human must keep.
- B** Discernment, because she accepted a polished answer without checking any of its factual claims against a reliable source.
- C** Description, because the request was too short, and a longer, better-worded prompt would have produced a usable answer on the first try.
- D** Diligence, because patient follow-up is regulated work and a consumer AI service was never approved for use with clinical records.

5. A lecturer pastes student names, grades, and disciplinary notes into a consumer AI service the university never approved. The feedback it drafts is excellent. What does the course conclude?

- A** The use is irresponsible whatever the output quality, and the failure is creation diligence: the wrong data went into an unapproved tool.
- B** The use is acceptable because he reviewed every single comment before sending it, which satisfies deployment diligence.
- C** The failure is transparency diligence, and disclosing the AI's role to the students afterwards would resolve the problem completely.
- D** The failure is discernment, because he should have judged whether the drafted comments were accurate before recording them.

6. Where does the course place Diligence in the 10-80-10 Rule, and why does the placement matter?

- A** In the final 10%, because verifying and vouching for the work is the last gate before anything is published, sent, or acted on.
- B** In the first 10%, because decisions about data, tools, and approvals all have to be settled before any AI is involved at all.
- C** Across all 100%, because it surrounds the whole workflow rather than sitting at any one stage, and is never a final checkbox.
- D** In the middle 80%, because it is exercised continuously while you steer the AI and give feedback on each round of output.

7. Ayesha's team tests a reconciliation agent against past reconciliations the firm already trusts. An accountant also reviews some matches that appear to have succeeded, not only the failures. Why review the successes?

- A** To measure how much faster the agent is than the manual process, which is the number partners want before approving rollout.
- B** Because a system can look safe when it fails without saying so, and a silent, repeated error hides among successes that look fine.
- C** Because reviewing successes satisfies transparency diligence, requiring clients be shown examples of correct AI-assisted work.
- D** Because accuracy on the cases it got right decides whether the escalation threshold sits at the right level.

8. The framework defines agency as configuring AI to independently perform future tasks, including for others, on their behalf. Which detail separates agency from automation most sharply?

- A** Agency needs a more capable model, while automation runs well enough on the smaller and cheaper models most people use.
- B** Agency means the AI picks its own goals instead of the goals a person set, which makes it hard to supervise.
- C** Agency always connects the AI to outside tools and live data, while automation only works on text pasted into the chat.
- D** You are not in the room while work happens, and the person served may not be you, so judgment must be built in beforehand.

9. An AI recommends Vendor A, and the recommendation is sound. On inspection it turns out the AI believed Vendor A ships a feature that does not exist. What does the course say about this situation?

- A** The recommendation stands, because discernment judges the usefulness of the output, and the output was the right call.
- B** This is a hallucination and therefore a Diligence failure, since checking vendor claims sits with whoever deploys the decision.
- C** The AI's stated reasoning is a true record of internal computation, so a false assumption proves the conclusion was random.
- D** A right answer resting on a wrong assumption will not stay right, so judging the result includes judging the case made for it.

10. You tell an assistant: 'do not agree with me just to be polite, and flag anything you are uncertain about.' The course says this small habit is a scaled-down version of what?

- A** Product description, since you are specifying the qualities the final output must have before you accept it as finished.
- B** Process description, since you are prescribing the order of the steps the AI should follow while it works through your request.
- C** Context engineering, since you are adding standing information to the conversation that the AI will keep using on later turns.
- D** Performance description, since you are defining behavior you have not seen yet, and at scale it becomes standing behavior rules.

11. A team ships an AI tutor. Every individual answer it gives is factually correct, yet students perform worse than before. Which kind of discernment would have caught this, and what does it look at?

- A** Product discernment, which checks each output for factual errors, missing requirements, contradictions, and doubtful claims.
- B** Process discernment, which asks whether working alongside the AI is still paying off for the team that builds and steers it.
- C** Performance discernment, which asks whether the AI's independent behavior serves users, showing up only across many cases.
- D** Creation discernment, which checks that training data and connected systems suited the audience the product is aimed at.

12. A reader compares two requests. The first is 'Summarize this report.' The second names the audience, the time they have, the topics to cover, the format, and what to leave out. What is the lesson?

- A** The longer request works because length itself signals seriousness, so adding more words reliably improves the answer.
- B** The second one is not more intelligent, it is more complete, and completeness usually matters far more than clever wording does.
- C** The second one works because it uses the vocabulary of the field, tapping the model's deep knowledge of financial reporting.
- D** The pair shows that summarizing is a poor use of AI, and that the reader should have given the reading to a human colleague.

13. An analyst is drafting an industry overview. She asks the AI to suggest 5 possible structural outlines, selects one, writes the core thesis herself, and instructs the AI to expand the section drafts. Which competency is being executed?

- A Automation bias, because she relies on the AI to structure her thoughts before she has written the prose.
- B Task delegation, because she intentionally divided the workflow into parts that fit machine breadth versus human judgment.
- C Performance description, because she instructed the AI on how to behave during an analytical drafting session.
- D Creation diligence, because she prevented confidential company research from leaking into public training sets.

14. A software agency team leader needs to select an AI tool for a complex legacy refactoring project. Instead of picking the fastest chat assistant, he chooses an agent-capable coding tool that reads project file trees. What aspect of Delegation does this reflect?

- A Problem awareness, because he recognized that code refactoring involves architectural debt and high financial stakes.
- B Platform awareness, because he understood tool capabilities and matched the multi-file software task to an agent environment.
- C Process description, because he specified the order in which test suites must be written before modifying code.
- D Deployment diligence, because he implemented an automated CI/CD pipeline to gate pull requests.

15. A marketing executive asks an AI: 'Tell me which marketing campaign strategy we should adopt for our Q4 budget.' The model outputs a definitive strategy which the executive immediately approves. What core delegation error occurred?

- A The executive failed problem awareness by delegating an executive business decision that requires organizational accountability.
- B The executive failed platform awareness because marketing strategies must only be drafted by search-enabled reasoning models.
- C The executive failed process description by not requesting a SWOT matrix before receiving the strategic recommendation.
- D The executive failed transparency diligence because customers will not know AI helped design the marketing campaigns.

16. The course contrasts the script writer, the co-creator, and the director. In which scenario is a professional acting strictly as a co-creator?

- A An operations lead prompts an AI to extract invoice numbers, vendor names, and tax totals from 50 scanned receipts into CSV.
- B A founder spends an hour exchanging ideas with an AI, challenging assumptions, testing counter-arguments, and sharpening a pitch.
- C A product manager configures a customer support bot with instructions and rules to resolve returns autonomously overnight.
- D A researcher writes a prompt script that scrapes ten academic papers and compiles an unedited bibliography table.

17. A team sets up an autonomous invoice-chasing agent. In their specification, they mandate: 'Invoices exceeding \$5,000 or accounts disputed by a client must be escalated to Sarah in Finance.' In Agent Factory terms, what is this rule?

- A A prompt engineering trick designed to prevent token exhaustion during accounting runs.
- B An escalation boundary established during Task Delegation that defines the limits of the Digital FTE's authority.
- C A process discernment checkpoint to evaluate whether the accounting agent is saving sufficient manual labor hours.
- D A product description constraint that regulates the stylistic tone and length of payment reminder emails.

18. Why does the course emphasize that 'domain expert first, AI delegator second' is the fundamental operating posture?

- A Because domain experts are the only individuals permitted by compliance officers to write system prompts.
- B Because without domain expertise, one cannot define what success looks like, where risks lie, or what must never leave human hands.
- C Because modern foundation models will refuse to answer technical prompts unless academic credentials are verified in the context.
- D Because prompt engineering has been fully automated, leaving domain verification as the only task AI cannot execute.

19. A manager requests: 'Review this pull request for correctness first, security second, and naming style last. Do not spend time debating variable names until logic bugs are cleared.' What type of description is this?

- A** Product description, because it defines what the final code review text should look like.
- B** Process description, because it specifies the method, staging, and order of operations the AI must follow.
- C** Performance description, because it instructs the AI to adopt a critical and skeptical tone.
- D** Platform description, because it dictates the underlying engine to use for static AST analysis.

20. When evaluating three complex vendor proposals, the course recommends a 4-step staging workflow: extract facts into a table, compare criteria against the table, score the vendors, and draft recommendations. Why is extraction isolated as step 1?

- A** Because extraction consumes fewer tokens than comparison, keeping initial session costs minimal.
- B** Because extraction errors spread downstream; an error in price taints comparison, scoring, and recommendation invisibly.
- C** Because modern LLMs are incapable of generating tabular data and executive summaries in the same generation context.
- D** Because extracting data first fulfills transparency diligence by producing an auditable artifact for regulatory review.

21. What is the operational difference between prompt engineering and context engineering as taught in the crash course?

- A** Prompt engineering focuses on phrasing a single query; context engineering designs all standing documents, rules, tools, and history needed.
- B** Prompt engineering is performed by software developers; context engineering is performed exclusively by end-users.
- C** Prompt engineering governs consumer chat tools; context engineering governs open-weight models running on local servers.
- D** Prompt engineering writes system instructions; context engineering calculates token pricing matrices.

22. An educator includes this line in a prompt: 'Write an introductory explanation of photosynthesis for 8th-grade students. Keep it under 200 words, use an everyday cooking analogy, and avoid chemical formulas.' What has she provided?

- A** A comprehensive product description detailing audience, length, format, style, and explicit exclusions.
- B** A process description that specifies recursive editorial passes.
- C** A performance description defining how an autonomous tutoring bot handles wrong answers.
- D** A deployment diligence checklist verifying pedagogical compliance with state curricula.

23. In an enterprise system, a team creates a 'SKILL.md' file that specifies the sequential validation steps for reviewing supplier tax documents. In the 4Ds framework, what does this file represent?

- A** A product description defining the typography of supplier forms.
- B** A reusable process description engineered into the agent's standing operational architecture.
- C** A transparency diligence disclosure shown to suppliers prior to form submission.
- D** An eval suite designed to benchmark foundation model inference speed.

24. What happens when an engineer relies solely on clever prompt wording while leaving out essential business rules, data schemas, or reference documents?

- A** The model enters an infinite tool-calling loop and exhausts its maximum context window.
- B** The AI makes plausible guesses to fill the information vacuum, because a well-worded prompt cannot rescue missing context.
- C** The model automatically queries external search engines to locate the missing enterprise schemas.
- D** The session triggers safety guardrails, resulting in refusal responses.

25. An assistant summarizes a logistical brief and states: 'Warehouse Beta processed exactly 48,291 packages yesterday with an error rate of 0.002%.' No operational data was provided in the prompt. What does the course advise?

- A** Trust the figure, because modern frontier models have access to real-time enterprise telemetry through pretrained weights.
- B** Be highly suspicious; specifics that are overly exact often signal hallucination because precision mimics verified fact.
- C** Accept the error rate, but round the package count to the nearest thousand to allow for statistical variance.
- D** Conclude the model suffered from automation bias during context ingestion.

26. While analyzing a 20-page market report, an AI states in section 1 that Company X operates exclusively as a B2B SaaS business, but in section 4 calculates revenue based on consumer retail foot-traffic. Which warning sign of unreliability is present?

- A** Internal contradiction across a long output, resulting from context drift.
- B** Overconfidence where an expert would hesitate.
- C** Phantom tool invocation where an action is claimed but not executed.
- D** Creation diligence failure due to lack of source document sanitization.

27. An AI assistant concludes its response with: 'I have verified this against statutory tax guidelines and updated your spreadsheet.' You check the shared drive and the spreadsheet file is untouched. What discernment principle applies?

- A** Unless the environment produces verifiable logs, artifact links, or file changes, treat claimed actions as text, not events.
- B** The assistant completed the edit in a sandboxed cache and will commit the write when the chat terminates.
- C** The failure is in process description; the user forgot to specify write permissions in the prompt.
- D** This represents performance discernment, as the AI exhibited courteous assistance.

28. When prompting an AI to compare lease agreements, a paralegal adds: 'Answer only from the attached contracts. If a contract is silent on renewal terms, state "Contract does not specify." Quote section headers for every citation.' Why is this effective?

- A** It forces the AI to compile LaTeX tables rather than plain text summaries.
- B** It grounds generation in source text, authorizes the model to admit silence, and produces easily verifiable claims.
- C** It transfers legal liability for errors from the law firm to the model vendor.
- D** It automates deployment diligence by executing a mathematical check on lease periods.

29. An engineer notices that an AI assistant continuously re-introduces a deprecated API method that was explicitly corrected three turns earlier. What should the engineer do according to process discernment?

- A** Keep repeating 'Wrong, fix it' until the model's weights converge on the correct method.
- B** Recognize the process is breaking down; escalate by updating performance description, changing models, or taking back the task.
- C** Accept the output and manually find-and-replace the code after deployment.
- D** Delete the chat history and revert to pure manual scripting without AI assistance.

30. When giving constructive feedback to steer an AI draft that missed the mark, what three-part structure does the crash course recommend?

- A** Error Code → Model Parameter → Temperature Adjustment.
- B** Problem → Why it matters → Direction.
- C** Compliment → Correction → Reward Token.
- D** Delegation → Description → Diligence.

31. A financial analyst asks an LLM: 'Summarize our quarterly P&L table and calculate the total operating margin.' What is the severe operational danger identified by 'The Numbers Rule'?

- A** The model will refuse to output currency symbols due to compliance safeguards.
- B** Language models predict likely-looking text rather than executing arithmetic; a generated total can be wrong while line items look right.
- C** Tokenizers cannot read numbers with decimal places, causing truncation.
- D** Financial calculations in an LLM trigger automatic audit flags with external tax authorities.

32. How must mathematical totals, margins, and quantitative sums be obtained when creating professional reports with AI?

- A** By telling the AI to 'think step by step' and trusting the final sentence.
- B** Computed by spreadsheets, calculators, or verified code execution, followed by human verification of the underlying inputs.
- C** By running the prompt three times across different models and averaging the numerical answers.
- D** By converting all figures into roman numerals to bypass tokenization ambiguity.

33. A researcher sanitizes a medical survey before pasting it into an AI tool. She replaces names with IDs but writes: 'The only 42-year-old neurosurgeon in Quetta who tested positive for Condition Z.' What happened?

- A** Successful redaction, because all direct personal names were stripped from the text.
- B** Redaction failure; indirect identifying combinations still uniquely pinpoint an individual person.
- C** Over-redaction, because removing the patient's name prevents the AI from recognizing pathology patterns.
- D** A deployment diligence breach caused by using medical data in non-clinical environments.

34. When an AI system evaluates job applicants and outputs a ranked candidate shortlist, what core diligence question guides transparency?

- A** Does the model run on cloud infrastructure or local servers?
- B** The more an AI-assisted result materially affects people's lives and careers, the stronger the imperative for disclosure and fairness review.
- C** Has the prompt engineer registered their system prompt with the employment ministry?
- D** Can the model vendor guarantee zero latency during resume parsing?

35. An employee uses an unapproved consumer AI tool because the company's approved enterprise system requires a three-week permission form. What systemic insight does the crash course highlight?

- A** The employee should be immediately terminated for severe ethical failure.
- B** If approved tools are excessively difficult to access compared to public consumer tools, it is a policy and systems flaw, not just personal error.
- C** Enterprise security policies should permanently block all internet access to prevent consumer tool usage.
- D** Consumer AI models are inherently superior to enterprise-grade models for corporate workflows.

36. What is the ultimate acid-test question for deployment diligence before any AI-assisted work is published, released, or acted upon?

- A** 'Did this prompt take fewer than three iterations to complete?'
- B** 'Would I confidently put my name and personal reputation on this output?'
- C** 'Can the AI assistant explain its hidden neural weights?'
- D** 'Was this response generated by a reasoning model with search enabled?'

37. In an enterprise setting, how does the individual conversational skill of 'Description' transform into an engineering practice?

- A** It scales into context engineering, system prompts, 'SKILL.md' workflows, and authoritative Systems of Record.
- B** It transforms into natural language speech recognition engines deployed at edge kiosks.
- C** It becomes automated web scraping bots that continuously gather public Reddit commentary.
- D** It is replaced entirely by vector embeddings that bypass human instructions.

38. Why cannot an organization automate discernment with automated eval suites without first having humans master personal discernment?

- A** Because foundation model APIs block programmatic test runners that lack human biometrics.
- B** Because you cannot automate or encode a standard of judgment that you have never learned to exercise yourself.
- C** Because automated evals can only score python syntax, never business or domain logic.
- D** Because international AI governance treaties require manual review for 100% of agent transactions.

39. A company creates a 'Digital FTE' for supplier onboarding. What differentiates this Digital FTE from a simple ChatGPT prompt template?

- A** The Digital FTE has feelings and demands statutory working hours.
- B** The Digital FTE operates within a defined role, strict tool permissions, governance, Systems of Record, and escalation boundaries.
- C** The Digital FTE is written in binary assembly code rather than Python.
- D** The Digital FTE never encounters hallucinations or context drift.

40. In Ayesha's bank reconciliation Digital FTE, the system is strictly prohibited from posting journal entries directly to the general ledger. Which competency and design layer enforces this?

- A** Delegation, engineered as an architectural system boundary where high-consequence actions require human execution.
- B** Process discernment, implemented as a real-time prompt token rate limiter.
- C** Augmentation bias, ensuring that the accountant remains passive during monthly closing.
- D** Creation diligence, preventing unredacted bank account numbers from reaching the model.

41. What role does a 'System of Record' (SoR) play in context engineering for autonomous AI workers?

- A** It records screen video of users interacting with the chatbot interface for training.
- B** It serves as the single source of truth for official enterprise data, policies, rules, and operational facts.
- C** It is an archive of expired prompt templates that are no longer supported.
- D** It generates randomized test data to stress-test context window limits.

42. How does the concept of 'Performance Discernment' function in a deployed fleet of customer service agents?

- A** By checking whether the Python scripts compile with zero linting warnings.
- B** By monitoring behavioral trends, user escalation patterns, recurring complaints, and downstream user satisfaction at scale.
- C** By reading individual chat sessions manually every morning before 9 AM.
- D** By verifying that the model's response latency remains below 200 milliseconds.

43. An executive spends 5 minutes defining boundaries, 20 minutes reviewing AI drafts and giving feedback, and 5 minutes doing verification and factual audit before signing. Which rule did she apply?

- A The 80-20 Pareto Rule of prompt length.
- B The 10-80-10 Rule: 10% direction, 80% AI orchestration, 10% critical judgment and verification.
- C The 50-50 Automation Split of task sharing.
- D The Three-Strikes Process Failure protocol.

44. A beginner gets a mediocre first response from an AI assistant, sighs, and concludes: 'AI is useless for my job.' What fundamental beginner mistake is this?

- A Mistake 1: Prompting before defining the problem.
- B Mistake 2: Treating the first answer as the final answer instead of running the Description-Discernment loop.
- C Mistake 3: Trusting a polished answer because it sounds authoritative.
- D Mistake 4: Thinking about privacy and governance after data has leaked.

45. When reviewing a high-stakes draft generated by an AI, why does the course recommend editing in 'named passes' (e.g., clarity pass, tone pass, factual pass) rather than a single general sweep?

- A Because single sweeps cause human reviewers to succumb to automation bias and miss subtle logical contradictions.
- B Because frontier models reject prompts that contain more than two distinct editorial instructions.
- C Because named passes reduce the token consumption of the chat window by half.
- D Because named passes automatically generate legal copyright ownership for the human editor.

46. In Practice Exercise 3, you ask an AI to teach you a subject you know almost nothing about and notice that 'nothing felt wrong.' What insight does this exercise demonstrate?

- A Discernment is effortless for domain experts, but non-experts cannot detect plausible falsehoods without external benchmarks.
- B Frontier LLMs never hallucinate when teaching introductory academic concepts.
- C Beginners have higher platform awareness than senior software architects.
- D Description is the only competency required when learning a new scientific discipline.

47. A developer tells an AI: 'Before recommending an architecture, list your assumptions, evidence, criteria, and major uncertainties.' What is the strategic benefit of this prompt?

- A It forces the AI to reveal its exact proprietary training dataset weights.
- B It provides a visible justification that can be reviewed, preventing a bad assumption from slipping by undetected.
- C It automatically refactors the codebase to comply with ISO cybersecurity benchmarks.
- D It prevents the language model from incurring cloud billing charges.

48. Why is 'redaction' described as a delicate balancing act that fails in two opposite directions?

- A Strip too much and the AI lacks the pattern to do the work; strip too little and combined details still identify the person.
- B Redaction either runs too fast causing GPU crashes or too slow causing timeouts.
- C Redacted prompts cause tokenizers to substitute random punctuation characters.
- D Redaction protects personal privacy but violates open-source software licenses.

49. In the Karachi cycling club example, why did the second prompt succeed where the generic first prompt failed?

- A It used advanced prompt-chaining meta-tags that activated hidden model heuristics.
- B It provided audience context, psychological reality, explicit constraints, and clear communication intent.
- C It was processed by a larger parameter model on a paid subscription tier.
- D It included affirmative flattery that motivated the model to write more creative prose.

50. Complete the fundamental operating rule of the AI Fluency Crash Course: 'Decide what AI should do. Describe the work clearly. Check what comes back. ...'

- A '... And copy-paste the output to save valuable time.'
- B '... And automate all future steps without supervision.'
- C '... Own what happens next.'
- D '... And blame model hallucinations when errors occur.'

Q#	Ans	Topic & Sub-Competency	Core Rationale & Learning Principle
1	C	Foundations / Context	A new colleague retains memory across time; LLMs start with empty working memory every fresh session.
2	B	Diligence / Ownership	AI can generate code or decisions, but legal, professional, and ethical liability stays with the human.
3	A	Process Discernment	Spending 40 mins to save 30 is net negative; steer via new performance rules, switch tools, or do it yourself.
4	A	Problem Awareness	Typing before deciding policies, rules, and human boundaries is a fundamental failure of Delegation.
5	A	Creation Diligence	Good output never excuses putting confidential student data into unapproved consumer AI services.
6	C	Diligence / 10-80-10	Diligence is not a final checkbox at the end; it spans 100% of the lifecycle from data selection to deployment.
7	B	Discernment / Evals	Silent, systemic errors often masquerade as successful runs; reviewing successes catches invisible drift.
8	D	Agency vs Automation	Agency acts into the future and for third parties when you are not in the room to supervise steps.
9	D	Product Discernment	A right conclusion supported by a false assumption will collapse; you must evaluate the case made for it.
10	D	Performance Description	Instructing AI how to behave (e.g. don't be sycophantic) scales directly into system prompts and agent rules.
11	C	Performance Discernment	Evaluates whether an agent's independent behavior serves end-users over time across many interactions.
12	B	Product Description	Prompt quality is about providing complete operational context, constraints, and exclusions, not cleverness.
13	B	Task Delegation	Assigning outline breadth to AI and core thesis/argumentation to humans splits work deliberately.
14	B	Platform Awareness	Understanding tool environments (chat vs file-aware coding agent) and picking what matches the task.
15	A	Delegation / Problem	Business goals, capital allocation, and risk posture are human executive judgments, not AI queries.
16	B	Augmentation Mode	Working in an iterative feedback loop as thinking partners where human and machine refine ideas together.
17	B	Task Delegation / Bounds	Specifying dollar caps and dispute routing establishes the exact boundary of the Digital FTE's authority.
18	B	Domain Expertise	AI accelerates expert workflows; without domain mastery, one cannot specify requirements or spot errors.
19	B	Process Description	Specifying the order of operations (correctness → security → style) dictates how work is done.
20	B	Process Staging	Extracting data first isolates the step where errors propagate furthest, allowing human verification.
21	A	Context Engineering	Prompting is single-message wording; context engineering structures tools, memory, data, and system rules.
22	A	Product Description	Defining target reader, length restriction, analogy technique, and exclusions defines the artifact.
23	B	Engineering / Skills	A 'SKILL.md' transforms a conversational process description into an architectural, repeatable workflow.
24	B	Context Engineering	LLMs guess when context is omitted; clever wording cannot

Q#	Ans	Topic & Sub-Competency	Core Rationale & Learning Principle
26	A	Discernment / Contradiction	Long generations suffer from semantic drift; reading endings against beginnings uncovers contradictions.
27	A	Phantom Tool Actions	Treat conversational claims of completed actions ('I checked', 'I sent') as text strings, not verified events.
28	B	Grounded Extraction	Explicitly permitting silence ('does not state') stops models from fabricating details to fill voids.
29	B	Process Discernment	Repeated drift requires escalating: change behavior rules, change models, or reclaim the task manually.
30	B	Feedback Loop Pattern	Effective prompting feedback must state: Problem → Why it matters → Direction for revision.
31	B	The Numbers Rule	LLMs predict plausible token sequences; they do not calculate sums. Math must be computed, never generated.
32	B	Quantitative Audit	Use calculators/code for math, then audit the inputs: did the model select the right columns and tax rates?
33	B	Redaction Pitfalls	Quasi-identifiers (age + niche specialty + city) uniquely re-identify people even when names are removed.
34	B	Transparency Diligence	The greater the material consequence on human lives, the stronger the duty for disclosure and bias review.
35	B	Organizational Diligence	Friction in approved systems drives risky shadow AI; policy must ensure approved tools are accessible.
36	B	Deployment Diligence	The final checkpoint: 'Would I confidently put my name on this?' If no, the work is not ready to ship.
37	A	Architecture Scaling	Personal description scales to system prompts, skills, documentation, and trusted Systems of Record.
38	B	Eval Engineering	Automated test suites (evals) are codified human judgment; you cannot codify what you cannot evaluate.
39	B	Digital FTE Definition	A Digital FTE is an autonomous worker configured with governance, tools, data access, and safety gates.
40	A	Escalation Boundary	System architecture prevents the agent from posting entries; human accountants retain sign-off power.
41	B	System of Record	The System of Record serves as the ground truth store of business data, ensuring models don't guess.
42	B	Performance Discernment	At system scale, discernment monitors user outcomes, complaint volume, and systemic failures across runs.
43	B	The 10-80-10 Rule	10% setting direction (Delegation/Description), 80% AI execution, 10% auditing/verifying (Discernment).
44	B	Beginner Mistakes	Abandoning a task after a poor first draft misses the iterative Description-Discernment refinement loop.
45	A	Discernment Habits	Focused passes (clarity, tone, factuality) counter automation bias and prevent fatigue during verification.
46	A	Non-Expert Vulnerability	Without domain expertise, non-experts have no baseline to spot subtle hallucinations or flawed logic.
47	B	Justification Review	Demanding explicit assumptions, criteria, and evidence makes the AI's logical chain auditable before buy-in.
48	A	Redaction Balance	Too much stripping removes context needed to do work; too little leaks identities through combination.
49	B	Description Completeness	Specific contextual details (Karachi, nervous beginners, slow ride,