

AI Prompting in 2026

50 Scenario-Based Multiple-Choice Questions

A Zia Tutor AI practice set, built from every concept in *The AI Agent Factory's* “AI Prompting in 2026: A Crash Course”

These questions test the thirteen concepts from the crash course through realistic scenarios, not bare definitions — the way a power user actually runs into them: briefing AI well, knowing where its knowledge is strong or thin, choosing a retrieval mode, managing context and memory, using reasoning, neutralizing sycophancy, running the brainstorm-iterate loop, working with images and audio, building small apps, analyzing data, using desktop apps safely, choosing the right model, and checking models against each other. Read each scenario, pick the single best answer, and check your work against the answer key at the end. No peeking before you've committed to an answer — that's how the score means something.

Part 1 — How AI Knows Things

1. Ayesha wants to know which car to buy and simply types “which car should I buy?” into an AI chatbot. Compared with a power-user prompt, what is missing from hers?
 - A) Nothing — a short question is the ideal way to prompt AI
 - B) Context such as her budget, dealer quotes, commute length, and family needs
 - C) A request for an image of the car
 - D) A request to turn off reasoning mode
2. A manager asks AI to “write my self-review for my boss.” A power-user version of this same request would additionally:
 - A) Ask the AI to keep the answer under 50 words
 - B) Attach a project-tracker screenshot, recent project docs, and a voice memo of notes
 - C) Disable web search
 - D) Ask the AI to use only pretrained knowledge
3. An entrepreneur asks AI to “critique my mobile tie-dyeing business idea” with no further detail, and the AI responds with mostly praise. What caused this, according to the course?
 - A) The AI's context window was full
 - B) The prompt was bait for sycophancy — it gave the AI no objective standard to check against
 - C) The AI does not know what tie-dyeing is
 - D) The model was in reasoning mode
4. A user asks an AI about the exact traditional rules of a folk game played only in their grandmother's small village. The AI answers confidently but turns out to be mostly wrong. What best explains this?
 - A) The AI intentionally lied
 - B) The topic was “quiet” in training data, so the AI generalized from similar games elsewhere
 - C) The AI was using deep research mode
 - D) The user's spelling errors confused the model

5. Which question would deserve the LOWEST trust when answered from pretrained knowledge alone?
- A) “How do I make a roux?”
 - B) “What is the plot of a top-1000 movie?”
 - C) “What did our company decide last quarter?”
 - D) “Why do cats stare at walls?”
6. A user’s prompt contains several misspelled words (“definately,” “recieve”). What does the course say about this?
- A) The AI will likely misunderstand the whole prompt
 - B) AI was trained on internet text full of typos, so it handles misspellings gracefully
 - C) The user should always run spellcheck before prompting
 - D) Typos increase the chance of sycophancy
7. A neighborhood volunteer wants a report on local water-quality issues over the last 24 months, built from EPA data, city utility reports, and peer-reviewed studies, while avoiding forums. Which retrieval mode is she implicitly requesting?
- A) Pretrained only
 - B) Web search
 - C) Deep research
 - D) Code execution
8. A user asks, “why do cats stare at walls?” Which retrieval mode is most likely to answer this, and why?
- A) Deep research, because animal behavior needs many sources
 - B) Pretrained knowledge, because this is a stable, commonly discussed topic
 - C) Web search, because the answer changes daily
 - D) Code execution, because it requires computation
9. A friend asked AI “where to run in Henderson, Nevada,” and it cited a 20-year-old webpage recommending a school no longer open to the public. What does this illustrate about web search?
- A) Web search always verifies whether a source is current
 - B) Web search does not automatically check whether a cited source is current
 - C) The AI should have used deep research instead
 - D) The AI made this up without searching at all
10. A user wants to find the official IRS page for form 1040. Which tool is better suited to this specific task, according to the course’s Google-vs-AI comparison?
- A) AI, because it synthesizes information
 - B) Google, because the user wants to land on a specific known site
 - C) Deep research, because it takes longer
 - D) Neither — this cannot be answered by either tool

Part 2 — Talking to AI Well

11. A teacher asks AI to draft a lesson plan on the water cycle with no other context and gets a generic plan. The next day she attaches her syllabus, last week’s worksheets with grades, and her school’s test

format. Why was the second plan much better?

- A) The AI switched to deep research mode
- B) She put more relevant context into the context window, which is all the model has to work with for that answer
- C) She asked the AI to “think hard”
- D) She used a different AI model entirely

12. In one long-running chat, a user planned a workout, then debugged a spreadsheet, then asked for a thank-you note to their aunt. Answers about the note start oddly referencing the workout. What is this called?

- A) Sycophancy
- B) Context rot
- C) A diffusion error
- D) Model ladder drift

13. A user notices the AI keeps assuming they still work in retail, even though they changed careers months ago, and this wrong assumption shows up in unrelated new chats. What is the likely cause and fix?

- A) The AI is hallucinating; nothing can be done
- B) An outdated note in the tool's memory layer is loading into new chats; the user should correct or edit that memory
- C) The context window is too small; the user should upgrade their plan
- D) The system prompt is broken and cannot be fixed

14. A tax preparer repeatedly pastes the same client documents, filer assumptions, and instructions into a new chat every time she has a tax question. What should she do instead?

- A) Keep repeating the context in each chat, since that is unavoidable
- B) Set up a project with the standing files and instructions, since the task, background, and output shape recur
- C) Switch to using only pretrained knowledge
- D) Ask the AI to remember everything permanently with no review

15. Which of the following is NOT one of the six things that can be present in the model's context window for a given answer?

- A) The system prompt
- B) Tool descriptions
- C) The user's uploaded files
- D) The AI company's stock price

16. A user is choosing between two job offers and has attached offer letters, a cost-of-living comparison, and a commute analysis. They want the AI to weigh many trade-offs carefully rather than respond instantly. What should they add to the prompt?

- A) “Keep it under one sentence”
- B) “Think hard” or turn on extended thinking
- C) “Only use pretrained knowledge”

D) "Do not use any files I attached"

17. According to the course, when is it best to leave thinking mode off?

- A) When choosing between two competing job offers
- B) For quick lookups, one-paragraph summaries, and casual idea-gathering
- C) When analyzing a large spreadsheet
- D) When comparing multiple contractor quotes

18. A 2025 METR study cited in the course found the length of task a leading AI model could reliably finish was roughly doubling on what timescale?

- A) Every month
- B) Every seven months
- C) Every five years
- D) It has stayed constant since 2023

19. A user writes: "Don't you think remote work is obviously better than office work?" What effect does this phrasing tend to have on the AI's answer?

- A) It makes the AI more likely to disagree, to seem balanced
- B) It signals a preferred conclusion, so the AI tends to agree and list supporting reasons
- C) It has no measurable effect on the answer
- D) It automatically triggers deep research mode

20. Which of the following is the BEST neutral rewrite of "Find evidence that this marketing strategy will work"?

- A) "Confirm that this marketing strategy will work"
- B) "Prove that this marketing strategy is a good idea"
- C) "Evaluate this marketing strategy. List the strongest arguments for and against it."
- D) "Tell me why this marketing strategy is smart"

21. A manager wants an honest score on a draft memo rather than empty praise. Why does forcing a numeric score (1–10) on named criteria work better than asking "is this good?"

- A) Numbers take less time to generate than words
- B) A specific number forces the AI to commit, which makes it look more carefully instead of offering cheap, vague praise
- C) Scores are always more accurate than any other output
- D) Numeric scores disable sycophancy entirely with no other steps needed

22. Why should a user explicitly tell the AI that "I don't know" is an acceptable answer?

- A) It makes responses shorter
- B) Without that permission, a model under pressure to answer is more likely to invent an answer than leave a gap open
- C) It automatically enables deep research
- D) It prevents the AI from using any file uploads

23. A user wants exercise ideas and simply asks "ways to exercise at home." Why is the answer likely to be generic (squats, push-ups, planks)?

- A) The AI intentionally withholds good ideas
- B) Most of the internet describes common ideas, and the AI learned from the internet, so the average answer is also average
- C) The AI cannot generate exercise advice at all
- D) The user needs to enable code execution

24. Put the brainstorm-iterate loop steps in the correct order: (1) ask for 3–5 new options built on feedback, (2) load full context, (3) expand the chosen option into a full draft, (4) ask for 3–5 options without expanding any, (5) give explicit feedback on what you like/dislike.

- A) 2, 4, 5, 1, 3
- B) 1, 2, 3, 4, 5
- C) 4, 2, 5, 1, 3
- D) 2, 1, 4, 5, 3

25. A marketer asks AI for five campaign ideas but pins the exact tagline, a required three-part structure, the tone, and a 40-word cap — and gets five near-identical rewordings of one idea. What is the fix?

- A) Ask for fifty ideas instead of five
- B) Loosen the structural limits temporarily so the AI can generate truly different directions, then tighten them once one is chosen
- C) Switch entirely to deep research mode
- D) Add even more constraints to force originality

26. According to the “four kinds of ask” idea, which element should be loosened (not stated) when a user wants genuine research that maps the ground rather than confirms a belief?

- A) The sources
- B) The scope
- C) The answer/conclusion — the user should never say what they expect to find
- D) The exact question

Part 3 — Beyond Text

27. A user uploads a photo of a whiteboard where their own head blocks part of a word in a diagram, and the AI still guesses the missing word correctly from the surrounding context. What does this best illustrate?

- A) AI is strong at zooming into fine printed detail
- B) AI is strong at filling in gaps from the overall sense of an image, but weak at zooming into fine detail
- C) AI cannot read whiteboards at all
- D) AI can only process photos, never diagrams

28. In the course's diagram-building recipe, which tool is recommended for the FIRST step — drawing the raw structure of a diagram as SVG?

- A) An audio transcription tool
- B) Claude, because deciding boxes/arrows/labels is a reasoning task it is especially strong at
- C) A spreadsheet program
- D) A video generator

- 29.** A doctor records a 45-minute patient consultation (with consent) and asks AI to produce a structured clinical note, flagging anything it could not understand confidently. What best describes the value this provided?
- A) It replaced the doctor's clinical judgment entirely
 - B) It removed most of the typing work while the doctor still checked and finished the note herself
 - C) It generated the note faster only because thinking mode was off
 - D) It could not process the audio because AI cannot handle audio input
- 30.** Why can't a user typically interrupt an image generation partway through to save time, the way they can stop a long text answer?
- A) Image generators are always slower than text models
 - B) An image generator (a diffusion model) produces the whole image at once by removing noise step by step, rather than writing part by part
 - C) Image tools do not have a stop button in the interface
 - D) Images always take longer than videos to generate
- 31.** A parent with no coding background wants a simple typing game for their child and describes the goal, what the player provides, and what they should see. What is this three-part prompt structure called?
- A) Rubric, Score, Iterate
 - B) Goal, Input, Output
 - C) Context, Tool, Memory
 - D) Draft, Grade, Fix
- 32.** What is the "artifact" that appears in a side panel when an AI builds a small app from one prompt?
- A) A static screenshot with no functionality
 - B) A working object from the conversation that can be edited, shared by link, embedded, or downloaded
 - C) A file that can only be viewed by the person who created it, with no sharing option
 - D) A permanent record stored only on Anthropic's servers with no export option
- 33.** Which of the following is explicitly named as something still HARD for a one-prompt AI app to handle?
- A) A Pomodoro timer
 - B) A bill-splitting calculator
 - C) Multiplayer functionality over the internet, with accounts and matchmaking
 - D) A simple fireworks click simulator
- 34.** A shop owner uploads a year of sales data and asks, "Which drinks had the biggest changes in sales?" without mentioning code. What risk does this create?
- A) The AI will always run code regardless of wording
 - B) The AI might skip running actual code and instead answer from a quick glance, producing a confident but uncomputed answer
 - C) The file upload will be rejected
 - D) The AI will refuse to answer without a subscription

35. What is the single most direct way to make sure the AI actually computed an answer about your spreadsheet rather than guessed?

- A) Ask it to answer faster
- B) Explicitly ask it to write and run code, and show you the code it ran
- C) Attach the file twice
- D) Turn off extended thinking

36. A user asks AI to state the exact row count, column names, and date range of an uploaded file before analyzing anything. What is the purpose of this technique?

- A) To slow down the response for dramatic effect
- B) It's a checkable fact — if the model is really reading the file these will be right; if it's inventing, they'll look plausible but be wrong
- C) To trigger deep research mode
- D) To disable sycophancy

37. Even when the AI's code genuinely ran on a spreadsheet, what should a user still double-check?

- A) Nothing — code execution is always fully correct
- B) Final totals, graph labels, and whether a column was interpreted correctly (e.g., what "TXN_AMT" actually means)
- C) Only the file's upload date
- D) Whether the AI used reasoning mode

Part 4 — Working Safely and Choosing Tools

38. A consultant asks an AI desktop app to reorganize a folder of 240 client PDFs. Before any files are actually moved or renamed, what should happen first?

- A) The app should immediately execute the reorganization
- B) The app should propose a plan (e.g., a tree structure) for the user to review and approve first
- C) The user should grant full disk access up front
- D) The user should delete the original folder as a backup step

39. What does the course warn happens to files an AI desktop app deletes, in many cases?

- A) They always go to the recycle bin, fully recoverable
- B) They often do NOT go to the recycle bin — they are simply gone
- C) They are automatically backed up to the cloud
- D) They cannot be deleted by AI apps at all

40. Per the "permission ladder," what should a user typically allow in their very first sessions with a new AI desktop app?

- A) Full read/write access to the entire computer
- B) Read-only access to a single small folder
- C) Open-ended "do whatever you need" permission
- D) Access to the desktop and documents root folders

41. According to the cost/speed comparison in the course, which output type generally costs the most and is slowest to produce?

- A) Text
- B) Speech/audio
- C) Images
- D) Video

42. A user always defaults to one single AI tool for every task, assuming it is simply “the best one.” What does the course say about this habit?

- A) It is a smart, efficient habit, since switching tools wastes time
- B) AI ability is “jagged” — different models lead on different tasks, so relying on only one tool is likely wrong for most tasks
- C) It is fine as long as thinking mode is always on
- D) It only matters for image generation, not text

43. A user picks the cheapest, fastest rung of a model ladder to run a complex, high-stakes financial analysis, and separately runs a one-line casual question on the most expensive flagship model. What is wrong with both choices?

- A) Nothing is wrong — model choice doesn't affect quality
- B) Both waste something: the hard analysis is undersupported on a cheap rung, and the flagship model wastes time/budget on a task that didn't need it
- C) The model ladder only applies to image generation
- D) Only the second choice is a problem

44. What does the course recommend doing “monthly” to keep up with which AI model currently leads on a given task?

- A) Reading a single company's marketing page
- B) Checking a public leaderboard like Arena and running one regular task through a few different tools
- C) Re-reading the entire crash course from the start
- D) Disabling memory in all tools

45. A user has Claude review a draft, then opens a second Claude conversation and asks it to review the same draft again. Does this count as genuine cross-model checking?

- A) Yes, because it is a “second opinion”
- B) No, because both are from the same model family and are too likely to share the same blind spots
- C) Yes, but only if extended thinking is enabled
- D) No, because Claude cannot review its own writing style

46. A consultant's memo scores 9/10 from her main AI model's self-critique. She pastes it into a model from a different company, which scores it 7.5 and raises several new issues. What does this illustrate?

- A) The first model was lying
- B) A self-graded 9 from one model isn't the same as a 9 confirmed by an independent model from a different family — cross-checking surfaces blind spots the first model couldn't see in itself
- C) The second model is always more accurate than the first
- D) Scores from different models should never be compared

47. For which kind of task is the FULL multi-model cross-checking loop LEAST necessary?
- A) A contract you are about to sign
 - B) A chapter that will be published
 - C) A quick casual list of ideas or a short email
 - D) A memo your boss will read and act on
48. What caveat does the course give about cross-model checking as a way to catch errors?
- A) It guarantees a fully correct answer every time
 - B) Models can still share the same wrong belief, since they were trained on overlapping data — the score is a progress signal, not proof, and humans must still check high-stakes facts
 - C) It replaces the need for any human review, always
 - D) It only works for coding tasks

Bringing It Together

49. A user wants to check whether an AI truly searched the web for very recent information rather than answering from memory. What single underlying fact about the model explains why this distinction matters?
- A) The model is multimodal
 - B) The model is stateless and normally knows only what is inside its context window, plus anything a tool like web search actively fetches into it
 - C) The model always prefers deep research
 - D) The model cannot process any text after training
50. Across the whole crash course, which one sentence is repeatedly offered as the single idea underlying almost every technique taught?
- A) "Always use the most expensive model available."
 - B) "Get the right context in, keep the wrong context out."
 - C) "Never trust an AI's output under any circumstances."
 - D) "Longer prompts are always better than shorter ones."

Answer Key

Score yourself: 45–50 correct means the thirteen concepts are solid — you're briefing AI like a power user. 35–44 means a couple of concepts need a second pass. Below 35, reread the crash course a part at a time and retake that part's questions.

1. B	2. B	3. B	4. B	5. C
6. B	7. C	8. B	9. B	10. B
11. B	12. B	13. B	14. B	15. D
16. B	17. B	18. B	19. B	20. C
21. B	22. B	23. B	24. A	25. B
26. C	27. B	28. B	29. B	30. B
31. B	32. B	33. C	34. B	35. B
36. B	37. B	38. B	39. B	40. B
41. D	42. B	43. B	44. B	45. B
46. B	47. C	48. B	49. B	50. B

Source: "AI Prompting in 2026: A Crash Course," The AI Agent Factory (agentfactory.panaversity.org/docs/ai-prompting-2026), read via the Agent Factory System of Record. Generated by Zia Tutor AI.