

Claude & ChatGPT 101: Certification Assessment

Scenario-Based Evaluation • Applied Workplace Competencies • 50 Questions

FOUNDATIONS SERIES

MODULE 1: COCKPIT ARCHITECTURE & PLATFORM AWARENESS

Question 1 • Two Cockpits, One Discipline

A financial consultant has used ChatGPT daily for twelve months and refuses to open Claude for an urgent client review, claiming that adopting a second product requires starting his operational training from scratch. What does the course say is wrong with that reasoning?

- A Claude is inherently stronger for consulting, so his earlier year in ChatGPT was largely wasted effort.
- B Most of what he learned is the transferable discipline of delegation; the new part is merely where switches live.
- C Cross-platform verification is mandated by ISO AI standards before drafting financial reports.
- D Interface buttons change weekly, so vendor loyalty guarantees total functional obsolescence within 90 days.

Question 2 • Model Tiers vs. Brand Names

An executive directs his staff to "always select Opus on Claude and Pro on ChatGPT" for every single prompt across the enterprise. Which core principle of model tier selection does this executive order violate?

- A Flagship models must be paired exclusively with custom GPTs or skills rather than direct chat.
- B Choose the capability level by the specific task rather than habit; routine drafting wastes expensive tiers.
- C Frontier reasoning flagships do not accept uploaded tabular spreadsheets or text attachments.
- D Heavy flagships produce verbose outputs that cannot be translated into downstream code blocks.

Question 3 • Choosing the Thinking Tier

An operations analyst leaves his workspace on the fast default tier and complains that the assistant consistently miscalculates tiered volumetric shipping margins across a three-step fulfillment supply chain. What is the course's diagnosis?

- A Multi-step logic and careful computation belong on the thinking tier, which plans before outputting.
- B The prompt was missing explicit markdown tables, which force fast default models to calculate properly.
- C Fast tiers disable Python execution, meaning basic math is completely impossible in standard chat.
- D His project storage exceeded capacity limits, corrupting numerical reasoning across conversation threads.

Question 4 • The Context Attachment Gap

Two corporate associates prompt: *"Summarize this non-compete clause and flag any high-risk liabilities."* One receives an actionable, cited breakdown; the other receives vague generic definitions. Assuming identical settings, what explains the gap?

- A One associate enabled extended thinking, which automatically fetches legal contracts from company clouds.
- B One associate physically attached the document file, whereas the other forced the model to hallucinate.
- C One associate had custom instructions specifying corporate law terminology in system settings.
- D The failed assistant was routed through workplace search rather than public web search.

Question 5 • Handling Context Window Drift

After twelve hours in a single chat drafting an annual audit, an assistant begins re-introducing banned bullet points and contradicting earlier agreed figures. What tactical recovery does the course recommend?

- A Immediately delete the conversation history and start over entirely from memory.
- B Request a concise state summary of decisions and rules, then paste it into a fresh chat with files attached.
- C Toggle the thinking switch on and off to purge stale tokens from the active inference buffer.
- D Switch from Claude to ChatGPT immediately, as chat degradation only occurs on Anthropic architectures.

Question 6 • The Three-Time Project Rule

A copywriter asks whether she should build a dedicated Project workspace for a one-off 45-minute assignment rewriting a speaker biography that will never recur. Applying the course's threshold test, what should she do?

- A** Yes; projects provide baseline encryption necessary for all professional client work.
- B** Yes; projects are mandatory whenever custom instructions differ from the global account profile.
- C** No; project workspaces are intended for recurring streams where context is explained or uploaded three times.
- D** No; creating a project requires enterprise administrator approval on both platforms.

Question 7 • ChatGPT Project-Only Memory Isolation

An HR investigator manages two highly sensitive disciplinary reviews in separate ChatGPT projects. She fears that witness details from Review Alpha might bleed into Review Beta. Which configuration addresses this?

- A** Custom instructions commanding the model to ignore cross-project tokens during generation.
- B** ChatGPT's project-only memory toggle, which walls off memories and prevents cross-project leakage.
- C** Claude's retrieval-augmented generation mode, which automatically partitions company knowledge.
- D** Running all chats in temporary incognito tabs to prevent database indexing.

Question 8 • Claude Project Retrieval Scaling

A municipal planning team uploads 35 comprehensive zoning ordinances into a single Claude Project, surpassing standard context window limits. How does Claude manage this knowledge volume on paid plans?

- A** It automatically compresses all documents into low-resolution vector graphics.
- B** It shifts automatically to retrieval-augmented search over project knowledge, scaling capacity tenfold.
- C** It deletes the oldest uploaded files to fit within the physical context boundary.
- D** It requires the user to manually convert all PDF files into raw SQL database tables.

Question 9 • Stale Project Knowledge Hazard

A marketing director leaves last year's Q1 product pricing matrix inside an active Project. While generating a new proposal, the assistant quotes outdated volume discounts without error. What failure mechanism does this illustrate?

- A** Vector collision within the model's public embedding layer.
- B** Stale context fails quietly: the assistant cannot detect that supplied source truth is obsolete.
- C** Fast default models systematically prioritize project files over direct conversational prompts.
- D** Model context protocol connectors corrupting local workspace attachments.

Question 10 • The Persistence Triad Allocation

A supply chain director pastes the note *"Currently resolving the Acme supplier tariff dispute"* into his account's global custom instructions. Six months after resolution, the assistant still references the dispute. What rule was violated?

- A** Custom instructions are for stable rules; evolving context belongs in memory, and scoped work in projects.
- B** Legal disputes must always be routed through temporary incognito chats.
- C** Tariff negotiations require custom GPT development rather than conversational prompting.
- D** Instructions should never contain business entity names due to copyright governance.

Question 11 • Incognito & Temporary Chats Reality

A healthcare executive discusses confidential merger prospects in an incognito/temporary chat, assuring her board that *"no digital trace exists anywhere on vendor infrastructure."* How does the course correct her?

- A** She is completely correct; temporary chats bypass all server logging mechanisms.
- B** Temporary chats prevent history and memory logging, but vendors retain copies for ~30 days for safety review.
- C** The chat is wiped immediately from vendor servers, but remains permanently visible in browser cache.
- D** Both platforms publish incognito chat logs directly to public debugging indexes.

Question 12 • Enterprise Default Memory Posture

A newly hired project manager on a corporate Claude Enterprise account assumes memory is actively recording her formatting preferences across projects. Why are her preferences failing to persist?

- A** Enterprise accounts disable custom instructions permanently across all business domains.
- B** Claude disables memory by default on Team and Enterprise tiers until an organization administrator enables it.
- C** Claude requires users to install an external MCP memory server before storing user attributes.
- D** Memory is exclusive to ChatGPT and does not exist in Anthropic's ecosystem.

Question 13 • Evolution of ChatGPT Canvas to Blocks

A junior developer reads an older 2024 workflow guide advising him to click the "Canvas" interface button in ChatGPT to co-edit a Python script, but cannot find the control. What platform change occurred?

- A** Canvas was gated strictly behind enterprise developer API tokens.
- B** Canvas was retired on modern models in favor of inline writing blocks and code blocks inside the conversation.
- C** OpenAI merged Canvas into the Model Context Protocol directory.
- D** Code editing can now only be performed inside Claude Artifacts.

Question 14 • Artifacts vs. Chat Thread Scroll

A lead designer needs Claude to generate an interactive SVG state diagram and a standalone product release document that can be iteratively revised over several hours. Why is an Artifact ideal here?

- A** Artifacts bypass the model's standard inference pricing token calculation.
- B** Artifacts isolate substantial work products beside the chat, preventing finished assets from getting lost in scroll.
- C** Artifacts execute directly on local desktop operating system threads without internet access.
- D** Artifacts automatically translate English prompts into compiled C++ binary executables.

Question 15 • Computed Numbers vs. Written Text

A controller asks an assistant to calculate quarterly EBITDA margins across seven regional subsidiaries for an audited SEC filing. What critical prompting discipline should she enforce?

- A** Ask for the numbers to be computed and executed via code tools, and verify the specific columns used.
- B** Demand the answer in prose paragraphs, as text tokenization handles arithmetic more reliably than code.
- C** Instruct the assistant to estimate margins using web search consensus.
- D** Switch the model to fast default mode to prevent recursive numerical rounding loops.

Question 16 • The Fundamental Separation: What vs. How

A risk consultancy maintains a library of proprietary compliance audit checklists and applies them across forty banking clients. How should they structure this work across projects and skills?

- A** Package the audit checklist as a reusable skill (the how) and store each bank's files in its project (the what).
- B** Build forty custom GPTs, hardcoding the audit checklists into each separate system instruction.
- C** Merge all forty client documents into a single global project knowledge folder.
- D** Paste the complete audit checklist into account memory so it applies invisibly to all prompts.

Question 17 • The Open Agent Skills Standard

An IT director chooses to package engineering runbooks following the Agent Skills open standard (agentskills.io). What strategic advantage does this provide over building vendor-locked features?

- A** The skills automatically bypass enterprise firewall authentication rules.
- B** It consists of folders of Markdown files without server runtimes, enabling skills to travel across tools.
- C** It guarantees 100% mathematical accuracy without human verification.
- D** It converts natural language prompts directly into cloud infrastructure deployments.

Question 18 • Anatomies: Custom GPTs vs. Portable Skills

A support department deliberates between building a Custom GPT in ChatGPT versus writing an Agent Skill. Which distinction correctly characterizes the two mechanisms?

- A** Custom GPTs require Python code servers; skills only run on Anthropic infrastructure.
- B** A custom GPT is a separately opened assistant persona; a skill is a portable method loaded on demand.
- C** Skills can access web search, whereas Custom GPTs are strictly barred from external browsing.
- D** Custom GPTs share context globally across organizations, while skills remain private to one user.

Question 19 • Deconstructing Plugins

A logistics manager reads that Claude and ChatGPT support plugins. What components constitute a plugin under current architectural definitions?

- A** An enterprise billing agreement tied to commercial API endpoint usage.
- B** A bundle packaging skills, external connectors/apps, and commands tailored for an entire job function.
- C** A specialized browser extension that injects scraped DOM elements into chat boxes.
- D** A proprietary vector database engine hosted exclusively inside Microsoft Azure.

Question 20 • The Model Context Protocol (MCP) Standard

An infrastructure architect evaluates MCP support across Claude and ChatGPT. Why is MCP widely referred to as the "USB-C for AI tools"?

- A** It physicalizes cloud data pipelines into local encrypted hardware drives.
- B** It provides an open, uniform connection protocol allowing one integration to connect across assistants.
- C** It restricts all AI tool integrations to hardware peripherals manufactured after 2025.
- D** It automatically revokes third-party access if sensitive tokens are transmitted.

Question 21 • The USB-C Shorthand Fallacy

A junior administrator connects an MCP integration to the corporate Jira database without IT review, arguing that *"since MCP is just a standardized plug, security risks are non-existent."* What does this overlook?

- A** The plug is standardized, but granted access permissions govern what the assistant can read and execute.
- B** Jira blocks MCP connections unless the organization maintains an on-premise Anthropic mainframe.
- C** Standardized protocols expose all user screens directly to public web search crawlers.
- D** MCP connectors can only read text files and cannot interact with relational databases.

Question 22 • Pre-Connection Governance Audit

Before approving an employee's request to connect an AI app/connector to the department's Microsoft Outlook workspace, which specific security question must leadership ask first?

- A** Does the connector operate on a Python 2.7 runtime environment?
- B** What can it read, can it execute write/send actions in the outside world, and is the source trusted?
- C** Has the connector been certified to run on fast default models without latency?
- D** Will connecting Outlook automatically disable the user's custom instructions?

Question 23 • Daily Delegation Routing: Quick Facts

A financial trader triggers an intensive 10-minute autonomous Research run simply to verify whether the European Central Bank adjusted its benchmark interest rate this morning. What operational defect did this cause?

- A** Stale memory saturation across enterprise project caches.
- B** Misrouting: one current fact is handled by web search in seconds, making deep research an expensive waste.
- C** Token hallucination, because research mode cannot parse financial macroeconomics.
- D** Violation of the Agent Skills open standard protocols.

Question 24 • Daily Delegation Routing: Deep Analysis vs. Web Search

An executive attaches three confidential product roadmaps and toggles Web Search to determine which roadmap carries the greatest operational bottleneck. Why is this choice flawed?

- A** Web search publishes local confidential attachments directly to public search engines.
- B** The task requires multi-step deductive reasoning over known context, which belongs on the thinking tier.
- C** Fast default search bars cannot read uploaded text attachments.
- D** Roadmaps must be processed by Custom GPTs rather than conversational models.

Question 25 • The Nature of Research Mode

An associate receives a comprehensive, 14-page cited research report generated by an assistant's Research mode and immediately submits it to a state regulator without reading the sources. What error was made?

- A** Research mode reports are invalid unless rendered through Claude Artifacts.
- B** Research mode extends investigative reach but never transfers professional accountability; citations must be checked.
- C** Autonomous research violates enterprise data privacy mandates across OECD nations.
- D** Research mode answers cannot be quoted in regulatory filings under any circumstances.

Question 26 • The 5-Step Proving Loop Baseline

A hospital billing manager wants an assistant to audit complex insurance claim denial codes. Before deploying it on tomorrow's claims, how should she initiate the proving loop?

- A** Have the assistant reproduce last month's completed audits where ground truth is already established.
- B** Deploy it immediately on live patient accounts and review denials when billing disputes arise.
- C** Rely on vendor benchmark scores published on AI leaderboards.
- D** Restrict prompt instructions to fast default models to prevent calculation errors.

Question 27 • The Ground Truth Imperative

Why does the proving discipline insist that an evaluator test an assistant against known past work rather than an unfamiliar, newly encountered problem?

- A** Modern frontier models refuse to generate responses to novel prompts without benchmark training.
- B** Without known truth, you cannot verify whether the model correctly solved the problem or merely sounded plausible.
- C** Stored memory caches only activate when processing retrospective historical data.
- D** It guarantees that the assistant will never hallucinate on any future tasks.

Question 28 • Earned Confidence vs. Transferred Accountability

A team runs historical payroll records through an assistant, achieving 100% numerical concordance across twelve pay cycles. The CFO proposes firing the external payroll auditing firm. What does the course say?

- A** The proposal is sound; passing historical evaluations proves zero percent future error rates.
- B** The test earned confidence for similar future work, but passing a test never transfers legal accountability.
- C** External auditing firms cannot be replaced unless the model is hosted on an on-premise MCP server.
- D** Concordance tests on historical data only prove the competence of prompt writers, not models.

Question 29 • Value of Failed Proving Runs

After four iterative attempts refining instructions, an assistant repeatedly fails to correctly parse handwritten shipping bills of lading. What conclusion should the logistics team draw?

- A** The proving loop failed; all earlier prompt refinement was wasted engineering time.
- B** A definitive finding that a task cannot be delegated reliably is a valuable result: the task stays human.
- C** The model must be re-trained from scratch on proprietary warehouse datasets.
- D** The organization should abandon AI adoption across all other departments.

Question 30 • Platform Awareness Through Comparative Runs

An investment analyst evaluates an intricate bond restructuring scenario. Instead of picking one tool, she runs identical inputs in both Claude and ChatGPT. What practical discipline does this demonstrate?

- A** Platform awareness: recognizing that models exhibit distinct strengths and spotting divergences builds rigor.
- B** Token hedging: exhausting subscription quotas evenly across enterprise accounts.
- C** MCP synchronization: forcing ChatGPT apps to calibrate against Claude's context window.
- D** Workspace redundancy mandated by corporate business continuity policies.

Question 31 • Model Interface Hallucinations

A user pastes: *"List every control I can see in this workspace right now, and tell me what each is for."* The assistant names a feature not present on screen. What lesson does the course draw from this?

- A** The browser's graphics rendering cache failed to display the missing UI button.
- B** Assistants are not always reliable about their own interfaces; checking confident claims is an essential habit.
- C** The account was silently downgraded to an unauthorized testing tier.
- D** Anthropic and OpenAI have deprecated conversational self-inspection prompts.

Question 32 • Three-Part Beginner Habit Formulation

A legal assistant submits: *"Draft a memorandum on our latest arbitration."* The response is unusable. Which foundational prompt structure from the course was completely omitted?

- A** The three-part habit: set the stage (who/goal), define the task (exact action), and set the rules (tone/constraints).
- B** Requesting that the memorandum be compiled as an executable Python artifact.
- C** Pre-loading the complete American Bar Association case law database via MCP.
- D** Mandating that the output be rendered strictly within a two-column markdown table.

Question 33 • Stored Context Governance Schedules

A consulting team experiences recurring prompt anomalies where projects quote retired brand guidelines. What recurring governance ritual does the reference mandate to prevent this failure?

- A** Review stored memory, standing instructions, and project files monthly: keep, correct, or delete.
- B** Delete all user accounts and purchase new enterprise licenses at the start of every fiscal year.
- C** Switch all conversations to temporary chats to avoid context retention entirely.
- D** Restrict project membership exclusively to employees with computer engineering degrees.

Question 34 • Structured Data vs. Unstructured Prose

A procurement officer asks an assistant to evaluate bids from 18 competing suppliers, requesting the comparison in narrative prose. Why does the course advise firmly against prose for this job?

- A** Language models charge double token costs when generating natural language paragraphs.
- B** Data requires structured output (tables/spreadsheets) so colleagues can sort, filter, or paste into Excel.
- C** Narrative paragraphs trigger automatic safety filters across enterprise LLM firewalls.
- D** Thinking mode models cannot format narrative text without hallucinating.

Question 35 • The Code-Executed Output Invariant

An assistant executes Python code in a code block to calculate a critical portfolio value at risk (VaR). The code executes without error. Why must the portfolio manager still verify the execution?

- A** Python interpreters running on cloud sandboxes randomly inject trigonometric rounding errors.
- B** Code execution is exact, but the assistant wrote the code and could have aggregated the wrong column.
- C** Code execution blocks do not support floating-point arithmetic.
- D** Running code inside chat automatically wipes the active project knowledge cache.

Question 36 • Claude Extended Thinking Mechanics

A software architect notices that when she selects Claude's Sonnet model, she can toggle "Thinking" on or off, whereas ChatGPT lists distinct model names like Instant and Pro. What does this reveal?

- A Claude only supports thinking on the Haiku model family.
- B Claude treats thinking as an explicit toggleable switch, while ChatGPT organizes reasoning effort across labeled tiers.
- C ChatGPT models do not perform step-by-step reasoning under any circumstances.
- D Anthropic requires enterprise API users to submit formal written petitions to unlock reasoning.

Question 37 • Workplace Enterprise Search Scope

A director prompts: *"Find the approved executive travel reimbursement guidelines."* Rather than searching the public web, the assistant queries internal connected tools. Which capability is being used?

- A Workplace Search (such as Claude's Enterprise Search or connected apps), querying company systems.
- B Autonomous Research mode scanning public corporate registries.
- C Temporary incognito workspace retrieval.
- D The Agent Skills open standard directory at agentskills.io.

Question 38 • Context Priming vs. Prompt Engineering

An executive spends two days crafting a 400-word prompt instructing an assistant on how to analyze a balance sheet, but provides no files. A junior clerk pastes a 20-word prompt with the audit attached. Who succeeds?

- A The executive, because prompt length and lexical complexity outweigh source documentation.
- B The clerk, because high-quality context attachments improve answers far more than clever prompt styling.
- C Neither, because financial audits must be processed exclusively via external MCP servers.
- D Both will produce identical results because frontier models infer missing balance sheet figures.

Question 39 • Action Permissions vs. Capability Boundaries

A sales manager grants an AI app full administrative permission to write and delete records in Salesforce, then asks it to dispatch personalized video messages. The assistant fails. Why?

- A Permissions select among existing tool capabilities; they cannot grant a tool actions it was never built to perform.
- B Salesforce automatically disables video rendering when accessed via Claude connectors.
- C The sales manager forgot to run the command inside a temporary incognito chat.
- D Video rendering requires upgrading the entire organization to ChatGPT Pro licenses.

Question 40 • The Industrialized Proving Loop

When a professional team takes the five-step proving loop, automates it across fifty historical test cases, and runs it before every deployment, what engineering discipline have they adopted?

- A Model Context Protocol routing.
- B Eval-Driven Development (an industrialized evaluation suite).
- C Retrieval-augmented vector hashing.
- D Global custom instruction compilation.

Q41 • System of Context Architecture

A team builds a Project with chats, files, and rules. In the framework of this book, what did they build?

- A:** An autonomous self-training foundational model.
- B:** The smallest System of Context they will ever build.
- C:** A compiled external MCP server runtime.
- D:** A transient browser-cached memory index.

Q42 • Directory Discovery Before Building

A team plans to spend weeks writing custom connectors for HubSpot. What should they do first?

- A:** Check the public directory (e.g. claude.ai/directory) to see if it exists.
- B:** Submit an architectural patent with national IP registries.
- C:** Train a specialized Haiku model from raw customer data.
- D:** Restrict all HubSpot data to temporary incognito chats.

Q43 • Stale Chat State Summary Protocol

What three elements must an explicit state summary contain when rescuing a stale conversation?

- A:** API token counts, latency logs, and server IDs.
- B:** What was decided, what is still open, and every rule given.
- C:** User passwords, billing tier, and model version hashes.
- D:** Memory variables, MCP endpoints, and vector scores.

Q44 • Screen Sharing Memory Risks

Why does the course recommend running a memory audit prompt before presenting live on screen?

- A:** Auditing memory clears the local browser cache.
- B:** To ensure sensitive/embarrassing facts do not appear on screen.
- C:** Model memory automatically shuts down when screens are shared.
- D:** Memory tokens double in cost during active screen casts.

Q45 • The Prompt Interview Pattern

When creating a Project for a complex stream of work, what prompting technique is recommended?

- A:** Demand the model guess the rules without human input.
- B:** Have the assistant interview you briefly, then generate instructions.
- C:** Paste unformatted raw desktop directories into memory.
- D:** Write raw C++ scripts inside the system instruction box.

Q46 • Vendor Model Name Volatility

Why is anchoring mental models around tier functions better than specific product names?

- A:** Model names rotate rapidly, while the tier trade-off persists.
- B:** Model names are trademarked and illegal to mention in prompts.
- C:** Vendor names do not correspond to technical performance.
- D:** Using model names automatically triggers enterprise billing.

Q47 • Research Mode Duration Expectation

How does authentic Research mode differ in execution profile compared to web search?

- A:** Research mode answers in milliseconds without citation.
- B:** Research plans and reads across sources, taking minutes.
- C:** Research mode only operates inside desktop terminal windows.
- D:** Research mode cannot access public web indices.

Q48 • The Role of Colleagues in Prompting

What mental heuristic does the course advise adopting before sending a serious request?

- A:** Treat the assistant like an adversarial penetration tester.
- B:** What would a capable human colleague need to see to answer well?
- C:** Pretend the assistant has zero knowledge of grammar or syntax.
- D:** Assume the model knows internal company plans without telling it.

Q49 • Verification of Generated Artifacts

When an assistant generates a 2,000-word contract audit inside an Artifact, what is true?

- A:** The artifact is pre-certified by vendor legal counsel.
- B:** Artifacts isolate the work product; human review is still essential.
- C:** Artifact text cannot be edited or copied into external apps.
- D:** Outputting into an artifact guarantees complete absence of bias.

Q50 • The Enduring Core Discipline

As both products release desktop agents and new tools, what principle remains constant?

- A:** The discipline is the constant; the tool is the variable.
- B:** Prompt engineering will be replaced by automated Python bots.
- C:** Fast default models will eliminate the need for reasoning tiers.
- D:** Proprietary models will fully replace open protocols like MCP.

Official Answer Key & Explanatory Rationales

Scoring Reference • Questions 1 through 50

ANSWER KEY

#	Key	Core Concept & Reference Explanation	#	Key	Core Concept & Reference Explanation
1	B	The cockpit discipline transfers; only switch locations change between vendors.	26	A	Proving loops require known past cases where ground truth is already verified.
2	B	Choose capability by task; avoid defaulting to expensive flagships for simple work.	27	B	Without known truth, plausible-sounding hallucinations cannot be detected.
3	A	Multi-step logic and calculation belong on the thinking tier, not fast defaults.	28	B	Passing historical tests earns confidence; it never transfers accountability.
4	B	Attachments provide grounding context; the model cannot read unattached files.	29	B	Learning that a task cannot be delegated reliably is a valuable proving outcome.
5	B	Rescue drifting chats by prompting for a state summary and pasting into a new chat.	30	A	Platform awareness means comparing outputs across models to evaluate strengths.
6	C	The project test: repeated background or re-uploaded files 3 times justifies a project.	31	B	Models hallucinate about their own UI; confident claims must be verified.
7	B	ChatGPT's project-only memory isolates context from wider account memories.	32	A	The beginner habit: Set the stage, define the task, and set the standing rules.
8	B	Claude switches to retrieval search over project files, expanding capacity ~10x.	33	A	Stored context compounds; review memory, instructions, and files monthly.
9	B	Stale context fails quietly; the assistant cannot know your files are outdated.	34	B	Data requires structured output (tables/sheets) for sorting and processing.
10	A	Instructions are for stable rules, memory for evolving facts, projects for scoped work.	35	B	Code execution is exact, but the code written by the model may sum wrong columns.
11	B	Incognito/temporary chats are memory-free, not trace-free (~30 day vendor retention).	36	B	Claude exposes thinking as an on/off switch; ChatGPT exposes named tiers.
12	B	Claude turns memory off by default on Team/Enterprise until admins enable it.	37	A	Workplace/Enterprise search queries internal permissioned connected apps.
13	B	OpenAI retired Canvas in current models, replacing it with inline writing/code blocks.	38	B	Rich context attachments improve outputs far more than elaborate prompt wording.
14	B	Artifacts keep distinct work products separate from conversational thread clutter.	39	A	Permissions govern existing actions; they cannot grant actions a tool lacks.
15	A	Critical numbers must be code-executed, and the columns verified by humans.	40	B	Automating proving loops across test suites is Eval-Driven Development.
16	A	Projects store knowledge (the what); skills perform repeatable tasks (the how).	41	B	A project holding chats, knowledge, and instructions is a System of Context.
17	B	Agent Skills standard uses serverless Markdown folders, making methods portable.	42	A	Check directories (claude.ai/directory) first before building custom tools.
18	B	Custom GPTs are configured standalone personas; skills are modular task methods.	43	B	A state summary records decisions made, open items, and active constraints.
19	B	Plugins bundle skills, external connectors/apps, and commands for a job function.	44	B	Audit memories to prevent embarrassing or private context appearing on screen.
20	B	MCP is an open standard connecting assistants to outside tools via a shared interface.	45	B	The prompt interview pattern has the assistant ask questions to surface context.
21	A	The plug is standard, but permissions govern sensitive data access and actions.	46	A	Model names rotate frequently; the 3-tier capability pattern endures.
22	B	Security audit questions: what can it read, what can it write/send, who built it?	47	B	Research mode conducts multi-step autonomous investigations taking minutes.
23	B	Misrouting: using deep research for a quick factual lookup wastes time and resources.	48	B	Provide the material a capable human colleague would need before answering.
24	B	Complex comparative analysis over existing files requires thinking mode, not search.	49	B	Artifacts isolate outputs, but the professional still owns accountability and review.
25	B	Research mode extends reach, but accountability remains human; verify citations.	50	A	Interface switches and tools vary over time; the core discipline is the constant.