

Skills & Connectors: Advanced Assessment

FOUNDATIONS CERTIFICATION & COMPREHENSIVE DIAGNOSTIC

Instructions: This examination tests the application of Skills, Connectors, Model Context Protocol (MCP), and the Agent Skills open standard. Questions reflect realistic operational scenarios. Where marked with **[Select 2 answers]** or **[Select 3 answers]**, select all correct options. The complete answer key is located on Pages 14–15.

1. An executive at an accounting firm is training new associates on AI workflow architecture. She notes: "When you type a single prompt asking for an analysis, you're shouting an order across the counter. The AI delivers it once and forgets your preferences. But when we configure our operating layer properly, AI acts like a complete commercial kitchen." **In the course's professional kitchen analogy, what do Connectors, Skills, and the Human User represent?**

- A. Connectors are the recipe cards, Skills are the ingredients in the pantry, and the Human is the customer waiting across the counter.
- B. Connectors are the kitchen itself (stove, knives, and real stocked ingredients), Skills are the recipe cards, and the Human is the cook who tastes and evaluates the dish.
- C. Connectors are the waitstaff delivering orders, Skills are the kitchen appliances, and the Human is the restaurant owner managing billing tokens.
- D. Connectors are storage freezers, Skills are the raw vegetables, and the Human is an observer who never reviews or touches the final output.

2. A digital content creator wants to build a Skill that checks standard English grammar and fixes missing commas in short paragraphs. A senior colleague advises him that this task fails the "market-of-one test." **What is the market-of-one test, and why does standard grammar checking fail it?**

- A. A task earns a skill only if it requires write access to at least three external applications; grammar checking requires only read access.
- B. A task earns a skill only if its underlying instruction set costs less than one dollar per execution across consumer models.
- C. A skill earns its place when it captures your own repeated, specific way; a task an application already does for everybody is just a good prompt.
- D. A skill earns its place only after it has been reviewed and certified by the Agentic AI Foundation.

3. A corporate attorney wants AI to always reply in formal British English, use metric units for foreign litigation, and avoid slang across *every single chat* she opens. However, she only wants her firm's proprietary 8-part merger checklist used when she specifically evaluates an acquisition agreement. **According to the four-feature framework, how should she split these requirements?**

- A. Place the British English and metric preferences in Custom Instructions, and build the 8-part merger checklist as a dedicated Skill.
- B. Bundle both requirements into a single master Skill, because Skills have automatic override authority over all settings.
- C. Place both requirements into Custom Instructions, because Custom Instructions handle all procedural checklists across workspaces.
- D. Configure an interactive Connector for British English, and create a Project workspace for the merger checklist.

4. A management consultant oversees three distinct enterprise clients: Alpha Corp, Beta LLC, and Gamma Inc. Each client has 25 permanent policy PDFs, historical contracts, and account guidelines that must load immediately into every conversation regarding that client. However, all three clients require the exact same four-section "Executive Risk Audit" format whenever a risk review is requested. **How should the consultant structure this environment?**

- A. Build three separate Skills (one for each client) and place the Executive Risk Audit template into global Custom Instructions.
- B. Create three separate Projects (one per client) to hold the standing policy files, and create one single Skill for the Executive Risk Audit format that runs across all of them on demand.
- C. Store all client files inside a single Project workspace, relying on negative triggers in the prompt to separate client data.
- D. Use Connectors to store the permanent policy PDFs, and configure Projects to format the risk audits on demand.

5. An enterprise IT manager discovers that a power user has installed 42 custom skills in Claude.ai. The manager worries that opening any conversation will become sluggish because the AI must process the full instructions for all 42 skills inside the context window at all times. **Why does progressive disclosure prevent this performance degradation?**

- A. The AI keeps only the short description of each skill in view (Level 1) and opens the full instructions (Level 2) only when your request matches.
- B. The Model Context Protocol compresses all 42 skills into local binary vectors that consume zero context tokens.
- C. Claude automatically disables all but the 5 most recently created skills in personal lists.
- D. Skills are offloaded to an external proxy server that executes commands without passing instructions to the language model.

6. A clinical research director identifies two recurring operational bottlenecks: (1) Associates constantly copy-paste lab tables out of Google Drive and message threads out of Gmail into chat; (2) Associates constantly re-explain their required four-part SOAP clinical note format, allergy sorting rules, and medication grouping. **According to the course diagnostic guide, what solution does each bottleneck require?** [Select 2 answers]

- A. Bottleneck 1 (fetching lab tables and message threads from Drive and Gmail) indicates the need for Connectors.
- B. Bottleneck 1 indicates the need for Custom Instructions with write access.
- C. Bottleneck 2 (re-explaining the four-part SOAP clinical format and sorting rules) indicates the need for a Skill.
- D. Bottleneck 2 indicates the need for an Interactive Connector configured for load-on-demand.

7. A product manager in Claude.ai types: "Create a PowerPoint about our Q3 logistics results," without typing a slash command or naming any skill. The PowerPoint skill fires automatically. **How does skill triggering operate in Claude.ai compared to desktop environments like Cowork?**

- A. Claude.ai scans your request against each enabled skill's description and fires the match automatically, whereas Cowork and Office add-ins allow typing `/` to browse and pick a skill from a menu.
- B. Claude.ai requires slash commands for all custom skills, but built-in document skills are hardcoded exceptions.
- C. PowerPoint skills only trigger automatically if the user pins them as their primary active application.
- D. Claude.ai cannot fire skills automatically; the user must have configured a background browser script.

8. A marketing specialist creates a skill named `brand-voice` to format social media copy. In a busy conversation, she asks Claude: "Write a short post about our quarterly software release." Claude outputs a generic post containing three exclamation marks and the word "game-changer," ignoring her rules. **Without editing the skill or starting a new conversation, what universal override works across all tools?**

- A. Type `OVERRIDE: brand-voice` at the very end of the prompt in capital letters.
- B. Explicitly name the skill in plain English: "Use my brand-voice skill to draft this caption."
- C. Disable "Allow network egress" under Capabilities in Settings.
- D. Delete the previous three conversation messages to reset the context scanner.

9. A financial analyst opens Settings > Capabilities in Claude.ai and observes that the toggle labeled "Code execution and file creation" is switched off. **What happens if she requests a formatted financial spreadsheet while this toggle remains disabled?**

- A. Claude will generate the spreadsheet using the Google Sheets connector instead.
- B. The Excel skill will fail to produce a finished spreadsheet file, because this toggle is the underlying engine required for document skills.
- C. Claude will save the raw data to her local hard drive as a `.csv` file automatically.
- D. Claude will prompt the analyst to install the `mcp-builder` skill from the public directory.

10. A new user opens Customize > Skills from Claude's left sidebar. She expects to see the built-in Word, Excel, and PowerPoint skills listed there, but finds only `skill-creator`. **Why are the document creation skills absent from this list?**

- A. Document skills were deprecated and replaced by the Agent Skills open standard in late 2025.
- B. Document skills live in the capabilities engine behind a single toggle, not in the personal skills list; `skill-creator` is there to help build custom skills.
- C. Document skills only appear after linking an active Microsoft 365 custom connector.
- D. The user is on a free tier, and built-in document skills only display on Enterprise plans.

11. A hospital administrator connects Google Drive to Claude so staff can summarize patient intake forms. Board members worry that connecting Drive will grant the AI visibility into confidential hospital personnel records and financial statements. **What does the course establish regarding connector access boundaries?**

[Select 2 answers]

- A. AI strictly inherits the permissions the user already possesses; if an account cannot open a file, neither can the connector.
- B. Connecting an app copies its entire directory structure onto Anthropic's model training servers.
- C. You control the scope per conversation; scoping access to a single designated folder is safer than connecting an entire drive.
- D. Connectors automatically elevate permissions to administrator status to allow index searching.

12. A logistics manager connects Claude to Slack and Linear. An assistant suggests immediately granting write access so the AI can automatically close tickets, update records, and send team-wide announcements. **What does the safe-use checklist recommend regarding permissions for newly connected applications?**

- A. Start read-only (such as searching and summarizing), granting write access only after observing that the tool behaves reliably.
- B. Always grant write access immediately so the model does not trigger permission errors during live runs.
- C. Keep third-party apps write-only, reserving read-only permissions exclusively for built-in Anthropic tools.
- D. Grant full access only if the connector was compiled using a terminal CLI.

13. An author connects Google Drive to Claude and issues the command: "Open Chapter 3 in my Drive, replace the second paragraph with this revised dialogue, and save the changes directly inside that file." Claude explains it cannot edit the document in place, but offers to save a new file to Drive. **What explains Claude's response?**

- A. The author's account ran out of token credits required for real-time document writing.
- B. This is the shape of the tool: the Google Drive connector reads files and saves new files, but cannot edit existing documents in place.
- C. The author failed to include the `SKILL.md` frontmatter trigger for file updating.
- D. The Google Drive API forbids AI models from accessing version history unless granted Enterprise Admin privileges.

14. A software architect has enabled 14 different connectors simultaneously in a single conversation, including GitHub, Jira, Confluence, Slack, Asana, Drive, Notion, and Figma. During complex architectural chats, she notices the conversation feels sluggish and tool choices occasionally misfire. **What setting manages this load?**

- A. Toggle off "Code execution and file creation" to compress connector bandwidth.
- B. Adjust the "Tool access" setting to "load-on-demand," which loads an action only when a request specifically calls for it.
- C. Convert all 14 connectors into Projects to isolate their token memory.
- D. Add negative triggers in Custom Instructions banning more than two connectors per prompt.

15. A design team integrates Figma with Claude. When a designer asks Claude to analyze an interface component, an interactive visual panel renders directly inside the chat window, allowing live component manipulation. **What is this feature called, and what distinguishes it from standard connectors?**

- A. An Interactive Connector, which renders a live interface panel inside the chat instead of returning text alone.
- B. A Progressive Skill, which executes graphical code in the user's terminal.
- C. A Custom Instruction Widget, which bypasses the Model Context Protocol.
- D. A Client-Side Script, which downloads raw design assets to the local hard drive.

16. A clinical team wants to automate consultation summaries using a SOAP-format skill and a Google Drive connector. Before deploying this workflow on real medical charts, they conduct a pilot test. **What best practices does the course outline for safely testing workflows involving regulated data (such as healthcare, legal, or financial records)?** [Select 2 answers]

- A. Run the entire end-to-end loop first using completely invented, fake demo files (such as DEMO-Okafor-2026-03).
- B. Test the workflow on real patient records, but use a personal Gmail account rather than a clinic account.
- C. When handling real records later, operate strictly within accounts and workspaces your organization has formally approved.
- D. Strip the description field from the SKILL.md so the skill cannot access the internet during live clinical trials.

17. An accountant handles the monthly close for DEMO Trading. Previously, she spent two hours downloading the ledger, formatting currency, grouping line items by expense head, flagging items above the \$50,000 withholding threshold, and drafting a four-section report. **When she pairs a "client-summary" Skill with a Google Drive Connector, what does each part do?**

- A. The Skill fetches the ledger from Drive, while the Connector formats currency and expense heads.
- B. The Connector fetches the ledger from Drive, the Skill formats amounts and enforces the four-section layout, and the accountant reviews the final result in two minutes.
- C. The Connector writes the accounting formulas, while the Skill manages the cloud storage authentication tokens.
- D. The Skill creates a Custom GPT, and the Connector runs a Python script in terminal to email the client directly.

18. A marketer wants to draft social captions for three scheduled announcements stored in a Notion calendar. Her company enforces a strict brand voice (no exclamation marks, question-hook openings, and a ban on buzzwords like "game-changer"). **What is the most efficient configuration for this weekly batch?**

- A. Export Notion to CSV, paste the text into chat, and re-type the brand voice guidelines every Monday.
- B. Connect Notion via a Connector, install a brand-voice Skill, and type a single prompt asking Claude to draft captions for the scheduled Notion posts in the brand voice.
- C. Place the Notion API token in Custom Instructions alongside the brand rules.
- D. Store the Notion calendar in a Project and use the `/` slash command in Claude Code to execute a data scrape.

19. A high school teacher wants to build a skill that formats raw class notes into topic headings, bold key terms, and three review questions. The teacher has zero coding or programming experience. **What is the recommended path for building this skill?**

- A. Enroll in a software development course to learn how to write Agent Skills in Python.
- B. Use the built-in `skill-creator` skill by describing the formatting rules in plain English, answering its clarifying questions, and clicking "Save skill."
- C. Download an unverified `SKILL.md` from a public developer forum and edit the hexadecimal tokens.
- D. Submit a support ticket requesting that Anthropic engineers hardcode the template into the Capabilities switch.

20. An operations director asks `skill-creator` to build an inventory variance skill. The AI outputs a finished skill folder containing a `SKILL.md` file, an asset template, and a small Python script for mathematical grouping. The director never reads or writes a line of code. **The course calls this "Code you never write." What is the director's role in this workflow?**

- A. The director is the contractor, responsible for debugging script syntax when errors occur.
- B. The director is the client: writing a clear brief, checking the output against measurable benchmarks, and directing AI to make corrections.
- C. The director is a passive observer who must accept whatever output the model generates without validation.
- D. The director is a systems administrator who must re-license the code under the Agentic AI Foundation.

21. A financial consultant is building a comprehensive `audit-helper` skill folder. The skill requires step-by-step reporting instructions, an official 40-page regulatory guideline document that AI should only read when needed, and a blank spreadsheet template. **According to the anatomy of a skill, where should each component be stored?** [Select 2 answers]

- A. The 40-page regulatory guideline document belongs in the `references/` subfolder (progressive disclosure Level 3).
- B. The blank spreadsheet template belongs in the `assets/` subfolder.
- C. The blank spreadsheet template belongs in the `scripts/` subfolder alongside executable programs.
- D. The 40-page regulatory guideline document must be pasted entirely into the frontmatter description.

22. A junior analyst packages a skill folder and uploads it to Claude. The upload is rejected with a validation error. **Which configuration violates the strict naming rules established in the course?**

- A. The folder is named `client-monthly-summary` and contains a file named `SKILL.md`.
- B. The folder is named `claude-tax-calculator` and its description contains XML angle brackets (` `).
- C. The frontmatter uses three dashes (`---`) at the top and bottom of the metadata block.
- D. The instructions block is written in plain English markdown under `# Instructions`.

23. A developer writes a skill containing 12 pages of thorough clinical instructions. However, in the frontmatter, he writes: `description: Helps doctors with medical paperwork.` In practical testing, the skill never activates when clinicians paste patient notes. **Why did this skill fail to fire?**

- A. The instructions exceeded the token capacity limit of progressive disclosure Level 2.
- B. AI does not read instructions to judge relevance; it reads only the description, so a vague description without trigger words means the skill never fires.
- C. Clinical skills require an active write connector to Google Drive before they can load into memory.
- D. Claude automatically suppresses skills that contain the word "paperwork" in their metadata.

24. A sales director wants to create a skill that converts raw call transcripts into CRM lead summaries. **Which description follows the recommended formula from the course?**

- A. `description: Handles sales stuff and notes.`
- B. `description: A machine-learning tool that interfaces with enterprise CRM platforms to streamline communications.`
- C. `description: Converts sales call transcripts into structured CRM lead summaries with deal size, key decision-makers, and next steps. Use when the user asks for a 'lead summary', 'call write-up', or 'CRM note'`.`
- D. `description: Use whenever the user mentions sales, leads, revenue, deals, clients, or calls.`

25. An accountant creates an `expense-audit` skill. She wants to verify that Claude understands exactly when it should and should not load the skill. **What debugging prompt does the course recommend asking Claude right after installing the skill?**

- A. "Print the compiled memory address of my expense-audit skill."
- B. "When would you use my expense-audit skill, and when would you NOT use it?"
- C. "Run an automated security scan against the frontmatter of my skill."
- D. "Compare my expense-audit skill against the top 10 skills in the directory."

26. A software firm creates a skill called `release-notes-writer`. However, engineers discover that whenever they ask Claude for a quick one-sentence commit message for a bug fix, the skill takes over and generates a massive, four-page public marketing document. **How should the author modify the `SKILL.md` file?**

- A. Delete the `scripts/` folder from the skill repository.
- B. Add a negative trigger to the description: "Do NOT use for one-off commit messages or short bug summaries. Use only for full version release notes."
- C. Convert the skill into a Project workspace and lock it with a password.
- D. Switch the conversation tool access setting to "read-only."

27. An educator drafts a skill for generating weekly biology review questions. Before rolling it out, she conducts testing using the course's recommended build loop. **What actions are critical components of testing a new skill?** [Select 3 answers]

- A. Read the draft instructions and ask: "Would two colleagues who read this do the exact same thing?"
- B. Test triggering by trying phrases that should start it, and test unrelated requests to confirm it does not take over.
- C. Test awkward edge cases, such as an input with zero data or transactions exactly on a boundary threshold.
- D. Export the skill immediately to an OpenAI Custom GPT to verify server synchronization.

28. A user builds a `travel-itinerary` skill. In a brand-new conversation, on the very first message, she pastes flight details and requests an itinerary. The skill fires, but the output omits departure times and uses the wrong time zone. **According to the troubleshooting timing list, what does failure on the first reply indicate?**

- A. The chat context window has degraded due to token overflow.
- B. The problem is in the skill file itself; an instruction is loose, missing, or assumed context that was never written down.
- C. The user's browser cache is corrupt and must be cleared.
- D. The Model Context Protocol dropped the connection during generation.

29. A researcher uses a `survey-formatter` skill. For the first four survey files, the output matches the required schema perfectly. By the eighteenth survey file in that same single conversation, the AI begins dropping column headers and mangling labels. **What is the cause and recommended fix?**

- A. The `SKILL.md` file contains a syntax error; edit the file immediately.
- B. The skill's description was too broad; add a negative trigger to the frontmatter.
- C. The chat got long and the AI lost track; the skill itself is fine, so do not edit the file—simply start a fresh chat and restate the task.
- D. The connector reached its daily API call quota.

30. A payroll manager creates a skill to calculate overtime pay based on local labor regulations. Even after rewriting the instructions four times to make them as clear as possible, the AI consistently makes math errors on split-shift calculations. **Why did tightening the text fail, and what is the proper solution?**

- A. English sentences are interpreted probabilistically; exact calculations belong in a small program inside the `scripts/` folder, because code is deterministic.
- B. The payroll manager must switch from Claude.ai to Gemini Gems.
- C. The user must type `/calculate` before every prompt to force strict mathematical processing.
- D. The skill must be converted into a Project to provide permanent mathematical memory.

31. A legal department's `contract-check` skill operated with 100% accuracy for five months. Suddenly, in August, the skill begins missing mandatory compliance flags on supplier agreements, despite no changes being made to the `SKILL.md` file. **What does the timing list indicate when a skill "used to work and now does not"?**

- A. The AI model's training weights expired during a routine infrastructure maintenance window.
- B. The skill did not change, so something around it did—such as the firm's compliance threshold, a vendor template, or team processes.
- C. The Agentic AI Foundation revoked the skill's security license.
- D. The skill exceeded its cumulative six-month token quota.

32. A COO builds a company-wide operational reporting skill using ``skill-creator``. She wants all 60 regional managers to use this exact skill so reports across the enterprise remain uniform. **How does skill sharing function on Team and Enterprise plans?**

- A. Sharing is on by default; any team member can edit the original file directly at any time.
- B. The COO must export the skill as a zip file and email it to each manager to install manually.
- C. The COO publishes the skill to the organization's directory; shared skills are view-only for employees and auto-update whenever the COO edits the original.
- D. Shared skills run only while the author's local workstation is connected to the corporate network.

33. A technology firm employs non-technical copywriters working in web browsers, senior analysts working on desktop machines with local files, and engineers working in terminals. **According to the "Five Tools, One Skill" framework, how does the course map these user groups to tools supporting the Agent Skills standard?** [Select 2 answers]

- A. Non-coders and general knowledge workers start and stay in Claude.ai (web/mobile).
- B. Knowledge workers who need local desktop file integration without terminal coding can use Cowork (or the open-source OpenWork).
- C. Terminal-based engineers are required to use ChatGPT Custom GPTs because CLI tools cannot parse markdown.
- D. The Agent Skills open standard is strictly exclusive to Claude Code and cannot run on desktop applications.

34. A data engineering team writes five complex skills in Claude.ai. When corporate procurement adopts a multi-model strategy, the team asks whether these skills will become obsolete if they migrate to terminal tools from other AI vendors. **What makes a basic ``SKILL.md`` portable across different vendors' developer tools, such as OpenAI's Codex CLI and Google's Gemini CLI?**

- A. The Agent Skills open standard (agentskills.io), published in December 2025, established an open format that multiple vendors agreed to read.
- B. OpenAI and Google acquired the rights to Claude's proprietary backend architecture.
- C. All AI CLI tools run on the same central server managed by Anthropic.
- D. Model Context Protocol forbids proprietary skill formats under international copyright law.

35. An enterprise architect evaluates ChatGPT's "Custom GPTs", Google's "Gems", and Anthropic's "Agent Skills" for the company's long-term workflow automation strategy. **What crucial architectural difference does the course emphasize regarding cross-vendor portability?**

- A. Custom GPTs and Gems are open standards that run everywhere, whereas Agent Skills only run inside Claude.ai.
- B. Custom GPTs and Gems are locked to their respective vendors' consumer apps, whereas Agent Skills are built on an open standard that travels across multi-vendor tools.
- C. Custom GPTs require write access to Google Drive, while Gems require write access to GitHub.
- D. There is no difference; all three represent identical implementations of the MCP specification.

36. A chief risk officer asks who governs the plumbing beneath Connectors (MCP) to ensure the company is not locking its systems into closed, single-vendor infrastructure. **According to the course, what entity provides neutral, cross-vendor governance for MCP?**

- A. The World Wide Web Consortium (W3C).
- B. The Agentic AI Foundation, formed under the Linux Foundation.
- C. A private joint venture consisting solely of OpenAI and Microsoft.
- D. The Federal Communications Commission (FCC).

37. An employee downloads an unverified skill called `auto-expense-tracker.zip` from a social media forum and uploads it into his company workspace. **What are the two primary security risks associated with running unvetted skills from strangers?**

- A. Token inflation and automated billing tier upgrades.
- B. Prompt injection (hidden instructions pushing AI to unauthorized acts) and data exfiltration (secretly transmitting data to external servers).
- C. Accidental deletion of the company's domain name and DNS spoofing.
- D. Automatic conversion of read-only connectors into paid Enterprise licenses.

38. An administrator configures a Google Drive connector for an executive assistant who only handles incoming event invoices stored in a single folder called `/2026-Invoices`. **What principles regarding "Scope" does the course emphasize for safe connector usage?** [Select 2 answers]

- A. Scope is safety: always scope access to the smallest folder or dataset the task actually needs.
- B. Scope is a Skill property, meaning you set the folder permissions inside the `SKILL.md` frontmatter.
- C. Scope is a Connector idea: narrowing a connector to one folder prevents the AI from reaching the rest of your drive.
- D. Connectors do not support scoping; you must disconnect and reconnect the app after every chat.

39. A developer attempts to add a line to her `SKILL.md` frontmatter: `scope: /Restricted/PayrollOnly`. She discovers the platform ignores this line completely. **Why does a Skill have no scope dial?**

- A. The syntax required `kebab-case`, such as `scope: payroll-only`.
- B. A Skill has no scope dial because it executes with whatever access its chat already has; its reach is governed by the connectors enabled in that session.
- C. Scope dials on skills are only enabled after purchasing an Enterprise API key.
- D. Skills automatically default to full-disk read-write permissions regardless of settings.

40. A project management team plans to give an AI connector write access to update a Jira board and edit files in a shared SharePoint document library. **What does the safe-use checklist mandate verifying before enabling write access on mission-critical assets?**

- A. Verify that the skill author has a verified profile on GitHub.
- B. Check recovery mechanisms (such as version history and undo) because some write actions bypass standard recovery paths.
- C. Confirm that the skill has executed at least 1,000 error-free runs in read-only mode.
- D. Ensure that the AI model is running on local hardware rather than cloud servers.

41. An associate is handed a skill folder from an external contractor. Before enabling it, she wants to audit its contents without learning how to inspect Python or JSON code. **What specific prompt does the course advise pasting into Claude to audit the skill?**

- A. "Run a malware scan on this attachment and send the log to my antivirus software."
- B. "Read this skill and tell me, in plain language, exactly what it instructs you to do. Flag anything that contacts an external server, handles credentials, or could send my data anywhere."
- C. "Translate this skill into a Custom GPT and verify if OpenAI approves of the security policies."
- D. "Execute this skill in debug mode and report the memory address allocations."

42. A university student uses the free plan of Claude.ai. He wants to know what limitations exist regarding connectors. **What are the connector entitlements on Claude's free tier?**

- A. Zero connectors are allowed on the free plan; all connectors require a Pro subscription.
- B. The free tier allows ready-made directory connectors and includes one custom connector.
- C. The free tier offers unlimited custom connectors, but blocks Google Drive.
- D. Connectors can only be accessed through terminal CLI commands on the free tier.

43. A systems architect explains to her team how Claude manages token economy and efficiency when working with complex skills. **Which statements correctly describe the three levels of progressive disclosure in Agent Skills?** [Select 3 answers]

- A. Level 1: Frontmatter name and description (always loaded in view to judge relevance).
- B. Level 2: Instructions body of `SKILL.md` (loaded dynamically only when the description matches).
- C. Level 3: Optional subfolders such as `references/`, `assets/`, and `scripts/` (loaded or executed only as needed).
- D. Level 4: The user's entire local hard drive (indexed continuously in the background).

44. A user is writing a single, one-off congratulations message to a retiring colleague who enjoys fly fishing. The user wonders whether to build a "Retirement Note Skill" or connect their Outlook app. **According to the friction diagnostic framework, what does this task need?**

- A. A Skill, because tone and warmth must be standardized.
- B. A Connector, because Outlook holds past emails from the retiree.
- C. Neither a Skill nor a Connector; a one-off task with no repeated rules and no outside data is just a good prompt.
- D. Both a Skill and a Connector, because writing personal messages carries reputational risk.

45. An office administrator connects Google Calendar to Claude. She is frustrated because Claude can look up upcoming events, but refuses to reschedule meetings when requested. She files a bug report claiming permissions are broken. **What does the course suggest doing before building a recurring habit around a connector?**

- A. Ask the connector in plain language: "What can you actually do through this connector, and what can't you do?"
- B. Upgrade immediately to an Enterprise seat to unlock hidden API endpoints.
- C. Edit the `SKILL.md` file to force calendar rescheduling.
- D. Clear browser cookies and reinstall the calendar plugin.

46. A corporate security officer compares the governance model for Skills versus Connectors. **Why is "reading the file first" the primary control for Skills, while "starting read-only and scoping tightly" is the primary control for Connectors?**

- A. Skills have no scope dials and run with all the permissions of the chat session, whereas Connectors explicitly control external system access and action boundaries.
- B. Skills can execute low-level operating system commands, while Connectors only format text.
- C. Connectors run entirely in plain text, while Skills are compiled into binary bytecode.
- D. Skills are open source, whereas Connectors are strictly closed-source commercial APIs.

47. A student creates four different folders on their computer to store newly authored skills. **Which folder name strictly adheres to the formatting rules required for Agent Skills?**

- A. `Weekly Study Notes`
- B. `weekly\_study\_notes`
- C. `weekly-study-notes`
- D. `WeeklyStudyNotes!`

48. A user has been interacting with a complex data analysis skill for 90 minutes. The conversation has accumulated 45 messages, multiple file uploads, and repeated revisions. The AI suddenly begins repeating previous mistakes and hallucinating columns. **According to the course troubleshooting advice, what is the most cost-effective solution?**

- A. Spend the next hour prompting the AI to identify where its memory broke down.
- B. Re-write the `SKILL.md` frontmatter to double its length.
- C. Abandon the tangled chat, start a fresh conversation, and restate the task in two lines.
- D. Reinstall the browser extension and delete the connector.

49. A team lead completes Project 4 ("Make it portable, or hand it off") from the crash course. She wants to prove that her new `code-refactor` skill is truly robust and self-contained. **What is the definitive test of success when handing off a skill to a colleague?**

- A. The colleague successfully reads the raw markdown file and understands the grammar.
- B. The colleague runs the skill cold with no verbal explanation from the author, and it produces the correct, standardized output.
- C. The colleague purchases a paid Anthropic subscription to run the file.
- D. The colleague imports the skill into a spreadsheet without errors.

50. A business analyst reflects on how his work has evolved over the year. In 2024, he treated AI like a vending machine—typing a fresh prompt every morning, pasting the same instructions, and re-copying spreadsheets. In 2026, he manages automated pipelines. **What core transformation do Skills and Connectors achieve together?**

- A. They replace human critical thinking by allowing AI to make final financial and legal decisions without oversight.
- B. They transform AI from an isolated chat box you type into, into an operating layer that safely interacts with your real tools and reliably applies your specific standards.
- C. They turn everyday users into professional software engineers who write production code in terminal interfaces.
- D. They eliminate the need for cloud security protocols by storing all company data in markdown text files.

Examination Complete

You have completed all 50 diagnostic scenario questions. Turn to Page 14 and Page 15 for the comprehensive answer key and explanatory breakdowns derived directly from the foundational curriculum.

Answer Key & Explanations (Part 1)

QUESTIONS 1 TO 25 REFERENCE GUIDE

1. **B** — Connectors represent the kitchen and pantry (providing tools and real ingredients), Skills are the recipe cards (teaching your exact method), and the human cook tastes and reviews the dish.
2. **C** — The market-of-one test states that a skill earns its place only when it captures your specific repeated way; general tasks that an application already performs for everyone are just prompts.
3. **A** — Custom instructions govern global preferences across all chats ("always British English"), while Skills encapsulate on-demand procedures ("use this 8-part merger checklist").
4. **B** — Standing client documents belong in dedicated client Projects (loaded permanently per workspace), while the repeatable audit method belongs in a reusable Skill that runs across them.
5. **A** — Progressive disclosure keeps only Level 1 metadata (frontmatter name and description) in active view, opening Level 2 instructions only when a prompt triggers a match.
6. **A, C** — Fetching external data from Drive and Gmail requires Connectors (A); enforcing procedural format and clinical rules requires a Skill (C).
7. **A** — In Claude.ai, skills fire automatically when prompt wording matches their description. In Cowork and Office add-ins, users can type `/` to browse and manually select a skill.
8. **B** — Explicitly naming the skill in plain English ("Use my brand-voice skill...") is the universal override supported across all tools and web interfaces.
9. **B** — The Word, Excel, and PowerPoint skills depend entirely on the "Code execution and file creation" toggle under Capabilities; without it, document generation fails.
10. **B** — Document skills live directly inside the capabilities engine, not the user customization list. The Customize > Skills menu displays user-installed skills and `skill-creator`.
11. **A, C** — Connectors strictly inherit user account permissions (A) and can be restricted to specific folder scopes per conversation (C). They never mirror entire cloud drives or elevate privileges.
12. **A** — The safe-use checklist advises starting read-only (search/summarize), granting write access only after establishing operational trust with the tool.
13. **B** — The Google Drive connector is engineered to read files and save new files, but cannot perform in-place edits on existing cloud documents.
14. **B** — Setting Tool access to "load-on-demand" ensures connector actions are only loaded when specifically needed, preventing cognitive and performance degradation.
15. **A** — Interactive connectors draw live, interactive interface panels inside the chat window rather than returning plain text strings.
16. **A, C** — Regulated workflows should be piloted end-to-end on invented demo data (A), and live processing must occur only within organization-approved accounts (C).
17. **B** — Connectors fetch data, Skills format and structure the output, and the human performs the critical two-minute review.
18. **B** — Pairing a Notion Connector (access) with a brand-voice Skill (expertise) allows batch content generation via a single plain-English prompt.

Answer Key & Explanations (Part 2)

QUESTIONS 26 TO 50 REFERENCE GUIDE

26. B — A negative trigger explicitly instructs the AI when NOT to load a skill (e.g., "Do NOT use for one-off commit messages"), preventing over-triggering.

27. A, B, C — Robust validation requires reading instructions for ambiguity (A), testing triggering and non-triggering phrases (B), and pressure-testing edge cases (C).

28. B — A skill failing on the very first response indicates an internal file flaw: loose rules, missing guidelines, or assumed context.

29. C — Performance degradation over an extended session stems from context bloat. The file is fine; start a fresh conversation to reset context.

30. A — Deterministic tasks like exact arithmetic or split-shift payroll belong in executable programs inside the `scripts/` folder, not probabilistic text instructions.

31. B — When an unedited skill suddenly misbehaves, external factors have shifted (updated templates, new data schemas, or altered business rules).

32. C — Organization directory sharing provides view-only access to colleagues and automatically propagates updates whenever the author alters the master skill.

33. A, B — Claude.ai serves web/mobile knowledge workers (A), while Cowork/OpenWork provides desktop integration for local files without terminal coding (B).

34. A — The Agent Skills open standard (agentskills.io) allows standard `SKILL.md` files to be executed across multi-vendor tools, including OpenAI Codex CLI and Google Gemini CLI.

35. B — Custom GPTs and Gems are proprietary and vendor-locked, whereas Agent Skills provide portable cross-vendor compatibility.

36. B — MCP governance was established under the Agentic AI Foundation within the Linux Foundation to guarantee neutral, open-source infrastructure.

37. B — Unvetted third-party skills present severe security risks: prompt injection (hijacking behavior) and data exfiltration (secretly routing data outward).

38. A, C — Scope is a Connector concept (C), and constraining access to the narrowest necessary folder is an essential security posture (A).

39. B — Skills have no scope dials; they operate with whatever permissions and connector access the active chat session possesses.

40. B — Before granting write access, check that version history and undo paths exist, as automated write actions can bypass normal file recovery.

41. B — Having AI read the `SKILL.md` and explicitly flag external server calls, credential handling, or data transmission is the fastest non-technical safety audit.

42. B — Claude's free plan provides access to ready-made directory connectors and permits one custom MCP connector.

43. A, B, C — Progressive disclosure comprises Level 1 metadata (A), Level 2 instructions (B), and Level 3 resource folders (`references/`, `assets/`, `scripts/`) (C).