

CodeVerse Soban

Like · Subscribe · Stay Sharp

AI Agent Factory — Level 1 Exam

Topic: General Agents on the Web — A Crash Course

100 Tricky & Hard MCQs · 4 Options Each · Correct Answer + Explanation

Q1. The course opens by saying 'when you stop asking, the conversation stops doing.' What is the ONE test that separates a chat box from an agent surface?

- A) Whether the interface has a file-upload button
- B) Whether the response includes structured headings or bullet points
- C) 'If I stop typing, does the work stop?' — chat stops, agent surface keeps working
- D) Whether the tool can run Python code or execute commands

Correct Answer: C) 'If I stop typing, does the work stop?' — chat stops, agent surface keeps working

Explanation: This is the course's first concept: 'One test tells the two apart at any address, from any company: If I stop typing, does the work stop? Chat box: yes. Agent surface: no.'

Q2. Before July 2026, agent capabilities existed in browser AI tools. What specifically changed in mid-2026 that this course says made the difference?

- A) AI models became significantly smarter and more capable of autonomous work
- B) Sessions, files, connectors, scheduling, approvals, and file creation were assembled into one work surface beside chat — the change was packaging, not invention
- C) Browser security restrictions were lifted allowing AI to access local files directly
- D) Government regulations required AI companies to offer unattended automation features

Correct Answer: B) Sessions, files, connectors, scheduling, approvals, and file creation were assembled into one work surface beside chat — the change was packaging, not invention

Explanation: 'What changed in mid-2026 was packaging. The major assistants pulled sessions, files, connectors, scheduling, approvals, and file creation into one work surface beside chat.' Browser agents didn't appear from nowhere.

Q3. The course asks 'two questions instead of one' for understanding the August 2026 changes. What are the two questions, and why does this beat memorizing product buttons?

- A) Which vendor is better? And which plan should I buy? These beat product memorization by focusing on value
- B) Where does the agent loop run? And where does each tool execute? These beat button memorization because they reveal the actual architecture regardless of product changes
- C) Is the data in the cloud? And is the device online? These beat buttons because they explain security
- D) What is the model? And what is the context window? These beat buttons because capability determines what is possible

Correct Answer: B) Where does the agent loop run? And where does each tool execute? These beat button memorization because they reveal the actual architecture regardless of product changes

Explanation: 'This course asks two questions instead of one: Where does the agent loop run? Where does each tool execute? Those two questions beat memorizing a product button.'

Q4. A synchronous task runs while you wait and then waits for your next turn. What does 'delegated' mean in contrast, and why does this distinction matter?

- A) Delegated means the task is given to a human assistant rather than AI

- B) Delegated means the task keeps moving toward an outcome after you start it — it can run later, watch for a change, and does not need your next input to continue
- C) Delegated means the AI uses more computing resources and therefore costs more per run
- D) Delegated means the task requires pre-approval from your organization's IT department

Correct Answer: B) Delegated means the task keeps moving toward an outcome after you start it — it can run later, watch for a change, and does not need your next input to continue

Explanation: 'An agent run is delegated. Once you start it, it keeps moving toward an outcome. It can run later, or watch for a change.' This distinguishes it from synchronous chat which is turn-by-turn.

Q5. Someone types 'Every Monday, summarize these emails' into the chat box and a recurring task is set up. The course warns about a specific mistake here. What is it?

- A) Using the chat box for scheduling is not supported and will produce an error
- B) The words went into the chat box, but what was set up is a delegated task — watch what the work DOES, not where you typed it. Many beginners confuse the input location with the task type
- C) Email summarization requires a connector that must be set up separately before using chat commands
- D) Monday schedules require Manual approval mode and the beginner forgot to set it

Correct Answer: B) The words went into the chat box, but what was set up is a delegated task — watch what the work DOES, not where you typed it. Many beginners confuse the input location with the task type

Explanation: 'You often start a delegated run from the chat box, in plain words... So watch what the work does, not where you typed it.' The instruction goes into chat; the result is delegated work.

Q6. The course says the 'old picture' of cloud vs local is 'no longer enough.' What TWO-AXIS model replaces it?

- A) Axis 1: cost per run; Axis 2: quality of output — these together determine value
- B) Axis 1: where the agent loop runs (cloud vs local); Axis 2: where the tool executes (which may be different from where the loop runs)
- C) Axis 1: which vendor provides the service; Axis 2: which subscription plan covers the feature
- D) Axis 1: whether data is personal or business; Axis 2: whether the output is text or file

Correct Answer: B) Axis 1: where the agent loop runs (cloud vs local); Axis 2: where the tool executes (which may be different from where the loop runs)

Explanation: 'Use two axes instead. Axis 1: where does the loop run? (cloud or local) Axis 2: where does the tool execute? A cloud session may call a connector or, through a bridge, a browser or file tool on your computer.'

Q7. What happens to a cloud-plus-desktop-bridge workflow when the laptop closes?

- A) The entire workflow stops because the agent loop needs the laptop to function
- B) The cloud session continues with cloud-reachable tools, but loses the bridged tools that were borrowed from the desktop — the loop and the bridged tool have different availability
- C) The workflow pauses and automatically resumes when the laptop reopens
- D) A notification is sent to your phone asking whether to continue or stop the workflow

Correct Answer: B) The cloud session continues with cloud-reachable tools, but loses the bridged tools that were borrowed from the desktop — the loop and the bridged tool have different availability

Explanation: 'If the desktop goes offline, the cloud session continues with cloud-reachable tools, but loses the tools it was borrowing.' The bridge does not turn a cloud session into a local session.

Q8. The course lists four consequences worth memorizing about the cloud-desktop relationship. Which of these is one of them?

- A) Cloud sessions are always more expensive than local sessions due to vendor infrastructure costs

- B) A local file opened through a desktop bridge is processed in the vendor's cloud — 'the file came from my laptop' does not mean 'the file stayed on my laptop'
- C) Closing a tab always stops the work because the tab contains the runtime
- D) Local sessions are always more secure because data never leaves your machine

Correct Answer: B) A local file opened through a desktop bridge is processed in the vendor's cloud — 'the file came from my laptop' does not mean 'the file stayed on my laptop'

Explanation: 'A local file can be processed in the cloud. The file came from my laptop does not mean the file stayed on my laptop.' This is one of four key consequences about cloud-desktop architecture.

Q9. The course names six parts that every serious agent surface is built from. What are they?

- A) Model, prompt, context, output, feedback, memory
- B) Heartbeat, reach, run-until-done loop, state spine, human gate, body
- C) Schedule, connector, task, session, approval, file
- D) Design, brief, plan, execute, review, deploy

Correct Answer: B) Heartbeat, reach, run-until-done loop, state spine, human gate, body

Explanation: The six parts: 'Heartbeat (what starts a run), Reach (what it can read and act on), Run-until-done loop (how it keeps moving), State spine (what persists), Human gate (where autonomy stops), Body (where work executes).'

Q10. Claude Cowork's August 2026 changes expanded the 'state spine' significantly. Which of these correctly describes the August change to Cowork's memory?

- A) Memory was added for the first time — Cowork had no memory before August 2026
- B) Cloud Cowork tasks now use the same memory as Claude chat, so what Claude remembers from chat can reach a cloud Cowork task — inspectable and editable topic by topic
- C) Memory was restricted to Projects only — global memory was removed for privacy reasons
- D) Local Cowork sessions gained access to the same cloud memory as web sessions

Correct Answer: B) Cloud Cowork tasks now use the same memory as Claude chat, so what Claude remembers from chat can reach a cloud Cowork task — inspectable and editable topic by topic

Explanation: 'The important August change is memory inside Cowork. Cloud Cowork tasks now use the same memory as Claude chat... You can inspect memory topic by topic and edit or delete it. Local Cowork sessions do not use this cloud memory.'

Q11. The six-part lens includes 'body.' What does 'body' mean, and why does the course say asking where the loop runs AND where each tool executes matters?

- A) Body means the physical device running the AI — you must ask about both to estimate hardware requirements
- B) Body is where work actually executes — the loop can run in vendor cloud while a specific tool executes locally through a bridge, so they can be in different places
- C) Body means the size of the context window — larger body means more can be processed at once
- D) Body refers to the main instructions file (like CLAUDE.md) that defines how the agent behaves

Correct Answer: B) Body is where work actually executes — the loop can run in vendor cloud while a specific tool executes locally through a bridge, so they can be in different places

Explanation: 'Body: where the work actually executes.' A 'cloud sandbox, plus an optional local session and bridged desktop or browser tools.' A remote loop can borrow a local tool.

Q12. The course warns: 'Same shape does not mean same trust.' What example does it give?

- A) Two agents with the same model can produce different results on identical tasks
- B) A connector call and a browser click can reach the same business system with different permissions, visibility, and exposure to hidden instructions

- C) Two vendors with the same pricing can have different feature sets
- D) The same task can cost different amounts depending on which model tier is used

Correct Answer: B) A connector call and a browser click can reach the same business system with different permissions, visibility, and exposure to hidden instructions

Explanation: 'A connector call and a browser click can reach the same business system with different permissions, visibility, and exposure to hidden instructions.' Same shape, different trust.

Q13. The persistent context stack has five layers. What is the difference between 'session history' and 'semantic memory'?

- A) Session history is public; semantic memory is private and encrypted
- B) Session history records what happened in one specific task; semantic memory carries facts and preferences across cloud sessions until edited, reset, or turned off
- C) Session history is stored locally; semantic memory is stored in the cloud
- D) Session history is free; semantic memory requires a paid subscription

Correct Answer: B) Session history records what happened in one specific task; semantic memory carries facts and preferences across cloud sessions until edited, reset, or turned off

Explanation: 'Session history: What happened in this particular task? (one workstream) Semantic memory: What facts or preferences should carry across cloud sessions? (across sessions, until edited, reset, or turned off)'

Q14. An old Cowork session is not the same as memory. A Project is not memory. The course explains why people confuse these. What is the correct distinction?

- A) Old sessions and Projects cannot be searched, while memory can be queried with natural language
- B) All three persist but for different reasons — session history preserves task records, Project groups related work, and semantic memory carries facts across sessions. Calling them all 'memory' is the common beginner mistake
- C) Old sessions and Projects are free features; memory requires an enterprise plan
- D) Old sessions and Projects can be deleted by the vendor; only memory is permanently yours

Correct Answer: B) All three persist but for different reasons — session history preserves task records, Project groups related work, and semantic memory carries facts across sessions. Calling them all 'memory' is the common beginner mistake

Explanation: 'That fixes a common beginner mistake. An old session is not memory. A Project is not memory either. A global instruction is not memory. They all persist, for different reasons.'

Q15. The course gives FIVE layers of persistent context. A user asks: 'Should I put my important recurring instructions in the session history?' What is the correct answer and why?

- A) Yes — session history is the most reliable persistence layer
- B) No — session history records what happened in one task. Recurring instructions belong in standing instructions (for broad rules) or project instructions (for scoped work) — the narrowest layer that reliably applies
- C) Yes — but only if you are using a paid plan with extended history retention
- D) No — put recurring instructions in semantic memory so they transfer across sessions automatically

Correct Answer: B) No — session history records what happened in one task. Recurring instructions belong in standing instructions (for broad rules) or project instructions (for scoped work) — the narrowest layer that reliably applies

Explanation: 'Put each durable fact in the narrowest layer that reliably applies.' Standing instructions are for 'stable rules that apply again and again.' Session history is only 'the record of what happened in one task.'

Q16. What is Tier 1 in the three-tier file system, why is it 'never the home of anything you need tomorrow,' and what goes there?

- A) Tier 1 is the highest tier — most important files go here for permanent safe storage

- B) Tier 1 is the task filesystem — a temporary working space on vendor servers for drafts, calculations, and part-built files. It is cleaned up when work is done and never treated as storage
- C) Tier 1 is your local machine — files here are yours but risky because the device could fail
- D) Tier 1 is the backup tier — files here are copies of tier 2 and exist for disaster recovery only

Correct Answer: B) Tier 1 is the task filesystem — a temporary working space on vendor servers for drafts, calculations, and part-built files. It is cleaned up when work is done and never treated as storage

Explanation: 'Tier 1 is the task filesystem. While a remote session works, it uses temporary working space on the vendor's servers... It is cleaned up when the work is done... never treat it as storage. It is never the home of anything you need tomorrow.'

Q17. Tier 2 is 'safe today and hostage tomorrow.' What does this mean exactly?

- A) Tier 2 files are protected by encryption today but will expire after a fixed period
- B) Tier 2 files survive the task and the tab, but they are in the vendor's custody in the vendor's format behind the vendor's login — safe to access now, but not under your own retention rules
- C) Tier 2 files are accessible today but require a premium plan renewal to access tomorrow
- D) Tier 2 files are cached locally today but deleted from local storage and moved to vendor cloud overnight

Correct Answer: B) Tier 2 files survive the task and the tab, but they are in the vendor's custody in the vendor's format behind the vendor's login — safe to access now, but not under your own retention rules

Explanation: 'Tier 2 is platform storage... safe today, and it is also the vendor's custody, in the vendor's format, behind the vendor's login.' The 'hostage' metaphor means it exists only under the vendor's terms.

Q18. The course gives ONE LINE to memorize from the entire topic of file tiers. What is it?

- A) 'Always save to the cloud; local files are unsafe.'
- B) 'Finished work exits the platform. Everything else may stay.'
- C) 'Three tiers, three choices: temporary, platform, yours.'
- D) 'Never trust tier 1; always prefer tier 3.'

Correct Answer: B) 'Finished work exits the platform. Everything else may stay.'

Explanation: 'The discipline is one line, and it is the line to memorize from this whole course: Finished work exits the platform. Everything else may stay.'

Q19. In Ayesha's invoicing example, she saves the invoice to tier 3 TWICE — to Drive AND by email. Why does the course present this double-saving as deliberate and correct?

- A) Because tier 3 has a size limit and splitting large files across two destinations is best practice
- B) Because 'one day someone will ask, where is the March invoice? The answer must be the firm's records, and tier 3 is the only tier that reaches them' — two destinations ensure the firm's records are reliably populated
- C) Because double-saving prevents vendor lock-in by keeping copies in two different systems
- D) Because saving to Drive creates a backup and emailing provides an audit trail for compliance

Correct Answer: B) Because 'one day someone will ask, where is the March invoice? The answer must be the firm's records, and tier 3 is the only tier that reaches them' — two destinations ensure the firm's records are reliably populated

Explanation: 'Why twice? Because one day someone will ask, where is the March invoice? The answer must be the firm's records, and tier 3 is the only tier that reaches them.'

Q20. The course recommends ending every deliverable-producing brief with one specific instruction. What is it and what problem does it solve?

- A) 'Summarize what you did in three sentences' — solves the problem of not knowing what was accomplished
- B) 'End by listing every file you created and where each one landed: temporary working space, platform storage, or a system I control' — turns an invisible default into a visible decision about file tiers
- C) 'Ask me before completing the final step' — solves the problem of unwanted autonomous actions
- D) 'Save all intermediate files to tier 2 before completing' — prevents loss of partial work if the session ends unexpectedly

Correct Answer: B) 'End by listing every file you created and where each one landed: temporary working space, platform storage, or a system I control' — turns an invisible default into a visible decision about file tiers

Explanation: 'End every deliverable-producing brief with one instruction... End by listing every file you created and where each one landed: temporary working space, platform storage, or a system I control (connector save, download, local write, or repo commit).'

Q21. The course describes a 'reach ladder' with four rungs. What is the DEFAULT RULE for which rung to use?

- A) Always use the highest rung available to maximize capability and automation
- B) Use the most structured tool that can do the job — connector before browser, browser before full computer use — because the ladder ranks exposure, not just capability
- C) Always start at rung one and add higher rungs only when the task clearly requires them
- D) Choose based on speed — browser is fastest, computer use is slowest, connector is most reliable

Correct Answer: B) Use the most structured tool that can do the job — connector before browser, browser before full computer use — because the ladder ranks exposure, not just capability

Explanation: 'The default rule is simple: Use the most structured tool that can do the job. Connector before browser. Browser before full computer use. That is not because connectors are magically safe. The ladder ranks exposure, not safety.'

Q22. What is prompt injection, and how does it relate to the reach ladder?

- A) Prompt injection is when you add too many instructions to a prompt, making it too long to process effectively
- B) Prompt injection means hidden instructions placed in a webpage, email, document, or connector result to redirect the agent — higher rungs of the reach ladder (browser, computer use) increase exposure to this risk
- C) Prompt injection is an optimization technique where multiple prompts are combined to reduce API costs
- D) Prompt injection is when the agent adds context from previous sessions into new prompts without your knowledge

Correct Answer: B) Prompt injection means hidden instructions placed in a webpage, email, document, or connector result to redirect the agent — higher rungs of the reach ladder (browser, computer use) increase exposure to this risk

Explanation: 'Prompt injection means hidden instructions placed in a webpage, email, document, or connector result, written to redirect the agent.' The reach ladder increases exposure as you go up rungs.

Q23. The built-in Cowork browser and Claude in Chrome are DIFFERENT tools. What is the key distinction between them?

- A) The built-in browser is more powerful; Claude in Chrome is for simpler tasks only
- B) Built-in browser is for 'go handle this website task for me' (separate from your personal browser state, for delegated web work); Claude in Chrome is for 'help me with this page' (works in your Chrome with accounts you already have open)
- C) The built-in browser requires a paid plan; Claude in Chrome is available for free

D) Built-in browser works on any website; Claude in Chrome only works on Google services

Correct Answer: B) Built-in browser is for 'go handle this website task for me' (separate from your personal browser state, for delegated web work); Claude in Chrome is for 'help me with this page' (works in your Chrome with accounts you already have open)

Explanation: 'Use it (built-in browser) when the work is go handle this website task for me. Use the built-in browser when the work is go handle this website task for me... Use Claude in Chrome when the work is help me with this page.'

Q24. The course says browser login state in Cowork should be treated as a PERMISSION. Why?

- A) Because browsers save passwords and login state is a security vulnerability if exposed
- B) Because anything you sign into inside Claude's browser can reach future Cowork sessions on that computer — login state is access that persists across tasks, not just a convenience for one session
- C) Because each website charges fees based on session count, and logged-in sessions cost more
- D) Because browser login state is shared between Claude's browser and your personal browser

Correct Answer: B) Because anything you sign into inside Claude's browser can reach future Cowork sessions on that computer — login state is access that persists across tasks, not just a convenience for one session

Explanation: 'Anything you sign into inside Claude's browser can reach future Cowork sessions on that computer. Treat browser login state as a permission.' It persists and has downstream effects.

Q25. Anthropic establishes a specific rule about following instructions found in external content (emails, web pages). What is the rule?

- A) External instructions are treated equally to user instructions if they come from verified sources
- B) Following instructions found inside emails or web content is prohibited in every permission mode — your brief is authority, outside content is evidence, never a new boss
- C) External instructions can be followed if the user explicitly approves each one during the session
- D) External instructions are followed only in Auto mode, not in Manual mode, to reduce user burden

Correct Answer: B) Following instructions found inside emails or web content is prohibited in every permission mode — your brief is authority, outside content is evidence, never a new boss

Explanation: 'Anthropic goes further for Claude's browser surfaces. Following instructions found inside emails or web content is prohibited in every permission mode. Your brief is authority. Outside content is evidence, never a new boss.'

Q26. A Cowork plugin can bundle skills, connectors, AND sub-agents. What are sub-agents, and what are hooks?

- A) Sub-agents are backup models that activate when the primary model fails; hooks are notifications sent when tasks complete
- B) Sub-agents are helper workers the main agent hands a smaller job to; hooks are small rules that run automatically at set points in a task (e.g., before a tool is used)
- C) Sub-agents are separate accounts that can be added to share access; hooks are connector integrations that trigger on events
- D) Sub-agents are previous conversation sessions referenced by the current task; hooks are permission escalation requests

Correct Answer: B) Sub-agents are helper workers the main agent hands a smaller job to; hooks are small rules that run automatically at set points in a task (e.g., before a tool is used)

Explanation: 'A Claude plugin can bundle skills, connectors, and sub-agents, which means helper workers the main agent hands a smaller job to. It can also bundle hooks, which means small rules that run automatically at set points in a task.'

Q27. Why does 'read access' to a mail connector NOT mean the connector should also send emails as you?

- A) Because read access is technically incapable of sending — they are separate API endpoints
- B) Because write access (sending) is a separate increase in blast radius — grant the smallest scope that completes the workflow. A permission label is not a formality: read is not send, draft is not publish, view is not edit
- C) Because sending requires a separate 'send as user' connector that must be purchased separately
- D) Because read access expires faster than write access and must be renewed more frequently

Correct Answer: B) Because write access (sending) is a separate increase in blast radius — grant the smallest scope that completes the workflow. A permission label is not a formality: read is not send, draft is not publish, view is not edit

Explanation: 'A permission label is not a formality. Read is not send, draft is not publish, and view is not edit. If a mail connector can read threads, that does not mean it should send as you. Write access is a separate increase in blast radius.'

Q28. Cowork has three approval modes: Manual, Auto, and Skip. What is the CORRECT description of Auto mode?

- A) Auto mode automatically approves all actions without any safety checking — the fastest and cheapest mode
- B) Auto mode lets Claude keep moving while a safety classifier screens each action, blocking or escalating unsafe ones — the side panel defaults here, but it uses more usage due to extra checks
- C) Auto mode requires your approval once per session instead of per action, reducing interruptions
- D) Auto mode switches between Manual and Skip depending on the detected risk level of each action

Correct Answer: B) Auto mode lets Claude keep moving while a safety classifier screens each action, blocking or escalating unsafe ones — the side panel defaults here, but it uses more usage due to extra checks

Explanation: 'Auto: Claude keeps moving and screens each action for safety, blocking or escalating unsafe ones. The side panel starts here by default, and the extra checks consume more usage.'

Q29. The course gives 'blast radius' as a key concept for gate design. What does blast radius mean?

- A) The number of files an agent touches in a single run — more files means larger blast radius
- B) The worst damage a wrong action can do — used to match gate tightness to consequence severity
- C) The range of users affected by a scheduled task — wider blast radius means more stakeholders involved
- D) The token cost of a single run — higher cost means larger blast radius to your budget

Correct Answer: B) The worst damage a wrong action can do — used to match gate tightness to consequence severity

Explanation: 'Blast radius means the worst damage a wrong action can do.' The gate should be matched to this: 'Match the gate to blast radius.'

Q30. The course lists a 'practical autonomy ladder' with four rungs. Which rung requires Manual or another explicit human control unless you have a 'mature governed workflow'?

- A) Read and summarize — because summarization can misrepresent important information
- B) Send, publish, purchase, delete, or change critical data — these have the highest blast radius and are irreversible or highly consequential
- C) Draft but do not send — because drafts can be accidentally sent if Auto mode is active
- D) Write to reversible records — because 'reversible' is hard to guarantee in practice

Correct Answer: B) Send, publish, purchase, delete, or change critical data — these have the highest blast radius and are irreversible or highly consequential

Explanation: 'Send, publish, purchase, delete, or change critical data. High consequence. Use Manual or another explicit human control unless you have a mature governed workflow.'

Q31. The course says 'choosing the gate when the prompt appears' means you are designing the safety system too late. What is the correct timing?

- A) The gate should be designed during the task's execution, based on what actions actually appear
- B) Choose the gate BEFORE the run — if you decide autonomy when the prompt appears, the safety system is designed too late. Set permission mode and boundaries during briefing
- C) The gate should be designed after the first run, once you know what the workflow actually does
- D) The gate is automatically set by Anthropic based on the connector types and sites accessed

Correct Answer: B) Choose the gate BEFORE the run — if you decide autonomy when the prompt appears, the safety system is designed too late. Set permission mode and boundaries during briefing

Explanation: 'The important habit is not always click approve. It is to choose the gate before the run. If you decide autonomy when the prompt appears, you are designing the safety system too late.'

Q32. The course recommends 'ringing the gate on purpose' before trusting an unattended workflow. What does this mean?

- A) Set the gate to Manual for all first runs, regardless of workflow complexity
- B) Create one harmless step that should need input, start the task, walk away, confirm the request reaches a surface you actually watch, then deny or approve it and confirm the run responds correctly — test the escalation path
- C) Test the gate by trying to bypass it — if it can be bypassed, it is not secure enough
- D) Ring means enabling phone notifications specifically — verify that Slack or phone alerts arrive before trusting overnight runs

Correct Answer: B) Create one harmless step that should need input, start the task, walk away, confirm the request reaches a surface you actually watch, then deny or approve it and confirm the run responds correctly — test the escalation path

Explanation: 'Ring the gate on purpose. Create one harmless step that should need your input, start the task, walk away, and confirm the request reaches a surface you actually watch. Then deny or approve it and confirm the run responds correctly. An untested escalation path is only an assumption.'

Q33. The Anthropic certification material mentions 'four review thresholds' separate from the autonomy ladder. What are the four thresholds, and how do they differ from the ladder?

- A) Speed, accuracy, cost, and compliance — these replace the ladder for enterprise decisions
- B) Stakes, reversibility, audience, and regulatory exposure — these decide whether a person must look at the OUTPUT before it counts, regardless of the action's rung on the ladder. The rung picks the gate on the action; the thresholds pick the review on the result
- C) Data sensitivity, user count, business impact, and legal jurisdiction — these override the ladder when external parties are involved
- D) Connector type, browser access, file tier, and session length — these automatically trigger different approval modes

Correct Answer: B) Stakes, reversibility, audience, and regulatory exposure — these decide whether a person must look at the OUTPUT before it counts, regardless of the action's rung on the ladder. The rung picks the gate on the action; the thresholds pick the review on the result

Explanation: 'The four thresholds: stakes, reversibility, audience, and regulatory exposure. Those decide whether a person must look at the output before it counts, whatever rung the action sits on. The rung picks the gate on the action. The four thresholds pick the review on the result.'

Q34. What is 'workflow discovery,' and why does the course say it is often better than a longer prompt?

- A) Workflow discovery means searching the connector directory for pre-built workflows instead of writing custom prompts
- B) Letting the agent interview you to find the procedure you cannot yet describe — 'the agent helps you discover the procedure instead of waiting for you to already be an automation designer'
- C) Discovering which model tier is best for a given workflow by testing multiple options in parallel
- D) Reviewing logs from previous sessions to find patterns that can be automated

Correct Answer: B) Letting the agent interview you to find the procedure you cannot yet describe — 'the agent helps you discover the procedure instead of waiting for you to already be an automation designer'

Explanation: 'Often the better first move is workflow discovery, which means letting the agent help you find the procedure you cannot yet describe.' The example prompt: 'Interview me until you understand what I am trying to accomplish...'

Q35. The four-step delegation loop is: Brief, Plan, Approve or Redirect, Review. Which step does the course call 'the cheapest intercept' and why?

- A) Brief — because changing the brief before any work starts costs no execution resources
- B) Plan — 'the plan is usually your only intercept between intent and finished work' since you will walk away. Fixing direction before execution is cheaper than cleaning up a flawless run of the wrong plan
- C) Review — because catching problems at the end prevents the workflow from running again incorrectly
- D) Approve — because a single approval decision is faster than iterating through multiple review cycles

Correct Answer: B) Plan — 'the plan is usually your only intercept between intent and finished work' since you will walk away. Fixing direction before execution is cheaper than cleaning up a flawless run of the wrong plan

Explanation: 'The plan is usually your only intercept between intent and finished work. Fix the plan, not the cleanup.' On a web surface you walk away, so the plan review is your key intercept before expensive execution.

Q36. The brief template for unattended work has eight components. Which component does the course say 'beginners routinely leave out,' and what does it prevent?

- A) Audience — because without audience specification the agent writes for the wrong reader
- B) 'If the necessary evidence is missing, say what is missing and stop rather than inventing it' — prevents confident inventions from traveling through several unattended steps before you see them
- C) Escalation — because without escalation rules the agent never pauses even when genuinely stuck
- D) Done definition — because without a clear finish condition the agent keeps iterating unnecessarily

Correct Answer: B) 'If the necessary evidence is missing, say what is missing and stop rather than inventing it' — prevents confident inventions from traveling through several unattended steps before you see them

Explanation: 'Then add one line beginners routinely leave out: If the necessary evidence is missing, say what is missing and stop rather than inventing it. That line matters more in unattended work, because a confident invention can otherwise travel through several steps before you see it.'

Q37. When reviewing the plan (not every click), the course gives four specific checks. What are they?

- A) Length, tone, format, and deadline — the four output quality checks
- B) Scope (touching only named work and data), Order (verify before act), Reach (unanticipated connectors/sends/writes), Assumptions (deciding audience/format/rules for you)
- C) Speed, cost, accuracy, and safety — the four performance dimensions
- D) Source quality, citation count, output format, and tier-3 destination — the four deliverable checks

Correct Answer: B) Scope (touching only named work and data), Order (verify before act), Reach (unanticipated connectors/sends/writes), Assumptions (deciding audience/format/rules for you)

Explanation: Four plan checks: 'Scope: is it touching only the work and data you named? Order: does it verify before it acts? Reach: is it proposing a connector, browser account, send, publish, write, or delete you did not intend? Assumptions: is it deciding audience, format, business rule, or a missing fact for you?'

Q38. The course warns about 'Plan-first is an instruction, not a guarantee of Manual mode.' What does this mean for Cowork users specifically?

- A) Planning mode automatically switches the gate to Manual for all consequential actions
- B) In the current Cowork side panel, Manual mode does not automatically create a plan for approval before starting — if you want a plan intercept, put it in the brief: 'Show me your plan first and wait for my approval before taking consequential actions'
- C) Requesting a plan first is only possible in Manual mode — Auto mode skips the planning step
- D) Plan-first is a ChatGPT Work feature only — Cowork users must use step-by-step prompting instead

Correct Answer: B) In the current Cowork side panel, Manual mode does not automatically create a plan for approval before starting — if you want a plan intercept, put it in the brief: 'Show me your plan first and wait for my approval before taking consequential actions'

Explanation: 'In the current Cowork side panel, Manual mode does not automatically create a plan for you to approve before starting. If you want the Concept 8 safety intercept, put it in the brief: Show me your plan first and wait for my approval before taking consequential actions.'

Q39. The course lists four heartbeat patterns for scheduled tasks. What are they?

- A) Daily, weekly, monthly, yearly — the four standard recurrence intervals
- B) Once (single run), on a schedule (clock-based), on a supported event (trigger-based), and monitoring (watch for a condition and report when true)
- C) Manual start, automatic start, conditional start, and scheduled start — the four task initiation types
- D) Web trigger, API trigger, connector trigger, and desktop trigger — the four integration entry points

Correct Answer: B) Once (single run), on a schedule (clock-based), on a supported event (trigger-based), and monitoring (watch for a condition and report when true)

Explanation: 'The four useful heartbeat patterns are once, on a schedule, on a supported event, and monitoring, which means checking repeatedly and reporting only when a condition becomes true.'

Q40. The course states: 'A schedule is truly device-independent only if every required tool is also reachable without your device.' What does this mean for a Cowork schedule that needs a file from a sleeping laptop?

- A) The schedule will wait for the laptop to wake up before running, then proceed normally
- B) You have rebuilt the laptop as a hidden dependency — the Monday brief that needs a local file is not a truly cloud-independent schedule, it secretly depends on the desktop bridge being available
- C) The schedule will run but skip the laptop-dependent steps and complete the rest of the workflow
- D) The schedule will run on a cached copy of the file from the last time the laptop was online

Correct Answer: B) You have rebuilt the laptop as a hidden dependency — the Monday brief that needs a local file is not a truly cloud-independent schedule, it secretly depends on the desktop bridge being available

Explanation: 'If your Monday brief needs a file on your sleeping laptop, you have rebuilt the laptop as a hidden dependency.' A cloud clock is not enough if tools require the device.

Q41. The course says 'ran is not the same as worked.' What must every recurring workflow have to prove the run actually succeeded?

- A) A completion timestamp and a total execution duration record
- B) An observable success signal — items processed, a dated file created, the tier-3 destination used, or an explicit 'no new items' result. Absence of output means something is unclear, not that the run succeeded

- C) A green status icon and no error messages in the run log
- D) A notification sent to at least one human reviewer confirming the output quality

Correct Answer: B) An observable success signal — items processed, a dated file created, the tier-3 destination used, or an explicit 'no new items' result. Absence of output means something is unclear, not that the run succeeded

Explanation: 'Give every recurring workflow an observable success signal: items processed, a dated file created, the tier-3 destination used, or an explicit no new items result.' A completed run status proves only that the mechanism ran.

Q42. The course says 'the empty case' must be handled. What does this mean and why does it matter for schedules?

- A) An empty case means a run with no files attached — the brief must specify what to do when no documents are provided
- B) 'No new items' must be a valid, explicit result — not an invitation to invent work. A schedule facing nothing to do must report 'nothing to do' rather than fabricating content to fill the run
- C) An empty case means a run that costs zero tokens — schedules should optimize for empty cases to reduce costs
- D) Empty case refers to a schedule that fires but finds the connector is temporarily unavailable

Correct Answer: B) 'No new items' must be a valid, explicit result — not an invitation to invent work. A schedule facing nothing to do must report 'nothing to do' rather than fabricating content to fill the run

Explanation: 'What happens on the empty or unclear case? No new items must be a valid result, not an invitation to invent work.' This prevents fabricated output when the schedule runs and finds nothing to process.

Q43. The course says for beginners, the safest scheduled agent 'reads, analyzes, and reports.' Why is this the ceiling for beginners, and what comes after?

- A) Beginners should only read because write connectors cost more and exceed free plan limits
- B) Reading and reporting changes nothing while you sleep — it is safe for beginners. Once a workflow has clean runs, a clear checker, bounded permissions, and reversible actions, later courses take you toward schedules that act. Capability and readiness are different things
- C) Reading is the ceiling because Cowork does not support scheduled write actions on any plan
- D) Reporting-only is required for the first 30 days of a new account before write permissions are unlocked

Correct Answer: B) Reading and reporting changes nothing while you sleep — it is safe for beginners. Once a workflow has clean runs, a clear checker, bounded permissions, and reversible actions, later courses take you toward schedules that act. Capability and readiness are different things

Explanation: 'For beginners, the safest scheduled agent reads, analyzes, and reports, so it changes nothing while you sleep. Once a workflow has clean runs, a clear checker, bounded permissions, and reversible actions, later courses take you toward schedules that act.'

Q44. The course says 'scheduling is photocopying the workflow into the future.' Why does this metaphor matter?

- A) Like photocopying, scheduling reduces the quality of each copy — schedules should only be used for non-critical tasks
- B) A mistake inside a scheduled workflow is not diluted — it is reproduced exactly each time. 'Walk it before you schedule it' — run it by hand twice, inspect everything, THEN attach the heartbeat
- C) Like photocopying, scheduling creates multiple parallel runs simultaneously that can conflict
- D) Scheduling preserves the workflow in a read-only state — changes require creating a new schedule rather than editing the existing one

Correct Answer: B) A mistake inside a scheduled workflow is not diluted — it is reproduced exactly each time. 'Walk it before you schedule it' — run it by hand twice, inspect everything, THEN attach the heartbeat

Explanation: 'Scheduling is photocopying the workflow into the future. A mistake inside it is not diluted. It is reproduced.' This is why you must walk the workflow manually before scheduling it.

Q45. The course's routing table has a row for 'regulated or contractually restricted data.' What does it say, and why can this row NOT be waved away?

- A) Regulated data can be used in Cowork if you enable enterprise mode and accept the terms
- B) Stop and verify — availability is not permission. Get the written compliance answer for the exact surface before uploading. For HIPAA specifically, Cowork is not BAA-covered as of the fact check
- C) Regulated data requires Manual approval mode to be active, which provides sufficient protection
- D) Regulated data should be anonymized and then uploaded — anonymization satisfies all regulatory requirements

Correct Answer: B) Stop and verify — availability is not permission. Get the written compliance answer for the exact surface before uploading. For HIPAA specifically, Cowork is not BAA-covered as of the fact check

Explanation: 'Regulated or contractually restricted data: Stop and verify. Availability is not permission. Get the written compliance answer for the exact surface.' And specifically: 'Cowork is not covered under its BAA for HIPAA use.'

Q46. The course mentions 'removing identifiers' as a route for regulated data. But it warns stripping can fail in two ways. What are they?

- A) Over-stripping (task cannot be done) and under-stripping (a combination of details still identifies a person — 'the only student who missed week 3 lab' identifies as surely as a name does)
- B) Stripping too early (data becomes useless before analysis) and stripping too late (data has already been transmitted to the AI)
- C) Over-stripping (reduces accuracy) and improper format (stripped data must be in CSV format for AI to process it correctly)
- D) Stripping personally identifies you as the data subject and under-stripping makes the AI fabricate replacement details

Correct Answer: A) Over-stripping (task cannot be done) and under-stripping (a combination of details still identifies a person — 'the only student who missed week 3 lab' identifies as surely as a name does)

Explanation: From the AI Fluency course referenced here: 'It fails in two directions. Take out too much and the task cannot be done... Take out too little and a combination of details still points at one person.' The course here references this principle: 'Removing what the task does not need is a route, not an evasion.'

Q47. The course describes a 'bridge' between cloud Cowork and desktop. What is the critical custody fact about a local file opened through this bridge?

- A) The file remains in local custody — the bridge only gives the cloud agent read access without copying the data
- B) The file's contents are processed in the vendor's cloud, even though it came from the laptop — 'the file came from my laptop does not mean the file stayed on my laptop'
- C) The file is encrypted before leaving the desktop and decrypted only within the agent's secure context
- D) The bridge creates a temporary mirror of the file that is deleted from the cloud immediately after the task completes

Correct Answer: B) The file's contents are processed in the vendor's cloud, even though it came from the laptop — 'the file came from my laptop does not mean the file stayed on my laptop'

Explanation: 'If a cloud Cowork session reads a local file through Claude Desktop, its contents are processed in Anthropic's cloud. So do not route sensitive work to local file and assume local custody. The first question is where the processing happens.'

Q48. The open path section describes two paths for running agent work without a vendor cloud. What are they?

- A) Fine-tuning your own model and running it locally; building a local API wrapper around open models
- B) OpenWork with remote/shared workspaces (open-source desktop coworker connecting to infrastructure you control); OpenCode with a scheduler you own (coding agent fired by your own clock, e.g., cron or GitHub Actions)
- C) Running ChatGPT in airplane mode for local execution; using Claude's local mode with API keys
- D) Deploying an open-weight model on a home server; connecting it to your existing productivity tools

Correct Answer: B) OpenWork with remote/shared workspaces (open-source desktop coworker connecting to infrastructure you control); OpenCode with a scheduler you own (coding agent fired by your own clock, e.g., cron or GitHub Actions)

Explanation: 'Two open paths exist today. OpenWork, remote and shared workspaces... that is this course's core idea, on infrastructure you or your organization control. OpenCode, with a scheduler you own. For repository work, the open answer is a coding agent fired by your own clock.'

Q49. The course says the vendor 'sells you a spine' and the open path 'makes you build one.' What does 'spine' refer to here?

- A) The AI model weights and training data that power intelligent responses
- B) All the managed infrastructure: sessions, files, sync, the phone gate, and the scheduler — the things that make the agent surface work without setup, but that also mean your data lives in the vendor's custody and format
- C) The backbone API that connects different parts of the product together
- D) The system of rules and policies governing how the agent behaves

Correct Answer: B) All the managed infrastructure: sessions, files, sync, the phone gate, and the scheduler — the things that make the agent surface work without setup, but that also mean your data lives in the vendor's custody and format

Explanation: 'On the vendor surface, everything is handed to you and working on day one: sessions, files, sync, the phone gate, the scheduler. But it lives in their custody, in their format, at their price. On the open path, each of those is yours to set up, run, and fix.'

Q50. What is the key lesson from Concept 12 about the limits of the managed surface?

- A) Managed surfaces are inherently less capable than local alternatives for complex tasks
- B) More capability does not make weak work good — memory can preserve a bad assumption, a browser can execute a bad plan faster, Auto approval can reduce interruptions without making judgment correct, scheduling can repeat an error every Monday. The brief, plan review, permission boundary, checker, and exit are still the controls
- C) Managed surfaces will eventually become fully autonomous, reducing the need for human oversight
- D) The managed surface is limited by daily run caps that prevent its use for serious production workloads

Correct Answer: B) More capability does not make weak work good — memory can preserve a bad assumption, a browser can execute a bad plan faster, Auto approval can reduce interruptions without making judgment correct, scheduling can repeat an error every Monday. The brief, plan review, permission boundary, checker, and exit are still the controls

Explanation: 'The web surface still does not make weak work good. Memory can preserve a bad assumption. A browser can execute a bad plan faster. Auto approval can reduce interruptions without making the underlying judgment correct. Scheduling can repeat an error every Monday.'

Q51. A user sets up a cloud Cowork schedule for Monday mornings. It uses Google Drive connector (cloud), a Slack connector (cloud), and a browser task that requires Claude Desktop to be running. Is this a truly device-independent schedule?

- A) Yes — cloud connectors make the whole workflow cloud-independent regardless of other components
- B) No — the browser task requires Claude Desktop to be online. 'A schedule is truly device-independent only when every required tool is also reachable without your device.' The desktop requirement creates a device dependency

- C) Yes — cloud schedules automatically fall back to cloud-only tools when the desktop is offline
- D) It depends on the time zone — schedules run during business hours are device-dependent, others are not

Correct Answer: B) No — the browser task requires Claude Desktop to be online. 'A schedule is truly device-independent only when every required tool is also reachable without your device.' The desktop requirement creates a device dependency

Explanation: 'A schedule is truly device-independent only if every required tool is also reachable without your device.' The browser task needing Claude Desktop makes this schedule device-dependent despite the cloud connectors.

Q52. Which of the following correctly applies the six-part lens to a new agent product announcement?

- A) Read the marketing page and count the number of features — more features means the product covers more parts
- B) Find the heartbeat, reach, run-until-done loop, state spine, human gate, and body. 'If one is missing, ask how the product makes up for it. If a part is hidden, ask who controls it.' Do not start with the marketing name
- C) Compare pricing tiers across competitors to determine which product has the best value per part
- D) Test the product on five sample tasks and rate each part based on observed performance

Correct Answer: B) Find the heartbeat, reach, run-until-done loop, state spine, human gate, and body. 'If one is missing, ask how the product makes up for it. If a part is hidden, ask who controls it.' Do not start with the marketing name

Explanation: 'Run the lens over any new agent product. Do not start with the marketing name. Find the six parts. If one is missing, ask how the product makes up for it. If a part is hidden, ask who controls it.'

Q53. Project 6 (the portability drill) asks you to rebuild a workflow assuming your preferred vendor disappeared tomorrow. What does the course say the result reveals?

- A) How long vendor migration typically takes — used to plan contingency timelines
- B) What you could NOT rebuild — that list is your lock-in exposure. 'What travels when you move: your best assets are not vendor buttons. They are the outcome definition, portable instructions, source-of-truth locations, schema, checks, failure rules, permission boundaries, and examples of good output.'
- C) Which features are most used — used to justify the vendor subscription cost
- D) How complex your workflows have become — used to decide whether to hire additional staff

Correct Answer: B) What you could NOT rebuild — that list is your lock-in exposure. 'What travels when you move: your best assets are not vendor buttons. They are the outcome definition, portable instructions, source-of-truth locations, schema, checks, failure rules, permission boundaries, and examples of good output.'

Explanation: Project 6: 'Done when you have a written list of what you could not rebuild. That list is your lock-in exposure.' The goal is making invisible vendor dependency visible.

Q54. The course says Cowork looks less like 'two products called web and desktop' and more like 'one managed system with several execution surfaces.' What does this mean for how you should think about routing decisions?

- A) Routing should now focus entirely on which plan tier gives access to which surface
- B) Use the two-axis model — where the loop runs and where each tool executes — rather than assuming 'web = no local access' or 'desktop = local only.' A cloud session can borrow local tools through a bridge, so the surfaces overlap
- C) Routing decisions are no longer necessary — the managed system automatically selects the right surface
- D) Both web and desktop surfaces are now equivalent, so you should always use whichever is more convenient

Correct Answer: B) Use the two-axis model — where the loop runs and where each tool executes — rather than assuming 'web = no local access' or 'desktop = local only.' A cloud session can borrow local tools through a bridge, so the surfaces overlap

Explanation: 'Cowork looks less like two products called web and desktop and more like one managed system with several execution surfaces. That is why Concept 2's two-axis model matters.'

Q55. The harness war appendix describes three fronts of competition. What is the durable signal the course says to track rather than model rankings?

- A) Number of parameters — larger models always win the harness war
- B) Vendors compete on the whole agent system: memory, tools, browser reach, schedules, plugins, approvals, runtime, identity, and governance — that is the harness war. The durable signal is not a model ranking
- C) Integration count — the vendor with the most connectors wins the harness war
- D) Speed benchmarks — faster inference wins because users value real-time responses most

Correct Answer: B) Vendors compete on the whole agent system: memory, tools, browser reach, schedules, plugins, approvals, runtime, identity, and governance — that is the harness war. The durable signal is not a model ranking

Explanation: 'The durable market signal is not a model ranking. It is that major vendors compete on the whole agent system: memory, tools, browser reach, schedules, plugins, approvals, runtime, identity, and governance. That is the harness war.'

Q56. What are the eight questions from Concept 3 and the August update that the course says let you read ANY new agent product before the marketing page finishes loading?

- A) Price, features, integrations, models, plans, languages, security, support — the eight purchase evaluation questions
- B) The six-part lens questions (heartbeat, reach, loop, spine, gate, body) PLUS the two August questions (where does the loop run? where does each tool execute?) — eight total
- C) Who built it, when was it released, what model does it use, what data does it train on, what connectors, what plans, what security, what compliance — eight due diligence questions
- D) Name, version, model, context window, connectors, pricing, availability, roadmap — eight market comparison questions

Correct Answer: B) The six-part lens questions (heartbeat, reach, loop, spine, gate, body) PLUS the two August questions (where does the loop run? where does each tool execute?) — eight total

Explanation: 'Answer those eight questions and you can read a new agent product before the marketing page finishes loading.' The six parts of the shape plus the two runtime questions from Concept 2.

Q57. The course describes a 'security ceiling' for browser agents. What creates this ceiling and what reduces (but cannot eliminate) the risk?

- A) Browser memory limits create a ceiling; clearing browser cache regularly reduces risk
- B) What makes browser agents useful (reading arbitrary pages while signed in and taking actions) also creates a path for hidden instructions. Safety classifiers (automatic checkers on each proposed action) reduce the risk but cannot make untrusted content trusted
- C) Browser agents have a processing ceiling due to DOM parsing complexity; using simplified pages reduces risk
- D) The number of open browser tabs creates security risks; limiting to one tab per session reduces risk

Correct Answer: B) What makes browser agents useful (reading arbitrary pages while signed in and taking actions) also creates a path for hidden instructions. Safety classifiers (automatic checkers on each proposed action) reduce the risk but cannot make untrusted content trusted

Explanation: 'What makes them useful — reading arbitrary pages while signed in and taking actions — also creates a path for hidden instructions. Screening reduces risk. It cannot make untrusted content trusted.'

Q58. The course says 'persistent context creates both power and lock-in.' What is the practical recommendation to manage this tension?

- A) Avoid using persistent context entirely — always start fresh sessions to prevent lock-in
- B) Use the vendor's spine for active work, and put finished work and portable instructions in places you control — 'separate working continuity from system-of-record custody'
- C) Export all context to a neutral format weekly to prevent lock-in from accumulating
- D) Use only open-source tools for persistent context so the format is always yours

Correct Answer: B) Use the vendor's spine for active work, and put finished work and portable instructions in places you control — 'separate working continuity from system-of-record custody'

Explanation: 'The safest response is not to reject persistence. It is to separate working continuity from system-of-record custody. Use the vendor's spine for active work. Put finished work and portable instructions in places you control.'

Q59. In the Leaving the Laptop appendix, the course says to classify a workflow by four facts rather than 'the screen you started it from.' What are the four facts?

- A) Start time, end time, cost, and success rate — the four operational facts
- B) Where the agent loop runs, which tools are required and where they execute, who must keep a device or service alive, and who holds and processes the data
- C) Which vendor provides it, which plan covers it, which team owns it, and which compliance standard applies
- D) When it was created, when it was last run, when it was last modified, and when it is next scheduled

Correct Answer: B) Where the agent loop runs, which tools are required and where they execute, who must keep a device or service alive, and who holds and processes the data

Explanation: 'Classify it by four facts: Where does the agent loop run? Which tools are required, and where do they execute? Who must keep a device or service alive? Who holds and processes the data?'

Q60. The course says 'what travels when you move' to a different platform. What travels well, and what does NOT automatically travel?

- A) Everything travels if you use standard connectors — lock-in is a myth in the modern MCP ecosystem
- B) Travels well: outcome definition, portable instructions, source-of-truth locations, schema, checks, failure rules, permission boundaries, examples of good output. Does NOT travel: native memory, session history, Projects, plugin implementations, browser login state, approval labels, vendor-specific schedule mechanics
- C) Models travel well (via API); only proprietary fine-tunes don't travel. Everything else is platform-neutral
- D) Files travel if exported to PDF; structured data travels in CSV format. Everything else requires manual recreation

Correct Answer: B) Travels well: outcome definition, portable instructions, source-of-truth locations, schema, checks, failure rules, permission boundaries, examples of good output. Does NOT travel: native memory, session history, Projects, plugin implementations, browser login state, approval labels, vendor-specific schedule mechanics

Explanation: 'Your best assets are not vendor buttons. They are the outcome definition, portable instructions, source-of-truth locations, schema, checks, failure rules, permission boundaries, and examples of good output.' vs. 'Native memory, session history, Projects, plugin implementations, browser login state... Treat these as conveniences around the workflow, not the workflow's only copy.'

Q61. What does the course mean by 'owned enforcement versus rented enforcement' in the Harness Engineering appendix?

- A) Owned means you built the AI model; rented means you use a commercial API
- B) Rented enforcement means Anthropic owns the managed runtime and safety layer — you shape it through instructions and modes but cannot fully inspect or replace the mechanisms. Owned means you

write, deploy, inspect, and can reproduce the guardrails yourself

C) Owned enforcement means you paid for it with a subscription; rented means you are on a free plan

D) Owned enforcement uses your organization's identity provider; rented enforcement uses the vendor's default authentication

Correct Answer: B) Rented enforcement means Anthropic owns the managed runtime and safety layer — you shape it through instructions and modes but cannot fully inspect or replace the mechanisms. Owned means you write, deploy, inspect, and can reproduce the guardrails yourself

Explanation: 'The deeper distinction is who owns the enforcement layer.' Managed Cowork gives you 'shaping power' but 'What you cannot do is treat a managed hook as identical to a guardrail enforced in a runtime you control.'

Q62. The Trusting the Checker appendix says 'demos are the weakest evidence.' Why?

A) Demos use special demo data that is cleaner than real production data

B) A demo is one run, on a task chosen because it works, watched by someone hoping it does — 'grade the rate, not the run.' A scheduled report that ran well once is one run, not proof

C) Demos are always scripted and therefore cannot represent actual AI behavior

D) Demos run in a separate environment with different model parameters than production

Correct Answer: B) A demo is one run, on a task chosen because it works, watched by someone hoping it does — 'grade the rate, not the run.' A scheduled report that ran well once is one run, not proof

Explanation: 'Demos are the weakest evidence. It worked when I tried it is one run, on a task chosen because it works, watched by someone hoping it does.' The course applies this: 'The most useful idea to carry: An agent is a distribution, not a function. Grade the rate, not the run.'

Q63. A Cowork plugin adds a hook that runs before a specific tool is used. How does this compare to a hook in a coding harness you own?

A) They are functionally identical — hooks in both environments enforce the same guarantees

B) A Cowork plugin hook runs at supported lifecycle points in Anthropic's managed environment. A coding harness hook runs in your own runtime, tied to local tools, repository state, tests, and CI — you own and can inspect the enforcement mechanism

C) Coding harness hooks are more limited because they cannot access cloud connectors

D) Plugin hooks are more reliable because they are maintained by Anthropic rather than requiring individual maintenance

Correct Answer: B) A Cowork plugin hook runs at supported lifecycle points in Anthropic's managed environment. A coding harness hook runs in your own runtime, tied to local tools, repository state, tests, and CI — you own and can inspect the enforcement mechanism

Explanation: 'Plugins can attach hooks that run in Cowork at supported points in the managed workflow. Your hook code runs in the harness you operate, tied to local tools, repository state, tests, and CI.' The difference is who owns the enforcement.

Q64. The course says Cowork's August changes 'strengthen the overlap' with Loop Engineering but Loop Engineering remains a separate course. What is the key distinction?

A) Loop Engineering teaches Python coding; Cowork teaches non-coding automation

B) Managed Cowork gives you a loop. Loop Engineering teaches you to own the loop's contract — the explicit stopping and retry logic, a spine you own and can version, independent checker design, isolation, re-runnable orchestration, and portable unattended runtime

C) Loop Engineering requires more expensive hardware; Cowork runs on standard consumer devices

D) Loop Engineering is for teams; Cowork is for individual users

Correct Answer: B) Managed Cowork gives you a loop. Loop Engineering teaches you to own the loop's contract — the explicit stopping and retry logic, a spine you own and can version, independent checker design, isolation, re-runnable orchestration, and portable unattended runtime

Explanation: 'A managed surface gives you a loop. Loop Engineering teaches you to own the loop's contract.' Key differences: you write stopping/retry logic, own the spine (not vendor-managed), design the checker with your own rubric, and the loop definition can move across runtimes you operate.

Q65. The course gives a 'freshness rule' at the end of the sources section. What does it say, and what is always considered 'lasting'?

- A) Check all sources annually; only the model names are lasting because capability is constant
- B) Everything mechanical was checked on 31 August 2026 — recheck before relying on it. The lasting concepts are: six-part harness, two-axis runtime model, persistence stack, three file tiers, reach ladder, human gate, delegation loop, and routing questions
- C) The course is updated quarterly; subscribe to update notifications for the latest mechanical details
- D) All concepts should be reverified monthly; nothing in AI is truly lasting given the pace of change

Correct Answer: B) Everything mechanical was checked on 31 August 2026 — recheck before relying on it. The lasting concepts are: six-part harness, two-axis runtime model, persistence stack, three file tiers, reach ladder, human gate, delegation loop, and routing questions

Explanation: 'Freshness rule. Everything mechanical on this page was checked on 31 August 2026... Recheck the rest before relying on it.' The lasting concepts are explicitly listed as the architectural ideas, not the product specifics.

Q66. What is the difference between the 'built-in Cowork browser' and 'Claude in Chrome' in terms of who OWNS the browser state?

- A) Both use your Chrome profile — the distinction is only in where the side panel appears
- B) Built-in Cowork browser: Claude gets its own separate browser state (not your personal browser). Claude in Chrome: uses YOUR Chrome browser with your logins and accounts already open — you own and are responsible for that existing session state
- C) Cowork browser is owned by Anthropic; Claude in Chrome is owned by Google
- D) Both are owned by Anthropic — the difference is only the visual interface

Correct Answer: B) Built-in Cowork browser: Claude gets its own separate browser state (not your personal browser). Claude in Chrome: uses YOUR Chrome browser with your logins and accounts already open — you own and are responsible for that existing session state

Explanation: 'Built-in Cowork browser: Claude gets its own browser for delegated web work' (separate state). 'Claude in Chrome: Works in your Chrome browser, on the page and accounts you already have open.' The five-path comparison table distinguishes these clearly.

Q67. A connector scope can be 'too broad.' What additional risk beyond scope width does the course name for connectors?

- A) Connectors with broad scope cost more API credits per action
- B) Outside content arriving through a connector can carry hidden instructions (prompt injection) — 'the ladder ranks exposure, not safety. Rung one is still a real door'
- C) Broad connector scopes require admin approval that delays workflow setup
- D) Connectors with too-broad scope slow down the agent because it must search through more data

Correct Answer: B) Outside content arriving through a connector can carry hidden instructions (prompt injection) — 'the ladder ranks exposure, not safety. Rung one is still a real door'

Explanation: 'A connector scope can be too wide, and content arriving through it can carry hidden instructions.' The reach ladder puts connectors first, but that is about structure and exposure, not absolute safety.

Q68. What does 'HIPAA BAA' mean, and why does the course say the compliance row 'cannot be waved away' even when Cowork has enterprise controls?

- A) HIPAA BAA means High Impact Protocol for Automated Agents — a performance certification required for enterprise use

B) HIPAA is the US health privacy law; a BAA (business associate agreement) is the contract a vendor must sign before handling protected health information (PHI). Enterprise controls and Compliance API coverage do not by themselves make Cowork BAA-covered for HIPAA — that requires a specific vendor contract

C) HIPAA BAA means the interface has been audited; enterprises receive automatic BAA coverage when purchasing Cowork Team plans

D) BAA stands for Business Automation Approval — an internal policy document required before deploying any AI tool in a healthcare organization

Correct Answer: B) HIPAA is the US health privacy law; a BAA (business associate agreement) is the contract a vendor must sign before handling protected health information (PHI). Enterprise controls and Compliance API coverage do not by themselves make Cowork BAA-covered for HIPAA — that requires a specific vendor contract

Explanation: 'HIPAA is the US health privacy law. PHI means protected health information. A BAA is the contract a vendor signs before it may handle PHI for you. Cowork is not covered under its BAA for HIPAA use' (as of the fact check).

Q69. The course's dogfooding section reveals that the book's own loops do NOT run on Cowork. Why, and what principle guided this routing decision?

A) The book uses Cowork for writing but not for technical work — a best practice for content workflows

B) The book is a repository and its loops are attached to it — 'what the work touches routed it to a different surface.' Concept 10's routing question (repository work → coding agent or terminal) directed the choice away from the web surface

C) Cowork was not available when the book's workflows were first built, so they were never migrated

D) Using Cowork for the book would create a conflict of interest since Anthropic created both

Correct Answer: B) The book is a repository and its loops are attached to it — 'what the work touches routed it to a different surface.' Concept 10's routing question (repository work → coding agent or terminal) directed the choice away from the web surface

Explanation: 'The book is a repository, and its loops are attached to it. Its review gate lives beside its merge gate in CI. What the work touches routed it to a different surface.'

Q70. What TWO ways does the web surface 'earn its keep' for the book despite the book's loops running elsewhere?

A) Writing the book content and managing the book's financial accounts

B) Chat-side design work before every loop (drafting briefs, testing reviewer prompts, running one beat by hand) AND connector and document tasks around the book (the reporting work Concept 9 called the course's ceiling)

C) Hosting the book's website and providing customer support for readers

D) Training the AI models used in the book's examples and managing the book's code repository

Correct Answer: B) Chat-side design work before every loop (drafting briefs, testing reviewer prompts, running one beat by hand) AND connector and document tasks around the book (the reporting work Concept 9 called the course's ceiling)

Explanation: 'This surface still earns its keep in two ways. First, the chat-side design work before every loop: drafting briefs, testing reviewer prompts, running one beat by hand. Second, connector and document tasks around the book, which is the reporting work Concept 9 called this course's ceiling.'

Q71. The Trusting the Checker appendix describes what does NOT transfer to the managed surface in the course's 'engineered form.' What are the five things listed?

A) Speed, accuracy, cost, compliance, and user satisfaction — the five evaluation dimensions

B) The golden set (version-controlled case files in a repository), the runner (headless shell loop with grading), calibrating the judge (blind protocol and four-cell table), regression suite with CI gate, and drift watching with sealed hold-outs — all require infrastructure you own

C) The model, the training data, the inference engine, the safety classifier, and the compliance documentation

D) The prompt, the context, the tools, the memory, and the output format — the five elements of a complete agent specification

Correct Answer: B) The golden set (version-controlled case files in a repository), the runner (headless shell loop with grading), calibrating the judge (blind protocol and four-cell table), regression suite with CI gate, and drift watching with sealed hold-outs — all require infrastructure you own

Explanation: The appendix lists: 'The golden set... The runner, claude -p or opencode run in a shell loop with jq... Calibrating the judge... The regression suite and CI gate... Drift watching on a schedule... Goodhart's law defenses, sealed hold-outs.' These require a system you own.

Q72. What does the course mean by 'switching cost' in relation to persistent context, and what is 'the safest response'?

A) Switching cost is the monthly subscription price — the safest response is to choose the cheapest plan

B) Switching cost is the price of moving to another product later — the more persistence you accumulate, the harder it is to leave. The safest response is not rejecting persistence but separating working continuity (use vendor spine) from system-of-record custody (put finished work in places you control)

C) Switching cost is the time it takes to learn a new interface — the safest response is to standardize on one vendor

D) Switching cost is token overhead from importing historical context — the safest response is to start fresh sessions regularly

Correct Answer: B) Switching cost is the price of moving to another product later — the more persistence you accumulate, the harder it is to leave. The safest response is not rejecting persistence but separating working continuity (use vendor spine) from system-of-record custody (put finished work in places you control)

Explanation: 'There is also a business consequence. The same persistence that makes these products convenient creates switching cost, which means the price of moving to another product later. The safest response is not to reject persistence. It is to separate working continuity from system-of-record custody.'

Q73. The course says a Cowork task can run its agent loop in Anthropic's cloud 'while a browser action runs through a connected desktop.' What does this specific example illustrate?

A) That the agent loop and tool execution can run in different locations simultaneously — the loop is remote, the tool is local, through a bridge

B) That Cowork requires both cloud and desktop simultaneously for all browser tasks

C) That browser actions are always local while the agent loop is always remote

D) That the agent pauses its cloud execution and temporarily moves to the desktop to run the browser action

Correct Answer: A) That the agent loop and tool execution can run in different locations simultaneously — the loop is remote, the tool is local, through a bridge

Explanation: This is the two-axis model: 'A Cowork task can run its agent loop in Anthropic's cloud while a browser action runs through a connected desktop.' Loop location and tool location are separate axes.

Q74. The course says 'same label does not mean same implementation.' What example does it give for the label 'memory'?

A) Memory can mean either short-term or long-term retention depending on the vendor's implementation

B) 'Memory,' 'plugin,' 'approval,' and even 'browser' can mean different mechanisms across vendors. Claude's semantic memory, project memory, and session history are all 'memory' but with different scopes, lifetimes, and controls — use the name to find the part, then inspect the implementation

C) Memory in Claude Code means CLAUDE.md while memory in Cowork means the cloud knowledge base

D) Memory can mean either the user's stated preferences or the AI's learned behaviors, and vendors implement both differently

Correct Answer: B) 'Memory,' 'plugin,' 'approval,' and even 'browser' can mean different mechanisms across vendors. Claude's semantic memory, project memory, and session history are all 'memory' but with different scopes, lifetimes, and controls — use the name to find the part, then inspect the implementation

Explanation: 'Same label does not mean same implementation. Memory, plugin, approval, and even browser can mean different mechanisms across vendors. Use the name to find the part. Then inspect the implementation.'

Q75. The course says 'the chat box also kept its job.' What are chat's remaining jobs even after the agent surface exists?

A) Handling simple questions that don't need files or tools

B) Designing tasks, drafting briefs, testing ideas by hand — 'every assignment here begins there.' Chat is where you prepare what you hand over to the agent surface

C) Providing a fallback when the agent surface is unavailable or over its run limit

D) Handling conversations that involve sensitive data that cannot go to the agent surface

Correct Answer: B) Designing tasks, drafting briefs, testing ideas by hand — 'every assignment here begins there.' Chat is where you prepare what you hand over to the agent surface

Explanation: 'The chat box also kept its job. Designing a task, drafting a brief, testing an idea by hand: that is chat work, and every assignment here begins there.' Chat and the agent surface are partners, not replacements.

Q76. What is a 'user-owned context file' in the persistent context stack, and why does the course recommend it even when native memory exists?

A) A file that stores your API key and authentication credentials, kept local for security

B) Your own copy of important rules, portable and readable outside any vendor — 'keep a portable copy of anything critical.' Even when native memory and instructions exist, a file like CLAUDE.md provides vendor-neutral portability. 'Use the file for portability'

C) A log file that records all your AI interactions for personal auditing and analysis

D) A JSON configuration file that sets model parameters and temperature settings for your preferred output style

Correct Answer: B) Your own copy of important rules, portable and readable outside any vendor — 'keep a portable copy of anything critical.' Even when native memory and instructions exist, a file like CLAUDE.md provides vendor-neutral portability. 'Use the file for portability'

Explanation: 'A user-owned context file: What do I want portable, readable, and vendor-neutral? (as long as you maintain the file).' 'CLAUDE.md makes an excellent portable context file... Use the file for portability.' The file is the portability layer.

Q77. The course says 'three habits keep the spine clean.' What are the three habits?

A) Export weekly, review monthly, and delete annually — the three data hygiene habits

B) Name sessions like work products; keep one workstream per session or Project; put durable instructions in the right layer instead of repeating them in every prompt

C) Use Manual mode always, export to tier 3 always, and test with known data always

D) Check memory monthly, review instructions quarterly, and audit Projects annually

Correct Answer: B) Name sessions like work products; keep one workstream per session or Project; put durable instructions in the right layer instead of repeating them in every prompt

Explanation: 'Three habits keep the spine clean. Name sessions like work products. Keep one workstream per session or Project. Put durable instructions in the right layer instead of repeating them in every prompt.'

Q78. What is the 'review thresholds vs autonomy ladder' distinction, and why does the course say 'keep the two sets apart'?

- A) Review thresholds are for automated systems; autonomy ladder is for human decisions — they should never be mixed
- B) The autonomy ladder picks the gate on the ACTION (what the agent may do without asking). The review thresholds (stakes, reversibility, audience, regulatory exposure) pick the review on the RESULT (whether a human must check the output before it counts). A rung-one summary can still need review if it goes to a client or involves regulated data
- C) Review thresholds are for senior staff; the autonomy ladder is for junior staff — org hierarchy determines which applies
- D) The two frameworks apply to different vendors — ladder is for Cowork, thresholds are for ChatGPT Work

Correct Answer: B) The autonomy ladder picks the gate on the ACTION (what the agent may do without asking). The review thresholds (stakes, reversibility, audience, regulatory exposure) pick the review on the RESULT (whether a human must check the output before it counts). A rung-one summary can still need review if it goes to a client or involves regulated data

Explanation: 'The rung picks the gate on the action. The four thresholds pick the review on the result.' A rung-one (read and summarize) action can still require review if it meets a threshold condition like going to a client.

Q79. What four questions does the course recommend asking BEFORE scheduling a workflow?

- A) What is the budget? Who approves the cost? What is the retention policy? When should it stop permanently?
- B) What starts the run (clock/event/watch)? What does it touch (every source and destination)? Can it reach all of that with your devices off? How will you know it succeeded? (Plus two safety questions: what may it do without approval? what happens on the empty case?)
- C) Which model should run it? What is the expected output quality? Who reviews the results? When should it be discontinued?
- D) Is the data sensitive? Is the connector reliable? Is the schedule necessary? Can a human do this better?

Correct Answer: B) What starts the run (clock/event/watch)? What does it touch (every source and destination)? Can it reach all of that with your devices off? How will you know it succeeded? (Plus two safety questions: what may it do without approval? what happens on the empty case?)

Explanation: 'The four answers before you schedule: What starts the run? What does it touch? Can it reach all of that with your devices off? How will you know it succeeded? Then add two safety questions: What may it do without approval? What happens on the empty or unclear case?'

Q80. The course says that when an expected output is missing after a schedule runs, 'silence is unclear.' What is the diagnostic sequence?

- A) First rewrite the prompt, then switch to a different model, then try a different product
- B) First check whether the run is waiting for input. Then check usage limits, permissions, connector or browser availability, and the destination. Do not rewrite the prompt or switch products until you know whether it failed to start, failed to reach a tool, paused at a gate, or produced the wrong result
- C) First check the run log for error messages, then restart the workflow, then escalate to vendor support
- D) First compare the output to previous runs, then adjust the schedule timing, then reduce the scope

Correct Answer: B) First check whether the run is waiting for input. Then check usage limits, permissions, connector or browser availability, and the destination. Do not rewrite the prompt or switch products until you know whether it failed to start, failed to reach a tool, paused at a gate, or produced the wrong result

Explanation: 'If an expected output is missing, silence is unclear. First check whether the run is waiting for input. Then check usage limits, permissions, connector or browser availability, and the destination. Do not rewrite the prompt or switch products until you know whether the workflow failed to start, failed to reach a tool, paused at a gate, or produced the wrong result.'

Q81. The course describes 'a block of five rules' for when to use Manual approval specifically. What triggers should make you use Manual?

- A) Tasks over 30 minutes, tasks involving more than 3 connectors, any financial data, any email communication, and any file deletion
- B) New workflows (unfamiliar tools/sites), sensitive or hard-to-reverse actions, new plugins or connectors, actions that send/spend/delete/publish/change important records, and whenever learning the workflow. These are the same cases the course lists throughout
- C) Any task over 1,000 tokens, any task involving external APIs, any task accessing personal data, any task that runs longer than 10 minutes, and any scheduled task
- D) Tasks where the agent disagrees with your instructions, tasks involving multiple steps, tasks that produce files, tasks that access the internet, and tasks that involve calculations

Correct Answer: B) New workflows (unfamiliar tools/sites), sensitive or hard-to-reverse actions, new plugins or connectors, actions that send/spend/delete/publish/change important records, and whenever learning the workflow. These are the same cases the course lists throughout

Explanation: From the course: 'For unfamiliar tools, signed-in sites, new plugins, or high-consequence actions, switch to Manual while you learn the workflow.' And the gate lab: 'Use Manual approval for unfamiliar sites, new plugins or connectors, and actions that send, spend, delete, publish, or change important records.'

Q82. Why does the course say 'connectors come first' on the reach ladder, even though the course also acknowledges connectors are 'a real door' with risks?

- A) Connectors are cheaper than browser actions so they always come first for cost efficiency
- B) The ladder ranks exposure, not safety — connectors give structured permissions (clearer scope, clearer actions, less ambiguity than screen-driving). Being first on the ladder does not mean connectors are risk-free, but they offer the most structured, least ambiguous access path
- C) Connectors are faster to set up so beginners should always start there
- D) Connectors are audited by Anthropic for security before being listed in the directory

Correct Answer: B) The ladder ranks exposure, not safety — connectors give structured permissions (clearer scope, clearer actions, less ambiguity than screen-driving). Being first on the ladder does not mean connectors are risk-free, but they offer the most structured, least ambiguous access path

Explanation: 'That is not because connectors are magically safe. The ladder ranks exposure, not safety. Rung one is still a real door. Connectors come first because structure gives you clearer permissions, clearer actions, and less ambiguity than screen-driving.'

Q83. What does the course say makes a Cowork plugin 'a role-specific mini-harness'?

- A) Plugins are named after job roles (analyst, writer, coder) and restrict the agent to tasks matching that role
- B) A plugin can bundle skills (know-how), connectors (structured access), and sub-agents (specialist workers) AND hooks (automatic rules at lifecycle points) — making it a packaged role with tools, methods, and enforcement baked in together
- C) Plugins enforce a specific output format that matches the role's professional standard
- D) Plugins restrict the AI to information relevant to the role, reducing hallucination on specialized topics

Correct Answer: B) A plugin can bundle skills (know-how), connectors (structured access), and sub-agents (specialist workers) AND hooks (automatic rules at lifecycle points) — making it a packaged role with tools, methods, and enforcement baked in together

Explanation: 'A Claude plugin can bundle skills, connectors, and sub-agents... It can also bundle hooks... both run in Cowork. So a plugin can be a role-specific mini-harness.' It packages the role's know-how, reach, and automatic rules.

Q84. The course says 'the habit to build' for file custody requires a specific question before every important question. What is it?

- A) 'What tier will this file land in?' — ask this before every file creation action
- B) 'What would a capable human colleague need to see before they could answer this well?' — then give that material. Applied to custody: 'What would they need to access?' determines what sources to connect, not just what to upload
- C) 'Is this data regulated?' — always check compliance status before accessing any file
- D) 'Does this need to exit the platform?' — the tier-3 decision should come before every task brief

Correct Answer: B) 'What would a capable human colleague need to see before they could answer this well?' — then give that material. Applied to custody: 'What would they need to access?' determines what sources to connect, not just what to upload

Explanation: 'Before you ask an important question, ask yourself: What would a capable human colleague need to see before they could answer this well? Then give the assistant that material.' Applied to the surface: this drives what sources to connect and what to upload.

Q85. The course introduces the concept that the Cowork side panel 'defaults to Automatically approve.' What is the course's explicit stance on this default for beginners?

- A) The default is appropriate for all beginners since Anthropic designed it to be safe out of the box
- B) 'That is a product default, not a recommendation for every first run. For unfamiliar tools, signed-in sites, new plugins, or high-consequence actions, switch to Manual while you learn the workflow.' The default is a starting setting, not a safety certification
- C) Beginners should immediately change the default to Skip for maximum efficiency
- D) The default is only appropriate for users on enterprise plans who have IT oversight — free plan users should always use Manual

Correct Answer: B) 'That is a product default, not a recommendation for every first run. For unfamiliar tools, signed-in sites, new plugins, or high-consequence actions, switch to Manual while you learn the workflow.' The default is a starting setting, not a safety certification

Explanation: 'The current Cowork side panel defaults to Automatically approve. That is a product default, not a recommendation for every first run. For unfamiliar tools, signed-in sites, new plugins, or high-consequence actions, switch to Manual while you learn the workflow.'

Q86. What does the Loop Engineering appendix say is the key difference between a 'golden set' in an eval suite and using a managed surface for testing?

- A) Golden sets require 1,000+ cases; managed surfaces only support 10-20 test cases
- B) A golden set is version-controlled case files in a repository reviewed and dated like code, with a runner for headless execution and a committed baseline for CI gates. A managed web task can run repeated checks but cannot reproduce the same sequence headlessly in CI or another runtime you operate
- C) Golden sets use real production data; managed surfaces must use synthetic data for safety
- D) Golden sets are created by QA engineers; managed surface tests can be created by non-technical users

Correct Answer: B) A golden set is version-controlled case files in a repository reviewed and dated like code, with a runner for headless execution and a committed baseline for CI gates. A managed web task can run repeated checks but cannot reproduce the same sequence headlessly in CI or another runtime you operate

Explanation: 'The golden set, a version-controlled folder of case files... Cases are JSON files in a repository, reviewed and dated like code. A managed web task is not a substitute for an eval repository you own.'

Q87. The course gives a 'one-line summary' at the very end. What is it?

- A) 'Brief the work, check the plan, land files in tier 3, and match the gate to the stakes.'
- B) 'Chat is where you converse. The agent surface is where you delegate. A cloud agent keeps working after you leave, and it may borrow tools through a desktop bridge. Reach the world through the narrowest tool and the tightest gate the job requires.'

C) 'The six parts are the shape. The two axes are the runtime. The three tiers are the custody. The four checks are the plan review.'

D) 'Build your brief, inspect your plan, protect your data, and own your results.'

Correct Answer: B) 'Chat is where you converse. The agent surface is where you delegate. A cloud agent keeps working after you leave, and it may borrow tools through a desktop bridge. Reach the world through the narrowest tool and the tightest gate the job requires.'

Explanation: The course's own one-line summary: 'Chat is where you converse. The agent surface is where you delegate. A cloud agent keeps working after you leave, and it may borrow tools through a desktop bridge. Reach the world through the narrowest tool and the tightest gate the job requires.'

Q88. What does it mean that 'Codex is not selectable on web or mobile' but 'supported desktop Codex sessions open from the Remote tab in the ChatGPT mobile app'?

A) Codex only works on specific Apple devices that support the Remote tab feature

B) Codex stays a desktop experience (not available directly on web/mobile), but if you have an active Codex session running on your desktop, you can view and interact with it through the mobile app's Remote tab — the session runs on the desktop, the mobile is a window onto it

C) Codex is gradually being deprecated and the Remote tab is a legacy access path

D) The Remote tab is a ChatGPT Work feature that allows web access to coding agents — it is only in beta as of the course date

Correct Answer: B) Codex stays a desktop experience (not available directly on web/mobile), but if you have an active Codex session running on your desktop, you can view and interact with it through the mobile app's Remote tab — the session runs on the desktop, the mobile is a window onto it

Explanation: 'Codex is OpenAI's separate coding agent, its sibling to Claude Code. It stays a desktop experience and is not selectable on web or mobile, although supported desktop Codex sessions open from the Remote tab in the ChatGPT mobile app.'

Q89. When briefing unattended work, the course says to include an 'escalation' component. What should trigger escalation, and what should the agent do when it triggers?

A) Any action that costs more than a set token threshold should trigger escalation to show a cost warning

B) 'What uncertainty should stop the run and come back to you?' — when the agent cannot determine the right path forward due to missing evidence, ambiguous instructions, or a situation not covered by the brief, it should stop and report back rather than guess and continue

C) Any action involving sensitive data should trigger escalation to your designated data protection officer

D) Escalation should trigger after every 10 steps to provide regular progress updates to the user

Correct Answer: B) 'What uncertainty should stop the run and come back to you?' — when the agent cannot determine the right path forward due to missing evidence, ambiguous instructions, or a situation not covered by the brief, it should stop and report back rather than guess and continue

Explanation: 'Escalation: what uncertainty should stop the run and come back to you?' Combined with the missing-evidence rule: 'If the necessary evidence is missing, say what is missing and stop rather than inventing it.'

Q90. The Cowork and OpenWork appendix says 'transfers only with a desktop bridge' for four specific capabilities. Which of these four is correctly described?

A) Connectors transfer only with a bridge because MCP requires local routing

B) The built-in Cowork browser transfers only with a bridge — 'the browser lives in Claude Desktop. Web and mobile can steer it while the desktop app is online.' Without the desktop being online, the built-in browser is unavailable to the cloud session

C) Projects transfer only with a bridge because project files are stored locally

D) Scheduled tasks transfer only with a bridge because the scheduler runs on the local machine

Correct Answer: B) The built-in Cowork browser transfers only with a bridge — 'the browser lives in Claude Desktop. Web and mobile can steer it while the desktop app is online.' Without the desktop being online, the built-in browser is unavailable to the cloud session

Explanation: 'Built-in Cowork browser: The browser lives in Claude Desktop. Web and mobile can steer it while the desktop app is online.' This is explicitly listed under 'Transfers only with a desktop bridge.'

Q91. The course says the managed harness is 'rented' and gives five specific things the vendor still owns. What are they?

- A) The model weights, training data, API keys, pricing, and terms of service
- B) The runtime, sandbox, lifecycle, model routing, and enforcement machinery — 'on a managed surface you do not fully control' these. Cowork plugins add extension points but do not transfer ownership of the core mechanisms
- C) The hardware, cooling systems, network infrastructure, data centers, and power supply
- D) The UI, documentation, support, update schedule, and deprecation timeline

Correct Answer: B) The runtime, sandbox, lifecycle, model routing, and enforcement machinery — 'on a managed surface you do not fully control' these. Cowork plugins add extension points but do not transfer ownership of the core mechanisms

Explanation: 'The vendor still owns important parts of the harness. On a managed surface you do not fully control the runtime, sandbox, lifecycle, model routing, or enforcement machinery.'

Q92. What is the practical meaning of 'the vendor spine for active work, your own system of record for finished work' as a working strategy?

- A) Use the vendor's tools to draft and iterate actively; move completed deliverables to tier-3 destinations (Drive, email, repository, download) where your retention rules apply and the work survives vendor changes
- B) Use the vendor for paid subscriptions and your own hardware for free operations
- C) Run all active tasks in the cloud and archive completed tasks on your local machine
- D) Trust the vendor for security during active work and apply your own encryption after completion

Correct Answer: A) Use the vendor's tools to draft and iterate actively; move completed deliverables to tier-3 destinations (Drive, email, repository, download) where your retention rules apply and the work survives vendor changes

Explanation: This is the combination of Concept 4 (vendor spine for working context) and Concept 5 (tier-3 for finished work): 'Keep operational state in the vendor spine while work is active, and give critical state a portable source of truth you control.'

Q93. The course says 'a connection permission is not a formality' specifically about connectors. What should you check beyond the permission name?

- A) Check that the connector developer is a verified partner in the official directory
- B) Inspect what each permission actually allows — 'a mail connector can read threads' does not authorize sending, even if the UI groups read and send together. Grant the smallest scope that completes the workflow, and verify what each specific permission allows in that connector's implementation
- C) Check that the connector has been updated within the last 90 days for security patches
- D) Verify that your organization's IT team has approved this specific connector for business use

Correct Answer: B) Inspect what each permission actually allows — 'a mail connector can read threads' does not authorize sending, even if the UI groups read and send together. Grant the smallest scope that completes the workflow, and verify what each specific permission allows in that connector's implementation

Explanation: 'A permission label is not a formality. Read is not send, draft is not publish, and view is not edit.' You must inspect the specific permissions, not just their categories.

Q94. The course describes ChatGPT Work's 'Plan mode' as an example of Concept 8 (brief, plan, approve, review) built into a product. Why is it NOT a 'seventh part' of the six-part harness?

- A) Plan mode is a paid feature that not all users have access to, so it cannot be a universal part

B) Plan mode is an implementation of the delegation loop (Concept 8) within the run-until-done part of the harness — it gathers context, proposes an approach, and waits for approval. It is a human-gate mechanism, not a new structural part of the harness

C) Plan mode is too new to be considered a permanent part of the harness architecture

D) Plan mode is vendor-specific and therefore cannot be part of a product-neutral harness framework

Correct Answer: B) Plan mode is an implementation of the delegation loop (Concept 8) within the run-until-done part of the harness — it gathers context, proposes an approach, and waits for approval. It is a human-gate mechanism, not a new structural part of the harness

Explanation: 'ChatGPT Work Plan mode is a current example. It gathers context, asks questions, creates a step-by-step plan, and lets you approve it before execution. That is not a seventh part. It is Concept 8 built into a product.'

Q95. The course says 'walk it before you schedule it.' What does 'walking' a workflow mean, and what do you inspect during the walk?

A) Walking means reading the workflow brief aloud to check for clarity before submitting it

B) Running the workflow by hand at least twice and inspecting the sources, plan, permissions, output, empty case, and tier-3 destination. 'Scheduling is photocopying the workflow into the future. A mistake inside it is not diluted. It is reproduced.'

C) Walking means having a human teammate review the workflow brief before it is first run by AI

D) Walking means running the workflow in verbose mode to see every intermediate step before trusting the final output

Correct Answer: B) Running the workflow by hand at least twice and inspecting the sources, plan, permissions, output, empty case, and tier-3 destination. 'Scheduling is photocopying the workflow into the future. A mistake inside it is not diluted. It is reproduced.'

Explanation: 'Walk it before you schedule it. Run the workflow by hand at least twice. Inspect the sources, plan, permissions, output, empty case, and tier-3 destination. Only then attach a heartbeat.'

Q96. What is the 'state summary' approach recommended when a long chat has gone stale, and what signal tells you a state summary is needed?

A) A state summary is only needed when the AI refuses to answer — it is a reset mechanism for blocked conversations

B) When signs of staleness appear (forgotten rules, contradicted figures, generic drift), ask for a short record of decisions made, open items, and rules given — designed to paste into a fresh chat with files reattached. If you find yourself pasting the same summary every week, that signals a Project is needed

C) State summaries are generated automatically every 50 turns — the platform sends them to your email

D) A state summary is a token-usage report — request it when you suspect you are approaching your context limit

Correct Answer: B) When signs of staleness appear (forgotten rules, contradicted figures, generic drift), ask for a short record of decisions made, open items, and rules given — designed to paste into a fresh chat with files reattached. If you find yourself pasting the same summary every week, that signals a Project is needed

Explanation: From Concept 3 (referenced from the earlier Quick Reference course): 'Ask for a state summary first, a short record of where the work stands... paste it into a fresh chat with the files attached again.' Three staleness signals: forgetting rules, contradicting figures, drifting to generic advice.

Q97. The course gives a specific example of why 'capability and readiness are different things' for scheduling. What example does it use?

A) The fact that cloud Routines can send emails does not mean you should schedule mass email campaigns on day one

B) 'This is a teaching ceiling, not a claim that the products cannot act' — beginners should start with reporting-only schedules (read, analyze, report) even though the product supports acting. Schedules that

act 'wait for Loop Engineering' and require clean runs, a clear checker, bounded permissions, and reversible actions first

C) The product can process 1,000 files per hour but most users only need 10 — capability exceeds typical readiness

D) Cloud sessions can run 24/7 but most users are only ready for 8-hour windows — time-of-day readiness matters

Correct Answer: B) 'This is a teaching ceiling, not a claim that the products cannot act' — beginners should start with reporting-only schedules (read, analyze, report) even though the product supports acting. Schedules that act 'wait for Loop Engineering' and require clean runs, a clear checker, bounded permissions, and reversible actions first

Explanation: 'For beginners, the safest scheduled agent reads, analyzes, and reports, so it changes nothing while you sleep. This is a teaching ceiling, not a claim that the products cannot act. Capability and readiness are different things.'

Q98. The course describes how the 'old boundary' between web and desktop has changed. What specific example shows this most clearly?

A) The ChatGPT web interface can now run Python code directly without any desktop component

B) A cloud Cowork session can steer the built-in browser (which lives in Claude Desktop) while the desktop app is online — the loop is remote, the browser tool is local, and they cooperate in one task. The old 'web is remote, desktop is local' became too simple

C) Local Claude Desktop sessions now automatically sync to the cloud when Wi-Fi is connected

D) Desktop apps now run entirely in the browser through WebAssembly, eliminating the need for local installation

Correct Answer: B) A cloud Cowork session can steer the built-in browser (which lives in Claude Desktop) while the desktop app is online — the loop is remote, the browser tool is local, and they cooperate in one task. The old 'web is remote, desktop is local' became too simple

Explanation: 'A cloud Cowork task can run its agent loop in Anthropic's cloud while a browser action runs through a connected desktop. The old picture, web is remote, desktop is local, became too simple.'

Q99. The course recommends checking COMPLIANCE before uploading regulated data, not after. What does the compliance row in the routing table specify for the check?

A) Submit a ticket to your IT department and wait for approval before any AI use with business data

B) 'Stop and verify. Availability is not permission. Get the written compliance answer for the exact surface.' And: the course specifies that for HIPAA, an event-triggered Work task in ChatGPT for Healthcare is not BAA-covered and must not process PHI — the exact surface and configuration determine coverage, not just the product category

C) Review the vendor's general privacy policy and terms of service before uploading any data

D) Ask the AI assistant whether it can handle the data type — if it says yes, that constitutes compliance verification

Correct Answer: B) 'Stop and verify. Availability is not permission. Get the written compliance answer for the exact surface.' And: the course specifies that for HIPAA, an event-triggered Work task in ChatGPT for Healthcare is not BAA-covered and must not process PHI — the exact surface and configuration determine coverage, not just the product category

Explanation: 'Regulated or contractually restricted data: Stop and verify. Availability is not permission. Get the written compliance answer for the exact surface.' The course gives HIPAA as a concrete example of why general enterprise features do not equal specific regulatory coverage.

Q100. The course ends with 'the habits, in one place.' How many habits are listed, and which habit is the FIRST one recommended for when the workflow is unclear?

A) 7 habits; the first is 'always use Manual mode for new workflows'

B) Multiple habits; the first is 'Start with the frustration when the workflow is unclear, then let the agent interview you' — workflow discovery before workflow execution is the recommended starting point when you do not yet know the procedure

C) 10 habits; the first is 'always export to tier 3 before starting any task'

D) 5 habits; the first is 'choose the right routing before anything else'

Correct Answer: B) Multiple habits; the first is 'Start with the frustration when the workflow is unclear, then let the agent interview you' — workflow discovery before workflow execution is the recommended starting point when you do not yet know the procedure

Explanation: 'Start with the frustration when the workflow is unclear, then let the agent interview you.' This is listed first in the habits section, matching the workflow discovery concept from Concept 8.

CodeVerse Soban

Like · Subscribe · Stay Sharp

End of Exam — Best of Luck!

AI Agent Factory | General Agents on the Web | 100 MCQs