

Workflow Design & Diagnosis Assessment

Question 1 [SINGLE-ANSWER]

Scenario: A cross-functional supply chain committee issues a project charter for an automated inventory rebalancing workflow across regional distribution centers. The charter states: *"The system must reorder SKU buffers promptly whenever regional stock levels experience abnormal depletion, distributing stock evenly without causing stockouts in feeder depots."* During requirements review, an operations analyst points out that the statement cannot yet be built into a workflow.

Which element of the charter makes this requirement non-decidable, and what must happen FIRST before configuring the workflow?

- A) The requirement lacks an assigned technical owner; the engineering team must first select whether a Skill or standing instruction carries the logic.
- B) "Promptly" and "abnormal depletion" lack objective verification thresholds; the analyst must obtain specific quantitative trigger rules and timeframe bounds via clarifying questions.
- C) The model's context window cannot ingest multi-depot stock data; the data pipeline must first be split across independent prompt sessions.
- D) The requirement involves multiple distribution centers; the workflow must first be classified as human-retained due to multi-stakeholder variance.

Question 2 [SELECT ALL THAT APPLY]

Scenario: A senior logistics manager receives an informal message from the VP of Operations: *"Can we have regular visibility on carrier SLA breaches so we don't get blindsided at quarterly business reviews?"* The manager instructs an AI assistant to convert this request into a formal task definition rather than an executive summary.

According to the task definition framework, which of the following represent the mandatory settled components required to prevent the model from closing unspecified parameters with arbitrary defaults? (Select all that apply.)

- A) What exact artifact or deliverable is produced
- B) Who specifically the output is for
- C) How often the workflow runs
- D) What data sources the workflow draws upon
- E) In what format the output arrives
- F) Which commercial LLM tier provides the lowest token cost

Question 3 [SINGLE-ANSWER]

Scenario: A financial operations team uses an AI assistant to analyze quarterly cloud infrastructure spending across 40 cost centers. The prompt asks the assistant to: *"Analyze our AWS billing CSV, calculate the quarter-over-quarter percentage variance for each department, identify the top three cost overruns, and write an executive briefing explaining the drivers."* The resulting table displays cleanly formatted figures, percentages, and cost variance totals that seem completely plausible. However, an internal audit reveals that while the prose analysis was sound, three department totals were mathematically incorrect by 4% to 7%.

What is the root cause of this failure, and what is the strict design rule to prevent it?

- A) The prompt was under-specified; the user should have provided 3-shot few-shot examples of exact arithmetic in prose.

Question 4 [SINGLE-ANSWER]

Scenario: A hospital administration team is designing an automated clinician scheduling workflow. The proposed workflow ingests historical patient load, scheduled surgeries, contracted clinician hours, and regulatory fatigue limits. Step 4 of the workflow instructs the AI to: *"Recommend shift assignments and decide whether to approve or deny vacation requests submitted by senior attending physicians during upcoming holiday weekends."* Historical shift data and vacation forms are uploaded to the Project knowledge base. However, during the holidays, staffing tensions between departments are informally negotiated among department heads based on unspoken seniority privileges and unwritten departmental political trade-offs.

Applying the Window Test, how should Step 4 be evaluated and structured?

- A) Step 4 passes the Window Test because historical scheduling files and vacation forms fit easily within the context window; AI can own the approvals.
- B) Step 4 fails the Window Test because the decisive inputs (unwritten political trade-offs, interpersonal negotiations) cannot be supplied into the window; the approval decision must remain human-retained.
- C) Step 4 is AI-appropriate provided code execution is configured to optimize the shift distribution algorithm across all physicians.
- D) Step 4 is a synthesis step that should be delegated to a Skill with a standing instruction to enforce seniority rules.

Question 5 [SELECT ALL THAT APPLY]

Scenario: A corporate market intelligence team is configuring an automated competitor tracking workflow. The workflow requires retrieving three types of information:

1. The standard theoretical mathematical formulas for calculating Customer Lifetime Value (CLV) and Churn Rate;
2. The current subscription pricing tiers announced by a key SaaS rival yesterday; and
3. The effective regulatory compliance deadline for a recently enacted regional data privacy statute passed last week.

Under the research grounding principles of Workflow Design & Diagnosis, which of the following statements correctly direct this workflow? (Select all that apply.)

- A) Standard theoretical methods and established formulas may safely draw upon the model's internal parametric knowledge without external link verification.
- B) Current competitor pricing announced yesterday sits outside the model's temporal window and requires a live external source that can be opened and verified, never a confident parametric answer.
- C) Recent statutory compliance deadlines require an external source link because asking the model directly risks a plausible, confident hallucination based on typical statutory language.
- D) If an external web search tool is unavailable, the model should be prompted to estimate competitor pricing using logical synthesis of past industry trends.
- E) Passing the window test automatically makes a research step fully owned by AI without requiring human review of the retrieved citations.

Question 6 [SINGLE-ANSWER]

Scenario: During an executive design sprint, a fintech engineering manager suggests: *"Step 3 involves calculating the penalty interest rate on delinquent commercial loans. Because we have developed a deterministic Python script via code execution that guarantees zero calculation errors, Step 3 should be classified as an AI-owned step that dispatches notices directly to borrowers."*

Which structural flaw has the engineering manager committed according to the three-decision delegation framework?

Question 7 [SINGLE-ANSWER]

Scenario: A pharmaceutical company is designing a regulatory submission workflow. Step 8 involves verifying whether batch release test documentation conforms to the Ministry of Health's mandatory safety threshold. An error at this step could result in an unapproved drug batch being distributed to pharmacies, triggering severe legal penalties, patient harm, and product recalls. However, the verification is performed on a staging database where flagging a record as "non-compliant" can be reversed with a single click before physical dispatch.

An analyst argues that because the staging database flag is completely reversible, Step 8 should be classified as AI-appropriate. Why is this argument incorrect under the ownership framework?

- A) Reversibility is irrelevant in regulated industries; all steps must be classified as human-retained.
- B) The criteria are evaluated together, and where they disagree, the strictest criterion wins: the catastrophic bad-case stakes and ultimate regulatory accountability mandate human ownership.
- C) Reversibility only applies to prose drafting; database operations are inherently irreversible.
- D) The step should be classified as AI-appropriate, but only if executed using code execution rather than a Skill.

Question 8 [SINGLE-ANSWER]

Scenario: In a corporate commercial leasing firm, a junior associate argues: *"We cannot allow AI to draft the redline and financial penalty exposure table for tenant lease negotiations because our Senior Legal Partner is ultimately accountable for every term in the lease."*

Evaluating this argument against the accountability criterion in the delegation framework, which assessment is most accurate?

- A) The associate is correct; because accountability never moves onto the machine, no step in a legal workflow can be touched by AI.
- B) The associate is incorrect; drafting a redline and computing exposure generates intermediate inputs that the Partner reviews and judges before execution, so delegation is appropriate as long as the Partner retains the approval decision.
- C) The associate is correct; financial penalty exposure calculations are inherently high stakes and therefore cannot be carried by code execution.
- D) The associate is incorrect; accountability moves to the machine once the AI redlining Skill demonstrates 99% accuracy across benchmark evaluations.

Question 9 [SINGLE-ANSWER]

Scenario: An enterprise SaaS company has deployed an AI assistant to handle Tier-1 customer billing inquiries. For four consecutive months, the assistant has generated exceptionally polite, highly accurate explanatory draft emails reconciling complex prorated subscription invoices. Encouraged by these results, the customer support director updates the delegation map to allow the AI assistant to autonomously issue refunds up to \$500 directly into customer bank accounts without human supervisor approval.

Which mapping error has occurred, and what fundamental principle does it violate?

- A) The unstaffed gate, because the support director failed to schedule a quarterly review date for the billing workflow.
- B) Mapping the tool instead of the work, because the team built a refund Skill that forced the workflow into an automated shape.
- C) Halo delegation, because drafting quality is evidence about the draft, not evidence about the decision to disburse funds.

Question 10 [SELECT ALL THAT APPLY]

Scenario: A medical device manufacturing company established a collaborative workflow for reviewing supplier quality certificates. The delegation map explicitly specifies: *"Step 4 (Collaborative): AI cross-checks supplier test certificates against ISO compliance specs; Lead Quality Engineer reviews and signs off on certificate acceptance."* Over an eight-week period, supplier volume increased from 15 certificates per week to 140 per week. Facing severe backlogs, the Lead Quality Engineer begins approving the entire queue with a single bulk click every Friday afternoon without opening the underlying reports or inspection logs.

Which of the following statements accurately characterize this workflow's operational reality under the delegation framework? (Select all that apply.)

- A) Step 4 has functionally decayed from a collaborative step into an automated step in practice, despite what the written delegation map states.
- B) The failure emits no immediate software alarm or system error message because the map records what was designed, not what is running.
- C) This issue represents an unstaffed gate, where review becomes a glance and then a click due to capacity strain and overconfidence.
- D) Because the Lead Quality Engineer's digital signature is logged on every batch, regulatory accountability has safely transferred to the software system.
- E) The appropriate governance remedy is to conduct a scheduled read of the map and verify how long the named reviewer actually spends reviewing individual cases.

Question 11 [SINGLE-ANSWER]

Scenario: A marketing team builds an advanced prompt Skill that excels at generating 3-bullet executive summaries with sentiment tags from raw customer transcripts. To maximize the utility of this new capability, the project manager inserts an additional mandatory approval step into their weekly newsletter workflow: requiring all newsletter drafts to first be reduced to a 3-bullet sentiment summary, even though the newsletter editorial team never uses sentiment tags and previously drafted articles directly from long-form research notes.

Which mapping pathology does this workflow revision illustrate, and why is it notoriously difficult to identify from within the team?

- A) Halo delegation; difficult to identify because the newsletter editors are legally accountable for publications.
- B) The unstaffed gate; difficult to identify because the approval gate runs asynchronously without a fixed review timestamp.
- C) Mapping the tool instead of the work; difficult to identify because the Skill itself is genuinely high-performing and executed well.
- D) Context overload; difficult to identify because sentiment tagging consumes excessive tokens across long transcripts.

Question 12 [SELECT ALL THAT APPLY]

Scenario: A healthcare operations committee is reviewing an automated patient discharge documentation workflow that was deployed six months ago. The workflow was designed with collaborative review gates and standing compliance instructions.

According to the delegation map governance framework, which of the following elements are strictly required to ensure the map functions as an operational document rather than an obsolete static decision? (Select all that apply.)

- A) A single named individual assigned as the map owner who is answerable for whether the map still describes operational reality.

Question 13 [SINGLE-ANSWER]

Scenario: A compliance auditor investigates an enterprise credit-scoring workflow. Over a nine-month period, three separate controls degraded:

1. A credit manager review gate became a 5-second cursory click;
2. A standing instruction to filter out jurisdiction-specific restricted postal codes was silently dropped during model generation; and
3. The written delegation map still showed a legacy review pipeline that was abandoned three months prior.

What critical structural vulnerability do these three decaying controls share under the workflow diagnosis model?

- A) All three failures are caused by rate limits and context window truncation in the API.
- B) All three controls produce zero system alarms or software error signals at the exact moment they fail, making a scheduled read the only effective defense.
- C) All three controls can be fully restored by upgrading the foundational model tier from a fast tier to a reasoning tier.
- D) All three failures indicate that the underlying tasks failed the initial data eligibility gate.

Question 14 [SINGLE-ANSWER]

Scenario: A business intelligence developer is building an interactive weekly executive briefing generator. The developer attempts to prototype the tool across eight separate, unorganized ad-hoc chat sessions with an AI model. After two weeks of work, the team finds that each draft incorporates contradictory tone constraints, formatting fluctuates wildly, and earlier agreed-upon metric exclusions are missing.

Why does the build loop (Ideate → Prototype → Feedback → Refine) require execution inside a persistent Project rather than scattered one-off chats?

- A) Because one-off chats automatically disable web browsing and Python execution environments.
- B) Because across scattered chats, each round restarts from a slightly different baseline understanding, producing disconnected drafts rather than progressive refinement.
- C) Because only Projects have the technical capability to export code execution results as web artifacts.
- D) Because the build loop is an engineering procedure that requires formal API token authentication on every prompt.

Question 15 [SINGLE-ANSWER]

Scenario: During Cycle 2 of developing an executive dashboard artifact, a sales analytics team notices that while the customer segmentation charts render accurately, the overall "*Gross Margin Summary*" row shows numbers that look convincing and formatted correctly, but are mathematically inconsistent with the column subtotals. The junior prompt engineer spends three days rewriting and expanding the system prompt with verbose instructions like: "*Make sure your calculations are 100% accurate, double-check all arithmetic, and sum every column carefully.*"

Why is the prompt engineer's intervention fundamentally misguided under the build loop framework?

- A) The issue is a feature problem, not a description problem; totals are calculations that must be assigned to code execution, which no amount of prompt description can resolve.
- B) The prompt engineer failed to provide few-shot examples of arithmetic balance sheets within the Project knowledge base.
- C) The issue is a context overload problem; resetting the session would have permanently resolved the mathematical variance.

Question 16 [SINGLE-ANSWER]

Scenario: A single pricing analyst created a prompt-and-iterate spreadsheet analysis tool to help draft quotes for complex enterprise sales. Nine months later, eight regional sales teams rely on this tool daily, senior account executives use its output during client negotiations, and the Chief Commercial Officer includes its margin forecasts in quarterly board decks. However, the tool still runs inside the original analyst's personal chat session.

What does this situation signal, and what is the proper governance response?

- A) It signals that the prompt should be converted into a personal Skill so the analyst can execute it more quickly.
- B) It represents the escalation signal; other people now depend on it as infrastructure, requiring developer or architect expertise for uptime, access control, and engineering guarantees.
- C) It indicates that the task has become an expectation mismatch, and the workflow should be terminated immediately.
- D) It signals that the tool is ready for automated manufacturing via 'From One-Off to Worker' because the prompt wording is stable.

Question 17 [SINGLE-ANSWER]

Scenario: An HR specialist notices that an automated employee exit interview summarization tool has begun producing flawed outputs. When asked to troubleshoot, the specialist immediately concludes: *"The LLM has become obsolete; we need to switch from Claude 3.5 Sonnet to an expensive reasoning model."* The lead workflow architect stops the specialist and asks: *"When did the symptom first appear?"* The specialist confirms: *"The prompt used to produce flawless summaries throughout January and February, but starting in late March, its outputs started missing our new company severance policies."*

Based on the timing of the symptom, what is the first diagnostic hypothesis and its corresponding cheap confirmation test?

- A) Under-specification; re-read the prompt to see if the severance policy was omitted from the prompt text.
- B) Context overload; restart the prompt in a fresh session and check if it holds.
- C) Wrong feature or model tier; test the prompt once using a higher-tier reasoning model.
- D) Stale configuration; open the Project knowledge base and instruction files to check if the severance policy documents are outdated.

Question 18 [SELECT ALL THAT APPLY]

Scenario: A team of operational analysts is building an incident troubleshooting guide for AI workflow failures across their business unit.

Which of the following pairings correctly link the symptom's timing to its first hypothesis and cheap confirmation test? (Select all that apply.)

- A) Wrong from the very first response → Under-specification → Re-read the prompt to verify if the missing requirement was actually stated.
- B) Started fine, then degraded over a long session → Context overload → Restate the instruction in a fresh session and observe if performance holds.
- C) One repeatable error type across runs → Wrong feature or model tier → Run the step once with the proper tool (e.g., code execution) or a stronger tier.
- D) Used to work reliably, now does not → Stale configuration → Open the saved instructions, Skills, or knowledge sources and inspect their dates.
- E) Output contains subtle arithmetic errors → Expectation mismatch → Rewrite the prompt asking for higher

Question 19 [SINGLE-ANSWER]

Scenario: A portfolio manager asks an AI assistant: *"Predict the exact closing stock price of Company XYZ on December 31 of this year down to the cent, based on today's macroeconomic news."* The assistant generates a confident prediction of \$142.50 with a five-paragraph justification. At year-end, the actual closing price is \$98.20. The portfolio manager files an issue report claiming the assistant suffered from *"under-specification"* because it failed to incorporate Q4 interest rate cuts.

How should this failure be classified under the diagnostic hierarchy?

- A) Under-specification; the prompt should have included a 5-variable financial regression constraint.
- B) Stale configuration; the macroeconomic news uploaded to the knowledge base had expired.
- C) An expectation mismatch; the task demands something the tool cannot do (predicting a precise point future market value), requiring the task to be reshaped into scenarios or sensitivity drivers.
- D) Context overload; the prompt was evaluated on a model tier optimized for speed rather than reasoning.

Question 20 [SINGLE-ANSWER]

Scenario: An executive assistant receives a disappointing meeting briefing from a newly configured project assistant: the summary is completely disorganized, omits participant action items, and blends unrelated topics together. The executive assistant's immediate reaction is to discard the tool and conclude that AI is fundamentally incapable of preparing executive briefings.

According to the cheapest-first diagnostic sequence, what is the very FIRST action the assistant must take before concluding the task is not a fit?

- A) Switch immediately to the most expensive, frontier reasoning model available on the market.
- B) Re-read the prompt and evaluate it against its five core components: role, context, task, constraints, and format.
- C) Clear all Project knowledge base files and upload new corporate style guides.
- D) Check whether the API context window has exceeded 100,000 tokens.

Question 21 [SELECT ALL THAT APPLY]

Scenario: A compliance analyst is diagnosing an under-specified prompt that generated a vague, unformatted executive summary. In Step 1 of the diagnostic sequence, the analyst evaluates the prompt against the five essential components defined in the framework.

Which of the following represent the five mandatory components evaluated during Step 1 of the cheapest-first diagnostic sequence? (Select all that apply.)

- A) Role (the operational lens or persona the assistant adopts)
- B) Context (the background information the model could not have known)
- C) Task (the clear, actionable instruction for what to produce)
- D) Constraints (explicit boundaries, exclusions, or rules on the output)
- E) Format (the structural layout, schema, or styling of the deliverable)
- F) Temperature (the mathematical entropy parameter of token sampling)

Question 22 [SINGLE-ANSWER]

Scenario: A corporate communications director reviews an AI-generated draft of an internal re-organization announcement and complains: *"This is completely wrong—the tone is far too casual and flippant for such a serious corporate change."*

Under the "turn a reaction into an instruction" framework, which action represents the correct instruction and its corresponding architectural lever?

- A) Reaction: "Wrong tone" → Add a formal tone constraint that applies to every draft → Lever: Standing instruction
- B) Reaction: "Wrong tone" → Re-run the generation three times until a better mood is generated → Lever: The prompt
- C) Reaction: "Wrong tone" → Upload the CEO's personal biography to the knowledge base → Lever: The knowledge
- D) Reaction: "Wrong tone" → Convert the announcement generator into a deterministic Python script → Lever: Code execution

Question 23 [SINGLE-ANSWER]

Scenario: A legal tech team needs to configure three separate operational elements for their new M&A due diligence Project:

1. An updated list of sanctioned offshore jurisdictions that must be consulted;
2. A mandatory constraint stating: *"Never disclose target enterprise valuation in executive previews"*; and
3. A sequential 7-step checklist detailing exactly how to extract, categorize, and cross-reference intellectual property warranties from vendor contracts.

According to the "Rule, Reference, or Procedure" test, where does each fix belong?

- A) (1) Standing instruction, (2) Knowledge base, (3) Skill
- B) (1) Knowledge base, (2) Standing instruction, (3) Skill
- C) (1) Skill, (2) Standing instruction, (3) Knowledge base
- D) (1) Knowledge base, (2) Skill, (3) Standing instruction

Question 24 [SINGLE-ANSWER]

Scenario: A procurement specialist discovers that when generating vendor comparison sheets, the AI model frequently includes expired quotes from inactive suppliers. To fix this, the specialist types *"filter out expired supplier quotes"* into the chat prompt every Monday. For six months, this produces accurate tables. In July, the specialist takes a three-week medical leave, and an associate covers the task without typing the reminder. The resulting report contains three expired quotes, leading to a contract dispute.

What systemic breakdown occurred, and what was the required intervention?

- A) The associate lacked AI prompting certification; the associate should have completed advanced prompt training.
- B) The model suffered from context overload; the session should have been cleared prior to July.
- C) The fix was never promoted into configuration; because it was left in one person's conversational habit, it failed to survive a change of operator.
- D) The supplier quotes should have been classified as human-retained due to commercial reversibility.

Question 25 [SELECT ALL THAT APPLY]

Scenario: A software team lead considers relying on an AI assistant's automatic "*Memory*" feature rather than configuring standing instructions and Project knowledge bases for the team's release notes workflow.

Why does Workflow Design & Diagnosis reject model memory as an acceptable substitute for formal configuration in a team workflow? (Select all that apply.)

- A) Memory is per-user; what the model learns about one person's preferences is not accessible to a colleague covering the workflow.
- B) Memory is best-effort and opaque; you cannot inspect, verify, audit, or edit it with deterministic certainty.
- C) Memory cannot be exported or handed to an external auditor or new team member.
- D) Memory features automatic encryption data, preventing team members from viewing past responses.
- E) Formal configuration in a Project provides a shared, readable home that operates identically for anyone running the workflow.

Question 26 [SINGLE-ANSWER]

Scenario: Five customer account managers independently draft monthly account summaries using an AI assistant. Each manager's summaries are completely consistent from month to month. However, when the Vice President of Customer Success reviews all five submissions, the formatting, tone, and analytical frameworks differ radically across the accounts, forcing the VP to spend two hours every month manually standardizing the drafts into a unified deck.

Which friction signal does this scenario illustrate, and why was it invisible to the individual account managers?

- A) Repetition; invisible because managers paste their prompts quickly.
- B) Correction; invisible because the assistant rarely makes grammatical mistakes.
- C) Variance; invisible because each individual's output was internally consistent, meaning the cost exists only between people and surfaces only at the reviewer.
- D) Stale configuration; invisible because the customer accounts are active.

Question 27 [SINGLE-ANSWER]

Scenario: An operations audit reveals that a weekly reporting process requires three separate sequential prompts:

1. Ingests data and formats tables;
2. Takes the table and calculates KPI percentages; and
3. Drafts an executive summary from the table, requiring the user to re-paste company background context each time.

A junior analyst proposes immediately creating three separate Skills to automate each prompt individually.

Why does the workflow optimization framework reject this proposal in favor of consolidating steps before promoting them?

- A) Skills can only be created by external software vendors, not internal analysts.
- B) Consolidation reduces the number of steps by combining prompts that share background context; promoting a step you were about to merge away is unnecessary wasted effort.
- C) Promoting three Skills simultaneously causes severe context window fragmentation.
- D) Consolidation must always be preceded by a formal code execution security audit.

Question 28 [SINGLE-ANSWER]

Scenario: A project manager redesigns a customer onboarding workflow with AI assistance and claims to executive leadership: *"Our new AI process cut onboarding time from 90 minutes down to 30 minutes."* When the Chief Operating Officer asks to see the documentation for the original 90-minute figure, the project manager admits that the time was estimated from memory after the AI tool was already built.

What foundational principle of workflow measurement did the project manager violate?

- A) Baselines must always be recorded using automated biometric tracking software.
- B) You cannot reconstruct a baseline afterwards; your retrospective estimate of the old way will inevitably be biased in the direction that flatters you.
- C) Onboarding workflows should optimize for variance rather than time savings.
- D) A baseline requires at least twelve months of historical timesheet records before any workflow can be modified.

Question 29 [SELECT ALL THAT APPLY]

Scenario: A business process improvement director is establishing target performance metrics across diverse departmental AI workflows.

According to the metric selection matrix in Workflow Design & Diagnosis, which of the following pairings correctly match the workflow type to its primary optimization metric and rationale? (Select all that apply.)

- A) Internal draft → Optimize for Time → Speed is the point, and small errors are cheap.
- B) Customer-facing report → Optimize for Consistency → Varying formats erode client trust faster than minor delays.
- C) Compliance or finance output → Optimize for Accuracy → One wrong figure costs far more than an hour saved.
- D) Work done by several people across a team → Optimize for Variance → The cost lives between people, not within an individual person.
- E) Regulatory audit filing → Optimize for Speed → Submitting ahead of the statutory deadline eliminates legal liability.

Question 30 [SINGLE-ANSWER]

Scenario: A department head wants to immediately terminate an existing manual vendor invoicing verification process on Monday morning and switch 100% of transactions to a newly configured AI workflow. The lead risk architect objects, insisting on running both the legacy manual process and the AI workflow in parallel for three weekly cycles, despite the duplicated operational labor.

What three vital business benefits does a parallel run purchase that justify this temporary duplication of effort?

- A) It doubles API token throughput, trains the underlying foundational model weights, and satisfies ISO-9001 copyright requirements.
- B) It buys empirical evidence that the new method is superior on your chosen metric, provides a working fallback if it fails, and secures the genuine buy-in of staff whose daily work is changing.
- C) It eliminates the need for a named map owner, automates code execution, and removes review gates.
- D) It proves that the workflow is fully automated, allowing management to eliminate human oversight permanently.

Question 31 [SELECT ALL THAT APPLY]

Scenario: A financial operations director presents a business case to defend the budget for an AI-assisted invoice auditing workflow that has been running for six months.

Under the economic evaluation model in Workflow Design & Diagnosis, which of the following statements must be true for the business case to survive financial and risk scrutiny? (Select all that apply.)

- A) Total workflow cost must include both the software subscription/token costs and the human labor time required to staff retained review gates.
- B) Projected savings must be calculated by multiplying measured baseline time savings by execution frequency and the number of analysts running the process.
- C) The business case must explicitly state the operational and error risks that survive the human review gates, rather than claiming zero residual risk.
- D) Human review gates should be excluded from cost accounting because internal employee salaries are fixed overhead.
- E) An AI workflow redesign that hides residual risk ceases to be credible the moment an unexpected failure occurs.

Question 32 [SELECT ALL THAT APPLY]

Scenario: An enterprise risk management committee mandates that all newly proposed generative AI workflows include a predefined *"First-Hour Incident Protocol"* before being approved for production.

According to Concept 17 and Slide 25, which four questions must the workflow designer settle at design time to govern when a bad output escapes to external stakeholders? (Select all that apply.)

- A) How the workflow is stopped immediately, by whom, and whether they can halt it without prior managerial permission.
- B) How the blast radius / scope is checked to determine how many outputs were affected and where they were sent.
- C) Who must be notified (e.g., customer, internal owner, compliance), and within what strict deadline.
- D) What specific change will be recorded in the delegation map, along with a formal revision date.
- E) Which insurance underwriter will absorb third-party indemnification claims.

Question 33 [SINGLE-ANSWER]

Scenario: A financial advisory firm deploys an automated portfolio rebalancing alert tool. On a volatile trading morning, an associate notices that the tool is dispatching incorrect margin call warning emails to high-net-worth clients based on a corrupted pricing feed. The associate wants to halt the automated workflow immediately, but company policy requires scheduling an emergency sign-off meeting with the VP of Engineering, who is currently on a transatlantic flight.

What critical failure in incident design does this operational bottleneck expose?

- A) The prompt was under-specified regarding margin requirements.
- B) The workflow designer failed to ensure that the designated responder could halt the workflow immediately without requiring permission.
- C) The workflow should have been executed using code execution rather than a Skill.
- D) The associate lacked administrative root access to the cloud server repository.

Question 34 [SINGLE-ANSWER]

Scenario: A regional bank's loan pre-qualification workflow relied on an uploaded regulatory credit-tier policy PDF in the Project knowledge base. On October 15, the bank discovers that the regulatory policy had changed on August 1, rendering two and a half months of issued pre-qualification letters non-compliant. The compliance officer asks: *"Did this one output generated today fail the check?"*

Why does this question reflect a fundamental misunderstanding of stale configuration failures, and what question should the officer have asked instead?

- A) The officer should have asked which model tier was used, because higher tiers automatically override stale knowledge files.
- B) A stale configuration does not produce a single isolated failure; it silently corrupts every output generated since the change occurred, requiring an inquiry into: *"How many total outputs were affected since August 1, and where did they go?"*
- C) The officer should have asked whether the applicant's credit score can fit inside the context window.
- D) The officer should have asked why the loan letters were not formatted using a Skill.

Question 35 [SELECT ALL THAT APPLY]

Scenario: A project manager is drafting a governance charter and presentation slides to explain a new AI-assisted vendor procurement workflow to internal risk and external audit committees.

According to Concept 18 and Slide 24, which of the following phrases dangerously overstate capabilities and must be strictly avoided to preserve stakeholder credibility? (Select all that apply.)

- A) *"Fully automated"*
- B) *"AI handles vendor compliance screening"*
- C) *"The assistant is basically as good as a human paralegal at identifying indemnification risks"*
- D) *"Drafting is AI-assisted, subject to mandatory review and approval by a certified procurement officer"*
- E) *"The system automatically flags deviations from the playbook, but human reviewers retain final sign-off"*

Question 36 [SINGLE-ANSWER]

Scenario: A lead systems architect is defending an automated customer credit evaluation workflow before a regional banking regulator. To convince the regulator that the workflow has sound operational boundaries and will not act recklessly during anomalous market conditions, the architect adds a third clause beyond stating what the tool does and naming the human gate.

What is this crucial third element, and what does it prove to stakeholders?

- A) Stating the prompt temperature and top-p sampling rates, proving the model is deterministic.
- B) Explicitly naming what the workflow refuses (the edge cases, unusual contracts, or out-of-bounds requests it is built to decline and route to humans), proving the boundary was designed in advance rather than discovered via an incident.
- C) Promising that human intervention will be eliminated within twelve months, proving aggressive operational efficiency.
- D) Providing a warranty that the model architecture will never experience hallucinations.

Question 37 [SINGLE-ANSWER]

Scenario: A corporate workflow designer needs to explain an AI-assisted commercial lease review workflow to three different stakeholders:

1. The Lead Real Estate Attorney;
2. The Chief Operating Officer (Practice Executive); and
3. The Enterprise Risk & Compliance Committee.

Which of the following descriptions accurately matches the specific information needs and framing required for the Practice Executive (COO)?

- A) *"It extracts lease covenants, marks deviations from our leasing handbook, and writes the initial redline. Approving changes is still yours; watch out for implied tenant obligations, which it can miss."*
- B) *"Turnaround on a standard lease has dropped from four days to under six hours, with the same approval standard. A licensed attorney still signs off on every modification before it leaves the building."*
- C) *"Drafting is AI-assisted. A qualified attorney reviews and approves every individual term, and nothing is executed without that formal sign-off."*
- D) *"The AI handles commercial lease reviews end-to-end, liberating lawyers from tedious manual review."*

Question 38 [SINGLE-ANSWER]

Scenario: An executive sponsor asks an AI implementation team: *"Our senior marketing manager uses AI every day and gets fantastic results. Why do we need to spend time formalizing delegation maps, naming step owners, and establishing persistent Projects across the department?"*

How does the core philosophy of Workflow Design & Diagnosis answer the executive's question?

- A) Personal usage is illegal under regional enterprise labor regulations.
- B) *"I use AI"* is a personal habit that lives in one person's head, varies by mood, and disappears when they take leave; *"our workflow uses AI"* is an institutionalized team process where named steps are executed consistently whoever is at the keyboard.
- C) Individual prompting consumes significantly more electricity and server bandwidth than team Projects.
- D) Personal usage only works with open-source models, whereas enterprise workflows require commercial APIs.

Question 39 [SINGLE-ANSWER]

Scenario: In the standard contract review delegation map (Concept 6), the step *"Compute financial exposure of a penalty clause"* is classified as AI-owned (carried by code execution), citing the criterion *"Reversible, and checked at the approval gate below."*

What happens if executive management decides to streamline the workflow by eliminating the downstream lawyer approval gate?

- A) The step becomes more efficient because code execution guarantees mathematical certainty.
- B) The penalty-exposure step immediately ceases to be AI-appropriate and must become human-retained, even though the underlying arithmetic has not changed.
- C) The step remains AI-appropriate as long as the prompt is promoted into a standing instruction.
- D) The step automatically transfers its legal accountability to the code execution sandbox.

Question 40 [SINGLE-ANSWER]

Scenario: In the worked delegation map for employee onboarding (Concept 6), Step 3 is: *"Personalise the welcome note"* (Owner: Collaborative, Criterion: *"The manager's voice cannot be supplied"*, Carried by: Named reviewer).

Why cannot the personalization of the welcome note be classified as fully AI-appropriate under the Window Test and ownership framework?

- A) Because welcome notes are legally binding employment contracts.
- B) Because the hiring manager's authentic personal voice, interpersonal context, and unique tone cannot be fully supplied into the context window, requiring the manager to judge and refine the draft.
- C) Because AI models are strictly prohibited from generating informal text.
- D) Because the step is carried by code execution, which cannot output prose.

Question 41 [SINGLE-ANSWER]

Scenario: A workflow designer compares two workflows:

1. In contract review, *"Flag departures from company playbook"* is classified as AI-appropriate;
2. In employee travel expense reimbursement, *"Check expenses against travel policy"* is classified as Collaborative.

Why does the identical analytical action (comparing text against a written policy) receive two different ownership classifications?

- A) Travel expense receipts are images, whereas commercial contracts are text documents.
- B) The action did not change; what changed is whether a human is standing after it. In contract review, a lawyer judges every departure at the redline step; in travel reimbursement, nothing downstream re-checks the flags, so the check itself must be collaborative.
- C) Contract review has lower financial stakes than employee travel expenses.
- D) Travel expense rules cannot be converted into a Skill.

Question 42 [SINGLE-ANSWER]

Scenario: A solutions architect uses a two-prompt sequence to process a complex client RFP: Prompt 1 extracts all requirements into a table with source section citations; Prompt 2 asks the model to attack its own extracted list by identifying ambiguities, multiple interpretations, and indirect requirements.

What is the primary economic rationale for executing Prompt 2 (Attack) immediately after extraction?

- A) It automatically rewrites the client's RFP into legally binding contract language.
- B) It shifts verification leftward: an ambiguity caught at the brief stage costs a clarifying question, whereas catching it after the workflow is built costs the entire work.
- C) It doubles the context window size of the underlying Project workspace.
- D) It eliminates the need for human lawyers to review the winning bid.

Question 43 [SINGLE-ANSWER]

Scenario: In the capacity planning case study (Concept 4), an operations lead asked an AI assistant to analyze support ticket trends, calculate quarterly volume growth, and recommend next quarter's headcount. The assistant produced a well-structured, highly confident plan recommending that the company hire three new analysts. However, the plan was useless because the company had announced an internal hiring freeze two weeks earlier during an executive meeting.

What critical operational lesson does this case illustrate about model behavior when faced with unavailable inputs?

- A) When critical inputs sit outside the context window, a model does not stop or refuse the request; it fills the information gap with typical industry defaults at full confidence.
- B) The model experienced under-specification because the prompt did not specify Python pandas syntax.
- C) Capacity planning is inherently an expectation mismatch that AI can never support.
- D) The operations lead should have run the prompt on a speed tier model.

Question 44 [SELECT ALL THAT APPLY]

Scenario: A software company's customer complaint delegation map contains the following steps:

1. *Pull complaint and ticket history* → AI (Mechanical retrieval)
2. *Summarize into issue template* → AI (Stated reason: "Our summarizing Skill is very good at this")
3. *Classify severity (low, medium, high)* → AI (Stated reason: "The summaries have been accurate for three months")
4. *Draft response* → Collaborative (AI drafts, duty agent reviews before dispatch)
5. *Send response* → AI (Stated reason: "Low severity only, and drafts have been strong")
6. *Log outcome in CRM* → AI (Mechanical write)

Which of the following diagnoses correctly identify the structural flaws in this delegation map? (Select all that apply.)

- A) Step 3 is halo delegation: summarizing competence is real but belongs to Step 2; severity classification is a distinct decision that determines escalation or neglect.
- B) Step 5 is over-delegation: severity was assigned in Step 3 by the same automated system, allowing AI to grant itself permission to skip the human review gate on an irreversible send.
- C) Step 2 is mapping the tool instead of the work: the step exists because the team built a good summarizing Skill, not because the business workflow inherently required that structure.
- D) Step 4 is an unstaffed gate waiting to happen: having the duty agent review every draft without establishing shift capacity limits will cause review to decay into a perfunctory click.
- E) Step 1 is invalid because mechanical retrieval must always be human-retained.

Question 45 [SELECT ALL THAT APPLY]

Scenario: A newly appointed map owner conducts a 20-minute quarterly scheduled read of a team's customer refund delegation map.

Which of the following represent the four specific questions the owner must evaluate during the scheduled read? (Select all that apply.)

- A) Has a step been added or removed in practice that is not reflected on the map?
- B) For every collaborative step: who reviewed it last week, and how long did they actually spend?
- C) Has anything changed the volume or the stakes of the workflow?
- D) Is every configuration the map depends on (Skills, Project knowledge, standing instructions) still current?
- E) Which individual programmer originally authored the underlying LLM weights?

Question 46 [SINGLE-ANSWER]

Scenario: When an AI output disappoints, a common behavioral instinct among team members is to immediately jump to the most capable (and expensive) reasoning model tier or abandon the initiative entirely by declaring the task impossible for AI.

How does the diagnostic sequence characterize these two common reactions?

- A) They are sophisticated engineering moves that eliminate operational friction.
- B) They are moves from the very bottom of the ladder, executed before cheaper, more common root causes (prompt under-specification, context length, feature choice, stale configuration) have been investigated.
- C) They are mandatory compliance procedures under ISO-27001 AI risk management.
- D) They reflect appropriate escalation signals under Concept 10.

Question 47 [SELECT ALL THAT APPLY]

Scenario: A business analyst trains end-users on how to convert subjective complaints about AI output into disciplined configuration changes.

According to Concept 13 and Slide 21, which of the following pairings correctly link the user's reaction to the actionable instruction and its governing lever? (Select all that apply.)

- A) *"Too generic"* → Name the specific audience and the single action they must take → Lever: The prompt
- B) *"Wrong tone"* → Add an enduring tone constraint that applies across all drafts → Lever: The standing instruction
- C) *"Missed the point"* → State the single question the output must answer up front → Lever: The prompt
- D) *"It's using old data"* → Replace the outdated source document in the Project knowledge base → Lever: The knowledge
- E) *"Numbers are wrong"* → Write an emotional plea in the prompt demanding accuracy → Lever: The prompt

Question 48 [SINGLE-ANSWER]

Scenario: A financial reporting team audits an analyst workflow that currently takes 45 minutes per run with high output variance. The audit reveals three distinct friction points:

1. Analysts repeatedly re-paste the same company background context and accounting definitions;
2. Analysts reformat raw tables by hand following a fixed 5-step styling order; and
3. Different analysts check for different accounting anomalies based on personal habits.

Applying the Concept 14/15 framework, how should these three friction points be configured to cut cycle time down to 25 minutes?

- A) (1) Standing instruction, (2) Knowledge base, (3) Skill
- B) (1) Shared Project Knowledge Base (Reference), (2) Skill (Procedure), (3) Standing Instruction (Rule)
- C) (1) Code execution, (2) Prompt few-shots, (3) Review gate
- D) (1) Model Memory, (2) Artifact, (3) Parallel run

Question 49 [SINGLE-ANSWER]

Scenario: Before implementing an AI-assisted compliance auditing redesign, the project lead is instructed to execute one single, completely unchanged run of the legacy manual workflow to establish an empirical baseline.

Which three specific numbers must be recorded during this unchanged run?

- A) Total tokens consumed, API server latency, and cloud infrastructure cost.
- B) Time per run, rounds of revision required, and number of manual steps performed.
- C) Employee heart rate, keystrokes per minute, and document word count.
- D) Number of stakeholders attending the kick-off meeting, budget variance, and Net Promoter Score.

Question 50 [SINGLE-ANSWER]

Scenario: A B2B consulting firm uses AI to assist in drafting high-stakes regulatory compliance audits for enterprise clients. The client's Chief Risk Officer inquires about the firm's AI usage and demands to understand what safeguards protect against AI errors.

Which description adheres strictly to the framing principles of Workflow Design & Diagnosis to maximize trust with the client's risk function?

- A) *"Our workflow is 100% automated using proprietary deep learning agents that guarantee error-free compliance audits."*
- B) *"Our AI models are essentially as competent as human compliance auditors, so errors are exceptionally rare."*
- C) *"Drafting is AI-assisted. A qualified professional reviews and approves every finding, and nothing is released to stakeholders without that formal sign-off."*
- D) *"We run an iterative build loop inside a Project, which eliminates hallucination risk through code execution."*

Answer Key

The correct answers for all 50 questions are listed below.

Q1: B	Q2: A, B, C, D, E	Q3: C	Q4: B	Q5: A, B, C
Q6: B	Q7: B	Q8: B	Q9: C	Q10: A, B, C, E
Q11: C	Q12: A, B, C, E	Q13: B	Q14: B	Q15: A
Q16: B	Q17: D	Q18: A, B, C, D	Q19: C	Q20: B
Q21: A, B, C, D, E	Q22: A	Q23: B	Q24: C	Q25: A, B, C, E
Q26: C	Q27: B	Q28: B	Q29: A, B, C, D	Q30: B
Q31: A, B, C, E	Q32: A, B, C, D	Q33: B	Q34: B	Q35: A, B, C
Q36: B	Q37: B	Q38: B	Q39: B	Q40: B
Q41: B	Q42: B	Q43: A	Q44: A, B, C, D	Q45: A, B, C, D
Q46: B	Q47: A, B, C, D	Q48: B	Q49: B	Q50: C