

Official Examination Question Bank: Code You Never Write (Q1–Q50)

Comprehensive Assessment Protocol · Non-Technical Client Craft & Verification Doctrine

Question 1

SINGLE-ANSWER

A logistics coordinator at a regional depot pastes a 45-row delivery log containing timestamps and fuel consumption into an AI chat window, submitting the following prompt:

"Look over this week's route logs and give me a quick, rough sense of which drivers consistently missed their delivery windows, and highlight any routes that look problematic."

Within seconds, the AI outputs an organized bulleted summary listing five driver IDs, exact percentage delays, and an authoritative concluding assessment. No code block or execution indicator appears in the response.

Based strictly on the principles of the chapter, what fundamental failure has occurred, and which concept explains why this output cannot be trusted for operational decisions?

- A. The coordinator failed to specify Python as the required execution language, causing the AI to default to an unverified text-parsing heuristic.
- B. The coordinator fell into the conversational estimation trap by using glance-inviting phrasing on a task that tripped the Precision and Files signals, triggering silent estimation formatted identically to computation.
- C. The coordinator violated the blast radius protocol by uploading operational data to a cloud sandbox without first establishing a local backup copy and requesting a dry run.
- D. The task is an Answer Problem requiring subjective human judgment; therefore, the AI correctly avoided code execution and reasoned from first principles.

Question 2

SINGLE-ANSWER

A university administrator is preparing to reconcile 500 student registration records against bank payment receipts. Drawing on advice from a developer friend, she begins her prompt with:

"Write and run Python code to analyze these two files and identify every student who has not cleared their balance."

While the AI executes code and outputs a table, the Chief Examiner notes that explicitly dictating "Python" violates a key tenet of the client-side craft taught in the chapter.

Why does the chapter strictly insist that a client command the AI to "write and run code", rather than specifying "Python"?

- A. Browser-based execution sandboxes restrict external libraries, so naming Python triggers environment dependency errors that the AI cannot self-repair.
- B. Technical language selection is an architect-level judgment that must follow the intended output; dictating the language forces the AI's hand on a technical guess the client is not equipped to make and risks steering it wrong when another language is better suited.
- C. Specifying Python prompts the AI to output raw code snippets into the chat window for human review rather than executing the code automatically in the sandbox.
- D. Terminal and desktop agent surfaces only support shell scripts; dictating Python overrides the agent's automated plan-then-approve mechanism.

Question 3

SELECT ALL THAT APPLY

An operations manager at a retail chain is structuring a five-section Markdown brief (Goal, Input, Output, Rules, Edge cases) to analyze multi-branch transaction files. The chapter highlights that beginners frequently omit the two starred sections—*Rules* and *Edge cases*—which are responsible for preventing nearly every incorrect analysis.

To maintain clean separation between domain logic and defensive data handling, which of the following statements belong strictly under the **Rules** section of the brief, rather than the **Edge cases** section?

- A. The fiscal business day closes at 04:00 AM rather than midnight to account for late-night store hours.
- B. If an incoming transaction row contains a corrupted, unparseable timestamp, skip the line without guessing and record its row number at the end of the summary.

Question 4

SINGLE-ANSWER

A clinical researcher intends to deploy an AI-generated script to audit 5,000 patient dosage records against pharmacy inventory logs. Before running the script on the live dataset, she introduces a single deliberate discrepancy of 15 units into a sample copy of the data.

When the script executes, it flags her planted discrepancy along with six unexpected real discrepancies.

Based on the chapter's discussion of verification, mutation testing, and the case of the pharmacist who salted her own data (Concept 06), what fundamental testing principle does this maneuver validate?

- A. It establishes mathematical proof that the AI's underlying neural weights did not hallucinate the syntax of the data-parsing libraries.
- B. It proves the sensitivity of the audit; if a testing mechanism stays silent when presented with a known defect, it proves that the test proves nothing.
- C. It safely verifies file system write permissions and OS-level execution locks, allowing the researcher to bypass the dry-run requirement.
- D. It provides empirical evidence that the script is executing inside a sandboxed environment rather than operating directly on the local machine.

Question 5

SINGLE-ANSWER

A marketing lead runs an AI-generated script on an annual campaign spreadsheet containing multiple quarterly tabs. The script finishes execution with a green *"ran successfully"* indicator—producing zero red error text—and reports an annual customer acquisition total of 14,200. The lead knows from audited board presentations that the true figure was approximately 55,000.

Twice, the lead prompted: *"This is wrong, fix it."* Both times, the AI apologized, tweaked a minor date filter, and returned slightly different figures that remained around 15,000.

According to the debugging protocols outlined in Concept 07 (*When it breaks*), what disciplined diagnostic response should the lead execute next?

- A. Re-prompt with the exact phrase *"write in Python without third-party libraries"* to eliminate multi-sheet workbook driver conflicts.
- B. Recognize the *wrong-but-running* failure mode and the onset of *context rot*; stop digging, state the specific symptom alongside the expected value, and instruct the AI to inspect row counts, active sheets, and sample rows (or pivot to two completely fresh approaches in a clean conversation) before writing further code.
- C. Classify the spreadsheet as an *Answer Problem*, abandon programmatic execution, and instruct the AI to reason through the discrepancy conversationally.
- D. Conclude that browser-based sandboxes cannot handle multi-tab files and immediately migrate the project to OpenCode in the terminal.

Question 6

SINGLE-ANSWER

A systems administrator at a private hospital needs to oversee 35 diagnostic workstations connected to an internal clinical network. Every weekday morning at 06:00 AM—two hours before medical staff arrive—an automated routine must query each workstation across the local subnet, extract disk storage telemetry, identify drives with less than 10% capacity, and write a consolidated report directly into a shared network drive.

The administrator is currently attempting to run this workflow by manually uploading log files to Claude.ai each morning.

According to the surface taxonomy and graduation criteria detailed in Part Three (Concepts 09–11), which surface should the administrator transition to, and why?

- A. Remain on Claude.ai in the browser, because the isolated cloud sandbox provides a zero-risk perimeter that guarantees hospital patient data cannot be affected by file-system deletion bugs.
- B. Transition to Cowork or OpenWork, because their visual *plan-then-approve* interface is required to execute unattended cron schedules while the computer is locked.
- C. Transition to Claude Code or OpenCode, because the operational requirements demand direct local system access, communication with network devices beyond a browser tab's reach, and unattended execution on a schedule.
- D. Transition to Gemini or ChatGPT Plus, because modern browser-based sandboxes automatically open secure local SSH sockets when network subnet addresses are specified in the prompt.

Question 7

SELECT ALL THAT APPLY

A social sciences researcher is preparing to use Claude Code in the terminal to reorganize, rename, and deduplicate 400 audio recordings (.wav and .mp3 files) from a multi-year fieldwork study.

The chapter notes that when code runs directly on local storage, two harsh technical realities apply: *deleted files often skip the recycle bin*, and *edited files keep no undo history*.

Which of the following operational precautions directly embody the four safety rules and principles of **Blast Radius** management defined in Concept 12?

- A. Restrict the agent's working context exclusively to the specific, isolated project directory rather than initiating the agent at the user's home directory or drive root.
- B. Instruct the agent to bypass dry runs as long as it commits to sending deleted files to the operating system's standard Recycle Bin / Trash folder for one-click recovery.
- C. Demand a full dry-run plan listing every proposed file operation (showing old filenames alongside proposed new names and destinations) and require human approval before any file is touched.
- D. Instruct the agent to write all reorganized outputs and indexes to distinct new files or subfolders, leaving the raw source files completely untouched.
- E. Duplicate the entire folder to an independent backup location before running any script that has not yet established an audited track record of reliability.

Question 8

SINGLE-ANSWER

An executive steering committee at a commercial bank is reviewing four proposed automation initiatives. The Head of Operations cites the "**Edge of the Map**" principles from Concept 13 to explain which projects can be safely tackled using one-prompt client-commissioned scripts, versus those that cross the boundary into multi-user product engineering, unattended agency, or human judgment.

Which of the following initiatives lies legitimately **INSIDE** the operational territory of client-commissioned scripts described in the chapter?

- A. A customer-facing portal that allows retail depositors to log into accounts, initiate real-time international wire transfers, and maintain continuous high availability.
- B. An unattended background daemon (*Digital FTE*) that monitors live communications and autonomously issues legally binding loan default notices to borrowers without human verification.
- C. A monthly workflow that reconciles raw internal transaction ledgers against external clearinghouse CSV files, highlights date/amount mismatches, and outputs a visual HTML audit report for an accountant to review.
- D. An automated HR evaluation tool that analyzes employee ticket resolution times and independently decides which staff members should be promoted or placed on performance improvement plans.

Question 9

SINGLE-ANSWER

An admissions coordinator at a business school uses an AI session to build a custom validation script that reconciles applicant prerequisite transcripts against institutional exemption criteria. Six months later, the faculty senate approves a revised grading scale for transfer credits.

Under the asset-preservation methodology detailed in Concept 08 (*Keep the script: a solved problem becomes a button you press forever*), how should the coordinator have structured this automation asset to enable seamless updates without re-engineering the task from scratch?

- A. Delete the legacy Python script and initiate a fresh conversational chat, re-explaining the admission logic from memory to avoid polluting the new prompt with deprecated syntax.
- B. Maintain a dedicated task directory (`my-scripts/transcript-check/`) containing the executable script (`check.py`), the original Markdown brief (`brief.md`) serving as the *Intent Layer*, and a known-answer sample dataset (`sample-transcripts/`); to update the logic, supply both the brief and script to an AI and state the rule change in plain English.
- C. Compile the Python script into an uneditable binary executable, ensuring that administrative staff cannot tamper with the institutional evaluation logic.
- D. Leave the script unarchived within the browser chat history, relying on the model's server-side conversational memory to recall the original rules when the new term begins.

Question 10

SINGLE-ANSWER

An event coordinator is evaluating whether to commission a custom AI script or subscribe to commercial off-the-shelf SaaS software to manage several operational challenges for an upcoming community festival.

Under the chapter's **"Off-the-Shelf-App Test"** and the concept of a **"Market of One"**, which of the following scenarios represents the clearest, most legitimate candidate for commissioning custom AI code?

- A. Dividing an equal dinner bill evenly among six colleagues and distributing automated payment collection links.
- B. Calculating statutory provincial and federal sales tax on standard commercial retail invoices using published government tax schedules.
- C. Reconciling festival ticket dues against informal mobile-wallet receipts where memos contain unstandardized free text (e.g., separating personal meal repayments from ticket dues), certain lump-sum transfers cover multiple family members under unwritten arrangements, and the baseline attendance total is known only because the coordinator counted badges at the gate.
- D. Synchronizing client availability across Google Calendar to automatically book standard 30-minute introductory consultation slots.

Question 11

SINGLE-ANSWER

During an operational review, a project manager describes their automation process as follows:

"Whenever I need to process messy data, I prompt an AI model to generate a Python script, copy the provided code block into a local text editor, ask an in-house developer to troubleshoot any syntax or module import errors, and execute the script in a local development environment."

Based on Concept 01 (*Code is no longer gated by coding*), why does the chapter consider this workflow an obsolete artifact of the *"2023 era"*, and what structural transformation characterizes the modern client-developer relationship?

- A. Modern LLMs now synthesize binary machine instructions directly within their neural weights, rendering high-level languages like Python and JavaScript obsolete for business tasks.
- B. Three capabilities matured concurrently—the AI generates working code for small scoped tasks, executes the code autonomously within an integrated runtime seeing the output, and feeds execution errors back into its own loop to self-repair, completely eliminating human code copying, environment configuration, and syntax management.
- C. Enterprise data governance frameworks now legally mandate that all data transformations occur strictly through conversational estimation rather than programmatic computation.
- D. Modern client-side development requires non-programmers to write formal pseudocode specifications rather than plain-English descriptions before code can be generated.

Question 12

SINGLE-ANSWER

An accountant uploads a 1,420-row file named `q3_statement.csv` with specific headers: `Txn_ID`, `Posting_Date`, `Vendor_Entity`, and `Cleared_Amount`. The prompt submitted is:

"Analyze this bank statement and tell me if there are any duplicate charges exceeding \$500."

The AI immediately responds with an articulate, professional summary:

"I have thoroughly examined your statement of approximately 1,500 transactions across the Date, Vendor, and Amount columns. I identified 14 duplicate charges over \$500, totaling \$9,250."

No code block, script trace, or execution badge accompanies the reply.

Applying the diagnostic safeguards of Concept 04 (*Make the AI's hands move*), what specific evidence proves conclusively that the AI fabricated this analysis, and which protective measure was omitted?

- A. The total sum of \$9,250 contradicts double-entry accounting rules, proving that the cloud sandbox suffered an unhandled Python memory overflow.
- B. The absence of a code block combined with a rounded row count (*"approximately 1,500"*) and generic, hallucinated column names (*"Date, Vendor, Amount"*) fails the *cheapest lie detector* test, confirming the AI produced a silent, glance-based estimate because the prompt omitted explicit instructions to write and execute code.
- C. The presence of commas in row counts indicates that the underlying CSV parser encountered a fatal codec error, requiring an immediate plain-English replay to debug.
- D. Financial statements inherently trip the Repetition signal, meaning browser chats reject execution unless the

Question 13

SINGLE-ANSWER

A business analyst who recently completed a short introductory Python coding course is tasked with reconciling two messy customer export files. In their prompt to the AI, the analyst specifies:

"Import pandas as pd, loop through the rows using an iterrows() generator, write a regex pattern with re.sub() to clean phone numbers, create a dictionary of sets, and export the output using to_csv(index=False)."

Despite the granular instructions, the script silently drops 85 international customer records containing leading plus signs and misclassifies names with trailing whitespace.

According to Concept 05 (*Brief the problem, not the program*), why did this approach fail, and what disciplined sequence should have been used instead?

- A. The prompt specified pandas rather than vanilla Python lists, which causes cloud execution sandboxes to suppress regex error logs.
- B. The analyst fell into the trap of people who half-know programming: micromanaging the technical "how" while under-specifying the domain "what"; the correct workflow is to write briefs containing zero technical jargon, describe the desired outcome, and execute a discovery prompt first (*"examine the files and tell me what could be ambiguous, inconsistent, or surprising"*) before defining the rules and edge cases.
- C. Local data cleaning tasks require terminal agents like Claude Code, as browser-based models cannot interpret Python regular expressions without external browser extensions.
- D. The prompt failed because phone numbers trip the Precision signal, requiring the user to hand-code the cleaning script rather than delegating it to an AI developer.

Question 14

SINGLE-ANSWER

A corporate payroll officer is finalizing year-end performance bonuses totaling \$2.4M across 800 staff members. Recognizing the severe financial stakes, the officer applies the highest tier of the Verification Ladder—a **Cross-Model Check** (Level 5)—by submitting the identical Markdown brief and employee spreadsheet into both Claude.ai and ChatGPT.

Both independent models generate separate code, execute in their respective sandboxes, and return identical results: exactly 760 employees qualify, totaling \$2,310,000. Confident in this cross-provider consensus, the officer submits the disbursement schedule to the bank without further checks.

Three days later, leadership discovers that 40 eligible employees from a newly acquired regional office were omitted from the bonus pool because their records used a non-standard branch code.

According to Concept 06 (*Verifying work you cannot read*), what structural limitation of the Cross-Model Check explains why two competing AI families agreed on a flawed outcome?

- A. Competing AI providers share upstream Python container kernels, meaning low-level data-frame indexing bugs replicate across both platforms.
- B. Because both AI models operate strictly from the client's brief, any unstated institutional rule, omitted domain constraint, or unhandled data edge case is a shared blind spot that both models will execute identically and agree upon.
- C. A cross-model check is mathematically invalid unless one model runs Python and the other is forced to execute JavaScript in a browser tab.
- D. Cross-model verification only functions when conducted across local terminal agents like OpenCode and Claude Code, because browser chats lack cross-instance entropy.

Question 15

SINGLE-ANSWER

An academic department head is preparing to organize and re-index ten years of departmental evaluation files, course syllabi, and faculty portfolios stored locally across multiple subdirectories. The department head does not want to use a command-line terminal and fears that an automated script might misinterpret folder hierarchies and delete critical tenure files.

When evaluating desktop agent platforms (Claude Cowork and OpenWork) against both browser-based chat sandboxes (Claude.ai) and terminal agents (Claude Code / OpenCode), what core architectural mechanism distinguishes Cowork and OpenWork, and how does it redefine the verification process?

- A. They eliminate underlying programmatic execution entirely, replacing Python scripts with computer-vision GUI automation that mimics manual mouse clicks.

Question 16

SELECT ALL THAT APPLY

A chief operating officer is mentoring team leads on applying the **VPRF framework** (Concept 03: *The dividing line*) to determine whether an operational task is an **Answer Problem** (a question for the AI's mind) or a **Code Problem** (a job for the AI's hands). The rule states: *Trip none of the four signals, it is an Answer Problem; trip even one, it is a Code Problem.*

Which of the following business tasks are strictly categorized as **Code Problems** under the VPRF framework?

- A. "Draft a diplomatic email declining an invitation to keynote an industry conference due to scheduling conflicts."
- B. "Examine 120 vendor service-level agreements and identify every contract containing a termination notice period of fewer than 30 days or specifying an unfamiliar dispute jurisdiction."
- C. "Critique our product pitch deck's value proposition from the perspective of an early-stage venture capital investor and recommend narrative improvements."
- D. "Every Friday afternoon, extract line items from three regional branch sales CSV files, reconcile refund debits, and merge them into a consolidated weekly revenue table."
- E. "Inspect our company's annual expense ledger export and determine our exact total expenditure on cloud hosting infrastructure."

Question 17

SINGLE-ANSWER

An administrative associate uses an AI-generated script to re-sort and archive 600 vendor invoice PDFs into quarterly folders based on the invoice dates. Because the prompt did not specify how to interpret ambiguous numerical dates (04/05/2025 being treated as April 5th instead of May 4th), the script swiftly routed 220 invoices into incorrect fiscal quarter folders within three seconds.

The associate complains to the team: *"The AI's code suffered a major glitch and broke down."*

Applying the kitchen recipe analogy and the core property of code articulated in Concept 02 (*What code actually is*), what is the accurate analytical diagnosis of this outcome?

- A. The script suffered an asynchronous thread collision because the AI mistakenly chose JavaScript instead of Python for file-system manipulation.
- B. The code did not fail; code possesses the fundamental property of being exact in both directions—it executes instructions flawlessly at scale, meaning a brief containing an unstated assumption produces an exact, catastrophic error just as reliably and rapidly as it produces the right answer.
- C. The execution environment encountered *context rot*, which occurs whenever a batch exceeding 500 files is processed in a single run.
- D. The cloud sandbox suppressed an underlying syntax error, which converts file-renaming procedures into random shuffling operations.

Question 18

SINGLE-ANSWER

A media coordinator at an advertising agency intends to clean up an unorganized 60GB directory containing thousands of high-resolution photoshoot images. The coordinator asks: *"Can the AI visually inspect every image, judge which shots are redundant, and automatically purge the duplicates to free up disk space?"*

Based on the mechanics and safety protocols established in **Project 4 (Photo Gallery Rescue)**, how does the AI actually execute image deduplication, and why is human review of the pre-deletion plan indispensable?

- A. The AI uploads the full-resolution photos to a multimodal vision model to evaluate framing and lighting, and automatically relocates low-scoring duplicates to the OS trash folder.
- B. The AI never visually "looks" at the photos or exercises artistic judgment; it generates standard code that condenses each image into a mathematical fingerprint (such as a perceptual hash) to compare matches, but because two different moments (such as similar landscape shots) can yield identical hashes, a human must audit a dry-run manifest before deletion to prevent losing distinct images.
- C. Image deduplication is fundamentally an *Answer Problem* that cannot be executed via Python, requiring the coordinator to manually inspect all images in a photo-viewing application.
- D. The AI relies strictly on operating system file-creation timestamps; any photos captured within three seconds of each other are automatically overwritten by new files under Rule 4 of the Blast Radius framework.

Question 19

SINGLE-ANSWER

A junior software engineer criticizes a non-technical department manager who commissions data scripts via AI without ever reading or auditing the generated Python code, arguing:

"Not reading the code is a dangerous beginner's shortcut. Professional engineering demands line-by-line code review. Skipping syntax review is fundamentally reckless."

How does the chapter refute this skepticism using the prominent case of Robert C. Martin (*Uncle Bob*, author of *Clean Code*), and what vital operational principle governs this stance?

- A. It concedes that avoiding code review is an amateur compromise, but justifies it by claiming that non-developer business data never carries significant financial or legal stakes.
- B. It explains that not reading code is the very frontier where the most seasoned practitioners have converged, provided the foundational sequence is respected: one does not abandon reading code and simply hope for the best; rather, one surrounds the AI with rigorous verification checks first (such as known-answer tests and mutation testing), and those independent checks are precisely what *earn and permit* the freedom of not reading the code.
- C. It argues that reading code causes neural network interference, and that modern sandboxes enforce cryptographic obfuscation on all generated scripts to block human inspection.
- D. It demonstrates that Python syntax is mathematically identical to English prose, making line-by-line syntax auditing redundant for anyone who can read English.

Question 20

SELECT ALL THAT APPLY

A clinic administrator is piloting an AI-assisted reconciliation between daily patient attendance registers and specialist fee billing exports. During the initiative, three distinct challenges arise: the files contain confidential patient identifiers, the script crashes abruptly with a block of red text referencing an unrecognized codec error, and a subsequent run completes without errors but reports zero visits for two major clinical departments known to be active.

According to the diagnostic workflows in Concept 07 (*When it breaks*) and the project execution guidelines, which of the following actions represent correct, disciplined client responses?

- A. For the red-text codec crash, avoid attempting to research or comprehend text encoding standards; copy the entire red error block verbatim and prompt: *"It stopped and showed this error. Diagnose it, fix the code, and run it again."*
- B. To safeguard sensitive records, create a sanitized duplicate dataset with patient names and national identification numbers stripped out before uploading, or instruct the AI to generate an anonymized copy prior to computation.
- C. For the silent omission of active clinic departments (*wrong-but-running*), avoid generic prompts like *"This is wrong, fix it"*; instead, state the specific symptom alongside the expected departmental figures, and require the AI to inspect row counts, active categories, and the first five parsed rows before modifying code.
- D. Whenever an unhandled Python exception occurs, immediately delete the conversation history and start over from memory, because cloud sandbox environments permanently corrupt their local execution states upon any runtime crash.
- E. If an identical fix fails three times in succession, stop digging; instruct the AI to step back, restate the core problem fresh, and propose two entirely different approaches (or initiate a clean chat carrying over only the brief and key lessons learned).

Question 21

SINGLE-ANSWER

A university student undertakes **Project 1 (Money Detective)** to analyze four months of personal mobile-wallet exports. The student prompts the AI to identify recurring monthly charges, same-day duplicate transactions, and overall spending leaks.

The AI outputs a report stating that food delivery represents the largest cumulative expense, flags a streaming subscription billed four times (\$796 total), and identifies an identical duplicate debit on July 8th.

Under the specific execution methodology of Project 1, what required procedure must the student execute to *"earn trust backstage first"*, and what defines a legitimate *"catch"*?

- A. The student must convert the dataset into an interactive web game using JavaScript to ensure cross-platform arithmetic consistency across devices.
- B. The student must verify the output against at least two simple figures calculated independently by hand (such as verifying total income deposits and the exact four-month total of a known subscription); once the script

Question 22

SINGLE-ANSWER

An undergraduate student is preparing for semester finals. The professor's course syllabus contains a unique grading policy: the two lowest homework assignments are dropped, only the top 8 of 10 pop quizzes are averaged, and if the score on the cumulative final exam exceeds the midterm exam, the final exam score replaces the midterm score completely.

The student attempts to track their progress using a commercial mobile grade-tracking app, but the app cannot accommodate these stipulations.

Based on the core lessons of **Project 2 (What's My Grade, Really)** and Concept 05 (*Brief the problem, not the program*), why does commercial software fail in this scenario, how is the brief structured, and how does the student establish initial trust in the output?

- A. Commercial GPA apps require direct OAuth integration with university registrar databases; the student must write custom SQL commands in an agent terminal to override the app's internal schema.
- B. Commercial applications exist only when rules are standardized across an entire market, whereas this bespoke syllabus represents a "market of one" whose rules exist nowhere else in the world; by translating the syllabus criteria into plain English within the brief's **Rules** section and hand-checking a single component (such as the dropped-homework average) before reviewing the whole, the student establishes backstage trust in the computed standing and target score.
- C. Grading calculations represent an *Answer Problem* that demands academic intuition rather than numerical computation, meaning programmatic approaches inevitably trigger *context rot*.
- D. The student should have instructed the AI to execute the script in JavaScript, because browser-based Python interpreters cannot execute conditional grade-replacement logic.

Question 23

SINGLE-ANSWER

A neighborhood committee treasurer collects monthly building maintenance dues of PKR 2,000 from 20 resident families, expecting a known collection baseline of PKR 40,000. When auditing the incoming mobile-wallet export (JazzCash/Easypaisa), the statement contains messy user memos: one transaction of PKR 4,000 has the memo "Tariq + brother Asad", a PKR 500 entry notes "cricket match share", and a PKR 2,000 payment arrives from an unregistered mobile number.

If the treasurer simply instructs the AI to "reconcile this digital transaction file against the roster of 20 residents", what failure occurs, and what core principle from **Project 3 (The Books Don't Match)** provides the necessary bridge to achieve accurate reconciliation?

- A. The script will crash with a fatal runtime codec error, because digital banking exports require binary cryptographic decryption prior to sandbox ingestion.
- B. C. Local financial reconciliations can only be processed using Claude Code in a terminal, because browser-based sandboxes automatically delete unmapped transaction entries under Blast Radius Rule 3.
- D. Financial reconciliation is an *Answer Problem* that prohibits deterministic calculation, requiring the treasurer to let the AI estimate missing payments using qualitative judgment.

Question 24

SELECT ALL THAT APPLY

A data research supervisor is preparing to transition an analytics team from browser-based chat sandboxes (such as Claude.ai) to local folder-based agents (**Claude Code** and **OpenCode**). The supervisor needs to brief the non-technical staff on the precise operational transformations that occur when moving to a terminal agent surface.

According to Concept 10 (*Claude Code & OpenCode: the agent in your folder*), which of the following statements correctly characterize the technical and structural differences between terminal-based agents and browser chat interfaces?

- A. The agent has direct access to files already residing in the target folder, eliminating the friction of manual multi-file browser uploads ("40 files means 40 uploads").
- B. Terminal agents require non-technical users to master command-line shell syntax and memory-pointer allocation before the agent can interpret instructions.
- C. Generated automation scripts and output deliverables persist permanently on the local machine rather than expiring when the browser session or cloud sandbox terminates.
- D. The agent executes in a tight self-healing loop on the local machine, capturing runtime errors directly and iteratively repairing its own code without requiring the user to manually copy and paste red error traces.

Question 25

SINGLE-ANSWER

A warehouse operations lead reviews an AI-generated inventory reconciliation report derived from a 2,500-row stock ledger. Upon scanning the generated summary, the lead identifies two immediate anomalies:

1. The raw intake file contained 2,500 line items, but the processed report accounts for only 2,412 items.
2. The report lists a discontinued legacy widget as the highest-turnover SKU for the month, directly contradicting the lead's direct knowledge that a newly released product line led warehouse shipments.

Under the Verification Ladder established in Concept 06 (*Verifying work you cannot read*), which level of verification is the lead practicing, and what is the prescribed diagnostic follow-up prompt?

- A. The lead is executing an *Adversarial Pass* (Level 4); the required follow-up is: *"Rate your confidence 1–10 and explain why Python failed to index the legacy SKU."*
- B. The lead is applying *Reality Questions* (Level 2)—checking rows in versus rows out and verifying known reference items against professional ground truth; upon detecting figures that contradict domain reality, the required follow-up prompt is: *"Your report says X. Walk me through in plain English how you computed that number, and show me three of the rows behind it."*
- C. The lead is executing a *Known-Answer Test* (Level 1); the required follow-up is to inject three corrupted rows into the database and rerun the script.
- D. The lead has encountered an unrecoverable sandbox memory leak; the required response is to abandon automated computation and process the 2,500 rows manually in spreadsheet software.

Question 26

SINGLE-ANSWER

An audio archivist is preparing to use an AI agent to clean up duplicate recordings across multiple client folders. Remembering that agentic deletions typically bypass the operating system's recycle bin and leave no undo history, the archivist reviews the safety differences between using a command-line terminal agent (such as Claude Code) versus a desktop delegation app (such as Cowork or OpenWork).

According to Concept 11 (*Cowork & OpenWork*) and Concept 12 (*Blast radius: Rule 2*), what fundamental operational distinction exists between these two surfaces regarding dry runs, and what must the terminal user explicitly enforce?

- A. Desktop apps like Cowork and OpenWork automatically enforce a *plan-then-approve* rhythm into their interface by default, whereas in a command-line terminal agent, the user must explicitly command and enforce a dry run (*"Don't change anything yet. Show me every operation you would perform and wait for approval"*) to inspect the proposed manifest while it is still harmless words rather than irreversible actions on disk.
- B. Terminal agents automatically route deleted files to a temporary cloud sandbox for 30 days of safe retrieval, whereas desktop apps permanently shred deleted files immediately.
- C. Command-line terminal agents refuse to execute any file deletions unless the user manually approves the binary bytecode in a hex editor.
- D. Desktop apps only support read-only operations on CSV files, meaning destructive file deletions can only occur within terminal environments.

Question 27

SINGLE-ANSWER

A newly hired records officer at an educational institute is preparing to use Claude Code in the terminal to parse, validate, and compile graduation eligibility lists from raw departmental student CSVs. The officer asks where on their machine they should open the terminal to run the agent, and how the script should handle output generation. Following **Rule 3 (Scope the territory)** and **Rule 4 (Output to new files)** of the Blast Radius framework (Concept 12), what mandatory instruction set must the officer follow?

- A. Open the terminal at the root drive level (`/` or `C:\`) or the top-level user home directory so the agent can scan all system paths, and instruct the script to overwrite the raw departmental CSVs in-place to minimize disk usage.
- B. Point the agent strictly at the smallest, isolated subfolder that directly contains the target task (e.g., `graduation-2026/`)—never the home directory, Documents, or the whole drive—and mandate in the brief that all results be written to a newly named file while leaving the source files completely untouched, ensuring every run is reversible by deletion.
- C. Launch the terminal within the operating system's root system folder, but instruct the agent to append output records directly into the computer's BIOS log.
- D. Execute the script exclusively from an external read-only disk drive, because modern AI agents lack read

Question 28

SELECT ALL THAT APPLY

A chief information officer is drafting an institutional AI governance policy based on Concept 13 (*The edge of the map, and what's across it*). The policy must clearly delineate between tasks that sit safely *inside* the boundaries of one-prompt client-commissioned scripts versus systems that lie **Across the Edge**, requiring full software engineering, unattended agent architectures (Digital FTEs), or human governance.

Which of the following proposed initiatives fall strictly **ACROSS THE EDGE** of one-prompt client scripting?

- A. A multi-user customer portal where external clients create accounts, authenticate via passwords, execute credit card billing transactions, and rely on 24/7 service uptime.
- B. A weekly script executed by an operations assistant that merges three regional sales spreadsheets, highlights duplicate customer entries, and outputs an HTML summary chart for team review.
- C. An unattended background daemon (*Digital FTE*) that monitors hospital supply inventory and autonomously issues binding purchase orders and financial disbursements to vendors at 3:00 AM without human oversight.
- D. An automated HR decision engine designed to finalize employee promotions and contract terminations by executing mathematical formulas on performance ratings without human managerial review (*judgment wearing a computation costume*).
- E. A client-commissioned Python script that aggregates annual financial data and drafts an initial corporate tax filing package, while mandating that a certified accountant manually audit the claims and execute the final submission button.

Question 29

SINGLE-ANSWER

A senior operations analyst at an e-commerce firm uses an AI chat tool in March to successfully reconcile warehouse inventory valuation against sales ledgers. In April and May, rather than saving the generated Python script to a local folder, the analyst simply opens a fresh chat session and prompts from memory:

"Here are this month's warehouse exports. Calculate our ending inventory valuation using the same valuation logic we used last month."

During a mid-year financial audit, external accountants discover that the May report computed cost-of-goods-sold using a slightly different weighted-average formula than the March report, creating an unverified variance of \$85,000.

According to Concept 08 (*Keep the script: a solved problem becomes a button you press forever*), what exact operational breakdown occurred, and how does the disciplined *Keep-the-Script* habit prevent it?

- A. Cloud-based AI sandboxes experience silent memory corruption across multiple months, requiring analysts to reboot their router before prompting.
- B. Failing to preserve the script caused **subtle rule drift**—re-prompting the AI from memory each month introduces slight variations in natural-language description, leading the AI to generate subtly different logic every session; by archiving the script file alongside its original brief and known-answer sample in a permanent task folder (`my-scripts/`), the institutional rules are frozen, guaranteeing identical, reproducible execution month after month.
- C. The analyst failed to execute an Adversarial Pass (Level 4 verification) on the browser tab's JavaScript engine prior to uploading the May spreadsheet.
- D. Financial valuation tasks must run inside Cowork desktop apps; running them in a browser violates the permission ladder and causes automatic parameter randomization.

Question 30

SINGLE-ANSWER

A digital marketing specialist spends every Monday morning downloading four platform export CSV files (Google Ads, Meta, LinkedIn, TikTok) into a local folder. The specialist opens Claude.ai in a web browser, manually uploads all four files, prompts the AI to clean and consolidate the metrics into a cost-per-acquisition table, downloads the generated HTML summary, and closes the browser tab.

As the number of weekly campaigns expands, the specialist asks whether they should remain on browser-based chat sandboxes or transition to a local agent surface (such as Claude Code or OpenCode).

According to Concept 09 (*Claude.ai: the home surface*) and the graduation path in Part Three (Slide 32), what is the defined **"graduation signal"**, and what core reality dictates moving beyond browser chat?

- A. The graduation signal is triggered whenever an uploaded dataset exceeds 1,000 rows, because cloud-based sandboxes automatically terminate if execution time exceeds three seconds.

Question 31

SINGLE-ANSWER

A clinical director at a regional healthcare center reviews an AI-generated schedule that allocates emergency room nurse staffing based on seasonal patient admission logs. The director has never written or read a line of computer code, but needs to rigorously verify the operational validity of the staffing logic before presenting it to hospital administration.

Applying Level 3 of the Verification Ladder (**Plain-English Replay** from Concept 06), why is instructing the AI to *"explain, step by step in plain English as if to a colleague, what the code did and how it handled edge cases"* so effective at exposing critical vulnerabilities?

- A. The plain-English replay command forces the cloud sandbox to recompile the script into an immutable enterprise ledger that external hospital auditors can cryptographically verify.
- B. Most real-world failures in data automation are procedural errors—such as incorrect matching tolerances, mistaken classification rules, or unstated assumptions—rather than syntax or typographical errors; because flawed procedural logic sounds obviously wrong when stated in plain English, a domain expert can immediately detect and intercept invalid operational rules.
- C. The plain-English prompt overrides the Python execution kernel, forcing the AI to re-estimate all patient counts using conversational heuristic intuition.
- D. Plain-English explanations automatically purge any lines of code that contain non-standard external dependencies or deprecated module imports.

Question 32

SINGLE-ANSWER

A project lead at an infrastructure consultancy prepares an AI brief to audit subcontractor milestone billing files, detect duplicated expense entries, and present the final findings to the executive board. The lead adheres to the foundational design pattern emphasized throughout the chapter: **"Markdown In, HTML Out"**.

According to Concept 05 (*Brief the problem, not the program*) and the crash course's foundational architecture, what is the specific operational rationale behind submitting the brief as a **Markdown specification** and requesting the final output as an **HTML document**?

- A. Markdown text automatically bypasses cloud sandbox safety filters, while HTML documents execute self-extracting JavaScript routines on client machines.
- B. Structuring the brief in Markdown with clear headings and bullet points systematically strips away guesswork for the AI by compartmentalizing every constraint (particularly *Rules* and *Edge cases*), while demanding HTML for the output produces a structured, visually styled, and self-contained document (incorporating charts, tables, and narrative summaries) optimized for human consumption without exposing raw code syntax.
- C. Markdown is required because Python cannot parse raw ASCII text, whereas HTML guarantees that financial figures comply with international statutory tax standards.
- D. Markdown forces the AI to restrict execution to command-line terminals, while HTML restricts execution to isolated cloud sandboxes.

Question 33

SINGLE-ANSWER

An administrative clerk at a distribution warehouse runs an AI-generated automation script directly against the company's sole live inventory directory containing hundreds of master stock records, without making a backup copy. When the operations manager confronts the clerk about the risk of irreversible data corruption, the clerk defends their choice:

"I ran this using Claude Code from Anthropic, a premier AI lab with state-of-the-art safety alignments and benchmarks. Because the underlying AI technology is proven and exceptionally intelligent, taking manual backups is an unnecessary, paranoid waste of time."

Applying the safety doctrine and the **Permission Ladder** established in Concept 12 (*Blast radius: four rules*), what critical conceptual fallacy has the clerk committed?

- A. The clerk failed to realize that frontier AI models automatically restrict file permissions to read-only whenever a production database is detected.
- B. The clerk mistakenly assigned trust to the general technology and vendor brand rather than to the specific script; under the permission ladder, trust is never granted globally to AI technology, but must be earned incrementally by individual, specific scripts that have established a proven track record of safe execution on non-destructive copies first.

Question 34

SINGLE-ANSWER

A quantitative analyst attempts to use an AI chat session to parse a legacy financial tariff table from an exported PDF. Over three consecutive turns, the following sequence occurs:

1. *Turn 1*: The generated script crashes with an unhandled character encoding error; the analyst pastes the red error text, and the AI modifies a regular expression.
2. *Turn 2*: The script executes without crashing, but outputs zero parsed records; the analyst types *"It returned an empty table, fix it"*, and the AI tweaks the pattern matching.
3. *Turn 3*: The script terminates with a buffer allocation exception; the AI apologizes and attempts another minor patch to the exact same regex logic.

As the analyst prepares to submit a fourth prompt asking the AI to *"try fixing that regex once more"*, what procedural rule from Concept 07 (*When it breaks*) must be applied immediately?

- A. Proceed with incremental regex adjustments up to seven iterations, because stochastic LLM kernels require multiple consecutive runs to reach probabilistic convergence.
- B. Apply the **Three-Strikes Rule** and acknowledge **context rot**—three successive failures prove that the fundamental technical approach is wrong (not merely the typing); the analyst must stop digging, abandon the cluttered conversation window (*"abandoning beats rescuing"*), and initiate a clean chat containing the foundational brief and the lessons learned to evaluate two completely different architectural approaches.
- C. Instruct the AI to switch the execution runtime from Python to CSS styling, which natively bypasses document buffer restrictions.
- D. Permanently delete the source PDF document under Rule 4 of the Blast Radius protocol to clear contaminated memory buffers.

Question 35

SELECT ALL THAT APPLY

A clinical audit supervisor is reviewing prompt drafts written by junior clinic administrators who need to extract billing totals from quarterly medical ledger exports. The supervisor must ensure that the team's prompts consistently eradicate the *"silent failure mode"* (glance-based estimates dressed up as precise facts) and reliably compel the AI to execute deterministic code.

According to Concept 04 (*Make the AI's hands move: automatic vs. explicit*), which of the following prompt formulations and operational techniques successfully force computation and provide verification that code actually ran?

- A. *"Roughly how did our clinic pharmacy expenditures trend across the last two quarters? Provide a quick, high-level summary of the primary cost categories."*
- B. *"Write and run code to calculate total pharmacy expenditures by category and month. Show me the code you ran. Before analyzing anything, tell me the exact row count, the column names, and the date range of the file."*
- C. *"What does this pharmacy ledger show? Summarize the spending figures in an organized Markdown table and ensure the numbers appear credible."*
- D. *"Are you running code on this file, or estimating? If estimating, stop and write and run code instead to compute the exact category totals."*
- E. *"Write and run Python 3.12 code using pandas DataFrames and NumPy array operations to compute the monthly spending sums."*

Question 36

SINGLE-ANSWER

A business analyst runs an AI prompt to compute customer churn margins across a portfolio of corporate accounts. The AI generates a 25-line Python script inside a visible code box, runs it in the sandbox, and outputs a customer churn projection table. The analyst immediately copies the figures into an executive briefing deck, telling a colleague:

"The AI showed the exact Python script it ran in the chat window, so that proves the logic is sound and the resulting churn figures are 100% accurate."

Applying the **mechanic's open engine bay analogy** from Concept 02 and the verification doctrine from Concept 06, what critical distinction has the analyst misunderstood?

- A. Displaying an open code block indicates that the cloud sandbox suffered an environment failure and dumped raw syntax rather than executing the script.
- B. A visible code block merely proves that programmatic work occurred (just as a mechanic's open engine bay proves

Question 37

SINGLE-ANSWER

In the chapter's case study of the school administrator auditing overdue tuition payments, an AI review reported that eight families were behind on payments. A subsequent programmatic audit revealed the actual figure was eleven families—meaning three delinquent accounts were silently missed.

Synthesizing the lessons from Concept 03 (*The dividing line*) and Concept 05 (*Where the brief is won*), what were the two distinct root causes behind this discrepancy, and what operational corrections would have prevented it?

- A. The administrator attempted to process tuition files using an open-source terminal agent rather than a commercial browser sandbox, which suppressed accounting exceptions.
- B. First, the administrator prompted conversationally ("look at this spreadsheet and tell me..."), causing the AI to answer from a superficial skim where partial payments were glossed over as cleared; second, the brief omitted an unstated domain rule: "any family carrying any unpaid balance, regardless of partial payment, counts as behind."
- C. The administrator ran the script without setting an external 90-minute timer, causing the AI's internal memory buffer to drop the final three rows of the CSV.
- D. Tuition rosters inherently represent an *Answer Problem*, meaning numerical discrepancies cannot be caught unless audited by a human accountant line-by-line.

Question 38

SINGLE-ANSWER

A clinical quality and compliance auditor at a teaching hospital is reviewing AI-assisted data workflows across three distinct operational settings:

1. *Setting 1 (Personal curiosity)*: A staff nurse uses an AI script to total personal commute fuel receipts for private monthly budgeting.
2. *Setting 2 (Colleague reliance)*: A clinic coordinator uses an AI script to aggregate daily outpatient consultation counts so a colleague can schedule housekeeping staff for the following week.
3. *Setting 3 (High stakes / signed release)*: A chief pharmacist uses an AI script to audit pediatric medication reconstitution logs and dosage concentrations before signing the legal clinical dispensary release.

Under the risk-calibrated **Verification Ladder** in Concept 06 (*The verification ladder: rising with stakes*), what is the minimum required verification protocol for each respective setting?

- A. Setting 1 mandates a Cross-Model Check; Setting 2 requires only Reality Questions; Setting 3 requires an Adversarial Pass without human sign-off.
- B. Setting 1 requires **Reality Questions**; Setting 2 requires a **Known-Answer Test** and a **Plain-English Replay**; Setting 3 requires **all five checks** (Known-Answer Test, Reality Questions, Plain-English Replay, Adversarial Pass, and Cross-Model Check), plus mandatory direct human review of the critical claims.
- C. All three settings require identical, maximum-tier verification protocols, because healthcare institutions are prohibited from establishing tiered verification levels.
- D. Setting 1 requires code inspection; Setting 2 requires JavaScript compilation; Setting 3 requires deploying an unattended *Digital FTE* to bypass human verification.

Question 39

SINGLE-ANSWER

At the end of every academic term, a secondary school teacher spends an entire weekend calculating weighted final exam scores, converting numerical totals into letter grades, and drafting an individual feedback comment for 200 students. Historically, the teacher never hired a programmer to automate this workflow, nor did they attempt to learn programming.

Reflecting on the field case studies in Concept 01 (*Six professions, zero lines of code read*), why was this task historically characterized as "**always describable but never commissionable**", and how does the modern AI client-developer relationship eliminate that barrier?

- A. Educational data was legally classified as non-computable under historical copyright laws, requiring manual calculation until recent statutory amendments.
- B. Commissioning a human developer for small, bespoke, or one-off administrative tasks incurred prohibitive coordination costs: writing technical specifications, negotiating developer fees, and waiting weeks for code that often barely fit the actual workflow—making the overhead of commissioning greater than doing the chore by hand; modern AI provides an on-demand developer that works in seconds, at near-zero marginal cost, and accommodates endless iterative revisions, making even modest routine chores economically commissionable.
- C. Prior to 2023, computer programming languages like Python lacked mathematical formulas for calculating

Question 40

SELECT ALL THAT APPLY

A management consultant is explaining to department leads how AI models autonomously select programming languages based on the requested output format (Concept 02: *The AI picks the language*). The core rule states: "A good client describes the building and lets the architect choose steel or timber. The choice follows the output—and it isn't yours to make."

Which of the following deliverables will naturally compel the AI to write and execute **Python**, rather than JavaScript or TypeScript?

- A. A script that ingests 500 vendor invoice records, calculates net tax liabilities after deductions, and outputs a cleaned CSV ledger.
- B. An interactive, clickable mortgage amortization tool embedded in a web browser where clients can manipulate dynamic interest-rate sliders and toggle repayment options in real time.
- C. A file-system automation script that scans a folder of 1,200 PDF documents and renames each document according to the date and invoice number extracted from within the text.
- D. A browser-based interactive training game featuring clickable buttons, countdown timers, and live leaderboard animations for company onboarding.
- E. A quarterly financial review script that aggregates three departmental budget exports and produces a standalone, static HTML report containing summary tables and an expense trend chart.

Question 41

SINGLE-ANSWER

A risk management officer at an insurance underwriting firm receives an AI-generated risk-tiering analysis derived from 3,000 policyholder records. The numbers appear plausible on the surface, but before presenting the conclusions to the risk committee, the officer decides to execute **Level 4 of the Verification Ladder: an Adversarial Pass** (Concept 06).

According to Concept 06 (*Verifying work you cannot read*), what is the core methodology of an Adversarial Pass, and what prompt instruction operationalizes it?

- A. The officer hires an external ethical hacker to attempt prompt-injection attacks against the cloud sandbox runtime environment.
- B. The officer prompts the AI to aggressively interrogate and attack its own findings: "Find any way this analysis could be wrong or misleading. What assumptions did the code make? What is the weakest link? Score your own confidence from 1 to 10 per claim and justify each score."
- C. The officer commands the AI to translate the Python script into JavaScript to verify if the two language compilers produce bitwise-identical output.
- D. The officer deletes the raw insurance files from their drive and instructs the AI to reconstruct the policy records conversationally.

Question 42

SINGLE-ANSWER

An external consultant proposes an automation solution to a corporate HR director:

"We will implement an automated script that computes a single Composite Productivity Score for every employee based on daily attendance, ticket resolution velocity, and peer feedback forms. The script will independently determine which employees receive annual bonuses and immediately dispatches the updated salary orders directly into the bank's automated clearinghouse without needing managerial sign-off."

Evaluating this proposal through the boundary framework of Concept 13 (*The edge of the map, and what's across it*), what two distinct architectural and governance violations are occurring simultaneously?

- A. The script violates OpenCode licensing terms by using proprietary arithmetic formulas, and fails to package the deliverable as a .docx file.
- B. First, it commits the error of **Judgment in disguise** (*judgment wearing a computation costume*)—while the arithmetic of composite weighting is simple, determining promotions and merit is fundamentally a human policy decision (the AI computes, but the human must decide); second, it executes a **High-stakes act with no undo**, violating the core doctrine that code should *prepare* high-consequence actions while a human must verify and *fire* them.
- C. The proposal breaches Blast Radius Rule 1 by failing to compile the employee payroll database into raw C-language binaries.
- D. Employee evaluation is classified as an automated *Digital FTE* task, which strictly forbids the use of Python scripts.

Question 43

SINGLE-ANSWER

A clinic administrator builds a patient appointment no-show analysis script in March and organizes the output into a task folder containing `brief.md`, `noshow.py`, and a verified `sample-march/` data folder. By August, running the monthly analysis takes 30 seconds. When a new branch opens in Clifton, the administrator simply hands over the directory to the incoming branch manager, who executes the workflow on day one without technical hurdles.

According to Concept 08 (*Keep the script: a solved problem becomes a button you press forever*), what fundamental shift in the economics of knowledge work has occurred, and what specific function does the **Intent Layer** (`brief.md`) serve when business rules change?

- A. The clinic has transitioned into a commercial software vendor, requiring the administrator to submit the Python script to open-source software repositories for copyright protection.
- B. The administrator has crossed the threshold from merely *using* AI on an ad-hoc basis to *accumulating owned operational assets*, reducing the marginal cost of future recurring cycles to near zero; the **Intent Layer** (`brief.md`) preserves the "*what and why*" in human language, allowing any colleague or future AI session to update the operational logic effortlessly without ever reading or reverse-engineering Python syntax.
- C. Storing the script locally converts the operating system into a self-contained *Digital FTE*, which permanently locks the file system against further modifications.
- D. The handover proves that clinical appointment scheduling is an *Answer Problem*, meaning future executions will run as conversational memory summaries rather than programmatic computations.

Question 44

SELECT ALL THAT APPLY

An IT security officer is conducting a risk briefing on executing agentic file-system tasks using terminal environments (such as Claude Code or OpenCode). The officer explains why automated agent scripts present fundamentally different data-loss profiles compared to standard desktop graphical operations (like dragging files in Windows Explorer or macOS Finder).

According to Concept 12 (*Blast radius: four rules*), which of the following statements accurately characterize the technical risks of agentic file operations and the specific protections provided by the four blast-radius rules?

- A. Unlike manual deletions in graphical desktop environments that typically stage discarded files in the Recycle Bin or Trash, programmatic deletions executed by terminal scripts frequently bypass the recycle bin entirely, resulting in immediate, permanent disk erasure.
- B. When an automated script modifies a file in-place, the file system does not preserve an application undo stack; therefore, **Rule 4 (Output to new files)** preserves original source records and ensures every erroneous execution remains fully reversible simply by deleting the newly generated file.
- C. Launching an agent from a broad directory level (such as the user's home directory or root drive) grants a misunderstood script the latitude to damage irreplaceable files across the entire machine; **Rule 3 (Scope the territory)** neutralizes this by confining the agent's working path strictly to the smallest folder containing the problem.
- D. Demanding a dry run is unnecessary when using terminal agents, because modern command-line operating systems automatically execute file manipulations within an internal preview sandbox before committing changes to physical storage.
- E. Creating an initial duplicate backup of the target directory (**Rule 1: Work on copies**) guarantees that if an unproven script executes an erroneous procedure, the failure results in a harmless operational hiccup rather than an unrecoverable business catastrophe.

Question 45

SINGLE-ANSWER

A tax consultant is working under tight deadline pressure late at night to finalize an annual corporate tax calculation derived from 12,000 general-ledger transactions across 40 bank accounts. Because time is running out, the consultant considers skipping all post-execution verification checks and submitting the return directly, reasoning that the AI generated a clean Python script that executed without any syntax errors.

According to the non-negotiable principle established in Concept 06 (*Verifying work you cannot read*), what is the single firewall check that a client must **NEVER skip** before acting on precision-critical figures, and how is it efficiently executed?

- A. The consultant must transmit the script to an external multi-modal vision model to visually evaluate the formatting of the numbers.
- B. The consultant must execute a **Known-Answer Test** against a small, pre-verified slice of data whose correct

Question 46**SINGLE-ANSWER**

During a workflow automation seminar, an operations director notes that while any one of the four VPRF signals (*Volume, Precision, Repetition, Files*) is sufficient to trigger programmatic code execution, the chapter explicitly designates **Repetition** as "**the strongest signal**" (Concept 03 and Concept 08).

Why does the chapter single out **Repetition** as offering a unique strategic payoff that the other three signals do not possess?

- A. Repetition automatically elevates the task to an enterprise API tier, granting priority access to cloud GPU clusters.
- B. While Volume, Precision, and Files justify writing code to solve an immediate, isolated data problem, Repetition converts a solved problem into a permanent, reusable operational asset; once captured in a saved script, the marginal cost of performing the work in every future cycle drops to near zero, marking the exact moment an individual transitions from merely *using* AI to *accumulating* long-term organizational leverage.
- C. Repetitive data workflows are exempt from the *Known-Answer Test*, allowing users to bypass verification after the initial run.
- D. Tasks that repeat on a weekly cadence legally obligate the AI to write self-deleting shell scripts.

Question 47**SINGLE-ANSWER**

An operations director is establishing an internal assessment task based on the project curriculum in the presentation slides (Slide 37: *From rehearsal to real, on your own data*), specifically targeting **Project 3: The Handover**.

According to the chapter's specifications, what is the precise operational objective and definitive benchmark of a successful **Handover**?

- A. Deploying a scalable, multi-tenant cloud application that handles 500 simultaneous external logins with payment gateway processing.
- B. Converting a recurring operational workflow into a self-contained, portable directory asset that a colleague can run independently and successfully—relying solely on the files stored within that directory (brief.md, the script, and sample data)—with the original creator completely out of the room.
- C. Sharing an active, authenticated browser tab session with a coworker so they can continue a live conversation inside an ongoing cloud sandbox.
- D. Transferring operating system administrative root privileges to an unattended *Digital FTE* background daemon.

Question 48**SELECT ALL THAT APPLY**

An IT infrastructure architect is preparing an internal technology evaluation matrix comparing the **Five Surfaces** presented in Part Three of the chapter (Concepts 09–11 and Slide 32):

- Browser chat sandboxes (*Claude.ai, ChatGPT, Gemini*)
- Terminal folder agents (*Claude Code, OpenCode*)
- Desktop delegation apps (*Cowork, OpenWork*)

According to the chapter's taxonomy, graduation path, and architectural trade-offs, which of the following statements correctly describe these surfaces?

- A. Browser chat sandboxes execute code inside an isolated, temporary cloud container on the provider's side where local storage is never touched, creating a zero-risk perimeter ideal for one-off and exploratory tasks.
- B. In both the terminal agent pair and the desktop application pair, the first tool in each pair (Claude Code, Cowork) represents a commercial offering, while the second tool (OpenCode, OpenWork) represents an open-source equivalent.
- C. Desktop delegation apps like Cowork and OpenWork discard code generation completely, executing all desktop transformations via operating-system GUI mouse-macro automation.
- D. Only code operating directly on the local machine (such as via Claude Code or OpenCode) can execute on an unattended schedule while the user sleeps or reach network devices beyond a browser tab's sandbox.
- E. Non-programmers should bypass browser chat sandboxes immediately on day one, because browser sandboxes are architecturally incapable of processing CSV files.

Question 49

SINGLE-ANSWER

An administrative clerk uses a local agent to rename and standardize 800 intake forms in an active clinic folder. In their prompt, the clerk specifies:

"Rename all files in this directory so that the patient registration date is at the front. If you make an error on any filename, I will simply hit Ctrl+Z in the operating system to undo the change."

According to the safety doctrine in Concept 12 (*Blast radius: four rules*), why is relying on Ctrl+Z a dangerous fallacy, and which specific Blast Radius rule ensures true reversibility?

- A. File systems require typing Ctrl+Shift+Z to reverse batch programmatic changes, whereas standard Ctrl+Z is restricted to keyboard text entry.
- B. File system modifications executed by automated scripts do not maintain an operating-system undo stack; once a file is overwritten or renamed in-place by code, keyboard undo shortcuts cannot restore the original state. Therefore, **Rule 4 (Output to new files)** must be enforced—writing transformed results to distinct new files while leaving source files untouched, ensuring any run is reversible simply by deleting the new output.
- C. Ctrl+Z only functions if the script was written in JavaScript; Python scripts permanently disable OS hotkeys upon execution.
- D. The clerk should have executed the task inside a temporary browser sandbox, which automatically takes disk-level system restore points every 60 seconds.

Question 50

SINGLE-ANSWER

In the closing synthesis of the course (*The whole course in one line*, *The promise stated precisely*, and Slide 39: *One identity · One skill*), the chapter articulates the permanent transformation of knowledge work and defines the exact relationship between the learner, the AI, and code.

What is the definitive summary of this new identity, how are the "**freedom**" and the "**responsibility**" harmonized without tension, and what is the single essential skill that defines a great client for code?

- A. The learner's ultimate goal is to transition into a professional full-stack software engineer; skipping code review is merely a temporary training compromise that must be discarded once Python syntax is mastered.
- B. You are not learning to write code; you are learning to be a **good client for code**. Good clients do not lay bricks—they write clear briefs, check the work against what they can independently measure, and keep what they paid for. Not having to write code or read its syntax is the **freedom**; rigorously verifying what the code produced against independent ground truth is the non-negotiable **responsibility**. The two are not in tension: you achieve earned confidence in work you cannot see because you tested it against something you already knew.
- C. The learner abdicates both execution and verification to autonomous *Digital FTEs*, eliminating human oversight and relying on AI vendor alignment benchmarks for accuracy.
- D. The learner uses code exclusively for recreational and gaming tasks, while restricting serious enterprise financial and operational calculations to manual hand entry.

Official Master Answer Key & Reference Solution (Q1–Q50)

Question	Type	Correct Answer	Question	Type	Correct Answer
Q1	Single-Answer	B	Q26	Single-Answer	A
Q2	Single-Answer	B	Q27	Single-Answer	B
Q3	Select All That Apply	A, C, D	Q28	Select All That Apply	A, C, D
Q4	Single-Answer	B	Q29	Single-Answer	B
Q5	Single-Answer	B	Q30	Single-Answer	B
Q6	Single-Answer	C	Q31	Single-Answer	B
Q7	Select All That Apply	A, C, D, E	Q32	Single-Answer	B
Q8	Single-Answer	C	Q33	Single-Answer	B
Q9	Single-Answer	B	Q34	Single-Answer	B
Q10	Single-Answer	C	Q35	Select All That Apply	B, D
Q11	Single-Answer	B	Q36	Single-Answer	B
Q12	Single-Answer	B	Q37	Single-Answer	B
Q13	Single-Answer	B	Q38	Single-Answer	B
Q14	Single-Answer	B	Q39	Single-Answer	B
Q15	Single-Answer	B	Q40	Select All That Apply	A, C, E
Q16	Select All That Apply	B, D, E	Q41	Single-Answer	B
Q17	Single-Answer	B	Q42	Single-Answer	B
Q18	Single-Answer	B	Q43	Single-Answer	B
Q19	Single-Answer	B	Q44	Select All That Apply	A, B, C, E
Q20	Select All That Apply	A, B, C, E	Q45	Single-Answer	B
Q21	Single-Answer	B	Q46	Single-Answer	B
Q22	Single-Answer	B	Q47	Single-Answer	B
Q23	Single-Answer	B	Q48	Select All That Apply	A, B, D
Q24	Select All That Apply	A, C, D, E	Q49	Single-Answer	B
Q25	Single-Answer	B	Q50	Single-Answer	B