

CCAO-F Practice MCQs

132 scenario questions in 4 sections, answer and Roman Urdu explanation under each question. Cover the green answer line with a paper while practicing.

Section	Questions
Section A: PCAO-F Panaversity Mock Test (41 questions)	41
Section B: Output Evaluation and Validation Practice (5 questions)	5
Section C: Domain-wise Question Bank (deduplicated from your files)	61
Section D: New Questions Built From the 10 Crash Courses (25 questions)	25

How to practice: Read the last line of the question first. Note whether it says FIRST, BEST, or Select TWO/THREE. Eliminate extreme and 'magic fix' options. Target: 80%+ in each section before exam day.

Roman Urdu: Pehle sawal ka aakhri hissa parhein. 'Select TWO/THREE' ka khayal rakhein. Har section mein 80% se zyada ka hadaf rakhein.

Section A: PCAO-F Panaversity Mock Test (41 questions)

Official-style Panaversity mock. Attempt in 80 minutes without looking at answers, then check.

Q1. A reporting team wants to redesign its monthly inventory report around an AI assistant. It plans two changes (consolidating three prompts into one, moving format rules into standing instructions) and wants to report an improvement figure to management. What should the team do first? [Workflow]

- A. Move the repeated format rules into configuration first, so every later run benefits
- B. Consolidate the three prompts into a single prompt, since removing steps is the biggest saving
- C. Pick 'time saved' as the figure to report, so improvement can be shown quickly
- D. Run the current process once unchanged and record its time, revision rounds, and manual steps

Correct Answer: D

Easy Explanation: Pehle baseline napo. Tabdeeli ke baad purana waqt dobara nahi napa ja sakta.

Q2. An operations associate must reconcile fourteen months of technician-shift exports against site sign-in sheets. They opened only one file; a colleague warned the exports are 'messy'. The associate plans to write rules for every imaginable data problem. What should the first prompt be? [Prompting]

- A. The full set of rules for every irregularity the associate can imagine
- B. Produce the reconciliation now and fix data problems after the first output
- C. Profile the files and list anything ambiguous or irregular before any rule is written
- D. Request the code in a named language so the associate can read the logic first

Correct Answer: C

Easy Explanation: Pehle data ko samjho (profile). Andaze se rules likhna ghalat hai.

Q3. A proposal's control section says 'All replies will be reviewed by a person before they are sent.' A colleague says it is fine because it keeps a human in the loop. Most appropriate improvement? [Governance]

- A. Change to 'light human review' so reviewers are not overloaded
- B. Add that a reviewer samples 10% of replies after sending
- C. Name the role that reviews, the specific risk the review checks for, and keep that it happens before sending
- D. Remove the requirement since an unverifiable rule is worse than none

Correct Answer: C

Easy Explanation: Gate mukammal ho: kaun, kya check karega, kab. Sirf 'insaan dekhega' kafi nahi.

Q4. An associate asks Claude for next year's event calendar with budget per event. The confident plan uses last year's budget; last week the CFO halved it in a meeting recorded in no document. What best explains the failure? [Workflow]

- A. The prompt was under-specified; clearer instructions would have made the model hedge
- B. Wrong model tier; a more capable model would have flagged uncertainty
- C. A step depended on an input never in front of the model, so it filled the gap at full confidence
- D. The attached data was too large, so the context window dropped the budget part

Correct Answer: C

Easy Explanation: Window test: jo maloomat model ke saamne hi nahi, woh usay andaze se pooray yaqeen ke sath bhar deta hai.

Q5. An associate must apply 140 payroll corrections to named employee records; identifiers are the substance of the task. Policy classes such records as restricted. Most appropriate action? [Governance]

- A. Redact the identifiers anyway and apply as many corrections as possible
- B. Confirm with the compliance owner which route is approved for this data and purpose, then use that route
- C. Upload to the enterprise Claude workspace since enterprise is approved for internal work
- D. Build a custom connector to the HR system yourself

Correct Answer: B

Easy Explanation: Jab identifiers hi kaam ka hissa hon to redaction nahi chalega; approved route owner se confirm karo.

Q6. A strategy associate wants a proposal for redesigning the returns process and plans to ask for the full proposal straight away with extended thinking. Which first request most improves the eventual proposal? [Prompting]

- A. Ask for the complete proposal now with extended thinking and iterate
- B. Ask Claude first to map what is known: existing approaches, other firms' fixes, prior attempts
- C. Ask for three complete proposals in different styles and pick one
- D. Ask Claude to grade the associate's outline out of ten first

Correct Answer: B

Easy Explanation: Pehle maloomat ka naqsha banao, phir draft. Loop context se shuru hota hai.

Q7. A Google Drive tender folder is connected read-only. A new workflow needs Claude to move closed tenders into an archive folder. Which understanding is correct? [Configuration]

- A. Moving is already possible because the connector inherits everything the account can do
- B. Moving can be enabled by adding a Skill that instructs Claude how to archive
- C. Moving needs a second connector because read and write cannot share one connection
- D. Moving files is a write action to be granted deliberately, once the connector is confirmed to offer it

Correct Answer: D

Easy Explanation: Move karna write action hai; soch samajh kar ijazat do. Skill ijazat nahi deti.

Q8. Which THREE features of an AI market summary are signs a claim may be fabricated? (Select THREE.) [Output Evaluation]

- A. A statistic to two decimal places with no source
- B. Hedged phrasing such as 'this may vary by region'
- C. A statement that the assistant 'checked the supplier's website' with no page or quote
- D. A summary longer than expected
- E. A figure on page one contradicting the same figure near the end
- F. A structure that follows the associate's report template

Correct Answer: A, C, E

Easy Explanation: Bohat exact number baghair source, jhoota claimed action, aur andar ka tazad: teenon jhoot ki nishaniyan.

Q9. Two hours into a supplier-audit chat, Claude drops the page numbers it was told to include, contradicts an earlier score, and drifts into general advice. Most appropriate move? [Troubleshooting]

- A. Restate the rule and correct score every turn and keep going
- B. Open a blank chat and start the audit from the very beginning
- C. Ask for a summary of decisions, open points and rules; correct it; paste it into a fresh chat with the files
- D. Move the whole conversation into a Project so the drift stops

Correct Answer: C

Easy Explanation: Context overload: summary banao, theek karo, nayi chat mein files ke sath shuru karo.

Q10. First draft of a request: 'Tell me this shortlist is defensible.' Which TWO rewrites most likely produce an honest evaluation? (Select TWO.) [Prompting]

- A. 'Rate the shortlist 1 to 10 on coverage, cost realism, and risk, with the one change that most improves it.'
- B. 'Defend the shortlist against finance partners' likely objections.'
- C. 'Give the strongest case for and against sending the shortlist to the committee as it stands.'
- D. 'Isn't this shortlist good enough? Point out anything minor to tidy.'
- E. 'Rate your confidence, as a percentage, that it would survive the committee.'
- F. 'Rewrite the cover note so it reads more persuasively.'

Correct Answer: A, C

Easy Explanation: Sycophancy hatane ke liye: har criteria par score + behtari, aur dono taraf ke dalail.

Q11. A support team runs meeting-notes tidy-ups on Opus with extended thinking, and runs its annual pricing-model rebuild on Haiku 'to save budget'. Which statement best describes the problem? [Model Selection]

- A. Only the pricing choice is wrong
- B. Neither is wrong; quality differences are marginal
- C. Only the tidy-up choice is wrong
- D. Both are wrong in opposite directions: flagship wasted on simple work, cheapest rung wastes the rebuild

Correct Answer: D

Easy Explanation: Model kaam ke mutabiq: chhote kaam par Opus paisa zaya, ahem kaam par Haiku kaam zaya.

Q12. An associate asks Claude to 'write up the inspection'. The insurer's loss adjuster needs only the three open defects, each with date found and photo reference. A colleague suggests a much longer prompt. Most appropriate improvement? [Prompting]

- A. State the reader, the purpose, the length, the items to cover, and what to leave out
- B. Describe the warehouse and inspection day in full
- C. Rerun on Opus with extended thinking
- D. Attach last year's narrative so Claude copies it

Correct Answer: A

Easy Explanation: Lamba nahi, mukammal prompt: kaun parhega, kyun, kitna, kya shamil, kya nahi.

Q13. A delegated venue search: two venues tied on price and Claude silently dropped one; it silently widened dates to 'late March'. Which element of the request would have prevented the silent choices? [Prompting]

- A. Show the plan of steps before starting
- B. More detail about the event and audience
- C. List every assumption after finishing
- D. A line saying what Claude may decide alone and that it must say when it chose

Correct Answer: D

Easy Explanation: Brief ka 'Authority' hissa: AI khud kya decide kar sakta hai aur kab batana lazmi hai.

Q14. An associate proposes running supplier-pricing chats in Claude incognito (and ChatGPT temporary chat) 'so there is nothing on record if audit asks'. Which statement is accurate? [Governance]

- A. Both modes discard the conversation at once and nothing is stored
- B. Both keep chats out of history and memory, but each vendor retains a copy for a period and company records may hold it
- C. Claude incognito leaves no trace but ChatGPT temporary chat is retained
- D. Neither mode changes anything; both are stored in history

Correct Answer: B

Easy Explanation: Incognito ka matlab zero retention nahi; audit se chhupana bhi ghalat niyat hai.

Q15. In every new chat Claude addresses an associate as lead of a project that closed months ago. They have been adding a correction line to every prompt. Most appropriate fix? [Configuration]

- A. Keep adding the correction line
- B. Use incognito for everything
- C. Switch to a different model
- D. Read the stored memory, correct or delete the closed-project note, then start a fresh chat

Correct Answer: D

Easy Explanation: Stale memory: memory parho aur purana note theek/delete karo.

Q16. Badge numbers were replaced with coded values, but the workbook's second tab still holds the lookup mapping codes back to employees. Most appropriate action? [Governance]

- A. Proceed; data is now anonymised
- B. Upload in incognito so the lookup is not retained
- C. Remove the lookup tab and unneeded fields, confirm no mapping travels with the file, then run the analysis
- D. Abandon the analysis

Correct Answer: C

Easy Explanation: Mapping mojud ho to data pseudonymised hai, anonymised nahi. Mapping hatao.

Q17. Claude cites a 2023 consultancy blog for the threshold at which a data-protection rule currently applies. The post exists and states that figure. Most appropriate next step?

[Output Evaluation]

- A. Find the regulator's own current page, confirm the figure and date there, and cite that
- B. Use the figure since the citation was opened
- C. Ask Claude for two more sources and use it if they agree
- D. Drop the figure entirely

Correct Answer: A

Easy Explanation: Maujooda qanoon ke liye primary aur taaza source (regulator) se tasdeeq.

Q18. For which TWO tasks is the associate's usual read-through NOT enough, because it needs a second reviewer or must be held and verified before action? (Select TWO.)

[Governance]

- A. An agenda for their own one-to-one
- B. A routine reply to one customer they send personally
- C. A press statement published on the company website
- D. Cancelling forty duplicate purchase orders in the ERP from the assistant's list
- E. Five title ideas for the internal newsletter

Correct Answer: C, D

Easy Explanation: Public statement aur na-qabil-e-wapsi ERP action: dono par zyada check.

Q19. For which TWO is extended thinking worth its cost? (Select TWO.) [Model Selection]

- A. Converting a bulleted agenda into a calendar invite
- B. Choosing a rollout sequence across four warehouses under staffing, downtime and budget constraints
- C. Finding an acronym in an attached glossary
- D. A staffing decision across three sites with real consequences, seasonality and overtime rules
- E. Ten quick name ideas for a chat channel

Correct Answer: B, D

Easy Explanation: Thinking sirf kai pabandiyon wale ahem faislon ke liye.

Q20. To 'save context', an associate pastes a one-page summary of a 70-page agreement and asks for a risk analysis; two indemnity clauses not in the summary are missed. Most appropriate change? [Configuration]

- A. Upload the full agreement, since analysis quality depends on what is in front of the model
- B. Paste a longer, more careful summary
- C. Switch to Opus
- D. Ask Claude which sections it needs and paste only those

Correct Answer: A

Easy Explanation: Model sirf wohi dekhta hai jo window mein hai. Poora document do.

Q21. Wholesale and retail pricing files sit in one Project. In a retail quote, Claude applied wholesale volume discounts. Most appropriate fix? [Configuration]

- A. Add 'unit: retail' to every prompt
- B. Create one Project per unit with its own instructions and knowledge, and move the files
- C. Use incognito for retail quotes
- D. Delete wholesale files after each retail chat

Correct Answer: B

Easy Explanation: Har unit/client ka alag Project, taake context aapas mein na mile.

Q22. Which THREE pair the request with the right feature? (Select THREE.) [Model Selection]

- A. 'What did the port authority announce this morning?' -> web search
- B. 'Pull last quarter's approved rate card from the finance shared drive' -> web search
- C. 'Which lease is riskier to exit early, given the break clauses?' -> extended thinking
- D. 'What is today's VAT rate on imported machinery?' -> research mode
- E. 'Cited comparison of cold-chain providers in East Africa, with funding and key customers' -> research mode
- F. 'Rewrite this three-line status update friendlier' -> extended thinking

Correct Answer: A, C, E

Easy Explanation: Aaj ki khabar = web search; mushkil muqabla = thinking; cited multi-source report = research. Company drive = connector.

Q23. Configuring Claude for ongoing work with one client. Which THREE placements are correct? (Select THREE.) [Configuration]

- A. Firm-wide writing tone rule for every chat -> account-level instructions
- B. Client's recurring reference files and past deliverables -> client Project knowledge
- C. 'Negotiation paused until end of month' -> account-level instructions
- D. 'Treat figures absent from the client's files as unconfirmed' -> client Project instructions
- E. A one-off question about an unrelated memo -> its own new Project
- F. Preference for short replies -> memory, by mentioning it once

Correct Answer: A, B, D

Easy Explanation: Har chat ka usool account par, client files Project knowledge mein, client ka usool Project instructions mein.

Q24. A junior analyst with no customs background drafts a customs-compliance briefing with Claude. It meets every requirement and matches the regulation extract. Which check most directly addresses the remaining gap? [Output Evaluation]

- A. Ask a colleague with customs experience whether it is credible by the field's standards
- B. Ask a colleague to proofread tone and clarity
- C. Ask Claude to rate accuracy and revise below 9/10
- D. Run a formatting pass

Correct Answer: A

Easy Explanation: Teesra reference: field ke expert ka meyaar. Junior yeh khud check nahi kar sakta.

Q25. A team lead cannot tell whether recording calls for AI summaries needs consent in a new jurisdiction; it must go to the legal owner. Which escalation gives the owner the most useful basis? [Governance]

- A. The workflow, who it reaches, the control already added, and the one specific question policy leaves open
- B. A risk matrix with a total score and a recommendation to approve
- C. 'I think this is fine, can you approve today?'
- D. A request to pause until the policy is updated

Correct Answer: A

Easy Explanation: Achhi escalation: poori tasveer + ek wazeh sawal.

Q26. The CFO version of a workflow description reads 'AI clears the disputes, cutting the backlog from ten days to two.' What is the problem? [Workflow]

- A. No published source for the backlog figure
- B. It drops the human checkpoint, so the CFO believes there is no gate
- C. It is too brief for the CFO
- D. No problem: executives want outcomes

Correct Answer: B

Easy Explanation: 'AI handles X' jaisa jumla human gate ghaib kar deta hai.

Q27. An analyst asks for: a quick explanation of 'demurrage'; a nine-broker comparison finance will load into a sheet; a one-page onboarding checklist circulated to new hires. Which THREE format choices are correct? (Select THREE.) [Configuration]

- A. Broker comparison as structured output (table or spreadsheet)
- B. Demurrage explanation as an inline chat reply
- C. Checklist as an inline reply new hires copy out
- D. Checklist as a document artifact
- E. Broker comparison as a prose paragraph
- F. Demurrage explanation as a document artifact

Correct Answer: A, B, D

Easy Explanation: Chhoti wazahat chat mein, data table mein, baanta jane wala document artifact mein.

Q28. A Claude-drafted pricing notice with a serious error went to 300 customers an hour ago. Which THREE actions belong in the first hour? (Select THREE.) [Governance]

- A. Stop the spread: halt further sends and scheduled runs
- B. Write down what went out, to whom, when, and what was wrong
- C. Delete the chat and the draft
- D. Wait until full scope is understood before telling anyone
- E. Report promptly through the incident path, stating facts without guessing
- F. Have Claude draft a correction and send it yourself within the hour

Correct Answer: A, B, E

Easy Explanation: Roko, likho, report karo. Saboot mitana ya khud correction bhejna ghalat.

Q29. Claude pre-screens 900 warranty claims and routes 120 for fast-track. The lead reviews all 120 carefully and finds them sound. What is the most important missing step? [Governance]

- A. Add a note that AI did the first screening
- B. Tell the 780 others that AI was used
- C. Sample the 780 not fast-tracked, to check whether a group was routed away systematically
- D. Rerun with a stricter prompt and approve only claims in both sets

Correct Answer: C

Easy Explanation: AI ne jo bahar kiya use bhi check karo; sirf chune hue nahi.

Q30. A customer-onboarding workflow ran well for four months. Which TWO changes leave Claude with more authority than the step's risk allows? (Select TWO.) [Workflow]

- A. Let Claude approve KYC exceptions directly because summaries were accurate
- B. Have the compliance analyst sign off exceptions from the summary line without opening any case
- C. Move document extraction into a reusable Skill
- D. Compute risk scores with code and keep the approval gate right after
- E. Add a named reviewer and checkpoint for the rationale step

Correct Answer: A, B

Easy Explanation: A halo/over-delegation hai; B unstaffed gate (review sirf click).

Q31. Call transcripts collected with consent for quality assurance sit in an approved workspace. The product team wants them mined for upsell leads for sales. Most appropriate action? [Governance]

- A. Proceed since they are already in an approved workspace
- B. Strip names, then run the upsell analysis
- C. Run it in incognito
- D. Take it to the data owner, since a new purpose needs fresh permission

Correct Answer: D

Easy Explanation: Data jis maqsad ke liye liya, naye maqsad ke liye owner ki ijazat.

Q32. In a fresh chat, the first-ever tariff-change notice omits the legal footer and uses the wrong currency. Which action fits the symptom? [Troubleshooting]

- A. Switch to Opus
- B. Start another fresh chat with the same request
- C. Treat it as under-specification: re-read the request against its parts and add what it left out
- D. Check whether a stored instruction has gone out of date

Correct Answer: C

Easy Explanation: Pehle jawab se ghalat + pehli dafa = prompt adhoora.

Q33. A finding must reach the audit committee and the warehouse crew whose process changed. Most appropriate way to produce the two versions? [Output Evaluation]

- A. One rewrite per reader, each shaped around what that reader must decide or do, then edit in named passes
- B. Make it 'more professional' and send one version to both
- C. Same document with different cover notes
- D. Three committee drafts and forward the best to the crew

Correct Answer: A

Easy Explanation: Har parhne wale ke faislay/amal ke gird alag version.

Q34. A personal quote generator in a Project is now a CRM step; two branches quote customers from it; the associate patches the prompt whenever a branch reports an error. Most appropriate action? [Workflow]

- A. Keep improving the Project prompt
- B. Rebuild it with the API yourself
- C. Escalate to engineering: others depend on it, so it needs an owner, uptime and access control
- D. Package it as a shared Skill

Correct Answer: C

Easy Explanation: Doosre log depend karen = escalation signal, engineering ke hawale.

Q35. Which TWO correctly turn a reaction into an instruction aimed at the right part of the setup? (Select TWO.) [Troubleshooting]

- A. 'Keeps writing for engineers' -> add an audience constraint to the Project's standing instructions
- B. 'Cites the discontinued product line' -> replace the outdated catalogue in Project knowledge
- C. 'Too long every time' -> switch to a more capable model
- D. 'Missed this week's campaign point' -> put this week's goal into standing instructions
- E. 'Buries the ask' -> tell Claude to remember it via memory

Correct Answer: A, B

Easy Explanation: Rule = instructions, purana reference = knowledge badlo. Haftay ka goal arzi hai, prompt mein do.

Q36. A Project instruction says 'Flag unusual expenses.' Each run flags different items and colleagues get different results. Most appropriate fix? [Troubleshooting]

- A. Add 'be careful and consistent'
- B. Switch to a more capable model
- C. Rewrite the rule with a named amount threshold and expense category
- D. Add three examples and let the model infer the threshold

Correct Answer: C

Easy Explanation: Mubham usool ko check hone wale criteria (threshold + category) mein badlo.

Q37. An assistant-written script prepares the monthly supplier payment run and passed a known-answer check. A colleague suggests it release payments automatically. Most appropriate arrangement? [Governance]

- A. Script prepares the run and a named person releases it, since a released payment cannot be recalled
- B. Release automatically since it is tested
- C. Treat a clean confirmation report as approval
- D. Stop using the script at all

Correct Answer: A

Easy Explanation: Na-qabil-e-wapsi action se pehle naam wala insaan.

Q38. After twelve clean runs, a vendor proposes outputs be signed 'AI-generated' with no named human owner. Which statement is correct? [Governance]

- A. A track record earns similar jobs, but each output still carries a person's name and that person answers for it
- B. Accountability moves to the vendor
- C. The label replaces the need for an owner
- D. An owner is needed only if a client challenges an output

Correct Answer: A

Easy Explanation: Zimmedari kabhi machine ya vendor par nahi jati.

Q39. A discount is computed with code; next a supervisor releases or holds each payment. Someone proposes removing the release step 'because code is always right'. Most appropriate response? [Workflow]

- A. Agree: the step is redundant
- B. Keep the step: computation is safe to delegate only because a wrong figure can be reversed and a person decides next
- C. Replace the supervisor with a second AI pass
- D. Move the calculation back to the supervisor by hand

Correct Answer: B

Easy Explanation: Computation isliye delegate hua kyunke aage insaan hai. Gate hatane se yeh qadam ghair mehfooz.

Q40. After four rounds, a scorecard's columns and formatting are right, but the weighted-score column is still subtly wrong despite ever more precise wording. Most likely root cause? [Troubleshooting]

- A. A feature problem: weighted scores written as prose belong in code execution
- B. A description problem: one more precise attempt
- C. A model problem: needs Opus
- D. A context problem: four rounds overloaded the chat

Correct Answer: A

Easy Explanation: Ek hi qism ki hisaabi ghalti = wrong feature; code execution.

Q41. Sonnet with extended thinking handled a 20-page contract well, but on a 180-page agreement the answer is fluent but shallow and misses obligations spread across sections. Which change best fits? [Model Selection]

- A. Move to Opus, built for holding a long, complex structure coherent at once
- B. Switch to Haiku
- C. Turn thinking off
- D. Split into ten separate chats

Correct Answer: A

Easy Explanation: Lamba aur pecheeda structure = top tier (Opus).

Section B: Output Evaluation and Validation Practice (5 questions)

Domain 2 is 21% of the exam. These are harder, reasoning-style items.

Q42. A pharmacy manager plants one fake missing-units discrepancy; the script catches it plus four unknown flags. She calls the script 'proven'. Past real discrepancies came from double-counted ward transfers. What does catching the planted fault entitle her to conclude?

- A. It proves the script is reliable for all discrepancy types
- B. It proves nothing because a planted fault is artificial
- C. It supports flags of that class (missing-unit gaps); double-counted transfers are another class, so she should plant one of those too
- D. It entitles her to action the four extra flags now

Correct Answer: C

Easy Explanation: Nakli ghalti sirf usi qism ki tasdeeq karti hai; har qism ke liye alag test.

Q43. A controller gives the identical brief and files to two AI assistants from different companies; both return the same 23 mismatched invoices. How much assurance does agreement give?

- A. Weak, because they share training data
- B. Adequate once each rates its confidence
- C. Strong, two vendors are unlikely to invent the same mismatches
- D. Useful but bounded: both read the same brief, so an omitted rule (e.g. netting credit notes) is an error both can make

Correct Answer: D

Easy Explanation: Ek hi brief ki kami dono mein hogi. Faida hai magar mehdood.

Q44. Claude recommends Supplier B and lists assumptions; one reads 'Supplier B guarantees delivery within 48 hours', but no proposal contains that clause. Which action correctly uses the listed assumptions?

- A. Open extended thinking to see the true reasoning
- B. Ask Claude whether it is sure
- C. Accept, since the chain is coherent
- D. Treat the list as a justification offered for review and check each assumption; the 48-hour claim has no source, so the recommendation cannot stand

Correct Answer: D

Easy Explanation: Assumptions ko documents se milao; baghair source claim par recommendation nahi tik sakti.

Q45. Claude runs code on a loan book and reports PKR 4.31bn outstanding principal; the manager's manual review expected about 3.94bn. Best next step?

- A. Accept the code result
- B. Switch to the highest tier and rerun
- C. Check what was computed: which column and rows, with counts and samples; exact code can still sum gross instead of outstanding
- D. Regenerate the total in prose

Correct Answer: C

Easy Explanation: Code ka chalna sahi hone ki zamanat nahi; kya calculate hua woh dekho.

Q46. Which TWO outputs need the full four-step verification (identify, trace, challenge, decide) and a named reviewer? (1) internal call recap (2) demurrage-exposure figure for a counterparty (3) first-pass email tags (4) draft response to a regulator. (Select TWO.)

- A.** Output 3, email tagging
- B.** Output 1, call recap
- C.** Output 4, the regulator response
- D.** Output 2, the demurrage-exposure figure

Correct Answer: C, D

Easy Explanation: Regulator aur financial raqam wali cheezon par sakht review.

Section C: Domain-wise Question Bank (deduplicated from your files)

Merged from the Smart Practice Test and the CCAO-F Domain Question Bank; duplicates removed.

Q47. An analyst must find where Claude could improve a 12-step invoice-review process; some steps copy data, others need finance judgment. What should the analyst do FIRST?

[Workflow]

- A. Automate the full process, then check judgment steps
- B. Map the current process, separate repetitive from judgment-heavy steps, and identify suitable opportunities
- C. Ask Claude to redesign it without documenting the current workflow
- D. Start with the longest step even if it includes approvals

Correct Answer: B

Easy Explanation: Pehle mojudda process ka naqsha, phir faisla.

Q48. A marketing team has a strong approval process but slow synthesis of interview notes. Which use case is MOST appropriate? [Workflow]

- A. Claude approves the final brief
- B. Claude replaces customer research with generated reactions
- C. Claude synthesizes research into draft themes and evidence; keep the human approval process
- D. Claude selects the campaign budget

Correct Answer: C

Easy Explanation: Jahan rukawat hai (synthesis) wahan AI; approval insaan ke paas.

Q49. A pilot cuts drafting time 40% but reviewers still spend the same time correcting missing source details. BEST next action? [Workflow]

- A. Declare success
- B. Scale immediately
- C. Measure only satisfaction
- D. Investigate why source details are missed and refine before expanding

Correct Answer: D

Easy Explanation: Asal bottleneck theek kiye baghair scale na karo.

Q50. Routine account-access questions plus unusual cases needing specialist judgment. Which workflow is MOST appropriate? [Workflow]

- A. Claude answers everything; specialists handle complaints
- B. Keep everything manual
- C. Specialists write templates and stop using Claude
- D. Claude drafts routine responses; unusual cases route to specialists

Correct Answer: D

Easy Explanation: Routine AI, exceptions insaan.

Q51. A process owner cannot explain what success would look like after integrating Claude. What should the team do? [Workflow]

- A. Define measurable success criteria before finalizing the workflow
- B. Implement first, define later
- C. Use time saved as the only metric
- D. Let Claude propose criteria and score itself

Correct Answer: A

Easy Explanation: Kamyabi ka paimana pehle tay karo.

Q52. Different departments want different detail levels from the same weekly report. BEST design response? [Workflow]

- A. Force one output
- B. Create audience-specific output requirements while keeping the underlying process consistent
- C. Let each user rewrite manually
- D. Claude drafts only after employees do everything

Correct Answer: B

Easy Explanation: Process ek, output har audience ke liye mukhtalif.

Q53. A manager proposes Claude automatically decide the supplier watch list, currently decided by analyst research and a management review. BEST first design? [Workflow]

- A. Claude decides; analysts check its picks
- B. Remove the review meeting
- C. Avoid Claude entirely
- D. Claude summarizes evidence and surfaces risk indicators; analysts and managers keep the decision

Correct Answer: D

Easy Explanation: AI saboot jama kare, faisla insaan kare.

Q54. A manager says Claude will 'eliminate all errors and remove the need for review'. BEST response from the analyst? [Workflow]

- A. Agree for adoption
- B. Avoid discussing limits until later
- C. Explain expected time savings while stating outputs still need validation appropriate to the task
- D. Reject the project

Correct Answer: C

Easy Explanation: Faida bhi batao aur hudood bhi; overclaim na karo.

Q55. Analysts use a media tool that collects articles efficiently. Which integration is MOST sensible? [Workflow]

- A. Keep the tool for collection; Claude synthesizes collected material for analyst review
- B. Replace the tool
- C. Claude only rewrites the final brief
- D. Stop collecting and let Claude imagine themes

Correct Answer: A

Easy Explanation: Jo theek chal raha use na chheero; AI wahan lagao jahan faida ho.

Q56. Analysts prompt Claude differently for monthly commentary and outputs are inconsistent. Which improvement is BEST? [Workflow]

- A. Keep personal prompts
- B. Move to the most capable model
- C. Define a shared workflow with common inputs, instructions, output expectations and review steps
- D. Stop using Claude

Correct Answer: C

Easy Explanation: Variance ka hal: mushtarka workflow/configuration.

Q57. Outputs got worse after many one-off exceptions were added to shared instructions. What to troubleshoot FIRST? [Troubleshooting]

- A. Add more exceptions
- B. Use a cheaper model
- C. Simplify shared instructions and move temporary exceptions out of persistent configuration
- D. Replace with research mode

Correct Answer: C

Easy Explanation: Arzi exceptions mustaqil configuration mein na hon.

Q58. Claude keeps adding speculative explanations when summarizing incidents; source notes contain only confirmed facts. BEST prompt-level fix? [Troubleshooting]

- A. Require separating confirmed facts from unknown causes and prohibit unsupported inference
- B. Upload unrelated reports
- C. Use the highest model
- D. Reviewers delete speculation

Correct Answer: A

Easy Explanation: Wazeh hidayat: tasdeeq shuda aur na-maloom alag, andaza mana.

Q59. Which measure is MOST useful to know whether a Claude workflow improved after a redesign? [Troubleshooting]

- A. Tokens in the new prompt
- B. A defined combination of output quality, reviewer effort, and completion time
- C. How advanced it feels
- D. Longer responses

Correct Answer: B

Easy Explanation: Asal nateeja napo: quality, review effort, waqt.

Q60. Reviewers check every sentence of monthly reports though only a few sections are high consequence. BEST optimization? [Troubleshooting]

- A. Remove review
- B. Match review intensity to risk, keeping stronger review for consequential sections
- C. Claude reviews itself
- D. Identical review everywhere

Correct Answer: B

Easy Explanation: Review ki shiddat risk ke mutabiq.

Q61. A pilot works for normal cases but fails when source documents are missing sections. BEST optimization? [Troubleshooting]

- A. Ignore them
- B. Invent defaults
- C. Stop using Claude
- D. Add an exception path that detects incomplete inputs and routes them for handling

Correct Answer: D

Easy Explanation: Adhoore input ke liye alag raasta.

Q62. Only 2% of first-pass classifications need correction, but the team reclassifies every item manually. BEST optimization? [Troubleshooting]

- A. Stop all review
- B. Claude corrects itself
- C. Keep full reclassification
- D. Redesign review around sampling or exception checks while monitoring quality

Correct Answer: D

Easy Explanation: Sampling + monitoring; review khatam nahi.

Q63. Reviewers spend 20 minutes because Claude includes raw source excerpts under every finding, rarely needed. BEST optimization? [Troubleshooting]

- A. Keep all excerpts
- B. Reviewers collapse them
- C. Use a faster model
- D. Include concise evidence references by default and full excerpts only when needed

Correct Answer: D

Easy Explanation: Zaroorat ke mutabiq output.

Q64. A head asks: 'brief but comprehensive five-page overview with all details in one paragraph.' Result is inconsistent. BEST improvement? [Troubleshooting]

- A. Resolve the conflicting requirements by specifying the real priority, length and structure
- B. Say all are mandatory
- C. Add more adjectives
- D. Let Claude pick the most reasonable one

Correct Answer: A

Easy Explanation: Aapas mein takrati hidayat ko khud wazeh karo.

Q65. Claude gives a shallow analysis when asked 'Tell me what matters' about a complete long report. BEST troubleshooting step? [Troubleshooting]

- A. Break the analysis into stages and define what 'matters' means for the decision
- B. Upload again
- C. Switch to Opus first
- D. Ask for twice the words

Correct Answer: A

Easy Explanation: 'Ahem' ki tareef karo aur marhalon mein toro.

Q66. Users spend 15 minutes reformatting every weekly output into the same table. BEST optimization? [Troubleshooting]

- A. Keep it, drafting improved
- B. Build copy-paste shortcuts
- C. Use a more capable model
- D. Make the required table structure part of the recurring workflow output

Correct Answer: D

Easy Explanation: Baar baar ki correction = configuration mein daalo.

Q67. Win/loss notes: Claude says customers 'consistently prefer lower prices', but reasons are mixed and a minority mention price. Conclusion? [Output Evaluation]

- A. It likely overgeneralizes and should be checked against the interview set
- B. Acceptable, price is common
- C. Change 'consistently' to 'often' without review
- D. Ignore qualitative notes

Correct Answer: A

Easy Explanation: Overgeneralization: source se milao.

Q68. A comparison table shows one vendor's support as 24/7; the proposal says 8am-8pm weekdays. Rest is correct. What to do? [Output Evaluation]

- A. Correct the field and recheck other high-impact fields against the proposals
- B. Keep the table
- C. Ask Claude where 24/7 came from
- D. Remove the row

Correct Answer: A

Easy Explanation: Ek ghalti mile to baqi ahem fields bhi dobara check.

Q69. A weekly project-status table will be imported into another tool. MOST important check before import? [Output Evaluation]

- A. Varied wording
- B. Required fields, allowed status values and row-to-project mappings match the expected structure
- C. Visual attractiveness
- D. Claude can explain each status

Correct Answer: B

Easy Explanation: Structured output: schema, allowed values, mapping.

Q70. Claude recommends Target A but cites a growth figure that belongs to Target B. What problem is this? [Output Evaluation]

- A. Formatting
- B. Bias
- C. An internal inconsistency that undermines the reasoning
- D. Completeness

Correct Answer: C

Easy Explanation: Daleel hi ghalat to recommendation kamzor.

Q71. A summary of an internal control standard will train employees on a mandatory procedure. Review level? [Output Evaluation]

- A. None if clear
- B. Claude reviews itself
- C. An authorized reviewer validates it against the approved standard before training use
- D. Publish and fix later

Correct Answer: C

Easy Explanation: Lazmi procedure = authorized reviewer.

Q72. Claude says fulfillment improved 18%; the report shows regional improvements but no company-wide percentage. What to do? [Output Evaluation]

- A. Use it
- B. Replace with 'improved significantly'
- C. Ask Claude's confidence
- D. Verify whether 18% is supported by the source; remove or correct if not

Correct Answer: D

Easy Explanation: Baghair source ka exact number: tasdeeq ya hatao.

Q73. Policy allows exceptions only with written approval from the policy owner; Claude's summary says managers may approve whenever business needs justify. What occurred?

[Output Evaluation]

- A. An unsupported rule that materially changes the policy's meaning
- B. Simplification
- C. Improvement
- D. Style issue

Correct Answer: A

Easy Explanation: Maani badal dena sangeen ghalti hai.

Q74. Board update says a supplier missed the SLA 'three consecutive months'; source shows two confirmed misses and one under review. BEST revision? [Output Evaluation]

- A. Keep it
- B. State two confirmed misses and one month under review
- C. Say 'repeatedly'
- D. Remove the issue

Correct Answer: B

Easy Explanation: Jitna saabit utna hi likho.

Q75. Claude found three differences between a procedure and a standard; a human finds a fourth, more consequential one. Next? [Output Evaluation]

- A. Keep Claude's comparison
- B. Add the fourth and treat it as evidence the comparison needs rechecking for completeness
- C. Ignore unless it ranks higher
- D. Replace with only the fourth

Correct Answer: B

Easy Explanation: Ek choot mile to completeness dobara check.

Q76. A checklist is complete and accurate but steps are out of the required sequence. Evaluation? [Output Evaluation]

- A. Accept
- B. Style issue
- C. A correctness problem when order is required for safe or valid execution
- D. Shorten it

Correct Answer: C

Easy Explanation: Jahan tarteeb lazmi ho wahan tarteeb durusti hai.

Q77. Survey: 92 support, 78 oppose, 30 neutral out of 200. Claude says 'Most employees support the change.' Evaluation? [Output Evaluation]

- A. Inaccurate: fewer than half expressed support
- B. Accurate
- C. Acceptable if 'most' means largest group
- D. Neutrals don't count

Correct Answer: A

Easy Explanation: $92/200 = 46\%$, 'most' ghalat.

Q78. Top-5 audit findings exclude the only finding rated 'critical'. BEST evaluation? [Output Evaluation]

- A. Valid
- B. Incomplete for the stated purpose because it omitted a clearly material finding
- C. Needs shorter explanation
- D. Ask for ten

Correct Answer: B

Easy Explanation: Sab se ahem cheez chhoot jaye to output adhoora.

Q79. A survey summary highlights one department's workload complaints but barely mentions similar comments from two others. FIRST step? [Output Evaluation]

- A. Compare against the full comment set to see whether perspectives were underrepresented
- B. Make it positive
- C. Keep it
- D. Remove department references

Correct Answer: A

Easy Explanation: Bias ka shak: poore data se milao.

Q80. A variance table has correct numbers but inconsistent department names, breaking downstream matching. What to do? [Output Evaluation]

- A. Accept
- B. Convert to prose
- C. Normalize labels to approved names and validate the structured output before downstream use
- D. Let users interpret

Correct Answer: C

Easy Explanation: Structured data ke naam bhi durust hon.

Q81. A sensitive company-wide message on major change; nuance matters more than cost or speed; a human approves. MOST appropriate model? [Model Selection]

- A. Haiku, human review reduces need
- B. Sonnet for all drafting
- C. Haiku for speed
- D. Opus, maximum quality and nuance

Correct Answer: D

Easy Explanation: Quality sab se ahem = Opus.

Q82. Thousands of straightforward, low-risk comments need first-pass categories; staff sample-check; speed and cost matter. BEST fit? [Model Selection]

- A. Haiku
- B. Sonnet
- C. Opus
- D. Avoid Claude

Correct Answer: A

Easy Explanation: Bulk, sasta, tez = Haiku.

Q83. Daily routine alert summaries plus a weekly synthesis of ambiguous signals into an executive outlook. BEST approach? [Model Selection]

- A. Opus for everything
- B. Haiku for routine alerts and a more capable model for the weekly synthesis
- C. Haiku for all
- D. Sonnet for all

Correct Answer: B

Easy Explanation: Kaam ke hisaab se model.

Q84. A recurring low-risk internal analysis needs solid quality, but cost and responsiveness matter. Choice? [Model Selection]

- A. Sonnet, balanced quality, cost and responsiveness
- B. Opus
- C. Haiku
- D. Opus for subtle tradeoffs

Correct Answer: A

Easy Explanation: Darmiyana default = Sonnet.

Q85. A researcher needs a current overview, then a polished briefing revised several times. BEST sequence? [Model Selection]

- A. Artifact then research mode
- B. Research mode for current evidence, then an Artifact for iterative refinement
- C. A Project for everything
- D. Ordinary chat and restart each draft

Correct Answer: B

Easy Explanation: Taaza saboot = research; baar baar edit = artifact.

Q86. A monthly report reuses the same internal documents but one section needs current industry developments. BEST setup? [Model Selection]

- A. Only research mode
- B. A Project for reusable internal context and research mode when current external evidence is needed
- C. Only a Project
- D. An Artifact as store and research tool

Correct Answer: B

Easy Explanation: Project + research, dono apni jagah.

Q87. A business analyst needs one quick explanation of a chart already in chat, no follow-up. BEST choice? [Model Selection]

- A. Research mode
- B. Create a Project
- C. Create an Artifact
- D. Continue in chat: simple, self-contained, one-time

Correct Answer: D

Easy Explanation: Chhota ek dafa ka kaam = chat.

Q88. A long policy communication will be refined over several rounds, viewed separately from chat; no reusable knowledge. BEST feature? [Model Selection]

- A. Research mode
- B. A Project
- C. An Artifact alongside the conversation
- D. A fresh chat each revision

Correct Answer: C

Easy Explanation: Barhta hua draft = Artifact.

Q89. A Project has a stable glossary and process guide plus a threshold spreadsheet that changes monthly. BEST maintenance? [Configuration]

- A. Replace everything monthly
- B. Leave old thresholds
- C. Restate thresholds manually
- D. Update the threshold source on schedule; keep stable sources as needed

Correct Answer: D

Easy Explanation: Badalne wali file waqt par update, mustaqil wali waisi.

Q90. Users repeatedly add the same three approved exclusions. BEST improvement? [Configuration]

- A. Add them to persistent Project instructions
- B. Keep them manual
- C. Store in an old report
- D. Add at end of every prompt

Correct Answer: A

Easy Explanation: Baar baar ka usool = instructions.

Q91. Output changed from narrative to a five-column table, but the Project still says write paragraphs. BEST update? [Configuration]

- A. Convert manually
- B. Ask for a table every chat
- C. Update persistent instructions with the new table structure
- D. Stronger model

Correct Answer: C

Easy Explanation: Stale configuration update karo.

Q92. A Project contains a current leave policy and a superseded one; Claude mixes them. FIRST action? [Configuration]

- A. Tell Claude to prefer the newest
- B. Remove or archive the obsolete version
- C. Switch to Opus
- D. Users type the date each time

Correct Answer: B

Easy Explanation: Purana version retire karo.

Q93. A Project instruction says 'always be complete' and another says 'never exceed 100 words'. BEST improvement? [Configuration]

- A. Users decide
- B. Add 'use best judgment'
- C. Clarify priority and conditions so rules no longer conflict
- D. Stronger model

Correct Answer: C

Easy Explanation: Takrao khud hal karo.

Q94. Stable brand guidance plus launch details that change every week. BEST setup? [Configuration]

- A. Store every launch detail permanently
- B. Brand rules in the Project; temporary launch facts per task
- C. Restate tone every prompt
- D. One Project per announcement

Correct Answer: B

Easy Explanation: Mustaqil Project mein, arzi chat mein.

Q95. Recurring instructions say only 'make it useful' and 'use the right format'; outputs vary. BEST improvement? [Configuration]

- A. Keep short
- B. Replace with explicit audience, structure, content and uncertainty rules
- C. Add more documents
- D. Stronger model

Correct Answer: B

Easy Explanation: Mubham hidayat ko wazeh karo.

Q96. A shared Project holds one analyst's personal shorthand that conflicts with the approved team standard. BEST action? [Configuration]

- A. Keep both
- B. Delete all shared instructions
- C. Add more blocks
- D. Keep the approved shared standard and remove conflicting personal guidance

Correct Answer: D

Easy Explanation: Team standard ghalib.

Q97. A marketing Project has a revised style guide but old sample copy in the previous tone. What to do? [Configuration]

- A. Keep samples
- B. Delete the style guide
- C. Update or replace examples that conflict with the current guide
- D. Tell users to ignore

Correct Answer: C

Easy Explanation: Examples bhi hidayat ki tarah asar karti hain.

Q98. Policy lets Claude summarize approved legal guidance but requires counsel for individualized advice. A manager wants legal advice for an employee sent without review. What should happen? [Governance]

- A. Answer directly with guidance in context
- B. Employee decides
- C. Give a confidence score
- D. Use Claude only to summarize approved guidance and escalate individualized advice to counsel

Correct Answer: D

Easy Explanation: Policy ki had mein raho.

Q99. A client has not approved AI use for confidential materials; prompts drafted, nothing uploaded. Next? [Governance]

- A. Proceed
- B. Upload partial documents
- C. Ask Claude to sign a confidentiality statement
- D. Pause until the client authorizes an appropriate environment or use

Correct Answer: D

Easy Explanation: Ijazat ke baghair nahi.

Q100. A recruiter wants to rank applicants using disclosed medical information; policy prohibits it. What to do? [Governance]

- A. Exclude it and use only approved hiring criteria
- B. Use only for ties
- C. Use but omit from explanation
- D. Use if a human reviews

Correct Answer: A

Easy Explanation: Mamnoo data kisi soorat nahi.

Q101. Policy allows trend analysis only after regulated identifiers are removed from a customer spreadsheet. What should the employee do? [Governance]

- A. Upload since it stays internal
- B. Upload and say don't remember
- C. Skip the analysis
- D. Remove or anonymize the identifiers before using the approved workflow

Correct Answer: D

Easy Explanation: Identifiers hatao, phir approved raasta.

Q102. An employee wants to paste an API key to troubleshoot; policy prohibits credentials in prompts. What to do? [Governance]

- A. Remove last four characters
- B. Keep credentials out and use a safe redacted example
- C. Ask Claude to delete it
- D. Mark chat confidential

Correct Answer: B

Easy Explanation: Credentials kabhi prompt mein nahi.

Q103. Marketing wants Claude to generate fictional testimonials shown as if written by customers. MOST responsible response? [Governance]

- A. Generate, copy is fiction
- B. Generate without names
- C. Do not present fabricated testimonials as real customer experiences
- D. Use if manager approves

Correct Answer: C

Easy Explanation: Jhoot ko sach bana kar pesh karna ghalat.

Q104. A public statement about an incident: facts confirmed, cause not established. MOST responsible approach? [Governance]

- A. Infer the likely cause
- B. Limit to confirmed facts and clearly identify unresolved questions
- C. Publish a likely cause labelled as Claude's
- D. Delay all communication

Correct Answer: B

Easy Explanation: Sirf tasdeeq shuda baat, baqi sawal khule batao.

Q105. Claude occasionally includes internal information in drafts for external audiences. BEST response? [Governance]

- A. Keep publishing
- B. Review the workflow, source scope and human approval controls before continuing external publication
- C. Ask Claude to promise
- D. Switch models

Correct Answer: B

Easy Explanation: Control aur scope theek karo.

Q106. A support team wants Claude to draft account-dispute replies; policy requires human approval of any credit or refund promise. BEST workflow? [Governance]

- A.** Send automatically
- B.** Claude drafts; route any credit or refund commitment for required human approval
- C.** Avoid Claude
- D.** Claude approves small refunds

Correct Answer: B

Easy Explanation: Paise ka wada = insaan ki manzoori.

Q107. A manager asks Claude to assess employees' emotional stability from writing samples; no approved process or HR authorization. What to do? [Governance]

- A.** Proceed
- B.** Use as informal signal
- C.** Do not use Claude for it without an approved process and HR authorization
- D.** Anonymize first

Correct Answer: C

Easy Explanation: Ghair manzoor assessment nahi.

Section D: New Questions Built From the 10 Crash Courses (25 questions)

Written from the course concepts (delegation map, four causes, data tiers, file tiers, skills, 4Ds) in exam style.

Q108. A team's expense workflow has a step that totals the report and computes the amount over policy limits; an approver reviews next. How should that step be classified and carried? [Workflow]

- A. Human-retained because it decides reimbursement
- B. AI-appropriate (reversible, checked at approval), carried by code execution
- C. Collaborative with a human re-adding every line
- D. AI-appropriate, carried by a Skill

Correct Answer: B

Easy Explanation: Ownership 3 criteria se, implementation (code) alag faisla.

Q109. A proposal team's drafts have been strong for four months, so the lead proposes AI send low-value responses directly. Which mapping error is this? [Workflow]

- A. Mapping the tool instead of the work
- B. Unstaffed gate
- C. Halo delegation: sending judged by how well drafting went
- D. None, four months is good evidence

Correct Answer: C

Easy Explanation: Draft ki achhai bhejne ke faislay ki daleel nahi.

Q110. A refund step is 'AI drafts, manager reviews'. After six weeks the manager batch-approves every Friday without opening cases. The step is now: [Workflow]

- A. Collaborative as designed
- B. Automated in practice: a collaborative step with no real reviewer is automated
- C. Human-retained
- D. Correctly optimised

Correct Answer: B

Easy Explanation: Unstaffed gate.

Q111. A coordinator drafts a supplier shortlist which the procurement manager reviews line by line before contact. Someone says drafting must stay human because the manager is accountable. Is that right? [Workflow]

- A. Yes, any input to her decision must be human
- B. Yes, high stakes rule out delegation
- C. No: accountability lives at the step that is the decision; drafting is an input she judges
- D. No: accountability moved to the tool

Correct Answer: C

Easy Explanation: Jawabdehi faislay wale qadam par hai.

Q112. An analyst types 'exclude cancelled orders' into chat every month. On leave, a colleague runs the report and cancelled orders appear. What failed? [Troubleshooting]

- A. The reminder was vague
- B. The colleague ignored notes
- C. Ordinary human error
- D. The fix was never promoted out of the conversation into configuration

Correct Answer: D

Easy Explanation: Fix ko standing instruction mein daalna tha.

Q113. Three coordinators each write consistent weekly updates, but the director spends an hour every Monday unifying their voice. Which friction signal and fix? [Troubleshooting]

- A. Correction: standing instruction
- B. Repetition: saved context
- C. Variance: a shared Skill so all produce one format
- D. None

Correct Answer: C

Easy Explanation: Farq logon ke darmiyan = variance.

Q114. A weekly report 'used to be right' but now quotes last quarter's figures with perfect formatting. No error appears. Most likely cause? [Troubleshooting]

- A. Under-specification
- B. Context overload
- C. Wrong model tier
- D. Stale configuration: an outdated knowledge file

Correct Answer: D

Easy Explanation: 'Pehle theek tha' = stale config, khamoshi se.

Q115. A fix consists of five steps that must run in a fixed order. Where does it belong?

[Configuration]

- A. A one-line standing instruction
- B. The knowledge base
- C. A Skill (procedure)
- D. Memory

Correct Answer: C

Easy Explanation: Procedure = Skill; ek line mein order kho jata hai.

Q116. A user asks Claude to predict next quarter's exact sales figure; the answer was wrong. After ruling out prompt, context, feature and configuration, what is the diagnosis and fix? [Troubleshooting]

- A. Wrong model; use Opus
- B. Expectation mismatch: ask for a range with stated assumptions or an adjustable driver model
- C. Under-specification
- D. Stale configuration

Correct Answer: B

Easy Explanation: Na-mumkin kaam: task ki shakal badlo.

Q117. Which requirement would lead two teams to build different things? [Workflow]

- A. Report delivered promptly after month-end
- B. Report in both PDF and CSV from the same data
- C. Header shows reporting period and generation date
- D. Password-reset email within five minutes

Correct Answer: A

Easy Explanation: 'Promptly' ka check nahi likha ja sakta.

Q118. A Cowork session produced a finished client invoice PDF saved only in platform storage. What does the file-tier rule say? [Configuration]

- A. It is fine: tier 2 is permanent
- B. Finished work must exit the platform to a system you control (connector save, download, email)
- C. Leave it in the task filesystem
- D. Delete it after reading

Correct Answer: B

Easy Explanation: Mukammal kaam tier 3 (aap ke record) mein jaye.

Q119. You keep re-explaining HOW to format a weekly report, and you also keep copy-pasting data from Google Drive. What do you need? [Configuration]

- A. Only a Skill
- B. Only a Connector
- C. Both a Skill (method) and a Connector (data access)
- D. A longer prompt

Correct Answer: C

Easy Explanation: Tareeqa = Skill, app ka data = Connector.

Q120. You built a custom skill but it never triggers when you ask for the task. Most likely cause? [Configuration]

- A. The model is too small
- B. The skill's description does not match how you phrase requests
- C. Skills only work via slash command
- D. Memory is off

Correct Answer: B

Easy Explanation: Skill description se chalti hai.

Q121. Before enabling a new third-party connector, what is the safest starting setup? [Governance]

- A. Full write access so it is useful
- B. Read-only with tight scope, widened deliberately later
- C. Any scope since it is approved
- D. Install via a Skill

Correct Answer: B

Easy Explanation: Pehle read-only, kam scope.

Q122. In the 4D framework, 'Is working this way with AI actually paying off, or am I repairing the same formatting every turn?' is which kind of discernment? [Output Evaluation]

- A. Product discernment
- B. Process discernment
- C. Performance discernment
- D. Creation diligence

Correct Answer: B

Easy Explanation: Session ka tareeqa parakhna = process discernment.

Q123. A lecturer pasted student names, grades and disciplinary notes into an unapproved consumer AI; comments were excellent and students were not told. Which competency failed? [Governance]

- A. Description
- B. Discernment
- C. Diligence (creation and transparency)
- D. Delegation only

Correct Answer: C

Easy Explanation: Achha output bhi ghair zimmedarana istemal ko sahi nahi banata.

Q124. Which description earns the most trust from a client's risk function? [Workflow]

- A. 'Our system reviews contracts automatically, cutting turnaround by half.'
- B. 'Drafting is AI-assisted. A qualified lawyer reviews and approves every term, and nothing is sent without that sign-off.'
- C. 'It extracts clauses with a Skill and computes exposure in code.'
- D. 'It is as reliable as a junior associate.'

Correct Answer: B

Easy Explanation: Risk team ko control chahiye: human gate wazeh.

Q125. A vendor comparison: you paste three proposals and ask for a recommendation in one go. What does process description advise instead? [Prompting]

- A. Ask for three recommendations and pick one
- B. Split into steps (extract to a table, compare, score, draft) and check each step, extraction first
- C. Use extended thinking only
- D. Ask Claude to rate its confidence

Correct Answer: B

Easy Explanation: Qadam ba qadam, pehle extraction check.

Q126. A department head says 'compare the old plan with the new one' but the files contain several plans and revisions; Claude compares the wrong versions. BEST correction? [Prompting]

- A. Ask Claude to guess which are old/new
- B. Remove dates
- C. Name the exact two versions and the specific changes to address
- D. Ask for a longer comparison

Correct Answer: C

Easy Explanation: Mubham hawala = wazeh naam do.

Q127. Sales notes are missing renewal dates for several customers. Which prompt approach is BEST for account plans? [Prompting]

- A. Estimate dates and label assumptions
- B. Omit fields silently
- C. Use supplied facts, mark missing dates as unknown, and list information needed before finalizing
- D. Generate likely dates

Correct Answer: C

Easy Explanation: Kami ko 'unknown' likho, bharo mat.

Q128. JSON exception records come back with different field names each run. Change MOST likely to improve consistency? [Prompting]

- A. Say the output must be reliable
- B. Provide exact field names, data types and a short example of the required JSON
- C. Ask it to review before answering
- D. Explain in prose first

Correct Answer: B

Easy Explanation: Schema + example.

Q129. A spreadsheet of customer names needs only spending trends. Policy restricts personal data. What is the right move? [Governance]

- A. Upload as is
- B. Upload with 'do not retain'
- C. Strip identifiers, confirm no mapping travels with the file, then analyse
- D. Skip the analysis

Correct Answer: C

Easy Explanation: Pattern chahiye, log nahi.

Q130. Which statement about Claude incognito chats is accurate? [Governance]

- A. Nothing is ever retained
- B. They skip history and memory but still follow a retention period; they are not permission to upload restricted data
- C. They are available inside Projects
- D. They are used for training on all plans

Correct Answer: B

Easy Explanation: Incognito = history/memory se bahar, ijazat nahi.

Q131. A script reconciles bank data. Before trusting it on live data, which habit from 'Code You Never Write' is MOST useful? [Output Evaluation]

- A. Ask the AI if the code is correct
- B. Run a known-answer test on a sample whose result you already know
- C. Read every line of Python
- D. Use the biggest model

Correct Answer: B

Easy Explanation: Maloom jawab wala test.

Q132. A desktop agent will rename and delete files in a folder. Which safety sequence is correct? [Governance]

- A.** Run on the whole drive to save time
- B.** Work on copies, dry run before anything is destroyed, smallest folder, output to new files
- C.** Give full permissions then review
- D.** Let it run unattended overnight

Correct Answer: B

Easy Explanation: Pehle copy, dry run, chhota folder.