

How to use this set

Questions are grouped by the seven CCAO-F domains, then a section of long case-study sets where one scenario carries two or three questions (the exam does this). Multiple-response items state how many to select. Attempt a whole section under time (about 2 minutes per item), then check the Answer Key at the end, which explains each answer in English and Roman Urdu.

Har section time laga kar hal karein (har sawal ~2 minute), phir aakhir mein Answer Key dekhein jahan har jawab ki wazahat English aur Roman Urdu mein hai.

Domain 1: Prompting and Task Execution (14%)

Prompting aur kaam karwana

Q1. A facilities associate attaches warehouse inspection notes and asks Claude to "write up the inspection." The result is a tidy general narrative, useless to the insurer's loss adjuster, who needs only the three open defects with the date each was found and its photo reference. A colleague says the fix is a much longer prompt. What is the most appropriate way to improve the request?

- A. Describe the warehouse and the inspection day in full detail so Claude has more context.
- B. State the reader, the purpose, the length, the items to cover, and what to leave out.
- C. Rerun on Opus with extended thinking so the model infers what the adjuster needs.
- D. Attach last year's internal inspection narrative so Claude copies its structure.

Q2. A hotel-group associate delegates: shortlist three venues with on-site AV for a 120-person offsite in the second week of March, in a table with links, under a stated budget. In the result two venues tied on price and Claude silently dropped one, and it widened the dates to "late March" to find a third option without saying so. Which element of the request would have prevented these silent choices?

- A. An instruction to show the plan of steps before starting.
- B. More detail about the event audience and purpose.
- C. A line requiring Claude to list every assumption after finishing.
- D. A line saying what Claude may decide alone and that it must say when it chose.

Q3. A finance analyst asks Claude to extract invoice exceptions into JSON for a downstream system, but field names differ on every run. Which change is MOST likely to improve consistency?

- A. Explain that the JSON feeds a business process, so it must be reliable.
- B. Provide exact field names, data types and a short example of the required structure.
- C. Ask Claude to review its planned JSON carefully before answering.
- D. Request a prose explanation of each exception first, then convert to JSON.

Q4. A Project instruction says "Flag unusual expenses in the monthly report." Each run flags different items and two colleagues running the same data get different results. What is the most appropriate fix?

- A. Add "be careful and consistent" to the instruction.
- B. Switch the Project to a more capable model.
- C. Rewrite the rule with a named amount threshold and expense category.
- D. Add three flagged examples from last quarter and let the model infer the threshold.

Q5. A procurement manager's first draft request reads "Tell me this shortlist is defensible." Which TWO rewrites are most likely to produce an honest evaluation? (Select TWO)

- A. "Rate the shortlist 1-10 on coverage, cost realism and risk, with the one change that most improves it."
- B. "Defend the shortlist against objections finance is likely to raise."
- C. "Give the strongest case for and against sending the shortlist to the committee as it stands."
- D. "Isn't this good enough? Point out anything minor to tidy up."
- E. "Rate your confidence as a percentage that it survives the committee."

Q6. A sales manager wants account plans, but several customers' renewal dates are missing from the notes. Which prompt approach is BEST?

- A. Estimate missing dates from account history and label them as assumptions.
- B. Build complete plans and silently omit unsupported fields.
- C. Use supplied facts, mark missing renewal dates as unknown, and list what is needed before finalising.
- D. Generate likely dates from similar customers, labelled low confidence.

Q7. A strategy associate needs a proposal to redesign the returns process and wants to ask for the full proposal straight away with extended thinking on. Which first request would most improve the eventual proposal?

- A. Ask for the complete proposal now and iterate.
- B. Ask Claude first to map what is known - existing approaches, other firms' fixes, prior attempts.
- C. Ask for three complete proposals in different styles and pick one.
- D. Ask Claude to grade your draft outline out of ten.

Q8. A product leader has fixed evaluation criteria but no solution ideas for improving a customer portal. Which approach is BEST?

- A. Use the criteria immediately to pick the one improvement most likely to succeed.
- B. Infer the ideal solution type from the criteria and propose one.
- C. Generate one conservative and one ambitious idea and compare them.
- D. Brainstorm a diverse unranked set first, then evaluate the shortlist against the criteria in a separate step.

Domain 2: Output Evaluation and Validation (21%)

Output ko parakhna aur verify karna

Q9. A hospital pharmacy manager asks Claude for a script that cross-checks the dispensing log against weekly stock counts. She plants one fake gap (ten units removed from the log with no stock movement). The script flags her planted gap plus four others. She tells the director the script is "proven" and the four extra flags can be actioned. She then recalls some past discrepancies came from ward transfers being counted twice. What does catching the planted fault entitle her to conclude?

- A. The script is reliable for all discrepancy types.
- B. Nothing, because a planted fault is artificial.
- C. It supports flags of the missing-units class; double-counted transfers are another class that needs its own planted fault.
- D. She may action the four extra flags now.

Q10. A controller gives the identical written brief and the same two files to two AI assistants from different companies. Both return the same 23 mismatched invoices. Her deputy says "independent confirmation, case closed." How much assurance does the agreement give?

- A. Weak, because both share training data.
- B. Adequate once each rates its confidence.
- C. Strong, because two vendors will not invent the same mismatches.
- D. Useful but bounded: both read the same brief, so an omitted rule such as netting credit notes is an error both can make.

Q11. Claude recommends Supplier B and lists five assumptions. One reads "Supplier B guarantees delivery within 48 hours." None of the three attached proposals contains a 48-hour clause. Which action correctly uses the listed assumptions?

- A. Open the extended-thinking display to confirm the model's true reasoning.
- B. Ask Claude whether it is sure about the 48-hour clause.
- C. Accept it because the chain of assumptions is coherent.
- D. Treat the list as a justification for review, check each assumption against the proposals; the unsourced claim means the recommendation cannot stand.

Q12. Claude runs code on a loan-book extract and reports PKR 4.31 billion outstanding principal. The branch manager's hand review last week expected about PKR 3.94 billion. What is the best next step?

- A. Accept it because executed code is exact.
- B. Rerun on the highest model tier.
- C. Check what was computed: which column and rows, with row counts and samples.
- D. Ask for the total in prose without code.

Q13. Claude summarises a survey: "Most employees support the proposed schedule change." Results: 200 responses, 92 support, 78 oppose, 30 neutral. What is the BEST evaluation?

- A. Inaccurate: fewer than half of respondents support.
- B. Accurate: support outnumbered opposition.
- C. Acceptable if "most" means the largest group.
- D. Retain it; neutral responses do not count.

Q14. An associate reading an AI-generated market summary before repeating it to a client. Which THREE features suggest a claim may be fabricated and should be checked? (Select THREE)

- A. A statistic to two decimals with no source.
- B. Hedged phrasing such as "may vary by region".
- C. A statement that the assistant "checked the supplier's website" with no page or quote.
- D. A summary longer than expected.
- E. A figure on page one contradicting the same figure near the end.
- F. Headings that match the usual template.

Q15. Claude picks "the five most important findings" from an audit report. All five are genuine, but it excluded the only finding rated critical. What is the BEST evaluation?

- A. Valid, since all five came from the report.
- B. Incomplete for the stated purpose because it omitted a clearly material finding.
- C. Needs a shorter explanation.
- D. Increase to ten and accept.

Q16. An insurer's claims team has Claude pre-screen 900 warranty claims and route 120 for fast-track. The claims lead reviews all 120 carefully, finds them sound, and prepares to approve. What is the most important review step still missing?

- A. Note on each file that AI did first screening.
- B. Tell the 780 claimants AI was used.
- C. Sample the 780 not fast-tracked to check whether a group was routed away systematically.
- D. Rerun with a stricter prompt and approve only overlaps.

Q17. A legal-operations lead decides which AI-assisted outputs need full four-step verification and a named reviewer. Outputs: (1) internal recap of yesterday's call; (2) a demurrage-exposure figure going to a counterparty; (3) first-pass tags on incoming emails; (4) a draft reply to a regulator's information request, signed by a director. Which TWO cross the threshold? (Select TWO)

- A. Output 3
- B. Output 1
- C. Output 4
- D. Output 2

Q18. A junior analyst with no customs background drafts a customs-compliance briefing with Claude. It meets every requirement and matches the uploaded regulation extract. Which check most directly addresses the remaining gap?

- A. Ask a colleague with customs experience whether it is credible by the field's standards.
- B. Ask a colleague to proofread tone and clarity.
- C. Ask Claude to rate its accuracy and revise below 9/10.
- D. Run a formatting pass.

Domain 3: Product and Model Selection (12%)

Product feature aur model ka chunao

Q19. A support team runs every meeting-notes tidy-up on Opus with extended thinking, and runs its annual pricing-model rebuild on Haiku "to save budget." Which statement best describes the problem?

- A. Only the pricing choice is wrong.
- B. Neither is wrong; quality differences are marginal.
- C. Only the tidy-up choice is wrong.
- D. Both are wrong in opposite directions: the flagship is wasted on trivial work and the cheapest rung wastes the rebuild.

Q20. A legal-ops associate uses Sonnet with extended thinking to analyse a 180-page services agreement for cross-referenced obligations. The same request worked on a 20-page contract; here the answer is fluent but shallow. Which change best fits?

- A. Move to Opus, built for holding a long, complex structure coherent at once.
- B. Switch to Haiku for speed.
- C. Turn extended thinking off.
- D. Split the agreement into ten separate chats.

Q21. An associate has six requests. Which THREE pair the request with the right feature? (Select THREE)

- A. "What did the port authority announce this morning?" -> web search
- B. "Pull last quarter's rate card from the finance shared drive" -> web search
- C. "Which lease in this thread is riskier to exit early?" -> extended thinking
- D. "Today's VAT on imported machinery" -> research mode
- E. "Cited comparison of East African cold-chain providers with funding and customers" -> research mode
- F. "Rewrite a three-line status update friendlier" -> extended thinking

Q22. An associate proposes running supplier-pricing chats in Claude incognito and ChatGPT temporary chat "so there is nothing on record if audit asks." Which statement is accurate?

- A. Both discard everything immediately.
- B. Both keep chats out of history and memory, but each vendor retains a copy for a period and company records may hold it.
- C. Only Claude incognito leaves no trace.
- D. Neither changes anything; both go to history.

Q23. In one session: explain "demurrage"; compare nine customs brokers that finance will load into a sheet; a one-page onboarding checklist for new hires. Which THREE format choices are correct? (Select THREE)

- A. Broker comparison as a table or spreadsheet file
- B. Demurrage explanation inline in chat
- C. Checklist inline for hires to copy out
- D. Checklist as a document artifact
- E. Broker comparison as a prose paragraph
- F. Demurrage explanation as a document artifact

Q24. A consultant has a clean, focused, short chat with current requirements for a one-hour analysis. Nothing is outdated. What should they do next?

- A. Continue in the current chat.
- B. Restart; fresh context is always safer.
- C. Create a Project.
- D. Summarise the conversation first.

Domain 4: Workflow Integration and Solution Design (16%)

Workflow mein AI fit karna

Q25. A reporting team plans to consolidate three prompts into one and move format rules into standing instructions, and wants to report an improvement figure to management. What should it do first?

- A. Move the format rules into configuration first.
- B. Consolidate the prompts first.
- C. Choose "time saved" as the figure.
- D. Run the current process once unchanged and record time, revision rounds and manual steps.

Q26. An associate must reconcile fourteen months of messy technician-shift exports against sign-in sheets. They have opened one file and plan to write rules for every problem they can imagine. What should the first prompt be?

- A. The full set of imagined rules.
- B. Produce the reconciliation now and fix issues later.
- C. Profile the files and list anything ambiguous or irregular before any rule is written.
- D. Request code in a named language first.

Q27. A proposal's control section says "All replies will be reviewed by a person before they are sent." A colleague says that is fine. What is the most appropriate improvement?

- A. Change to "light human review".
- B. Sample 10% after sending.
- C. Name the role that reviews, the specific risk it checks, and confirm it happens before sending.
- D. Remove the requirement.

Q28. A pilot cuts drafting time 40%, but reviewers still spend the same time correcting missing source details. What is the BEST next action?

- A. Declare success.
- B. Scale immediately.
- C. Measure only satisfaction.
- D. Investigate why source details are missed and refine before expanding.

Q29. An associate built a quote generator inside a Claude Project for personal use. The sales director made it a CRM step and two branches now quote customers from it. The associate still patches the prompt when a branch reports a wrong figure. Most appropriate action?

- A. Keep improving the prompt.
- B. Rebuild with the API yourself.
- C. Escalate to engineering: others depend on it, so it needs an owner, uptime and access control.
- D. Package it as a shared Skill.

Q30. An AP desk computes early-payment discounts with code; the next step is a supervisor who releases or holds each payment. A team member proposes removing the release step because "code is always right." Most appropriate response?

- A. Agree; the step is redundant.
- B. Keep it: delegating the computation is safe only because a person decides next.
- C. Replace the supervisor with a second AI pass.
- D. Move the calculation back to manual.

Q31. A customer-onboarding workflow has run well for four months. Which TWO proposed changes give Claude more authority than the risk allows? (Select TWO)

- A. Let Claude approve KYC exceptions directly because summaries have been accurate.
- B. Have the analyst sign off exceptions from the summary line without opening any case, to hit SLA.
- C. Move document extraction into a reusable Skill.
- D. Compute risk scores with code and keep the approval gate right after.
- E. Add a named reviewer to the rationale-drafting step.

Domain 5: Configuration and Knowledge Management (12%)

Instructions aur knowledge ka intezaam

Q32. A consultant configures Claude for ongoing work with one client. Which THREE placements are correct? (Select THREE)

- A. Firm-wide tone rule for every chat -> account-level instructions
- B. Client's recurring reference files -> the client Project's knowledge
- C. "Negotiation paused until month end" -> account-level instructions
- D. "Treat figures absent from the client's files as unconfirmed" -> client Project instructions
- E. One-off question about an unrelated memo -> its own new Project
- F. Preference for short replies -> memory, by mentioning it once

Q33. Pricing files of wholesale and retail units sit in one Project. In a retail quote chat Claude applied wholesale volume discounts. Most appropriate fix?

- A. Add "unit: retail" at the top of every prompt.
- B. Create one Project per unit with its own instructions and files.
- C. Use incognito for retail.
- D. Delete wholesale files after each retail chat.

Q34. In every new chat Claude addresses an associate as lead of a project that closed months ago. The associate adds a correction line to each prompt. Most appropriate fix?

- A. Keep the correction line.
- B. Use incognito forever.
- C. Switch models.
- D. Read stored memory, correct or delete the closed-project note, then start a fresh chat.

Q35. A read-only Google Drive connector summarises tenders. A new workflow needs Claude to move closed tenders into an archive folder. Which understanding is correct?

- A. Already possible; connectors inherit your account's abilities.
- B. Add a Skill that tells Claude how to archive.
- C. A second connector is required.
- D. Moving files is a write action to grant deliberately, once the connector is confirmed to offer it.

Q36. An associate reviews Claude's campaign drafts. Which TWO reactions are turned into an instruction aimed at the right layer? (Select TWO)

- A. "Writes for engineers every time" -> audience constraint in standing instructions
- B. "Cites the discontinued product line" -> replace the outdated catalogue in knowledge
- C. "Too long every time" -> switch to a more capable model
- D. "Missed this week's campaign point" -> put this week's goal in standing instructions
- E. "Buries the ask" -> tell Claude to remember via memory

Domain 6: Governance, Risk, and Responsible Use (15%)

Governance, risk aur zimmedarana istemal

Q37. An associate must apply 140 payroll corrections to named employee records; identifiers are the substance of the task. Policy classes such records as restricted. Most appropriate action?

- A. Redact identifiers and do what remains.
- B. Confirm with the compliance owner which route is approved for this data and purpose, then use it.
- C. Upload to the enterprise workspace since it is approved for internal work.
- D. Build your own connector to the HR system.

Q38. A facilities analyst coded every badge number, but the workbook's second tab maps each code back to an employee. Most appropriate action?

- A. Proceed; it is anonymised.
- B. Upload in incognito.
- C. Remove the lookup tab and unneeded fields, confirm no mapping travels, then analyse.
- D. Abandon the analysis.

Q39. Call transcripts, recorded with consent for quality assurance, sit in an approved Claude workspace. Product asks you to mine them for upsell leads. Most appropriate action?

- A. Proceed; already in an approved workspace.
- B. Strip names, then proceed.
- C. Use incognito chats.
- D. Take it to the data owner: a new use needs fresh permission.

Q40. A Claude-drafted pricing notice with a serious error went to 300 customers an hour ago. Which THREE actions belong in the first hour? (Select THREE)

- A. Stop the spread: halt further sends and scheduled runs.
- B. Write down what went out, to whom, when, and what was wrong.
- C. Delete the chat and draft.
- D. Wait until full scope is understood before telling anyone.
- E. Report promptly through the incident path, plainly, without guessing.
- F. Have Claude draft a correction and send it yourself within the hour.

Q41. A team lead cannot tell from policy whether recording customer calls for AI summaries needs consent in a newly entered jurisdiction. Which escalation gives the legal owner the most useful basis for a decision?

- A. The workflow, who it reaches, the control already added, and the one specific question policy leaves open.
- B. A ten-criterion risk matrix with a recommendation to approve.
- C. "I think this is fine - can you approve today?"
- D. A request to pause until the policy is rewritten.

Q42. An associate writes the CFO version of a workflow description: "AI clears the disputes, cutting the backlog from ten days to two." What is the problem?

- A. No published source for the figure.
- B. It drops the human checkpoint, so the CFO believes there is no gate.
- C. Too brief for the CFO.
- D. No problem.

Domain 7: Troubleshooting and Optimization (10%)

Masle pakarna aur behtar banana

Q43. Two hours into a supplier-audit chat, Claude drops the page numbers it was told to include, contradicts an earlier supplier score, and drifts into general advice. Most appropriate move?

- A. Restate the rules every turn and continue.
- B. Open a blank chat and start from zero.
- C. Ask for a summary of decisions, open points and rules; correct it; paste it into a fresh chat with the files.
- D. Move the conversation into a Project.

Q44. After four rounds, a supplier scorecard's columns and filters are right, but the weighted-score column stays subtly wrong despite ever more precise wording. Most likely root cause?

- A. Feature problem: weighted scores are written as prose and belong in code execution.
- B. Description problem: one more precise attempt.
- C. Model problem: needs Opus.
- D. Context problem: the chat is overloaded.

Q45. In a fresh chat, the first draft of a first-ever tariff notice omits the legal footer and uses the wrong currency. Which action fits the symptom?

- A. Switch to Opus.
- B. Start another fresh chat with the same request.
- C. Treat it as under-specification: re-read the request and add what it left out.
- D. Check whether a stored instruction went stale.

Q46. To "save context," an associate pastes a one-page summary of a 70-page agreement and asks for a risk analysis. Two indemnity clauses never in the summary are missed. Most appropriate change?

- A. Upload the full agreement.
- B. Paste a longer summary.
- C. Switch to Opus.
- D. Ask which sections are needed and paste only those.

Q47. Outputs got worse after many one-off exceptions were added to shared instructions. What should the team troubleshoot FIRST?

- A. Whether an even longer instruction set is needed.
- B. Whether a cheaper model follows more literally.
- C. Whether to simplify shared instructions and move temporary exceptions out of persistent configuration.
- D. Whether research mode should replace the workflow.

Case Study Sets: Long Scenarios (2-3 questions each)

Lambi kahani, ek se zyada sawal - exam jaisa

Case A - Tyre-hire depot, morning log. A desk associate at a tyre-hire company drafts overdue-hire reminders in the company's approved AI workspace, owned by the operations manager and approved for customer data. The workspace shares its maker's consumer app, which she also uses at home on her own personal account. The company's AI policy treats near misses as well as incidents as reportable. Her notes: **08:47** - opened the assistant from the browser history bar to draft seven reminders; typed the request, saw the personal account name at the top, closed the tab before pasting anything. **08:52** - noticed the personal app history showed two chats drafted at home last month; nothing of the company's in them. **08:58** - opened the workspace from the company bookmark, signed in on the company account, pasted the overdue list (seven customers' names and phone numbers), drafted the seven reminders. **09:10** - one reminder carried a return date she did not recognise; she checked it against the hire system before sending and corrected it; sent all seven. Told the yard manager about the morning. Yard manager's view: "Nothing went in, so there is nothing to report."

Q48. What should the associate do next?

- A. Nothing; the yard manager is right that no data entered the personal account.
- B. Report the 08:47 event to the operations manager through the incident path as a near miss.
- C. Delete the personal app history to remove any risk.
- D. Ask the assistant in the company workspace whether a near miss occurred.

Q49. (Case A) Which detail in the log shows a control working as intended rather than a problem?

- A. The 08:47 tab being opened from browser history.
- B. The 09:10 check of the unfamiliar return date against the hire system before sending.
- C. The yard manager being told informally.
- D. Using the same maker's app at home.

Q50. (Case A) Which TWO changes would most reduce the chance of a repeat? (Select TWO)

- A. Always open the workspace from the company bookmark, not browser history.
- B. Stop using AI for reminders altogether.
- C. Make the account identity visible and checked before pasting customer data (a pre-paste check).
- D. Let the yard manager decide what counts as reportable.
- E. Use incognito for customer lists.

Case B - Monthly supplier report, time study. A procurement team produces a monthly supplier-performance report. Current process (recorded last month): export data 25 min; clean it by hand 70 min; draft narrative 90 min; manager review 40 min with 3 revision rounds; total about 3 h 45 min. The team pilots Claude: drafting falls to 20 min, but cleaning is still manual and manager review rises to 65 min because figures in the narrative do not always match the sheet. The manager wants to roll it out to four other teams next week and report "75% time saved on drafting".

Q51. What is the most accurate description of the pilot's result so far?

- A. A 75% drafting saving; ready to scale.
- B. End-to-end time fell only modestly (about 3h45 to about 3h00) and review effort rose, so value is unproven.
- C. No saving at all.
- D. The pilot failed and should be stopped.

Q52. (Case B) Figures in the narrative do not always match the sheet. What is the best fix?

- A. Ask Claude to be more careful with numbers.
- B. Compute figures with code execution from the sheet and have the narrative reference them; check samples.
- C. Switch to Opus.
- D. Let the manager keep correcting.

Q53. (Case B) Where should the team look next for the biggest remaining saving?

- A. The 70-minute manual cleaning step, starting by profiling the export.
- B. Shortening the manager review by removing it.
- C. The 25-minute export.
- D. Adding more narrative.

Case C - Clinic appointment summaries. A private clinic chain wants Claude to summarise doctors' consultation notes into patient-friendly after-visit letters. Notes contain names, diagnoses and medication. The clinic has an enterprise Claude plan; its data-protection officer (DPO) approved it for "internal administrative documents". A team lead proposes to start Monday, have Claude draft and a receptionist send the letters, and "review a sample later if complaints come in."

Q54. Before starting, what is the most appropriate step?

- A. Start Monday; the enterprise plan is approved.
- B. Confirm with the DPO whether patient clinical data for patient letters is within the approved route and purpose.
- C. Remove patient names and start.
- D. Use incognito to avoid retention.

Q55. (Case C) If approved, which review design is appropriate?

- A. Receptionist sends; sample later.
- B. A named clinician checks each letter's medical content before it is sent.
- C. Claude rates its own accuracy.
- D. Send, and correct on complaint.

Case D - Freight quote tool. An analyst built a freight-quote helper as a Claude Project: instructions hold the rate rules; knowledge holds the 2025 and 2026 rate cards. Three salespeople now use it daily. Last week a customer was quoted a 2025 rate. The analyst's fix was adding "use latest rates" to the instructions. This week a fourth salesperson quoted a rate with a surcharge nobody recognises.

Q56. What is the most direct fix for the 2025-rate problem?

- A. Keep "use latest rates".
- B. Retire the superseded 2025 rate card from the Project knowledge.
- C. Ask salespeople to mention the year in every prompt.
- D. Switch to a more capable model.

Q57. (Case D) The unrecognised surcharge most likely means what, and what is the next step?

- A. Nothing; trust the output.
- B. A possible unsourced figure: trace it against the rate card; if absent, treat the quote as wrong and report it to the owner.
- C. A model tier issue.
- D. Delete the chat.

Q58. (Case D) Several people now depend on the tool for customer quotes. What should happen?

- A. The analyst keeps patching it alone.
- B. Give it a named owner and change control, and escalate to engineering if it becomes part of a formal process.
- C. Make every salesperson build their own.
- D. Stop using it.

Answer Key with Explanations

Q1: B - Specific spec beats length. / Lamba prompt nahi, reader-purpose-length-items-exclusions wala specific prompt.

Q2: D - Decision rights + disclosure stops silent choices during the work. / Batayein kya khud decide kar sakta hai aur choice karte waqt bataye.

Q3: B - Schema + example. / Exact field names, types aur example dein.

Q4: C - Vague rule -> explicit threshold. / "Unusual" mubham hai, exact threshold likhein.

Q5: A, C - Criteria rating and both-sides case neutralise sycophancy. / Criteria par rating aur for/against dono.

Q6: C - Unknown stays unknown; list what is needed. / Andaza nahi, "unknown" likhein.

Q7: B - Map before you make. / Pehle jo maloom hai uska map.

Q8: D - Diverge, then converge. / Pehle bohat se ideas, phir alag step mein evaluation.

Q9: C - A seeded fault proves only its class. / Ek qisam ka test sirf usi qisam ko prove karta hai.

Q10: D - Same brief = shared blind spot. / Brief ki ghalti dono mein.

Q11: D - Check assumptions against source. / Har assumption ko asal document se milaein.

Q12: C - Executed code can sum the wrong column. / Column aur rows check karein.

Q13: A - 92/200 is 46%. / Aadhe se kam, "most" ghalat.

Q14: A, C, E - Unsourced precision, unverifiable check, internal contradiction. / Bina source figure, bina saboot "check kiya", contradiction.

Q15: B - Completeness failure. / Critical finding chhootna = incomplete.

Q16: C - Check what the AI excluded. / Chhore gaye cases ka sample.

Q17: C, D - Regulatory and financially material external outputs. / Regulator aur paise wali cheezen.

Q18: A - Expert judgment covers what a novice cannot see. / Domain expert se check.

Q19: D - Match model to task. / Chhote kaam par Opus, bare kaam par Haiku - dono ghalat.

Q20: A - Long cross-referenced structure -> Opus. / Splitting cross-references tor deta hai.

Q21: A, C, E - Quick fact=search; reasoning=thinking; multi-source=research. / Internal drive ke liye connector chahiye.

Q22: B - Private modes are not audit evasion. / Vendor kuch arsa copy rakhta hai.

Q23: A, B, D - Format follows use. / Istemal ke hisaab se format.

Q24: A - No problem, no fix. / Masla nahi to chat jari rakhein.

Q25: D - Baseline first. / Pehle baseline.

Q26: C - Profile before rules. / Pehle data profile.

Q27: C - Who, what risk, when. / Kaun, kis risk ke liye, kab.

Q28: D - Value includes review effort. / Review time kam nahi hua to kamyabi adhoori.

Q29: C - Dependency -> engineering. / Doosre depend karein to engineering.

Q30: B - Human gate is what makes delegation safe. / Insaani gate hi delegation ko mehfooz banata hai.

Q31: A, B - Direct approval and rubber-stamp review both hand authority to AI. / Rubber stamp review bhi AI ko authority deta hai.

Q32: A, B, D - Scope each rule to where it applies. / Har rule apni jagah.

Q33: B - Separate knowledge. / Alag Projects.

Q34: D - Fix the source: memory. / Memory theek karein.

Q35: D - Write permission is deliberate. / Write access soch kar diya jata hai.

Q36: A, B - Recurring = instructions; stale fact = knowledge. / Is hafte ka goal prompt mein.

Q37: B - Owner confirms route for data + purpose. / Compliance owner se route.

Q38: C - Mapping = not anonymised. / Lookup tab hatayein.

Q39: D - Purpose limitation. / Naya maqsad, nayi ijazat.

Q40: A, B, E - Contain, document, report. / Roko, likho, report karo.

Q41: A - Clear, specific escalation. / Wazeh aur ek sawal.

Q42: B - Every audience must see the gate. / Human checkpoint gaib.

Q43: C - Verified summary to fresh context. / Theek ki gayi summary nayi chat mein.

Q44: A - Arithmetic -> code. / Hisaab code se.

Q45: C - First try, missing parts = spec. / Pehli dafa - brief adhoori.

Q46: A - Model sees only what is given. / Poora document.

Q47: C - Clean the configuration. / Aarzi exceptions bahar.

Q48: B - Policy counts near misses; owner is the operations manager. / Near miss bhi report, owner ko, incident path se.

Q49: B - Verify before sending = good control. / Bhejne se pehle check karna sahi control hai.

Q50: A, C - Fix the route and add a check. / Sahi rasta aur paste se pehle check.

Q51: B - Measure end-to-end incl. review. / Poora process naapein, sirf drafting nahi.

Q52: B - Numbers -> code, then verify. / Figures code se.

Q53: A - Biggest repetitive step. / Sab se bara repeat hone wala step.

Q54: B - Approval is per data AND purpose. / Route aur maqsad dono approve hon.

Q55: B - Clinical guidance needs clinical review before release. / Doctor pehle check kare.

Q56: B - Remove obsolete source. / Purana rate card hatayein.

Q57: B - Trace to source. / Source mein na ho to ghalat.

Q58: B - Dependency needs ownership. / Owner aur control.