

Exam Overview and Strategy

CCAO-F (Claude Certified Associate - Foundations): 60 items, 120 minutes, scaled cut score 720/1000.

Multiple-choice and multiple-response (each item says how many to select). No per-domain minimum; only the total counts. PCAO-F (Panaversity) follows the same blueprint and weights, then adds vendor-neutral judgment (Claude and ChatGPT side by side).

60 sawal, 120 minute, passing 720/1000. Kuch sawalon mein 2 ya 3 jawab select karne hotay hain - sawal khud batata hai kitne. Har domain ka alag pass nahi, sirf total score.

#	Domain	Weight	Main courses
1	Prompting and Task Execution	14%	AI Prompting 2026, Just Delegate It, AI Fluency
2	Output Evaluation and Validation	21%	AI Fluency, Just Delegate It, Governance
3	Product and Model Selection	12%	Claude/ChatGPT Quick Ref, What AI Actually Is
4	Workflow Integration and Solution Design	16%	Workflow Design and Diagnosis, Governance
5	Configuration and Knowledge Management	12%	What AI Actually Is, Skills and Connectors
6	Governance, Risk, and Responsible Use	15%	Governance, Just Delegate It, Code You Never Write
7	Troubleshooting and Optimization	10%	Workflow Design, AI Prompting, General Agents

52% of the exam = Output Evaluation + Workflow Design + Governance. Study AI Fluency, Workflow Design and Governance hardest.

Sab se zyada number: Evaluation (21%), Workflow (16%), Governance (15%). In teen par sab se zyada mehnat karein.

How to read long scenario questions (Lambay sawal kaise parhein)

- Read the LAST line (the actual question) first, then the story. / Pehle aakhri line parhein - asal sawal kya hai.
- Underline: who owns the data/decision, what policy says, what already happened (timeline), is it reversible, who is the audience. / Owner, policy, timeline, reversible ya nahi, audience - ye cheezen mark karein.
- Distractors usually: switch to a bigger model, ask Claude if it is sure, add "be careful", delete evidence, wait until everything is known, remove the human. / Ghalat options aksar yehi hotay hain.
- Correct answers are usually specific, proportionate, name an owner/role, and fix the ROOT cause (input, spec, feature, context, data route). / Sahi jawab specific hota hai aur root cause theek karta hai.
- Log/timeline questions (e.g. 08:47, 08:58, 09:10) test incident handling: near misses are REPORTED through the incident path to the owner, plainly, even if "nothing went in". / Timeline wale sawal: near miss bhi report hota hai.
- Time budget: 2 minutes per item. Mark hard ones and return. / Har sawal ~2 minute.

1. Just Delegate It - Giving AI a Job, and Checking the Result

Treat Claude like a capable new colleague: give a complete brief, decide what it may decide alone, and check the result in proportion to the stakes. Delegation is a skill, not a single prompt.

Claude ko ek qabil naye colleague ki tarah treat karein: poori brief dein, batayein woh khud kya decide kar sakta hai, aur result ko risk ke hisaab se check karein.

A complete brief

A good delegation states: what should exist (the deliverable), why it is needed (purpose), who reads it (audience), constraints (budget, dates, length), format, what to leave out, and how you will check it.

Roman Urdu: Achhi brief mein: kya banana hai, kyun, kis ke liye, limits (budget/dates), format, kya شامل nahi karna, aur aap kaise check karenge.

Decision rights - "say when you chose"

The brief must say what Claude may decide alone and that it must SAY when it made a choice (e.g. dropped a tied option, widened a date range). This prevents silent choices. Listing assumptions afterwards or approving a plan beforehand does not catch silent choices during the work.

Roman Urdu: Brief mein likhein Claude khud kya decide kar sakta hai aur jab bhi koi choice kare to bataye. Is se "silent choices" ruk jati hain.

Missing inputs are filled confidently

If a step depends on information never put in front of the model (a meeting decision, a verbal budget cut), it will fill the gap at full confidence. This is not a prompt-wording or model-tier problem; it is a missing-input problem.

Roman Urdu: Agar koi zaroori maloomat model ke saamne thi hi nahi (jaise meeting mein budget half hua), to model poore confidence se gap bhar dega. Hal: woh input dena.

Check in proportion to stakes

Low-stakes, reversible, internal output: your own read-through is enough. High-stakes, external, financial, regulatory or irreversible actions: a second reviewer or a hold-and-verify step before the action.

Roman Urdu: Kam risk, internal kaam: khud parhna kaafi. Bahar jaane wala, paisa, regulator ya na-wapas hone wala kaam: doosra reviewer ya pehle rok kar verify.

Accountability never moves to the AI

A track record earns the assistant similar jobs, but every output still carries a named human owner who answers for it. "AI-generated" labels do not replace an owner; the vendor does not become accountable.

Roman Urdu: Kitne bhi clean runs hon, har output ka ek naam wala insaan zimmedar hota hai. "AI-generated" label owner ki jagah nahi leta.

Exam traps / Yaad rakhein

- Silent choices -> fix with "what you may decide alone + say when you chose".
- Confident wrong budget -> the input was never given (not model tier).
- Clean track record does not remove the named owner.

2. What AI Actually Is - What Claude Is and Is Not

Claude is a language model: it predicts plausible text from what is in its context window. It has no live knowledge unless given tools (search, connectors), it can be fluent and wrong, and its self-reported confidence is not evidence.

Claude ek language model hai jo context window mein maujood cheezon se text banata hai. Tools ke baghair taaza maloomat nahi, fluent hote hue bhi ghalat ho sakta hai, aur uska "confidence" saboot nahi.

Context window = what the model sees

Quality depends on what is in front of the model. Pasting your own 1-page summary of a 70-page contract means clauses you left out cannot be analysed. Upload the full source when analysis depends on it.

Roman Urdu: Jo model ke saamne hai wahi analyse hoga. 70 page contract ki 1 page summary di to missing clauses kabhi nazar nahi ayengi. Poora document dein.

Hallucination signs

Precise statistics with no source; claims like "I checked the website" with no page or quote; internal contradictions (a figure on page 1 differs from the end); fabricated citations. Hedging and length are NOT signs of fabrication.

Roman Urdu: Fabrication ki nishaniyan: bina source ke exact figure, "maine website check ki" bina quote ke, andar se contradictions. Hedging ya lambai nishani nahi.

Asking "are you sure?" is not verification

Self-rated confidence, re-asking the same model, or asking it to review itself does not verify facts. Verify against the source document, a primary source, or independent computation.

Roman Urdu: "Kya aap sure hain?" poochna verification nahi. Asal source ya primary document se check karein.

Shown reasoning is a justification, not a window

Listed assumptions or visible "thinking" are a justification offered for review. Check each assumption against the documents; an unsourced assumption means the conclusion cannot stand.

Roman Urdu: Model jo reasoning dikhati hai woh review ke liye justification hai, asli andar ka process nahi. Har assumption ko documents se milaein.

Memory, context rot, and fresh chats

Long chats drift: rules given early are forgotten and answers contradict earlier ones. Fix: ask for a summary of decisions, open points and rules, correct it, paste into a fresh chat with the files. Stale stored memory (an old project) is fixed by reading and correcting/deleting that memory, then starting fresh.

Roman Urdu: Lambi chat mein model purane rules bhool jata hai. Hal: decisions aur rules ki summary banwa kar, theek karke, nayi chat mein files ke saath paste karein. Purani memory ho to memory edit/delete karein.

Exam traps / Yaad rakhein

- Summary instead of source -> upload the full document.
- Drift in long chat -> verified summary into a fresh chat (not "restart from zero", not "keep repeating").
- Old wrong assumption every chat -> fix stored memory, not a correction line each time.

3. AI Fluency - The 4Ds (Delegation, Description, Discernment, Diligence)

AI Fluency is the habit set for working with AI effectively, efficiently, ethically and safely. The 4Ds: Delegation (what to hand over), Description (how to ask), Discernment (judging output), Diligence (responsibility for what you use and share).

AI Fluency ke 4D: Delegation (kya kaam dena hai), Description (kaise maangna hai), Discernment (output ko parakhna), Diligence (zimmedari aur sachai).

Delegation

Decide which tasks or steps suit AI: repetitive synthesis, drafting, first-pass classification. Keep judgment, approval and final decisions with people. Automation (AI does it), augmentation (AI helps a human), agency (AI acts within bounds).

Roman Urdu: Kaunsa kaam AI ko dena hai: repeat hone wala, draft, pehla classification. Faisle aur approval insaan ke paas.

Description

Describe the product (what), the process (how to approach), and the performance (how Claude should behave, e.g. challenge me). Clear audience, purpose, length, items to cover and exclude.

Roman Urdu: Product (kya chahiye), Process (kaise karna), Performance (kaisa behave kare) - teeno wazeh karein.

Discernment - judging the output

Check accuracy against sources, completeness (was anything material left out, e.g. a critical finding or a new regulation), overgeneralisation ("most" when 92 of 200 support), framing bias, internal consistency, and sequence where order matters.

Roman Urdu: Output ko parakhna: sahi hai? kuch ahem chhoota? "most" jaisa overgeneralization? bias? andar se consistent? order sahi?

Sycophancy and leading prompts

If your prompt hints at the answer you want ("tell me this is defensible", "isn't this good enough?"), the model tends to agree. Neutralise it: ask for the strongest case for AND against, or rate on named criteria with the one change that most improves it.

Roman Urdu: Agar prompt mein apni pasand jhalakti hai to model haan mein haan milata hai. Hal: for aur against dono maangein, ya criteria par rating.

Diligence

Be transparent about AI use where it matters, own the output you send, and never present generated content as real evidence (e.g. fake testimonials, invented customer research).

Roman Urdu: AI ka istemal jahan zaroori ho wahan batayein, apne bheje hue kaam ki zimmedari lein, fake testimonials kabhi nahi.

Exam traps / Yaad rakhein

- "Tell me this is good" = sycophancy trap; correct rewrites ask for both sides or criteria-based rating.
- 92 support / 78 oppose / 30 neutral of 200 -> "most support" is inaccurate.
- Missing critical finding or new regulation = completeness failure, even if every stated fact is correct.

4. AI Prompting in 2026 - Getting Useful Output

Modern prompting is less about magic words and more about a complete specification: goal, context, audience, constraints, format, examples, and how uncertainty must be handled. Longer is not better; specific is better.

Aaj prompting ka matlab jaadui alfaaz nahi balke mukammal specification hai: maqsad, context, audience, limits, format, misaalen, aur uncertainty kaise handle ho. Lamba nahi, specific.

The five-part fix for generic output

State the reader, the purpose, the length, the items to cover, and what to leave out. A write-up for an insurer's loss adjuster needs 3 open defects with dates and photo refs, not a tidy narrative.

Roman Urdu: Generic output ka hal: reader, purpose, length, kya cover karna, kya chhorna.

Structured output

For JSON or tables feeding another system: give exact field names, types and a short example. Allowed values, required fields and row mapping must be validated before import.

Roman Urdu: JSON/table ke liye exact field names, types aur ek example dein. Import se pehle required fields check karein.

Classification consistency

Give definitions plus representative examples, including borderline cases. For "flag unusual expenses", write a named threshold and category so two people get the same result.

Roman Urdu: Consistency ke liye definitions + borderline examples. "Unusual" ki jagah exact threshold likhein.

Map before you make; diverge then converge

For big deliverables, first ask Claude to map what is known (existing approaches, prior attempts). For ideas, generate a broad unranked set first, then evaluate against fixed criteria in a separate step.

Roman Urdu: Bara kaam: pehle jo maloom hai uska map. Ideas: pehle bohat se unranked ideas, phir alag step mein criteria par evaluate.

Model and thinking choice

Haiku: fast, cheap, routine high-volume (tagging, simple summaries). Sonnet: balanced default. Opus: hardest, nuanced, long complex structures (180-page cross-referenced contract). Extended thinking: multi-constraint decisions with real consequences, not formatting or lookups. Using the flagship on trivial tasks and the cheapest rung on a critical rebuild are both wrong.

Roman Urdu: Haiku: tez aur sasta routine kaam. Sonnet: balanced. Opus: sab se mushkil aur nuanced kaam. Extended thinking: sirf complex faislon ke liye.

Missing data

Tell Claude to mark missing fields as unknown and list what is needed, never to estimate them silently. Blanks are unknown, not zero.

Roman Urdu: Missing data ko "unknown" likhwayein, andaze se na bharwayein. Blank ka matlab zero nahi.

Exam traps / Yaad rakhein

- First fix is the prompt/spec, not "switch to Opus".
- Longer prompt with scenery is not the fix; the five-part spec is.
- Change model when the task truly needs capability (very long, cross-referencing) after the prompt is right.

5. Claude 101 - ChatGPT and Claude Quick Reference

Know which product feature fits which job: chat, Projects, artifacts, memory, incognito/temporary chat, web search, research mode, extended thinking, code execution, file creation, connectors and skills.

Kaunsa feature kis kaam ke liye: chat, Projects, artifacts, memory, incognito, web search, research mode, extended thinking, code execution, connectors, skills.

Chat vs Project

One-off question in a clean focused chat: stay in chat. Recurring work with the same files and stable instructions: a Project (instructions + knowledge). One Project per client/business unit; mixing wholesale and retail files causes bleed.

Roman Urdu: Ek dafa ka kaam: chat. Baar baar same files aur instructions: Project. Har client/unit ka alag Project.

Where each instruction lives

Firm-wide rule for every chat -> account-level instructions. Client reference files -> that Project's knowledge. Client-specific rules ("treat figures absent from files as unconfirmed") -> Project instructions. Temporary facts (negotiation paused this month) -> per task, not permanent settings.

Roman Urdu: Har chat ka rule: account-level. Client files: Project knowledge. Client rules: Project instructions. Aarzi baat: sirf us task mein.

Web search vs research mode vs extended thinking

Web search: a quick current fact (this morning's announcement). Research mode: a multi-source cited investigation (compare providers with funding and customers). Extended thinking: reasoning over material you already have (which lease is riskier to exit). Friendly rewrite of 3 lines needs none. Internal shared drive files need a connector, not web search.

Roman Urdu: Web search: taaza ek fact. Research mode: bohat se sources wali cited report. Extended thinking: mushkil reasoning. Internal drive: connector.

Artifacts and formats

Quick explanation -> inline reply. Data that finance loads into a sheet -> table or spreadsheet file. A checklist to circulate -> document artifact. Iterated draft -> artifact beside the conversation.

Roman Urdu: Choti wazahat: chat mein. Sheet mein jaane wala data: table/spreadsheet. Share hone wala document: artifact.

Incognito / temporary chat

Both keep chats out of history and memory, but the vendor retains a copy for a period, and company records may hold it. They are not a way to avoid audit or to make restricted data safe.

Roman Urdu: Incognito chat history/memory mein nahi jati, lekin vendor kuch arsa copy rakhta hai. Audit se bachne ya restricted data ke liye nahi.

Code execution for numbers

Arithmetic, sums and weighted scores belong in code execution, not prose. But executed code can still sum the wrong column; check what was computed (column, rows, counts, samples).

Roman Urdu: Hisaab ke liye code execution. Lekin code bhi ghalat column jor sakta hai - check karein kya compute hua.

Exam traps / Yaad rakhein

- Weighted score wrong after many rewordings -> feature problem: use code execution.
- Incognito is not "nothing on record".
- "Clean short chat, nothing outdated" -> just continue in the same chat.

6. Skills and Connectors - Extending Claude

Connectors (MCP) give Claude access to your tools and data (Drive, Gmail, Calendar, CRM). Skills are reusable packaged instructions/know-how for a kind of task. A connector is capability and access; a skill is method.

*Connector Claude ko aapke tools aur data tak rasai deta hai. Skill ek dobara use hone wala tareeqa/instructions hai.
Connector = access, Skill = method.*

Permissions are deliberate

Read-only access does not allow moving or writing files. A write action (e.g. archive closed tenders) must be granted deliberately, after confirming the connector offers it at all. A Skill cannot add a permission the connector lacks.

Roman Urdu: Read-only connector se file move nahi hoti. Write access soch samajh kar diya jata hai, aur Skill permission nahi barha sakti.

Scope sources tightly

Point a connector at the relevant sources (one support program's mail, one approved folder). Do not let it search everything and infer.

Roman Urdu: Connector ko sirf zaroori source tak mehdood karein.

Maintain knowledge

Verify connected sources stay current, relevant and authoritative. Retire superseded editions; update changing sources (monthly thresholds) on schedule; keep stable ones (glossary). After a rename, confirm the Project still points to the authoritative file.

Roman Urdu: Sources ko updated rakhein, purane editions hatayein, rename ke baad link check karein.

When to make a Skill

When the same method must run the same way for everyone (document extraction on an onboarding desk), package it as a Skill. A Skill does not replace an owner, uptime or access control for tools others depend on.

Roman Urdu: Jab ek hi tareeqa sab ke liye same chalna ho to Skill banayein.

Reaction -> right layer

"Always writes for engineers" -> audience constraint in standing instructions. "Cites discontinued product" -> replace the outdated knowledge file. "Missed this week's goal" -> put it in this week's prompt, not standing instructions.

Roman Urdu: Har shikayat ko sahi jagah theek karein: audience = instructions, purani info = knowledge file, is hafte ka goal = prompt.

Exam traps / Yaad rakhein

- Connector does not automatically inherit everything your account can do.
- Old and new manual both in Project -> retire the old one.
- Too many one-off exceptions in shared instructions -> simplify and move temporary ones out.

7. General Agents on the Web - Cowork in a Browser

General agents (Claude in Chrome, Cowork, ChatGPT agent) take multi-step actions: browse, fill forms, move files, send messages. More autonomy means more need for clear bounds, checkpoints and approval before irreversible steps.

General agents khud multi-step kaam karte hain: browse, form bharna, files move karna, message bhejna. Jitni zyada azaadi, utni zaroorat wazeh hudood aur approval ki.

Bound the task

Tell the agent the goal, allowed sites/tools, what it must not do (no purchases, no sends), and where it must stop and ask. Approve irreversible actions yourself.

Roman Urdu: Agent ko batayein kya karna, kaun se tools, kya bilkul nahi karna, aur kahan ruk kar poochna hai.

Watch for untrusted content

Web pages and documents can contain instructions (prompt injection). An instruction inside a page is data, not your command; the agent should not follow it without you.

Roman Urdu: Web page ke andar likhi hidayat aap ka hukm nahi. Agent ko us par amal nahi karna chahiye.

Session hygiene

For long agent runs, give a verified brief, check intermediate results, and restart with a clean summary if the run drifts.

Roman Urdu: Lambhe runs mein beech beech mein check karein, drift ho to clean summary se dobara shuru.

Account and workspace separation

Use the company-approved workspace on the company account for company data. Personal consumer accounts or history from home must not mix with customer data; if it nearly happened, report it as a near miss.

Roman Urdu: Company data sirf company ke approved workspace mein. Ghar ka personal account mix ho jaye ya hone lage to near miss report karein.

Exam traps / Yaad rakhein

- Agent autonomy does not remove approval before irreversible actions.
- Prompt injection: follow the user, not the page.

8. Workflow Design and Diagnosis - Fitting AI In, and Fixing It

Map the current process first, measure a baseline, place AI where it removes repetitive work, keep human gates at decisions, and diagnose failures by root cause: input, description, feature, model, or context.

Pehle maujooda process ka map aur baseline naap, phir AI ko repeat wale kaam par lagayein, faislon par insaani gate rakhein, aur masla root cause se pakrein.

Baseline before change

Run the current process once unchanged and record time, revision rounds and manual steps. Without a baseline you cannot report an improvement. Success = quality + reviewer effort + time, not time alone.

Roman Urdu: Tabdeeli se pehle ek dafa purana process chala kar time, revisions aur manual steps note karein. Kamyabi = quality + review effort + time.

Profile before rules

With unknown, "messy" data, first ask Claude to profile the files and list anything ambiguous before writing rules. Do not invent rules for imagined problems.

Roman Urdu: Ganda data ho to pehle profile karwayein, phir rules likhein.

Human-in-the-loop that actually works

Name the role that reviews, the specific risk the review checks, and when it happens (before sending). "Reviewed by a person" is too vague. Match review intensity to risk; sample routine outputs, fully check consequential ones.

Roman Urdu: Review ka role, kis risk ko check karna hai, aur kab - teeno likhein. Review risk ke mutabiq.

Check what was routed away

If AI routes 120 of 900 claims to fast-track, reviewing the 120 is not enough; sample the 780 not selected to check whether a group was systematically excluded.

Roman Urdu: Jo cases AI ne chhor diye unka sample bhi check karein - kahin koi group systematically bahar to nahi hua.

Diagnose by root cause

Omitted footer / wrong currency on first try -> under-specification. Numbers subtly wrong after many rewordings -> feature (use code). Long doc shallow -> model/capability. Drift after hours -> context. Different results each run -> vague rule. Output depended on missing info -> input.

Roman Urdu: Pehli dafa footer/currency ghalat: brief adhoori. Hisaab baar baar ghalat: code use karein. Lamba document: model. Ghanton baad drift: context. Har run alag: rule mubham.

When a prompt tool becomes a system

When others depend on it (a CRM step, branches quote from it), escalate to engineering: it needs an owner, uptime and access control.

Roman Urdu: Jab doosre log us par depend karne lagein to engineering ke paas le jayein.

Exam traps / Yaad rakhein

- "First" questions: baseline, profile, map - before changing anything.
- Fast drafts but same correction time -> investigate missed source details before scaling.
- Human gate stays after code computes; delegation is safe because a person decides next.

9. Governance, Risk and Responsible Use

Know what is allowed, where data may go, who owns decisions, and what to do when something goes wrong. Policy and data owners decide routes; you escalate clearly and report incidents and near misses promptly.

Kya allowed hai, data kahan ja sakta hai, faisla kis ka hai, aur ghalti par kya karna hai. Route policy/data owner decide karta hai; aap wazeh escalate aur fauran report karti hain.

Approved route for data

Restricted data (payroll, PII, confidential client material) goes only through the route the compliance/data owner approves for that data AND purpose. "Enterprise plan" alone is not approval. No approval -> pause. Credentials never go into prompts.

Roman Urdu: Restricted data sirf approved route se, aur us maqsad ke liye. Approval nahi to ruk jayein. Password/API key kabhi prompt mein nahi.

Anonymisation must be real

Coded IDs with a lookup tab in the same file is not anonymised. Remove the mapping and unneeded fields. Minimise data: drop names the task does not need.

Roman Urdu: Code wale IDs ke saath lookup tab ho to anonymized nahi. Mapping hatayein, ghair zaroori fields hatayein.

Purpose limitation

Data collected for one purpose (call recordings for QA) needs fresh permission from the data owner for a new use (upsell mining), even if it already sits in an approved workspace.

Roman Urdu: Ek maqsad ke liye liya gaya data naye maqsad ke liye nayi ijazat mangta hai.

Protected attributes and prohibited uses

Never use medical info, age, family status or other protected attributes where policy prohibits, even with human review. No emotional-stability assessments without an approved process. No fake testimonials.

Roman Urdu: Medical, age, family status jaisi cheezen jahan mana hain wahan bilkul nahi - human review ke saath bhi nahi.

Incidents and near misses

First hour after an error: stop the spread (halt sends/scheduled runs), write down what went out, to whom, when, what was wrong, and report promptly through the incident path plainly without guessing. Never delete evidence or wait for full scope. A near miss (e.g. customer data almost entered a personal consumer account, or a wrong account was open) is reported to the owner through the same path, even if "nothing went in".

Roman Urdu: Ghalti ke baad: phelao rokein, likhein kya gaya kis ko kab, aur fauran incident path se report karein. Saboot delete na karein. Near miss bhi report hota hai chahe kuch na gaya ho.

Escalating a policy gap

Give the owner: the workflow, who it reaches, the control already added, and the one specific question policy leaves open. Not a vague "approve today" note.

Roman Urdu: Owner ko: workflow, kis tak pohanchta hai, kya control laga hai, aur ek wazeh sawal.

Honest descriptions to each audience

Every version (ops, CFO, auditor) must keep the human checkpoint. "AI clears the disputes" misleads the CFO into thinking there is no gate.

Roman Urdu: Har audience ke version mein human checkpoint zikar hona chahiye.

Exam traps / Yaad rakhein

- Enterprise workspace does not equal approval for every data and purpose.
- Near miss -> report through incident path to the owner.
- Accuracy of a novice-written expert briefing -> ask a domain expert, not a proofreader or Claude's self-rating.

10. Code You Never Write - What Is Buildable Without an Engineer

Claude can write and run code for you (analysis, reconciliations, scripts). You still own verification: test with known answers and planted faults, check what was computed, and keep humans on irreversible releases. Hand off to engineering when others depend on it.

Claude aap ke liye code likh aur chala sakta hai. Verification phir bhi aap ki zimmedari: known answer se test, jaan boojh kar ghalti daal kar test, aur na-wapas hone wale kaam par insaan.

Executed is not the same as correct

A number from executed code can still be wrong: wrong column (gross instead of outstanding), wrong rows, duplicates. Check column, row counts, filters and samples.

Roman Urdu: Code chala hai iska matlab sahi nahi. Kaun sa column aur rows use hue, check karein.

Planted faults prove only their class

Seeding a missing-units gap proves the script catches that class. Other classes (double-counted transfers) need their own planted fault. Extra flags are leads to verify, not proven findings.

Roman Urdu: Ek qisam ki nakli ghalti pakri to sirf usi qisam par bharosa. Har qisam ke liye alag test.

Agreement is not independence

Two assistants given the identical brief can share the same omission (e.g. forgetting to net credit notes). Agreement is useful but bounded; independence needs a different brief or method.

Roman Urdu: Do AI ek hi brief par same jawab dein to bhi brief ki ghalti dono mein ho sakti hai.

Profile, test, then trust

Profile the data, test on a past period with a known correct answer, then use. Keep a named person to release payments or other irreversible actions.

Roman Urdu: Pehle profile, phir purane known month par test, phir istemal. Payment release insaan kare.

Hand-off point

Profile-and-prompt tools are fine for your own use. Once other teams depend on it, it needs engineering ownership, uptime, authentication and access control.

Roman Urdu: Apne liye theek; jab doosre depend karein to engineering.

Exam traps / Yaad rakhein

- Tested script -> still a named human releases payments (cannot recall).
- Two vendors agree -> "useful but bounded".
- Big gap vs manual figure -> check which column/rows were summed.

Last-Day Revision: Decision Rules

Situation	Correct move	Roman Urdu
Output looks generic	Add reader, purpose, length, items, exclusions	<i>Spec theek karein</i>
Output uses info you never gave	Supply the missing input	<i>Missing input dein</i>
Silent choices	State decision rights + "say when you chose"	<i>Decision rights likhein</i>
Numbers subtly wrong repeatedly	Use code execution	<i>Code execution</i>
Code result differs from expected	Check column, rows, counts, samples	<i>Kya compute hua check</i>
Long doc shallow	Upload full doc; Opus for long complex structure	<i>Poora doc / Opus</i>
Long chat drifts	Summary -> correct -> fresh chat with files	<i>Summary se nayi chat</i>
Wrong assumption every chat	Fix stored memory	<i>Memory theek karein</i>
Recurring work	Project with instructions + knowledge	<i>Project</i>
Files of two units bleed	One Project per unit	<i>Alag Projects</i>
Restricted data	Owner-approved route for data and purpose	<i>Approved route</i>
New use of existing data	Data owner permission	<i>Nayi ijazat</i>
Error already sent	Stop spread, document, report	<i>Roko, likho, report</i>
Near miss	Report via incident path to owner	<i>Near miss report</i>
Irreversible action	Named human approves/releases	<i>Insaan approve kare</i>
Claude recommends with unsourced claim	Check assumptions against sources	<i>Source se milaein</i>
Leading prompt	Ask for case for and against	<i>Dono taraf maangein</i>
AI screened many, you reviewed selected	Sample the rejected set	<i>Rejected ka sample</i>