# AI Prompting in 2026 Practice MCQs

## 10 situation-based questions with answer key

**Made by Taha Shabbir**

Based on: AI Prompting: Describing, Evaluating and Validating Outputs - The AI Agent Factory by Panaversity

Q1 · Practice

A student asks an AI for the rules and customs of a folk festival from her grandmother's village. Three confident paragraphs come back. Her aunt, who grew up there, says most of it is wrong. What explains this, and what is the right habit?

**A** The AI lied on purpose, so she should switch to a different tool and trust whatever it says next.

**B** The topic is sparse in training data, so the AI blended similar traditions; confidence is not correctness, so check with someone who knows or a primary source.

**C** The question was too short; adding more adjectives to the prompt would have produced the true rules.

**D** She should turn on a reasoning mode, since thinking longer guarantees facts about niche topics.

A runner in Lahore asks an AI "where should I run this week?" It recommends a park that has been closed for years and cites an old web page. What is the best fix for the next prompt?

**A** Switch to a different model family, because this one is clearly weak at local knowledge.

**B** Use deep research for every question from now on, so the answer is always checked.

**C** Turn web search off so the AI answers only from what it already learned in training.

**D** Name the sources you trust, ask for inline citations with dates, and ask it to flag anything it could not verify.

Q3 · Practice

Every week an analyst pastes the same three-page brand guide and the same audience description into a brand-new chat before asking for copy. What is the best way to stop repeating this?

**A** Create a project that holds the brand guide and audience, so every chat inside it inherits that context.

**B** Save the guide and the client's name into memory so the tool always knows them.

**C** Keep one endless chat open forever and never start a new one.

**D** Cut the guide down to five keywords so pasting it is quicker.

Q4 · Practice

One long chat has covered a home workout, a budget spreadsheet and a thank-you note. Now the AI is vague and suggests squats, even though the first message said "bad knee, no squats". What is the best move?

**A** Write one more very detailed message in the same chat that re-explains everything from the start.

**B** Wait for a model with a bigger context window and try again then.

**C** Start a new chat for this topic and paste in only the two facts that matter.

**D** Decide the tool is broken and switch to another product.

A founder types: "Help me defend my decision to hire X." The AI replies with warm, one-sided support. Which rewrite removes the bait?

**A** "Give me the strongest reasons I am right, as quickly as you can."

**B** "Here is my decision and the context. What is the strongest counter-argument?"

**C** "Confirm that hiring X was wise, and tell me how confident you are."

**D** "Explain again, in more detail, why hiring X was the right call."

Q6 · Practice

A writer pastes an essay and says "critique it". The AI answers with warm praise and no real criticism. Which follow-up works best?

**A** Tell the AI to "be harsh" and run it again.

**B** Ask it to rewrite the whole essay so it is better.

**C** Ask a yes/no question: "Is this ready to publish?"

**D** Give a rubric: score 1 to 10 on named criteria with one sentence each, then name the change that would raise each score most.

Q7 · Practice

A manager has four tasks and wants to use a "read everything, think hard" reasoning mode on only one. Which task is the best fit?

**A** Weighing two supplier contracts, cost data and delivery history to recommend one, with the conditions under which the recommendation flips.

**B** Looking up the capital city of a country for a slide.

**C** Summarizing a two-paragraph email from a colleague.

**D** Brainstorming a few casual names for a team lunch.

Q8 · Practice

An analyst asks: "Analyze our churn data and confirm that the price rise caused the drop." The analysis comes back agreeing with the price theory. What is the right way to set up an analysis ask?

**A** Loosen the data and definitions, and tighten the conclusion you expect.

**B** Tighten everything, including the conclusion, so the AI cannot wander.

**C** Tighten the data, the definitions and the exact question; loosen the conclusion and never specify it.

**D** Loosen the exact question so the AI can pick a more interesting one.

Q9 · Practice

A user pastes 18 numbers and asks for the median, the average and the outliers. A confident paragraph of statistics comes back with no code block. What is the best response?

**A** Trust it, since 18 numbers is small enough for the AI to get right at a glance.

**B** Ask it to write and run code, show the code it ran, and give a verifiable specific such as the row count first.

**C** Tell it to "think step by step" and accept the new answer.

**D** Switch to a model with a larger context window.

Q10 · Practice

An assistant drafts a message to a regulator that cites "a 63 percent adoption rate, up from 41" with no source. The deadline is tomorrow. What should happen before it is sent?

**A** Require a citation you can open for each figure, verify it against a primary source, and have a human reviewer sign off before it ships.

**B** Send it, because the draft is fluent and professional-sounding.

**C** Keep it, because precise numbers are a sign the AI knows the topic.

**D** Ask the AI whether the number is right and accept its confident yes.

## Q1 · ANSWER · DOMAIN 1 — HOW AI KNOWS THINGS

**Correct: B**

**A** The AI lied on purpose, so she should switch to a different tool and trust whatever it says next.

**Wrong —** it did not lie. It generalized from sparse data.
Concept: frequency in training data is roughly reliability

**B** The topic is sparse in training data, so the AI blended similar traditions; confidence is not correctness, so check with someone who knows or a primary source.

**Correct.** Sparse data means confident blends, not facts.
Concept: Frequency in training data is roughly equal to reliability

**C** The question was too short; adding more adjectives to the prompt would have produced the true rules.

**Wrong —** length cannot create knowledge that was never in the data.
Concept: strong / sparse / absent

**D** She should turn on a reasoning mode, since thinking longer guarantees facts about niche topics.

**Wrong —** reasoning does not add facts the model never saw.
Concept: think hard is for multi-trade-off questions

**CONCEPT TESTED**

Frequency in training data is roughly equal to reliability. Cooking is strong, regional folk traditions are sparse. Ask "how would it know that?" and check anything important against a primary source.

## Q2 · ANSWER · DOMAIN 1 — RETRIEVAL MODES

**Correct: D**

| | | |
|---|---|---|
| **A** | Switch to a different model family, because this one is clearly weak at local knowledge. | **Wrong —** not a model problem. Web search cites popular sources first and does not check they are current.<br>Concept: diagnose before you escalate |
| **B** | Use deep research for every question from now on, so the answer is always checked. | **Wrong —** overkill for simple questions: slower, and it still needs trusted sources named.<br>Concept: match the mode to the ask |
| **C** | Turn web search off so the AI answers only from what it already learned in training. | **Wrong —** pretrained mode is blind to anything after the cutoff, so it is worse for "this week".<br>Concept: three retrieval modes |
| **D** | Name the sources you trust, ask for inline citations with dates, and ask it to flag anything it could not verify. | **Correct.** Name trusted sources, demand dated citations.<br>Concept: Your phrasing picks the retrieval mode |

**CONCEPT TESTED**

Your phrasing picks the retrieval mode. Web search does not check that sources are current, so name authoritative sources, ask for citations after each claim, and have the AI mark what it could not verify.

## Q3 · ANSWER · DOMAIN 2 — CONTEXT

**Correct: A**

| | | |
|---|---|---|
| **A** | Create a project that holds the brand guide and audience, so every chat inside it inherits that context. | **Correct.** Repeated context belongs in a project.<br>Concept: Pasted the same files or audience into two chats? That context belongs in a project, inherited by every chat inside it |
| **B** | Save the guide and the client's name into memory so the tool always knows them. | **Wrong —** memory is a short note about you; keep client and student details out of it.<br>Concept: memory vs project |
| **C** | Keep one endless chat open forever and never start a new one. | **Wrong —** a chat window is working memory, and long chats rot.<br>Concept: context rot |
| **D** | Cut the guide down to five keywords so pasting it is quicker. | **Wrong —** keywords remove exactly what a smart new colleague would need.<br>Concept: brief it like a new colleague |

**CONCEPT TESTED**

Pasted the same files or audience into two chats? That context belongs in a project, inherited by every chat inside it. You reset the chat, not the context.

## Q4 · ANSWER · DOMAIN 2 — CONTEXT

**Correct: C**

**A** Write one more very detailed message in the same chat that re-explains everything from the start.

**Wrong —** it piles more text onto a window that is already degrading.
Concept: the reset beats the rescue

**B** Wait for a model with a bigger context window and try again then.

**Wrong —** windows are large, not infinite, and recall degrades inside them.
Concept: context rot

**C** Start a new chat for this topic and paste in only the two facts that matter.

**Correct.** Topic changed? New chat, paste two facts.
Concept: Context rot: answers get vaguer and contradict earlier constraints in one long chat

**D** Decide the tool is broken and switch to another product.

**Wrong —** this is not a tool defect; the same thing happens in any long chat.
Concept: reset the chat, not the tool

**CONCEPT TESTED**

Context rot: answers get vaguer and contradict earlier constraints in one long chat. Start a new chat when the topic changes and paste in only what matters. The reset is almost always faster than the rescue.

## Q5 · ANSWER · DOMAIN 1 — SYCOPHANCY

**Correct: B**

**A** "Give me the strongest reasons I am right, as quickly as you can."

**Wrong —** it still hands the AI the conclusion first.
Concept: find, defend, confirm, prove are bait

**B** "Here is my decision and the context. What is the strongest counter-argument?"

**Correct.** Ask for the counter-argument, not the defense.
Concept: Models are trained on human thumbs-up, and agreeing earns more of them than disagreeing

**C** "Confirm that hiring X was wise, and tell me how confident you are."

**Wrong —** "confirm" is bait; a confidence score does not make it neutral.
Concept: replace with evaluate, compare, critique

**D** "Explain again, in more detail, why hiring X was the right call."

**Wrong —** asking for more detail on one side just repeats the flattery.
Concept: models are trained on thumbs-up

**CONCEPT TESTED**

Models are trained on human thumbs-up, and agreeing earns more of them than disagreeing. Find, defend, confirm and prove hand the AI a conclusion before the question. Replace them with evaluate, compare, critique.

## Q6 · ANSWER · DOMAIN 1 — RUBRICS

**Correct: D**

**A** Tell the AI to "be harsh" and run it again.

**Wrong —** harshness is a tone, not criteria; the AI still has nothing specific to look at.
Concept: without a rubric criteria collapse into great work

**B** Ask it to rewrite the whole essay so it is better.

**Wrong —** a rewrite replaces your work instead of evaluating it.
Concept: evaluate first, then fix

**C** Ask a yes/no question: "Is this ready to publish?"

**Wrong —** a yes/no lets the AI declare you finished; you decide when to stop.
Concept: force a number

**D** Give a rubric: score 1 to 10 on named criteria with one sentence each, then name the change that would raise each score most.

**Correct.** Rubric + number beats "be harsh".
Concept: Without a rubric, criteria collapse into "great work." With one, the AI has to actually look

**CONCEPT TESTED**

Without a rubric, criteria collapse into "great work." With one, the AI has to actually look. Choosing between a 6 and a 7 forces it to commit, and a 4 next to a 7 tells you what to fix first.

## Q7 · ANSWER · DOMAIN 1 — REASONING

**Correct: A**

| | | |
|---|---|---|
| **A** | Weighing two supplier contracts, cost data and delivery history to recommend one, with the conditions under which the recommendation flips. | **Correct.** Think hard for multi-trade-off decisions only.<br>Concept: Use "think hard" for the multi-input, multi-trade-off question you would hand to a colleague |
| **B** | Looking up the capital city of a country for a slide. | **Wrong —** a quick lookup; reasoning is slower and costs more of your budget.<br>Concept: skip it for quick lookups |
| **C** | Summarizing a two-paragraph email from a colleague. | **Wrong —** short summaries do not need extended reasoning.<br>Concept: skip it for short summaries |
| **D** | Brainstorming a few casual names for a team lunch. | **Wrong —** casual brainstorming is better served by the loop than by thinking longer.<br>Concept: skip it for casual brainstorming |

**CONCEPT TESTED**

Use "think hard" for the multi-input, multi-trade-off question you would hand to a colleague. Skip it for quick lookups, short summaries and casual brainstorming. "Think step by step" is obsolete.

## Q8 · ANSWER · DOMAIN 1 — MATCH THE GRIP TO THE ASK

**Correct: C**

| | | |
|---|---|---|
| **A** | Loosen the data and definitions, and tighten the conclusion you expect. | **Wrong —** backwards: a pinned conclusion gives you an analysis that flatters the ask.<br>Concept: if you get it backwards |
| **B** | Tighten everything, including the conclusion, so the AI cannot wander. | **Wrong —** specifying the conclusion is exactly what flattered the ask here.<br>Concept: never specify the conclusion |
| **C** | Tighten the data, the definitions and the exact question; loosen the conclusion and never specify it. | **Correct.** Analysis: tight data, open conclusion.<br>Concept: Analysis: tighten the data, the definitions and the exact question; loosen the conclusion |
| **D** | Loosen the exact question so the AI can pick a more interesting one. | **Wrong —** a loose question gives a loose, unfocused answer.<br>Concept: tighten the exact question |

**CONCEPT TESTED**

Analysis: tighten the data, the definitions and the exact question; loosen the conclusion. Backwards, you get an analysis that flatters the ask. Brainstorming loosens format; drafting tightens everything.

## Q9 · ANSWER · DOMAIN 3 — BEYOND TEXT

**Correct: B**

| | | |
|---|---|---|
| **A** | Trust it, since 18 numbers is small enough for the AI to get right at a glance. | **Wrong —** a number in prose is a guess in the shape of an answer.<br>Concept: the silent failure |
| **B** | Ask it to write and run code, show the code it ran, and give a verifiable specific such as the row count first. | **Correct.** Demand code that ran, and look for the block.<br>Concept: Code execution is a tool the model can call, and it chooses whether to |
| **C** | Tell it to "think step by step" and accept the new answer. | **Wrong —** step by step is obsolete and does not make the model call code.<br>Concept: code execution is a tool it chooses to call |
| **D** | Switch to a model with a larger context window. | **Wrong —** context size is not the issue; the code never ran.<br>Concept: did the code actually run? |

**CONCEPT TESTED**

Code execution is a tool the model can call, and it chooses whether to. Ask explicitly: "Write and run code. Show me the code you ran." Look for the code block, and still check totals and column meaning.

## Q10 · ANSWER · DOMAIN 4 — DISCERNMENT

**Correct: A**

| | | |
|---|---|---|
| **A** | Require a citation you can open for each figure, verify it against a primary source, and have a human reviewer sign off before it ships. | **Correct.** Regulator-bound? Verify sources, human sign-off.<br>Concept: Never ship on a draft alone: regulated or confidential work, public or legal communications, financially material figures |
| **B** | Send it, because the draft is fluent and professional-sounding. | **Wrong —** fluency and confidence cannot flag an error.<br>Concept: what a plausible error looks like |
| **C** | Keep it, because precise numbers are a sign the AI knows the topic. | **Wrong —** precision is the tell of a fabricated specific, not proof of accuracy.<br>Concept: fabricated specifics |
| **D** | Ask the AI whether the number is right and accept its confident yes. | **Wrong —** asking the same system to grade itself repeats the sycophancy problem.<br>Concept: check against sources, not feelings |

**CONCEPT TESTED**

Never ship on a draft alone: regulated or confidential work, public or legal communications, financially material figures. Real figures come with a citation you can open. The accountability does not transfer to the tool.

# Answer Key

Quick reference for all 10 questions

| | | |
|---|---|---|
| Q1 | B | Sparse data means confident blends, not facts. |
| Q2 | D | Name trusted sources, demand dated citations. |
| Q3 | A | Repeated context belongs in a project. |
| Q4 | C | Topic changed? New chat, paste two facts. |
| Q5 | B | Ask for the counter-argument, not the defense. |
| Q6 | D | Rubric + number beats "be harsh". |
| Q7 | A | Think hard for multi-trade-off decisions only. |
| Q8 | C | Analysis: tight data, open conclusion. |
| Q9 | B | Demand code that ran, and look for the block. |
| Q10 | A | Regulator-bound? Verify sources, human sign-off. |