AI Prompting: Describing, Evaluating and Validating Outputs

Description and Discernment

Made by Taha Shabbir

Based on: AI Prompting in 2026 - The AI Agent Factory by Panaversity

1. Describe: Brief the AI Well

Professional AI collaboration converges through iteration. It almost never succeeds in one shot.

Describe
State product, process, performance

>

AI produces
Output arrives, fluent and confident

>

Discern
Evaluate product, process, performance

>

Refine
Feedback that names, explains and directs

The five slots of a good prompt

Role - who the AI should be for this task.

Context - what it cannot know unless you say it: audience, situation, sources.

Task - one primary verb: summarize, compare, draft, identify.

Constraints - length, tone, what to include and leave out.

Format - a table, bullets, a memo, a draft email.

A quick question needs a task. A client deliverable needs all five.

Weak vs Strong prompt

Weak: "Write a summary of our quarterly operations." Result: three paragraphs that could describe any company.

Strong: You are an operations analyst. I am preparing a one-page update for our regional director, who cares about throughput and cost. Summarize the attached Q3 data, covering only the three metrics that moved more than 10 percent against target. Format as a headline and three one-sentence bullets.

The model did not get smarter. The prompt did.

Novice vs Power user

Novice: "which car should I buy?" Result: a generic three-model list, no use for the decision.

Power user: attaches insurance quotes, dealer quotes and a cost spreadsheet; adds context (30-minute commute, two kids); asks "What are the trade-offs? Read everything and think hard."

Result: a five-year cost comparison, a safety analysis, and a recommendation with the conditions under which it flips.

AI is like a really smart fresh college grad: highly motivated, but it does not know much about you yet. Brief it like one.

2. How AI Knows Things

It read the internet; it did not live in the world. Frequency in training data is roughly equal to reliability.

Pretrained knowledge

Strong: cooking, common medical advice, top movies, popular programming languages.

Sparse: quasars, Cantonese, regional history, niche professional knowledge.

Absent: your company's data, your calendar, anything after the cutoff.

Example: a folk game from a grandmother's village came back as three confident, almost entirely wrong paragraphs.

Three retrieval modes

Pretrained (seconds): fires on "what is X." Blind after the cutoff.

Web search (tens of seconds): fires on "this week" or a city name. Cites popular sources and does not check they are current.

Deep research (minutes): fires on "research thoroughly, with citations." Overkill for simple questions.

Context: six layers

1 System prompt (hidden) | 2 Memory | 3 Tool descriptions | 4 Your prompt | 5 Chat history | 6 Uploaded files

Before any non-trivial prompt: attach the document to match, state constraints it cannot infer, summarize prior context in one paragraph, name the output format, name the audience.

Context rot

One long chat across a workout, a spreadsheet and a thank-you note makes answers vaguer and contradictory. **Start a new chat when the topic changes** and paste in only the facts that matter.

Projects

Pasted the same files or audience into two chats? Put that context in a project, inherited by every chat inside it. **You reset the chat, not the context.**

Memory

Tools write a note about you and load it before every chat. Read it once, delete stale lines, use incognito when it would mislead. **Keep client and student details out of it.**

Confidence is not a signal of correctness. Ask "how would it know that?" and check anything important against a primary source.

3. Reasoning, Sycophancy and Rubrics

"Think step by step" is obsolete. Models are trained on thumbs-up, and agreeing earns more of them.

Reasoning: "think hard"

Modern models have built-in reasoning modes you invoke in plain language.

"Read everything. Think hard. Then tell me: 1. the three trade-offs that matter, 2. which you'd choose and why, 3. under what conditions that flips."

Use it for multi-input, multi-trade-off questions. **Skip it for** quick lookups, short summaries, casual brainstorming.

Longest reliable task: 7 minutes to 1 hour, doubling about every 7 months (METR, 2025).

Sycophancy: how to neutralize it

The model opened with "yes" or "correct" 10x more often than "no" (Washington Post, 2025).

Bait	Neutral rewrite
"Find evidence that this strategy will work."	"Evaluate this strategy. List the strongest arguments for and against."
"Why is approach A better than B?"	"Compare A and B. Score each on cost, risk, and time."
"Help me defend my decision to hire X."	"What is the strongest counter-argument?"
"Confirm that this code is correct."	"Find any bug, edge case or unstated assumption. If none, say so."

Use a rubric, and force a number

Without a rubric, criteria collapse into "great work." Score each criterion 1 to 10 with a one-sentence justification, then ask how to raise each score, even the high ones ("if something is at 9, how do I get to 9.5?"). Choosing between a 6 and a 7 forces the model to actually look. In the example, the same AI that praised a tie-dye party business idea gave it **8 out of 100** once it had a rubric.

Find, defend, confirm, prove hand the AI a conclusion. Replace them with evaluate, compare, critique. You decide when to stop.

4. Iterate, Then Go Beyond Text

The average AI answer is the average internet answer. The way past it is a loop, not a cleverer prompt.

1 Load all context, constraints, files, audience > 2 Ask for 3 to 5 options. Do not expand yet > 3 Give explicit feedback: reject what, accept why > 4 New options. Iterate two or three rounds > 5 Only then expand. Grade it out of 10; fix

Kind of ask	Tighten	Loosen	If backwards
Brainstorming	problem, audience, what to avoid	structure, tone, format, length	five rewordings of one idea
Research	scope, sources, what counts as evidence	the answer; never say what you expect	a map of your assumptions
Drafting	everything: voice, length, structure	almost nothing	a fluent draft of the wrong thing
Analysis	data, definitions, exact question	the conclusion; never specify it	an analysis that flatters the ask

Data analysis: did the code run?

The model chooses whether to call code. On small questions it may answer from a glance, which looks identical to real analysis. Ask: **"Write and run code. Show me the code."** Demand row count and column names first, and still check totals and labels.

Multimodal: images and audio

Image in: strong on scenes, whiteboards, handwriting; weak on fine detail and counting, so verify detail. Image out: let a text AI write the prompt, put sign text in quotes, name the aspect ratio. Audio: summarize decisions, open questions and action items by owner.

Small apps from one prompt

One prompt can produce a working tool (like a Pomodoro timer) in a side panel. This is an **artifact**: persistent, editable in place and shareable by link.

Specify loosely while diverging, precisely while converging. Outline before you draft. Treat AI analysis like a sharp junior analyst.

5. Discernment: Evaluating Output

Evaluation is a check against three references: your requirements, the source material, and your field's standards.

Fabricated specifics

"63 percent of firms adopted an AI tool, up from 41." Precise, plausible, no source. The precision is the tell.

Confident tone

"Is this clause enforceable?" Four assured sentences. Assurance is not evidence.

Internal contradiction

Page two says "$2 billion", page eight builds on "$2.6 billion". Long outputs need a consistency pass.

Claimed actions

"I've emailed that to your team." It performs no action it has no tool for. Treat it as unverified.

Missing pieces

A thorough-looking comparison that skipped the most important file. Nothing on screen shows what is absent.

Build the check into the prompt

Allow "I don't know." Restrict to attached sources. Require citations you can open. For numbers, require code that ran.

When human review is non-negotiable

Threshold	Ask
Stakes	What is the cost if this is wrong?
Reversibility	Can it be undone once sent?
Audience	External, executive or regulatory?
Regulatory exposure	Does a rule, contract or law govern this?

Never ship on a draft alone

Final client deliverables; financially material calculations; anything with regulated or confidential data; public or legal communications.

The slow creep: if rounds three to five of a proposal changed almost nothing, that flat curve is the signal to stop prompting and get a colleague's fresh read.

Cross-model checking: same-family models share blind spots, so use different families.

The one sentence to remember: get the right context in, keep the wrong context out. Accountability stays with you.

A marketer asks for five campaign concepts with a tight brief: exact tagline, a required three-part structure, strict tone, a 40-word cap. He gets five near-identical ideas. He adds more detail and asks for twenty — still neighbours of one idea. Best next move?

A Add two strong example concepts to the brief, so the model has a clearer picture of what a winning idea looks like.

B Switch to the most capable model tier, since a stronger model should find more original angles in the same brief.

C Ask for fifty concepts instead of twenty, then shortlist the few that are genuinely different from the rest.

D Keep the product, audience and what to avoid, but drop the format rules until one direction is chosen.

Quiz: Q5 - Prompting and Task Execution

Course: AI Prompting in 2026

Correct: D

Question

A marketer asks for five campaign concepts with a tight brief: exact tagline, a required three-part structure, strict tone, a 40-word cap. He gets five near-identical ideas. He adds more detail and asks for twenty - still neighbours of one idea. **Best next move?**

A	Add two strong example concepts to the brief.	**Wrong** - more grip, not less. Examples anchor the model even harder; you get variations on your two examples.
B	Switch to the most capable model tier.	**Wrong** - not a capability problem. Any tier returns near-copies when the brief has pinned the shape shut.
C	Ask for fifty concepts instead of twenty, then shortlist.	**Wrong** - volume isn't variety. Twenty already proved it: more of the same brief means more neighbours.
D	Keep the product, audience and what to avoid; drop the format rules until a direction is chosen.	**Correct.** Keep the full situation, loosen the shape of the answer to open up options, then tighten again to build the winner.

CONCEPT TESTED - MATCH THE GRIP TO THE ASK

Situation: always full. Output shape: **loose for brainstorming, tight for drafting**. (The grip dial: loose for options.)