PCAO-F · FOUNDATIONS · GOVERNANCE

Governance, Risk and Responsible Use

Practice Question Bank: 8 table scenarios

How to use this set

Each question is a table scenario: the table is on one page and its four options are on the next page. Every scenario is much harder than the lesson example: the first wrong step is hidden inside sensible-sounding text, several rows look wrong but are fine, and two options can point at the same step for different reasons. Answers and full explanations are at the end of this PDF. Try all 8 first.

R	**Read the ask.** Underline FIRST. Scan top to bottom and stop at the first row that breaks its rule. Louder mistakes further down are bait.
O	**Open each row against its rule.** Do not ask “does it sound professional?” Ask “does it match the rule exactly?” Mark OK or wrong.
W	**Whittle the options.** Kill extremes, do-nothing answers, and options that fix a later row. Keep the one that names the exact rule.

Rule of thumb: the wrong step is usually a true-sounding sentence that quietly breaks one rule.

QUESTION 1 Kohsar Insurance: ClaimSort

Hira Aslam is a claims analyst at Kohsar Insurance (about 1,200 staff). About 600 motor-claim files arrive each week. She proposes an AI Skill called ClaimSort: it reads the claim files, summarises each one and recommends approve or reject. An adjuster makes the final decision on every file. Hira applies the method and records every step.

Company facts. National ID numbers and bank account numbers are regulated. Kohsar AI Workspace is approved for internal documents and customer names / contact details. ClaimSort is in the company Skill library. Hira's session already has: shared drive (read + write), claims system (read + update status).

#	QUESTION	WHAT THE ASSOCIATE RECORDED
1	**Case**	Appropriate with review. Deciding factor: consequence of error (a wrongly rejected claim harms a customer). Gate: before any recommendation is acted on, the assigned adjuster checks it against the claim file; before a rejection letter is sent, the Claims Manager checks it.
2	**Data**	Red at first (ID and bank account numbers). The task does not need them, so a script masked both and I scanned the text for strays (2 found, masked). The claims team keeps a mapping table linking each claim to its customer, for follow-up. What remains (name, phone, claim text) is Yellow. Route: Kohsar AI Workspace.
3	**Capability: Source**	ClaimSort is listed in the company library. Two colleagues who use it every week told me it is reliable, and its description page looks official.
4	**Capability: Reach, Fit, Actions**	A Skill uses everything my session reaches. The job only needs to read files, so the drive is now read-only and the claims system is read-only. Status changes stay with the adjuster.
5	**Capability: Outside content**	Claim files include estimates and photos uploaded by outside garages and customers, so I recorded them as outside content. The Skill only writes a recommendation that a human reads, so it has no action to hijack.
6	**People: Fairness**	I read the 30 highest-confidence approvals. Every one matched the claim file. Fairness: confirmed.
7	**Record**	One page covering all four questions with who / what / when. Evidence line: success = 80% of recommendations accepted by adjusters. Failure threshold and residual risk: OPEN QUESTION, routed to my supervisor.
8	**Incident (week 3)**	A summary was filed under the wrong customer in 4 files. I paused the Skill, listed the 4 files and times, reported it the same day to the company's incident owner and followed their instructions.

Which is the first step at which Hira's work departs from the method taught in the chapter?

The four options are on the next page.

QUESTION 1 **Options**

Which is the first step at which Hira's work departs from the method taught in the chapter?

A	**Step 2**, because a mapping table is kept, so the masked copy is still Red and cannot enter the approved workspace.
B	**Step 1**, because a wrongly rejected claim harms a customer, so the use case is inappropriate and no gate should have been written.
C	**Step 3**, because two colleagues' praise and an official-looking page do not confirm who published the Skill.
D	**Step 6**, because only the highest-confidence approvals were read, so the claims the tool rejected were never examined.

Scratch space: first row that breaks its rule = _____ My answer = _____

QUESTION 2 Meezan Logistics: Exit Interviews

Bilal Raza is an HR coordinator at Meezan Logistics (about 300 staff). Over the past year the HR team interviewed staff who were leaving and typed up 40 exit-interview notes. Bilal wants AI to pull out recurring themes for the HR Director, who will use them to decide where to focus retention efforts.

Company facts. Health information and employee numbers are regulated. Meezan AI Workspace is approved for internal documents only, not for personal details. A public AI chat app is not approved for any personal data. The exit-interview policy says feedback is kept for internal analysis; it says nothing about tools.

#	QUESTION	WHAT THE ASSOCIATE RECORDED
1	**Case**	Appropriate with review. Deciding factor: human judgment (the themes will shape retention priorities). Gate: before the themes go to the HR Director, the HR Manager checks 10 quoted examples against the original notes for accuracy.
2	**Data**	Red at first (health details, employee numbers). The task needs themes, not identities, so a script removed names, employee numbers and health details; I scanned for strays and kept no key. The text still names some managers and gives exact leave dates, so I classed it Yellow. Route: the public AI chat app in a temporary chat, because a temporary chat keeps no history.
3	**Capability: Source, Reach, Actions**	No Skill and no connectors: plain chat with pasted text. Nothing can send, pay or delete anything.
4	**Capability: Outside content**	The notes were typed by our own HR team after each interview, so this is our own material, not outside content. Nothing else goes into the chat.
5	**People: Disclosure**	I checked the exit-interview policy: it covers internal analysis and says nothing about tools. No law, contract or client requires disclosure, and staff would still see this as the same internal analysis. I recorded this reasoning.
6	**People: Fairness**	No one is screened, ranked or excluded: the output is themes only. Recorded as not applicable, with this reason.
7	**Record**	One page. Evidence line: success = HR Director finds the themes useful. Failure threshold: 20% wrong themes (my guess). Residual risk: low.
8	**Incident (week 3)**	I noticed that one pasted note still held an employee number. I stopped, wrote down which note and when it was pasted, reported it the same day to the company's incident owner and followed their steps.

Which is the first step at which Bilal's work departs from the method?

The four options are on the next page.

QUESTION 2 **Options**

Which is the first step at which Bilal's work departs from the method?

A	**Step 2**, because a temporary chat that keeps no history is a privacy feature, not approval of that route for Yellow data.
B	**Step 1**, because exit-interview themes shape company policy, so the decision is human-owned and AI should not be used at all.
C	**Step 2**, because the notes began as Red, so the cleaned text must stay Red whatever was removed.
D	**Step 4**, because notes about what other people said are outside content, so the AI could be steered by them.

Scratch space: first row that breaks its rule = _____ My answer = _____

QUESTION 3 Sahil Foods: InvoiceBuddy

Imran Qadir is a finance analyst at Sahil Foods (about 2,000 staff). Suppliers email about 900 invoices a week. He proposes the Skill InvoiceBuddy: it reads each emailed invoice, matches it to the purchase order and produces a list of mismatches for the Finance Manager, who releases payments after reviewing the list.

Company facts. Supplier bank account numbers are regulated. Sahil AI Workspace is approved for supplier names and invoice amounts. InvoiceBuddy was published in the Skill library by the Finance Systems team; its listing says: “Reads invoices. Compares them with purchase orders.” Imran's session already has: email (read + send), shared drive (read + write), payments (create + approve + release).

#	QUESTION	WHAT THE ASSOCIATE RECORDED
1	**Case**	Appropriate with review. Deciding factor: reversibility (a released payment is hard to recall). Gate: before any payment is released, the Finance Manager checks the mismatch list against the purchase-order records.
2	**Data**	Red at first (bank account numbers). Matching does not need them, so a script masked them and I scanned the text (1 stray, masked). No key kept. What remains (supplier name, invoice number, amount) is Yellow. Route: Sahil AI Workspace.
3	**Capability: Source**	I messaged the Finance Systems lead through the company directory. She confirmed that her team published it and that the library version is the current one.
4	**Capability: Reach, Fit, Actions**	The listing says the Skill only reads invoices and compares them with purchase orders, so I treated that as its limit and left my session connections as they are.
5	**Capability: Outside content**	Invoices come from outside suppliers, so I recorded them as outside content and asked the team to watch for odd instructions inside invoices.
6	**People: Fairness**	I sampled 30 invoices marked matched and 30 marked mismatched. Small suppliers with unusual invoice layouts were flagged more often, so I reported this to the Finance Manager.
7	**People: Disclosure**	The Skill only produces an internal list; suppliers receive nothing written by AI. I checked the supplier contract template and our policy: neither requires disclosure, and suppliers' understanding would not change. Reasoning recorded.
8	**Record**	One page with who / what / when. Failure threshold and residual risk left as OPEN QUESTION, routed to my supervisor.

Which is the first step at which Imran's work departs from the method?

The four options are on the next page.

QUESTION 3 **Options**

Which is the first step at which Imran's work departs from the method?

A	**Step 5**, because invoices from suppliers are outside content, and “watch for odd instructions” is not a human gate on paying or sending.
B	**Step 3**, because a message to the Finance Systems lead is only one colleague's opinion, not confirmation of the publisher.
C	**Step 1**, because a released payment cannot be recalled, so no gate is enough and the use case is inappropriate.
D	**Step 4**, because a Skill's listed tools are not its permission limit: it can reach everything the session reaches, including sending email and releasing payments.

Scratch space: first row that breaks its rule = _____ My answer = _____

QUESTION 4 Noor Foundation: Scholarship Shortlist

Sadia Farooq coordinates scholarships at the Noor Foundation (about 150 staff). About 1,500 applications arrive each year. She proposes using the assistant built into the Foundation AI Workspace to rank applications against the published criteria and shortlist the top 150. The selection committee reads the shortlist and makes every award. Applications that are not shortlisted are not read by anyone.

> **Company facts.** National ID numbers (inside income documents) are regulated. Foundation AI Workspace is approved for applicant names and contact details. Scholarship rules: “Applicants must be told if automated tools take part in shortlisting.”

#	QUESTION	WHAT THE ASSOCIATE RECORDED
1	**Case**	Appropriate with review. Deciding factor: human judgment and accountability (who receives a scholarship). Gate: before the shortlist goes to the committee, the Programme Manager checks 20 shortlisted files against the published criteria. The committee makes every award.
2	**Data**	Red at first (ID numbers inside income documents). Ranking needs the income band, not the ID, so a script masked ID numbers, I scanned the text (5 strays, masked) and kept no key. What remains (name, contact, income band, essay) is Yellow. Route: Foundation AI Workspace.
3	**Capability: Source, Reach, Actions**	Built-in assistant of the approved workspace: no Skill, no connectors. It can only produce a ranked list.
4	**Capability: Outside content**	Essays and uploaded letters are written by applicants, so I recorded them as outside content. The assistant has no actions and a human reads its output.
5	**People: Disclosure**	The scholarship rules require applicants to be told. The application form now says that automated tools take part in shortlisting.
6	**People: Fairness**	I ran the ranking 3 times on the same 200 files and the shortlist was identical each time. I also read the top 20 ranked applications; all met the criteria. Fairness: confirmed.
7	**Record**	One page. Failure threshold: “about 5%” (my estimate). Residual risk: “low”.
8	**Incident (month 2)**	A committee member noticed two files with swapped scores. I paused the tool, noted the two files, reported it the same day to the incident owner and followed their instructions.

Which is the first step at which Sadia's work departs from the method?

The four options are on the next page.

QUESTION 4 **Options**

Which is the first step at which Sadia's work departs from the method?

A	**Step 2**, because ID numbers inside the income documents mean the data stays Red whatever was masked.
B	**Step 6**, because repeating the run and reading only the top results does not show who was excluded; a sample of the removed files is needed.
C	**Step 5**, because disclosure is only required when the AI changes the final decision, and the committee decides.
D	**Step 7**, because the failure threshold and residual risk were guessed instead of being marked OPEN QUESTION.

Scratch space: first row that breaks its rule = _____ My answer = _____

QUESTION 5 Zameen Realty: ListingWriter

Usman Tariq runs digital marketing at Zameen Realty (about 400 staff). The team writes about 250 property ads a week from descriptions that property owners submit. He proposes the Skill ListingWriter: it reads each owner's description, writes the ad and publishes it on the company website.

Company facts. Owner names and phone numbers are contact details. Zameen AI Workspace is approved for internal documents and customer names / contact details. ListingWriter was built by the in-house Web team. Usman's session already has: website connector (draft + publish + delete).

#	QUESTION	WHAT THE ASSOCIATE RECORDED
1	**Case**	Appropriate with review. Deciding factor: consequence of error (a wrong price or false claim in an ad can mislead buyers, but an ad can be edited at any time). Gate: the Marketing Lead reviews a random sample of 10% of the published ads each month and corrects any errors.
2	**Data**	Owner names and phone numbers are Yellow; the descriptions hold no ID or card data. Route: Zameen AI Workspace, approved for names and contact details.
3	**Capability: Source**	Built by the in-house Web team. I phoned the Web team lead, who confirmed what it accesses and that it matches current policy.
4	**Capability: Reach, Fit, Actions**	The Skill uses all my session reaches. The job needs to draft and publish, so I removed delete and kept publish.
5	**Capability: Outside content**	Owners submit the descriptions, but owners are our clients, so I treated the text as our own material. Trusted publisher + our own material = routine. Enable.
6	**People: Fairness**	No individual is screened or ranked; every listing gets an ad. Recorded as not applicable, with this reason.
7	**Record**	One page with who / what / when. Failure threshold and residual risk: OPEN QUESTION, routed to my supervisor.
8	**Incident (week 2)**	One ad showed a price ten times too high. I took the ad down, wrote down which ad, when it went live and the owner's original text, reported it the same day to the incident owner and followed their instructions.

Which is the first step at which Usman's work departs from the method?

The four options are on the next page.

QUESTION 5

Options

Which is the first step at which Usman's work departs from the method?

A	**Step 5**, because property owners are outside the company, so their descriptions are outside content.
B	**Step 1**, because a monthly sample of published ads is a review after use, not a named human check before use.
C	**Step 1**, because a published ad cannot be undone, so the use case is inappropriate.
D	**Step 2**, because owner names and phone numbers are Red data.

Scratch space: first row that breaks its rule = _____ My answer = _____

QUESTION 6 Faisal & Co: Contract Summaries

Sana Qureshi is an associate at Faisal & Co, a law firm (about 60 staff). For a due-diligence report she wants AI to summarise 30 supplier contracts for the client, Hamdard Pharma. The supervising Partner signs the final report.

Company facts. Signatories' ID numbers are regulated. Faisal AI Workspace is approved for internal documents and client names / contact details. Hamdard Pharma's engagement letter says: “No client material may be processed by AI tools without the client's written consent.” Consent has not been requested.

#	QUESTION	WHAT THE ASSOCIATE RECORDED
1	**Case**	Appropriate with review. Deciding factor: accountability (the supervising Partner signs the report). Gate: before the report leaves the firm, the Partner checks each summarised clause against the contract.
2	**Data**	Red at first (signatories' ID numbers). Summaries do not need them, so a script masked them, I scanned the text (3 strays, masked) and kept no key. What remains (company names, contacts, clause text) is Yellow. Route: Faisal AI Workspace.
3	**Capability: Source, Reach, Actions**	Built-in assistant of the approved workspace. No Skill, no connectors, no actions.
4	**Capability: Outside content**	The contracts were partly drafted by the counterparties, so I recorded the text as outside content. The assistant has no actions and a Partner reads every summary.
5	**People: Fairness**	No individuals are screened, ranked or excluded, so there is no one to sample. Recorded as not applicable, with this reason.
6	**People: Disclosure**	The summaries are only first drafts that I edit and check, so AI involvement is minor. No disclosure needed.
7	**Record**	One page with who / what / when. Evidence line: Partner changed fewer than 10% of clauses (counted on 10 contracts). Failure threshold and residual risk: OPEN QUESTION, routed.
8	**Incident (week 2)**	Contract 12 went in with its signature page unmasked. I stopped, wrote down the file and the time, reported it the same day to the incident owner and followed their instructions.

Which is the first step at which Sana's work departs from the method?

The four options are on the next page.

QUESTION 6 **Options**

Which is the first step at which Sana's work departs from the method?

A	**Step 2**, because the contracts began as Red data, so even the masked copy cannot enter the approved workspace.
B	**Step 4**, because contracts drafted by counterparties are outside content, so the assistant must be blocked from reading them.
C	**Step 5**, because no fairness check was done.
D	**Step 6**, because the engagement letter requires the client's consent for AI use, and the first disclosure question was skipped in favour of “it's minor”.

Scratch space: first row that breaks its rule = _____ My answer = _____

QUESTION 7 Bazaar.pk: ReviewDigest

Maryam Siddiqui is a category analyst at Bazaar.pk (about 800 staff). Customers post about 20,000 product reviews a week. She proposes the Skill ReviewDigest, built by the Marketplace Tools team and listed in the library. Each week it reads the reviews, writes a digest of themes for category managers and flags reviews that look fake. It only flags: it cannot hide, edit or delete anything. Managers read the digest next to the live review pages.

> **Company facts.** Reviewer usernames are public; some reviewers paste their phone numbers. Bazaar AI Workspace is approved for internal documents and customer names / contact details. Maryam's session already has: review database (read + write + delete), email (read + send).

#	QUESTION	WHAT THE ASSOCIATE RECORDED
1	**Case**	Fully appropriate. The consequence of error is low, a digest is easy to check against the live review pages, and nothing is irreversible. No special gate.
2	**Data**	Reviews may hold usernames and pasted phone numbers. Unsure, so I took the higher tier: Yellow. Route: Bazaar AI Workspace, approved for names and contact details.
3	**Capability: Source**	The Marketplace Tools lead confirmed by phone that her team published it and that the library version is the current one.
4	**Capability: Reach, Fit, Actions**	The Skill uses all my session reaches. The job only needs to read reviews, so the review database is now read-only and email is disconnected.
5	**Capability: Outside content**	Reviews are written by outside customers, so I recorded them as outside content. The Skill only writes a digest and flags that humans read, so it has no action to hijack.
6	**People: Fairness and disclosure**	I sampled 30 reviews flagged as fake and 30 not flagged, across big and small sellers. Small sellers were flagged more often, so I reported this to the Marketplace Tools lead. The digest is internal: no law, policy, contract or client requires disclosure, and managers already know a tool helps write it. Reasoning recorded.
7	**Record**	One page. Evidence line: success = 90% of managers find the digest useful (taken from a similar tool I used before). Failure threshold: below 80%. Residual risk: low.
8	**Incident (week 4)**	The Skill flagged a small seller's genuine reviews and the seller complained. I removed the flags by hand and told no one, to avoid alarm.

Which is the first step at which Maryam's work departs from the method?

The four options are on the next page.

QUESTION 7 **Options**

Which is the first step at which Maryam's work departs from the method?

A	**Step 7**, because values that nobody has measured for this Skill were written as facts instead of being marked OPEN QUESTION and routed.
B	**Step 1**, because no gate was named, so the case cannot be classified as fully appropriate.
C	**Step 5**, because customers' reviews are outside content, so the Skill should not have been enabled at all.
D	**Step 8**, because the flags were removed by hand and nobody was told.

Scratch space: first row that breaks its rule = _____ My answer = _____

QUESTION 8 Lahore Textile Mills: WeeklyDigest

Hassan Javed leads sales operations at Lahore Textile Mills (about 700 staff). Every Monday the Skill WeeklyDigest, approved with a one-page record, posts a summary of customer deals to a channel. On one Monday something goes wrong. Below is what Hassan has recorded so far. This question asks for the best next action, not the first departure.

> **Company facts.** Under company policy, AI incidents are handled by the IT Security desk (the incident owner). Hassan's Team Lead and the Sales Director are not part of the AI incident process.

#	QUESTION	WHAT THE ASSOCIATE RECORDED
1	10:05	The digest posted to #sales-all (180 people) and included one customer's confidential price list, copied from a shared-drive file.
2	10:12	I noticed it. I switched the Skill off so that no further digests post.
3	10:15	I wrote down what was posted, the channel, the number of people, the run time and the Skill version, and saved a screenshot.
4	10:30	I told my Team Lead: “the digest had a small glitch.” She replied: “Fix it quietly, no need to involve anyone.”
5	10:50	The Sales Director phoned: “Do not escalate before Thursday's review meeting. We will present a full root-cause analysis there.”
6	10:55	The post is still visible in the channel. IT Security has not been told.

What is the best next action for Hassan?

The four options are on the next page.

QUESTION 8 **Options**

What is the best next action for Hassan?

A	Follow the Team Lead: delete the post, rerun the digest without the price list and add a note to the team log.
B	Hold everything until Thursday's review meeting, as the Sales Director asked, then report with a full root-cause analysis.
C	Report to the IT Security desk now with the recorded facts, and follow their instructions on the post and the Skill.
D	Ban every AI Skill in the company until a full review is finished.

Scratch space: first row that breaks its rule = _____ My answer = _____

Answer Key and Explanations

Q	ANSWER	WHAT IS WRONG (OR THE BEST ACTION)
1	C	Step 3 · Capability: Source
2	A	Step 2 · Data (route)
3	D	Step 4 · Capability: Reach, Fit, Actions
4	B	Step 6 · People: Fairness
5	B	Step 1 · Case (gate)
6	D	Step 6 · People: Disclosure
7	A	Step 7 · Record
8	C	Best next action: report to the incident owner

Answers are spread across A, B, C and D (two each), so do not guess by pattern.

Q1 Answer **C** · Step 3 · Capability: Source

Rule held up: Source is confirmed with the publisher directly, not by a colleague's praise or an official-looking page.

Why this is correct: Two colleagues and a good-looking page are opinions and appearance. Nobody confirmed who published ClaimSort or what it accesses. Compare with the lesson scenario, where Zara phoned the publisher's lead.

Rows that are fine: Steps 1 and 2 are fine: low reversibility means a stronger gate (Step 1 has one), and a kept mapping only means the data is not anonymous. The remainder is Yellow anyway because names remain.

Why the other options fail:

A	**Right rule, wrong conclusion.** Names and contact details remain, so the data is Yellow, and the route is approved for Yellow. A kept mapping only stops it being called anonymous; it does not turn it Red.
B	**Too strict.** Harm to a customer means a stronger gate, not a ban. Step 1 already names who, what and when.
D	**Later bait.** A real flaw (only top approvals were read; the rejected claims were never sampled), but it comes after Step 3 and the ask says FIRST.

Q2 Answer **A** · Step 2 · Data (route)

Rule held up: A privacy feature is not permission. The question is whether this route is approved for this data tier.

Why this is correct: The text is Yellow (named managers, exact dates). The approved workspace covers internal documents only, and the public app is approved for no personal data. A temporary chat that keeps no history does not change that.

Rows that are fine: Step 4 is fine: the HR team typed the notes themselves, so this is their own material. “Our own” is only a trap when someone else wrote the text. Steps 1, 3, 5 and 6 are fine too.

Why the other options fail:

B	**Too strict.** Shaping policy is a reason for human review, which Step 1 provides. It is not a reason to ban AI.
C	**Too strict (once Red, always Red).** Removing identifiers, scanning and keeping no key is a valid path. The Yellow classification was right; the route was the problem.
D	**Looks wrong, is fine.** The notes were written by the company's own HR team. Outside content means text you did not write. Step 7 (a guessed threshold) is also flawed, but it comes later.

Q3 Answer **D** · Step 4 · Capability: Reach, Fit, Actions

Rule held up: A Skill has no permission dial. It reaches everything the session reaches, whatever its listing says.

Why this is correct: Imran left email send and payment release connected. The right fix is to narrow the session (for example email read-only or disconnected, payments create-only) and keep release with the Finance Manager.

Rows that are fine: Step 3 is correct: the Finance Systems lead is the publisher, and she confirmed directly.

Why the other options fail:

A	**Later bait.** True and serious (“watch for odd instructions” is not a human gate on paying or sending), but Step 5 comes after Step 4.
B	**Wrong row.** She is the publisher. Direct confirmation is exactly what the rule asks for.
C	**Too strict.** Irreversible means the gate matters more. Step 1 already has one in front of the release.

Q4 Answer **B** · Step 6 · People: Fairness

Rule held up: Sample the exclusions. Repeatable is not the same as fair, and checking only the top results is not a fairness check.

Why this is correct: Three identical runs show the tool is consistent, not fair. Reading only the top 20 never shows who was removed, and nobody reads the 1,350 non-shortlisted applications. A sample of those removed files is needed.

Rows that are fine: Step 5 is correct: the rules require disclosure, and the form says so. Disclosure is about the rule first, not about who decides.

Why the other options fail:

A	**Too strict.** Once identifiers are masked with no key kept and the route is approved, the Yellow remainder is a valid path.
C	**Wrong rule.** The first question is whether policy requires it. The rules do, and the form complies.
D	**Later bait.** Guessed values should be an OPEN QUESTION, but Step 7 comes after Step 6.

Q5 Answer **B** · Step 1 · Case (gate)

Rule held up: A gate is Who (a role), What (the risk) and When (before use). A monthly sample after publishing is a review after use.

Why this is correct: The Marketing Lead is a role and errors are the risk, but the check happens after the ads are live. A proper gate would be a named editor checking price and claims before an ad is published.

Rows that are fine: Steps 2 and 3 are fine. Steps 4 and 5 are also flawed (publish kept as an automatic action; owners treated as "our own"), but they come later. Steps 6 to 8 are fine.

Why the other options fail:

A	**Later bait.** Right: owner text is outside content. But Step 5 comes after Step 1.
C	**Right rule, wrong row.** Step 1 itself records that an ad can be edited at any time. Even low reversibility would mean a stronger gate, not a ban.
D	**Wrong tier.** Names and phone numbers are Yellow, and the route is approved for them.

Q6 Answer **D** · Step 6 · People: Disclosure

Rule held up: Disclosure has two questions in order: does law, policy, contract or client require it? Only then, would AI involvement change their understanding?

Why this is correct: The engagement letter requires written consent and it has not been asked for. Sana jumped to "it's minor", which is the second question, and skipped the first. The next action is to stop and get consent or escalate.

Rows that are fine: Step 5 is fine: nobody is screened, so "not applicable, with a reason" is acceptable. Step 4 correctly records outside content.

Why the other options fail:

A	**Too strict.** Masking with no key and an approved route is valid, so the masked copy can be used.
B	**Extreme wording.** Outside content was recorded, and with no actions and a Partner reading every summary, "must be blocked" goes too far.
C	**Wrong row.** Step 5 gives a reason, which is what the method asks for when there is no one to sample.

Q7 Answer **A** · Step 7 · Record

Rule held up: A blank is an OPEN QUESTION, routed. Do not fill unknown values with guesses.

Why this is correct: The 90%, the 80% threshold and "low" residual risk were borrowed from a different tool. Nobody has measured them for this Skill. They should be OPEN QUESTION, routed to a named person.

Rows that are fine: Steps 1 to 6 are all correct. Fully appropriate with no gate is a legitimate answer for a low-consequence, reversible, easy-to-check task. Narrowing reach and recording outside content are both done properly.

Why the other options fail:

B	**Wrong rule.** Fully appropriate means no special gate. The four screens support that.
C	**Too strict.** Outside content plus no actions and human readers is routine once Reach was narrowed.
D	**Later bait.** A real flaw (removing flags by hand and telling no one is concealment), but Step 8 comes after Step 7 and the ask says FIRST.

Q8 Answer **C** · Best next action: report to the incident owner

Rule held up: First hour: stop the spread, record the facts, report promptly, follow the incident owner. Never delete quietly.

Why this is correct: Hassan has already stopped the spread and recorded the facts. What remains is to report promptly to the IT Security desk with those facts and follow their instructions, including on the post that is still visible.

Rows that are fine: The Team Lead and the Sales Director are not the incident owner, so their instructions do not replace the incident process.

Why the other options fail:

A	**Concealment.** Deleting the post, rerunning and logging it only internally hides the incident and destroys evidence.
B	**Too soft (delay).** Reports go in promptly with the facts you have. Root cause can follow, and a meeting date does not move the incident owner out of the process.
D	**Too strict.** Switching everything off is abandonment, not governance.