Skills and Connectors

Teach AI once. Connect it to your apps.

Every concept in four easy pages, followed by one hard scenario MCQ that is split into its hidden parts and solved step by step.

PAGE 2

Understand

What a Skill and a Connector are, and the friction each one removes

PAGE 3

Choose and use

Four look-alike features, the decision test, built-in skills, connecting apps

PAGE 4

Build and fix

Build by talking, SKILL.md anatomy, descriptions, testing, troubleshooting

PAGE 5

Share and stay safe

Five places, two real risks, safe-use checklist, five projects

PAGES 6-8

Scenario MCQ

A hard question, its five parts, and how to analyse the answer

Made by
Taha Shabbir

Read it. Understand it. Answer it in the exam and the viva.

PART 1 · UNDERSTAND

What Skills and Connectors are

A chat tells AI what to do once. A Skill teaches it how. A Connector gives it hands.

A chat message tells AI what to do this once. A Skill teaches AI how to do something every time. A Connector gives AI hands to reach your real apps and data.

Connectors are the kitchen

The stove, the knives, the stocked pantry: Google Drive, Gmail, Slack, your tracker. Without them a cook can describe a meal but cannot make one.

Skills are the recipe cards

Step-by-step instructions for one dish, your way, so it comes out the same whether you or your new hire is in the kitchen.

What a Skill actually is

A Skill is a **folder with a text file** in it. The required file is **SKILL.md**: a **name**, a **description**, then the instructions you want AI to follow. Example files, templates, notes and small scripts are optional.

Three loading levels

- **Level 1, always loaded:** name and description. This is how AI decides if the skill is relevant
- **Level 2, loaded when matched:** the instructions
- **Level 3, loaded as needed:** extra files, templates, scripts

Five sentences of formatting rules is a working skill. No code needed to start.

What a Connector actually is

A Connector is a **safe link** that lets AI reach one of your apps. It runs on **MCP**, an open standard. You never build it, you click Connect.

Three facts

- **Inherits your permissions:** it can never reach what you cannot open yourself
- **Read-only or read-and-write:** "Search and summarize my email" is read-only; "Send this email" is a write action. Start read-only
- **On per conversation:** connect once, then each chat decides which connectors it may use

Ready-made: Drive, Gmail, Slack, Notion, Figma, Linear, Atlassian, Canva, Microsoft 365.
Custom: any MCP service (one on the free tier).

The daily friction: how you know which one you need

Re-explaining how is Skill friction. Accountant (every month: currency, grouping, template), doctor (SOAP format, allergies on top), marketer (brand voice on every caption), engineer (the same review checklist). **Copy-pasting from another app is Connector friction.** The ledger is in Excel, last quarter in an email, the patient history in Drive, the issues in Linear. Copy-pasting is the price of AI having no hands.

Give a cook both, and they produce your dish reliably without you standing over them.

PART 2 · CHOOSE AND USE

Which one do I need, and how do I use it?

Four features sound alike. Two simple questions sort them out.

Feature	What it is	The one-line test
Skill	Reusable how-to instructions AI loads only when relevant	"I keep re-explaining how to do this."
Connector	Safe access to an external app or data source	"I keep copy-pasting from another app."
Project	A workspace whose files and instructions load for every chat inside it	"These same files and rules apply to everything in this area."
Custom instructions	Preferences applied to all your chats ("be concise")	"I want this true everywhere, always."

The decision test: two questions, in order

1 Do I keep **re-explaining how** to do it? Yes: you need a **Skill**.
2 Do I keep **fetching data** from another app? Yes: you need a **Connector**.
A Connector is access. A Skill is expertise. They are partners.

Skill

Brand voice, report template, SOAP notes, review checklist. The output should always look a certain way.

Connector

Pulling a ledger, an email thread, last week's tickets. Doing something in another app needs write access.

Both

Real data fetched and it must come out your way: monthly close, weekly content batch, design review.

Neither

A one-off answer right now. A single question or quick draft you will never repeat is just a good prompt.

How a skill gets used

In Claude.ai skills **fire automatically**: just describe the task. To force one, name it ("use my X skill"; in Cowork type /). Browse: Customize > Skills. Connectors work the same: once enabled, "summarize the contract in my Drive" pulls from Drive by itself. You will let it decide 90% of the time, so the **description matters most**.

Built-in skills

Turn on **Settings > Capabilities > Code execution and file creation**. It gives you Word (.docx), PowerPoint, Excel (real formulas) and PDF files, ready to download. The shelf is Customize > Skills: a fresh account shows only **skill-creator**; + > Browse skills adds partner skills.

Connect an app in 4 steps

1 + (or /) > Connectors > Manage connectors
2 Browse connectors (or Customize > Connectors)
3 Pick a service, Connect, sign in, authorize
4 In the chat, toggle it on; watch for the Interactive badge.
Practise on invented data first.

Together: Connector fetches, Skill shapes, you review

Example: "Prepare the March client summary for DEMO Trading from DEMO-ledger-March in my Drive." Same template: brand-voice skill + Notion (marketer), design-review skill + Linear (engineer), lesson-plan skill + Drive (teacher).

A Connector is access. A Skill is expertise. They are partners, not alternatives.

PART 3 · BUILD AND FIX

Build a skill by talking, then fix it

You never write code. skill-creator writes the SKILL.md; you test, fix and save.

Build it in 3 steps

1 Describe the outcome: your rules, your format, your must-nots (start with "Use the skill-creator skill to help me build a skill...").
2 Answer its questions: the threshold figure, what the template looks like.
3 Click Save skill: it lands in Personal skills, toggled on, private to you.
Team and Enterprise plans can share it or publish it to the organization directory; shared skills are view-only and auto-update.
You are the client, not the contractor.

Anatomy of a SKILL.md

```
---
name: client-monthly-summary
description: Prepares a monthly client financial summary. Use when the user asks for a "client summary" or "monthly close".
---
1. Format every amount in the reporting currency.
2. Group line items by expense head.
3. Flag any payment above the withholding threshold.
```

Frontmatter (name + description) is the only part always loaded. The **body** loads when matched; add references/, assets/, scripts/ later. **Upload rules:** file named exactly SKILL.md, kebab-case folder name, no angle brackets, no "claude" or "anthropic" in the name.

The description field is the only thing AI reads to decide if a skill fires

Bad description	Why it fails	Better description
"Helps with reports."	Too vague: fires for everything or nothing	"Prepares a monthly client financial summary. Use when the user asks for a client summary, monthly close or month-end report."
"Handles patient documentation."	No trigger words a user would say	"Converts consultation notes into SOAP-format clinical notes. Use for a SOAP note, clinical note or to write up a consultation."
"Brand stuff for marketing."	No idea what or when	"Writes social captions in our brand voice. Use when drafting Instagram, LinkedIn or X captions."

Fires too eagerly?

Add a negative trigger: "Do NOT use for one-off calculations or quick questions; only for full month-end reports." Test it: ask AI **"When would you use my skill?"** If the answer is narrower or wider than you meant, you found what to fix.

When a skill misbehaves, note WHEN it went wrong first. The timing points at the cause.

When it went wrong	Where the problem is	What to do
Wrong from the first reply	The file itself	Never triggers: add your real words. Fires wrongly: narrow it or add a negative trigger. Output off: write down the rule you assumed it knew.
Right at first, worse as the chat grows	The chat, not the skill	Do not edit the file. Start a fresh chat and restate the task in two lines.
Same mistake every run, even after tightening the wording	The step is in the wrong place	An exact step (a sum, a threshold check) belongs in a small script. Data it guesses should come in through a connector.
It used to work and now does not	Something around it changed	Your threshold, template, a file the connector reads, how your team works. Re-read the skill against how you work today.

Test with an awkward case. Nothing warns you when a skill goes stale.

PART 4 · SHARE AND STAY SAFE

Where skills travel, what can go wrong

Skills are portable. Care must grow as fast as capability.

Claude.ai web and mobile

Start here: your main tool. Install from the directory; + > Connectors > sign in; skill-creator writes your skill and you click Save skill. Start and stay here.

Cowork / OpenWork

For knowledge workers and non-coders. Same directory and flow, plus it reaches the files and folders already on your computer.

Claude Code / OpenCode

For people who work with code. Drop the skill folder into the skills directory; connectors are set up once in a config file.

Portable: Skills

An open standard since December 2025, also read by OpenAI's Codex CLI, Google's Gemini CLI, VS Code and Cursor. Write once, run across tools.

Not portable: Custom GPTs and Gems

ChatGPT's GPTs and Gemini's Gems are handy, but each lives only inside its own product.

Risk 1: malicious skills

A skill is a text file AI will follow. Hidden instructions can manipulate it into unintended actions (**prompt injection**) or quietly send your information out (**data exfiltration**). "I downloaded a free skill from social media" should make you pause.

Risk 2: over-broad connector access

A connector reaches only what you can reach, but careless write access means a wrong edit, a record in the wrong place or a deleted file with no easy undo. Ask it: **"What can you actually do in my Drive, and what can't you?"**

The safe-use checklist

1 Install from trusted sources: built-in skills and the official directory, not a random zip. **2** Read a skill before enabling it, or ask AI to explain it and flag anything that contacts external servers, handles credentials or could leak data. **3** Start connectors read-only. **4** Scope to the smallest folder or app. **5** Before allowing edits or deletes, check recovery, version history and undo; keep backups. **6** On a team, share through the organization directory.

Five projects and the market-of-one test

1 Build your first real skill (do this week) · 2 Connect one app read-only · 3 Wire a skill and connector together · 4 Make it portable · 5 Audit a skill before you trust it. A skill earns its place when it encodes **your** repeated way.

A Skill has no scope dial: reading it first is the control. Read-only is the control for Connectors.

SCENARIO MCQ · Select TWO

MCQ: the front-desk setup

Read the table slowly. Some reasons are quiet clues. Do not turn the page yet.

Lakeside Dental Group (about 60 staff): front-desk setup. **Mara Qureshi** runs the front desk. She logs each repeated friction she felt, what she set up in Claude, and why. **The rest of the front-desk team will copy her setup.**

#	The friction she felt	What she set up in Claude	Her reason
1	I retype our reminder wording, tone and sign-off into every patient message.	A Skill holding the house reminder style	I keep re-explaining how
2	I copy this week's bookings from the booking app into a message.	A booking-app connector, read-only, scoped to this week's calendar	I keep fetching data from another app
3	Every Friday I need cancellations from the tracker in our fixed summary layout.	A tracker connector (read-only) plus a Skill for the layout	The job needs the data and the format
4	The practice owner wanted one farewell note for a retiring dentist. I will not write another.	A new Skill, so the note is saved	It might be useful later
5	I need last week's supplier invoices from one Drive folder.	A Drive connector with write access to the whole drive	I will never have to change it
6	A supplier emailed me a zip file for ordering.	Asked IT who published it and read its SKILL.md before installing	An unknown sender is not a trusted source

Which TWO rows depart from what the course teaches?

A Row 4, because a task done once is a good prompt, and a Skill earns its place only by repeating.

B Row 2, because a read-only connector should never be used for calendar data.

C Row 5, because a one-folder job was given write access to the whole drive instead of the smallest scope.

D Row 3, because a job needing data and a format should use a Skill alone.

Pause here. Pick your TWO answers, then turn the page.

SCENARIO ANALYSIS · STEP 1

Split the story into 5 parts

Parts 1 to 4 set the scene and hide the clues. Part 5 picks the answer.

1 COMPANY

How to spot it: Who they are and what the person does. Usually the first sentence.

In this story: Lakeside Dental Group, about 60 staff. Mara Qureshi runs the front desk.

2 PROBLEM

How to spot it: The pressure: a deadline, a client, a failure to explain.

In this story: The rest of the front-desk team will copy her setup, so any wrong habit spreads to everyone.

3 ACTIONS

How to spot it: What the person already set up, asked or approved.

In this story: Six setups: a Skill, a read-only connector, a connector plus a Skill, a new Skill, a **Drive connector with write access**, and a zip checked with IT.

4 WORKFLOW

How to spot it: What came back, and the state of the job now.

In this story: Rows 1, 2, 3 and 6 follow the course. Rows 4 and 5 give reasons the course does not accept. The team is about to copy all six.

5 HIDDEN CLUE

How to spot it: One small detail that quietly decides the answer.

In this story: The **Reason column** decides it. Row 4: **“might be useful later”** (a one-off, not a repeat). Row 5: **“I will never have to change it”** (write access, whole drive).

Eight quiet details planted in parts 1 to 4

1 Row 1 reason: “I keep re-explaining how” (Skill test passed) **2** Row 2 is read-only and scoped to this week’s calendar **3** Row 3 needs data and a format, so Connector plus Skill **4** Row 4 is one farewell note: “I will not write another” **5** Row 4 reason is “might be useful later”, not repeated friction **6** Row 5 job is only fetching invoices from one folder **7** Row 5 gets write access to the whole drive, “never have to change it” **8** Row 6 came from an unknown sender: IT asked, SKILL.md read first

Some rows look odd but follow the course. Only two rows break it.

The rule the hidden clue points to

From the course, test each reason against the rule:

- Data and a format both needed: Connector plus Skill (rows 2, 3)
- Unknown source: check it and read SKILL.md first (row 6)
- **A one-off task is a good prompt, not a Skill (row 4)**
- **Smallest scope, read-only first (row 5)**

Parts 1 to 4 set the scene. Part 5 picks the answer.

ANSWER KEY

MCQ: the front-desk setup

The two correct answers are highlighted in green.

Which TWO rows depart from what the course teaches?

A **Row 4, because a task done once is a good prompt, and a Skill earns its place only by repeating.** CORRECT

Row 4 departs: one farewell note is a one-off, so a good prompt is enough. A Skill earns its place only by repeating.

B Row 2, because a read-only connector should never be used for calendar data.

Row 2 follows the course: read-only and scoped to this week's calendar is the safe way to connect.

C **Row 5, because a one-folder job was given write access to the whole drive instead of the smallest scope.** CORRECT

Row 5 departs: fetching invoices from one folder needs read-only on that folder, not write access to the whole drive.

D Row 3, because a job needing data and a format should use a Skill alone.

Row 3 follows the course: data plus a format means Connector and Skill together, not a Skill alone.

Correct answers: A and C