

CodeVerse Soban

■ Like · Subscribe · Stay Sharp

AI Agent Factory — Level 1 Exam
Topic: The Roles This Book Trains

50 MCQs · Part A: 25 Normal · Part B: 25 Detailed · Answers + Explanations

PART A — 25 Normal Questions

Q1. What does the chapter say should replace the old education answer of 'teach them to code'?

- A) Teach data science and ML engineering
- B) Train people who specify, supervise, and verify AI Workers — judgment over syntax
- C) Focus on communication and creativity
- D) Teach cloud infrastructure management

✓ **Answer: B) Train people who specify, supervise, and verify AI Workers — judgment over syntax**

Why: Syntax can be automated; judgment, specification, and deployment move up the ladder as the machine climbs it.

Q2. In the three-level map, what does 'the Unit' refer to?

- A) The team of engineers on a project
- B) The AI Worker (Digital FTE) built by the practitioner
- C) A company's revenue unit
- D) A single SKILL.md file

✓ **Answer: B) The AI Worker (Digital FTE) built by the practitioner**

Why: The three levels are: the Person (you), the Unit (the AI Worker you build), and the Enterprise (the AI-Native Company).

Q3. What is the Outcome Architect's primary responsibility?

- A) Writing code for AI Workers
- B) Designing the cloud infrastructure
- C) Owning intent — deciding what to build and authoring the spec
- D) Running evals on deployed agents

✓ **Answer: C) Owning intent — deciding what to build and authoring the spec**

Why: The Outcome Architect owns intent, not execution. It decides what a Worker should achieve and sets what 'correct' means.

Q4. What is MCP, as mentioned in the Digital FTE Builder role description?

- A) Model Compute Protocol — the billing standard for AI API calls
- B) The standard way an agent plugs into outside tools and data
- C) Multi-Channel Pipeline — a data engineering framework
- D) Managed Cloud Platform — Anthropic's hosting service

✓ **Answer: B) The standard way an agent plugs into outside tools and data**

Why: The chapter defines MCP as 'the standard way an agent plugs into outside tools and data' — it appears in the Digital FTE Builder's training spine.

Q5. Which role is described as the book's 'primary graduate'?

- A) AI-Native Company Architect
- B) Cloud AI Engineer
- C) Digital FTE Builder
- D) Forward Deployed Engineer

✓ **Answer: C) Digital FTE Builder**

Why: The chapter explicitly states: 'This is the book's primary graduate' when describing the Digital FTE Builder.

Q6. Palantir originally called Forward Deployed Engineers by what name?

- A) Embedded Architects
- B) Solution Operators
- C) Deltas
- D) Field Engineers

✓ **Answer: C) Deltas**

Why: The chapter states: 'Palantir... created this role in the early 2010s, originally calling them Deltas.'

Q7. By approximately how much did FDE job postings grow in nine months in 2025?

- A) 150%
- B) 400%
- C) Over 800%
- D) 1,200%

✓ **Answer: C) Over 800%**

Why: The chapter cites Fast Company: 'FDE job postings grew by more than 800% in the first three quarters of 2025.'

Q8. What is the median salary range for FDE roles across the market according to the chapter?

- A) \$80,000–\$120,000
- B) Near \$190,000, ranging roughly \$160,000–\$220,000
- C) \$300,000–\$500,000
- D) \$500,000–\$1,000,000

✓ **Answer: B) Near \$190,000, ranging roughly \$160,000–\$220,000**

Why: The chapter: 'Across the market, the median sits near \$190,000, ranging roughly \$160,000 to \$220,000.' It also notes top-of-ladder figures exist but are the ceiling, not the door.

Q9. What did the MIT Media Lab Project NANDA study find about enterprise AI pilots?

- A) 95% showed measurable return on investment
- B) 95% showed no measurable return
- C) 67% failed before reaching deployment
- D) 33% were cancelled due to regulatory concerns

✓ **Answer: B) 95% showed no measurable return**

Why: The MIT study found ~95% of custom enterprise AI pilots showed no measurable return — not because AI doesn't work, but because fitting it into real-world systems is hard.

Q10. What is the Evals Engineer's role compared to the person who built the AI Worker?

- A) They are the same person — the Builder also runs evals

- B) The Evals Engineer is a senior Builder who reviews junior work
- C) The Evals Engineer designs tests for correctness and failure — crash-testing the Worker before it goes live
- D) The Evals Engineer approves governance rules written by the Governance Officer

✓ **Answer: C) The Evals Engineer designs tests for correctness and failure — crash-testing the Worker before it goes live**

Why: Evals Engineer: designs tests — does the Worker give the right answer? Does it fail gracefully? Does it stay within given lines? Like crash-testing a car before shipping.

Q11. What does the AI Governance Officer decide, and what does the AI-Native Company Architect do with those decisions?

- A) The Governance Officer builds enforcement systems; the Architect writes the rules
- B) The Governance Officer writes the rules; the Architect builds the system that enforces them
- C) Both roles jointly write and enforce governance rules
- D) The Governance Officer handles legal compliance; the Architect handles technical compliance

✓ **Answer: B) The Governance Officer writes the rules; the Architect builds the system that enforces them**

Why: The chapter is precise: 'The AI-Native Company Architect builds the system that enforces these rules; the Governance Officer decides what the rules should say.'

Q12. Kevin Bai's hiring test for an FDE requires two conditions simultaneously. What are they?

- A) Five years of experience AND a consulting background
- B) A customer-facing software engineer — someone you'd hire on engineering bar alone AND trust in front of a customer
- C) Python expertise AND product management certification
- D) Cloud deployment skills AND sales quota attainment history

✓ **Answer: B) A customer-facing software engineer — someone you'd hire on engineering bar alone AND trust in front of a customer**

Why: Bai compresses the FDE profile: 'a customer-facing software engineer — a person you would hire onto the engineering team on the engineering bar alone, and would also trust in front of a customer.' Both halves must be true.

Q13. How does Palantir's average contract value compare to other public SaaS companies serving Fortune 500, according to Bai?

- A) Palantir averages ~\$200,000 per contract, same as Workday
- B) Palantir averages ~\$4M per contract, far above ServiceNow (~\$1.2M) and Workday (~\$600K)
- C) Palantir averages ~\$10M per contract, driven by government defense contracts
- D) Palantir's contract values are comparable to other SaaS companies at ~\$500,000

✓ **Answer: B) Palantir averages ~\$4M per contract, far above ServiceNow (~\$1.2M) and Workday (~\$600K)**

Why: Bai's figures: Palantir ~\$4M average contract value, ServiceNow ~\$1.2M, Workday ~\$600K. Selling outcomes prices differently from selling seats.

Q14. Who is Pauline Brunet, and what does she run?

- A) CEO of Palantir's enterprise division
- B) VP of Forward Deployed Engineering at Cursor, with ten years of enterprise AI deployment
- C) Head of Applied AI at Anthropic, overseeing FDE hiring
- D) Founder of OpenWork, the open-source alternative to Cowork

✓ **Answer: B) VP of Forward Deployed Engineering at Cursor, with ten years of enterprise AI deployment**

Why: The chapter identifies Brunet as 'VP of Forward Deployed Engineering at Cursor after ten years of enterprise AI deployment' who laid out her playbook at the AI Engineer World's Fair in June 2026.

Q15. What fraction of the market's FDE profiles on platforms like Upwork does the chapter say are actually filled by genuine FDEs?

- A) About half — the market is evenly split between real FDEs and generalists
- B) Most — the category is well-stocked with qualified candidates
- C) Almost none — the marketplace built the shelf before anyone stocked it; listed profiles are generalists and full-stack developers
- D) About a quarter — enough to meet current demand but supply is growing fast

✓ **Answer: C) Almost none — the marketplace built the shelf before anyone stocked it; listed profiles are generalists and full-stack developers**

Why: 'Browse the profiles Upwork lists as Forward Deployed Engineers and you find capable generalists, full-stack developers, DevOps engineers... none of whom describe FDE work.' The marketplace built the shelf before the supply arrived.

Q16. Boris Cherny created which tool at Anthropic, and what were his five archetypes about?

- A) He created TutorClaw and described five curriculum stages
- B) He created Claude Code and described five archetypes observed on his team: Prototyper, Builder, Sweeper, Grower, Maintainer
- C) He created OpenCode and described five types of AI governance roles
- D) He created the Agent Skills open standard and described five deployment patterns

✓ **Answer: B) He created Claude Code and described five archetypes observed on his team: Prototyper, Builder, Sweeper, Grower, Maintainer**

Why: The chapter: 'Boris Cherny, creator of Claude Code... looked at his own team and asked what roles become when engineering, product, design, and data science melt into a new kind of role.' His five archetypes: Prototyper, Builder, Sweeper, Grower, Maintainer.

Q17. What cloud platforms does the chapter name for the Cloud AI Engineer's production deployment work?

- A) AWS Lambda, Google Cloud Run, and Heroku
- B) Azure Container Apps to ship, Inngest for durable execution, Dapr and Kubernetes to scale
- C) Docker Desktop, GitHub Actions, and Vercel
- D) Render, Railway, and Fly.io

✓ **Answer: B) Azure Container Apps to ship, Inngest for durable execution, Dapr and Kubernetes to scale**

Why: The chapter lists: 'Azure Container Apps to ship, Inngest for durable execution, Dapr and Kubernetes to scale' as the Cloud AI Engineer's production infrastructure.

Q18. What is the Digital FTE Supervisor's relationship to the audit trail when something goes wrong?

- A) The Supervisor investigates errors but is not personally named in the audit trail
- B) The Supervisor's name is on the audit trail — they are the reviewer, approver, and point of accountability
- C) The Builder who created the Worker is named in the audit trail, not the Supervisor
- D) The Governance Officer holds audit trail responsibility, not the Supervisor

✓ **Answer: B) The Supervisor's name is on the audit trail — they are the reviewer, approver, and point of accountability**

Why: The chapter: 'the name the audit trail points to when something goes wrong' is the Supervisor — 'the reviewer who checks the work, the manager who approves the output.'

Q19. Why does the LLMops Engineer role sit outside the book's scope when it comes to foundation models?

- A) Pre-training foundation models from scratch is commoditizing and not the job the book trains
- B) Foundation model training requires hardware the book's audience doesn't have access to
- C) The book deliberately avoids all machine learning topics to stay accessible
- D) Foundation models are trained by Anthropic internally and cannot be customized

✓ **Answer: A) Pre-training foundation models from scratch is commoditizing and not the job the book trains**

Why: The chapter: 'The hard stop is building the model itself: pre-training a foundation model from scratch stays out of scope, because that capability is commoditizing.' The book trains fine-tuning and ops around the model, not building foundation models.

Q20. Sanjeev Aggarwal founded which company before co-founding Fundamentum with Nandan Nilekani?

- A) Infosys
- B) Wipro
- C) Daksh — one of the pioneers of India's BPO industry
- D) TCS

✓ **Answer: C) Daksh — one of the pioneers of India's BPO industry**

Why: The chapter identifies Aggarwal as 'who founded Daksh, one of the pioneers of India's BPO industry, before co-founding the venture firm Fundamentum with Infosys's Nandan Nilekani.'

Q21. What is the 'bus factor of one' test the chapter applies to a pod of one?

- A) Whether the engineer can commute to client sites using public transport
- B) Whether another graduate could pick up the engagement from the repository alone if the engineer were unreachable for two weeks
- C) Whether the engineer can handle all client meetings without a business partner
- D) Whether the system can run for one week without any human intervention

✓ **Answer: B) Whether another graduate could pick up the engagement from the repository alone if the engineer were unreachable for two weeks**

Why: 'If you were unreachable for two weeks, could another graduate of this book pick up your engagement from your repository alone? If not, you do not have a pod of one. You have a bus factor of one.'

Q22. What are Anthropic's posted salary bands for FDE roles, according to the chapter?

- A) \$80,000–\$120,000
- B) \$120,000–\$180,000
- C) \$200,000–\$300,000
- D) \$400,000–\$600,000

✓ **Answer: C) \$200,000–\$300,000**

Why: The chapter cites: 'Anthropic's own FDE postings run \$200,000 to \$300,000.'

Q23. What does the chapter say a vendor-neutral FDE carries that a vendor's FDE cannot carry to a new client?

- A) Access to proprietary model weights trained on client data
- B) A portfolio of code repositories from previous engagements
- C) The discipline itself — the method and the profession — which live in the person, not in any vendor's platform
- D) A network of enterprise contacts from previous deployments

✓ **Answer: C) The discipline itself — the method and the profession — which live in the person, not in any vendor's platform**

Why: A vendor's FDE leaves the platform behind when they leave the vendor. A vendor-neutral FDE's discipline lives in the person — they can join the client directly and keep doing the same work, because the method travels with them.

Q24. According to the chapter, what happened when Microsoft was pressed to compare itself to Palantir?

- A) Microsoft refused to comment on the comparison
- B) Microsoft argued it supports more models, connectors, and open systems — but the chapter notes this is a vendor measuring its own lock-in against another's, calling a looser cage the feature
- C) Microsoft agreed Palantir's model was superior and announced a partnership
- D) Microsoft claimed its FDEs are more technically skilled due to a stricter hiring bar

✓ **Answer: B) Microsoft argued it supports more models, connectors, and open systems — but the chapter notes this is a vendor measuring its own lock-in against another's, calling a looser cage the feature**

Why: The chapter: 'the executive running Frontier argued that Microsoft supports more models, more connectors... Notice what that defense is: a vendor measuring its own lock-in against another vendor's and calling the looser cage the feature.'

Q25. What does Brunet say she is 'now building' that she never planned, because clients keep demanding it?

- A) A certification program for clients' internal AI engineers
- B) Help reorganizing the company itself — who to hire, what job descriptions say, how ways of working change once agents arrive
- C) A model fine-tuning service using the client's proprietary data
- D) An open-source alternative to Cursor's proprietary FDE platform

✓ **Answer: B) Help reorganizing the company itself — who to hire, what job descriptions say, how ways of working change once agents arrive**

Why: The chapter: 'she named an offering she never planned and is now building, because clients keep demanding it: help reorganizing the company itself, who to hire, what the job descriptions say, how the ways of working change once the agents arrive.'

PART B — 25 Detailed Questions

Q26. The chapter describes the 'execution vs. intent' gap created by agentic coding tools. What specific ratio does it mention, and what does the widening gap between the two lines on the graph actually measure?

- A) A 1-to-5 ratio of PMs to engineers, showing that product teams are already over-staffed relative to AI-augmented engineering output.
- B) The old ratio was roughly 1 PM per 8 engineers; now one person effectively feeds the work of twenty. The widening gap measures that execution scaled with agentic coding while intent capacity per decision-maker stayed flat — turning the 'deciding of what to build' into the bottleneck.

C) A 1-to-3 ratio of Outcome Architects to Digital FTE Builders, showing the staffing structure of an AI-Native Company.

D) A 10-to-1 ratio of AI Workers to human supervisors, measuring the automation density of a mature AI-Native Company.

✓ **Answer: B) The old ratio was roughly 1 PM per 8 engineers; now one person effectively feeds the work of twenty. The widening gap measures that execution scaled with agentic coding while intent capacity per decision-maker stayed flat — turning the 'deciding of what to build' into the bottleneck.**

Why: Old ratio ~1 PM per 8 engineers; now one person feeds the work of twenty. The graph shows two lines: execution output per engineer (steep rise) vs. intent capacity per decision-maker (flat). The widening gap = the new bottleneck.

Q27. Bai describes what happens structurally when an FDE function lacks shared primitives. What term does he use, what does it mean, and what does the chapter offer as the vendor-neutral answer to his warning?

A) He calls it a 'feature factory' — a team that ships many small features with no strategic coherence. The chapter answers with the Two-Layer Model that ensures Workers align to company strategy.

B) He calls it a 'dev shop' — each client gets entirely custom code, creating unmaintainable repositories and engineers who quit rather than learn them. The chapter's answer: two Systems of Record that function as shared primitives, so the engineer never starts from scratch.

C) He calls it a 'consulting trap' — a firm that sells advice without building anything. The chapter answers with the capstone requirement that every graduate must ship at least one deployed Worker.

D) He calls it a 'platform dependency' — over-reliance on one vendor's infrastructure. The chapter answers with the instruction to use only open-source tools across all client engagements.

✓ **Answer: B) He calls it a 'dev shop' — each client gets entirely custom code, creating unmaintainable repositories and engineers who quit rather than learn them. The chapter's answer: two Systems of Record that function as shared primitives, so the engineer never starts from scratch.**

Why: Bai: without shared primitives 'you do not have an FDE function, you have a dev shop.' Maintenance costs eat P&L; and engineers resign. The chapter's answer: the method (given, identical at every client) and the profession (built by the graduate, one vertical) both pass Bai's test — the engineer never starts from scratch.

Q28. The chapter explains that vendor-neutrality forces specialization to 'land somewhere.' Where does it land, and why can't it land on a platform instead?

A) It lands on geography — a vendor-neutral FDE specializes in a specific region's enterprise ecosystem. Platform specialization fails because regional laws change too quickly.

B) It lands on the profession — one vertical, one jurisdiction. Platform specialization fails because taking the platform away and becoming a generalist consultant means nothing travels from client two to client three. The profession is what a buyer actually pays for: which standard governs, which regulator, whose signature.

C) It lands on a programming language — a vendor-neutral FDE specializes in Python or JavaScript. Platform specialization fails because models increasingly generate code in any language.

D) It lands on company size — a vendor-neutral FDE specializes in mid-market clients. Platform specialization fails because enterprise clients always prefer vendor-locked solutions for support

reasons.

✓ **Answer: B) It lands on the profession — one vertical, one jurisdiction. Platform specialization fails because taking the platform away and becoming a generalist consultant means nothing travels from client two to client three. The profession is what a buyer actually pays for: which standard governs, which regulator, whose signature.**

Why: The chapter: 'Take the platform away and specialization has to land somewhere, or you are a generalist consultant with nothing to reuse.' The profession is what travels client to client. The method travels free (nobody pays extra). What a buyer pays for: which standard, which version, which regulator, which personal signature — all professional and jurisdictional.

Q29. What is Brunet's 'ROI in three questions' framework, and what specific client example does she use to demonstrate a company that measured cost but never measured return?

A) The three questions are: Did the pilot finish on time? Did it stay on budget? Did the client renew? Her example: a retail client who measured project timeline but ignored inventory cost reduction.

B) The three questions are: Did we increase revenue, decrease costs, or mitigate risk? Her example: a client alarmed that an agent cost \$2,000 a day, until she asked what the agent was doing — dispatching the right technician to failing equipment — and the client conceded the agent was cheap.

C) The three questions are: Did user adoption increase? Did support tickets decrease? Did NPS improve? Her example: a SaaS company that measured feature adoption but ignored churn reduction.

D) The three questions are: Was the data pipeline stable? Were model outputs accurate? Was compliance maintained? Her example: a financial firm that measured model accuracy but ignored regulatory penalty avoidance.

✓ **Answer: B) The three questions are: Did we increase revenue, decrease costs, or mitigate risk? Her example: a client alarmed that an agent cost \$2,000 a day, until she asked what the agent was doing — dispatching the right technician to failing equipment — and the client conceded the agent was cheap.**

Why: Brunet's three questions: increase revenue, decrease costs, mitigate risk? Example: client alarmed by \$2,000/day agent. Brunet asked what it was doing. Answer: dispatching the right technician to failing equipment. Client: 'the agent was cheap.' They measured cost, never return.

Q30. How does the chapter describe the 'second System of Record thickening' over time, and what rule does Bai state about what should be generalized vs. kept client-specific?

A) The second System of Record thickens by adding new programming frameworks with each engagement. Bai says all bespoke client code should eventually be open-sourced as shared libraries.

B) It thickens as what generalizes across clients goes into the vertical System of Record (a standard that governed three clients, a procedure signed everywhere, an invariant across a jurisdiction) while what is truly bespoke stays client-specific. Bai's rule: anything generalizable should be generalized over time — for a vendor, that goes into the product; for the vendor-neutral FDE, it goes into the vertical SoR.

C) The second System of Record thickens by accumulating case studies and client testimonials that are repurposed as marketing material for future engagements.

D) It thickens by expanding into new industries with each engagement. Bai's rule: any code that works for one industry should be tested in a second before being declared generalizable.

✓ **Answer: B)** It thickens as what generalizes across clients goes into the vertical System of Record (a standard that governed three clients, a procedure signed everywhere, an invariant across a jurisdiction) while what is truly bespoke stays client-specific. Bai's rule: anything generalizable should be generalized over time — for a vendor, that goes into the product; for the vendor-neutral FDE, it goes into the vertical SoR.

Why: Bai's rule: anything bespoke stays per-client; anything generalizable gets generalized. For a vendor, that goes into the platform. For a vendor-neutral FDE, it goes into the vertical SoR. This is 'what thickens engagement by engagement' means mechanically — the second client costs less than the first.

Q31. The decomposition case study is described as 'the filter' in FDE interviews. What exact sequence does the chapter say scoring well requires, and why does the chapter call it 'spec-driven development performed verbally'?

A) Scoring requires: propose three alternative architectures, select the best one, estimate the timeline, and list the required APIs. It's called spec-driven because you are writing an API specification during the interview.

B) Scoring requires: clarify the actual goal, name stakeholders and success metric, map existing data and ownership, decompose by risk, propose the thinnest end-to-end skeleton first, state assumptions out loud, surface failure modes unprompted, and narrate thinking continuously. It's called spec-driven because writing intent a Worker can be held to and scoping a vague enterprise problem are the same discipline at two volumes.

C) Scoring requires: identify the single biggest technical risk, propose a mitigation plan, estimate the compute cost, and name the first hire needed. It's called spec-driven because the interviewer is checking whether you can write a technical specification.

D) Scoring requires: ask three clarifying questions, then immediately begin coding a prototype, then explain the code to the interviewer. It's called spec-driven because you write the code specification before writing the implementation.

✓ **Answer: B)** Scoring requires: clarify the actual goal, name stakeholders and success metric, map existing data and ownership, decompose by risk, propose the thinnest end-to-end skeleton first, state assumptions out loud, surface failure modes unprompted, and narrate thinking continuously. It's called spec-driven because writing intent a Worker can be held to and scoping a vague enterprise problem are the same discipline at two volumes.

Why: The scoring sequence: clarify goal → name stakeholders/metric → map data → decompose by risk → thinnest skeleton first → assumptions out loud → failure modes unprompted → narrate continuously. 'Writing intent a Worker can be held to and scoping a vague enterprise problem out loud are the same skill at two different volumes' — that's why it's spec-driven development performed verbally.

Q32. What two things does the chapter say a vendor-neutral FDE can offer a client that a vendor's FDE cannot, and what asymmetry survives when the FDE is hired directly in-house?

A) The vendor-neutral FDE offers lower cost and faster deployment. The asymmetry: the FDE's personal domain expertise doesn't transfer to the client's internal team when they become an employee.

B) The vendor-neutral FDE offers freedom to switch platforms as better models appear and the ability to be hired directly. The asymmetry: the first SoR (the method) travels freely with the engineer, but the second SoR (the domain startup, held under expert licence and third-party terms) is a conversation about that business, not an automatic transfer.

C) The vendor-neutral FDE offers open-source tools and community support. The asymmetry: the client cannot access the FDE's professional network after they become a full-time employee.

D) The vendor-neutral FDE offers multi-cloud deployment and vendor diversity. The asymmetry: the client cannot renegotiate the FDE's salary once they're hired directly, since it's set by market rates.

✓ **Answer: B) The vendor-neutral FDE offers freedom to switch platforms as better models appear and the ability to be hired directly. The asymmetry: the first SoR (the method) travels freely with the engineer, but the second SoR (the domain startup, held under expert licence and third-party terms) is a conversation about that business, not an automatic transfer.**

Why: Two things a vendor-neutral FDE offers: (1) optionality — switch platforms as better models appear; (2) can be hired for keeps. The asymmetry: SoR 1 (the method) travels freely because it belongs to the ecosystem. SoR 2 (the domain startup) is held by the graduate and her expert under licence — a direct hire is a conversation about that business, not an automatic transfer.

Q33. What does the chapter say about the Connector and Plugin Engineer's work being 'one job at two addresses'? What are the two addresses and what is the single through-line?

A) Address 1: building connectors for enterprise software like Salesforce; Address 2: building connectors for consumer apps like Spotify. Through-line: OAuth authentication used in both.

B) Address 1: a connector-native app that extends the chat app (claude.ai) for end users, with remote MCP server, tools, stored state, real sign-in, fail-closed gate; Address 2: a plugin that extends the coding agent (Claude Code/OpenCode) for builders, with skills, subagents, hooks, MCP servers. Through-line: both ship a unit a host loads, and you own the extension while the host owns the loop.

C) Address 1: building for web browsers; Address 2: building for mobile apps. Through-line: the same REST APIs power both surfaces with minimal adaptation.

D) Address 1: building for Anthropic's products; Address 2: building for OpenAI's products. Through-line: the Agent Skills open standard works across both vendors.

✓ **Answer: B) Address 1: a connector-native app that extends the chat app (claude.ai) for end users, with remote MCP server, tools, stored state, real sign-in, fail-closed gate; Address 2: a plugin that extends the coding agent (Claude Code/OpenCode) for builders, with skills, subagents, hooks, MCP servers. Through-line: both ship a unit a host loads, and you own the extension while the host owns the loop.**

Why: Two addresses: connector-native app extends claude.ai for end users (remote MCP server, tools, stored state, real sign-in, fail-closed session gate); plugin extends Claude Code/OpenCode for builders (skills, subagents, hooks, MCP servers behind one install). Through-line: 'you ship a unit a host loads, and you own the extension while the host owns the loop.'

Q34. Why does the chapter say Bai's structural hypothesis is the best explanation for WHY the FDE role exploded in 2026, rather than the industry 'finally noticing' what Palantir had been doing for a decade?

A) Because the MIT study was published and immediately read by every major tech company, triggering a coordinated hiring response across the industry.

B) Because nearly every platform became agentic, meaning customizable, pushing nearly every vendor into the one cell of Bai's 2x2 that requires forward deployment — the cell where a technically complex product is sold to a buyer who cannot implement it. The agentic turn created the structural demand that made 2026 different from 2016.

C) Because a regulatory change required on-site technical staff for all enterprise AI deployments, creating mandatory demand that didn't exist in Palantir's early days.

D) Because the emergence of open-source AI models made it economically viable for small companies to afford FDEs, expanding the market beyond large enterprises.

✓ **Answer: B) Because nearly every platform became agentic, meaning customizable, pushing nearly every vendor into the one cell of Bai's 2x2 that requires forward deployment — the cell where a technically complex product is sold to a buyer who cannot implement it. The agentic turn created the structural demand that made 2026 different from 2016.**

Why: Bai's hypothesis: the industry didn't 'finally notice.' What changed structurally: nearly every platform became agentic = customizable = the customer no longer knows what the product does. This pushed nearly every vendor into the one cell requiring FDEs. The timing is explained by the agentic shift, not by discovery.

Q35. The chapter describes the live-build interview format that is 'spreading' in FDE hiring. What are its three phases, and what does the chapter say notably does NOT appear at any stage?

A) Phase 1: whiteboard architecture (30 min); Phase 2: code review of existing codebase (90 min); Phase 3: executive presentation (60 min). No LeetCode, but data structures knowledge is still tested verbally.

B) Phase 1: turning a vague use case into requirements under cross-examination (30 min); Phase 2: building the working solution with an AI coding assistant while validating every suggestion and debugging live (90 min); Phase 3: presenting the result to a role-played stakeholder in pure business language (60 min). No LeetCode appeared at any stage.

C) Phase 1: take-home problem set (48 hours); Phase 2: code walkthrough with engineering team (60 min); Phase 3: customer simulation with acting coaches (30 min). No system design questions appear.

D) Phase 1: product requirements writing (45 min); Phase 2: SQL and data analysis (45 min); Phase 3: business case presentation (30 min). No coding assessment appears.

✓ **Answer: B) Phase 1: turning a vague use case into requirements under cross-examination (30 min); Phase 2: building the working solution with an AI coding assistant while validating every suggestion and debugging live (90 min); Phase 3: presenting the result to a role-played stakeholder in pure business language (60 min). No LeetCode appeared at any stage.**

Why: The live-build loop: 30 min (vague use case → requirements under cross-examination) + 90 min (building with AI coding assistant, validating every suggestion, debugging live) + 60 min (presenting to role-played stakeholder in business language). 'No LeetCode appeared at any stage.'

Q36. The chapter says the SME as Skill Author and the vendor-neutral FDE 'need each other.' How does the chapter describe what each brings, and what does it call this pairing?

A) The SME brings technical credentials; the FDE brings client relationships. Together they form a 'delivery partnership' that the chapter calls a vertical integration.

B) The SME's second SoR can't exist without the expert's voice and real files; the Skill Author needs the FDE to build the governed home her judgment lives in. Neither is junior: the expert brings 20 years and the licence; the engineer brings the method and the build. The chapter calls this pairing a 'vertical' in the FDE AF Model.

C) The SME brings domain data for model training; the FDE brings the infrastructure to host the fine-tuned model. Together they form a 'data-to-deployment pipeline' in the FDE AF Model.

D) The SME brings regulatory expertise; the FDE brings engineering speed. Together they form a 'compliance-first development team' that the chapter calls a governance vertical.

✓ **Answer: B) The SME's second SoR can't exist without the expert's voice and real files; the Skill Author needs the FDE to build the governed home her judgment lives in. Neither is junior: the expert brings 20 years and the licence; the engineer brings the method and the build. The chapter calls this pairing a 'vertical' in the FDE AF Model.**

Why: The chapter: 'the expert brings twenty years and the licence. The engineer brings the method and the build. That pairing is the unit the FDE AF Model calls a vertical.' Neither is junior. And Choosing Your Vertical 'refuses to launch one without a committed expert.'

Q37. How does the chapter describe Aggarwal's compression arithmetic, and what does it say his conclusion generalizes beyond just India?

A) 500 FDEs replace 5,000 traditional IT workers at 50% margins. His conclusion: Southeast Asia will become the global FDE hub due to lower labor costs.

B) ~100 FDEs replace 2,000–2,500 traditional IT staff for a \$100M business at 70–90% gross margins — ~25x compression. His conclusion: 'India can be the FDE factory for the world,' but the chapter generalizes: the IT-services decades taught embedding in messy systems; that muscle lives in Karachi and Lahore as much as Bangalore. What converts it is training, available to anyone who does the work.

C) 10 FDEs replace 100 traditional IT staff for a \$10M business at 40–50% margins. His conclusion: the FDE model only works at large scale and is unsuitable for SMB-focused service firms.

D) 50 FDEs replace 1,000 traditional IT staff for a \$50M business at 60% margins. His conclusion: the FDE model requires significant upfront investment before it becomes profitable.

✓ **Answer: B) ~100 FDEs replace 2,000–2,500 traditional IT staff for a \$100M business at 70–90% gross margins — ~25x compression. His conclusion: 'India can be the FDE factory for the world,' but the chapter generalizes: the IT-services decades taught embedding in messy systems; that muscle lives in Karachi and Lahore as much as Bangalore. What converts it is training, available to anyone who does the work.**

Why: Aggarwal: ~100 FDEs, \$100M business, 70–90% gross margins vs. 2,000–2,500 in traditional IT. ~25x compression. His claim: 'India can be the FDE factory for the world.' The chapter generalizes: 'the IT-services decades taught South Asia how to embed inside a client's messy systems — that muscle lives in Karachi and Lahore as much as Bangalore. What converts it is training.'

Q38. What does Brunet's 'staff augmentation red flag' look like in practice, and what is her counter-move that tests whether an engagement is genuine FDE work or body-shopping in disguise?

A) Red flag: the client asks for a fixed-price contract. Counter-move: insisting on time-and-materials with weekly billing to maintain flexibility.

B) Red flag: the client says 'we're understaffed' — requesting to rent hours rather than transfer capability. Counter-move: asking 'who will be the working team?' — if they cannot name the people who will build alongside you, the engagement is body-shopping in disguise.

C) Red flag: the client wants NDA-protected deliverables. Counter-move: proposing that all IP created during the engagement remains with the client to avoid misaligned incentives.

D) Red flag: the client asks for more than two FDEs simultaneously. Counter-move: capping all engagements at one FDE to maintain quality and prevent the function from reverting to a dev shop.

✓ **Answer: B) Red flag: the client says 'we're understaffed' — requesting to rent hours rather than transfer capability. Counter-move: asking 'who will be the working team?' — if they cannot name the people who will build alongside you, the engagement is body-shopping in disguise.**

Why: Brunet's exact rule: 'we're understaffed' = request to rent hours, not transfer capability, she declines. Counter-move: 'who will be the working team?' If they can't name those people, 'the engagement is body-shopping in disguise.' This is the field test for the FDE vs. services pyramid distinction.

Q39. The chapter describes the 'first order of action' for someone who wants to enter the freelance FDE market. What is the first step, and why does the chapter say the empty marketplace shelf is an opportunity rather than a warning?

- A) The first step is to create a detailed profile on Upwork listing all technical skills. The empty shelf is an opportunity because the platform algorithm favors new profiles with no competition.
- B) The first step is to build first, sell second — produce deployed Workers, shipped plugins, live connector apps, and one governed slice of a profession as the portfolio/credential. The shelf being empty before it's stocked means the demand side is posting projects at published rates into a category supply hasn't learned to fill — a trained reader has the opening.
- C) The first step is to get hired by a vendor's FDE team for two years to build a track record before going independent. The empty shelf means the market is not yet mature enough for independent FDEs.
- D) The first step is to pass a freelance FDE certification exam offered by Anthropic or Palantir. The empty shelf is an opportunity because certification immediately differentiates you from uncertified generalists.

✓ **Answer: B) The first step is to build first, sell second — produce deployed Workers, shipped plugins, live connector apps, and one governed slice of a profession as the portfolio/credential. The shelf being empty before it's stocked means the demand side is posting projects at published rates into a category supply hasn't learned to fill — a trained reader has the opening.**

Why: 'Build first, sell second.' The chapter: 'the portfolio is the credential.' The empty shelf is the opening: 'the demand side is posting projects and paying published rates into a category the supply side has not yet learned to fill.' For most job titles an empty shelf is a warning; for a trained reader it is the opening.

Q40. Cherny says most people span two or three archetypes. How does the chapter connect this observation to the concept of a 'pod of one,' and what does it say about which supporting disciplines must be borrowed vs. genuinely held?

- A) Spanning multiple archetypes means one person can replace an entire team, so a pod of one is natural for anyone with broad skills. Supporting disciplines like evals and governance can be handled entirely within the pod.
- B) The spanning is the pod of one seen from the inside: Workers absorb the tasks, while the human's two or three archetypes decide which seats they can genuinely hold — and which supporting disciplines they must borrow rather than fake: evals for the Sweeper's subtraction, governance for the Maintainer's caution.
- C) Spanning archetypes means an FDE will naturally outgrow the pod of one and need to hire additional specialists within 12 months of starting client engagements.
- D) The five archetypes map to five different pod sizes: Prototypers work alone, Builders work in pairs, and Maintainers require full teams of 5–6 engineers.

✓ **Answer: B) The spanning is the pod of one seen from the inside: Workers absorb the tasks, while the human's two or three archetypes decide which seats they can genuinely hold — and which supporting disciplines they must borrow rather than fake: evals for the Sweeper's subtraction, governance for the Maintainer's caution.**

Why: The chapter: 'the spanning is the pod of one seen from the inside: the Workers absorb the tasks, while the human's two or three archetypes decide which seats they can genuinely hold, and which supporting disciplines they must borrow rather than fake — evals for the Sweeper's subtraction, governance for the Maintainer's caution.'

Q41. The chapter explains why the FDE title 'was never a description of the engineer or their skills.' What was it always describing, and what happens to the title when a client

hires the FDE directly?

- A) The title described the engineer's seniority level. When hired directly, the engineer becomes a Principal Engineer or Staff Engineer depending on experience.
- B) The title described where they work — embedded inside a client's company as an outsider, carrying the whole pipeline end-to-end. When hired directly, the client's company becomes their own company; they are no longer 'forward deployed,' just deployed — the title falls away and they fold back into the core as AI-Native Company Architect.
- C) The title described the engineer's specialization in deployment infrastructure. When hired directly, they become a Platform Engineer focused on the client's internal developer tooling.
- D) The title described the engineer's consulting background. When hired directly, they transition to a product management role since engineering skills are less needed once embedded full-time.

✓ **Answer: B) The title described where they work — embedded inside a client's company as an outsider, carrying the whole pipeline end-to-end. When hired directly, the client's company becomes their own company; they are no longer 'forward deployed,' just deployed — the title falls away and they fold back into the core as AI-Native Company Architect.**

Why: The chapter: "FDE" was never a description of the engineer or their skills. It describes where they work: embedded inside a client's company, as an outsider, carrying the whole line end to end.' When hired directly: 'the client's company becomes their own company... they are no longer forward deployed: just deployed. The title falls away.'

Q42. The chapter describes the 'Harness Engineer' as a deliberate stop in the book. What does the book train vs. what is explicitly out of scope for this role?

- A) The book trains building agent runtimes from scratch; out of scope is using commercial harnesses like OpenAI Agents SDK.
- B) The book trains the operator who uses any harness (OpenAI Agents SDK, Claude's managed agents) fluently and portably across runtimes; out of scope is building the harness/agent runtime itself.
- C) The book trains fine-tuning harnesses for specific domains; out of scope is general-purpose agent orchestration.
- D) The book trains deploying harnesses on cloud infrastructure; out of scope is local development and testing of agent runtimes.

✓ **Answer: B) The book trains the operator who uses any harness (OpenAI Agents SDK, Claude's managed agents) fluently and portably across runtimes; out of scope is building the harness/agent runtime itself.**

Why: 'Trains the operator who uses any runtime, not the engineer who builds one.' The harness is the agent runtime — OpenAI Agents SDK, Claude's managed agents, etc. The book trains fluent, portable use across runtimes. Building the runtime itself is not the job.

Q43. How does the chapter describe 'how the second System of Record thickens engagement by engagement,' and what does it say the second client costs compared to the first?

- A) The second SoR thickens by training junior engineers on each engagement, who then join the firm. The second client costs more due to onboarding new team members.
- B) What generalizes across clients (a standard governing three clients, a procedure signed everywhere, an invariant across a jurisdiction) goes into the vertical SoR, not a vendor's platform. The second client costs less to serve than the first because the generalizable work is already done.

C) The second SoR thickens by adding new programming skills with each engagement. The second client costs less because the engineer writes code faster with practice.

D) The second SoR thickens by expanding the number of industries covered. The second client costs less only if they are in the same industry as the first.

✓ **Answer: B) What generalizes across clients (a standard governing three clients, a procedure signed everywhere, an invariant across a jurisdiction) goes into the vertical SoR, not a vendor's platform. The second client costs less to serve than the first because the generalizable work is already done.**

Why: The chapter: 'What generalizes does not go into a vendor's platform. It goes into the vertical System of Record: the standard that turned out to govern three clients rather than one, the procedure the expert wrote once and now signs off everywhere.' 'Why the second client costs less to serve than the first' is the mechanical payoff of this thickening.

Q44. The chapter describes five FDE résumé signals that screeners specifically look for. What are the first three, and why does the chapter say résumés that bury these signals get rejected before the technical bar is even tested?

A) The first three: programming languages listed, GitHub star count, and years of experience. These get buried when candidates use generic job titles instead of specific FDE terminology.

B) The first three: shipped production systems (real deployments, not features on a team backlog), quantifiable impact (what the customer gained in numbers, not 'built feature X'), and direct customer exposure (you sat with the stakeholder, not behind a PM). These get buried when candidates write activity-based bullets instead of outcome-based ones, making screeners unable to tell if the candidate actually delivered.

C) The first three: ML/AI certifications, data engineering experience, and cloud platform expertise. These get buried when candidates list soft skills prominently instead of technical qualifications.

D) The first three: consulting firm alumni status, MBA or equivalent business education, and enterprise client name-dropping. These get buried when candidates focus too heavily on technical details at the expense of business context.

✓ **Answer: B) The first three: shipped production systems (real deployments, not features on a team backlog), quantifiable impact (what the customer gained in numbers, not 'built feature X'), and direct customer exposure (you sat with the stakeholder, not behind a PM). These get buried when candidates write activity-based bullets instead of outcome-based ones, making screeners unable to tell if the candidate actually delivered.**

Why: The chapter's first three signals: shipped production systems, quantifiable impact (numbers, not activity), direct customer exposure. The chapter: 'a profile that buries them gets rejected before the technical bar is ever tested.' Fix: reframe every bullet from activity to outcome; write 'I' not 'we'; cut algorithm trophies.

Q45. What is the 'one honest label' the chapter places on its own claims about vendor-neutral FDEs converting governed corpora into first clients?

A) The chapter labels these claims as fully verified — multiple pilot programs have been run and the conversion rate is documented at 73%.

B) The chapter says: 'no verified count yet exists of vendor-neutral graduates who converted a governed corpus into a first client, because the category is new and the shelf below is still empty. Read the order as this book's reasoning, in the same class as Aggarwal's arithmetic: sound, and not yet a measurement.'

C) The chapter labels these as projections based on the MIT study's 67% deployment rate for outside-partner-led initiatives.

D) The chapter says the claims are verified by the Certified Agentic AI Architect program's track record, which has placed 200+ graduates in vendor-neutral FDE roles since 2024.

✓ **Answer: B) The chapter says: 'no verified count yet exists of vendor-neutral graduates who converted a governed corpus into a first client, because the category is new and the shelf below is still empty. Read the order as this book's reasoning, in the same class as Aggarwal's arithmetic: sound, and not yet a measurement.'**

Why: The chapter is explicit: 'no verified count yet exists of vendor-neutral graduates who converted a governed corpus into a first client, because the category is new and the shelf below is still empty. Read the order as this book's reasoning, in the same class as Aggarwal's arithmetic: sound, and not yet a measurement.'

Q46. How does the chapter map Cherny's Prototyper archetype to the Outcome Architect role, and how does it explain the Grower archetype being 'honestly unmapped'?

A) Prototyper maps to Outcome Architect because both focus on creating new ideas without worrying about execution quality. Grower is unmapped because no role in the book focuses purely on growth metrics.

B) Prototyper maps to Outcome Architect because both generate new ideas and most never ship — the Outcome Architect churns out specs, most Workers never get built. Grower is unmapped because the chapter cannot cleanly assign iterating a built product toward PMF to any single role in its framework.

C) Prototyper maps to Outcome Architect because both work in the pre-specification phase before any code is written. Grower is unmapped because growth functions are handled externally by the Certified Agentic AI Business Strategist track.

D) Prototyper maps to Outcome Architect because both roles are temporary — they hand off to Builders once the intent is defined. Grower is unmapped because the book focuses on building AI Workers, not growing existing products.

✓ **Answer: B) Prototyper maps to Outcome Architect because both generate new ideas and most never ship — the Outcome Architect churns out specs, most Workers never get built. Grower is unmapped because the chapter cannot cleanly assign iterating a built product toward PMF to any single role in its framework.**

Why: The chapter: 'Prototyper → Outcome Architect, Builder → Digital FTE Builder, Sweeper → Evals Engineer, Maintainer → Cloud AI Engineer and Supervisor — with Grower left honestly unmapped.' The mapping is described as 'loud without being one-to-one.' The Grower's work (iterating toward PMF) doesn't cleanly map to a single seat in the framework.

Q47. The chapter cites the five client simulation red flags that FDE interviewers consistently name across sources. What are they?

A) Moving too slowly, using too much jargon, proposing expensive solutions, ignoring the timeline, and asking too many clarifying questions.

B) Solving before clarifying, ignoring cost and constraints, thin deployment stories (no account of what happened when an API failed in production), zero compliance vocabulary in regulated domains, and no customer instinct.

C) Overconfidence in AI accuracy, underestimating integration complexity, skipping the governance discussion, proposing open-source tools when vendors are required, and poor stakeholder communication.

D) Failing to mention specific AI frameworks, using non-technical language with a technical interviewer, proposing a waterfall delivery approach, underestimating team size, and ignoring data privacy considerations.

✓ **Answer: B) Solving before clarifying, ignoring cost and constraints, thin deployment stories (no account of what happened when an API failed in production), zero compliance vocabulary in regulated domains, and no customer instinct.**

Why: The chapter names five red flags in the client simulation: (1) solving before clarifying, (2) ignoring cost and constraints, (3) thin deployment stories — no account of what happened when an API failed in production, (4) zero compliance vocabulary in regulated domains, and (5) no customer instinct.

Q48. What does the chapter say distinguishes how Palantir, OpenAI, and Anthropic each emphasize different things in their FDE interview processes?

A) Palantir focuses on sales skills and enterprise relationship management; OpenAI focuses on research publications and academic credentials; Anthropic focuses on Python coding speed and LeetCode performance.

B) Palantir leans on data engineering, ontology thinking, and the decomposition round it invented; OpenAI leans on building and evaluating systems against its APIs with "how do you know it's actually working?" as the differentiator; Anthropic (which titles the role Applied AI Engineer) leans on production LLM systems, evals, and mission alignment.

C) All three companies use identical interview processes because FDE skills are standardized across the industry — the only difference is the domain of the case study used.

D) Palantir focuses on defense and government use cases; OpenAI focuses on consumer product experience; Anthropic focuses on safety research background and alignment philosophy.

✓ **Answer: B) Palantir leans on data engineering, ontology thinking, and the decomposition round it invented; OpenAI leans on building and evaluating systems against its APIs with "how do you know it's actually working?" as the differentiator; Anthropic (which titles the role Applied AI Engineer) leans on production LLM systems, evals, and mission alignment.**

Why: The chapter: "Palantir leans on data engineering, ontology thinking, and the decomposition round it invented; OpenAI leans on building and evaluating systems against its APIs, with 'how do you know it's actually working?' as the differentiator; Anthropic (which titles the role Applied AI Engineer) leans on production LLM systems, evals, and mission alignment."

Q49. Why does the chapter say the FDE interview's hardest rounds are not 'extra material bolted onto the curriculum' but rather the curriculum itself examined at a different volume?

A) Because the book's curriculum is specifically designed to prepare students for FDE interviews, so there is no difference between learning the material and preparing for the interview.

B) Because the decomposition round tests spec-driven development performed verbally, and the live-build tests Mode 1 discipline under observation — a candidate who has actually manufactured a Worker under spec discipline has rehearsed the decomposition round dozens of times, because writing intent a Worker can be held to and scoping a vague enterprise problem are the same skill at two volumes.

C) Because Anthropic designed the FDE interview rounds in collaboration with the Agent Factory curriculum authors, ensuring perfect alignment between what is taught and what is tested.

D) Because all engineering interviews test the same fundamental skills, and the FDE interview simply applies those skills in a business context rather than a technical one.

✓ **Answer: B) Because the decomposition round tests spec-driven development performed verbally, and the live-build tests Mode 1 discipline under observation — a candidate who has actually manufactured a Worker under spec discipline has rehearsed the decomposition round dozens of times, because writing intent a Worker can be held to and scoping a vague enterprise problem are the same skill at two volumes.**

Why: The chapter: 'the interview's hardest rounds — decomposition, the live build, the eval question — are not extra material bolted onto the curriculum. They are the curriculum, examined.' A candidate who has manufactured Workers under spec discipline walks into the decomposition round having rehearsed it dozens of times — the same discipline at two different volumes.

Q50. Yousuf Imran's story is used to illustrate the Subject Matter Expert's wager. What does the chapter say his decision reveals about the relationship between domain expertise and the manufacturing metaphor in the AI era?

- A) His story shows that high-earning professionals should always keep their salary and invest in AI startups on the side, rather than leaving stable employment for entrepreneurship.
- B) His 20 years of learning salespeople's problems is the asset; his bet is that this judgment encoded into AI products he owns is worth more than renting it out for a near-million-dollar salary. The mechanism: 'domain expertise became a manufacturing input, and the expert kept the factory' — the Skill Author's wager at the most visible price the market has yet published.
- C) His story is primarily about the risks of leaving enterprise employment, showing that even highly paid professionals fail at AI entrepreneurship more often than they succeed.
- D) His decision reveals that sales professionals with large networks have an unfair advantage in building AI products because they can immediately sell to their existing client base.

✓ Answer: B) His 20 years of learning salespeople's problems is the asset; his bet is that this judgment encoded into AI products he owns is worth more than renting it out for a near-million-dollar salary. The mechanism: 'domain expertise became a manufacturing input, and the expert kept the factory' — the Skill Author's wager at the most visible price the market has yet published.

Why: The chapter reads past the headline: Imran's asset was 20 years of learning salespeople's problems. His bet: 'this judgment, encoded into AI products he owns, is worth more than a near-million-dollar salary renting it out.' The chapter's mechanism: 'domain expertise became a manufacturing input, and the expert kept the factory.' That is the Skill Author's wager.

CodeVerse Soban

■ Like · Subscribe · Stay Sharp

End of Exam — Best of Luck!

AI Agent Factory | The Roles This Book Trains | 50 MCQs