Agent Factory Complete MCQ Bank

200 Comprehensive Multiple Choice Questions

200 Comprehensive Multiple Choice Questions

Based on:

- Agent Factory Thesis (agentfactory.panaversity.org)
- Agentic AI Orientation Guide (79-page PDF)
- Agentic AI Presentation (79-slide PPTX)

Document Structure:

- **Beginner Level**: 50 MCQs (Questions 1-50)
- **Intermediate Level**: 50 MCQs (Questions 51-100)
- **Expert Level**: 50 MCQs (Questions 101-150)
- **Ultra Level**: 50 MCQs (Questions 151-200)

Format: Each question includes 4 options, correct answer marked with Correct, and 2-line explanation with reference.

BEGINNER LEVEL (Questions 1-50)

Question 1

What is the fundamental difference between regular AI and Agentic AI?

A. Regular AI is faster, Agentic AI is more accurate

B. Regular AI requires internet connection, Agentic AI works offline

C. Regular AI tells you how to do something, Agentic AI actually does the task for you from start to finish

D. Regular AI uses neural networks, Agentic AI uses rule-based systems

****Answer: C Correct****

Explanation: Agentic AI plans, executes, and completes entire jobs autonomously, while regular AI provides answers or instructions that you must act on yourself. The analogy in the sources compares asking "how to make a sandwich" (regular AI tells you steps) versus "make me a sandwich" (agentic AI does the complete job).

***Reference:** Agentic-AI.pdf pages 8-10, Agentic-AI.pptx slides 8-10*

Question 2

According to the Agent Factory thesis, what percentage of work does AI handle in the 10-80-10 rhythm?

A. 10%

B. 80%

C. 50%

D. 100%

****Answer: B Correct****

Explanation: The 10-80-10 pattern allocates 10% to human intent-setting, 80% to AI execution, and 10% to human verification and approval. This rhythm ensures humans maintain control over direction and quality while AI handles the bulk of execution work.

***Reference:** Agentic-AI.pdf pages 49-52, Agentic-AI.pptx slides 49-52*

Question 3

What does "Digital FTE" stand for in the Agent Factory context?

A. Digital File Transfer Engine

B. Digital Full-Time Equivalent

C. Digital Fast Task Executor

D. Digital Framework for Technology Execution

Answer: B Correct

Explanation: Digital FTE (Full-Time Equivalent) refers to an AI agent that performs a complete job like a human employee would, working autonomously within defined guardrails. It's both a technical and economic construct, replacing traditional headcount with AI-powered capacity.

Reference: Agent Factory website core concepts, Agentic-AI.pdf page 76

Question 4

Which analogy does the orientation guide use to explain how AI learns and makes decisions?

A. A smart kitchen appliance (microwave)

B. A self-driving car

C. A chess-playing computer

D. A weather prediction system

Answer: A Correct

Explanation: The guide uses a smart microwave analogy: while a regular microwave uses fixed time and power, a smart one can look at the food, recognize what it is, and automatically choose the best settings. This illustrates AI's pattern recognition and decision-making abilities.

Reference: Agentic-AI.pptx slides 4-5

Question 5

What are the four abilities that turn a chatbot into an agentic worker?

A. Speed, accuracy, memory, and learning

B. Planning, tool use, self-correction, and goal completion

C. Reading, writing, calculating, and communicating

D. Searching, filtering, sorting, and analyzing

Answer: B Correct

Explanation: The four defining abilities of agentic AI are: Planning (breaking goals into steps), Tool Use (accessing external capabilities), Self-Correction (fixing mistakes independently), and Goal Completion (working until the job is done, not just one answer).

Reference: Agentic-AI.pdf page 13, Agentic-AI.pptx slide 13

Question 6

In the Agent Factory model, who is always in charge?

A. The AI delegate

B. The AI workforce

C. The human

D. The system of record

Answer: C Correct

Explanation: Rule #1 of the 7 Golden Rules states "Human is always in charge" - every action traces back to a human and AI never sets its own goals. This ensures accountability and maintains human oversight over all AI operations.

Reference: Agentic-AI.pdf page 69, Agentic-AI.pptx slide 69

Question 7

What is Mode 1 in the Agent Factory framework?

A. Building AI workers for long-term use

B. Using AI to solve immediate problems right now

C. Training AI models from scratch

D. Deploying AI to production environments

Answer: B Correct

Explanation: Mode 1 (Problem-Solving) means using AI to help you complete tasks faster in the current session - nothing permanent is built. Example: "Help me analyze this spreadsheet." The session ends when the problem is solved.

Reference: Agentic-AI.pdf pages 47, Agentic-AI.pptx slides 46-47

Question 8

What is Mode 2 in the Agent Factory framework?

A. Using AI for personal tasks only

B. Upgrading from Mode 1 features

C. Building AI workers that continue working after you close your laptop

D. Running AI in offline mode

Answer: C Correct

Explanation: Mode 2 (Manufacturing) involves using general agents to build autonomous AI workers (Digital FTEs) that run continuously and independently. Example: "Build me a customer support bot that runs 24/7." The worker persists beyond the building session.

Reference: Agentic-AI.pdf pages 47, Agentic-AI.pptx slides 46-47

Question 9

According to the thesis, which tool is described as "frontier-first" with the deepest ecosystem integration?

A. OpenCode

B. ChatGPT

C. Claude Code

D. GitHub Copilot

Answer: C Correct

Explanation: Claude Code is positioned as Anthropic's frontier coding agent offering the deepest ecosystem integration and frontier model performance. It's the canonical tool for the Agent Factory methodology alongside OpenCode (the open-source alternative).

Reference: Agent Factory website tools section

Question 10

What is the primary difference between Claude Code and OpenCode?

A. Claude Code is free, OpenCode is paid

B. Claude Code is for experts, OpenCode is for beginners

C. Claude Code uses Anthropic models with deep integration, OpenCode is model-agnostic and open-source

D. Claude Code works offline, OpenCode requires internet

Answer: C Correct

Explanation: Claude Code is Anthropic's proprietary agent with frontier capabilities, while OpenCode is the open-source, model-agnostic alternative supporting dozens of providers (Claude, GPT, Gemini, DeepSeek, local models). Both implement identical SKILL.md patterns and specifications.

Reference: Agent Factory website two-tool philosophy section

Question 11

What is a "System of Record" in the AI-Native Company architecture?

A. A backup system for AI decisions

B. The company's official memory and single source of truth that AI workers read from and write to

C. A recording of all AI conversations

D. A log file for debugging AI errors

Answer: B Correct

Explanation: The System of Record is the authoritative data layer that AI agents must access to act correctly - it prevents AI from "making things up" by providing verified information. Without it, agents operate on probabilistic generation rather than facts.

*Reference:** Agentic-AI.pdf pages 70, 77, Agentic-AI.pptx slides 70, 77*

Question 12

What does TutorClaw do in the Agent Factory ecosystem?

A. Builds AI agents automatically

B. Acts as a 24/7 AI tutor teaching from the Agent Factory book via WhatsApp, Telegram, and web

C. Monitors AI agent performance

D. Translates code between different programming languages

Answer: B Correct

Explanation: TutorClaw is one of the four delivery channels - an AI tutor accessible via WhatsApp, Telegram, and web that teaches from the Agent Factory book rather than the open internet. It's the conversational learning interface for the ecosystem.

*Reference:** Agent Factory website delivery channels section*

Question 13

How many delivery channels does the Agent Factory ecosystem have?

A. Two: the book and online courses

B. Three: book, software, and consulting

C. Four: the Book, TutorClaw, Agent Factory Skillpack, and Derivative Books

D. Five: book, tutorials, forums, workshops, and certification

Answer: C Correct

Explanation: The Agent Factory delivers through four channels: (1) The Book (canonical source), (2) TutorClaw (AI tutor), (3) Agent Factory Skillpack (building guidance for Claude Code/OpenCode), and (4) Derivative Books (specialized by topic or audience).

*Reference:** Agent Factory website delivery system section*

Question 14

According to the presentation, what role does a personal AI delegate play?

A. It replaces your job entirely

B. It's your chief of staff that knows your preferences and manages AI workers for you

C. It only handles email responses

D. It generates reports at the end of each day

Answer: B Correct

Explanation: The personal AI delegate serves as your chief of staff - it knows your preferences, carries your authority, coordinates specialist AI workers, and reports back to you. You don't manage 20 AI workers directly; your delegate manages them for you.

*Reference:** Agentic-AI.pdf pages 62-64, Agentic-AI.pptx slides 62-64*

Question 15

What is the difference between General Agents and Personal Agents in terms of lifespan?

A. General agents are permanent, personal agents are temporary

B. General agents are sessional (summoned per job), personal agents are always-on and persistent

C. Both have the same lifespan

D. General agents live longer than personal agents

Answer: B Correct

Explanation: General agents are summoned for specific jobs and exist for that session only, while personal agents are always-on and persistent, maintaining long-term memory of you. General agents are your workforce operators; personal agents are your persistent delegates.

Reference: Agentic-AI.pdf page 26, Agentic-AI.pptx slide 26

Question 16

In the "Fish vs. Fishing Net" analogy, which mode does "building a net" represent?

A. Mode 1 - Problem-Solving

B. Mode 2 - Manufacturing

C. Mode 3 - Optimization

D. Training mode

Answer: B Correct

Explanation: "Build a net" represents Mode 2 (Manufacturing) - building something that catches fish every day, creating sustainable results. "Catch a fish" represents Mode 1 (Problem-Solving) - solving today's immediate problem. The analogy emphasizes building for tomorrow versus solving for today.

Reference: Agentic-AI.pptx slide 46

Question 17

What are the three things that always stay with humans according to the 10-80-10 framework?

A. Coding, testing, deploying

B. Intent, verification, and ownership

C. Planning, executing, reviewing

D. Reading, writing, approving

Answer: B Correct

Explanation: The three human responsibilities are: Intent (knowing what you want), Verification (checking if the work is good), and Ownership (being responsible for the result). AI handles the 80% execution, but humans maintain control over direction and judgment.

Reference: Agentic-AI.pdf page 51, Agentic-AI.pptx slide 51

Question 18

According to the presentation, what happens to you in the AI era?

A. You get replaced by AI

B. You get promoted from worker to director

C. You become unemployed

D. You need to learn to code

Answer: B Correct

Explanation: The thesis emphasizes "you're promoted, not replaced" - from typist to editor, from coder to architect, from worker to director. AI takes over the middle 80% execution work, while you keep the direction-setting and judgment roles that matter most.

Reference: Agentic-AI.pdf page 52, Agentic-AI.pptx slide 52

Question 19

What does SKILL.md represent in the Agent Factory technical architecture?

A. A resume format for AI agents

B. Plain-text files packaging reusable AI capabilities

C. A programming language for AI

D. A skill assessment test

Answer: B Correct

Explanation: SKILL.md files are plain-text specifications that package reusable AI capabilities. They're portable across Claude Code and OpenCode, allowing skills written for one tool to work unchanged in the other. They're the standardization layer for agent behaviors.

Reference: Agent Factory website technical architecture section

Question 20

What is an "AI-Native Company" according to the thesis?

A. A company that only sells AI products

B. A company where most workers are AI, with small human teams overseeing them

C. A company founded after 2020

D. A company using cloud computing

Answer: B Correct

Explanation: An AI-Native Company operates with many Digital FTEs alongside small human teams, selling results rather than tools. Example: 10 human supervisors + 190 AI workers replacing a traditional 200-person team, with AI handling routine work and humans managing exceptions.

Reference: Agentic-AI.pdf pages 59, 76, Agentic-AI.pptx slides 59, 76

Question 21

On June 1, 2026, what did NVIDIA and Microsoft unveil according to the presentation?

A. A new operating system

B. A petaflop-class chip (RTX Spark) to run agents locally on Windows

C. A quantum computer

D. A new programming language

Answer: B Correct

Explanation: NVIDIA and Microsoft unveiled RTX Spark, a petaflop-class chip enabling frontier-model AI agents to run locally on Windows devices with ~1 PF of AI performance, 128GB unified memory, and 20 Arm CPU cores - eliminating cloud round-trips for agent operations.

Reference: Agentic-AI.pdf page 18, Agentic-AI.pptx slide 18

Question 22

What does Jensen Huang mean by "you ask - and the PC does the work"?

A. Voice commands replace typing

B. The agent layer replaces human operation of the machine

C. Automation scripts run tasks

D. AI generates code for you

Answer: B Correct

Explanation: Huang's statement represents the paradigm shift where the agent operating layer takes over - you state intent, and the agent operates the machine on your behalf. The PC stops being an instrument you play and becomes an actor you direct.

***Reference:** Agentic-AI.pdf pages 18, 31, Agentic-AI.pptx slides 18, 31*

Question 23

According to the thesis, what is the "smaller death" in the agentic era?

A. Traditional programming languages die

B. SaaS dies - apps dissolve into function calls

C. Desktop computers die

D. Cloud computing dies

Answer: B Correct

Explanation: The "smaller event" is SaaS death: the agent replaces the app, unbundling login/navigation/UI into capabilities an agent calls via API or MCP. The app dissolves into a function call rather than a destination you visit.

***Reference:** Agentic-AI.pdf page 19, Agentic-AI.pptx slide 19*

Question 24

What is the "larger death" according to the agentic era thesis?

A. Microsoft Windows ends

B. The internet becomes obsolete

C. The PC as we know it dies - you stop operating the machine

D. Programming jobs disappear

Answer: C Correct

Explanation: The "larger event" is the death of the PC operating model: the agent replaces you at the controls, making the app-on-OS model driven by hand through a graphical shell obsolete. Not the silicon, but the human's job of driving it ends.

***Reference:** Agentic-AI.pdf page 19, Agentic-AI.pptx slide 19*

Question 25

In the SaaS unbundling model, which layer becomes "the prize"?

A. Workflow UI

B. Capabilities

C. System of Record

D. Marketing pages

Answer: C Correct

Explanation: When SaaS apps unbundle, the System of Record becomes the strategic prize - the authoritative data an agent must read to act. Workflow UI dies first, capabilities survive as function calls, but the data layer becomes the most valuable part of the stack.

***Reference:** Agentic-AI.pdf page 20, Agentic-AI.pptx slide 20*

Question 26

What happens to the Workflow UI layer when SaaS gets unbundled by agents?

A. It gets upgraded with better graphics

B. It dies first because agents operate capabilities directly and bypass screens entirely

C. It becomes the main interface

D. It merges with the capabilities layer

Answer: B Correct

Explanation: The Workflow UI dies first because it existed only so humans could operate the capabilities. When an agent operates capabilities directly via API or MCP, it bypasses the screens entirely, making the UI layer obsolete.

***Reference:** Agentic-AI.pdf page 20, Agentic-AI.pptx slide 20*

Question 27

What does the Agent Factory Skillpack provide?

A. A certification program for AI developers

B. Guided building agent with verified patterns and SKILL.md templates integrated into Claude Code/OpenCode

C. A marketplace for buying pre-built AI agents

D. A testing suite for AI models

Answer: B Correct

Explanation: The Agent Factory Skillpack is one of the four delivery channels - it's a building guidance agent equipped with verified patterns and SKILL.md templates that works as a partner inside Claude Code and OpenCode.

***Reference:** Agent Factory website delivery channels section*

Question 28

According to the three levels of AI, which level "predicts"?

A. Level 1 - Predictive AI

B. Level 2 - Generative AI

C. Level 3 - Agentic AI

D. All levels predict equally

Answer: A Correct

Explanation: Level 1 (Predictive AI) is "The Smart Filter" that gives you data and returns one prediction (spam or not spam, product recommendation). It has no conversation, no memory, just predictions based on patterns.

***Reference:** Agentic-AI.pptx slide 12*

Question 29

Which level of AI "creates and chats"?

A. Level 1 - Predictive AI

B. Level 2 - Generative AI

C. Level 3 - Agentic AI

D. None of the above

Answer: B Correct

Explanation: Level 2 (Generative AI) is "The Chat Partner" - you have a conversation, it remembers context, creates text/images/code, but still waits for your next instruction. Examples include ChatGPT and Claude in browser chat mode.

***Reference:** Agentic-AI.pptx slide 12*

Question 30

Which level of AI "plans and delivers"?

A. Level 1 - Predictive AI

B. Level 2 - Generative AI

C. Level 3 - Agentic AI

D. Level 4 - Superintelligence

Answer: C Correct

Explanation: Level 3 (Agentic AI) is "The Worker" - you give it a goal, and it plans, uses tools, fixes mistakes, and finishes the job without you guiding every step. It moves from conversation to actual work completion.

***Reference:** Agentic-AI.pptx slide 12*

Question 31

What is the role of general agents according to the two-layer model?

A. Your personal assistant

B. Your workforce/operators

C. Your data storage

D. Your security system

Answer: B Correct

Explanation: In the two-layer model, general agents serve as your workforce/operators - they are summoned per job to do specific work. Personal agents are your delegate/interface, while general agents are the workers that get dispatched to complete tasks.

***Reference:** Agentic-AI.pdf page 26, Agentic-AI.pptx slide 26*

Question 32

What does "governed capability" mean in the context of RTX Spark?

A. Government regulations on AI

B. Control over what an agent may touch and what may leave the machine

C. Speed limits on AI processing

D. User interface design guidelines

Answer: B Correct

Explanation: NVIDIA's OpenShell governs what an agent may touch and what may leave the machine on RTX Spark devices. The interesting engineering is not just the petaflop performance - it's the governed capability that ensures security and control.

***Reference:** Agentic-AI.pdf page 30, Agentic-AI.pptx slide 30*

Question 33

In the film director analogy for the 10-80-10 rhythm, what does the director NOT do?

A. Set the vision

B. Do every task themselves

C. Review the final cut

D. Guide the work

Answer: B Correct

Explanation: The film director analogy illustrates that directors don't do every task - they set the vision (first 10%), let the team work (middle 80%), and review the final cut (last 10%). This mirrors how humans work with AI agents.

***Reference:** Agentic-AI.pptx slide 50*

Question 34

What is the "carpenter's drill vs. hiring a contractor" analogy illustrating?

A. The difference between Mode 1 and Mode 2

B. The difference between AI as a TOOL vs. AI as a WORKER

C. The difference between hardware and software

D. The difference between local and cloud AI

Answer: B Correct

Explanation: This analogy contrasts AI as a tool (you hold it, you decide, you do the work - the drill makes one part faster) versus AI as a worker (you give the job, define what you want, it handles the task, you review the result).

Reference: Agentic-AI.pptx slide 54

Question 35

What are the three things a SaaS app is made of before agents unbundle it?

A. Frontend, backend, database

B. Workflow UI, Capabilities, and System of Record

C. Login, dashboard, reports

D. Users, data, algorithms

Answer: B Correct

Explanation: A SaaS app consists of three welded-together components: Workflow UI (dies first when agents bypass screens), Capabilities (survive as function calls agents reach for), and System of Record (the prize - authoritative data agents must read).

Reference: Agentic-AI.pdf page 20, Agentic-AI.pptx slide 20

Question 36

According to the "game" analogy in the presentation, when kids figure out how to get ready for the playground on their own, what concept does this demonstrate?

A. Teamwork

B. Agentic behavior - having a goal and making your own plan

C. Following instructions

D. Time management

Answer: B Correct

Explanation: The playground game demonstrates agentic behavior: given just the goal "get everyone ready to go play," the kids figure out their own plan (wear shoes, grab the ball, line up) without being told each step - exactly what agentic AI does.

Reference: Agentic-AI.pptx slide 9

Question 37

How many "Arm CPU cores" does RTX Spark include?

A. 10

B. 16

C. 20

D. 32

Answer: C Correct

Explanation: RTX Spark includes 20 Arm CPU cores alongside ~1 petaflop of AI performance and 128GB unified memory. This combination enables frontier-model agents to run entirely on-device with local, private computation.

Reference: Agentic-AI.pdf page 30, Agentic-AI.pptx slide 30

Question 38

What is the approximate AI performance of RTX Spark?

A. 1 teraflop

B. 1 petaflop

C. 1 gigaflop

D. 1 exaflop

Answer: B Correct

Explanation: RTX Spark delivers approximately 1 petaflop (PF) of AI performance on-device. This petaflop-class performance enables running frontier-model agents locally without cloud round-trips, making agent operations private and always-available.

***Reference:** Agentic-AI.pdf page 30, Agentic-AI.pptx slide 30*

Question 39

In the real example of Ecomm Mart, what is Step 1 in building the Support Worker?

A. Test the worker

B. Connect to company records

C. Write clear instructions (define goal, limits, budget)

D. Deploy to production

Answer: C Correct

Explanation: Step 1 is "Write Clear Instructions" - define the goal (read tickets, classify, write replies, escalate hard cases) and set constraints like budget ($200/month max). Clear specifications come first before any technical implementation.

***Reference:** Agentic-AI.pdf page 72, Agentic-AI.pptx slide 72*

Question 40

In building the Ecomm Mart Support Worker, what does Step 2 involve?

A. Writing code

B. Connecting to Company Records (Shopify orders, customer database, ticket system)

C. Hiring developers

D. Creating marketing materials

Answer: B Correct

Explanation: Step 2 is "Connect to Company Records" - linking the AI worker to Shopify orders, customer database, and ticket system so it can look up real data. This System of Record connection prevents the AI from making things up.

***Reference:** Agentic-AI.pdf page 72, Agentic-AI.pptx slide 72*

Question 41

What does Step 3 involve when building the Ecomm Mart Support Worker?

A. Launching to customers

B. Giving it skills (how to write professional replies, classify tickets, when to escalate)

C. Purchasing licenses

D. Training human staff

Answer: B Correct

Explanation: Step 3 is "Give It Skills" - teaching the worker how to write professional replies, how to classify tickets, and when to escalate issues. This is like creating a training manual that defines the worker's capabilities.

***Reference:** Agentic-AI.pdf page 72, Agentic-AI.pptx slide 72*

Question 42

What percentage of support tickets does the Ecomm Mart Digital FTE handle automatically after deployment?

A. 50%

B. 60%

C. 80%

D. 100%

Answer: C Correct

Explanation: The completed Digital FTE for Ecomm Mart handles 80% of support tickets automatically, working 24/7 within defined rules. The remaining 20% presumably require human judgment for complex or exceptional cases.

***Reference:** Agentic-AI.pdf page 73, Agentic-AI.pptx slide 73*

Question 43

According to the presentation, what is the traditional "Old Way" time to write a business report?

A. 1 hour

B. 2 hours

C. 4 hours

D. 8 hours

Answer: C Correct

Explanation: The "Old Way" example states you spend 4 hours writing a business report yourself. The "New Way" with AI reduces this to 45 minutes total (20 min instructions + 5 min AI writing + 20 min review) - one-tenth the time for same quality.

***Reference:** Agentic-AI.pdf page 52, Agentic-AI.pptx slide 52*

Question 44

In the "New Way" of writing a business report with AI, how long does the AI take to actually write it?

A. 1 minute

B. 5 minutes

C. 20 minutes

D. 45 minutes

Answer: B Correct

Explanation: With AI assistance, writing clear instructions takes 20 minutes, the AI writes the report in 5 minutes, and reviewing/fixing takes 20 minutes - total 45 minutes versus 4 hours the old way, achieving the same quality in one-tenth the time.

***Reference:** Agentic-AI.pdf page 52, Agentic-AI.pptx slide 52*

Question 45

What does "MCP" stand for in the Agent Factory technical architecture?

A. Multi-Cloud Platform

B. Model Context Protocol

C. Master Control Program

D. Machine Cognitive Processing

Answer: B Correct

Explanation: MCP Connectors are mentioned as standard interfaces integrating external tools and data in the technical architecture. MCP stands for Model Context Protocol, which enables agents to connect with external tools and data sources.

***Reference:** Agent Factory website technical architecture section*

Question 46

According to the 7 Golden Rules, what happens when a gap appears in the workforce?

A. Human managers must manually hire

B. The system shuts down

C. The system hires a new worker automatically within human-set limits (Rule #6)

D. Work gets queued indefinitely

Answer: C Correct

Explanation: Rule #6 states "Workforce can grow under rules" - when a gap appears, the system hires a new worker automatically within human-set limits. This enables dynamic scaling while maintaining human control over growth boundaries.

*Reference:** Agentic-AI.pdf page 70, Agentic-AI.pptx slide 70*

Question 47

What does Rule #7 (Company runs on a nervous system) mean?

A. AI monitors employee health

B. The system sends work between workers automatically and recovers from breaks on its own

C. The company uses neural networks

D. Human managers coordinate all tasks

Answer: B Correct

Explanation: Rule #7 compares the AI company to a body's nervous system - signals are sent automatically ("hand touched something hot, pull back!"). Work passes between workers automatically without human intervention, and the system self-recovers from failures.

*Reference:** Agentic-AI.pdf page 70, Agentic-AI.pptx slide 70*

Question 48

According to Rule #4, how should worker engines be selected?

A. Always use the most expensive engine

B. Always use the cheapest engine

C. Important work gets reliable engine, simple work gets cheap engine

D. Random selection based on availability

Answer: C Correct

Explanation: Rule #4 "Each worker uses the right engine" means matching the engine to the work's importance - critical work requires reliable, capable engines while routine tasks can use cheaper, lighter engines. This optimizes for both quality and cost.

*Reference:** Agentic-AI.pdf page 69, Agentic-AI.pptx slide 69*

Question 49

What does Rule #2 state about personal delegates?

A. Only executives get delegates

B. Delegates are optional features

C. Every human gets a delegate (identic AI)

D. Delegates cost extra

Answer: C Correct

Explanation: Rule #2 states "Every human gets a delegate (identic AI)" because you can't manage 20 workers by hand - your delegate coordinates them. This is a foundational requirement for AI-Native companies, not an optional feature.

*Reference:** Agentic-AI.pdf page 69, Agentic-AI.pptx slide 69*

Question 50

In the OSWorld benchmark mentioned, what was the agent success rate that crossed the human baseline?

A. 12%

B. 45%

C. 66%

D. 95%

Answer: C Correct

Explanation: Agent success on OSWorld (real desktops, no partial credit) jumped from ~12% to ~66% in a year, crossing the human baseline. This ~66% success rate is one of three convergent factors that make "this time different" for agentic computing.

*Reference:** Agentic-AI.pdf page 37, Agentic-AI.pptx slide 37*

INTERMEDIATE LEVEL (Questions 51-100)

Question 51

A startup wants to build an AI customer support system but has limited technical staff. They need to decide between building in Mode 1 (session-based problem solving) or Mode 2 (manufacturing autonomous workers). Which approach best fits their constraint?

A. Mode 1, because it requires no infrastructure and can be used immediately for customer queries

B. Mode 2, because building one Digital FTE that runs 24/7 reduces the need for ongoing human involvement

C. Mode 1, because Mode 2 requires a full engineering team to maintain

D. Neither mode works for startups with limited staff

Answer: B Correct

Explanation: Mode 2 (Manufacturing) is more suitable for resource-constrained teams because once built, the Digital FTE runs autonomously 24/7 with minimal oversight. Mode 1 requires humans to be in the loop for each session, which doesn't scale with limited staff. The initial investment pays off through autonomous operation.

*Reference:** Agentic-AI.pdf pages 47, Agentic-AI.pptx slides 46-47*

Question 52

An organization is evaluating where to place their strategic investment as agents become the operating layer. According to the thesis, which layer offers the most defensible moat?

A. Building a beautiful, sticky UI that users love

B. Being the layer the agent lives in, the capability it must call, or the governance it must obey

C. Having the fastest response times

D. Offering the lowest pricing

Answer: B Correct

Explanation: The thesis states "the agent has no eyes for your interface" - beautiful UIs no longer provide defense. The strategic ground shifts to being the layer agents operate in, the capabilities they must invoke, or the governance framework they must follow. System of Record becomes the prize.

*Reference:** Agentic-AI.pdf page 41, Agentic-AI.pptx slide 41*

Question 53

When building the Ecomm Mart Support Worker, why is Step 4 (Set the Guardrails) critical before deployment?

A. It improves response speed

B. It ensures the AI can never go outside defined limits like budget caps and approval requirements

C. It makes the AI more intelligent

D. It reduces cloud costs

Answer: B Correct

Explanation: Step 4 establishes guardrails (budget cap, approval for big refunds, audit trail) that the AI can never exceed. This ensures safety, compliance, and accountability - the AI operates within bounded, recoverable parameters rather than having unlimited authority.

***Reference:** Agentic-AI.pdf page 73, Agentic-AI.pptx slide 73*

Question 54

According to the thesis, what is the relationship between personal agents and general agents in the recursive layer model?

A. They are competitors for the same tasks

B. Personal agents are built using general agents, then at runtime the personal agent dispatches general agents to do work

C. General agents always run before personal agents

D. They operate on separate, non-interacting layers

Answer: B Correct

Explanation: The layer is recursive: you use developer-grade general agents (Claude Code, OpenCode) to build and configure your personal agent. Then at runtime, your personal delegate dispatches general agents and AI Workers to carry out work and reports back. The personal agent is your chief of staff; general agents are the staff.

***Reference:** Agentic-AI.pdf page 28, Agentic-AI.pptx slide 28*

Question 55

A company currently uses 200 human support agents. They're evaluating the AI-Native model. According to the presentation example, what would a transformed structure look like?

A. 200 AI workers with no humans

B. 100 humans and 100 AI workers

C. 10 humans supervising 190 AI workers

D. 190 humans with 10 AI assistants

Answer: C Correct

Explanation: The AI-Native Company example shows 10 human supervisors handling tricky cases and overseeing 190 AI workers that handle routine questions 24/7. AI handles the bulk of routine work while humans manage exceptions and complex cases - not full replacement, but dramatic transformation.

***Reference:** Agentic-AI.pdf page 59, Agentic-AI.pptx slide 59*

Question 56

Why does the thesis emphasize that Digital FTEs are "both a technical AND economic construct"?

A. Because they require expensive hardware

B. Because they replace headcount and enable scaling intelligence rather than scaling headcount

C. Because they need economic models to function

D. Because they cost money to run

Answer: B Correct

Explanation: Digital FTEs are economic constructs because they fundamentally change the cost structure - enabling organizations to "scale intelligence, not headcount" (as demonstrated by Klarna's ~700 FTE equivalent). They're not just technical tools but workforce replacements with direct P&L impact.

***Reference:** Agent Factory website core concepts, Agentic-AI.pdf page 41*

Question 57

According to the 7 Golden Rules, why must "every human get a delegate" rather than making delegates optional?

A. It increases software sales

B. Because you can't manage 20+ AI workers by hand - your delegate coordinates them for you

C. It's a nice-to-have feature for executives

D. To justify higher pricing

Answer: B Correct

Explanation: Rule #2 establishes delegates as mandatory infrastructure, not optional features. The operational reality is that humans cannot effectively manage 20+ AI workers directly - the delegate layer is necessary for coordination, not a luxury. It's structural, not cosmetic.

***Reference:** Agentic-AI.pdf page 69, Agentic-AI.pptx slide 69*

Question 58

What is the primary advantage of the "bounded, recoverable tasks go first" principle mentioned in the honest objections section?

A. It makes AI cheaper

B. It allows organizations to gain experience with agents in low-risk scenarios before high-stakes deployment

C. It requires less computing power

D. It avoids regulatory issues

Answer: B Correct

Explanation: With ~66% average success rates, starting with bounded (limited scope) and recoverable (fixable if wrong) tasks allows organizations to learn agent orchestration in contexts where failures aren't catastrophic. This staged adoption manages the reliability gap pragmatically.

***Reference:** Agentic-AI.pdf page 38, Agentic-AI.pptx slide 38*

Question 59

In the Agent Factory flow (Human -> Factory -> Workers -> Results), what is the factory's role?

A. It stores the final outputs

B. It builds and orchestrates AI workers based on human direction

C. It provides computing resources

D. It trains the AI models

Answer: B Correct

Explanation: The Agent Factory is the process that manufactures AI Workers and composes them into an AI-Native Company. It transforms general-agent capability (Claude Code, OpenCode) into specialized Digital FTEs, then orchestrates them into a coordinated system.

***Reference:** Agentic-AI.pptx slides 60-61*

Question 60

Why does the thesis state "Tools change. Rules stay" when discussing the 7 Golden Rules?

A. Because the rules are legally mandated

B. Because structural principles (like human oversight, system of record, right-engine-for-task) remain constant even as specific tools evolve

C. Because old tools are better

D. Because rules never need updating

Answer: B Correct

Explanation: The 7 Golden Rules are invariants - structural principles that hold regardless of which specific tools you use. Like building codes that remain constant while construction materials vary, these rules provide enduring architecture while allowing tool ecosystem evolution.

***Reference:** Agentic-AI.pdf page 70, Agentic-AI.pptx slides 68, 70*

Question 61

When the presentation states "The firm that orchestrates agents out-produces the firm that buys more seats," what core shift is being described?

A. Subscription software is getting more expensive

B. Competitive advantage moves from accumulating SaaS licenses to orchestrating AI workforce

C. Hiring is becoming more difficult

D. Remote work is becoming standard

Answer: B Correct

Explanation: This captures the fundamental competitive shift: buying seats of traditional software scales linearly with cost, while orchestrating agents scales intelligence without proportional headcount increase. Klarna's example (~700 FTE work, ~$40M profit lift) demonstrates this leverage.

***Reference:** Agentic-AI.pdf page 41, Agentic-AI.pptx slide 41*

Question 62

What does the "SaaSpocalypse" fundamentally represent?

A. The end of cloud computing

B. The unbundling of SaaS apps into capabilities agents call, making the bundled product obsolete

C. A security breach in SaaS systems

D. The transition from on-premise to cloud

Answer: B Correct

Explanation: The SaaSpocalypse is the unbundling of SaaS: apps dissolve into function calls as agents operate capabilities directly via API/MCP. The login, navigation, and seat-based UI die because agents bypass the wrapper. The bundle was the business; agents destroy that business model.

*Reference:** Agentic-AI.pdf page 20, Agentic-AI.pptx slide 20*

Question 63

According to the thesis, what becomes "invisible infrastructure" in the agentic era?

A. Cloud providers

B. The OS (Windows, macOS, Linux) - like TCP/IP or BIOS, humans never touch it again

C. Programming languages

D. Databases

Answer: B Correct

Explanation: The operating system "becomes plumbing" - surviving as invisible infrastructure like TCP/IP or BIOS. The desktop, dock, and app grid become legacy surfaces behind the agent layer. Humans interact with the agent layer, not the OS directly.

*Reference:** Agentic-AI.pdf page 24, Agentic-AI.pptx slide 24*

Question 64

Why is the hybrid approach (human + UI + agent) described as "the thesis in its transitional phase, not an alternative to it"?

A. Because hybrids are always temporary

B. Because high-stakes work currently requires the hybrid until reliability improves, but the direction remains full agent operation

C. Because humans prefer using UIs

D. Because agents can't work with humans

Answer: B Correct

Explanation: The honest objections section acknowledges that ~66% reliability means high-stakes work stays human+UI+agent for now. However, this is positioned as transitional (reliability will improve, bounded tasks will expand) rather than a permanent alternative architecture.

*Reference:** Agentic-AI.pdf page 38, Agentic-AI.pptx slide 38*

Question 65

What is the significance of RTX Spark being "local" rather than cloud-based?

A. It's faster for gaming

B. Your data stays private on-device, with no cloud round-trip, making agents always-available and unmetered

C. It uses less electricity

D. It's cheaper to manufacture

Answer: B Correct

Explanation: Local petaflop performance means private (your data never leaves), always-available (no internet required), and unmetered (no per-token costs) agent compute. This shifts agents from cloud subscription services to native device capabilities.

Reference: Agentic-AI.pdf page 30, Agentic-AI.pptx slide 30

Question 66

In the two-layer model, what distinguishes the EDGE LAYER from the AI WORKFORCE LAYER?

A. The edge is faster

B. The edge layer is your personal delegate (knows your preferences, manages workers); the workforce layer is specialized workers (support, finance, sales, legal)

C. The edge is more expensive

D. The edge handles storage

Answer: B Correct

Explanation: The EDGE LAYER contains your personal AI delegate that knows your preferences, carries your authority, and manages workers. The AI WORKFORCE LAYER contains specialized workers (Support, Finance, Sales, Legal). You interact with your delegate; the delegate coordinates the workforce.

Reference: Agentic-AI.pdf page 64, Agentic-AI.pptx slide 64

Question 67

What does "the agent replaces you at the controls" mean for the PC's operating model?

A. Agents control hardware directly

B. The app-on-OS model driven by hand through a graphical shell becomes obsolete - you state intent, agent operates the machine

C. Users can no longer use computers

D. Operating systems are eliminated

Answer: B Correct

Explanation: This describes the larger death: not the silicon, but the human's job of driving it. The forty-year stack flips - you stop launching apps and navigating UIs; instead you state goals and the agent layer reaches down to operate the machine on your behalf.

Reference: Agentic-AI.pdf pages 19, 24, Agentic-AI.pptx slides 19, 24

Question 68

Why does Rule #5 state "Every worker uses a system of record" rather than allowing probabilistic generation?

A. It's legally required

B. Without authoritative data to read, AI makes things up - the system of record prevents fabrication by providing verified facts

C. It improves speed

D. It reduces costs

Answer: B Correct

Explanation: Rule #5 prevents AI hallucination: workers must read/write the company's official memory (System of Record). Without this grounding in authoritative data, AI operates on probabilistic generation and invents information. The System of Record ensures factual accuracy.

***Reference:** Agentic-AI.pdf page 70, Agentic-AI.pptx slide 70*

Question 69

What is the strategic implication of "the bundle was the business" in the SaaSpocalypse?

A. SaaS companies should bundle more features

B. The business model of selling bundled UI+capabilities+data collapses when agents unbundle them and only need data access

C. Bundling will increase in importance

D. Companies should focus on packaging

Answer: B Correct

Explanation: SaaS companies monetized the bundle - you paid for the wrapper (UI) to access capabilities and data. Agents demolish this by accessing capabilities directly via API/MCP and only needing the System of Record. The bundle that was the business model becomes unbundled infrastructure.

***Reference:** Agentic-AI.pdf page 20, Agentic-AI.pptx slide 20*

Question 70

In the context of the three convergent factors that make "this time different," why is the OS vendor rebuild significant?

A. It creates new operating systems

B. Microsoft wiring agents into Windows behind every surface makes it platform infrastructure, not an app - the change is systemic

C. It increases competition

D. It reduces hardware costs

Answer: B Correct

Explanation: The third convergence factor is that OS vendors (Microsoft) wired agents into the platform itself behind every surface, with eight makers shipping RTX Spark PCs. This makes agents platform infrastructure rather than an optional app - a systemic architectural change.

***Reference:** Agentic-AI.pdf page 37, Agentic-AI.pptx slide 37*

Question 71

When Klarna's customer-service agent did the work of ~700 FTEs, why were humans "re-added for complex cases"?

A. To save costs

B. Because bounded tasks go first - AI handles routine work at scale, humans handle exceptions requiring nuanced judgment

C. Due to regulations

D. Because the AI failed

Answer: B Correct

Explanation: This exemplifies the "bounded, recoverable tasks first" principle: AI handles the 80% of routine, repetitive work at massive scale (~700 FTE equivalent), while humans are re-added for the 20% of complex cases requiring judgment beyond current agent reliability.

Reference: Agentic-AI.pdf page 41, Agentic-AI.pptx slide 41

Question 72

What does the Agent Factory's "four-pillar approach" consist of?

A. Planning, execution, verification, deployment

B. Structured specifications, domain expertise, engineering architecture, and human oversight

C. Frontend, backend, database, API

D. Development, testing, staging, production

Answer: B Correct

Explanation: The Agent Factory methodology rests on four pillars: clearly defined specifications (what to build), domain knowledge (expertise embedded), engineering infrastructure (how it runs), and human supervision (oversight and verification). These four together enable reliable Digital FTEs.

Reference: Agent Factory website architecture section

Question 73

In the Ecomm Mart example, why is Step 5 "Test, Then Deploy" sequenced last?

A. To save time

B. Testing on old tickets validates the worker before exposing it to real customers, catching problems in a safe environment

C. It's optional

D. Testing is always last

Answer: B Correct

Explanation: Step 5 runs the worker on 100 old tickets first to check results and fix problems before live customer deployment. This validates in a bounded, recoverable context (historical data) before the unbounded, risky context (live customers) - pragmatic risk management.

Reference: Agentic-AI.pdf page 73, Agentic-AI.pptx slide 73

Question 74

What does "scale intelligence, not headcount" mean in the context of Digital FTEs?

A. Hire smarter people

B. Increase capacity by adding AI workers rather than human employees, growing capability without proportional cost increase

C. Focus on IQ tests

D. Reduce team size

Answer: B Correct

Explanation: This phrase (from Klarna's case) captures the economic transformation: traditional scaling adds headcount linearly with costs; Digital FTE scaling adds intelligence/capacity with sub-linear cost growth. You grow capability without growing payroll proportionally.

***Reference:** Agentic-AI.pdf page 41, Agentic-AI.pptx slide 41*

Question 75

Why does the thesis position SKILL.md files as "portable" between Claude Code and OpenCode?

A. They use a standard programming language

B. They're plain-text specifications that work unchanged across both tools, making skills tool-agnostic

C. They're stored in the cloud

D. They're open-source

Answer: B Correct

Explanation: SKILL.md files package reusable AI capabilities in a portable format - write once, run on both Claude Code and OpenCode. This tool-agnosticism means your investment in skills transcends any single platform, reducing lock-in.

***Reference:** Agent Factory website technical architecture section*

Question 76

According to the presentation's example, in an AI-Native Company with 10 humans and 190 AI workers, what do the 10 humans primarily handle?

A. All customer interactions

B. Tricky cases and oversight of AI - the exceptions requiring human judgment

C. Marketing only

D. Software development

Answer: B Correct

Explanation: The 10 human supervisors handle tricky/complex cases that exceed current AI capability and oversee the AI workforce. The 190 AI workers handle routine questions 24/7. This division follows the bounded-tasks-first principle: AI scales the routine, humans handle the nuanced.

***Reference:** Agentic-AI.pdf page 59, Agentic-AI.pptx slide 59*

Question 77

What does "the interface is the agent" mean for the future of computing?

A. Agents will have better graphics

B. The agent layer becomes your primary interaction point, replacing app navigation and UI manipulation

C. Interfaces will be faster

D. Agents will design UIs

Answer: B Correct

Explanation: "The interface is the agent" captures the shift from humans operating apps through UIs to humans delegating to agents. Your interface becomes conversation with your delegate rather than clicking through application interfaces. The agent is the computer interface.

***Reference:** Agentic-AI.pdf page 42, Agentic-AI.pptx slide 42*

Question 78

Why does the thesis describe general agents as "summoned per job" rather than persistent?

A. To save memory

B. Because they're task-oriented tools invoked for specific work sessions, not always-on entities

C. It's a technical limitation

D. To reduce costs

Answer: B Correct

Explanation: General agents (Claude Code, OpenCode) have sessional lifespans - you summon them for a specific job, they complete it, then the session ends. They're workforce operators oriented toward tasks, not persistent delegates oriented toward you. This contrasts with personal agents that are always-on.

***Reference:** Agentic-AI.pdf page 26, Agentic-AI.pptx slide 26*

Question 79

In the context of Derivative Books as a delivery channel, what makes them "specialized"?

A. They cost more

B. They're organized by topic (e.g., Python in AI Era) or audience (profession-specific, age-appropriate)

C. They're only available in hardcover

D. They require certification

Answer: B Correct

Explanation: Derivative Books are specialized editions of the Agent Factory knowledge organized by topic (focusing on specific technologies like Python) or by audience (tailored for different professions or age groups). They make the core methodology accessible to specific contexts.

***Reference:** Agent Factory website delivery channels section*

Question 80

What is the strategic significance of Postgres being described as the "System of Record" in the technical stack?

A. It's the most popular database

B. It's a single database holding relational data, documents, full-text search, and AI vectors - consolidating authoritative data

C. It's free and open-source

D. It has good performance

Answer: B Correct

Explanation: Postgres as the System of Record is strategic because it consolidates all authoritative data types (relational, documents, search, vectors) into one source of truth. This prevents data fragmentation and ensures agents read from verified, consistent information rather than scattered sources.

***Reference:** Agent Factory website technical architecture section*

Question 81

When the thesis says "doing it yourself stops being the default," what gap is this closing?

A. The skill gap

B. The gap the Siri decade never closed - voice assistants were novelties, not reliable workforce replacements

C. The cost gap

D. The speed gap

Answer: B Correct

Explanation: The "Siri decade" saw voice assistants remain novelties because reliability, capability, and integration weren't sufficient for real work delegation. The convergence of 66% OSWorld success, local petaflop compute, and platform integration finally closes this gap.

***Reference:** Agentic-AI.pdf page 37, Agentic-AI.pptx slide 37*

Question 82

What does the "building code" analogy in the Golden Rules section illustrate?

A. Programming standards

B. Structural principles (7 rules) stay constant like building codes while tools (materials) vary

C. Construction industry practices

D. Legal requirements

Answer: B Correct

Explanation: The analogy shows that like building codes (roof protects, walls structure, foundation supports) which remain constant while materials vary (bricks vs concrete), the 7 Golden Rules are structural invariants while specific tools evolve. Principles persist; implementations change.

***Reference:** Agentic-AI.pptx slide 68*

Question 83

Why does Rule #3 state "Workforce needs a management layer" rather than letting AI workers self-organize?

A. To create management jobs

B. Because someone must hire, assign tasks, control budgets, and fire underperformers - governance requires structure

C. It's traditional corporate hierarchy

D. To slow down automation

Answer: B Correct

Explanation: Rule #3 establishes that AI workforces require managed governance: hiring decisions, task assignment, budget control, performance management. Self-organizing systems without management lack accountability, resource control, and quality enforcement - the management layer provides operational discipline.

Reference: Agentic-AI.pdf page 69, Agentic-AI.pptx slide 69

Question 84

In the Ecomm Mart example, why is the $200/month budget specified in Step 1 rather than being unlimited?

A. To save money

B. Setting budget constraints in the spec ensures the AI operates within bounded, accountable limits

C. Company policy

D. Technical limitation

Answer: B Correct

Explanation: The budget constraint is a guardrail defined upfront in the specification. It ensures the AI worker can't exceed defined spending limits, maintaining cost control and accountability. Bounded operation (including budget) is part of the spec-driven approach's safety model.

Reference: Agentic-AI.pdf page 72, Agentic-AI.pptx slide 72

Question 85

What distinguishes "browser chat" general agents from "environment" general agents like Claude Code?

A. Browser chat is faster

B. Browser chat is conversation-only (ask questions, get answers); environment general agents work inside files, tools, commands, and projects

C. Browser chat is free

D. Browser chat is more accurate

Answer: B Correct

Explanation: The presentation distinguishes browser-based LLM chat (ChatGPT, Claude, Gemini - conversation in a window) from general agents (Claude Code, OpenCode) that move beyond chat to work inside environments with files, codebases, tools, commands, workflows, and projects.

Reference: Agentic-AI.pptx slide 15

Question 86

According to the thesis, what makes the OSWorld benchmark significant for measuring agent progress?

A. It's the largest dataset

B. It tests agents on real desktops with no partial credit - measuring actual task completion, not just API calls

C. It's the fastest benchmark

D. It's endorsed by governments

Answer: B Correct

Explanation: OSWorld's significance is that it tests agents on real desktop environments without partial credit - measuring end-to-end task completion on actual systems. The jump from ~12% to ~66% (crossing human baseline) represents genuine capability increase, not gaming of narrow benchmarks.

Reference: Agentic-AI.pdf page 37, Agentic-AI.pptx slide 37

Question 87

Why does the thesis position the agent factory as "a practice, not a product"?

A. It's not for sale

B. It's a repeatable methodology you learn (design -> train -> deploy -> verify), not something you buy off the shelf

C. It's still in development

D. It's too complex to productize

Answer: B Correct

Explanation: The Agent Factory is a practice - a repeatable way of building and managing AI workers, like a manufacturing process. It's a methodology you learn and apply, not a product you purchase. The emphasis is on capability development rather than software consumption.

Reference: Agentic-AI.pptx slide 58

Question 88

In the garment factory analogy, what corresponds to AI worker manufacturing?

A. Making clothes

B. Having a system (Cutting -> Stitching -> Quality Check -> Packing) that produces output reliably, just like Design -> Train -> Deploy -> Verify for AI workers

C. Hiring seamstresses

D. Selling garments

Answer: B Correct

Explanation: The garment factory analogy shows that just as a factory has a repeatable process (cutting, stitching, quality check, packing) for producing shirts, the Agent Factory has a repeatable process (design, train, deploy, verify) for producing AI workers. Both are systematic manufacturing approaches.

Reference: Agentic-AI.pptx slide 58

Question 89

What does the thesis mean by "the moat moves"?

A. Companies should relocate

B. Competitive advantages shift from UI stickiness to being the layer agents operate in, the capabilities they call, or the governance they obey

C. Costs are increasing

D. Markets are changing

Answer: B Correct

Explanation: "The moat moves" describes the strategic shift: traditional moats (beautiful UI, user habits, switching costs) become irrelevant when agents bypass interfaces. New moats are being the infrastructure agents require, the capabilities they must invoke, or the governance frameworks they must follow.

***Reference:** Agentic-AI.pdf page 41, Agentic-AI.pptx slide 41*

Question 90

Why is "128 GB unified memory" significant for RTX Spark's agent capabilities?

A. It enables gaming

B. Large unified memory allows frontier models with extensive context windows to run on-device without cloud constraints

C. It reduces power consumption

D. It's industry standard

Answer: B Correct

Explanation: 128GB unified memory enables running large frontier models entirely on-device with their full context windows. This memory capacity, combined with petaflop AI performance, allows sophisticated agent operations locally rather than requiring cloud round-trips for model inference.

***Reference:** Agentic-AI.pdf page 30, Agentic-AI.pptx slide 30*

Question 91

In the Agent Factory flow, why does "Results delivered" loop back to "Human" rather than going to storage?

A. For archiving

B. Because the 10-80-10 rhythm requires human verification of results before acceptance - humans close the loop

C. For billing purposes

D. For performance metrics

Answer: B Correct

Explanation: Results return to humans for the final 10% - verification and approval. The flow is circular: human sets goal -> factory builds workers -> workers execute -> results return to human for judgment. This maintains human-in-the-loop oversight per the 10-80-10 rhythm and Golden Rule #1.

***Reference:** Agentic-AI.pptx slide 60*

Question 92

What does "spec-driven development" mean in the Agent Factory methodology?

A. Writing fast code

B. Writing specifications first that define goal, limits, budget, and success criteria before building

C. Using specialized hardware

D. Following industry standards

Answer: B Correct

Explanation: Spec-driven development means writing clear specifications upfront (like the Ecomm Mart Step 1: define goal, limits, budget) before implementation. The spec defines what the AI should do, its boundaries, and how success is measured - then AI builds to that spec.

***Reference:** Agent Factory website architecture, Agentic-AI.pdf page 72*

Question 93

According to the thesis, why do "capabilities survive" when SaaS unbundles?

A. They're the most valuable part

B. They're demoted from product to function call - agents reach for them via API/MCP then set them down, making them tools rather than destinations

C. They're legally protected

D. They're easiest to build

Answer: B Correct

Explanation: Capabilities survive unbundling but are demoted: from being "the product" you buy and navigate to being function calls agents invoke via API or MCP. They become tools in the agent's toolkit rather than destinations users visit. The business model shifts from seat licenses to API calls.

***Reference:** Agentic-AI.pdf page 20, Agentic-AI.pptx slide 20*

Question 94

What is the significance of "eight makers ship RTX Spark PCs this fall" in the convergence story?

A. More competition

B. Wide hardware manufacturer adoption signals ecosystem-wide commitment to agent-native computing, not a single vendor experiment

C. Lower prices

D. Better marketing

Answer: B Correct

Explanation: Eight different PC manufacturers shipping RTX Spark indicates industry-wide platform shift rather than isolated experimentation. This broad adoption makes agent-native PCs a category, not a niche product, signaling the industry's commitment to the architectural transition.

***Reference:** Agentic-AI.pdf page 37, Agentic-AI.pptx slide 37*

Question 95

In the two-layer model, why is the personal delegate at the "edge" rather than in the cloud?

A. It's cheaper

B. Being at the edge means it's closest to you - knows your preferences, carries your authority, responds quickly to your needs

C. Edge computing is faster

D. Privacy regulations require it

Answer: B Correct

Explanation: "Edge layer" in this context means proximity to the user - your personal delegate is your interface, knows you intimately, and carries your authority. It's the layer you interact with directly, while the AI workforce layer operates behind it. Edge here is about architectural position, not necessarily physical location.

***Reference:** Agentic-AI.pdf page 64, Agentic-AI.pptx slide 64*

Question 96

What does the thesis mean by "the era is no longer 'humans use apps' but 'humans delegate work'"?

A. Apps are getting simpler

B. The fundamental interaction model shifts from operating software to assigning tasks to agents

C. Humans do less work

D. Apps are being deprecated

Answer: B Correct

Explanation: This captures the paradigm shift: from humans as operators of software (clicking, typing, navigating) to humans as directors who delegate to agents. The computer becomes something you direct rather than operate. Interaction moves from manipulation to delegation.

***Reference:** Agentic-AI.pdf page 42, Agentic-AI.pptx slide 42*

Question 97

Why does the Agent Factory emphasize "verified knowledge" over "probabilistic generation" for the System of Record?

A. Verification is faster

B. Because agents operating on facts from an authoritative source are reliable; agents generating probabilistically make things up

C. It's required by law

D. It uses less compute

Answer: B Correct

Explanation: The System of Record pattern prioritizes verified knowledge over probabilistic generation to ensure reliability. Agents that read from authoritative data sources produce accurate outputs; agents that generate probabilistically can hallucinate. This architectural choice prioritizes correctness over convenience.

***Reference:** Agent Factory website system of record concept*

Question 98

In the context of the three convergent factors, why is it significant that the change happened "together"?

A. It saves costs

B. Each factor alone was insufficient - models needed hardware, hardware needed platform support, platform changes needed capable models for viable agent-native computing

C. Easier to market

D. Regulatory approval

Answer: B Correct

Explanation: The thesis emphasizes convergence: capable models without local compute would be cloud-dependent; powerful hardware without OS integration would be unutilized; platform changes without capable models would be empty. All three together create the conditions for agent-native computing to actually work.

***Reference:** Agentic-AI.pdf page 37, Agentic-AI.pptx slide 37*

Question 99

What does the "nervous system" metaphor in Rule #7 specifically illustrate about AI-Native companies?

A. They're biological

B. Signals and work pass automatically between components without manual coordination, with self-recovery when things break

C. They're complex

D. They require medical expertise

Answer: B Correct

Explanation: The nervous system metaphor shows automatic signal routing (work passes between AI workers without human intervention) and autonomous recovery (system detects and fixes breaks itself, like reflexes). This contrasts with traditional companies where humans manually coordinate handoffs.

***Reference:** Agentic-AI.pdf page 70, Agentic-AI.pptx slide 70*

Question 100

According to the thesis, what is the relationship between "the book" and the other delivery channels (TutorClaw, Skillpack, Derivatives)?

A. They're separate products

B. The book is the canonical source; the other channels teach from it (TutorClaw), build from it (Skillpack), or specialize it (Derivatives)

C. They compete with each other

D. They're independent

Answer: B Correct

Explanation: The book serves as the canonical source/system of record for the ecosystem. TutorClaw teaches from the book, the Skillpack builds using the book's patterns, and Derivative Books are specialized editions of it. This mirrors the System of Record architectural pattern - one authoritative source feeding multiple interfaces.

***Reference:** Agent Factory website delivery system and system of record sections*

EXPERT LEVEL (Questions 101-150)

Question 101

A product team wants to preserve a beautiful, highly polished dashboard as their main competitive advantage. According to the thesis, why is this strategy increasingly fragile in the agentic era?

A. Because dashboards are expensive to maintain

B. Because agents do not perceive the UI as the primary layer; they interact with capabilities, system of record, and governance instead

C. Because users prefer voice interfaces

D. Because dashboards cannot display enough data

Answer: B Correct

Explanation: The thesis argues that the agent has no eyes for your interface, so a sticky UI is no longer a durable moat. Competitive advantage shifts to the layers agents actually use: capabilities, system of record, and governance.

Reference: Agentic-AI.pdf page 41, Agentic-AI.pptx slide 41

Question 102

A company wants to deploy an AI support worker that can operate continuously, but the leadership is concerned about hallucinations and off-policy actions. Which design choice best aligns with the Agent Factory approach?

A. Give the AI broad autonomy without a system of record

B. Connect the worker to the company's official memory, set guardrails, and test on historical tickets before production

C. Use browser chat only so the AI can't make changes

D. Skip testing because the AI is already intelligent

Answer: B Correct

Explanation: The correct approach combines system-of-record grounding, explicit guardrails, and staged validation. The Ecomm Mart example shows this flow: write specs, connect records, give skills, set guardrails, then test before deployment.

Reference: Agentic-AI.pdf pages 72-73, Agentic-AI.pptx slides 72-73

Question 103

Why does the thesis call the System of Record the "prize" when SaaS unbundles?

A. Because it is the most visually appealing part of the product

B. Because authoritative data becomes the most strategic layer once the UI dissolves and capabilities become callable functions

C. Because it is the easiest layer to monetize

D. Because it contains marketing analytics

Answer: B Correct

Explanation: When the workflow UI dies and capabilities become function calls, the authoritative data layer remains the most valuable asset. Agents need verified facts to act, so control of the System of Record becomes strategically dominant.

Reference: Agentic-AI.pdf page 20, Agentic-AI.pptx slide 20

Question 104

A manager wants to assign a persistent delegate to every employee, but also wants those delegates to be built and configured using developer-grade tools. Which recursive relationship best matches the thesis?

A. Human agents configure general agents, who then delegate to personal agents

B. General agents build personal agents, and personal agents then dispatch general agents and AI workers at runtime

C. Personal agents replace all general agents

D. Human managers should directly control the workforce without any delegates

Answer: B Correct

Explanation: The thesis explicitly describes a recursive architecture: you build and configure personal agents with general agents like Claude Code/OpenCode, and then the personal agent acts as chief of staff, dispatching general agents and AI workers at runtime.

Reference: Agentic-AI.pdf page 28, Agentic-AI.pptx slide 28

Question 105

A team is comparing two rollout strategies: (1) use AI to speed up current workflows, or (2) build a new AI worker that replaces a recurring process. Which statement best captures the strategic difference?

A. They are identical because both use AI

B. The first is Mode 1 and ends with the session; the second is Mode 2 and creates a durable worker

C. The second is always cheaper

D. The first is only for non-technical teams

Answer: B Correct

Explanation: Mode 1 is problem-solving in a session, while Mode 2 is manufacturing an AI worker that persists and keeps working. The strategic shift is from temporary assistance to durable capability creation.

Reference: Agentic-AI.pdf pages 47, 55, Agentic-AI.pptx slides 46-47, 55

Question 106

A CTO asks why the 10-80-10 model is framed as a "promotion" rather than a reduction in human role. What is the strongest interpretation?

A. Humans do less and therefore have lower responsibility

B. Humans move from execution to direction, verification, and ownership, which are higher-order responsibilities

C. Humans only do documentation now

D. Humans are removed from the loop entirely

Answer: B Correct

Explanation: The model repositions humans from doing the repetitive 80% to directing, reviewing, and owning outcomes. That is a promotion in responsibility and leverage, not a demotion or removal.

Reference: Agentic-AI.pdf page 52, Agentic-AI.pptx slide 52

Question 107

A company wants to add AI workers but is tempted to let them self-organize without management because "they are smart enough." Which invariant directly rejects this idea?

A. Every worker uses the right engine

B. Workforce needs a management layer

C. Company runs on a nervous system

D. Human is always in charge

Answer: B Correct

Explanation: Rule #3 states that a workforce needs a management layer to hire, assign tasks, control budgets, and fire underperformers. Without that layer, the system lacks governance and operational discipline.

***Reference:** Agentic-AI.pdf page 69, Agentic-AI.pptx slide 69*

Question 108

A policy team is deciding whether to let an AI worker access a broad dataset directly or only a curated source of truth. Which option best matches the source's reliability model?

A. Broad direct access is best because it maximizes flexibility

B. Curated source of truth is best because the system of record anchors the worker in verified facts

C. Access doesn't matter because AI can infer missing information

D. The worker should operate without data to avoid bias

Answer: B Correct

Explanation: The source repeatedly emphasizes the System of Record as official memory and single source of truth. Curated access reduces fabrication and keeps the agent's actions grounded in authoritative data.

***Reference:** Agentic-AI.pdf pages 70, 77, Agentic-AI.pptx slides 70, 77*

Question 109

Why is the "OS becomes plumbing" claim important to builders?

A. Because operating systems will stop working

B. Because the old operating surface is no longer where users or agents create value; strategic focus moves above the OS layer

C. Because apps will be free

D. Because cloud infrastructure will disappear

Answer: B Correct

Explanation: The thesis does not say the OS vanishes; it becomes invisible infrastructure. Builders should focus on the agent layer, callable capabilities, and governance rather than OS-centric user interaction.

***Reference:** Agentic-AI.pdf page 24, Agentic-AI.pptx slide 24*

Question 110

A team wants to deploy an AI worker for customer support but only has a limited monthly budget. Which element of the spec most directly enforces financial discipline?

A. The title of the worker

B. The budget constraint written into the spec

C. The worker's name

D. The user interface theme

Answer: B Correct

Explanation: In the Ecomm Mart example, the budget is defined in the specification as a hard guardrail. This ensures the AI worker is operating within explicit economic boundaries, not merely producing output.

***Reference:** Agentic-AI.pdf page 72, Agentic-AI.pptx slide 72*

Question 111

A platform team argues that the 66% OSWorld success rate means agents are "good enough" for all critical work. Which objection from the thesis directly counters this conclusion?

A. Cost

B. Reliability

C. The hybrid

D. Trust & governance

Answer: D Correct

Explanation: The trust & governance objection states that an agent with broad access can be misled, and the winning platform must solve governed capability. High success rates do not eliminate the need for guardrails and governance.

***Reference:** Agentic-AI.pdf page 38, Agentic-AI.pptx slide 38*

Question 112

A founder wants one assistant that can manage email, budgeting, hiring, and technical work. Which concept in the thesis most directly supports this coordination layer?

A. General agent

B. Personal AI delegate

C. Predictive AI filter

D. Workflow UI

Answer: B Correct

Explanation: The personal AI delegate is designed to be the coordination layer that knows your preferences, carries your authority, and manages the workforce. It is the chief of staff that routes work to specialist agents.

***Reference:** Agentic-AI.pdf pages 62-64, Agentic-AI.pptx slides 62-64*

Question 113

Why does the thesis distinguish "capabilities" from "workflow UI" when discussing SaaS unbundling?

A. To show that UI is more important than functions

B. Because capabilities survive as callable functions while workflow UIs die as unnecessary human surfaces

C. Because both are equally valuable

D. Because workflow UI becomes the data store

Answer: B Correct

Explanation: The workflow UI existed to let humans operate capabilities. Once agents operate those capabilities directly, the UI becomes redundant while the capabilities remain valuable as callable tools.

***Reference:** Agentic-AI.pdf page 20, Agentic-AI.pptx slide 20*

Question 114

A company has the option to use an expensive, highly reliable engine for every task or a cheaper engine for routine work. Which rule best guides the decision?

A. Rule #1: Human is always in charge

B. Rule #4: Each worker uses the right engine

C. Rule #6: Workforce can grow under rules

D. Rule #7: Company runs on a nervous system

Answer: B Correct

Explanation: Rule #4 explicitly says important work should use a reliable engine and simple work should use a cheap engine. This avoids overprovisioning and aligns compute quality with task criticality.

***Reference:** Agentic-AI.pdf page 69, Agentic-AI.pptx slide 69*

Question 115

Why does the thesis describe "the human as the source of intent" as something that survives?

A. Because AI cannot type

B. Because humans still define what to want and judge whether the outcome is good

C. Because humans must always approve every keystroke

D. Because intent is stored in the UI

Answer: B Correct

Explanation: The human role does not disappear; it shifts to intent and judgment. Humans decide what should be done and whether the result is acceptable, while agents perform the operational work.

***Reference:** Agentic-AI.pdf page 40, Agentic-AI.pptx slide 40*

Question 116

A company wants to move from a chatbot that answers questions to a worker that completes a monthly marketing campaign. Which capability is most essential for this transition?

A. Better typography

B. Planning, tool use, self-correction, and goal completion

C. More chat windows

D. A different color scheme

Answer: B Correct

Explanation: A chatbot only answers; an agentic worker must plan the campaign, use tools to execute it, correct mistakes, and continue until the campaign outcome is achieved. Those four abilities are the core bridge to worker behavior.

***Reference:** Agentic-AI.pdf page 13, Agentic-AI.pptx slide 13*

Question 117

What is the most accurate description of the relationship between the book and TutorClaw?

A. TutorClaw replaces the book entirely

B. TutorClaw is a tutoring interface that teaches from the book's canonical knowledge

C. TutorClaw is a separate product with unrelated content

D. TutorClaw is only for installation help

Answer: B Correct

Explanation: TutorClaw is one of the delivery channels that teaches from the Agent Factory book itself. The book is canonical; TutorClaw is the conversational learning surface that delivers that knowledge.

***Reference:** Agent Factory website delivery channels section*

Question 118

Which scenario best reflects Mode 2 rather than Mode 1?

A. Asking an AI to summarize a document during one session

B. Building a customer service worker that continues to answer routine tickets every day

C. Requesting a one-time email draft

D. Asking for a spreadsheet formula

Answer: B Correct

Explanation: Mode 2 is about manufacturing durable AI workers that continue operating after the session ends. A daily customer service worker fits this model, while one-time requests are Mode 1.

***Reference:** Agentic-AI.pdf page 47, Agentic-AI.pptx slide 47*

Question 119

What does the "agent layer becomes the floor" metaphor imply for system design?

A. The OS must be deleted

B. The agent layer becomes the default operating surface, with the old interface beneath it as plumbing

C. The browser becomes the main system

D. The app store becomes unnecessary

Answer: B Correct

Explanation: The metaphor means the agent layer is now the surface on which users stand and work. The OS and traditional interfaces remain underneath as infrastructure, not the main place where work is done.

Reference: Agentic-AI.pdf page 24, Agentic-AI.pptx slide 24

Question 120

A leadership team asks why "human oversight" is still necessary if AI can do 80% of the work. Which answer best fits the thesis?

A. Humans are required by policy only

B. Humans provide intent, verification, and ownership; those functions cannot be delegated away

C. Humans are slower but cheaper

D. Humans maintain the UI

Answer: B Correct

Explanation: The 10-80-10 rhythm preserves the human functions that matter most: knowing what to want, checking quality, and being accountable for outcomes. AI does the middle work, but humans retain responsibility.

Reference: Agentic-AI.pdf page 51, Agentic-AI.pptx slide 51

Question 121

A team wants to choose between building a personal agent first or a workforce of general agents first. What does the thesis imply?

A. Workforce first, because personal agents are unnecessary

B. Build the delegate with general agents, then let it orchestrate the workforce

C. Start with a full AI company immediately

D. Neither should be built

Answer: B Correct

Explanation: The thesis explicitly says to build the delegate first using general agents. Once configured, the personal agent orchestrates general agents and AI workers at runtime. That recursive layering is foundational.

Reference: Agentic-AI.pdf page 28, Agentic-AI.pptx slide 28

Question 122

Why does the thesis emphasize "verified patterns" in the Skillpack?

A. To make the skills look official

B. To ensure reusable agent behavior is grounded in proven specifications rather than ad hoc prompting

C. To increase file size

D. To make the skillpack exclusive

Answer: B Correct

Explanation: The Skillpack is built around verified patterns and SKILL.md templates so builders can reuse reliable structures. This is consistent with the broader thesis: repeatable, spec-driven methods outperform ad hoc prompting.

Reference: Agent Factory website Skillpack section

Question 123

What makes the "company runs on a nervous system" rule especially powerful at scale?

A. It eliminates the need for humans

B. It automates signal routing and recovery, reducing coordination overhead as the workforce grows

C. It makes the company feel organic

D. It increases brand identity

Answer: B Correct

Explanation: As AI workforce size grows, manually passing tasks becomes a bottleneck. The nervous-system model allows automatic signal transfer and recovery, enabling scale without proportional coordination burden.

Reference: Agentic-AI.pdf page 70, Agentic-AI.pptx slide 70

Question 124

Why is "connect to company records" a foundational step before giving skills in the Ecomm Mart example?

A. Skills are unnecessary if records exist

B. The worker must have real data access before it can apply skills correctly

C. It improves the UI

D. It reduces the budget

Answer: B Correct

Explanation: A worker without access to the company's records cannot act on real customer context. The sequence matters: connect to authoritative data first, then teach behavior/skills so the worker can apply them to actual cases.

Reference: Agentic-AI.pdf page 72, Agentic-AI.pptx slide 72

Question 125

Which statement best captures the thesis's view on what dies and what survives in the agentic era?

A. The OS and PC die completely; apps survive unchanged

B. The app, shell, SaaS destination, and human-as-operator die; the OS, capabilities, human intent, and judgment survive

C. Everything survives equally

D. Only cloud services die

Answer: B Correct

Explanation: The thesis is explicit: what dies is the app as the unit of human work, the graphical shell, SaaS as destination, and the human as operator. What survives is the OS as plumbing, capabilities as callable tools, and the human as intent/judgment source.

Reference: Agentic-AI.pdf page 40, Agentic-AI.pptx slide 40

Question 126

A bank wants to deploy an AI loan-processing worker. Which combination best reflects the thesis's requirements for a safe production rollout?

A. Broad access, no guardrails, and live rollout to maximize speed

B. Spec-first design, system-of-record access, engine selection by task criticality, and historical testing before production

C. Browser chat only, no persistent memory, and manual copy-paste into forms

D. A flashy UI and a larger prompt window

Answer: B Correct

Explanation: Safe rollout in the Agent Factory model is spec-driven, grounded in authoritative records, bounded by engine choice and guardrails, and validated on historical cases first. The goal is reliable work under human-set limits, not uncontrolled autonomy.

Reference: Agentic-AI.pdf pages 72-73, Agentic-AI.pptx slides 72-73

Question 127

A founder says, "I'll wait until agents are 100% reliable before using them for anything important." Which thesis-based rebuttal is strongest?

A. Reliability will never improve

B. The hybrid path is transitional; bounded and recoverable tasks are the correct starting point while reliability matures

C. Agents are already perfect

D. Human review is unnecessary

Answer: B Correct

Explanation: The thesis explicitly acknowledges reliability gaps (~66% average success) and recommends bounded, recoverable tasks first. The hybrid model is a transition, not a dead end, allowing practical adoption without pretending the system is flawless.

Reference: Agentic-AI.pdf pages 38, 41, Agentic-AI.pptx slides 38, 41

Question 128

A team wants to maximize the economic leverage of AI workers. According to the thesis, which outcome most directly represents success?

A. More dashboard views

B. More user logins

C. Scale intelligence, not headcount

D. Higher training-token usage

Answer: C Correct

Explanation: The economic objective is to replace linear headcount growth with AI worker capacity, producing larger output with fewer human staff additions. That is the meaning of "scale intelligence, not headcount."

Reference: Agentic-AI.pdf page 41, Agentic-AI.pptx slide 41

Question 129

Why does the thesis treat the personal delegate as "your chief of staff" rather than just another chatbot?

A. It produces better wording

B. It knows your preferences, carries authority, and coordinates a workforce, which makes it an organizational layer rather than a conversational toy

C. It runs faster

D. It uses more memory

Answer: B Correct

Explanation: A chief of staff doesn't merely answer questions; it interprets intent, delegates to specialists, and manages execution. That's exactly how the personal agent is positioned: persistent, authoritative, and orchestration-focused.

Reference: Agentic-AI.pdf pages 63-64, Agentic-AI.pptx slides 63-64

Question 130

A company wants to support both expensive high-stakes tasks and cheap routine tasks. What is the best engine strategy according to the invariants?

A. Use the same expensive engine for everything

B. Use the cheapest engine for everything

C. Match reliable engines to important work and cheap engines to simple work

D. Avoid engines and do everything manually

Answer: C Correct

Explanation: Rule #4 says each worker uses the right engine. Important work should use reliable engines; routine work can use cheaper engines. This preserves quality where it matters and efficiency where it doesn't.

***Reference:** Agentic-AI.pdf page 69, Agentic-AI.pptx slide 69*

Question 131

Why does the thesis elevate "judgment: what to want, and whether it's good" as a surviving human function?

A. Because AI cannot generate text

B. Because the core human role shifts to deciding goals and assessing outcomes, not operating the machine

C. Because judgment is needed for design only

D. Because all tasks still require manual approval

Answer: B Correct

Explanation: The human is no longer the operator; the human is the source of intent and judgment. That means deciding what should be pursued and evaluating whether the result is good remains a human responsibility even as agents do more execution.

***Reference:** Agentic-AI.pdf page 40, Agentic-AI.pptx slide 40*

Question 132

Which statement best explains why the book treats the Agent Factory as a "canonical source"?

A. It is the most popular book

B. It provides the authoritative methodology that other channels teach, build from, and specialize

C. It contains code examples only

D. It is meant to replace all other material

Answer: B Correct

Explanation: The book functions as the system of record for the ecosystem. TutorClaw teaches from it, the Skillpack operationalizes it, and derivative books specialize it. That makes it the canonical source of truth for the methodology.

***Reference:** Agent Factory website delivery structure*

Question 133

A team wants to preserve some human-in-the-loop control while still benefiting from AI autonomy. Which architecture from the thesis best supports that compromise?

A. Personal agent at the edge with AI workforce beneath it

B. No delegate, direct tasking of all workers

C. Only browser chat agents

D. A single prompt with no persistence

Answer: A Correct

Explanation: The two-layer model is designed for this compromise: your personal delegate sits at the edge, and the AI workforce operates underneath. The delegate provides human-aligned oversight while the workforce provides autonomous execution.

***Reference:** Agentic-AI.pdf page 64, Agentic-AI.pptx slide 64*

Question 134

A product team argues their app's "beautiful, sticky UI" is enough to defend against agents. What is the strongest thesis-based counterargument?

A. Beautiful UI is always valuable

B. Agents bypass the UI and interact with capabilities and records directly, so UI stickiness is no longer a strong moat

C. UI is only for mobile apps

D. UI will disappear immediately

Answer: B Correct

Explanation: The thesis says the agent has no eyes for your interface. Agents operate through tool calls and records rather than human navigation. Therefore, interface stickiness does not defend the product the way it once did.

Reference: Agentic-AI.pdf page 41, Agentic-AI.pptx slide 41

Question 135

A team wants to design a worker that can take action, observe results, and retry. Which agentic feature is most directly enabling this loop?

A. Self-correction

B. Branding

C. Faster chat responses

D. The UI theme

Answer: A Correct

Explanation: Self-correction is the capability that allows the agent to notice mistakes and try a different approach. That retry loop is essential for goal completion in real work environments.

Reference: Agentic-AI.pdf page 13, Agentic-AI.pptx slide 13

Question 136

Why is the "agent layer becomes the floor" concept more than just a UI redesign?

A. Because it changes the icon layout

B. Because it replaces the operating model of computing, moving humans from manual operation to intent delegation

C. Because it only affects desktop shortcuts

D. Because it reduces window clutter

Answer: B Correct

Explanation: The thesis frames this as a structural reversal: the OS sinks, the shell becomes legacy, and the agent layer becomes the default operational plane. This is an operating-model shift, not a skin-deep redesign.

Reference: Agentic-AI.pdf pages 24, 31, Agentic-AI.pptx slides 24, 31

Question 137

A company wants an AI system that keeps learning the team's preferences over time. Which concept best fits?

A. General agent

B. Personal agent with long-term memory

C. Predictive filter

D. Workflow UI

Answer: B Correct

Explanation: Personal agents are described as always-on, persistent, and long-term memory systems that "know you." They are built to remember preferences and adapt to the human they serve over time.

Reference: Agentic-AI.pdf page 26, Agentic-AI.pptx slide 26

Question 138

A team is deciding whether to create a company-wide AI workforce or just improve one employee's productivity. Which distinction should guide their strategy?

A. Browser chat vs code editor

B. Mode 1 vs Mode 2

C. UI vs API

D. Cloud vs on-device

Answer: B Correct

Explanation: Mode 1 helps an individual do a thing right now; Mode 2 builds a worker that keeps doing the job. The company strategy determines whether the initiative stops at productivity or becomes durable workforce capacity.

***Reference:** Agentic-AI.pdf pages 47, 55, Agentic-AI.pptx slides 46-47, 55*

Question 139

What is the most defensible interpretation of "Tools change. Rules stay." in the context of platform migrations?

A. Tools never matter

B. The implementation may move from Claude Code to OpenCode or future systems, but the architectural invariants remain constant

C. Rules are less important than tools

D. New tools invalidate old principles

Answer: B Correct

Explanation: The thesis insists that structural principles are invariant even as tool ecosystems evolve. You may switch engines, languages, or platforms, but the rules for governance, records, delegates, and management layers still apply.

***Reference:** Agentic-AI.pptx slides 68-70*

Question 140

A team wants the fastest possible route from goal to action on a local machine. Which hardware shift does the thesis identify as enabling this?

A. More browser tabs

B. RTX Spark bringing petaflop-class agent compute on-device

C. Larger monitors

D. Faster keyboards

Answer: B Correct

Explanation: RTX Spark is presented as the device-level compute breakthrough that allows frontier agents to run locally with no cloud round-trip. This makes agentic action immediate, private, and always-available on the machine.

***Reference:** Agentic-AI.pdf page 30, Agentic-AI.pptx slide 30*

Question 141

Why does the thesis say the AI-Native company runs on a "nervous system" rather than a command chain?

A. Because management is automated

B. Because tasks and signals route automatically between workers and the system can recover from breaks without manual routing

C. Because employees are remote

D. Because the company is biological

Answer: B Correct

Explanation: A nervous system transmits signals reflexively and adapts quickly. The thesis uses that metaphor to show how AI work can be routed and recovered automatically, reducing human handoffs and bottlenecks.

***Reference:** Agentic-AI.pdf page 70, Agentic-AI.pptx slide 70*

Question 142

A stakeholder claims the system-of-record requirement is just an implementation detail. What is the strongest rebuttal from the thesis?

A. It is only useful for analytics

B. It is the core grounding mechanism that prevents AI from making things up and therefore is strategic, not incidental

C. It is optional for advanced models

D. It is only for finance systems

Answer: B Correct

Explanation: The thesis positions the System of Record as the authoritative memory and prize layer. Without it, agents have no grounding and operate probabilistically; with it, they can perform reliable work on real data.

***Reference:** Agent Factory website system of record section, Agentic-AI.pdf page 70*

Question 143

What is the best reason the thesis stresses "digital, bounded, recoverable work" as the first stage of transformation?

A. Such work is always boring

B. It is the most suitable environment for current agent reliability and governance constraints

C. It requires no AI skills

D. It reduces the need for management

Answer: B Correct

Explanation: Bounded, recoverable tasks are the correct first targets because current agents are not universally reliable. Limited scope and reversible actions make it possible to adopt agents safely while learning how to govern them.

***Reference:** Agentic-AI.pdf page 19, page 38, Agentic-AI.pptx slides 19, 38*

Question 144

A team wants to understand whether the Agent Factory is mostly about technology, business, or operations. Which answer best matches the sources?

A. Only technology

B. Only business

C. It combines technical architecture, workflow operations, and economic transformation into one system

D. Only operations

Answer: C Correct

Explanation: The sources present Agent Factory as a complete system: it includes specs, tools, workflow patterns, governance, and economic models like Digital FTEs and AI-Native companies. It is simultaneously technical, operational, and economic.

***Reference:** Agent Factory website overview, Agentic-AI.pdf pages 76-78*

Question 145

What is the thesis's position on "selling tools" versus "selling results"?

A. Tools are always the better business model

B. AI-native organizations should sell results, because AI workers can own the middle of the workflow end-to-end

C. Results are less important than access

D. Neither model matters

Answer: B Correct

Explanation: The shift from tool to worker changes the business model. Instead of selling seat access or software usage, an AI-native company sells finished outcomes because its AI workers can own the workflow from intent to completion.

***Reference:** Agentic-AI.pdf page 55, page 59, Agentic-AI.pptx slides 55, 59*

Question 146

A company wants an AI worker that can classify tickets, draft replies, and escalate hard cases while respecting rules. Which combination of thesis elements is essential?

A. UI polish and faster internet

B. Clear instructions, company records, skills, guardrails, and testing

C. A larger prompt and more tokens

D. Multiple browser tabs

Answer: B Correct

Explanation: The Ecomm Mart example shows that reliable workers require the whole stack: clear specs, system-of-record access, skill definitions, guardrails, and pre-production testing. No single element is enough on its own.

***Reference:** Agentic-AI.pdf pages 72-73, Agentic-AI.pptx slides 72-73*

Question 147

Why does the thesis treat the 10-80-10 rhythm as more than a productivity hack?

A. Because it just reduces typing

B. Because it defines a durable operating model for human-AI collaboration with preserved accountability

C. Because it is a marketing slogan

D. Because it only applies to reports

Answer: B Correct

Explanation: The rhythm is an operating model: humans set direction, AI executes, humans verify and own the outcome. That's a stable way to structure collaboration, not merely a speed trick.

***Reference:** Agentic-AI.pdf pages 51-52, Agentic-AI.pptx slides 51-52*

Question 148

Which of the following best captures the role of the Skillpack in an AI-native organization?

A. A tool to replace the book

B. A guided partner for turning canonical patterns into practical builds in Claude Code or OpenCode

C. A monitoring dashboard

D. A training dataset

Answer: B Correct

Explanation: The Skillpack is the hands-on implementation layer: it uses verified patterns and templates to help builders create agents in the supported tools. It operationalizes the book's ideas rather than replacing them.

***Reference:** Agent Factory website delivery channels section*

Question 149

What most directly differentiates a "smart tool" from a "smart worker" in the thesis?

A. A smart worker has a better interface

B. A smart worker can take a goal, plan, act, use tools, correct itself, and finish without step-by-step human guidance

C. A smart worker is always more accurate

D. A smart tool is always cheaper

Answer: B Correct

Explanation: The move from tool to worker is the move from assistance to delegated execution. The worker handles the full job lifecycle, not just a narrow subtask, which is why the thesis frames it as a major shift.

***Reference:** Agentic-AI.pdf pages 53-55, Agentic-AI.pptx slides 53-55*

Question 150

Which statement best describes the thesis's overall economic and organizational conclusion?

A. AI is mostly a user interface upgrade

B. AI changes the economics of work by enabling organizations to manufacture digital labor and organize it into AI-native companies

C. AI is only useful for experimentation

D. AI only helps developers

Answer: B Correct

Explanation: The sources conclude that AI is not just a tool category; it changes work economics by manufacturing AI workers (Digital FTEs) and composing them into AI-native companies where humans set direction and verify outcomes.

***Reference:** Agentic-AI.pdf pages 76-78, Agentic-AI.pptx slides 76-78*

ULTRA LEVEL (Questions 151-200)

Question 151

A multinational enterprise wants to design an AI operating model where local teams can delegate work while headquarters retains control over policy, memory, and budgets. Which architecture from the thesis most directly supports this?

A. A single browser-based chatbot for all employees

B. An edge-layer personal delegate managing a layered AI workforce, grounded in a centralized system of record and governed by human-set rules

C. Multiple independent agents with no memory sharing

D. A manual ticketing queue with AI suggestions

Answer: B Correct

Explanation: The two-layer model is exactly this pattern: a personal delegate at the edge carries authority and preferences, while an AI workforce executes specialist tasks beneath it. Centralized records and rule-based governance preserve control across the organization.

***Reference:** Agentic-AI.pdf pages 64, 69-70, Agentic-AI.pptx slides 64, 69-70*

Question 152

A founder argues that because an AI worker can draft replies, the company no longer needs a system of record. Which response best reflects the thesis?

A. True, because the worker already has context

B. False, because without a system of record the worker loses grounding and can fabricate information or act on stale assumptions

C. True, as long as the prompt is long enough

D. True, if the worker uses a large model

Answer: B Correct

Explanation: The thesis makes the system of record non-negotiable because AI must read/write the company's official memory to avoid hallucination and stale context. Drafting capability alone does not create reliable operational behavior.

*Reference:** Agentic-AI.pdf pages 70, 77, Agentic-AI.pptx slides 70, 77*

Question 153

A platform team wants to know why the Agent Factory repeatedly stresses "design -> train -> deploy -> verify" instead of a one-shot build. What is the deepest reason?

A. It sounds more professional

B. It reflects a manufacturing model where repeatability, quality control, and bounded rollout are essential for producing reliable AI workers

C. It is required by the UI

D. It reduces prompt size

Answer: B Correct

Explanation: The garment-factory analogy is deliberate: you don't get production-grade output from a single accidental step. Manufacturing AI workers requires an explicit pipeline with quality gates and verification, not a one-time prompt.

*Reference:** Agentic-AI.pptx slide 58, Agentic-AI.pdf page 57*

Question 154

A security review team worries that an always-on delegate could become a high-risk point of failure. Which thesis-based control most directly mitigates that risk?

A. More UI training

B. Human-set guardrails and governed capability on what the delegate and workforce can access or do

C. Letting the delegate self-optimize without oversight

D. Using a bigger model

Answer: B Correct

Explanation: The thesis emphasizes governed capability: what an agent may touch and what may leave the machine is explicitly controlled. Human-set guardrails, budgets, and access rules keep delegated autonomy bounded and auditable.

*Reference:** Agentic-AI.pdf pages 30, 73, Agentic-AI.pptx slides 30, 73*

Question 155

Why does the thesis insist that "every action traces back to a human" even in a highly autonomous AI company?

A. To slow automation

B. To preserve accountability, legitimacy, and governance over AI behavior

C. Because AI cannot store logs

D. To keep managers employed

Answer: B Correct

Explanation: The first invariant anchors every action to a human source of intent and authority. This keeps AI within accountable governance structures and prevents the system from drifting into goal-setting autonomy without oversight.

Reference: Agentic-AI.pdf page 69, Agentic-AI.pptx slide 69

Question 156

A large team is considering replacing their internal workflow app with a direct API layer for agents. Which thesis claim best explains why this is strategically sound?

A. APIs are easier to design

B. Workflow UI dies first; agents operate capabilities directly and do not need the human navigation layer

C. APIs are always cheaper than UI

D. Apps are inherently obsolete

Answer: B Correct

Explanation: The thesis says workflow UI exists for humans, not agents. Once agents become the operators, they go straight to capabilities via API/MCP, bypassing the screens entirely. The strategy is to serve the agent, not the clicker.

Reference: Agentic-AI.pdf page 20, Agentic-AI.pptx slide 20

Question 157

A company wants to deploy multiple specialist workers: support, finance, sales, and legal. According to the two-layer model, what should coordinate them?

A. Direct human command over every worker

B. A personal delegate that understands the human's intent and routes tasks to specialist workers

C. A single general agent with no memory

D. A separate dashboard for each worker

Answer: B Correct

Explanation: The personal delegate is the orchestration layer that manages the AI workforce. It carries authority, knows preferences, and coordinates specialist workers so the human does not need to micromanage each worker directly.

Reference: Agentic-AI.pdf pages 63-64, Agentic-AI.pptx slides 63-64

Question 158

If agents become the operating layer, what becomes the most important thing for a builder to optimize: UI polish, agent ability, or governance?

A. UI polish

B. Agent ability only

C. Governance plus the capabilities and records agents require

D. Marketing copy

Answer: C Correct

Explanation: The thesis says the moat moves to where agents live, what capabilities they call, and what governance they must obey. UI polish is no longer the primary defense; robust capabilities and governed records matter more.

Reference: Agentic-AI.pdf page 41, Agentic-AI.pptx slide 41

Question 159

A team asks whether the OS still matters if the agent layer becomes dominant. What is the best thesis-based answer?

A. No, the OS disappears

B. Yes, but as invisible plumbing rather than the user-facing operating surface

C. No, all hardware disappears

D. Yes, because the OS becomes the app store

Answer: B Correct

Explanation: The thesis does not eliminate the OS; it demotes it. Windows, macOS, and Linux remain as foundational infrastructure like TCP/IP or BIOS, while the agent layer becomes the surface humans actually interact with.

***Reference:** Agentic-AI.pdf page 24, Agentic-AI.pptx slide 24*

Question 160

Why does the thesis elevate "trust & governance" as a core objection rather than a minor implementation detail?

A. Because governance is an afterthought

B. Because an agent with broad access can be misled, so the winning platform must govern capability rather than merely maximize FLOPs

C. Because governance slows all software

D. Because users prefer audits

Answer: B Correct

Explanation: The thesis explicitly identifies governance as a major limit on deployment. Capability without control creates risk; therefore, the platform that wins will solve governed capability, not just deliver the most compute.

***Reference:** Agentic-AI.pdf page 38, Agentic-AI.pptx slide 38*

Question 161

A builder says, "I'll just make a stronger prompt and that will solve the AI worker problem." Which thesis insight best refutes that claim?

A. Prompts are never useful

B. Reliable AI workers require specs, records, guardrails, skills, engine selection, and verification-not just prompts

C. Prompts only work in chat mode

D. Bigger prompts are always worse

Answer: B Correct

Explanation: The sources show that agentic systems are built through architecture, not prompt length. The worker needs a defined spec, system-of-record grounding, skills, guardrails, and testing to be reliable in production.

***Reference:** Agent Factory website architecture, Agentic-AI.pdf pages 72-73*

Question 162

What does the thesis mean by saying the personal delegate "carries your authority"?

A. The delegate can override your decisions

B. The delegate is empowered to act on your behalf within the intent and guardrails you define

C. The delegate owns the company

D. The delegate can ignore budgets

Answer: B Correct

Explanation: Carrying your authority means the delegate is not merely advisory; it can route work, coordinate workers, and act within the scope you set. That authority is bounded by human intent, budgets, and governance.

***Reference:** Agentic-AI.pdf pages 63-64, Agentic-AI.pptx slides 63-64*

Question 163

A company wants its AI support worker to be able to escalate hard cases. In the thesis, why is escalation part of the skill definition rather than an afterthought?

A. Escalation is rare

B. Skills include not only what the worker does, but also when it should hand off to humans or other systems

C. Escalation is only for legal teams

D. It makes the worker look smarter

Answer: B Correct

Explanation: A skill is more than a task checklist; it defines operational behavior, including boundaries and handoff points. The Ecomm Mart example explicitly includes when to escalate, because real workers must know their limits.

***Reference:** Agentic-AI.pdf page 72, Agentic-AI.pptx slide 72*

Question 164

A company has only a small budget and wants to move into agentic systems. Which thesis path is most realistic?

A. Start with a fully autonomous, unlimited agentic enterprise

B. Begin with bounded, recoverable tasks and a narrow worker spec, then expand under rules

C. Skip governance until scale arrives

D. Buy more SaaS seats first

Answer: B Correct

Explanation: The thesis recommends starting with bounded tasks and clear limits because current reliability is imperfect and governance matters. Once the system proves itself, the workforce can grow under rules.

***Reference:** Agentic-AI.pdf pages 38, 70, Agentic-AI.pptx slides 38, 70*

Question 165

How should a builder interpret the statement "the platform that wins solves governed capability, not the most FLOPs"?

A. Compute never matters

B. Raw compute is insufficient without controls over what the agent can access, change, and emit

C. Only local compute matters

D. FLOPs are the only thing that matters

Answer: B Correct

Explanation: The thesis is clear that compute power alone does not make a useful agent platform. Success depends on combining compute with policy, access control, and operational guardrails that govern behavior.

***Reference:** Agentic-AI.pdf page 38, Agentic-AI.pptx slide 38*

Question 166

Which of the following is the best description of the "general agents" category in the presentation?

A. Persistent personal assistants that know you

B. Tools oriented toward the work, summoned per job, and used inside files, commands, workflows, and projects

C. Browsers with chat windows

D. Model training pipelines

Answer: B Correct

Explanation: General agents are task-oriented worker tools like Claude Code and OpenCode. They are sessional, not persistent, and they operate within real environments such as codebases, command shells, and project workflows.

***Reference:** Agentic-AI.pdf page 26, page 15, Agentic-AI.pptx slides 15, 26*

Question 167

A team wants to build a durable AI support function. Which sequence best reflects the manufacturing mindset in the sources?

A. Prompt -> hope -> deploy -> fix later

B. Design -> train -> deploy -> verify

C. Chat -> copy -> paste -> send

D. Build -> launch -> ignore -> scale

Answer: B Correct

Explanation: The factory metaphor explicitly uses a repeatable pipeline: design the worker, train it with skills, deploy it under guardrails, and verify the results. That is the disciplined route to reliable AI labor.

Reference: Agentic-AI.pptx slide 58

Question 168

What does the thesis imply about the future of app-centric productivity tools?

A. They will become more central than ever

B. Their UI wrapper becomes less important as agents call capabilities directly, though the underlying functions still matter

C. They will instantly disappear

D. They will become free

Answer: B Correct

Explanation: App-centric tools lose the user as the direct operator because agents call their functions directly. The value shifts away from the wrapper and toward the callable capability and the data it accesses.

Reference: Agentic-AI.pdf pages 19-20, Agentic-AI.pptx slides 19-20

Question 169

Why is the "system of record" also a governance concept, not just a database concept?

A. Because databases are always governed

B. Because it defines the authoritative memory that workers must obey and that actions must trace back to

C. Because it improves search speed

D. Because it stores only logs

Answer: B Correct

Explanation: The system of record is not just storage; it is the official memory and source of truth that governs what an AI worker may know and do. It anchors both factual correctness and accountability.

Reference: Agent Factory website system of record section, Agentic-AI.pdf page 70

Question 170

A leadership team wants to know when an AI worker should use a cheaper engine. Which principle applies?

A. Rule #1

B. Rule #2

C. Rule #4

D. Rule #7

Answer: C Correct

Explanation: Rule #4 says each worker should use the right engine - important work gets a reliable engine, simple work gets a cheaper one. This is the core cost-performance matching principle.

Reference: Agentic-AI.pdf page 69, Agentic-AI.pptx slide 69

Question 171

Why is the "AI-Native Company" example of 10 humans and 190 AI workers more than a staffing anecdote?

A. It is a marketing slogan

B. It demonstrates the company operating principle of selling results with AI-driven scale and human supervision

C. It is a hypothetical number only

D. It shows that humans are unnecessary

Answer: B Correct

Explanation: The example illustrates the organizational model: AI handles routine work at scale, humans supervise, and the company delivers outcomes efficiently. It is evidence of an operational and economic transformation, not just a headcount statistic.

Reference: Agentic-AI.pdf page 59, Agentic-AI.pptx slide 59*

Question 172

A company wants to provide AI workers with "skills." In thesis terms, what are skills most analogous to?

A. Marketing slogans

B. Training manuals / operational instructions for how the worker should behave

C. UI themes

D. Hardware drivers

Answer: B Correct

Explanation: The Ecomm Mart example explicitly compares skill definition to a training manual. Skills teach the worker how to perform tasks, when to escalate, and how to operate within defined boundaries.

Reference: Agentic-AI.pdf page 72, Agentic-AI.pptx slide 72*

Question 173

What is the most accurate interpretation of "the app dissolves into a function call"?

A. The app is deleted from the store

B. The visible human-facing interface becomes unnecessary while the underlying capability remains callable by agents

C. The app becomes open-source

D. The app becomes a spreadsheet

Answer: B Correct

Explanation: In the agentic era, the workflow UI is removed from the human path, and agents call capabilities directly. The product is no longer experienced as an app destination but as a callable capability.

Reference: Agentic-AI.pdf page 19-20, Agentic-AI.pptx slides 19-20*

Question 174

Why does the thesis repeatedly compare AI systems to factory systems rather than purely software systems?

A. To make them sound industrial

B. Because reliable AI labor requires repeatable process, quality control, specialization, and workflow composition

C. Because software is obsolete

D. To reduce development costs

Answer: B Correct

Explanation: A factory is a useful metaphor because AI workers need repeatable production processes, quality gates, and orchestration into larger systems. That is closer to manufacturing than to one-off software execution.

Reference: Agentic-AI.pptx slides 58-61*

Question 175

A chief information officer wants to know what survives after the agentic transformation if the old app model disappears. Which answer is most aligned with the thesis?

A. Only the user interface survives

B. The OS as plumbing, capabilities as tools, human intent, and human judgment survive

C. Nothing survives

D. Only cloud APIs survive

Answer: B Correct

Explanation: The thesis is explicit about survival: the OS remains as invisible infrastructure, capabilities become callable tools, and humans remain the source of intent and judgment. The shell, app-as-destination, and human-as-operator are what fade away.

***Reference:** Agentic-AI.pdf page 40, Agentic-AI.pptx slide 40*

Question 176

A platform architect argues that because agents can call APIs directly, the old product surface no longer matters at all. Which subtle thesis-based correction is most accurate?

A. Correct; product surfaces are irrelevant in every case

B. Incorrect; the surface matters less as a human workflow layer, but the callable capability and the system of record become the critical assets

C. Incorrect; the surface becomes the only thing that matters

D. Correct only for SaaS, not for agents

Answer: B Correct

Explanation: The thesis does not say all surface area disappears. It says the human-facing workflow UI dies as the primary work surface, while capabilities and authoritative data become strategic. The product may still exist, but its value center shifts downward in the stack.

***Reference:** Agentic-AI.pdf pages 19-20, 41, Agentic-AI.pptx slides 19-20, 41*

Question 177

A company wants to let AI workers expand on their own when demand grows, but only within policy boundaries set by leadership. Which rule exactly captures this behavior?

A. Rule #2: Every human gets a delegate

B. Rule #4: Each worker uses the right engine

C. Rule #6: Workforce can grow under rules

D. Rule #7: Company runs on a nervous system

Answer: C Correct

Explanation: Rule #6 explicitly states that when a gap appears, the system can hire a new worker automatically within human-set limits. That is the thesis's model for controlled, policy-bound workforce growth.

***Reference:** Agentic-AI.pdf page 70, Agentic-AI.pptx slide 70*

Question 178

A CEO wants a delegate that can manage sales, finance, legal, and support workers but also wants the delegate to remember the CEO's style for years. Which architecture best fits?

A. General agent only

B. Personal agent with persistent memory and authority over specialist workers

C. Predictive filter with no memory

D. Static dashboard with no AI

Answer: B Correct

Explanation: The personal agent is always-on, persistent, and long-term memory oriented. It is the delegate layer that understands preferences and manages specialist workers over time, unlike general agents that are summoned per job.

***Reference:** Agentic-AI.pdf page 26, 63-64, Agentic-AI.pptx slides 26, 63-64*

Question 179

Why does the thesis insist that an AI worker without a system of record is fundamentally unreliable, even if it has advanced model capabilities?

A. Because models cannot answer questions

B. Because capability without grounding produces probabilistic outputs that can drift, hallucinate, or act on stale context

C. Because the model will be too slow

D. Because advanced models require a UI

Answer: B Correct

Explanation: The system of record is the authoritative memory that keeps the worker grounded in real data. Without it, the model may still sound convincing, but operational reliability suffers because the worker lacks verified facts.

***Reference:** Agent Factory website system of record section, Agentic-AI.pdf pages 70, 77*

Question 180

An enterprise wants to move from "pilot AI" to "production AI workers." Which sequence best reflects the thesis's recommended progression?

A. Deploy immediately, then define rules later

B. Write specs, connect records, teach skills, set guardrails, test on historical data, then deploy

C. Build a chatbot, then assume it will self-improve

D. Choose a larger model first, then decide the workflow

Answer: B Correct

Explanation: This is the Ecomm Mart pattern generalized into a production method. Production-grade AI workers are manufactured through explicit stages that create bounded, verified capability before live deployment.

***Reference:** Agentic-AI.pdf pages 72-73, Agentic-AI.pptx slides 72-73*

Question 181

A team wonders whether the Agent Factory is primarily about "AI prompts" or "organizational design." Which answer is most consistent with the sources?

A. It is mostly about prompts

B. It is mostly about organizational design and operating model, with prompts as one component of a larger system

C. It is only about code generation

D. It is only about UX design

Answer: B Correct

Explanation: The thesis frames Agent Factory as a methodology for building AI workers and AI-native companies. That is an organizational and operational design problem first, with specs, guards, records, skills, and delegation structures-not just prompting technique.

***Reference:** Agent Factory website architecture, Agentic-AI.pdf pages 57-61*

Question 182

A company believes that because their AI worker can browse the web, it no longer needs internal records. What is the best rebuttal?

A. Web browsing is always better than internal records

B. External browsing can supplement context, but the authoritative source for actions is still the company's system of record

C. Internal records are never useful

D. Browsing and records are identical

Answer: B Correct

Explanation: The thesis does not deny external tools; it prioritizes authoritative internal memory for action. Browsing is an input, but the worker must still read/write the company's official record to stay aligned with real operations.

***Reference:** Agentic-AI.pdf pages 70, 77, Agentic-AI.pptx slides 70, 77*

Question 183

Why is the "system of record" a better fit than a chat transcript for long-term AI workforce memory?

A. Chat transcripts are never useful

B. The system of record is structured, authoritative, and writable, while transcripts are conversational and not designed as operational truth

C. Chat transcripts are too short

D. The system of record is only for compliance

Answer: B Correct

Explanation: The thesis treats the system of record as official memory, which must support reliable reads and writes. A chat transcript may preserve conversation, but it is not the company's authoritative operational source of truth.

***Reference:** Agent Factory website system of record section, Agentic-AI.pdf page 70*

Question 184

A security architect wants to know why "what may leave the machine" matters as much as what the agent can touch. What is the thesis-aligned answer?

A. It doesn't matter if the agent is accurate

B. Both input access and output exfiltration must be governed to keep the agent's capability bounded and safe

C. Output controls slow down the UI

D. Only cloud systems need output controls

Answer: B Correct

Explanation: Governing what the agent may touch is not enough; the thesis explicitly mentions governing what may leave the machine. That means both access and emission must be controlled to maintain privacy, safety, and policy compliance.

Reference: Agentic-AI.pdf page 30, Agentic-AI.pptx slide 30

Question 185

A team asks whether a beautiful interface can still be useful in the agentic era. Which answer is most nuanced and correct?

A. No, interfaces are worthless

B. Yes, for humans in transitional or oversight workflows, but not as the main competitive moat against agentic systems

C. Yes, because UI is always the business

D. No, because agents cannot render graphics

Answer: B Correct

Explanation: The thesis does not claim UIs vanish immediately. It says the human-facing workflow UI is no longer the primary moat because agents bypass it. Interfaces can still exist for humans during transition and oversight, but the strategic center moves elsewhere.

Reference: Agentic-AI.pdf page 41, 38, Agentic-AI.pptx slides 41, 38

Question 186

Why does the thesis frame the AI revolution as "the agent is the operating layer" rather than "AI is a feature"?

A. To make the headline more dramatic

B. Because agents mediate intent-to-action across the whole computing stack, changing how work is operated, not just adding a feature to an app

C. Because features are obsolete

D. Because operating systems are disappearing tomorrow

Answer: B Correct

Explanation: The operating-layer framing means agents are not peripheral features; they become the layer through which work is performed. This changes the architecture of computing, from human-driven app use to goal-driven delegation.

Reference: Agentic-AI.pdf pages 16, 24, 31, Agentic-AI.pptx slides 16, 24, 31

Question 187

A company wants to reduce cost by using the cheapest engine for all work. Which invariant explains why this can be a bad idea?

A. Rule #1

B. Rule #2

C. Rule #4

D. Rule #5

Answer: C Correct

Explanation: Rule #4 says each worker uses the right engine. Using the cheapest engine for everything ignores task criticality and can degrade important work. The thesis favors cost-performance matching, not universal minimization.

***Reference:** Agentic-AI.pdf page 69, Agentic-AI.pptx slide 69*

Question 188

A human leader wants all AI workers to be autonomous, but also wants any major decisions to trace back to them. Which invariant is being balanced?

A. Rule #1 and Rule #6

B. Rule #3 and Rule #7

C. Rule #4 and Rule #5

D. None; autonomy and traceability cannot coexist

Answer: A Correct

Explanation: Rule #1 keeps human authority at the top, while Rule #6 allows the workforce to grow under human-set limits. Together they enable controlled autonomy: AI can act and scale, but every action still traces back to a human.

***Reference:** Agentic-AI.pdf pages 69-70, Agentic-AI.pptx slides 69-70*

Question 189

Why is "human reviews the result" not merely a quality-control step but a structural part of the architecture?

A. Because AI is always wrong

B. Because verification is one of the human responsibilities that cannot be delegated away in the 10-80-10 model

C. Because customers demand it

D. Because reports need signatures

Answer: B Correct

Explanation: Verification is one of the three surviving human responsibilities: intent, verification, and ownership. That makes review a structural architectural role, not an optional audit after the fact.

***Reference:** Agentic-AI.pdf page 51, Agentic-AI.pptx slide 51*

Question 190

A systems team wants the AI workforce to operate around the clock. Which thesis point makes this feasible at scale?

A. More managers are hired

B. AI workers are always-on and persistent, while humans supervise exceptions

C. Only browser chat can work at night

D. The OS becomes faster

Answer: B Correct

Explanation: AI workers in the AI-native model handle routine tasks 24/7. Their persistence and always-on nature enable continuous operation, while humans handle the tricky cases and oversight.

***Reference:** Agentic-AI.pdf pages 59, 26, Agentic-AI.pptx slides 59, 26*

Question 191

What is the strongest reason the thesis says the app as the unit of human work dies?

A. Apps are too expensive

B. Because work is delegated to agents that invoke capabilities directly, eliminating the app as the required human operating surface

C. Because apps are slow to install

D. Because app stores are disappearing

Answer: B Correct

Explanation: The app's role was to provide a human-friendly surface for operating capabilities. Once agents operate them directly, the app is no longer the unit of work; the capability plus record becomes the useful abstraction.

Reference: Agentic-AI.pdf pages 19-20, 40, Agentic-AI.pptx slides 19-20, 40

Question 192

A team wants to turn a general agent into a specialized AI worker. What is the thesis's most important transformation step?

A. Changing the UI theme

B. Encoding the worker's spec, skills, guardrails, and record access so it can do one job reliably

C. Increasing chat memory only

D. Training a new base model

Answer: B Correct

Explanation: The Agent Factory transforms general-agent capability into Digital FTEs by wrapping it in job-specific specifications, skills, records, and guardrails. The worker is a manufacturing output, not just a chat session.

Reference: Agentic-AI.pdf pages 57-61, 72-73, Agentic-AI.pptx slides 57-61, 72-73

Question 193

Why is the "human as operator" one of the things that dies, but the "human as source of intent" survives?

A. Because humans are slower than agents

B. Because the operational middle layer is delegated to AI, while humans remain responsible for goals and judgment

C. Because humans no longer need to think

D. Because only machines can use software

Answer: B Correct

Explanation: The thesis separates execution from intention. AI becomes the operator; humans remain the origin of goals and the evaluator of outcomes. That preserves human agency while eliminating the need for manual machine operation.

Reference: Agentic-AI.pdf page 40, 42, Agentic-AI.pptx slides 40, 42

Question 194

A leadership team wants a measure of success for their AI-native transformation. Which outcome most closely matches the thesis?

A. More app downloads

B. More seats purchased

C. More results delivered with fewer humans doing repetitive work

D. More chat sessions per user

Answer: C Correct

Explanation: The thesis emphasizes selling results, scaling intelligence, and using AI workers to do the routine work. Success is measured by output and leverage, not by software consumption or seat counts.

Reference: Agentic-AI.pdf pages 41, 55, 59, Agentic-AI.pptx slides 41, 55, 59

Question 195

An operations team wants to ensure their AI workforce can recover when a task breaks or fails. Which invariant directly addresses this?

A. Rule #2

B. Rule #5

C. Rule #7

D. Rule #1

Answer: C Correct

Explanation: Rule #7 says the company runs on a nervous system and should recover on its own if something breaks. That reflexive, self-healing behavior is exactly what the operations team needs.

***Reference:** Agentic-AI.pdf page 70, Agentic-AI.pptx slide 70*

Question 196

A CTO is designing a role for the AI delegate and wonders whether it should focus on direct execution or orchestration. What does the thesis imply?

A. Direct execution only

B. Orchestration first: the delegate understands intent and dispatches workers; it is the management layer

C. Neither; just chat with users

D. It should replace the workforce entirely

Answer: B Correct

Explanation: The delegate is described as the chief of staff. Its primary job is orchestration - understanding the human's intent, coordinating workers, and returning results. It is not meant to be the lone executor of all tasks.

***Reference:** Agentic-AI.pdf pages 28, 63-64, Agentic-AI.pptx slides 28, 63-64*

Question 197

Why is the Agent Factory model compatible with both Claude Code and OpenCode?

A. They have identical user interfaces

B. Because the core methodology is expressed in portable specs and skills rather than platform-specific behavior

C. Because both are browser tools

D. Because both use the same cloud provider

Answer: B Correct

Explanation: The methodology is tool-agnostic: SKILL.md and specification patterns can move between tools without changing the core system. That portability protects the architecture from lock-in and platform churn.

***Reference:** Agent Factory website two-tool philosophy, technical architecture section*

Question 198

A policymaker is evaluating agent transactions and asks why capability layers and governance are emphasized alongside protocols. Which answer best fits the thesis?

A. Protocols alone solve everything

B. Protocols matter, but they must be embedded in governed agent systems and authoritative records to support real work

C. Protocols replace workforce management

D. Protocols only matter for payments

Answer: B Correct

Explanation: The broader thesis is about operating work, not just transferring data. Even when protocols exist, they must be embedded in a governed architecture with records, delegates, and rules so agents can act reliably.

***Reference:** Agentic-AI.pdf pages 69-70, Agentic-AI.pptx slides 69-70*

Question 199

Why does the thesis repeatedly use analogies like a microwave, a fishing net, a film director, and a garment factory?

A. To make the content longer

B. To map abstract agentic concepts to familiar operational models that reveal how goal-setting, process, and verification work

C. To avoid technical explanations

D. To make it easier for children only

Answer: B Correct

Explanation: The analogies are teaching devices that translate abstract agentic ideas into concrete mental models: decision-making, durable production, direction-setting, and manufacturing. They help readers internalize the architecture rather than memorizing terms.

***Reference:** Agentic-AI.pptx slides 4, 46, 50, 58*

Question 200

What is the most complete synthesis of the thesis's final message?

A. AI is a better chatbot

B. AI is software that learns and makes decisions, agentic AI completes jobs, and the Agent Factory organizes AI workers into AI-native companies where humans set direction and verify outcomes

C. AI is mostly about better prompts

D. AI only benefits developers

Answer: B Correct

Explanation: This is the thesis in one sentence: AI becomes agentic when it can plan, use tools, self-correct, and complete work; the Agent Factory turns that capability into Digital FTEs and AI-native companies; humans keep intent, judgment, and accountability.

***Reference:** Agentic-AI.pdf pages 76-78, Agentic-AI.pptx slides 76-78*