

CodeVerse Soban

AI Agent Factory — Level 1 Exam | MCQ Practice Paper
Topic: The AI Agent Factory — Core Concepts & Ecosystem

Total Questions: 30 | Type: Tricky MCQs with Answers & Explanations

Q1. What does the term 'Digital FTE' stand for, and what does it represent in the Agent Factory framework?

- A) Digital Functional Task Executor — a bot that answers customer queries
- B) Digital Full-Time Equivalent — an AI agent that performs one employee's worth of structured work inside an organisation
- C) Distributed File Transfer Engine — a cloud storage protocol for AI data
- D) Dynamic Fine-Tuning Entity — a model that continuously retrains on live data

✓ **Correct Answer: B) Digital Full-Time Equivalent — an AI agent that performs one employee's worth of structured work inside an organisation**

Explanation: FTE is the HR term for one full-time human employee's work capacity. A Digital FTE is its AI equivalent performing real structured work.

Q2. According to the Agent Factory, which of the following CORRECTLY describes an 'AI-Native Company'?

- A) A company that uses AI only for marketing and customer support
- B) A company where one founder runs a real business with many AI Workers instead of a large human staff
- C) A company that develops its own proprietary AI models from scratch
- D) A company that exclusively hires AI researchers and data scientists

✓ **Correct Answer: B) A company where one founder runs a real business with many AI Workers instead of a large human staff**

Explanation: An AI-Native Company = a founder plus a handful of people plus many AI Workers — not a large human staff.

Q3. In the Five-Layer Cake of AI, which layer is described as 'where AI stops being technology and starts becoming useful'?

- A) Infrastructure
- B) Models
- C) Chips
- D) Applications

✓ **Correct Answer: D) Applications**

Explanation: The Applications layer (Layer 5, the top) is where AI becomes useful — the focus of the entire Agent Factory book.

Q4. A Digital FTE is explicitly described as NOT being 'just a model with a prompt.' Which combination makes it a TRUE system?

- A) Hardware, software, training data, and an internet connection
- B) Structured Specifications, Domain Expertise, Engineering Architecture, and Human Oversight

- C) Cloud storage, a REST API, a database, and a front-end interface
- D) A fine-tuned model, a vector database, a chatbot UI, and an SLA agreement

✓ **Correct Answer: B) Structured Specifications, Domain Expertise, Engineering Architecture, and Human Oversight**

Explanation: The book's four critical elements: Structured Specifications, Domain Expertise, Engineering Architecture, Human Oversight.

Q5. The Agent Factory states that the 'next competitive gap between organisations' will come from which factor?

- A) Owning the largest GPU cluster
- B) Having exclusive access to the latest foundation model
- C) Turning intelligence into repeatable, governable execution at the workforce layer
- D) Spending the most on AI research and development

✓ **Correct Answer: C) Turning intelligence into repeatable, governable execution at the workforce layer**

Explanation: The book is explicit: competitive advantage comes from who can turn AI into repeatable execution, not who has the best model or biggest hardware.

Q6. Which statement about the terms 'Digital FTE,' 'Digital Worker,' and 'AI Worker' is TRUE according to the book?

- A) They refer to three distinct types of AI systems with different capabilities
- B) 'AI Worker' is the consumer-facing term; 'Digital FTE' is the technical term
- C) They are used interchangeably and all name the same thing: a role-based AI agent performing structured work under human oversight
- D) 'Digital Worker' refers to human remote workers; the other two refer to AI systems

✓ **Correct Answer: C) They are used interchangeably and all name the same thing: a role-based AI agent performing structured work under human oversight**

Explanation: The book explicitly states: 'Digital FTE, Digital Worker, and AI Worker are used interchangeably.'

Q7. Jensen Huang's argument about 'systems of record' in the Agent Era concludes that these systems become:

- A) Obsolete because AI agents replace them with dynamic real-time reasoning
- B) More essential, not less — agents need authoritative ground truth to read from, write to, and verify against
- C) Replaced by large language model weights which serve as the new memory
- D) Optional for small companies but mandatory for enterprises

✓ **Correct Answer: B) More essential, not less — agents need authoritative ground truth to read from, write to, and verify against**

Explanation: Huang argues agents reinforce the need for systems of record — without ground truth, agents hallucinate; with it, they execute.

Q8. TutorClaw operates on which platforms and what distinguishes it from a generic AI assistant?

- A) Runs only on desktop apps; distinguished by supporting 100+ programming languages

- B) Runs 24/7 on WhatsApp, Telegram, and web; teaches from the canonical book rather than the open internet
- C) Runs inside Claude Code only; distinguished by having access to live GitHub repositories
- D) Runs on iOS and Android only; distinguished by voice interaction capability

✓ **Correct Answer: B) Runs 24/7 on WhatsApp, Telegram, and web; teaches from the canonical book rather than the open internet**

Explanation: TutorClaw runs 24/7 on WhatsApp, Telegram, and web — teaching from the verified canonical book, not probabilistic internet generation.

Q9. The Agent Factory book states that the method is 'the constant' and the tool is 'the variable.' What does this mean in practice?

- A) You must use the same AI model for every project to maintain consistency
- B) Skills, specifications, and architectural patterns written for Claude Code work identically in OpenCode — the discipline outlives any specific tool
- C) The book's methodology changes with every new AI model release
- D) Only Claude Code can implement the Agent Factory method correctly

✓ **Correct Answer: B) Skills, specifications, and architectural patterns written for Claude Code work identically in OpenCode — the discipline outlives any specific tool**

Explanation: A SKILL.md written for Claude Code drops into OpenCode unchanged. The method is portable — it must outlive any specific vendor's tool.

Q10. Which of the following is the CORRECT order of the 10-80-10 Rule as defined in the Agent Factory?

- A) AI sets intent (10%) → Human executes (80%) → AI verifies (10%)
- B) Human training (10%) → Model fine-tuning (80%) → Human testing (10%)
- C) Human sets intent (10%) → AI executes (80%) → Human verifies with professional judgment (10%)
- D) Planning (10%) → Building (80%) → Deploying (10%)

✓ **Correct Answer: C) Human sets intent (10%) → AI executes (80%) → Human verifies with professional judgment (10%)**

Explanation: 10-80-10: humans set intent (first 10%), AI handles execution (80%), human professional judgment closes with verification (final 10%).

Q11. What does OpenCode offer that makes it suitable for 'daily coursework, learning, and experimentation'?

- A) It is faster than Claude Code because it uses less memory
- B) It connects to dozens of model providers and allows switching between them based on economics, latency, and task complexity
- C) It comes pre-installed on all major operating systems without setup
- D) It only uses local models, ensuring complete data privacy

✓ **Correct Answer: B) It connects to dozens of model providers and allows switching between them based on economics, latency, and task complexity**

Explanation: OpenCode is model-agnostic — Claude, GPT, Gemini, DeepSeek, Qwen, local models via Ollama — giving flexibility, cost control, and vendor independence.

Q12. According to the Agent Factory, why does the book use TWO tools (Claude Code and OpenCode) instead of committing to one?

- A) Because neither tool alone is capable of building production-grade agents
- B) To increase revenue by partnering with two companies simultaneously
- C) Because binding the discipline to a single vendor inherits risks (pricing, access, strategy) and silently excludes readers with different constraints
- D) Because Claude Code handles design tasks and OpenCode handles coding tasks exclusively

✓ **Correct Answer: C) Because binding the discipline to a single vendor inherits risks (pricing, access, strategy) and silently excludes readers with different constraints**

Explanation: Two tools, one discipline — the design ensures portability and avoids vendor lock-in, pricing risk, and exclusion of readers with economic or regulatory constraints.

Q13. What is the PRIMARY role of 'Human Oversight' as one of the four critical elements of a Digital FTE?

- A) To retrain the AI model when it makes mistakes
- B) To write all the code that the AI then executes
- C) To provide feedback loops that maintain accountability and governance
- D) To replace the AI when tasks become too complex

✓ **Correct Answer: C) To provide feedback loops that maintain accountability and governance**

Explanation: Human Oversight = feedback loops that maintain accountability and governance — ensuring Digital FTEs remain auditable and trustworthy.

Q14. The Agent Factory describes three 'Eras of AI Tools.' Which matches the correct progression?

- A) Era 1: Applications → Era 2: Models → Era 3: Harnesses
- B) Era 1: Models as product → Era 2: Harnesses as product → Era 3: The discipline/methodology that runs across harnesses becomes the product
- C) Era 1: Chips → Era 2: Infrastructure → Era 3: Models
- D) Era 1: Chatbots → Era 2: Digital FTEs → Era 3: AI-Native Governments

✓ **Correct Answer: B) Era 1: Models as product → Era 2: Harnesses as product → Era 3: The discipline/methodology that runs across harnesses becomes the product**

Explanation: Era 1 = model as product. Era 2 = harness as product. Era 3 = the portable discipline/SKILL.md methodology that outlives both.

Q15. Sam Altman's prediction about 'ten-person billion-dollar companies' is used in the book to illustrate:

- A) Why only large tech companies will benefit from AI
- B) The economic leverage AI-Native Companies gain through Digital FTEs, making massive value creation possible with tiny teams
- C) The risk that AI will create monopolies and eliminate competition
- D) Why investors should fund AI hardware companies

✓ **Correct Answer: B) The economic leverage AI-Native Companies gain through Digital FTEs, making massive value creation possible with tiny teams**

Explanation: Altman's quote demonstrates the economic thesis: AI Workers replace large staffs, enabling small teams to build billion-dollar businesses.

Q16. A 'Derivative Book' in the Agent Factory ecosystem differs from the canonical source in which way?

- A) It uses a completely different methodology tailored to its specific audience
- B) It keeps the same methodology but is specialised by topic or audience — while inheriting updates from the canonical source
- C) It is written by different authors with no connection to the original
- D) It replaces the canonical source once published

✓ **Correct Answer: B) It keeps the same methodology but is specialised by topic or audience — while inheriting updates from the canonical source**

Explanation: Derivative books specialise along topic or audience axes but inherit the same vocabulary, architecture, and standards — and all updates — from the single canonical source.

Q17. In the Agent Factory's database recommendation, Postgres is preferred because it can hold which combination of data types in a single system?

- A) Images, audio, video, and binary files
- B) Relational data, documents, full-text search, and AI vectors
- C) Only structured SQL tables and stored procedures
- D) Only vector embeddings for AI similarity search

✓ **Correct Answer: B) Relational data, documents, full-text search, and AI vectors**

Explanation: 'One Postgres holding relational data, documents, full-text search, and AI vectors together' — consolidate by default, specialise deliberately.

Q18. The book states that 'Foundations plus Mode 1 is a serious skill set on its own.' Mode 1 is BEST described as:

- A) Building custom AI Workers that operate autonomously after you close the laptop
- B) Using AI to do your OWN work faster — you remain the doer, AI is the power tool
- C) Deploying AI agents on cloud-native infrastructure with Kubernetes
- D) Writing SKILL.md specifications for domain-specific knowledge encoding

✓ **Correct Answer: B) Using AI to do your OWN work faster — you remain the doer, AI is the power tool**

Explanation: Mode 1: you stay the doer, AI is the power tool for your own tasks. Mode 2 is where you build AI Workers that keep running independently.

Q19. Which reader type is described as 'The Translators' in the Agent Factory framework, and what is their primary function?

- A) Domain Experts — they translate expertise into SKILL.md files
- B) Enterprise Executives — they translate business strategy into AI budgets
- C) AI Product Managers & Solutions Architects — they decompose complex business processes into automated agent tasks
- D) Department Leaders — they translate internal playbooks into repeatable workflows

✓ **Correct Answer: C) AI Product Managers & Solutions Architects — they decompose complex business processes into automated agent tasks**

Explanation: Product Managers & Solutions Architects are 'The Translators' — mapping workflows into skills and defining boundaries between automated reasoning and human decisions.

Q20. Dario Amodei (Anthropic CEO) identified which THREE categories as most likely to produce the first single-person billion-dollar company?

- A) Gaming, social media, and e-commerce
- B) Developer tools, automated customer service, and proprietary trading
- C) Healthcare diagnostics, legal research, and financial auditing
- D) Education platforms, logistics automation, and renewable energy

✓ **Correct Answer: B) Developer tools, automated customer service, and proprietary trading**

Explanation: Amodei named developer tools, automated customer service, and proprietary trading as the most likely categories for the first solo-founder billion-dollar company.

Q21. The Agent Factory describes the solo telehealth founder who built hundreds of millions in first-year revenue. This example demonstrates which architectural shape?

- A) A large enterprise with AI as a supporting tool for existing human teams
- B) A canonical source the founder owns + AI agents doing team-level work + rented infrastructure (harnesses, models, messaging)
- C) A startup that trained its own proprietary AI model from scratch
- D) A consulting firm that uses AI only for report generation

✓ **Correct Answer: B) A canonical source the founder owns + AI agents doing team-level work + rented infrastructure (harnesses, models, messaging)**

Explanation: The telehealth founder exemplifies the AI-Native Company shape: owned canonical source + AI agents replacing teams + rented infrastructure.

Q22. 'Structured Specifications' as a critical element of Digital FTEs specifically refers to:

- A) The hardware specifications of the servers running the AI
- B) Clear definitions of what agents must do — written before AI builds to them (spec-driven development)
- C) Database schemas used to store agent conversation history
- D) Legal compliance documents required for enterprise AI deployment

✓ **Correct Answer: B) Clear definitions of what agents must do — written before AI builds to them (spec-driven development)**

Explanation: Structured Specifications = clear definitions of what agents must do. The Agent Factory method: write the spec first, then have AI build to it.

Q23. What makes the Agent Factory Book a 'system of record for agentic AI education' rather than just another AI textbook?

- A) It is the longest AI book ever published, with more than 2,000 pages
- B) It was written by the founders of Anthropic and OpenAI
- C) TutorClaw and building agents read from it as a verified canonical source — not the open internet — and human judgment verifies what is taught and built
- D) It includes access to a proprietary AI model available only to purchasers

✓ **Correct Answer: C) TutorClaw and building agents read from it as a verified canonical source — not the open internet — and human judgment verifies what is taught and built**

Explanation: The book acts as a verified canonical source: TutorClaw reads from it, building agents use its specs, and humans verify — mirroring the enterprise system-of-record pattern.

Q24. The book describes two TIERS of content for readers. Which pairing is CORRECT?

- A) Beginner Chapters (theory) and Expert Chapters (practice)
- B) Getting Started: Crash Courses (fast path for daily skills) and Deep Dive Chapters (comprehensive reference opened when real work surfaces a gap)
- C) Free Content (overview) and Paid Content (detailed implementation)
- D) Online Modules (video) and Offline Chapters (text)

✓ **Correct Answer: B) Getting Started: Crash Courses (fast path for daily skills) and Deep Dive Chapters (comprehensive reference opened when real work surfaces a gap)**

Explanation: Two tiers: Crash Courses = fast, high-leverage primers; Deep Dive Chapters = comprehensive reference you reach for on demand, not front-to-back.

Q25. An enterprise executive wants to move from 'isolated AI experiments' to 'reliable enterprise deployment.' The Agent Factory recommends which strategic path?

- A) Buy a pre-built AI platform and avoid custom development entirely
- B) Hire an army of AI researchers to develop proprietary models
- C) Establish governance models, implement human-in-the-loop supervision, and execute phased adoption from pilots to enterprise-wide scale
- D) Replace all human employees with Digital FTEs as quickly as possible

✓ **Correct Answer: C) Establish governance models, implement human-in-the-loop supervision, and execute phased adoption from pilots to enterprise-wide scale**

Explanation: For executives, the roadmap is: governance models + risk controls, human-in-the-loop supervision, phased adoption from pilot to enterprise scale.

Q26. 'Domain Expertise' as a critical element means that the MOST VALUABLE AI systems depend on:

- A) The largest training datasets and highest model parameter counts
- B) Deep domain knowledge from professionals (accounting, law, finance, supply chain) that serves as the guiding structure for AI behaviour
- C) Real-time internet access so the AI can look up current information
- D) Proprietary hardware optimised specifically for that domain's calculations

✓ **Correct Answer: B) Deep domain knowledge from professionals (accounting, law, finance, supply chain) that serves as the guiding structure for AI behaviour**

Explanation: Domain expertise is the 'knowledge engine' — professionals' judgment becomes the guiding structure. AI executes routine reasoning; experts provide oversight and accountability.

Q27. How does the Agent Factory distinguish between a 'general agent' (like Claude Code or ChatGPT) and a 'Digital FTE'?

- A) General agents are smarter; Digital FTEs are simpler rule-based bots
- B) A general agent is the all-purpose assistant you direct; a Digital FTE is a custom role-based AI Worker built to perform a specific structured job continuously
- C) General agents require human supervision; Digital FTEs operate with complete autonomy
- D) General agents only work in the browser; Digital FTEs only work on desktop

✓ **Correct Answer: B) A general agent is the all-purpose assistant you direct; a Digital FTE is a custom role-based AI Worker built to perform a specific structured job continuously**

Explanation: General agent = the all-purpose assistant you direct or use to build. Digital FTE = the manufactured role-based worker it produces that keeps running independently.

Q28. The MCP (Model Context Protocol) in the Agent Factory ecosystem serves what purpose?

- A) A training protocol for fine-tuning models on domain-specific data
- B) The standard way an agent plugs into outside tools and data
- C) A security encryption standard for agent-to-agent communication
- D) A file format for storing conversation history across sessions

✓ **Correct Answer: B) The standard way an agent plugs into outside tools and data**

Explanation: MCP = the standard connector protocol that lets agents plug into outside tools, APIs, and data sources.

Q29. Which description BEST captures Claude Code's primary strength compared to OpenCode?

- A) Claude Code is free and open-source; OpenCode requires a subscription
- B) Claude Code connects to more model providers than OpenCode
- C) Claude Code runs Anthropic's most capable models with the deepest Claude ecosystem integration — best for complex, production-critical work where frontier model performance is the constraint
- D) Claude Code works offline; OpenCode requires an internet connection

✓ **Correct Answer: C) Claude Code runs Anthropic's most capable models with the deepest Claude ecosystem integration — best for complex, production-critical work where frontier model performance is the constraint**

Explanation: Claude Code is 'frontier-first' — best for complex multi-file refactors and production-critical work where frontier model performance matters most.

Q30. The Agent Factory concludes that the people who 'gain leverage no previous generation of knowledge worker has had' are those who:

- A) Master the most programming languages and AI frameworks
- B) Work at the largest technology companies with the most resources
- C) Learn to identify work that becomes a Digital FTE, specify the agent, deploy the architecture, and govern the workforce that emerges
- D) Publish the most research papers on agentic AI systems

✓ **Correct Answer: C) Learn to identify work that becomes a Digital FTE, specify the agent, deploy the architecture, and govern the workforce that emerges**

Explanation: The final thesis: those who can identify, specify, deploy, and govern Digital FTEs gain unprecedented leverage as knowledge workers.

CodeVerse Soban

End of Paper — All the Best!

AI Agent Factory | Level 1 | 30 Tricky MCQs