

THE AGENT FACTORY THESIS

MCQ Study Guide & Test Preparation
75 Questions with Answers & Explanations

Topic	The Agent Factory Thesis — Architectural Argument
Total Questions	75 MCQs with Full Explanations
Format	Single Best Answer (A/B/C/D)
Coverage	All Major Sections of the Thesis

SECTION 1: Core Concepts & Definitions

1. What is the primary claim of the Agent Factory Thesis about the most valuable companies in the AI era?

- A) They will sell advanced software subscriptions
- B) They will manufacture AI employees (Digital FTEs) that deliver outcomes at scale
- C) They will replace all human workers with robots
- D) They will focus on building better AI models

✓ **Correct Answer: B**

Explanation: The thesis states: 'the most valuable companies won't sell software — they'll manufacture AI employees (Digital FTEs): role-based systems that compose tools, spawn specialist agents, and deliver outcomes at scale.'

2. In the Agent Factory era, what do customers do instead of 'buying' from AI-Native companies?

- A) Subscribe to services
- B) License software tools
- C) Hire them
- D) Partner with them

✓ **Correct Answer: C**

Explanation: The thesis explicitly states: 'You don't buy from these companies. You hire them.'

3. What are 'Digital FTEs' in the context of the Agent Factory?

- A) Remote human freelancers hired digitally
- B) Role-based AI systems that compose tools, spawn specialist agents, and deliver outcomes

- C) Software tools with full-time equivalent licenses
- D) Digital financial tracking employees

✓ **Correct Answer: B**

Explanation: Digital FTEs (Full-Time Equivalents) are AI employees — role-based systems that form the operating substrate of AI-Native companies.

4. What is the Agent Factory itself defined as in the thesis?

- A) A product you buy from Anthropic
- B) A physical factory building AI hardware
- C) A spec-driven, human-supervised, agent-tool-powered process for designing, manufacturing, and deploying AI Workers
- D) A management software platform

✓ **Correct Answer: C**

Explanation: The Agent Factory is the PROCESS — it uses Claude Code/OpenCode and is a practice you adopt, not a product you buy.

5. What is the 'AI-Native Company' in the Agent Factory framework?

- A) The process that builds AI Workers
- B) The output — a running enterprise staffed by AI Workers, coordinated by a management layer, directed by humans at the edge
- C) The AI Workers themselves
- D) The system of record

✓ **Correct Answer: B**

Explanation: The AI-Native Company is the OUTPUT of the Agent Factory process — also called the Agentic Enterprise in the book.

SECTION 2: Economic Actors & Payment Protocols

6. Which payment protocol was developed by OpenAI and Stripe to power ChatGPT's Instant Checkout?

- A) AP2
- B) x402
- C) ACP
- D) MPP

✓ **Correct Answer: C**

Explanation: ACP (OpenAI + Stripe) powers ChatGPT's Instant Checkout — when an agent buys something inside the chat, ACP authorizes and clears the transaction.

7. Which payment protocol uses cryptographically signed mandates and is backed by 60+ companies?

- A) ACP
- B) AP2
- C) x402
- D) MPP

✓ **Correct Answer: B**

Explanation: AP2 (Google) is a cross-vendor standard backed by 60+ companies. It uses digitally signed permission slips proving the human authorized the agent to spend.

8. What is the x402 protocol and who developed it?

- A) A micropayment protocol by Stripe
- B) A crypto-native payment protocol by Coinbase
- C) An enterprise payment system by Google
- D) A subscription billing protocol by OpenAI

✓ **Correct Answer: B**

Explanation: x402 is a crypto-native payment protocol by Coinbase. Version 2 launched in late 2025; Stripe integrated with it on Coinbase's Base blockchain in early 2026.

9. The MPP protocol is specifically designed for:

- A) Large one-time enterprise payments
- B) Micropayments — enabling AI agents to pay pennies per second under a preset cap
- C) Cross-border currency exchange
- D) Subscription management

✓ **Correct Answer: B**

Explanation: MPP (Stripe/Tempo) is built for micropayments. An AI agent streaming a service can pay pennies per second under a preset cap — enabling consumption-based commerce.

10. What major shift does the Agent Factory era represent in terms of resource acquisition?

- A) Humans still procure all tools and services manually
- B) Agents autonomously buy compute, data, and services within a budget envelope set by humans
- C) Companies outsource procurement to third parties
- D) AI models decide all company spending without human approval

✓ **Correct Answer: B**

Explanation: In the Agent Factory era, agents buy compute, data & services autonomously — a fundamental shift from the SaaS era where humans procured tools and services.

SECTION 3: SaaS Era vs Agent Factory Era

11. What is the value metric in the SaaS Era vs the Agent Factory Era?

- A) SaaS: Per-Outcome Results; Agent Factory: Per-Seat Subscriptions
- B) SaaS: Per-Seat Subscriptions; Agent Factory: Per-Outcome Results
- C) Both use Per-Seat Subscriptions
- D) Both use Per-Outcome Results

✓ **Correct Answer: B**

Explanation: The paradigm shift table clearly shows: SaaS Era = Per-Seat Subscriptions; Agent Factory Era = Per-Outcome Results.

12. What standard replaces 'rigid, point-to-point APIs' in the Agent Factory era?

- A) REST APIs
- B) GraphQL
- C) Model Context Protocol (MCP)
- D) OpenAPI Specification

✓ **Correct Answer: C**

Explanation: The paradigm shift table shows: Integration shifts from 'Rigid, point-to-point APIs' to 'Model Context Protocol (MCP)' in the Agent Factory era.

13. What is the human role in the SaaS Era vs the Agent Factory Era?

- A) SaaS: Supervisor; Agent Factory: Operator
- B) SaaS: Operator; Agent Factory: Supervisor & Verifier
- C) SaaS: Designer; Agent Factory: Manager
- D) Both eras have the same human role

✓ **Correct Answer: B**

Explanation: SaaS Era: Human Role = Operator. Agent Factory Era: Human Role = Supervisor & Verifier.

14. The 'focus' shift in the Agent Factory era changes from:

- A) Who does the work → What work is done
- B) How the work is done → That the work is done (verifiably correct)
- C) When work is done → Why work is done
- D) Quality of work → Speed of work

✓ **Correct Answer: B**

Explanation: The focus shifts from 'How the work is done' (SaaS) to 'That the work is done — verifiably correct' (Agent Factory).

SECTION 4: The Production Engine

15. What are the three core elements that power the Agent Factory production engine?

- A) Models, APIs, and Databases
- B) Specs, Skills, and Feedback loops
- C) Humans, Robots, and Algorithms
- D) Intent, Execution, and Revenue

✓ **Correct Answer: B**

Explanation: Three things power the factory: Specs (written instructions), Skills (packaged abilities in portable folders), and Feedback loops (how the system improves).

16. In the production engine analogy, what does 'intent' correspond to in a car factory?

- A) The finished car
- B) The assembly line workers
- C) The raw materials
- D) The factory building

✓ **Correct Answer: C**

Explanation: In the car factory analogy: raw materials = your intent (what you want done), specialized stations = AI Workers, finished car = verified outcome.

17. What is MCP (Model Context Protocol) compared to in the Agent Factory?

- A) The factory building
- B) A universal power outlet that lets every AI Worker talk to every tool
- C) The system of record
- D) The management layer

✓ **Correct Answer: B**

Explanation: MCP is 'a universal standard that lets every AI Worker talk to every tool, the same way every device in a real factory plugs into the same type of power outlet.'

18. What are 'Skills' in the Agent Factory framework?

- A) Human training programs
- B) Packaged abilities each AI Worker brings — captured as portable, version-controlled folders following the Agent Skills format
- C) Software licenses for AI tools
- D) API endpoints for data access

✓ **Correct Answer: B**

Explanation: Skills are packaged abilities following the open Agent Skills format (agentskills.io), originally released by Anthropic and now adopted across the agent ecosystem.

SECTION 5: The 10-80-10 Rule

19. The 10-80-10 Rule was originally attributed to which famous business leader?

- A) Elon Musk
- B) Jeff Bezos
- C) Steve Jobs
- D) Bill Gates

✓ **Correct Answer: C**

Explanation: Steve Jobs famously followed the 10-80-10 rule: 10% vision-setting, 80% team execution, 10% polish and perfect. Dan Martell also articulates it as 10% ideation, 80% execution, 10% refinement.

20. In the Agent Factory's 10-80-10 rule, what happens during the middle 80%?

- A) The human defines goals and constraints
- B) AI Workers execute: compose tools, spawn sub-agents, deliver outcomes
- C) The human reviews and approves outcomes
- D) The management layer hires new workers

✓ **Correct Answer: B**

Explanation: Middle 80% = AI Workers execute: compose tools, spawn sub-agents, and deliver outcomes. This is the autonomous execution phase.

21. What is the 'verification surface' changing to in the cloud-agent era?

- A) Line-by-line code diffs reviewed in editors
- B) Quickly reviewable artifacts — logs, video recordings, live previews — from agents working on cloud VMs
- C) Manual testing by QA engineers
- D) Automated unit test results

✓ **Correct Answer: B**

Explanation: In the cloud-agent era, agents work for hours on dedicated VMs and return artifacts (logs, video recordings, live previews) rather than line-level changes.

22. As of February 2026, what percentage of pull requests merged into Cursor's own product were produced by autonomous agents?

- A) 10%
- B) 20%
- C) 35%
- D) 50%

✓ **Correct Answer: C**

Explanation: Cursor reports that 35% of pull requests merged into its own product are produced by autonomous agents running on cloud VMs, as of February 2026.

23. According to the 10-80-10 rule, where is human attention 'irreplaceable'?

- A) In the middle 80% execution phase
- B) At the boundaries — the first 10% (intent) and final 10% (verification)
- C) Throughout all phases equally
- D) Only during the final 10%

✓ **Correct Answer: B**

Explanation: The pattern works because it allocates human attention where it is irreplaceable — at the boundaries — while letting execution scale without bottlenecks.

SECTION 6: Two-Layer Model & Identic AI

24. Who coined the term 'identic AI' and what does it mean?

- A) Elon Musk — AI that identifies threats
- B) Don Tapscott — a self-sovereign personal agent that carries your identity, judgment, preferences, and authority
- C) Sam Altman — AI with identical capabilities to humans
- D) Anthropic — AI that verifies user identity

✓ **Correct Answer: B**

Explanation: Don Tapscott coined 'identic AI' — a self-sovereign agent owned by the individual (not the platform) that carries their identity, judgment, preferences, and authority.

25. In the Two-Layer Model, what does the Edge Layer consist of?

- A) Role-based AI Workers executing enterprise tasks
- B) Personal identic agents serving the individual — translating human intent and delegating to AI Workers
- C) The system of record databases
- D) The nervous system event substrate

✓ **Correct Answer: B**

Explanation: Edge Layer = Personal identic agents, serving the individual. They translate human intent, delegate to AI Workers, and govern on behalf of the principal.

26. What is the AI Workforce Layer in the Two-Layer Model?

- A) Personal agents for individual users
- B) Role-based AI Workers serving the enterprise — executing tasks, coordinating workflows, delivering verified outcomes
- C) The management software platform
- D) The payment processing layer

✓ **Correct Answer: B**

Explanation: AI Workforce Layer = Role-based AI Workers serving the enterprise, executing tasks, coordinating workflows, and delivering verified outcomes.

27. Why do personal agents WITHOUT an industrialized workforce fail?

- A) They cost too much to operate
- B) They are just digital assistants with no one to command
- C) They cannot understand human language
- D) They require too much compute

✓ **Correct Answer: B**

Explanation: Personal agents without an industrialized workforce behind them are 'digital assistants with no one to command' — neither layer works alone.

SECTION 7: Two Modes & Seven Principles

28. What is Principle 1 (P1) of the Seven Principles of General Agent Problem Solving?

- A) Code as Universal Interface
- B) Bash is the Key — the agent can act, not just describe
- C) Verification as Core Step
- D) Persisting State in Files

✓ **Correct Answer: B**

Explanation: P1: Bash is the Key — 'The agent can act, not just describe.' This is the foundational principle that distinguishes action from description.

29. What does Principle 3 (Verification as Core Step) state about outputs?

- A) Outputs should be checked by another AI agent
- B) Every meaningful output is checked before it ships — 'Looks right' is the failure mode
- C) Outputs are automatically verified by the system
- D) Human verification is optional for routine tasks

✓ **Correct Answer: B**

Explanation: P3: Verification as Core Step — 'Every meaningful output is checked before it ships. Looks right is the failure mode.'

30. What does Principle 5 (Persisting State in Files) mean?

- A) All data should be stored in cloud databases
- B) The conversation is volatile; the filesystem is durable — what mattered lives in a file
- C) Agent state should be kept in memory for speed
- D) Files should be backed up to multiple locations

✓ **Correct Answer: B**

Explanation: P5: The conversation is volatile; the filesystem is durable. What mattered lives in a file — ensuring important state survives across sessions.

31. Which two principles are described as the 'operational wrapper' around the five core disciplines?

- A) P1 and P2
- B) P4 and P5
- C) P6 (Constraints & Safety) and P7 (Observability)

D) P3 and P4

✓ **Correct Answer: C**

Explanation: P6 defines what the agent is allowed to touch; P7 defines what it actually did. These wrap the five inner working disciplines (P1-P5).

32. What is the key difference between Problem-Solving engagements and Manufacturing engagements in terms of tooling?

A) Problem-solving always uses Claude Code; manufacturing can use any tool

B) Manufacturing always uses Claude Code/OpenCode; problem-solving splits by audience (engineers vs domain experts)

C) Both always use the same tools

D) Manufacturing uses only OpenWork; problem-solving uses Claude Code

✓ **Correct Answer: B**

Explanation: Manufacturing always anchors to Claude Code/OpenCode because building an AI Worker is fundamentally a coding task. Problem-solving splits: engineers use Claude Code/OpenCode, domain experts use Claude Cowork/OpenWork.

SECTION 8: Seven Invariants (1–6)

33. What is Invariant 1 of the Agent Factory?

A) Every human needs a delegate

B) The human is the principal — every legitimate chain of action originates with a human who sets intent, defines budget, and owns the outcome

C) The workforce needs a management layer

D) Each worker picks its own engine

✓ **Correct Answer: B**

Explanation: Invariant 1: The human is the principal. No exception. No delegation of this layer. Intent, judgment, values, budget authority, and outcome accountability are non-transferable.

34. What happens if Invariant 1 (Human as Principal) is absent?

A) The system runs slower

B) Costs increase

C) Unowned systems produce unaccountable outcomes — liability evaporates, alignment becomes impossible, budget has no owner

D) Workers cannot communicate

✓ **Correct Answer: C**

Explanation: Failure if absent: 'Unowned systems produce unaccountable outcomes. Liability evaporates. Alignment becomes impossible because there is no party whose alignment is being preserved.'

35. What is a 'delegate' in Invariant 2, and what happens without one?

A) A backup AI model; without it the system fails

B) A personal agent holding context and authority; without it the human becomes a bottleneck and scale collapses to human typing speed

C) A document approval system; without it papers are lost

D) A database administrator; without it data is corrupted

✓ **Correct Answer: B**

Explanation: The delegate holds the principal's context, represents their judgment, carries their authority envelope. Without it: 'The human becomes a bottleneck. Scale collapses to human typing speed.'

36. What is 'Paperclip' in the Agent Factory reference implementation?

- A) A document management tool
- B) The management layer — the AI-native company operating system that hires, assigns, governs, observes, and retires AI Workers
- C) An AI model for coding tasks
- D) The nervous system event substrate

✓ **Correct Answer: B**

Explanation: Paperclip is the management layer (Invariant 3 reference implementation) — the AI-native company OS exposing hire, assign, govern, observe, and retire as callable APIs.

37. Why must each AI Worker pick its own engine (Invariant 4)?

- A) To reduce costs by using the cheapest option always
- B) Because forcing the whole workforce onto one engine either over-pays for unneeded reliability or under-pays for required reliability
- C) To ensure all workers use the same programming language
- D) To maintain vendor independence

✓ **Correct Answer: B**

Explanation: Invariant 4: Uniform engine choice guarantees uniform trade-offs — the company either cannot afford its reliable workers or cannot trust its cheap ones.

38. What is the 'system of record' (Invariant 5) and why is it critical?

- A) The AI model's context window; critical for remembering conversations
- B) The authoritative store of state (databases, CRMs, ledgers) the workforce reads from and writes to; critical because context windows are transient but systems of record are permanent
- C) The agent's training data; critical for performance
- D) The specification documents; critical for planning

✓ **Correct Answer: B**

Explanation: System of record = authoritative durable state. Without it: 'agents hallucinate facts, double-write transactions, lose work between sessions.' A context window is transient; a system of record is permanent.

39. What does Invariant 6 (Workforce Expandable Under Policy) enable?

- A) Human managers to hire new employees faster
- B) An authorized agent to generate a prompt, provision a runtime, and register a new AI Worker without waking a human — within the authority envelope
- C) The system to automatically scale server infrastructure
- D) Workers to upgrade their own capabilities

✓ **Correct Answer: B**

Explanation: Invariant 6: The meta-layer exposes hiring as a callable capability. An authorized agent can generate, provision, and register a new AI Worker inside the authority envelope without human intervention.

SECTION 9: Invariant 7 — The Nervous System

40. What is the mathematical case for durable execution in Invariant 7?

- A) 100% reliability without durability, 50% with it
- B) A six-step Worker at 95% per-step reliability completes only 74% of runs without durable execution, but ~99.7% with step memoization and selective retry
- C) 80% reliability without durability, 95% with it
- D) No mathematical difference

✓ **Correct Answer: B**

Explanation: The thesis gives the math: $0.95^6 \approx 0.74$ (74%) without durability vs ~99.7% with step memoization and selective retry — 'the difference between a workforce that ships and one that drops one run in four on the floor.'

41. What is 'Inngest' in the Agent Factory reference implementation?

- A) The delegate personal agent
- B) The nervous system — carrying external triggers, internal events, durable execution, and flow control under one operational envelope
- C) The management layer
- D) An AI model provider

✓ **Correct Answer: B**

Explanation: Inngest is the nervous system reference implementation (Invariant 7). It carries: external triggers (schedules, webhooks), internal events (Worker handoffs), durable execution, and flow control.

42. What are the four failure modes if the nervous system (Invariant 7) is absent?

- A) High costs, slow processing, poor UI, no logging
- B) No external triggers (runs at typing speed), no internal events (no coordination), no durability (reliability fails), no flow control (traffic drowns system)
- C) Model hallucinations, data loss, security breaches, compliance failures
- D) Slow hiring, budget overruns, worker conflicts, poor outcomes

✓ **Correct Answer: B**

Explanation: Four failure modes: (1) Without external triggers → runs at human-typing speed. (2) Without internal events → Workers can't coordinate. (3) Without durability → reliability compounds against you. (4) Without flow control → one customer drowns the rest.

43. What distinguishes an 'AI-Native Company' from merely having copilot tools?

- A) It uses more expensive AI models
- B) The nervous system enables work to arrive and propagate WITHOUT a human routing it — making it a company, not an assistant
- C) It has no human employees
- D) It only hires AI-trained humans

✓ **Correct Answer: B**

Explanation: 'A company that only moves when a human prompts it is not a company; it is an assistant.' The nervous system gives the company autonomous trigger and execution capability.

SECTION 10: Runtime Engines

44. Which engine is recommended for 'can't fail' mission-critical AI Workers?

- A) Claude Managed Agents
- B) OpenClaw-native
- C) Dapr Agents wrapping an SDK

D) OpenAI Agents SDK

✓ **Correct Answer: C**

Explanation: Job profile 'Can't fail' → Engine: Dapr Agents wrapping an SDK. Reason: Durable execution, auto-recovery, full observability.

45. Claude Managed Agents is recommended when the job profile is:

- A) Can't fail, with own infrastructure
- B) Shouldn't fail, don't want to operate infrastructure
- C) Nice if it works, low priority
- D) Engineering fleet, parallel agents

✓ **Correct Answer: B**

Explanation: Shouldn't fail + don't want to operate → Claude Managed Agents. It is 'hosted and operated for you.'

46. What are the two planes every runtime engine has?

- A) Model plane and data plane
- B) Harness (control plane — agent loop, model calls, routing, recovery) and compute (execution plane — sandbox where code runs)
- C) Frontend plane and backend plane
- D) Hire plane and retire plane

✓ **Correct Answer: B**

Explanation: Every engine has two planes: Harness (control plane) = agent loop, model calls, tool routing, approvals, tracing, recovery. Compute (execution plane) = sandbox where model-directed code reads files, runs commands, writes artifacts.

47. What is 'OpenClaw' in the Agent Factory?

- A) The management layer
- B) The delegate — the personal agent that holds identity, context, and the authority envelope
- C) A payment protocol
- D) An engine for coding tasks

✓ **Correct Answer: B**

Explanation: OpenClaw is the reference implementation for the delegate (Invariant 2). It is the personal agent that holds context, represents the human's judgment, and brokers work to the management layer.

SECTION 11: Invariants vs Implementations

48. What is the key distinction between an 'invariant' and a 'reference implementation'?

- A) Invariants are software; implementations are hardware
- B) An invariant is a structural requirement that stays true regardless of product — a rule that never changes. A reference implementation is the specific product used today to realize it.
- C) Invariants change yearly; implementations are permanent
- D) Invariants are for humans; implementations are for machines

✓ **Correct Answer: B**

Explanation: 'An invariant is a structural requirement that stays true across every version of the system. A reference implementation is the concrete product used in 2026 to realize one.' The building stands even when furniture changes.

49. What does the thesis mean by 'the architecture was never the products. It was the invariants'?

- A) Products are more important than architecture
- B) Named products in the reference stack can all be replaced tomorrow and the architecture still holds — because the structural rules (invariants) remain constant
- C) The invariants can change but products stay the same
- D) Architecture and products are the same thing

✓ **Correct Answer: B**

Explanation: 'Seven invariants. One chain. Swap any named product in the middle column tomorrow and the architecture still stands — because the architecture was never the products. It was the invariants.'

50. In the reference stack, what can replace any named product?

- A) Nothing — the named products are mandatory
- B) Any product that meets the structural requirements of the corresponding invariant
- C) Only Anthropic-certified alternatives
- D) Only open-source alternatives

✓ **Correct Answer: B**

Explanation: The 'What can replace it' column in the reference stack shows functional requirements (e.g., 'Any MCP-speaking personal agent,' 'Any runtime meeting the job's reliability contract') — not specific brands.

SECTION 12: Workforce Opportunity & Infrastructure

51. According to the World Economic Forum's Future of Jobs Report 2025, what percentage of workers globally are expected to require reskilling by 2030?

- A) 30 out of 100
- B) 40 out of 100
- C) 59 out of 100
- D) 75 out of 100

✓ **Correct Answer: C**

Explanation: By 2030, '59 out of every 100 workers globally are expected to require reskilling or upskilling to adapt to new technologies and ways of working.'

52. By January 2026, what had US data center construction reached annually?

- A) \$10 billion annualized
- B) \$25 billion annualized
- C) \$42 billion annualized
- D) \$100 billion annualized

✓ **Correct Answer: C**

Explanation: By January 2026, US data center construction had reached \$42 billion annualized, while office construction plunged 35% from its peak.

53. How much copper does a single hyperscale AI facility require compared to a conventional data center?

- A) The same amount
- B) Twice as much

- C) Five times as much
- D) Up to 50,000 tons — up to ten times what a conventional data center needs

✓ **Correct Answer: D**

Explanation: 'A single hyperscale AI facility requires up to 50,000 tons of copper, up to ten times what a conventional data center needs.'

54. Meta, Google, Amazon, and Microsoft alone project how much in AI infrastructure spending for 2026?

- A) Over \$100 billion
- B) Over \$300 billion
- C) Over \$600 billion
- D) Over \$1 trillion

✓ **Correct Answer: C**

Explanation: Meta, Google, Amazon, and Microsoft alone project over \$600 billion in AI infrastructure spending for 2026.

SECTION 13: Business Functions & GTM

55. What does 'GTM' stand for and what does the AI workforce handle in GTM?

- A) Global Technology Management — handles server management
- B) Go-To-Market — handles lead enrichment, outreach sequencing, CRM hygiene, pipeline analysis, proposal generation, and demo customization
- C) Growth and Training Management — handles employee development
- D) General Task Management — handles all business tasks

✓ **Correct Answer: B**

Explanation: GTM = Go-To-Market (combined sales, marketing, revenue motion). AI Workers handle lead enrichment, outreach sequencing, CRM hygiene, pipeline analysis, proposal generation, and demo customization.

56. Which system of record do Finance AI Workers run against?

- A) The CRM system
- B) The general ledger
- C) The ticketing system
- D) The code repository

✓ **Correct Answer: B**

Explanation: Finance Workers run against the general ledger. GTM → CRM; Support → ticketing system; Engineering → code repository; Finance → general ledger.

SECTION 14: Future Trajectories

57. What are the three future trajectories the thesis identifies beyond its current scope?

- A) Better AI models, faster compute, cheaper storage
- B) Physical AI Workers (embodied), Fully autonomous economic agents, Cross-company workforce mobility
- C) More human jobs, better UX, lower costs

D) Quantum computing, blockchain integration, VR interfaces

✓ **Correct Answer: B**

Explanation: Three trajectories: (1) Physical AI Workers — embodied agents with actuators; (2) Fully autonomous economic agents — agents as economic actors with legal identity; (3) Cross-company workforce mobility — AI Workers hired across multiple companies.

58. How does the thesis describe physical AI Workers (robots)?

- A) A completely different category requiring new invariants
- B) AI Workers under the same authority envelope, hired through the same management layer, running on a runtime engine that happens to drive actuators instead of API calls
- C) Only relevant to manufacturing companies
- D) Not feasible with current technology

✓ **Correct Answer: B**

Explanation: 'The invariants don't change. The compute layer adds a body.' Physical AI Workers follow the same architecture with actuators instead of API calls.

59. By mid-2026, what category of company validates the thesis as 'no longer a forward-looking abstraction'?

- A) Fortune 500 companies with AI divisions
- B) Single-digit-headcount firms reporting billion-dollar annualized revenue against AI-operated workforces
- C) Government agencies using AI
- D) Universities with AI research programs

✓ **Correct Answer: B**

Explanation: 'By mid-2026, single-digit-headcount firms were reporting billion-dollar annualized revenue against AI-operated workforces — a category of company that did not exist in any meaningful form three years earlier.'

SECTION 15: Key Vocabulary & Advanced Concepts

60. What is the difference between 'AI Workers' and permanent fixtures like OpenClaw and Paperclip?

- A) No difference — all are equally permanent
- B) AI Workers are the workforce (hired, assigned, rostered, retired); OpenClaw and Paperclip are permanent fixtures that coordinate the workforce
- C) OpenClaw and Paperclip are temporary; AI Workers are permanent
- D) AI Workers are software; fixtures are hardware

✓ **Correct Answer: B**

Explanation: 'The delegate and Paperclip are permanent fixtures of the company; AI Workers are the workforce hired and retired through them.'

61. What is the 'system of record' described as in relation to the AI-Native Company?

- A) The AI model's training dataset
- B) The authoritative state the AI Workforce runs against — databases, ledgers, and platforms that hold the truth of the AI-Native Company
- C) The log files of agent activity
- D) The specification documents

✓ **Correct Answer: B**

Explanation: The system of record = 'the authoritative state the AI Workforce runs against — the databases, ledgers, and platforms that hold the truth of the AI-Native Company.'

62. What is an 'engagement' in the Agent Factory framework?

- A) A contract between companies
- B) A single bounded interaction between a human and a general agent
- C) A marketing campaign
- D) A performance review

✓ **Correct Answer: B**

Explanation: An engagement = 'a single bounded interaction between a human and a general agent.' It can be problem-solving (outcome to human) or manufacturing (AI Worker produced).

63. Which tool governs Manufacturing Engagements and which governs Problem-Solving Engagements?

- A) Manufacturing → Seven Principles; Problem-Solving → Seven Invariants
- B) Manufacturing → Seven Invariants; Problem-Solving → Seven Principles
- C) Both governed by the same rules
- D) Manufacturing → Ten Commandments; Problem-Solving → Seven Laws

✓ **Correct Answer: B**

Explanation: Manufacturing engagements → governed by Seven Invariants (architectural constitution). Problem-solving engagements → governed by Seven Principles (session conduct).

64. What analogy does the thesis use to describe the management layer's relationship to AI Workers?

- A) Like a CEO to a board of directors
- B) Like an operating system to a fleet of processes — it composes, schedules, accounts for, and terminates them on policy
- C) Like a teacher to students
- D) Like a factory floor to machines

✓ **Correct Answer: B**

Explanation: 'This layer is to the AI-Native Company what an operating system is to a fleet of processes: it composes them, schedules them, accounts for them, and terminates them on policy.'

65. In the 'worked trace' example, a customer writes in Bahasa Indonesia. What happens WITHOUT human intervention?

- A) The system rejects the request
- B) Paperclip detects the capability gap, calls its hiring API, manufactures a new Bahasa-speaking Worker, handles the customer, and logs the interaction
- C) A human translator is called in
- D) The request is queued for later

✓ **Correct Answer: B**

Explanation: The worked trace shows: Paperclip sees capability gap → calls hiring API → new Bahasa Worker manufactured → reads context from system of record → composes reply → writes log back → handoff through OpenClaw. No human woken.

66. What does the Cursor CEO's description of 'the third era' align with in the Agent Factory thesis?

- A) Traditional SaaS subscription model
- B) Autonomous agents working for hours on cloud VMs, with humans defining problems and reviewing artifacts — precisely the 10-80-10 pattern
- C) Full AI autonomy without human oversight
- D) AI-assisted coding with humans writing most code

✓ **Correct Answer: B**

Explanation: Truell describes 'autonomous agents working for hours on cloud VMs, with humans defining problems and reviewing artifacts' — matching the Agent Factory's 10-80-10 and architectural invariants.

67. What do 'harness' and 'compute' represent as two orthogonal choices?

- A) They are tightly coupled — changing one requires changing the other
- B) They are orthogonal (independent) — the compute plane can be swapped (E2B, Cloudflare, Daytona, Modal, Kubernetes) without rewriting the agent's harness
- C) Only the harness can be swapped
- D) Only the compute plane can be swapped

✓ **Correct Answer: B**

Explanation: 'Sandboxes are also orthogonal. Whichever engine a Worker runs on, the compute plane can be swapped — E2B, Cloudflare, Daytona, Modal, your own Kubernetes — without rewriting the agent.'

68. What does the thesis say will be the winners' measurement in the Agent Factory era?

- A) Number of AI models deployed
- B) Number of software seats sold
- C) Outcomes guaranteed
- D) Number of human employees retained

✓ **Correct Answer: C**

Explanation: 'Winners won't be measured by seats sold. They'll be measured by outcomes guaranteed.' — The fundamental shift from SaaS subscription economics.

69. What does the Three Industrialized Stack consist of?

- A) Models, Data, and Compute
- B) Intent, The Production Engine, and Outcome
- C) Hire, Execute, Retire
- D) Spec, Skill, and Deploy

✓ **Correct Answer: B**

Explanation: The Industrialized Stack has three layers: (1) Intent — goals, constraints, budgets, permissions; (2) The Production Engine — transforms intent into outcomes; (3) Outcome — high-fidelity actions and artifacts.

70. What critical insight does the 'trust layer' focus on in the economic actor trajectory?

- A) Building better AI capabilities
- B) Mandate enforcement, audit trails, and liability — not capability (since the agent can already do the work)
- C) Reducing transaction costs
- D) Improving payment processing speed

✓ **Correct Answer: B**

Explanation: 'The trust layer is where the action is now — mandate enforcement, audit trails, and liability — not capability, because the agent can already do the work; the real challenge is making sure we can trust it while it does.'

71. What does 'Small, Reversible Decomposition' (Principle 4) mean?

- A) Use small AI models for cost efficiency
- B) Break work into atomic steps where every step can be undone
- C) Decompose the company into smaller business units
- D) Write small, reversible contracts

✓ **Correct Answer: B**

Explanation: P4: Small, Reversible Decomposition — 'Work moves in atomic steps; every step can be undone.' This ensures safety and recoverability in agent execution.

72. The thesis states that 'Claude Code Routines remains the specialist trigger for coding-agent automation.' What broader substrate does it 'front'?

- A) Paperclip
- B) OpenClaw
- C) Inngest
- D) Dapr Agents

✓ **Correct Answer: C**

Explanation: 'Claude Code Routines remains the specialist trigger for coding-agent automation, fronting the same substrate when the event is code-shaped.' The same Inngest substrate lies beneath.

73. What is Day AI cited as in the thesis?

- A) A competitor to Anthropic
- B) A production AI-native CRM whose founding engineer describes Inngest as 'the nervous system' of their product
- C) A company that developed the x402 protocol
- D) The company behind Paperclip

✓ **Correct Answer: B**

Explanation: Day AI is a production AI-native CRM. Founding engineer Erik Munson calls Inngest 'the nervous system' of the product — production language from a live company corroborating the thesis.

74. What is the 'authority envelope' in the Agent Factory framework?

- A) A legal document granting AI legal rights
- B) The set of permissions, budget limits, and constraints within which agents are authorized to act on behalf of the human principal
- C) An envelope of encrypted agent communications
- D) The physical server boundaries for AI execution

✓ **Correct Answer: B**

Explanation: The authority envelope is the scope of what agents are permitted to do: budget, permissions, constraints. It originates with the human principal and governs all downstream agent actions.

75. According to the thesis, what is 'identic AI' NOT?

- A) A personal representative agent
- B) Something that carries your identity and judgment
- C) A generic assistant

D) Self-sovereign

✓ **Correct Answer: C**

Explanation: 'It's not a generic assistant. It's your representative.' Identic AI is specifically personal and identity-bearing — the opposite of a generic, one-size-fits-all assistant.

QUICK REFERENCE ANSWER KEY

Q	A	Q	A	Q	A	Q	A	Q	A
1	B	2	C	3	B	4	C	5	B
6	C	7	B	8	B	9	B	10	B
11	B	12	C	13	B	14	B	15	B
16	C	17	B	18	B	19	C	20	B
21	B	22	C	23	B	24	B	25	B
26	B	27	B	28	B	29	B	30	B
31	C	32	B	33	B	34	C	35	B
36	B	37	B	38	B	39	B	40	B
41	B	42	B	43	B	44	C	45	B
46	B	47	B	48	B	49	B	50	B
51	C	52	C	53	D	54	C	55	B
56	B	57	B	58	B	59	B	60	B
61	B	62	B	63	B	64	B	65	B
66	B	67	B	68	C	69	B	70	B
71	B	72	C	73	B	74	B	75	C

CodeVerse Soban

AI Learning Resources — Agent Factory Thesis MCQ Guide