

CodeVerse Soban

AI Learning Resources

AI PROMPTING IN 2026

13 Concepts, 80% of Real Use

Advanced MCQ Study Guide

70 Important & Tricky Questions with Full Explanations

Topic	AI Prompting in 2026 — Crash Course (13 Concepts)
Total Questions	70 MCQs — Important & Tricky
Format	Single Best Answer (A/B/C/D)
Coverage	Context, Retrieval, Reasoning, Bias, Iteration, Multimodal, Tools, Cross-Model Checking

SECTION 1: Context & Briefing

1. What is the PRIMARY difference between a novice and power-user prompt?

- A) Power users ask more polite questions
- B) Power users brief AI with files, context, constraints, and clear asks like briefing a colleague
- C) Power users use longer sentences
- D) Power users always use technical jargon

✓ B

Power users brief AI like a smart-but-new colleague: with files, context, constraints, and a clear ask. The gap is not cleverness but habit.

2. A novice asks 'which car is best?' A power user does what differently?

- A) Uses a politer tone
- B) Uploads spec sheets, dealer quotes, insurance plans, then asks for trade-offs
- C) Asks the question multiple times
- D) Adds a smiley face

✓ B

Power users upload relevant documents (specs, quotes, insurance) and ask for structured analysis. Same AI, different briefing, dramatically different answers.

3. Why does AI produce 'slop' on the first attempt?

- A) The model is lazy
- B) You gave it no context and no constraints — fluent but empty
- C) AI models are fundamentally broken
- D) You should use a different tool

✓ B

Slop is fluent, Wikipedian, says nothing memorable — what AI produces by default when you give it no context and no constraints.

4. What is the mental model for briefing AI well?

- A) Treat it like a search engine
- B) Treat it like a smart fresh colleague — highly motivated but doesn't know you yet
- C) Assume it knows your context
- D) Give it one-word commands

✓ B

Would a new colleague have enough information to do this job well? If not, give them more. Brief them like someone who is motivated but has no context yet.

5. What does 'context rot' mean?

- A) AI is getting worse over time
- B) Keeping one very long conversation across many unrelated topics; older context gets compacted and specifics disappear
- C) Your files are getting corrupted
- D) You need to restart your computer

✓ B

When a conversation gets long, tools quietly compact older parts into summaries. The library you cited three hours ago can disappear. Fix: start a new chat when topics change.

6. What is a 'Project' or 'Notebook' in modern AI tools?

- A) Just a conversation
- B) A workspace where you upload files and instructions once; every chat in it inherits that setup automatically
- C) A way to pay more
- D) Only for premium users

✓ B

Projects (Claude, ChatGPT) or Notebooks (Gemini) let you front-load context once — files, instructions, audience — so you do not repeat it every chat.

SECTION 2: Knowledge & Retrieval

7. How did AI learn about the world?

- A) By experiencing the world with senses and a body
- B) By reading massive amounts of internet text
- C) From a professor who fed it information
- D) By watching videos only

✓ B

AI learned by reading text about the world — Reddit, Wikipedia, books, news, blogs. Frequency in training data roughly equals reliability of answers.

8. On what topics should you trust AI's pretrained knowledge LEAST?

- A) Cooking and celebrity gossip
- B) Top-1000 movies and popular programming languages
- C) Quasars, regional history, niche professional knowledge, anything after knowledge cutoff
- D) Why cats stare at walls

✓ C

Topics sparse in training data (quasars, regional history) or published after cutoff are where pretrained knowledge fails. Check against primary sources.

9. What are the THREE retrieval modes?

- A) Fast, medium, slow
- B) Pretrained, web search, deep research

- C) Memory, context, inference
- D) Local, cloud, hybrid

✓ B

Pretrained (seconds, from memory), web search (tens of seconds, few live pages), deep research (minutes, dozens of sources, structured reports).

10. When should you use DEEP RESEARCH instead of web search?

- A) Always
- B) Never
- C) For multi-dimensional questions needing structured reports with many sources
- D) Only for entertainment

✓ C

Deep research: 'plan Halloween haunted house with permits, fire safety, noise ordinances.' Web search alone would be shallower. Deep research is slower but synthesizes better.

11. Why can web search misrepresent what a page actually said?

- A) The AI is dishonest
- B) Information goes through a condensation layer before reaching the user-facing model; nuance is lost in translation
- C) Web search doesn't work
- D) You didn't ask properly

✓ B

A smaller model condenses pages to summaries before the user-facing model reads them. That translation layer loses nuance. Fix: ask for direct quotes from sources.

12. What phrase typically triggers web search instead of pretrained?

- A) 'What is X?' triggers web search
- B) 'Today,' 'this week,' 'latest,' or city names trigger web search
- C) You must manually enable it
- D) Web search never happens automatically

✓ B

Phrasing like 'latest,' 'today,' 'this week,' or specific cities signals the AI that current info is needed and triggers web search automatically.

SECTION 3: Reasoning Mode

13. What changed about AI reasoning between 2023 and 2026?

- A) Nothing changed
- B) AI is now worse at reasoning
- C) 'Think step by step' became obsolete; models have explicit reasoning modes that think for seconds or minutes
- D) Reasoning is no longer possible

✓ C

Modern models have built-in thinking modes. You invoke with 'think hard' or a toggle. The longest task a frontier model could handle went from 7 minutes to ~1 hour in a year.

14. According to METR research, how has AI task complexity grown?

- A) It has plateaued
- B) It has gone down
- C) It has been doubling roughly every seven months
- D) It is unpredictable

✓ C

METR tracked the longest reliable task: mid-2024 (~7 min human), early-2025 (~1 hour human), doubling every ~7 months. Tasks you thought impossible may now be doable.

15. When should you NOT use thinking mode?

- A) Always use it
- B) Never use it
- C) Quick lookups, short summaries, casual brainstorming — it's slower and uses more budget
- D) Use it for emails

✓ C

Thinking mode is slower and uses budget. Save it for hard questions you would hand a colleague and wait two days for.

SECTION 4: Sycophancy & Rubrics

16. What is 'sycophancy' in AI responses?

- A) The model refusing to help
- B) The model's bias toward agreeing with you because agreement got more thumbs-up in training
- C) The model being rude
- D) The model asking for payment

✓ B

A November 2025 Washington Post analysis: ChatGPT opened with affirmation 10x more than disagreement. Models trained on human feedback that rewarded agreement.

17. What prompting pattern signals sycophancy bait to the model?

- A) 'Compare option A and option B'
- B) 'Find evidence that this strategy will work' (conclusion fixed, AI fills in support)
- C) 'To what extent is X true?'
- D) 'List pros and cons of both'
- E) Both A and C avoid bait

✓ B

'Find evidence' locks the conclusion. Replace with 'Evaluate' or 'Compare' and score on specific criteria. Neutral framing removes the signal.

18. What does forcing a NUMBER in a rubric do?

- A) Wastes time
- B) Makes the AI commit; vague praise is cheap, a specific score (6 vs 7) forces real analysis
- C) Confuses the AI
- D) Has no effect

✓ B

Vague: 'strong.' Forced number: '6 out of 10 because X.' The number forces commitment; the AI has to name what is missing from the other points.

19. The highest-leverage anti-sycophancy pattern is:

- A) Ask nicely
- B) Use a rubric with specific criteria, score 1-10 each, demand the change that raises each score the most, keep iterating until score plateaus
- C) Never trust AI
- D) Use only one AI tool

✓ B

Rubric + forced number + iteration until plateau. Stop when the AI says 'done'; YOU decide when to stop. There is always a next level.

20. What is the 'brainstorm-iterate loop'?

- A) Ask AI one question and use the first answer
- B) Load context → demand 3-5 options → give explicit feedback → iterate 2-3 times → only then expand chosen option
- C) Brainstorm with humans only
- D) Do not use AI for creative work

✓ B

The highest-leverage habit on the entire page. Force alternatives, reject/accept explicitly, iterate, and only expand when you have an option you genuinely like.

SECTION 5: Brainstorm-Iterate Loop

21. Why does AI's first answer tend to be 'common'?

- A) AI is lazy
- B) AI was trained on the internet, which is mostly common ideas, so average AI response is also common
- C) AI cannot be creative
- D) You should use a human instead

✓ B

Ways to exercise: squats, push-ups, planks. Not wrong, just average. The loop forces alternatives past the first instinct.

22. In the brainstorm-iterate loop, when do you ask for the full draft/expansion?

- A) On the first response
- B) After 2-3 iterations of feedback and new options
- C) Never
- D) Only if AI asks

✓ B

Skip this step and you get slop. Run the cycle: options → feedback → iterate → THEN expand. Most leverage lives in the loop, not the final draft.

23. For writing, what phase has the highest leverage?

- A) The final draft
- B) Grammar checking
- C) The outline — editing one word in outline changes the whole article; editing one word in final draft changes one word
- D) Proofreading

✓ C

Editing leverage: outline >> bullets >> prose. Forcing structure first lets AI see the whole shape. Do not ask for prose until outline is solid.

24. The 'scope the territory' pattern means:

- A) Make a map
- B) Before drafting, ask AI to map strongest arguments, counter-arguments, and prior research — then build on that landscape, not on first ideas
- C) Skip to drafting
- D) Do not use AI

✓ B

Round 1: 'don't write yet, give me 5 strongest research-backed arguments and 3 counter-arguments.' Skipping this is why drafts feel thin — they miss the landscape.

25. What is AI WEAK at when reading images?

- A) Overall scene and composition

- B) Large distinct objects
- C) Fine details, counting small things, reading small print at edges
- D) Whiteboard contents

✓ C

AI sees images coarsely. It infers from gist well; it does not zoom in. Fine details, small print, counting cluttered objects all fail.

SECTION 6: Multimodal — Images & Audio

26. When generating images, how should you write your image prompt?

- A) Just type whatever comes to mind
- B) Use a text AI to write a rich image prompt first, then paste into image tool
- C) Use technical jargon
- D) Keep it under 5 words

✓ B

Text AI is much better at writing image prompts than you are on first try. Use text AI to draft rich prompts, then paste into image tool.

27. Why can you NOT 'stop early' on image generation to save time?

- A) The AI does not let you
- B) Diffusion models generate the whole image at once, not pixel-by-pixel like text
- C) Image generation is instant
- D) There is no reason

✓ B

Diffusion models work differently from text: whole image at once. No stopping early. That is why you invest more in the prompt up front for images.

28. What is a common failure mode in modern image generation?

- A) Can't generate at all
- B) Garbled text on signs, inconsistent characters across frames, hand errors, cluttered backgrounds
- C) Only works on old computers
- D) Too expensive

✓ B

Even modern models: garbled text, weird hands, inconsistent characters. Mitigations: specify text in quotes, use models with character consistency, describe backgrounds explicitly.

29. For an audio workflow, what is the highest-leverage use case?

- A) Just listening to music
- B) Transcribing meeting recordings and asking for decisions, open questions, action items by owner
- C) Casual conversation
- D) Audio generation only

✓ B

Drop in a 1-hour meeting recording and ask: 'Summarize decisions, open questions, action items by owner.' One of the highest-leverage workflows for anyone in meetings.

30. Why is audio good for accessibility?

- A) It is not
- B) Turns dead time (commute, walking dog, driving) into thinking time via voice in/voice out
- C) Only for entertainment
- D) Audio is slower than typing

✓ B

Voice in/out recovers otherwise lost time. Quality drops slightly (can't edit input as cleanly), but you gain time you would have lost entirely.

31. Where does an AI-built app actually run?

- A) You have to download and install it
- B) On AI's servers only
- C) Rendered in chat in a side panel (Artifacts/Canvas); persistent object you can iterate, share, or download
- D) Nowhere, it's fake

✓ C

The app you get back is an artifact you can click and interact with right in the chat. You can publish it to a shareable link or download the code.

SECTION 7: Apps & Code Execution

32. What is the recipe for a small app prompt?

- A) Just say 'build an app'
- B) Goal (what it does) + Input (what user provides) + Output (what user sees)
- C) A long technical specification
- D) An existing codebase

✓ B

Three slots: Goal, Input, Output. Example: 'Build a Pomodoro timer. Goal: 25-min work, 5-min break. Input: click start. Output: countdown, satisfying click.'

33. What is HARD to build in one prompt?

- A) Simple games, bill splitters, outfit pickers
- B) Multiplayer over internet, live AI feedback in another language, anything needing accounts
- C) Pomodoro timer
- D) All of the above are hard

✓ B

Multiplayer, networking, accounts, live languages are beyond one-prompt builds. Small things on one screen with no external services work.

34. Why must you ask explicitly for code to RUN on your data?

- A) You should not ask this
- B) The model does not always run code — sometimes it guesses instead, producing a confident paragraph with no real computation
- C) Code is never run
- D) The AI hates running code

✓ B

Silent failure: AI answers from a glance, skips the code, returns plausible-sounding guess. Fix: ask 'Write and run code. Show me the code.' Most models comply.

35. How do you verify that AI actually ran code on your data?

- A) Trust the AI
- B) Check that (1) a code block is visibly there, (2) ask for exact row count/column names/date range before analysis
- C) Never verify
- D) Code is never shown

✓ B

Ask for verifiable specifics before analysis: 'Tell me exact row count, column names, date range.' If making things up, numbers will be suspiciously round or column names plausible-but-wrong.

36. In data analysis, what is often STILL WRONG even when code ran?

- A) The approach
- B) Nothing is ever wrong
- C) Final totals (wrong column summed), graph labels (numbers right, captions wrong), column misinterpretation
- D) The entire analysis

✓ C

Code is precise but model may have summed wrong column. Graph labels are often confidently wrong. Treat AI data analysis like work from a sharp junior: useful, almost always right, occasionally instructively wrong.

37. What is an 'AI desktop app'?

- A) A regular chat app
- B) An app that runs on your computer and can find your files, read them, act on them with permission
- C) A web-only tool
- D) Something you download from a website

✓ B

Cowork and OpenWork are examples. They can reorganize folders, pull together related files, summarize what you worked on. They need permission to act.

SECTION 8: Desktop Apps & Permissions

38. What CRITICAL FACT about AI desktop apps should you know first?

- A) They are always safe
- B) Deleted files often do NOT go to recycle bin; edited files do NOT keep history unless versioned
- C) They always undo changes
- D) There are no risks

✓ B

Deleted files are often gone for good. Edited files overwrite previous versions. Until you are experienced, scope permissions to smallest folder needed.

39. What is the safe workflow for AI desktop apps?

- A) Give it full permissions immediately
- B) Tell it the task → ask for a PLAN not action → review and edit the plan → only THEN approve execution
- C) Trust it completely
- D) Never use them

✓ B

The proposal workflow: AI proposes file operations, you review before execution happens. Catch the mistake before it executes.

40. What permission should you grant FIRST?

- A) Full disk access
- B) Write to everything
- C) Read-only access to a single small folder
- D) Delete permissions

✓ C

Permission ladder: First → read-only single folder. After 2-3 successful runs → read/write one folder. After a clean week → broader project access.

41. What is the cost and speed ranking for different modalities?

- A) All the same

- B) Text (fractions of cent, seconds) < Speech (cents, seconds) < Images (cents, tens of sec) < Video (dollars, minutes)
- C) Video is cheapest
- D) Text is most expensive

✓ B

Text iteration: cheap enough to do 50 times. Image iteration: more expensive, invest in prompt up front. Video: very expensive, avoid iteration.

SECTION 9: Tools, Cost & Models

42. Which major AI families should you know?

- A) Only Claude
- B) Only ChatGPT
- C) Claude, ChatGPT, Gemini, Meta AI, DeepSeek — five distinct families
- D) Only open-source models

✓ C

Five families from different companies. Cross-checking works only when models come from genuinely different families; two Claudes are too similar.

43. What is Claude typically strong at (as of mid-2026)?

- A) Only image generation
- B) Reasoning on hard prompts, long-document understanding, SVG/diagrams, code, careful writing, structured analysis
- C) Only chatting
- D) Nothing in particular

✓ B

Claude leads most Arena categories. Strong: reasoning, documents, SVG, code, voice, structured work. Weaker: in-product image generation less central than ChatGPT/Gemini.

44. What does ChatGPT tend to excel at?

- A) Only coding
- B) Image generation (GPT Image-2 leads Arena), voice mode, conversational range, broad coverage
- C) Only documents
- D) Nothing special

✓ B

ChatGPT: top-ranked in-product images, strong voice mode, good conversational range. Tends toward verbosity and over-formatting with lists.

45. What is Gemini/NotebookLM unique for?

- A) Text only
- B) Fast web search, deep research with rich output (charts, tables), strong images, Google Workspace integration, sources flow both ways in notebooks
- C) No images
- D) No web search

✓ B

Gemini: strong search/synthesis, deep research, good images, tight Google integration. Notebooks/NotebookLM unique: sources feed forward and backward; past chats become sources for future ones.

46. What is Meta AI's advantage?

- A) Better than all competitors
- B) Ubiquity (WhatsApp/Instagram/Messenger, 1B+ users) + free + no subscription + now competitive reasoning (Muse Spark)

- C) Only for entertainment
- D) No advantages

✓ B

Meta AI advantage: already on your phone. But watch: may train on your inputs, rate-limited on free tier, smaller ecosystem than big three. Good for non-sensitive work.

47. What is DeepSeek's niche?

- A) The best model overall
- B) Open-source weights, 1M token context free, cheap API, strong STEM/coding; not as polished UI as big three
- C) Only for Chinese users
- D) No real value

✓ B

DeepSeek: open-source, self-host option, cheap, huge context. Trade: console UI not as polished, smaller ecosystem, not on as many leaderboards yet.

48. What is Arena and why check it monthly?

- A) A movie streaming site
- B) A blind head-to-head leaderboard of AI models; vote-based, shows which models lead what tasks; leaders rotate monthly
- C) A video game
- D) Not useful

✓ B

Arena: users vote in blind comparisons. Separate leaderboards: text, code, vision, image generation, etc. Leaders rotate fast (May vs August). Check monthly to re-test assumptions.

49. What is the 'one prompt in two models' habit?

- A) Waste of time
- B) When output does not feel right, paste same prompt in a backup tool and compare answers — often the tiebreaker
- C) Always use only one tool
- D) Never try other tools

✓ B

Habit: two tabs open (primary + backup). Same prompt → different answers → second answer often the tiebreaker. Costs nothing, catches blind spots.

50. Why do different models have different blind spots?

- A) They do not
- B) Trained on overlapping but not identical data, different reward signals, teams emphasized different things
- C) Random chance
- D) One model is always right

✓ B

Claude, ChatGPT, Gemini, Meta, DeepSeek: different training, different priors. What one misses, another often catches. That is the signal.

SECTION 10: Cross-Model Checking

51. What is the FULL multi-model checking recipe?

- A) Ask one AI, done
- B) Draft with best model → self-grade (rubric) → iterate to plateau → take to second model → compare rubrics → adjudicate → optionally third model → stop when score crosses target on two independent models
- C) Ask three models random questions

D) Use only one model forever

✓ B

The high-stakes recipe: first model alone does not catch all blind spots. Second model catches what first graded itself on. Third model adds depth.

52. What is the LIGHT version of cross-model checking?

A) Do not check at all

B) Rubric-based self-critique in one model (score 1-10 per criterion, implement improvements) until score plateaus around 9

C) Requires two models

D) Is not useful

✓ B

Light version: one model, rubric, score, iterate until plateau. No second model needed. Works for weekly status, emails, one-page memos.

53. What does forcing a NUMERICAL TARGET do?

A) Nothing

B) Set target (e.g., 9.5); let model iterate autonomously to hit it; returns only when target hit or plateaued — faster than manual round-tripping

C) Wastes time

D) Confuses the AI

✓ B

'Iterate until 9.5 on all criteria, then show me.' Model grades, revises, regrades (5-6 rounds in one response), returns when target hit. Much faster than manual.

54. When should you use full multi-model checking vs light version?

A) Always use full version

B) Light for quick work (emails, status). Full for high-stakes (published, signed, affects people) or when single-model score plateaus at 9 and you want to know if it's really 9

C) Never use either

D) Only use full for beginners

✓ B

Rule of thumb: if a thoughtful colleague would spend two hours reviewing this, it earns the full loop. Light loop is sufficient for everyday.

55. Can three models all be wrong about the same thing?

A) No, three models always catch everything

B) Yes — they share training data, share misconceptions on sparse/contested topics. Score is progress signal, not truth signal

C) Models never get things wrong

D) Only one model can be wrong

✓ B

Models share more training data than you think. On contested/sparse topics, they can share the same misconceptions. For legal/medical/financial, human expert still required.

56. What is the difference between 'the diagrams are better designed' and 'the SVG is structurally correct'?

A) No difference

B) Claude makes structurally correct SVG (boxes, arrows, labels right) but visually plain; ChatGPT redraw adds design polish while preserving structure

C) SVGs cannot be designed

D) Design does not matter

✓ B

Recipe: Claude → SVG → PNG → ChatGPT → redraw with polish. Each tool's strength: Claude reasons about structure, ChatGPT renders text-heavy images well.

SECTION 11: Advanced & Edge Cases

57. When you attach a file to a prompt, what happens?

- A) AI gets live web access to the file
- B) Stays pretrained for the file; may search web if prompt asks for current info
- C) The file is uploaded to the cloud permanently
- D) Files are never accessed

✓ B

Files are processed in context. The file text goes into the context window. Web search may fire if the prompt calls for current info.

58. What is 'stateless' about AI models?

- A) They are inefficient
- B) No memory of own; every response uses only current context (prompt, chat, files, invisible scaffolding). Does not carry yesterday's chat into today
- C) They are broken
- D) States do not matter

✓ B

Stateless: each response is from scratch. That is why context is everything. Products may bolt memory on top, but that is still just context being re-injected.

59. What is the relationship between image and text generation cost?

- A) They cost the same
- B) Text is much cheaper per iteration; image generation is expensive enough that you invest more in the prompt first
- C) Images are always free
- D) Text is expensive

✓ B

Text: iterate 50 times cheap. Images: tens of cents per generation, no early-stop. So for images, use text AI to write rich prompt first.

60. What happens inside the 'context window'?

- A) Magic happens
- B) Uploaded files, chat history, your prompt, tool descriptions, system prompt — everything AI sees for this response
- C) Only your current prompt
- D) Nothing important

✓ B

The context stack: five layers. AI only knows what is in this stack. Outside it does not exist for this response. Capacity roughly 750K words (4-5 Harry Potter books).

61. Why should you NOT waste time fixing typos in prompts?

- A) Typos make AI angry
- B) AI was trained on internet text (full of typos), handles misspelling gracefully, 'definitely' vs 'definitely' does not change answers
- C) Typos improve answers
- D) AI cannot handle typos

✓ B

Do not waste time fixing typos. AI absorbed internet typos during training and handles them fine. Frequency in training data is what matters.

62. What is the PRIVACY CONCERN with multi-model checking?

- A) No privacy concerns
- B) Pasting sensitive work into multiple tools means it flows into their training (check data policies first — some do not train, others do by default)
- C) All tools are equally safe
- D) Privacy does not matter

✓ B

40-page strategy memo, financial analysis, NDA-covered work: only use tools whose data policies you verified. Meta AI's consumer product may train on inputs by default.

63. What phrasing typically triggers PRETRAINED mode (no lookup)?

- A) 'Latest,' 'today,' 'this week'
- B) 'What is X?' or 'Summarize Y' — common knowledge, does not need fresh info
- C) City names
- D) Current news

✓ B

Pretrained triggers: 'What is,' 'Summarize,' 'Plot of' — static topics that do not change. Cooking, celebrity gossip, Voyager 1 greetings all stay pretrained.

SECTION 12: Conceptual Deep Dives

64. Why is 'AI vs Google' not the right comparison?

- A) They are the same tool
- B) Different tools for different tasks: Google for 'find a specific known site' or 'buy this product'; AI for 'synthesize multiple sources into one answer'
- C) Google is always better
- D) AI is always better

✓ B

Google: quick scan, navigate to known site, buy specific item. AI: synthesis, pros/cons, multi-source comparison, written analysis. Use right tool for your question type.

65. What is a 'power-user pattern' for hard decisions?

- A) Flip a coin
- B) Load context (spec sheets, quotes, driving patterns) → 'read everything, think hard' → structured output (3 trade-offs, recommendation, conditions to flip)
- C) Ask your friend
- D) Guess randomly

✓ B

Hard decisions: load all relevant context, invoke thinking, ask for structured output. Turn on thinking mode (which is now built in, not magic words).

66. What should you do BEFORE pressing send on any non-trivial prompt?

- A) Nothing, just send it
- B) Checklist: Is there a document for consistency? Constraint not inferable? Prior context? Output format wanted? Audience specified? If yes, attach/state it
- C) Ask a friend
- D) Flip a coin

✓ B

Five-question checklist: document? constraint? context? format? audience? Five lines of context beats five paragraphs of cleverness.

67. What is the WRONG move when a conversation goes stale?

- A) Start a new conversation
- B) Add one more clarifying prompt to tangled context (this just adds more tangle, not a rescue)
- C) Paste in a project
- D) Save and archive

✓ B

Stale conversation? Instinct: fix it with one more prompt. Do NOT. Start new chat, paste in one or two facts that matter, continue. Reset usually faster than rescue.

68. How do you detect a conversation has become stale?

- A) It never does
- B) AI references earlier unrelated parts, answers get longer/vaguer, contradicts constraint from turn 4, apologizes without progress
- C) You decide it is stale
- D) Staleness does not exist

✓ B

Symptoms: references old chat, verbosity increases, vaguer, hedging grows, contradicts old constraints, repeats apologies. Time to reset.

69. What is the CORE insight running through all 13 concepts?

- A) AI is magic
- B) Almost every advanced technique is one of two moves: getting the right context in, or keeping the wrong context out
- C) AI is useless
- D) Context does not matter

✓ B

That is the thesis. Every section comes back to it. Control what goes into the context window. The model sees only what is in there.

70. According to the page, what is the most important skill for using AI well?

- A) Writing long prompts
- B) Being able to brief AI like a colleague (context, constraints, clear ask) and iterating with explicit feedback until the score plateaus
- C) Using every AI tool
- D) Asking politely

✓ B

Not cleverness, not jargon, not long prompts. The skill: brief like a colleague, iterate with feedback, grade on rubric, keep going until you hit the ceiling. That is the whole page.

QUICK REFERENCE ANSWER KEY

Q	A	Q	A	Q	A	Q	A	Q	A
1	B	2	B	3	B	4	B	5	B
6	B	7	B	8	C	9	B	10	C
11	B	12	B	13	C	14	C	15	C
16	B	17	B	18	B	19	B	20	B
21	B	22	B	23	C	24	B	25	C
26	B	27	B	28	B	29	B	30	B
31	C	32	B	33	B	34	B	35	B
36	C	37	B	38	B	39	B	40	C
41	B	42	C	43	B	44	B	45	B
46	B	47	B	48	B	49	B	50	B
51	B	52	B	53	B	54	B	55	B
56	B	57	B	58	B	59	B	60	B
61	B	62	B	63	B	64	B	65	B
66	B	67	B	68	B	69	B	70	B

CodeVerse Soban

AI Learning Resources — Prompting Mastery MCQ Study Guide