CodeVerse Soban

AI Agent Factory — Level 1 Exam

Topics: Code You Never Write · Skills & Connectors

---

70 Tricky MCQs · 4 Options Each · Correct Answer + Explanation

---

Q1. What is the main role of the user when working with AI-written code?

A) Write the code and hand it to the AI to run

B) Be a good client: brief clearly, verify the output, keep the result

C) Learn Python before using AI for any data task

D) Choose the programming language for every task

✔ Answer: B) Be a good client: brief clearly, verify the output, keep the result

Why: The course says: you are the client, not the contractor. Brief clearly, check the result, keep what you paid for.

Q2. Which FOUR signals tell you a task is a code problem?

A) Speed, Accuracy, Clarity, Format

B) Volume, Precision, Repetition, Files

C) Size, Style, Frequency, Output

D) Data, Logic, Memory, Time

✔ Answer: B) Volume, Precision, Repetition, Files

Why: VPRF — Volume, Precision, Repetition, Files. Any one signal is enough to make it a code problem.

Q3. A bookkeeper needs to find 23 mismatches between her ledger and bank statement. Which signal fires?

A) Style

B) Speed

C) Precision and Files

D) Creativity

✔ Answer: C) Precision and Files

Why: Precision (a wrong digit has consequences) and Files (the data lives in spreadsheets) both fire — a classic code problem.

Q4. Why does the AI almost always use Python for data problems?

A) It is the newest programming language

B) Python is built into all browsers

C) More Python exists in AI training data than any other language — it is the AI's mother tongue for data work

D) Python requires no installation on the user's computer

✔ Answer: C) More Python exists in AI training data than any other language — it is the AI's mother tongue for data work

Why: The course explains: Python is the default because more of it appears in AI training data, making the AI far more reliable when writing it.

Q5. You ask the AI to total your expenses and it gives a confident answer with no code block visible. What does this mean?

A) The AI used a faster internal engine

B) The AI estimated rather than computed — no code ran

C) The result is definitely correct

D) The AI used JavaScript instead of Python

✔ Answer: B) The AI estimated rather than computed — no code ran

Why: No code block = no code ran. The AI answered from a glance (estimated), not from computation. Treat it as an opinion, not a fact.

Q6. Which three-line phrase forces the AI to compute instead of estimate?

A) 'Calculate this please. Be accurate. Use numbers.'

B) 'Write and run code. Show me the code you ran. Tell me the exact row count first.'

C) 'Use Python. Import pandas. Print the result.'

D) 'Run a script. Skip errors. Give me totals.'

✔ Answer: B) 'Write and run code. Show me the code you ran. Tell me the exact row count first.'

Why: Line 1 forces computation. Line 2 gives proof. Line 3 is a lie-detector — if the AI is reading the file, it will give exact row counts.

Q7. Which part of the five-section brief captures your professional knowledge that the AI cannot guess?

A) Goal

B) Output

C) Rules

D) Input

✔ Answer: C) Rules

Why: Rules is where your domain expertise lives — e.g., 'the school year starts in August' or 'Pending means not yet submitted.' The AI knows Python; it does not know your system.

Q8. What should the Edge cases section of a brief specify?

A) Which programming language to use

B) What to do when data is imperfect — blank rows, duplicates, wrong formats

C) The expected output format

D) The total number of rows in the file

✔ Answer: B) What to do when data is imperfect — blank rows, duplicates, wrong formats

Why: Edge cases decide, in advance, how imperfect data is handled. The professional move: 'If you hit a row you can't interpret, skip it and list it at the end.'

Q9. What is the single most powerful verification move in this course?

A) Ask a second AI to re-check the answer

B) Run the code on a small slice whose correct answer you already know

C) Read every line of the Python code carefully

D) Ask the AI to score its own confidence

✔ Answer: B) Run the code on a small slice whose correct answer you already know

Why: The known-answer test: run the code on one slice you can verify yourself. If it matches, you have earned confidence for the full dataset.

Q10. When the AI's code crashes with red text, what is the correct response?

A) Start a new chat and try again from scratch

B) Learn to read the error message yourself

C) Paste the full error text back and say: diagnose, fix, rerun

D) Switch to a different AI tool

✔ Answer: C) Paste the full error text back and say: diagnose, fix, rerun

Why: Errors are dialogue. You forward the red text; the AI reads its own error messages and fixes the code. You never need to understand the error.

Q11. A task that trips the Repetition signal is especially valuable because:

A) It runs faster than one-time tasks

B) The script becomes an asset you own and can rerun forever for free

C) Repetitive tasks are simpler to code

D) AI gives discounts for repeated prompts

✔ Answer: B) The script becomes an asset you own and can rerun forever for free

Why: Repetition turns a script into a permanent asset. Describe once, run forever, hand to anyone. This is the keep-the-script habit.

Q12. What three things should you keep together when storing a reusable script?

A) The script, the AI chat log, and a screenshot

B) The brief, the script, and a known-answer sample input

C) The script, the output file, and an email to your manager

D) The Python file, the requirements.txt, and a README

✔ Answer: B) The brief, the script, and a known-answer sample input

Why: Brief (the what and why), script (the how), and a sample input (tomorrow's known-answer test) — three files, one folder.

Q13. Claude.ai runs code in a sandbox. What does 'sandbox' mean here?

A) A folder on your desktop

B) A paid cloud service requiring subscription

C) A temporary computer on Anthropic's side — your own machine is never touched

D) A Python environment you install once

✔ Answer: C) A temporary computer on Anthropic's side — your own machine is never touched

Why: The browser sandbox is zero-risk: uploads go in, results come out, and nothing on your computer is affected.

Q14. When is Claude Code or OpenCode better than Claude.ai for a code task?

A) When you want a prettier interface

B) When you need code to run on a schedule, touch your real files, or handle many files without uploading

C) When the task involves JavaScript instead of Python

D) When you are using a phone

✔ Answer: B) When you need code to run on a schedule, touch your real files, or handle many files without uploading

Why: Claude Code / OpenCode sit inside your folder on your machine — no uploads, persistent scripts, can run on a schedule and touch local files.

Q15. What makes Cowork and OpenWork different from Claude Code and OpenCode?

A) They only work on Windows

B) They are desktop apps with a built-in plan-then-approve safety step

C) They cannot write files to your computer

D) They require a paid account

✔ Answer: B) They are desktop apps with a built-in plan-then-approve safety step

Why: Cowork / OpenWork are desktop apps with plan-then-approve built in — the agent shows you its plan and waits for approval before acting on your files.

Q16. Rule 1 of the blast-radius rules says you should work on what before running a script?

A) A live production folder

B) A copy of the folder

C) A cloud backup

D) A ZIP archive

✔ Answer: B) A copy of the folder

Why: Rule 1: copy the folder first, then work only on the original. If something goes wrong, the backup makes it a shrug, not a disaster.

Q17. What is a 'dry run' in the context of file-touching scripts?

A) Running the code without any input data

B) A test that shows every planned operation without actually performing any changes

C) Running the script twice to confirm the result

D) A simulation that runs only on the first 10 rows

✔ Answer: B) A test that shows every planned operation without actually performing any changes

Why: A dry run lists every planned rename, move, or delete and then stops. You read the plan, catch mistakes, and approve before anything actually changes.

Q18. Why should you never point an agent at your entire home directory or Desktop?

A) It would run out of memory

B) AI cannot read folders with spaces in the name

C) A confused script in the wrong folder could damage many unrelated files

D) It costs more tokens to scan large directories

✔ Answer: C) A confused script in the wrong folder could damage many unrelated files

Why: Rule 3 — scope to the smallest folder that contains the problem. A mistake inside expenses-2025/ damages 12 CSVs; a mistake at the root can damage your entire drive.

Q19. Which task is an ANSWER problem, not a code problem?

A) Total 5,000 rows of sales data by region

B) Rename 300 photos using the date stored inside each file

C) Explain what a mutual fund is

D) Find which of 80 contracts have a non-standard cancellation clause

✔ Answer: C) Explain what a mutual fund is

Why: 'Explain what a mutual fund is' trips none of the four signals (VPRF). It is a knowledge question — an answer problem answered by the AI's reasoning, not code.

Q20. 'Roughly how did costs trend this year?' is risky because:

A) It is too long a question for the AI

B) The word 'roughly' signals that precision doesn't matter, so the AI may estimate instead of compute

C) Trend questions require JavaScript, not Python

D) The AI cannot access cost data without a connector

✔ Answer: B) The word 'roughly' signals that precision doesn't matter, so the AI may estimate instead of compute

Why: 'Roughly' tells the AI precision is not needed. On anything where a decision rests on the number, always use explicit language: 'Write and run code to compute...'

Q21. What is a CSV file?

A) A compressed zip archive of spreadsheets

B) A proprietary Excel format only Microsoft software can open

C) A plain text file with values separated by commas — the simplest spreadsheet format

D) A PDF report exported from a database

✔ Answer: C) A plain text file with values separated by commas — the simplest spreadsheet format

Why: CSV = comma-separated values. Plain text rows. Almost every system's 'export' button produces one, and Excel opens it like any other spreadsheet.

Q22. The plain-English replay verification asks the AI to:

A) Translate the Python code into English

B) Describe step-by-step what the code did, including how it handled blank rows and duplicates

C) Re-run the code and compare both outputs

D) List every variable name used in the script

✔ Answer: B) Describe step-by-step what the code did, including how it handled blank rows and duplicates

Why: You ask the AI to explain the code's behavior as a procedure a colleague can audit. Wrong logic shows up in English — most real errors are procedure errors.

Q23. After three failed attempts to fix the same error, the course advises:

A) Download the file and fix it in Excel manually

B) Give up and use a different AI tool

C) Stop and ask for two completely different approaches — three strikes means the approach is wrong

D) Ask the AI to try once more with the same method

✔ Answer: C) Stop and ask for two completely different approaches — three strikes means the approach is wrong

Why: Three strikes = wrong approach, not wrong typing. Stop digging, restate the problem fresh, and ask for two completely different strategies.

Q24. Which statement about code is TRUE according to the course?

A) Code is approximate — it may round numbers slightly

B) Code is exact in both directions: it computes perfectly, and also executes the wrong instruction perfectly at scale

C) Code automatically detects and skips errors in the input data

D) Code runs slower on larger files but is always accurate

✔ Answer: B) Code is exact in both directions: it computes perfectly, and also executes the wrong instruction perfectly at scale

Why: 'Code is exact, in both directions.' It will rename 300 files wrongly just as fast as rightly. Power and danger come from the same property.

Q25. When should you use the cross-model check (verification level 5)?

A) For every task, regardless of importance

B) Only when the AI returns red-text errors

C) For high-stakes outputs — reconciliations, auditor reports, board analyses

D) When the output is longer than one page

✔ Answer: C) For high-stakes outputs — reconciliations, auditor reports, board analyses

Why: Cross-model check = take the same brief to a second AI from a different family and compare numbers. Two independent programs agreeing is the strongest signal this technology offers.

Q26. According to the course, 'judgment wearing a computation costume' means:

A) The AI pretends to run code when it is actually estimating

B) Tasks like 'which employees to promote' look like code problems but the criteria are yours, not Python's

C) AI dresses up wrong answers to look like correct computations

D) Python code that produces graphs looks more convincing than plain numbers

✔ Answer: B) Tasks like 'which employees to promote' look like code problems but the criteria are yours, not Python's

Why: 'Code computes; you decide.' The arithmetic of a promotion decision is trivial; the criteria for who deserves promotion are a human judgment the AI cannot make.

Q27. What does saving a script with a plain-English description at the top achieve?

A) It makes the script run faster

B) It lets the AI read and modify the script in a future session without you re-explaining the rules

C) It is required by Claude.ai's file upload rules

D) It converts the Python file into a Markdown document

✔ Answer: B) It lets the AI read and modify the script in a future session without you re-explaining the rules

Why: A description at the top = the Intent Layer for that script. Six months later, hand any AI the brief plus the script and say 'the late-cutoff rule changed; update it.'

Q28. Which surface should you 'graduate to' when uploading files every week starts to feel tedious?

A) A different AI provider with higher limits

B) Claude Code or OpenCode — agents that live inside your folder and see files without uploading

C) A paid Claude.ai plan with larger file limits

D) The Cowork mobile app

✔ Answer: B) Claude Code or OpenCode — agents that live inside your folder and see files without uploading

Why: The graduation signal from the course: 'stay on Claude.ai until uploading becomes the annoying part of your week.' Claude Code / OpenCode remove that friction.

Q29. Why does the course say 'you never tell the AI which language to use'?

A) The AI ignores language instructions

B) All AI tools only support Python

C) Choosing the language is a technical judgment you are delegating — describing the outcome lets the AI pick correctly

D) Specifying Python slows down the response

✔ Answer: C) Choosing the language is a technical judgment you are delegating — describing the outcome lets the AI pick correctly

Why: 'Describe the job, demand code, leave the language to the AI.' Adding 'in Python' is micromanagement based on a guess you can't justify — the AI picks based on what the output needs.

Q30. A teacher's submission-check brief takes effort the first time. What is the marginal cost of running it next assignment?

A) The same — you re-explain from scratch each time

B) Higher — more students means more setup

C) Near zero — one sentence: 'run my submission-check script on this new class list'

D) Moderate — you must re-upload the script and re-explain the rules

✔ Answer: C) Near zero — one sentence: 'run my submission-check script on this new class list'

Why: Keep-the-script habit: once built, rerunning costs almost nothing. That multiplication — describe once, run forever — is the economic argument of the course.

Q31. In one sentence, what does a Skill do that a chat message cannot?

A) It gives the AI access to your apps and files

B) It teaches the AI how to do a task your specific way, every time, without re-explaining

C) It runs code on your computer automatically

D) It connects the AI to the internet for live data

✔ Answer: B) It teaches the AI how to do a task your specific way, every time, without re-explaining

Why: A chat tells AI what to do once. A Skill teaches AI how to do it every time — no re-explaining, always your way.

Q32. In the kitchen analogy, Connectors represent:

A) The recipe cards pinned to the wall

B) The chef's personal cooking style

C) The kitchen itself — the tools and stocked pantry with your real ingredients

D) The order a customer shouts across the counter

✔ Answer: C) The kitchen itself — the tools and stocked pantry with your real ingredients

Why: Connectors are the kitchen: the stove, the knives, the fridge full of your actual data. Without them, AI can describe a meal but cannot cook one.

Q33. What is the minimum required file inside a Skill folder?

A) A Python script named main.py

B) A file named SKILL.md containing a name, description, and instructions

C) A JSON config file listing all trigger phrases

D) An assets folder with a template document

✔ Answer: B) A file named SKILL.md containing a name, description, and instructions

Why: The minimum Skill is a folder with one file: SKILL.md. It needs a name, a description, and plain-English instructions. No code is required to start.

Q34. A Connector can ONLY reach data that:

A) Has been uploaded to Claude.ai

B) Is stored in a CSV format

C) Your own account can already access in the original app

D) Has been shared publicly on the internet

✔ Answer: C) Your own account can already access in the original app

Why: AI inherits your permissions. A connector cannot reach something you yourself can't reach — connecting Google Drive does not hand AI your whole company's drive.

Q35. How do Skills fire in Claude.ai by default?

A) You type a slash command to activate each skill manually

B) You click the skill name in the sidebar before every task

C) Automatically — AI reads your prompt and loads the matching skill on its own

D) You must name the skill explicitly in every message

✔ Answer: C) Automatically — AI reads your prompt and loads the matching skill on its own

Why: In Claude.ai, skills fire automatically. You describe your task; AI checks it against skill descriptions and loads the right one. No command needed 90% of the time.

Q36. Which field in a SKILL.md decides whether the skill ever fires?

A) The name field

B) The instructions section

C) The description field

D) The assets folder

✔ Answer: C) The description field

Why: The description is all AI keeps loaded at all times. It reads the description to decide relevance — a vague description means the skill never triggers.

Q37. What is 'progressive disclosure' in the context of Skills?

A) Skills are gradually revealed to users over time

B) AI keeps only the short description loaded always, and opens full instructions only when a request matches

C) Skills reveal more options the more you use them

D) The AI discloses how many skills you have installed

✔ Answer: B) AI keeps only the short description loaded always, and opens full instructions only when a request matches

Why: Progressive disclosure: AI holds a short summary of each skill. Full instructions are loaded only when your request matches. This is why dozens of skills don't slow the AI down.

Q38. Which is the correct folder name format for a Skill?

A) Client Monthly Summary

B) client_monthly_summary

C) client-monthly-summary

D) ClientMonthlySummary

✔ Answer: C) client-monthly-summary

Why: Skill folders use kebab-case: lowercase words separated by hyphens. No spaces, no underscores, no capital letters.

Q39. What is the biggest difference between a Skill and a Project in Claude?

A) Projects are free; Skills cost credits

B) A Skill fires on demand anywhere; a Project always loads its context for every chat inside it

C) Projects can contain code; Skills cannot

D) Skills are shared with the team; Projects are private

✔ Answer: B) A Skill fires on demand anywhere; a Project always loads its context for every chat inside it

Why: Project = always-on context for a body of work. Skill = dormant until a matching request fires it, from anywhere, even outside any project.

Q40. An accountant re-pastes the same formatting rules every month. What does she need?

A) A Connector to Google Drive

B) A Project with a standing context

C) A Skill that encodes those formatting rules

D) Custom Instructions in Settings

✔ Answer: C) A Skill that encodes those formatting rules

Why: Re-explaining how to do it every time = Skill. She needs to teach the AI once, and it applies the rules automatically whenever she asks for a client summary.

Q41. A marketer keeps copy-pasting content from Notion into the chat. What does she need?

A) A Skill with brand voice rules

B) A Connector to Notion so AI can reach the data directly

C) Custom Instructions with her brand guidelines

D) A Project that holds the Notion content

✔ Answer: B) A Connector to Notion so AI can reach the data directly

Why: Copy-pasting from another app every time = Connector. A Notion connector lets AI reach the calendar directly — no copy-paste needed.

Q42. Which is the BEST practice when first connecting an app as a Connector?

A) Grant full write access immediately so AI can do everything

B) Connect all your apps at once for maximum productivity

C) Start read-only and grant write access only after the tool has behaved well

D) Share the connector with your whole team on day one

✔ Answer: C) Start read-only and grant write access only after the tool has behaved well

Why: Start read-only, then expand to write access as trust is earned. 'Search and summarize' is low-risk. 'Send' or 'delete' are write actions that need a track record first.

Q43. Why is the description field more important than the instructions in a Skill?

A) Instructions are optional; descriptions are required

B) AI reads only the description to decide relevance — wrong description means the skill never fires

C) Descriptions are shorter and use fewer tokens

D) Instructions are only read by the skill-creator tool, not by Claude itself

✔ Answer: B) AI reads only the description to decide relevance — wrong description means the skill never fires

Why: The description is the gatekeeper. AI decides relevance from the description alone — a vague or wrong description means a skill that never triggers or fires on the wrong things.

Q44. Which tool exists specifically to help you build Skills by talking to it?

A) mcp-builder

B) web-artifacts-builder

C) skill-creator

D) brand-guidelines

✔ Answer: C) skill-creator

Why: skill-creator is an Anthropic skill pre-installed on Claude.ai whose whole job is to help you build Skills by describing what you want.

Q45. What does the acronym MCP stand for, and what does it power?

A) Multi-Channel Protocol — powers cross-platform messaging

B) Model Context Protocol — the open standard that Connectors run on

C) Machine Code Plugin — the format Skill scripts are written in

D) Managed Cloud Pipeline — the server infrastructure behind Claude.ai

✔ Answer: B) Model Context Protocol — the open standard that Connectors run on

Why: MCP = Model Context Protocol. It is the open standard under Connectors. You never build it; you click Connect. Like email delivery, you don't need to understand it to use it.

Q46. Before enabling a community skill from an unknown source, what should you do?

A) Install it and test it on dummy data first

B) Ask AI to read the skill and flag anything that contacts external servers, handles credentials, or could leak data

C) Email the skill author to confirm it is safe

D) Run it in Incognito mode to prevent data leaks

✔ Answer: B) Ask AI to read the skill and flag anything that contacts external servers, handles credentials, or could leak data

Why: Read a skill before enabling it. You can ask AI: 'Read this skill and tell me exactly what it instructs you to do. Flag anything that sends data externally.' Pay special attention to bundled scripts.

Q47. What is the difference between Custom Instructions and a Skill?

A) Custom Instructions can contain code; Skills cannot

B) Custom Instructions apply globally to all chats; a Skill fires only for one specific type of task

C) Custom Instructions are only available on paid plans

D) Skills are permanent; Custom Instructions reset each month

✔ Answer: B) Custom Instructions apply globally to all chats; a Skill fires only for one specific type of task

Why: Custom Instructions are always-on preferences (like 'be concise' or 'use metric units'). A Skill is task-specific and fires on demand — 'when I ask for a board paper, use this structure.'

Q48. A skill that fires too eagerly on unrelated requests needs which fix?

A) Delete the instructions section and rewrite it

B) Change the folder name to something more specific

C) Narrow the description or add a negative trigger phrase

D) Move the skill to a Project instead

✔ Answer: C) Narrow the description or add a negative trigger phrase

Why: Over-broad firing = description too wide. Narrow it, or add: 'Do NOT use for quick one-off questions — only for full month-end reports.'

Q49. Which publishing option makes a Skill portable across tools like OpenAI's Codex CLI and Google's Gemini CLI?

A) Uploading it to the Claude.ai directory

B) Sharing it as a ChatGPT Custom GPT

C) The Agent Skills open standard — a SKILL.md file runs across any skills-compatible tool

D) Saving it as a Gemini Gem

✔ Answer: C) The Agent Skills open standard — a SKILL.md file runs across any skills-compatible tool

Why: The Agent Skills open standard (published Dec 2025) means a SKILL.md travels across Claude, Codex CLI, Gemini CLI, VS Code, Cursor, and others. GPTs and Gems are vendor-locked.

Q50. When building a Skill with skill-creator, what loop does the course recommend for testing?

A) Build once, trust it immediately

B) Test triggering, test output, test hard cases, bring failures back to skill-creator to fix

C) Run the skill 100 times and average the quality

D) Send it to a colleague for independent review before using it yourself

✔ Answer: B) Test triggering, test output, test hard cases, bring failures back to skill-creator to fix

Why: The build loop: generate → read → test triggering → test output → test hard cases → bring failures back to skill-creator. Iterate until the skill is reliable.

Q51. Which statement correctly describes the safe-use rule for connector scope?

A) Connect every app you own for maximum convenience

B) Connect only the apps a specific workflow actually needs

C) Start with full write access, then restrict if problems occur

D) One connector is enough for all tasks

✔ Answer: B) Connect only the apps a specific workflow actually needs

Why: Scope is safety. Connect only the apps a workflow needs. Don't grant whole-drive access for a one-folder job. Unnecessary open doors are unnecessary risk.

Q52. What happens when you click 'Save skill' after skill-creator builds a skill?

A) The skill is uploaded to the public directory for others to use

B) The skill is saved privately to your account under Customize → Skills, toggled on and ready

C) A ZIP file downloads that you must manually upload later

D) The skill is saved to your Google Drive for backup

✔ Answer: B) The skill is saved privately to your account under Customize → Skills, toggled on and ready

Why: One click: Save skill puts it directly on your Personal skills shelf (Customize → Skills), toggled on, ready to fire. Private to your account.

Q53. What is the key combination that makes Skills and Connectors most powerful together?

A) Skills write data; Connectors read data

B) The Connector fetches real data; the Skill shapes the output your specific way

C) Skills provide AI intelligence; Connectors provide speed

D) Connectors handle text; Skills handle numbers

✔ Answer: B) The Connector fetches real data; the Skill shapes the output your specific way

Why: The pattern: Connector fetches real data from your apps, Skill formats and shapes the output your way. Together: one sentence in, a finished correctly-formatted result out.

Q54. A doctor who re-explains SOAP note format daily and also copies patient files from Drive needs:

A) Only a Skill — for the SOAP format

B) Only a Connector — for Drive access

C) Both — a Skill for the SOAP format and a Connector for Drive

D) Custom Instructions — global medical style rules

✔ Answer: C) Both — a Skill for the SOAP format and a Connector for Drive

Why: Re-explaining how = Skill (SOAP format). Fetching from Drive = Connector. Both frictions exist, so both solutions are needed together.

Q55. Which setting in Claude.ai must be ON for the built-in Word, PowerPoint, and Excel file skills to work?

A) Allow network egress

B) Code execution and file creation

C) Advanced reasoning mode

D) Multi-model access

✔ Answer: B) Code execution and file creation

Why: Settings → Capabilities → Code execution and file creation. This single toggle is the engine that powers the built-in document skills.

Q56. Why are ChatGPT Custom GPTs and Gemini Gems NOT portable in the same way as Skills?

A) They use a different file format that Claude cannot read

B) They are vendor-locked — they live only inside their respective products

C) They require a paid subscription to share

D) They cannot contain instructions, only knowledge files

✔ **Answer: B) They are vendor-locked — they live only inside their respective products**
Why: Custom GPTs live only in ChatGPT; Gems live only in Google. A SKILL.md is an open standard and travels across tools. Vendor-specific = not portable.

Q57. The 'market-of-one test' for deciding whether to build a Skill means:

A) A skill is worth building only if 1 person in a company will use it

B) Build a skill only when it encodes your specific repeated way that no app does for everybody

C) Skills should cost less than one market price to build

D) A skill is ready when one test case passes

✔ **Answer: B) Build a skill only when it encodes your specific repeated way that no app does for everybody**
Why: If an app already does it for everyone (Splitwise for splitting bills), that's a good prompt, not a skill. A skill earns its place encoding YOUR specific rules no app encodes.

Q58. Where does the Skill description appear in the SKILL.md file, and why is it critical?

A) At the bottom after all instructions — AI reads the whole file top to bottom

B) In a separate JSON file — the SKILL.md only holds instructions

C) In the frontmatter at the top (between --- lines) — it is the only part AI always keeps loaded to decide relevance

D) In the folder name — AI reads the folder name as the trigger

✔ **Answer: C) In the frontmatter at the top (between --- lines) — it is the only part AI always keeps loaded to decide relevance**
Why: Frontmatter (between the --- lines) holds the name and description. This is the always-loaded part — progressive disclosure level one. The body only loads when the description matches.

Q59. Which debugging question helps you catch a too-narrow or too-broad Skill description?

A) 'What language did you use for this skill?'

B) 'When would you use my [skill name] skill, and when would you NOT use it?'

C) 'How many tokens does this skill use?'

D) 'Can you show me the SKILL.md source code?'

✔ **Answer: B) 'When would you use my [skill name] skill, and when would you NOT use it?'**
Why: Ask AI when it would and would not use the skill. It paraphrases the description back — if the answer is wider or narrower than intended, you know exactly what to fix.

Q60. On a Team or Enterprise plan, what can you do with a Skill that you cannot do on a personal account?

A) Create skills with more than 100 instructions

B) Share a skill with colleagues or publish it to the organization's directory — shared skills auto-update when you change the original

C) Connect more than five apps at once

D) Access community skills from the public marketplace

✔ **Answer: B) Share a skill with colleagues or publish it to the organization's directory — shared skills auto-update when you change the original**
Why: Team/Enterprise plans allow sharing skills with specific colleagues or the org directory. Shared skills are view-only and auto-update when the owner edits the original — great for standardizing firm-wide output.

Q61. A kept script (from Code You Never Write) and a Skill (from Skills & Connectors) are BOTH examples of which broader principle?

A) Code that runs on a schedule without human input

B) Describe once, the work is done forever — commissioned once, reused permanently

C) AI replacing human judgment in repetitive tasks

D) Tools that require a paid subscription to share with colleagues

✔ Answer: B) Describe once, the work is done forever — commissioned once, reused permanently

Why: Both encode the same move: describe the work once precisely, let AI produce the artifact, and reuse it forever. The script keeps code; the Skill keeps instructions.

Q62. The 'Rules' section of a code brief and the 'description' field of a SKILL.md share which common purpose?

A) They both list the programming language to use

B) They both put your professional expertise into a place the AI can actually read and act on

C) They both determine which AI model runs the task

D) They both control how the output is formatted

✔ Answer: B) They both put your professional expertise into a place the AI can actually read and act on

Why: Rules = your domain knowledge in a brief. Description = your task knowledge in a skill. Both convert expert knowledge into structured signals the AI reads instead of guesses.

Q63. Using a Connector to fetch data + running a code script on that data + shaping the output with a Skill is an example of:

A) A basic chat workflow requiring no special setup

B) A multi-step pipeline where each layer handles what it does best — access, computation, formatting

C) A workflow only available on enterprise plans

D) A method that bypasses the need for human verification

✔ Answer: B) A multi-step pipeline where each layer handles what it does best — access, computation, formatting

Why: Connector (access) + code (computation) + Skill (formatting) = three layers working together. Each does what it does best — this is the direction the book builds toward with Digital FTEs.

Q64. The plan-then-approve habit in Cowork and the dry-run rule in Code You Never Write both exist for the same reason. What is it?

A) They reduce the number of tokens used per session

B) They let you catch a wrong plan in words before it becomes a wrong action on real files

C) They make the AI explain its code in plain English after running

D) They are required by data privacy laws

✔ Answer: B) They let you catch a wrong plan in words before it becomes a wrong action on real files

Why: Both habits: see the plan before execution. Whether it's a Cowork approval or a dry-run list of renames, you catch mistakes while they are still words — not after files are deleted.

Q65. A Skill's description is to a Skill as the three-line incantation is to a code task. Both serve to:

A) Format the output consistently

B) Force the AI to do the right thing rather than guess or improvise

C) Tell the AI which programming language to use

D) Limit the number of tokens in the response

✔ Answer: B) Force the AI to do the right thing rather than guess or improvise

Why: Description forces the right skill to fire at the right moment. The incantation forces computation instead of estimation. Both are explicit signals that override AI's default of guessing.

Q66. Which course concept explains why you should NOT skip verification just because the AI sounds confident?

A) The jagged frontier — AI is unreliable on all tasks equally

B) Idea 3 from What AI Actually Is — the model has no internal truth-checker; fluency and correctness are not the same

C) The knowledge cutoff — the AI's data may be outdated

D) Progressive disclosure — the AI only loads partial instructions

✔ Answer: B) Idea 3 from What AI Actually Is — the model has no internal truth-checker; fluency and correctness are not the same

Why: Idea 3: the model has no second faculty that audits predictions for truth. Confident tone is a learned style, not a truth signal. Verification in Code You Never Write exists because the AI cannot do it internally.

Q67. The 'Intent Layer' mentioned in both courses refers to:

A) A hidden system prompt Anthropic adds to every conversation

B) The human-written documents (briefs, SKILL.md files) that precisely encode your intent for an agent to act on

C) The AI's internal reasoning process before generating a reply

D) A folder where AI stores its memory between sessions

✔ Answer: B) The human-written documents (briefs, SKILL.md files) that precisely encode your intent for an agent to act on

Why: The Intent Layer = your written-down intentions: the Markdown brief in Code You Never Write, the SKILL.md in Skills & Connectors. Both encode human expertise so agents act on it reliably.

Q68. Both courses warn about a similar failure when a task is NOT described precisely. What is that failure?

A) The AI refuses to help and asks for clarification

B) The AI produces a fluent, confident result that looks correct but is based on guessing rather than your actual rules

C) The AI chooses the wrong programming language

D) The AI uses too many tokens and hits the context limit

✔ Answer: B) The AI produces a fluent, confident result that looks correct but is based on guessing rather than your actual rules

Why: Vague brief = AI guesses the rules (wrong totals, wrong formatting). Vague description = skill fires at wrong moments. In both cases the output looks confident but reflects improvisation, not your standards.

Q69. A pharmacist plants a fake discrepancy to test her script (Code course), and a developer asks AI 'When would you NOT use this skill?' (Skills course). Both actions are examples of:

A) Wasting time on unnecessary double-checking

B) Adversarial testing — deliberately probing for failure before trusting the tool in production

C) Using the cross-model verification method

D) The plan-then-approve safety ritual

✔ Answer: B) Adversarial testing — deliberately probing for failure before trusting the tool in production

Why: Both are adversarial tests: plant something you know is wrong and see if the system catches it; probe the skill's boundaries to confirm it doesn't over-fire. Testing failure modes before trusting is a consistent discipline across both courses.

Q70. According to both courses, what is the FIRST thing you should do before handing any AI-produced output to another person or system?

A) Ask the AI to rate its own confidence from 1 to 10

B) Verify the output against at least one thing you independently know to be true

C) Run the output through a second AI tool for confirmation

D) Format the output as an HTML report for readability

✔ Answer: B) Verify the output against at least one thing you independently know to be true

Why: Both courses share one principle: never act on an untested precision-critical result. In Code — the known-answer test. In Skills — test triggering and output before trusting with real client work. You are the verification layer.

CodeVerse Soban

End of Exam — Best of Luck!

Code You Never Write · Skills & Connectors | 70 MCQs