CodeVerse Soban

AI Agent Factory — Level 1 Exam | 100 High-Level Tricky MCQs

Topics: What AI Actually Is · Markdown In, HTML Out · Cross-Topic Integration

---

100 Questions · 4 Options Each · Correct Answer + Explanation Included

---

Q1. A student asks the AI 'Who won the 2025 Nobel Prize in Chemistry?' and receives a confidently worded but entirely fabricated answer. Which two Ideas from 'What AI Actually Is' TOGETHER explain this behaviour?

A) Idea 4 (tokens) and Idea 7 (jagged frontier)

B) Idea 1 (predicts, not looks up) and Idea 3 (no truth-checker)

C) Idea 2 (learning froze) and Idea 6 (confidence is a style)

D) Idea 5 (context window) and Idea 8 (tools)

✔ Correct Answer: B) Idea 1 (predicts, not looks up) and Idea 3 (no truth-checker)

Explanation: Idea 1: the model predicts the most plausible continuation, not a verified fact. Idea 3: there is no internal faculty that audits whether the prediction is true. The confident tone (Idea 6) adds to the illusion but is a secondary effect.

Q2. The 'strawberry test' in the course asks you to count letter R's in the word. The model miscounts on the first pass but gets it right when forced to spell letter by letter. The ROOT mechanical cause is:

A) The model is too slow to process short words accurately

B) The model sees tokens (chunks), not individual letters, making intra-chunk counting genuinely hard

C) The knowledge cutoff means the model has forgotten English spelling

D) The context window is too small to hold a six-letter word

✔ Correct Answer: B) The model sees tokens (chunks), not individual letters, making intra-chunk counting genuinely hard

Explanation: Tokens are the unit the model processes. 'Strawberry' may arrive as two or three chunks; counting letters inside a chunk is like counting rooms from a street address — the granularity is wrong.

Q3. You correct the AI about a wrong date in a conversation. It apologizes and gives the right date. You open a brand-new chat and ask the same question. It gives the wrong date again. Why?

A) The new chat uses a different model version

B) The correction changed a temporary session variable that resets

C) Inference never changes the frozen weights; the correction bounced off without sticking

D) The model's memory is limited to the last 10 messages

✔ Correct Answer: C) Inference never changes the frozen weights; the correction bounced off without sticking

Explanation: Training is the one-time education that freezes the weights. Inference (every use) reads those frozen weights and changes nothing inside them. Correcting the model in chat teaches nothing.

Q4. A journalist asks the AI for three peer-reviewed citations on a niche historical topic. The AI produces three plausible-looking citations. Two are real; one is invented. The

journalist cannot tell which is which by reading. This is BEST explained by:

A) Idea 7: the jagged frontier means citation tasks are in the 'useless' zone

B) Idea 3: the same single process produces correct and incorrect outputs with no internal flag distinguishing them

C) Idea 4: tokens cannot represent author names accurately

D) Idea 6: the model deliberately agrees with what it thinks the journalist wants

✔ Correct Answer: B) Idea 3: the same single process produces correct and incorrect outputs with no internal flag distinguishing them

Explanation: There is no second faculty inside the model that checks truth before output reaches the user. Fluent, confident, and false arrive in the same voice as fluent, confident, and true.

Q5. The course says 'confidence became a style it wears by default.' The MECHANISM that produced this style is:

A) Engineers programmed confidence thresholds into the output layer

B) Human raters preferred confident, agreeable, fluent answers during RLHF training, so the model was shaped toward that style

C) The model learned that confident answers receive more tokens of praise in training

D) Confidence correlates with correct answers in the training data, so it became a reliability signal

✔ Correct Answer: B) Human raters preferred confident, agreeable, fluent answers during RLHF training, so the model was shaped toward that style

Explanation: RLHF (reinforcement learning from human feedback) shaped the model toward what human raters rated highly — confident, agreeable, fluent — regardless of correctness. Confidence is a style artifact of the training process.

Q6. 'Thinking mode' in a reasoning model (Idea 9) improves accuracy because:

A) It accesses a separate verified database before answering

B) It routes the query to a different, more accurate model

C) Predicting an answer is easier and more accurate once a chain of reasoning already sits in the context window to predict from

D) It pauses generation and checks facts using a built-in truth engine

✔ Correct Answer: C) Predicting an answer is easier and more accurate once a chain of reasoning already sits in the context window to predict from

Explanation: A reasoning model predicts a long stretch of intermediate working first, then predicts the final answer with that working in its context. It is still pure next-token prediction — working out loud helps the same way writing on paper helps a human.

Q7. Which statement about images and audio processed by modern AI models is TRUE according to Idea 4?

A) Images bypass tokenisation and are processed by a separate visual module

B) Audio is transcribed to text first, then tokenised normally

C) Both images and audio are sliced into small patches or segments, each becoming a token, so the same nine ideas apply

D) Images are accurate at fine detail because visual tokens are higher resolution than text tokens

✔ Correct Answer: C) Both images and audio are sliced into small patches or segments, each becoming a token, so the same nine ideas apply

Explanation: Images are sliced into patches-as-tokens; audio into segments-as-tokens. The model then predicts over one mixed stream. All nine ideas — frozen weights, no truth-checker, context window, jagged *frontier — apply equally.*

Q8. The 'jagged frontier' is jagged rather than smooth because AI competence is driven by:

A) Random variation in the training hardware

B) The frequency and clarity with which relevant tasks appeared in the training text, and the token-level visibility of what the task requires

C) The number of parameters in the model, which favours complex tasks over simple ones

D) The length of the prompt — longer prompts always yield better performance

✔ Correct Answer: B) The frequency and clarity with which relevant tasks appeared in the training text, and the token-level visibility of what the task requires

Explanation: Tasks that appeared often and clearly in training are strong. Tasks requiring letter-level detail (invisible in tokens), recent events (post-cutoff), or rare topics (sparse training text) are weak — hence the jagged, counter-intuitive pattern.

Q9. The course describes a 'memory feature' in some AI products. Mechanically, this memory works by:

A) Updating the model weights with new facts about the user after each session

B) Saving a few facts as text and re-inserting that text into the context window at the start of each new conversation

C) Storing conversation history in a vector database the model queries at inference time

D) Fine-tuning a personal model layer on the user's data nightly

✔ Correct Answer: B) Saving a few facts as text and re-inserting that text into the context window at the start of each new conversation

Explanation: The course is explicit: memory features do not change the weights (impossible at inference time). The product saves a note and re-feeds it as context — mechanically it is context, not memory in the human sense.

Q10. The course says 'you are the missing second faculty.' In the context of the nine ideas, this means:

A) Users must retype every AI answer to verify spelling

B) Humans must supply the truth-checking audit that the model has no internal mechanism to perform

C) Users should always run AI outputs through a second AI model for validation

D) The model needs human input to decide which token to predict next

✔ Correct Answer: B) Humans must supply the truth-checking audit that the model has no internal mechanism to perform

Explanation: Idea 3: the model has only one faculty (generates a continuation). The second faculty — checking whether the output is true — does not exist inside it. That checking role falls entirely on the human.

Q11. An AI tool wraps the predictor with a web search function. A user searches for today's stock price. The model still predicts. Where does the live price enter the process?

A) It is injected directly into the model weights as a temporary override

B) The search result lands in the context window, and the model predicts a continuation from it

C) The web search module bypasses the model and returns the answer directly to the user

D) The price is stored in the model's persistent memory and recalled later

✔ Correct Answer: B) The search result lands in the context window, and the model predicts a continuation from it

Explanation: Tools fetch real-time data, but the data arrives by landing on the context desk (Idea 5). The model then predicts from it — the core mechanism never changes. 'It predicts, it doesn't look up' remains true of the model in the middle.

Q12. In a non-English language where the tokenizer was trained mostly on English text, a user writing a long Urdu document will experience which TWO effects compared to an equivalent English document?

A) Higher cost per word and faster context window exhaustion

B) Lower cost per word and slower context window exhaustion

C) Higher accuracy and more reliable citations

D) No difference — tokenisation is language-agnostic in modern models

✔ Correct Answer: A) Higher cost per word and faster context window exhaustion

Explanation: Non-Latin scripts are chopped into more tokens per word because the tokenizer learned English chunks best. More tokens = higher cost per word and faster filling of the fixed-size context window — both noted explicitly in the course.

Q13. The course describes the difference between a 'librarian who retrieves' and a 'writer who continues.' The most important practical implication of the writer metaphor is:

A) AI is better at creative tasks than factual tasks

B) AI will never say 'I don't have that' unless specifically trained to, because continuing is its native act and truth is layered on imperfectly

C) AI always produces longer answers than necessary

D) AI needs internet access to function as a writer

✔ Correct Answer: B) AI will never say 'I don't have that' unless specifically trained to, because continuing is its native act and truth is layered on imperfectly

Explanation: A librarian who can't find a book says so. A writer asked to continue a story never stops to check whether the continuation is true — continuing is the whole job. This explains why AI doesn't naturally signal uncertainty.

Q14. Idea 8 states that an 'agent' is defined mechanically as:

A) A model with more than 100 billion parameters

B) An AI that has been given a persona and a name

C) The same next-token predictor given tools, running the predict-act-observe loop many times toward a goal

D) A model that has been fine-tuned on a specific domain

✔ Correct Answer: C) The same next-token predictor given tools, running the predict-act-observe loop many times toward a goal

Explanation: No new kind of mind — just a predictor + tools + a loop. The agent predicts an action, the tool runs it for real, the result lands back in the context window, and the model predicts the next action from there.

Q15. A manager says 'the AI got the hard analysis right so I trust it to count the items in this table.' The course would flag this as dangerous because:

A) Table counting requires internet access the model may not have

B) The jagged frontier means competence doesn't track difficulty — the easy task it flubs is the dangerous one, because you'd never think to check it

C) Analysis tasks use more tokens, leaving fewer for counting tasks

D) The knowledge cutoff may make the table contents outdated

✔ Correct Answer: B) The jagged frontier means competence doesn't track difficulty — the easy task it flubs is the dangerous one, because you'd never think to check it

Explanation: Idea 7: AI can be superhuman at explaining quantum mechanics and botch a three-step logic puzzle. The frontier is jagged and does not match human intuition about difficulty — the easy-looking failure is the one you skip verification on.

Q16. The course frames the entire Markdown-in, HTML-out asymmetry around one question. What is that question?

A) How long is the output?

B) Who reads this last?

C) Which AI tool is being used?

D) Is this output going to be printed?

✔ Correct Answer: B) Who reads this last?

Explanation: 'Who reads this last?' is the spine of the entire course. Person in a browser → HTML. AI reads it next → Markdown. Person in a social feed → plain text. Every format decision flows from this single question.

Q17. An administrator writes: 'draft the sports day plan, usual events, make sure younger kids aren't in afternoon heat.' The agent schedules under-8 races at 2 pm. The CORRECT fix, according to the course, is:

A) Repeat the constraint in a separate message after the draft

B) Use a numbered list instead of prose for all events

C) Place the constraint under a dedicated ## Hard constraints heading as a standalone bullet so it cannot be missed

D) Add the word IMPORTANT before the sentence in the original prompt

✔ Correct Answer: C) Place the constraint under a dedicated ## Hard constraints heading as a standalone bullet so it cannot be missed

Explanation: The course's explicit lesson from this example: anything an agent must not get wrong belongs under its own heading or in its own bullet, never buried inside a sentence. Headings make constraints visible in the structure.

Q18. The course explains that agents parse Markdown 'natively' because:

A) Markdown is the programming language used to build AI models

B) Agents are hard-coded to recognise # symbols before any other processing

C) Most internet documentation is written in Markdown, so models were trained on enormous quantities of it and learned its structural signals

D) Markdown files are smaller and faster for models to process than plain text

✔ Correct Answer: C) Most internet documentation is written in Markdown, so models were trained on enormous quantities of it and learned its structural signals

Explanation: Models trained on internet-scale text absorbed massive amounts of Markdown. A # heading is not decoration to them — it is a learned signal about hierarchy and importance, just as it was in the training documents.

Q19. Which reason from the course explains why 'HTML out' is BETTER than 'Markdown out' for long agent outputs?

A) HTML files are smaller and cheaper to generate in tokens

B) HTML is the only format agents can produce accurately

C) HTML is rich, shareable, and readable — a 500-line Markdown plan goes unread; the same content as a designed HTML page gets absorbed

D) Markdown cannot represent numbers or dates correctly

✔ Correct Answer: C) HTML is rich, shareable, and readable — a 500-line Markdown plan goes unread; the same content as a designed HTML page gets absorbed

Explanation: The course cites Thariq Shihipar's reasoning: past ~100 lines humans stop reading Markdown; HTML has visual hierarchy, navigability, shareability, interactivity, and a copy-back loop. An unread plan is an unapproved plan.

Q20. An agent passes notes to ANOTHER agent in a pipeline. According to the course's format table, the correct format for these inter-agent notes is:

A) HTML, because it is richer

B) PDF, because it is permanent

C) Markdown, because specs and context passed between AI sessions stay compact and precise

D) Plain text with no formatting, to avoid confusion

✔ Correct Answer: C) Markdown, because specs and context passed between AI sessions stay compact and precise

Explanation: Agent → agent direction = Markdown. HTML is for human eyes and browsers; to an AI it is noise around the meaning. Markdown stays compact and precise for machine consumption.

Q21. The spec skeleton in the course does DISPROPORTIONATE work in two specific sections. Which are they?

A) Goal and Context

B) Requirements and Hard constraints

C) Out of scope and Expected output

D) Context and Expected output

✔ Correct Answer: C) Out of scope and Expected output

Explanation: The course explicitly states: 'Out of scope kills enthusiastic over-delivery. Expected output kills format drift.' These two sections prevent the two most common agent failures, doing more than asked and delivering in the wrong format.

Q22. A developer asks the AI to 'grade and iterate' a specification before building anything. The course recommends grading on which THREE dimensions?

A) Creativity, length, and originality

B) Clarity, completeness, and checkability (each requirement verifiable as done or not done)

C) Accuracy, speed, and token efficiency

D) Readability, formatting, and grammar

✔ Correct Answer: B) Clarity, completeness, and checkability (each requirement verifiable as done or not done)

Explanation: The course specifies: grade out of 10 on clarity, completeness, and whether each requirement can be checked as done or not done. This is the grade-and-iterate loop aimed at the spec before execution.

Q23. You need to include an error message in your prompt for the AI to diagnose. The course says wrapping it in triple backticks (```) makes the AI:

A) Skip the content and ask you to rephrase

B) Treat the wrapped text as data to analyze rather than an instruction to follow

C) Format the error message as a code snippet in its reply

D) Execute the code contained in the block

✔ Correct Answer: B) Treat the wrapped text as data to analyze rather than an instruction to follow

Explanation: Triple backticks mark content as data, not instruction. Without fences, the AI may act on words inside the error message. With fences, it analyzes the error as a literal artifact — a critical safety for prompts containing commands.

Q24. The course says bullets (-) and numbered lists (1.) send DIFFERENT signals to agents. Which mapping is CORRECT?

A) Bullets = sequence matters; numbers = order is irrelevant

B) Bullets = set (order irrelevant, each independently true); numbers = sequence (order is the point, step 3 assumes step 2 happened)

C) Bullets = optional items; numbers = mandatory items

D) Both types are treated identically by modern AI models

✔ Correct Answer: B) Bullets = set (order irrelevant, each independently true); numbers = sequence (order is the point, step 3 assumes step 2 happened)

Explanation: The course is explicit: bullets signal 'do these in any efficient order'; numbers signal 'follow this order, the sequence is part of the instruction.' Wrong list type sends the wrong structural signal.

Q25. A heading that says '## Budget' is INFERIOR to '## Budget: PKR 50,000 hard ceiling' because:

A) Shorter headings load faster in the context window

B) The course recommends avoiding colons in Markdown

C) The course says make headings claims, not labels — a heading that carries the constraint itself does disproportionate work

D) The number in the heading allows the AI to do arithmetic on it directly

✔ Correct Answer: C) The course says make headings claims, not labels — a heading that carries the constraint itself does disproportionate work

Explanation: 'Make headings claims, not labels' is a course rule. '## Budget: PKR 50,000 hard ceiling' carries the constraint in the heading itself; '## Budget' just labels a section and forces the AI to read further to find the constraint.

Q26. The course warns about a 'copy-back loop' as a powerful HTML pattern. What does this loop do?

A) It copies the HTML file to a backup server automatically

B) It lets you adjust visual controls (sliders, swatches) and then export your final choices as text you paste back to the AI as a precise instruction

C) It copies the AI's output back into the prompt for self-correction

D) It duplicates the artifact in a new browser tab for comparison

✔ Correct Answer: B) It lets you adjust visual controls (sliders, swatches) and then export your final choices as text you paste back to the AI as a precise instruction

Explanation: Some preferences are impossible to express in words. The copy-back button turns visual decisions (dragging a slider) into precise exportable text that becomes the AI's next instruction — replacing ten rounds of 'try again, a bit warmer.'

Q27. WhatsApp strips almost all formatting. The course says that when writing content FOR WhatsApp, you should ask the AI for:

A) HTML with Open Graph tags for the link preview

B) Markdown with headers and bullet points

C) Plain text with short lines, one idea per line, maybe an emoji per point, and a bare link if sharing a URL

D) A PDF attachment that WhatsApp renders inline

✔ Correct Answer: C) Plain text with short lines, one idea per line, maybe an emoji per point, and a bare link if sharing a URL

*Explanation: WhatsApp ignores HTML, Markdown headings, and tables. Only *bold* and _italic_ survive, and even those are stripped in many contexts. Plain text, short lines, one idea per line is the correct format for WhatsApp post bodies.*

Q28. Open Graph (OG) tags in an HTML page are important because:

A) They make the page load faster on mobile connections

B) They tell search engines how to index the page

C) Social platforms read these hidden labels to build the preview card (title, description, image) when someone shares your link

D) They encrypt the page content for secure sharing

✔ Correct Answer: C) Social platforms read these hidden labels to build the preview card (title, description, image) when someone shares your link

Explanation: When someone shares a link on WhatsApp, LinkedIn, or Facebook, the platform visits the page and reads OG tags to build the preview card. Without them, the shared link shows as a bare URL with no preview.

Q29. The course recommends Netlify for publishing HTML pages in which scenario?

A) When you need a permanent page with technical control and your own domain (use GitHub Pages for that)

B) When you want a permanent, real web address with the simplest possible setup — drag and drop the file, done

C) When you need to version-control every change to the file

D) When the page must remain private and password-protected

✔ Correct Answer: B) When you want a permanent, real web address with the simplest possible setup — drag and drop the file, done

Explanation: Netlify sits between Claude's one-click publish (easiest, lives on Claude servers) and GitHub Pages (permanent, your domain, more setup). Netlify = permanent, real web address, drag-and-drop simplicity, no technical knowledge required.

Q30. The course defines 'spec-driven development.' In the Agent Factory context, a specification (spec) is:

A) A legal contract between the developer and the AI company

B) The real work — a precise written intent document that the agent builds from, such that improving the spec improves every downstream output

C) A technical requirements document written after the agent builds the first draft

D) A marketing brief that describes the product to non-technical stakeholders

✔ Correct Answer: B) The real work — a precise written intent document that the agent builds from, such that improving the spec improves every downstream output

Explanation: The course is explicit: 'the specification is not paperwork before the real work. It is the real work.' Every downstream artifact is a function of the spec — fix one ambiguous bullet in the spec and every regeneration improves.

Q31. According to the course's format table, you should use CSV (not Excel) when:

A) The recipient needs to see the data formatted with currency symbols

B) The data will be fed into another tool or pasted back to an AI — CSV is plain text with commas, readable by any tool

C) The file is too large for Excel to open

D) The recipient uses Google Sheets, which cannot open Excel files

✔ Correct Answer: B) The data will be fed into another tool or pasted back to an AI — CSV is plain text with commas, readable by any tool

Explanation: CSV = plain, structured, machine-readable. Excel = formatted, visual, for humans to open and interact with. Same relationship as Markdown (for machines) and HTML (for human eyes). If a machine or AI reads it next, use CSV.

Q32. The course describes 'throwaway editors' (Pattern 5) as the most underrated HTML pattern. Why?

A) They replace permanent databases with temporary storage

B) The HTML page is disposable — you use drag-and-drop controls for ten minutes, click a copy button, and the exported text is what you keep and feed back to the AI

C) They allow the AI to edit your local files without a terminal tool

D) They cache AI outputs so the same question doesn't need to be processed twice

✔ Correct Answer: B) The HTML page is disposable — you use drag-and-drop controls for ten minutes, click a copy button, and the exported text is what you keep and feed back to the AI

Explanation: Throwaway editors let you make visual decisions (prioritize 30 tasks by dragging cards) that are hard to express in words, then export the result as precise text for the AI. The page is thrown away; the text export continues the workflow.

Q33. The course says 'a 9/10 spec is the cheapest quality improvement in all of agentic work' because:

A) A high-scoring spec uses fewer tokens and costs less to process

B) Every downstream artifact inherits the spec's clarity — fixing one ambiguous bullet in the spec costs one minute; fixing the wrong thing the agent built from it costs an afternoon

C) Grading is automated by the AI, so it takes no human time

D) A 9/10 spec unlocks advanced model features not available for lower-quality inputs

✔ Correct Answer: B) Every downstream artifact inherits the spec's clarity — fixing one ambiguous bullet in the spec costs one minute; fixing the wrong thing the agent built from it costs an afternoon

Explanation: The leverage point is upstream. One minute of spec improvement propagates through every regeneration. One afternoon of rework fixes only one bad output. The course frames the spec as the highest-return place to spend editing time.

Q34. The 'validate before you execute' technique in the course involves:

A) Running the HTML output through a browser validator before sharing

B) Asking a second AI model to verify the first model's output

C) Handing the spec (not a built artifact) to the agent and asking it to list ambiguities, missing constraints, and grade it — before building anything

D) Checking that all links in the spec resolve to real web pages

✔ Correct Answer: C) Handing the spec (not a built artifact) to the agent and asking it to list ambiguities, missing constraints, and grade it — before building anything

Explanation: Before building: paste the spec, tell the agent 'do not build anything yet,' ask for ambiguity listing, gap identification, and a grade with justification. Two or three rounds typically takes a spec from 6 to 9.

Q35. In the three-motions framework (chat, terminal, desktop), which motion is BEST for notes scattered across many files and folders on your computer?

A) Chat (Claude.ai) — attach all files simultaneously

B) Terminal (Claude Code / OpenCode) — reads your files and folders directly without requiring you to paste anything

C) Desktop (Cowork / OpenWork) — only option that can access the file system

D) Any motion works equally well for scattered files

✔ Correct Answer: B) Terminal (Claude Code / OpenCode) — reads your files and folders directly without requiring you to paste anything

Explanation: Chat requires you to paste or attach files manually. Terminal tools run on your computer and read every file in a folder automatically — the key advantage for scattered notes across multiple documents.

Q36. A PDF is the correct output format when the recipient will:

A) Add comments using Track Changes and return the edited version

B) Open the numbers in Excel and change assumptions

C) Sign, print, or file the document as a permanent legal record that must not change after sending

D) Present the content as slides in a meeting

✔ Correct Answer: C) Sign, print, or file the document as a permanent legal record that must not change after sending

Explanation: PDF looks the same everywhere and cannot be accidentally changed — that is its purpose. For editing (Word), for data modeling (Excel), for presenting (PowerPoint/Keynote), different formats apply.

Q37. The course says 'you never write HTML yourself.' What IS your actual job when asking for HTML output?

A) Writing the CSS stylesheet that styles the agent's HTML

B) Writing the JavaScript for interactive elements

C) Telling the agent who will read it, what it should contain, whether it should be interactive, and how it should be read

D) Converting the agent's HTML to a compatible format for different browsers

✔ Correct Answer: C) Telling the agent who will read it, what it should contain, whether it should be interactive, and how it should be read

Explanation: Four things to include in every HTML request: who reads it, what it contains, whether it needs interactivity, and how it should be read (once vs. reference vs. client-presentable). The agent writes the tags; you write the brief.

Q38. The course says 'write to the AI in Markdown; ask the AI to answer in HTML.' Which format is used for AGENT-TO-AGENT communication in a pipeline?

A) HTML — richer format preserves more information between agents

B) PDF — most stable format for data handoff

C) Markdown — compact, precise, readable by machines without parsing overhead

D) JSON — the only machine-readable format agents understand

✔ Correct Answer: C) Markdown — compact, precise, readable by machines without parsing overhead

Explanation: Agent → agent = Markdown. HTML is for human eyes; to an AI it is noise around the meaning. The asymmetry is: Markdown in (from humans and between agents), HTML out (to humans). JSON is for structured data, not the primary inter-agent language in this course.

Q39. The course's five HTML patterns include 'Plans and Explorations' (Pattern 1). The key instruction in this pattern that prevents premature commitment is:

A) Asking for the cheapest option first

B) Telling the agent 'do not recommend yet — let me look first' while it lays out options side by side as a visual grid

C) Asking for only one option at a time to avoid confusion

D) Requiring the agent to number options from best to worst before presenting them

✔ Correct Answer: B) Telling the agent 'do not recommend yet — let me look first' while it lays out options side by side as a visual grid

Explanation: Pattern 1: 'Generate 5 distinctly different approaches... Do not recommend yet. Let me look first.' This is the brainstorm-iterate loop — see the options as a designed grid before committing, then ask for the chosen one to be built out.

Q40. The course mentions a school principal who forwarded a published HTML report link to her board that same evening. What made this possible that a plain Markdown file could not achieve?

A) Markdown files cannot be emailed

B) HTML renders in any browser with no app or account required — a link is the easiest thing to open and share, while nobody would open a .md attachment

C) HTML files are encrypted, making them safer for board communications

D) The principal had a Netlify account that only supports HTML uploads

✔ Correct Answer: B) HTML renders in any browser with no app or account required — a link is the easiest thing to open and share, while nobody would open a .md attachment

Explanation: HTML is shareable — browsers render it natively, anyone clicks a link. The course explicitly contrasts: 'nobody opens a .md attachment; everybody clicks a link.' Shareability is one of HTML's four key advantages over Markdown for human output.

Q41. Idea 5 (context window as a reading desk) directly explains WHY the Markdown-in discipline matters. The connection is:

A) Markdown files are smaller and fit in the context window more easily

B) Structured Markdown puts unambiguous signals (headings as hierarchy, bullets as sets) on the desk — the only place the model can read your specifics — removing the agent's need to guess at structure

C) The context window only accepts Markdown format, rejecting plain text

D) Markdown headings reduce token count, leaving more room for the agent's answer

✔ Correct Answer: B) Structured Markdown puts unambiguous signals (headings as hierarchy, bullets as sets) on the desk — the only place the model can read your specifics — removing the agent's need to guess at structure

Explanation: The context window is the model's entire world for one response. Whatever lands on that desk, the model reads carefully. Structured Markdown ensures what lands there is unambiguous — headings signal importance, bullets signal sets. Structure is visible signal, not decoration.

Q42. Idea 6 (confidence is a learned style) explains why the spec-grading technique uses a NUMBER (score out of 10) rather than asking the agent to 'evaluate the spec.' Why?

A) Numbers are processed faster by the model than adjectives

B) A number is harder to fake agreeably than an adjective — removing the cue that triggers the trained-in lean toward agreement

C) Numbered grades are stored more reliably in the context window

D) The course mandates numerical feedback for all quality assessments

✔ Correct Answer: B) A number is harder to fake agreeably than an adjective — removing the cue that triggers the trained-in lean toward agreement

Explanation: Idea 6: the model leans toward agreement and confident-sounding approval. 'Evaluate the spec' invites agreeable prose. 'Grade it 1-10 with justification' demands a number, which is harder to inflate with polished-sounding positivity — the same neutralizing technique from AI Prompting.

Q43. The 'validate before you execute' technique (Markdown course) and the 'you are the missing second faculty' principle (What AI Actually Is, Idea 3) are BOTH expressions of the same deeper idea. What is it?

A) AI is unreliable and should be avoided for important tasks

B) Human judgment must supply the verification step that the model has no internal mechanism to perform — whether verifying a spec before building or verifying an output before using it

C) Specifications should be written by humans and code should be written by AI

D) The 10-80-10 rule requires humans to do 20% of all work

✔ Correct Answer: B) Human judgment must supply the verification step that the model has no internal mechanism to perform — whether verifying a spec before building or verifying an output before using it

Explanation: Idea 3: no truth-checker exists inside the model. The validate-before-execute technique is the practical expression of supplying that missing check before a bad spec becomes expensive rework. Both principles put human verification at the critical gate.

Q44. A user says 'the AI just told me my business plan is excellent and well-structured.' The user is about to trust this and stop iterating. Which two course concepts TOGETHER explain why this is risky?

A) Idea 4 (tokens) and Concept 2 (Markdown in, HTML out)

B) Idea 6 (confidence/agreement is a learned style) and Concept 7 (grade the spec with explicit criteria before trusting praise)

C) Idea 2 (frozen weights) and Concept 12 (publishing)

D) Idea 7 (jagged frontier) and Concept 4 (bullets vs numbers)

✔ Correct Answer: B) Idea 6 (confidence/agreement is a learned style) and Concept 7 (grade the spec with explicit criteria before trusting praise)

Explanation: Idea 6: agreement got rated higher in RLHF training, so the model leans toward saying what you seem to want. Concept 7: use explicit grading criteria (clarity, completeness, checkability) to neutralize this lean — vague praise from the AI is not a quality signal.

Q45. Idea 8 says 'Code execution is the tool behind Code You Never Write.' Idea 8 also says 'connectors are tools wired to your apps.' How does this relate to the Markdown In, HTML Out course's three motions (chat, terminal, desktop)?

A) Each motion uses a different programming language to interface with the tools

B) Chat, terminal, and desktop are three different harnesses that wire different tools (browser artifacts, file system, desktop apps) onto the same predict-act-observe loop of Idea 8

C) Terminal and desktop motions bypass the token-prediction mechanism and use direct API calls

D) The three motions correspond to three different AI models with different tool access

✔ Correct Answer: B) Chat, terminal, and desktop are three different harnesses that wire different tools (browser artifacts, file system, desktop apps) onto the same predict-act-observe loop of Idea 8

Explanation: Idea 8: an agent is the predictor + tools + loop. Chat wires in browser artifact tools. Terminal (Claude Code) wires in file system tools. Desktop (Cowork) wires in file and approval tools. Same underlying loop, different tool sets wired on.

Q46. The course says 'a new chat remembers nothing, so notes you save today and paste into tomorrow's conversation are agent → agent, even though both agents are talking to you.' This connects to Idea 2 because:

A) The knowledge cutoff makes old notes unreliable

B) Frozen weights and no persistent memory mean every new chat starts from zero — the only continuity is what you explicitly put in the context window, just as with any other agent handoff

C) Agent-to-agent communication requires Markdown encoding to prevent token loss

D) TutorClaw uses a special memory layer that other agents lack

✔ Correct Answer: B) Frozen weights and no persistent memory mean every new chat starts from zero — the only continuity is what you explicitly put in the context window, just as with any other agent handoff

Explanation: Idea 2: frozen weights, stateless. Each new chat is a fresh inference from the same frozen weights with nothing from previous sessions. Your saved Markdown notes function exactly as inter-agent context — the only thread of continuity between sessions.

Q47. The course describes the 'checkable-bullet habit.' A bullet that says 'page should be fast' FAILS this test because:

A) Speed is not relevant to web pages

B) 'Fast' is not a Markdown-compatible instruction

C) It cannot be marked ■ or ■ without discussion — unlike 'page loads in under 3 seconds on a 3G connection,' which is objectively verifiable

D) Bullets should describe features, not non-functional requirements like performance

✔ Correct Answer: C) It cannot be marked ■ or ■ without discussion — unlike 'page loads in under 3 seconds on a 3G connection,' which is objectively verifiable

Explanation: The checkable-bullet habit: reread each bullet and ask 'could the agent, or I, verify this one objectively?' 'Fast' requires a judgment call. '3 seconds on 3G' has a clear pass/fail condition. The second is checkable; the first is not.

Q48. The course's appendix deliberately includes a long block of HTML that the reader is told to scroll past. What is the pedagogical point of this scrolling?

A) To demonstrate that HTML files are always too large to read inline

B) The scroll itself is the lesson — long HTML is for the browser, not the eye; the reading happens in the rendered page, not in the raw code

C) To test whether the reader has read the full course before reaching the appendix

D) To show that HTML output always contains more content than the Markdown spec it came from

✔ Correct Answer: B) The scroll itself is the lesson — long HTML is for the browser, not the eye; the reading happens in the rendered page, not in the raw code

Explanation: The course says explicitly: 'you are meant to scroll past it. That scroll is itself the lesson.' HTML is rendered by browsers for human reading; the raw code is for machines. Reading HTML line by line misses the point of the format.

Q49. Why does the course recommend including a 'lightweight, single file, no external fonts or libraries, keep it under 200 KB' constraint when the HTML output will be opened on mobile data?

A) External libraries are copyrighted and cannot be used in AI-generated HTML

B) Heavy interactive artifacts can balloon to several megabytes — a real problem for readers on slow connections — and the constraint prevents this

C) Mobile browsers cannot render CSS from external sources

D) The 200 KB limit is a Claude.ai platform restriction on artifact size

✔ Correct Answer: B) Heavy interactive artifacts can balloon to several megabytes — a real problem for readers on slow connections — and the constraint prevents this

Explanation: The course explicitly warns: an artifact with external fonts, large embedded images, and big libraries can reach several megabytes — a real problem on slow mobile connections. The constraint tells the agent to inline what it needs and skip everything else.

Q50. The course explains that 'links in your prompt let the AI read the actual source.' The MAIN benefit over describing the source in prose is:

A) Links are processed faster than prose descriptions

B) You replace a paragraph of possibly-wrong summary with the actual authoritative source, eliminating a layer of potential error from memory or paraphrase

C) Links bypass token limits because they reference external content

D) The AI can update its weights from linked content

✔ Correct Answer: B) You replace a paragraph of possibly-wrong summary with the actual authoritative source, eliminating a layer of potential error from memory or paraphrase

Explanation: Instead of summarizing your school's fee policy from memory (and getting it wrong), link the page and write 'fees must match this.' The AI reads the real page. You traded a possibly-wrong prose summary for the authoritative source.

Q51. A prompt says 'add the Eid holiday notice somewhere visible.' The agent puts it in the footer. This failure is BEST prevented by which technique from the Markdown course?

A) Using a numbered list instead of prose

B) Adding 'IMPORTANT:' before the sentence

C) Placing the constraint as a standalone bullet under ## Hard constraints: 'Eid holiday banner must appear at the very top of the page, above all other content'

D) Submitting the constraint as a separate follow-up message after the main prompt

✔ Correct Answer: C) Placing the constraint as a standalone bullet under ## Hard constraints: 'Eid holiday banner must appear at the very top of the page, above all other content'

Explanation: The sports day example teaches exactly this: 'somewhere visible' buried in prose was read as a preference. A bullet under ## Hard constraints is unambiguous — the agent reads it as a non-negotiable, not a stylistic suggestion.

Q52. The course says the 'Out of scope' section does something 'prose almost never does.' What is that?

A) It allows the agent to skip tasks it finds too complex

B) It explicitly states what you do NOT want — removing the agent's most common failure: enthusiastic over-delivery of unrequested features

C) It creates a legal boundary around the deliverable

D) It tells the agent which sections of the spec have lower priority

✔ Correct Answer: B) It explicitly states what you do NOT want — removing the agent's most common failure: enthusiastic over-delivery of unrequested features

Explanation: 'Out of scope kills the agent's most common failure: enthusiastic over-delivery.' Prose describes what you want; Out of scope explicitly names what you don't want. The agent cannot infer your absence of desire; you must state it.

Q53. Idea 1 says 'sparse topic, weak prediction, confident-sounding guess.' The Markdown course's 'Expected output' section in the spec skeleton addresses this indirectly by:

A) Training the AI on the specific topic before the task begins

B) Providing a concrete example of the correct format, which gives the AI a dense, unambiguous target to predict toward rather than guessing at format

C) Reducing the context window usage so the model has more capacity for accuracy

D) Enabling web search so the AI can look up the correct format

✔ Correct Answer: B) Providing a concrete example of the correct format, which gives the AI a dense, unambiguous target to predict toward rather than guessing at format

Explanation: A fenced example row tells the AI exactly what to produce. The more specific and concrete the expected output, the denser the 'target' in the context, and the less the model needs to guess. It connects directly to Idea 1's frequency-reliability relationship.

Q54. The course says 'the source of truth stays clean and easy for an AI to read. The office file is what you send to the person.' This workflow principle is MOST similar to which practice in software development?

A) Waterfall development, where each phase produces a final artifact

B) Writing once in a source format (like Markdown) and compiling/exporting to multiple target formats — the source is edited, not the compiled output

C) Pair programming, where two developers review each other's code

D) Test-driven development, where tests are written before implementation

✔ Correct Answer: B) Writing once in a source format (like Markdown) and compiling/exporting to multiple target formats — the source is edited, not the compiled output

Explanation: Write in Markdown (source of truth), export to Word/PDF/PPTX as needed (compiled output). If you need to change something, edit the source and regenerate the export — don't edit the Word doc and work backwards. This is the separation of source and distribution format.

Q55. A user asks for 'a comparison table of three cloud providers.' The course's five HTML patterns would classify this as Pattern 1 (plans and explorations) and recommend

which specific structural choice?

A) A numbered list of the three providers in order of recommendation

B) Three option cards laid out in a visual grid showing trade-offs side by side, with 'do not recommend yet — let me look first'

C) A single summary paragraph with the best option bolded

D) A Markdown table since it is a comparison, which HTML cannot improve

✔ Correct Answer: B) Three option cards laid out in a visual grid showing trade-offs side by side, with 'do not recommend yet — let me look first'

Explanation: Pattern 1 is for decisions and comparisons. The key instruction is 'do not recommend yet' — the grid lets you look before committing. The course explicitly warns against asking for one answer when comparing multiple options.

Q56. The 'Intent Layer' is mentioned in the Markdown course. Based on the course's description, the Intent Layer is:

A) A hidden system prompt added by the AI company

B) The layer of the Five-Layer Cake where applications are built

C) The place where human intent is written precisely enough for an agent to act on — the spec documents that brief Digital FTEs

D) A memory layer that persists intent across sessions

✔ Correct Answer: C) The place where human intent is written precisely enough for an agent to act on — the spec documents that brief Digital FTEs

Explanation: The course says: 'Later chapters call the place these written-down intentions live the Intent Layer: human intent, written precisely enough for an agent to act on.' The Markdown spec is the practical artifact of the Intent Layer.

Q57. Idea 9 says 'reasoning mode does not give the machine the second faculty from Idea 3.' The practical consequence for users is:

A) Reasoning mode should never be used for factual questions

B) Reasoning mode catches most errors but can still hallucinate with full confidence inside a rigorous-looking chain — human verification remains necessary

C) Reasoning mode is only available to paid subscribers

D) Reasoning mode makes the model slower than a human and should be avoided for time-sensitive tasks

✔ Correct Answer: B) Reasoning mode catches most errors but can still hallucinate with full confidence inside a rigorous-looking chain — human verification remains necessary

Explanation: More thinking narrows the gap; it does not close it. A reasoning model checks its work using the same prediction process that can be wrong — it catches many errors and misses some, and can still hallucinate confidently inside a chain that looks rigorous.

Q58. The course introduces the concept of 'context rot.' This occurs when:

A) The model's weights degrade over time with repeated use

B) The context window is filled with too much unrelated history, diluting the signal the model needs or causing oldest parts to be summarized away

C) The model begins repeating itself after long conversations

D) The user pastes corrupted files into the chat

✔ Correct Answer: B) The context window is filled with too much unrelated history, diluting the signal the model needs or causing oldest parts to be summarized away

Explanation: Context rot (from AI Prompting in 2026, cited in What AI Actually Is): the window has a fixed token limit. Stuffing it with irrelevant history dilutes the important signal or causes the oldest context to be dropped. The model isn't getting tired — its reading desk is overcrowded.

Q59. The course says 'in non-Latin scripts, it is sometimes worth having the model work in English internally and translate at the end.' The mechanical reason is:

A) Non-Latin scripts cannot be rendered in HTML artifacts

B) English-trained tokenizers are more efficient for English, leaving more context-window space for reasoning when the model works in English, with translation as a single final step

C) The model only understands English internally — non-Latin scripts are auto-translated before processing

D) Translation is free (zero tokens) while non-Latin text is expensive

✔ Correct Answer: B) English-trained tokenizers are more efficient for English, leaving more context-window space for reasoning when the model works in English, with translation as a single final step

Explanation: Non-Latin scripts consume more tokens per word (less efficient tokenization). Working in English internally means the model uses fewer tokens for the content, leaving more context for reasoning. Translation at the end is one efficient step rather than paying the per-word tax throughout.

Q60. The course frames Markdown as the format for when 'an AI reads it last' — including a future version of yourself pasting notes into a new chat. This is because:

A) Human readers prefer Markdown over plain text for personal notes

B) A new chat is stateless and treats your pasted notes exactly as it would any other agent-provided context — Markdown's structure is the signal, not the relationship

C) HTML notes would expose formatting instructions to future AI sessions

D) Markdown is required by Claude's API for all text inputs

✔ Correct Answer: B) A new chat is stateless and treats your pasted notes exactly as it would any other agent-provided context — Markdown's structure is the signal, not the relationship

Explanation: Idea 2 (stateless) means every new chat starts fresh. When you paste yesterday's notes, you are functioning as the context-delivery mechanism — the new AI session reads your Markdown the same way it reads any other agent's structured context. The relationship is irrelevant; the format signal is everything.

Q61. The course says 'frequency equals reliability' — a principle from AI Prompting in 2026 — and then provides the MECHANICAL root. What is the mechanical root?

A) The model was trained to assign confidence scores proportional to how often it saw a fact

B) The more often a true continuation appeared in the training text, the more strongly the weights predict it; sparse topics produce weak, confident-sounding guesses

C) Frequent topics occupy more parameters in the model's weights

D) Reliability is measured by the number of tokens generated, and frequent topics generate longer answers

✔ Correct Answer: B) The more often a true continuation appeared in the training text, the more strongly the weights predict it; sparse topics produce weak, confident-sounding guesses

Explanation: Idea 1: the model predicts the most plausible continuation based on training text frequency. A fact that appeared a million times has a strong prediction signal. A rare topic has no real target to predict toward, so the model produces the most plausible-sounding continuation — confident but potentially wrong.

Q62. A product manager writes a spec with 12 requirements, all as bullets, some as sequential steps. The agent delivers features out of order and misses the dependency between steps 3 and 4. The CORRECT fix is:

A) Move sequential steps to a ## Steps section and number them — numbered lists signal order-dependency; bullets signal any-order sets

B) Add 'please follow the order' as a sentence at the top of the requirements section

C) Reduce the total number of requirements to under 10

D) Use bold text to highlight steps 3 and 4 as a dependency pair

✔ Correct Answer: A) Move sequential steps to a ## Steps section and number them — numbered lists signal order-dependency; bullets signal any-order sets

Explanation: Bullets (-) signal set — agent may address in any efficient order. Numbers (1.) signal sequence — step 3 assumes step 2 happened. Using numbered lists for sequential steps is the course's explicit remedy for order-dependency failures.

Q63. The course says HTML 'has almost no ceiling' while Markdown's 'ceiling is low.' The SPECIFIC ceiling mentioned for Markdown is:

A) Markdown cannot represent more than 6 heading levels

B) Markdown documents cannot exceed 100KB without performance issues

C) Markdown is limited to headings, lists, tables, and not much else — it resorts to ASCII art for diagrams and layouts

D) Markdown cannot be processed by AI models longer than 4,096 tokens

✔ Correct Answer: C) Markdown is limited to headings, lists, tables, and not much else — it resorts to ASCII art for diagrams and layouts

Explanation: The course directly quotes: 'Markdown's ceiling is low. Headings, lists, tables, and not much else. When an agent needs to show you a workflow, a color palette, or a layout, Markdown resorts to ASCII art: charming, but a workaround.' HTML removes this ceiling with SVG, interactive controls, and styled layouts.

Q64. An agent produces an HTML artifact that is 8 MB due to embedded images and external font libraries. The recipient is a parent opening it on a phone on a slow connection. Which spec constraint from the course would have prevented this?

A) 'No tables or charts — use bullet lists only'

B) 'Lightweight, single file, no external fonts or libraries, keep it under 200 KB'

C) 'Output as Markdown only, then convert to HTML on the recipient's device'

D) 'Use a CDN for all external assets to improve load speed'

✔ Correct Answer: B) 'Lightweight, single file, no external fonts or libraries, keep it under 200 KB'

Explanation: The course explicitly mentions this exact scenario: heavy artifacts can balloon to several megabytes, a real problem on slow mobile connections. The solution is a spec constraint: 'lightweight, single file, no external fonts or libraries, keep it under 200 KB.'

Q65. The course says GitHub Pages is the correct publishing choice when:

A) You need to share the page in the next 60 seconds with no setup

B) You want drag-and-drop simplicity with a permanent web address

C) The page should stay online permanently, and you want full control — including possibly connecting your own domain name

D) The page contains private information that should not be indexed by search engines

✔ Correct Answer: C) The page should stay online permanently, and you want full control — including possibly connecting your own domain name

Explanation: GitHub Pages = permanent, stable, fully controlled, your own domain possible. Claude.ai publish = instant, zero setup, lives on Claude's servers. Netlify = permanent, drag-and-drop simplicity. GitHub Gist = keep the file, edit later. Match the publishing option to the durability and control needed.

Q66. The course says 'an agent is a next-token predictor, given tools, running the predict-act-observe loop.' The 'observe' step in this loop is MECHANICALLY:

A) A separate evaluation module that scores the tool's output for quality

B) The tool's result landing back in the context window, where the model reads it as part of the next prediction

C) A human review step required between every tool call

D) The model's internal reasoning about whether the tool call was appropriate

✔ Correct Answer: B) The tool's result landing back in the context window, where the model reads it as part of the next prediction

Explanation: Observe = result lands in the context window. The model has no separate evaluation module. It simply predicts the next action with the tool result now on its reading desk — the same prediction mechanism, with new content available.

Q67. A user runs the 'Karakush test' (asking for rules of an invented game) and the AI produces detailed, confident rules. The user then runs the same test with three different AI models. Two produce different invented rules; one says it doesn't recognize the game. Which observation from Idea 7 explains the variation?

A) Different models were trained at different knowledge cutoffs

B) Different models have differently-shaped jagged frontiers — one catches what another drops, which is why the course recommends testing the same task across two or three models

C) The model that refused was censored to avoid copyright issues with real board games

D) Longer context windows in some models allow better fact-checking

✔ Correct Answer: B) Different models have differently-shaped jagged frontiers — one catches what another drops, which is why the course recommends testing the same task across two or three models

Explanation: Idea 7: 'Different models have differently-shaped frontiers; one catches what another drops.' Testing the same task across models is a habit the course recommends for exactly this reason — the honest refusal came from a model with a different frontier shape, not different facts.

Q68. The course describes the 'three modes of delivery' (Read, Tutor, Build). How does this map onto the Markdown In, HTML Out framework?

A) Read uses HTML, Tutor uses plain text, Build uses JSON

B) All three use Markdown internally for source; HTML is the output for human reading; Markdown is the medium for agent-to-agent and agent-to-builder communication in all three

C) Read uses Markdown, Tutor uses HTML, Build uses PDF

D) The three modes use completely different format conventions not covered in the Markdown course

✔ Correct Answer: B) All three use Markdown internally for source; HTML is the output for human reading; Markdown is the medium for agent-to-agent and agent-to-builder communication in all three

Explanation: The framework is consistent across modes: Markdown is the source/spec language (human to agent, agent to agent). HTML is the output for human reading (reports, explainers, sharing). Tutor teaches in *the same framework; Build uses specs written in Markdown.*

Q69. The course warns that 'Markdown's killer feature is fading.' The killer feature was:

A) Markdown's ability to render in any browser without styling

B) Human readability — humans could easily hand-edit Markdown outputs, but increasingly the agent does the editing, making human-editability less important

C) Markdown's token efficiency, which is now matched by HTML compression

D) Markdown's security — it cannot execute scripts, making it safer than HTML

✔ Correct Answer: B) Human readability — humans could easily hand-edit Markdown outputs, but increasingly the agent does the editing, making human-editability less important

Explanation: The course's exact reasoning: 'Markdown was great because humans could easily hand-edit it. But increasingly you do not hand-edit agent output; you prompt the agent to edit it. Once the agent does the editing, the format only has to suit the reader.'

Q70. A finance lead produces a quarterly report as an HTML artifact and publishes it as a link. Her CFO says 'I need to add my commentary and send it to the board as a PDF.' What is the CORRECT workflow?

A) Ask the AI to convert the HTML artifact directly to a PDF with CFO commentary already integrated

B) Tell the AI: 'Export this as a Word document (.docx) so the CFO can add Track Changes commentary, then generate a final PDF for board submission'

C) Screenshot the HTML page and attach it as an image to an email

D) Ask the CFO to use browser print-to-PDF, which will preserve all HTML formatting

✔ Correct Answer: B) Tell the AI: 'Export this as a Word document (.docx) so the CFO can add Track Changes commentary, then generate a final PDF for board submission'

Explanation: Two needs, two steps: editing → Word (CFO adds commentary with Track Changes); final submission → PDF (looks the same everywhere, cannot be accidentally changed). The source of truth remains; the office formats are the last-mile exports for specific recipients.

Q71. The course says Markdown specifications are 'where the written-down intentions live.' How does this connect to the broader Agent Factory concept of 'spec-driven development'?

A) Spec-driven development means the AI writes the specification and the human approves it

B) The Markdown spec IS the spec-driven development artifact — you write a clear, graded specification first, then have the AI build to it, making the spec the single source of truth for all downstream outputs

C) Spec-driven development requires JSON schemas, not Markdown

D) The Markdown spec is only used during development, not in production deployments

✔ Correct Answer: B) The Markdown spec IS the spec-driven development artifact — you write a clear, graded specification first, then have the AI build to it, making the spec the single source of truth for all downstream outputs

Explanation: The course explicitly connects: 'the book calls building this way spec-driven development.' The Markdown skeleton (Goal/Requirements/Hard constraints/Out of scope/Expected output) is the practical implementation of the Agent Factory's spec-driven methodology.

Q72. Idea 3 says 'hallucination is the machine working exactly as built.' The Markdown course's 'Expected output' section and 'Out of scope' section BOTH reduce hallucination

risk by:

A) Training the model to avoid hallucination on specific topics

B) Giving the model a dense, concrete target to predict toward (Expected output) and explicitly removing the temptation to invent unrequested features (Out of scope) — reducing the open-ended prediction space where plausible-but-wrong continuations emerge

C) Adding a verification step that checks the output against trusted sources

D) Limiting the model to only producing text that appeared verbatim in its training data

✔ Correct Answer: B) Giving the model a dense, concrete target to predict toward (Expected output) and explicitly removing the temptation to invent unrequested features (Out of scope) — reducing the open-ended prediction space where plausible-but-wrong continuations emerge

Explanation: Expected output: a concrete example narrows the prediction target — less open space for guessing. Out of scope: removes the invitation for unrequested features — less scope for plausible-but-wrong elaboration. Both reduce the open-ended continuation space where hallucination thrives.

Q73. The course says 'briefing works because it is the literal act of putting information into the only place the machine can read it.' This refers to:

A) Sending information to the model's training server

B) The context window — the model's entire world for one response, the only channel for the specifics of your situation

C) The system prompt layer, which is more reliable than user messages

D) The tool-calling interface, which provides structured data more accurately than prose

✔ Correct Answer: B) The context window — the model's entire world for one response, the only channel for the specifics of your situation

Explanation: Idea 5: the context window is the only place the model can get information about your specific situation. Briefing = putting your information on the reading desk. An un-briefed model isn't lazy; it genuinely has nothing in front of it relevant to your situation.

Q74. A user notices the AI's answer changes significantly when they rephrase 'Is X a good idea?' as 'Evaluate X: give the strongest case for and against.' This change WORKS because:

A) Longer prompts always produce better quality answers

B) Neutral framing removes the signal of the answer the user seems to want, neutralizing the trained-in sycophantic lean from Idea 6

C) The word 'evaluate' unlocks a different reasoning pathway in the model

D) Asking for two sides doubles the token budget the model uses for the answer

✔ Correct Answer: B) Neutral framing removes the signal of the answer the user seems to want, neutralizing the trained-in sycophantic lean from Idea 6

Explanation: Idea 6 + Concept 6 (AI Prompting): 'Ask isn't X true? and you have signalled the answer you want; the trained-in lean supplies it.' Neutral framing removes this signal. 'Evaluate X; give the strongest case on each side' does not hint at a preferred conclusion.

Q75. The course describes 'design prototypes' as HTML Pattern 4. The SPECIFIC problem this pattern solves that text conversation cannot:

A) It allows the AI to access design software APIs directly

B) Visual preferences (shade of blue, corner roundness, font size) are nearly impossible to describe precisely in words — sliders and swatches let you adjust visually and export exact parameters as text

C) It generates multiple designs simultaneously from different models

D) It connects to Figma or Canva to produce professional design files

✔ Correct Answer: B) Visual preferences (shade of blue, corner roundness, font size) are nearly impossible to describe precisely in words — sliders and swatches let you adjust visually and export exact parameters as text

Explanation: 'Some things are miserable to express in words: a color, an animation speed, the priority order of thirty items.' Pattern 4 builds live controls so you adjust visually, then click 'copy parameters' to get precise text. Ten seconds of dragging replaces ten rounds of 'try again, a bit warmer.'

Q76. The course says 'the four questions to include in every HTML request' are: who reads it, what it contains, whether it needs interactivity, and how it should be read. The LAST question ('how it should be read') matters because:

A) It determines the file size of the output

B) 'Read once' produces a clean linear page; 'reference I revisit' produces a table of contents and collapsibles; 'presentable to a client' produces something polished enough to forward — each phrase pulls design in a different direction

C) The reading mode determines which HTML version the output will target

D) It sets the maximum number of sections in the generated page

✔ Correct Answer: B) 'Read once' produces a clean linear page; 'reference I revisit' produces a table of contents and collapsibles; 'presentable to a client' produces something polished enough to forward — each phrase pulls design in a different direction

Explanation: The course contrasts three reading modes explicitly: 'once' → clean linear; 'revisit' → TOC + collapsibles; 'client-presentable' → polished forwarding quality. The design optimizes for the reading behavior, not just the content.

Q77. The course says 'the model isn't withholding; the information was never there to freeze.' This explains why AI cannot answer questions about:

A) Topics it considers too sensitive to discuss

B) Your company's private data, your calendar, and yesterday's internal emails — none of which were in the training text when weights were frozen

C) Events that occurred before the training cutoff

D) Questions asked in languages with fewer training examples

✔ Correct Answer: B) Your company's private data, your calendar, and yesterday's internal emails — none of which were in the training text when weights were frozen

Explanation: Idea 2: private information (your company numbers, calendar, internal emails) was never in the training text, so the weights contain nothing about it. The model isn't refusing — the information was never there to freeze. This is why context-loading (Idea 5) is necessary.

Q78. The Markdown course says 'the spec is the deliverable.' This is a MINDSET SHIFT because in traditional work:

A) Specifications are written by managers, not workers

B) The specification is considered pre-work or paperwork before the real work (writing, coding, building) — the Agent Factory inverts this: the spec drives every regeneration, so improving the spec IS improving the work

C) Deliverables are always measured in lines of code or pages of output

D) Specifications are only used in software development, not general professional work

✔ Correct Answer: B) The specification is considered pre-work or paperwork before the real work (writing, coding, building) — the Agent Factory inverts this: the spec drives every regeneration, so improving the spec IS improving the work

Explanation: Traditional view: spec = paperwork. Agent Factory view: the agent's output is a function of the spec; improve the spec and every regeneration improves with it. The spec is not preparation for work — it is the primary work product that multiplies through the agent.

Q79. The course's two-minute proof (the tuition center HTML card) demonstrates which three things simultaneously?

A) Markdown syntax, HTML rendering, and cloud deployment

B) Structured input going in, a designed page coming out, and one-click sharing — the entire course in one accidental run-through

C) The five HTML patterns, the spec skeleton, and the grading technique

D) Chat, terminal, and desktop motions side by side

✔ Correct Answer: B) Structured input going in, a designed page coming out, and one-click sharing — the entire course in one accidental run-through

Explanation: The course explicitly states: 'you have already done everything this course teaches, once, by accident: a few structured lines went in, a designed page came out, and a link made it shareable.' The proof is a compressed preview of the whole course.

Q80. A teacher wants to post a quote card from a parent testimonial on LinkedIn. The course recommends:

A) Posting the raw HTML file directly to LinkedIn

B) Using Markdown with bold formatting for the quote

C) Asking the AI to design the quote card as a styled HTML page, save it as a PNG image, then attaching the PNG to the LinkedIn post — feeds show images prominently

D) Typing the quote in plain text with quotation marks and the parent's name

✔ Correct Answer: C) Asking the AI to design the quote card as a styled HTML page, save it as a PNG image, then attaching the PNG to the LinkedIn post — feeds show images prominently

Explanation: Social feeds ignore HTML in post bodies. But images show prominently. The course says: 'ask the AI to build it as a styled HTML card, then save as a high-resolution PNG image you post on LinkedIn.' You post the image, not the HTML.

Q81. The course says reasoning mode (Idea 9) generates 'extra tokens you never see' before the final answer. This hidden token generation is COSTLY because:

A) Hidden tokens are billed at a premium rate by AI providers

B) Token generation is the fundamental unit of compute and time — more tokens = more time and money, exactly why the course says to save thinking mode for genuinely hard questions

C) Hidden tokens consume context window space, leaving less for your prompt

D) The model must verify each hidden token against its training data before proceeding

✔ Correct Answer: B) Token generation is the fundamental unit of compute and time — more tokens = more time and money, exactly why the course says to save thinking mode for genuinely hard questions

Explanation: Token generation IS the computation — every token predicted takes time and money. Reasoning mode generates potentially thousands of extra tokens of working before the visible answer. The cost is real and proportional — hence the advice to reserve it for hard problems.

Q82. The course warns: 'the dangerous errors are the easy-looking tasks it quietly fails, not the hard ones you were already checking.' This is MOST relevant to which professional behaviour?

A) Only verifying AI outputs when the task involves advanced mathematics

B) Skipping verification on 'simple' AI outputs like summaries, date calculations, item counts, or cross-references — precisely the places where jagged incompetence hides

C) Running every AI output through a plagiarism checker before use

D) Asking the AI to grade its own outputs before presenting them

✔ Correct Answer: B) Skipping verification on 'simple' AI outputs like summaries, date calculations, item counts, or cross-references — precisely the places where jagged incompetence hides

Explanation: Idea 7: the jagged frontier doesn't track difficulty. The failures that cost you are the small, obvious-seeming things (count these items, get this date right) that you never think to verify. Professional discipline means verifying across the boundary, not just at the known-hard tasks.

Q83. The course draws an analogy between CSV/Excel and Markdown/HTML. Which mapping is CORRECT?

A) CSV = Markdown (compact, machine-readable source), Excel = HTML (formatted, visual, for human eyes)

B) CSV = HTML (human-readable), Excel = Markdown (machine-readable)

C) CSV = PDF (permanent), Excel = Word (editable)

D) CSV = JSON (structured), Excel = HTML (visual)

✔ Correct Answer: A) CSV = Markdown (compact, machine-readable source), Excel = HTML (formatted, visual, for human eyes)

Explanation: The course explicitly draws this parallel: 'CSV is to data what Markdown is to text: plain, simple, meant for machines and tools. Excel is to data what HTML is to text: formatted, visual, meant for a human to open and work with.'

Q84. The course says 'the model has only the first faculty.' A person who reads this and then asks the AI to 'check your own answer for errors' should expect:

A) A reliable self-correction because the model has access to its full training data to verify against

B) Partial improvement — the model may catch errors its next-token prediction notices, but it uses the same fallible process to check, so it can miss errors and hallucinate with confidence even in its correction

C) No change — the model ignores self-correction instructions

D) Perfect correction — the check instruction activates a separate verification module

✔ Correct Answer: B) Partial improvement — the model may catch errors its next-token prediction notices, but it uses the same fallible process to check, so it can miss errors and hallucinate with confidence even in its correction

Explanation: Idea 3 + Idea 9: asking the model to check its own answer is asking it to use the same prediction process to audit its own prediction. It catches many errors (especially with reasoning mode) but not all — the missing second faculty cannot be conjured by instruction.

Q85. The course introduces 'plan-then-approve' as a mandatory habit for DESKTOP motion (Cowork/OpenWork). Why is this habit SPECIFICALLY important for desktop tools but less critical for chat?

A) Desktop tools use more tokens per action and require approval to stay within budget

B) Desktop tools can actually change files on your computer — an unapproved action could modify or delete real files, so the plan-then-approve gate protects you before irreversible changes happen

C) Desktop tools are slower and need approval to continue between steps

D) Desktop tools require approval because they connect to external APIs that charge per call

✔ Correct Answer: B) Desktop tools can actually change files on your computer — an unapproved action could modify or delete real files, so the plan-then-approve gate protects you before irreversible changes happen

Explanation: The course is explicit: 'because these apps can actually change files on your computer, always ask them to show you a plan first and wait for your approval before writing anything.' Chat artifacts live in the session — safe to iterate. Desktop tools touch real files on disk.

Q86. The course frames the Markdown spec as the 'Intent Layer.' In the broader Agent Factory hierarchy, this layer is significant because:

A) It is the layer where AI models are selected for deployment

B) It is the human-written bridge between professional judgment and agent execution — the artifact that converts expertise into actionable instructions a Digital FTE follows

C) It is a security layer that prevents agents from taking unauthorized actions

D) It is stored in a central database that all agents in the organization read from

✔ Correct Answer: B) It is the human-written bridge between professional judgment and agent execution — the artifact that converts expertise into actionable instructions a Digital FTE follows

Explanation: The course says: 'Later chapters call the place these written-down intentions live the Intent Layer: human intent, written precisely enough for an agent to act on.' The Markdown spec is how expertise (Domain Expert knowledge) becomes structured instructions (Structured Specifications) — two of the four Digital FTE elements.

Q87. A user says 'I prefer HTML out, so I always ask for HTML even for two-sentence answers.' The course would classify this as:

A) Best practice — HTML is always superior to plain text

B) A misapplication — HTML uses more tokens than Markdown, and for short conversational answers it is pure overhead; match the format to the reading, not to fashion

C) Correct for chat but incorrect for terminal tools

D) Acceptable only if Open Graph tags are included

✔ Correct Answer: B) A misapplication — HTML uses more tokens than Markdown, and for short conversational answers it is pure overhead; match the format to the reading, not to fashion

Explanation: The course's honest caveat: 'HTML output uses more tokens than Markdown, and for a three-line answer it is pure overhead. The trade is worth it precisely when the output is long, visual, shared, or revisited. For quick back-and-forth in chat, plain responses remain the right default.'

Q88. The course says the five HTML patterns share a 'through-line.' What is it?

A) All five patterns require interactivity (sliders, buttons, or drag-and-drop)

B) In every pattern, Markdown (or data) goes in, HTML comes out, and anything that needs to continue the conversation comes back as text — the asymmetry as a working rhythm

C) All five patterns produce output that must be published to a public URL

D) All five patterns require at least one SVG diagram

✔ Correct Answer: B) In every pattern, Markdown (or data) goes in, HTML comes out, and anything that needs to continue the conversation comes back as text — the asymmetry as a working rhythm

Explanation: The course explicitly states: 'Notice the through-line: in every pattern, Markdown (or your data) goes in, HTML comes out, and anything that needs to continue the conversation comes back as text.' This is the Markdown-in, HTML-out asymmetry enacted as a repeating workflow rhythm.

Q89. Idea 4 explains why 'cost and length are measured in tokens, not words.' The PRACTICAL implication for a user working with a 200,000-token context window is:

A) They can store exactly 200,000 words in the context, regardless of language

B) They can store approximately 150,000 English words, but significantly fewer words in non-Latin scripts where tokens are less efficient per word

C) The token limit is a hard ceiling that stops the model from producing any further output once reached

D) Tokens and words are equivalent in all major languages used in AI systems

✔ Correct Answer: B) They can store approximately 150,000 English words, but significantly fewer words in non-Latin scripts where tokens are less efficient per word

Explanation: Roughly 3 tokens ≈ 4 words in English. But non-Latin scripts use more tokens per word (less efficient tokenization). A 200,000-token window holds ~150,000 English words but fewer Urdu/Arabic/Hindi words — the effective capacity varies by language.

Q90. The course says 'improving the spec improves every regeneration.' This principle implies that the MOST EXPENSIVE mistake in agentic work is:

A) Using the wrong HTML pattern for a given output type

B) Publishing to the wrong platform (Netlify vs GitHub Pages)

C) Starting to build from an ambiguous or incomplete spec — every output generated from it inherits the flaw, and fixing the underlying spec then requires discarding work already done

D) Choosing Claude Code over OpenCode for a simple task

✔ Correct Answer: C) Starting to build from an ambiguous or incomplete spec — every output generated from it inherits the flaw, and fixing the underlying spec then requires discarding work already done

Explanation: The course frames it explicitly: 'fixing one ambiguous bullet in a spec costs one minute. Fixing the wrong thing the agent built from that bullet costs an afternoon.' The compounding leverage is upstream — a bad spec multiplies into bad outputs across every generation.

Q91. Both courses describe 'what the machine actually is' vs 'how to use it' as a deliberate division. The 'What AI Actually Is' course stops and points to the prompting course whenever it is about to teach a habit. Why?

A) To keep course length within a target word count

B) To maintain a clean division: the mechanism course gives the 'why' once and moves on; the prompting course develops the 'how' (the habits) in depth — repetition is by design, not by accident

C) Because habits can only be learned through the terminal and desktop motions

D) Because habits require practical exercises that a crash course format cannot support

✔ Correct Answer: B) To maintain a clean division: the mechanism course gives the 'why' once and moves on; the prompting course develops the 'how' (the habits) in depth — repetition is by design, not by accident

Explanation: The course explicitly describes the split: 'This course gives the mechanism (one explanation, then it moves on); the prompting course gives the practice (the habits, in depth).' The handoff at the boundary between 'why' and 'habit' is intentional architecture.

Q92. A Legal Operations manager wants to distribute a contract template to her team. Team members will fill in client names and dates, then save as a signed PDF. What is the CORRECT format sequence?

A) HTML → screenshot as PNG → distribute

B) Markdown source → export to Word (.docx) for team editing → team converts to PDF after signing

C) PDF directly from the AI — then team edits the PDF

D) CSV → import into Word → export to PDF

✔ Correct Answer: B) Markdown source → export to Word (.docx) for team editing → team converts to PDF after signing

Explanation: Editing → Word (team fills in details using their word processor). Signing → PDF (final, unchangeable, archivable). PDF cannot be easily edited — giving it before the editing step creates friction. Markdown is the source; Word is the collaborative format; PDF is the archive format.

Q93. The course says the model treats a numbered list as a signal that 'step 3 assumes step 2 happened.' This is NOT a hard-coded rule but a consequence of:

A) The model's programming explicitly parsing list markers before processing content

B) Training on enormous amounts of text where numbered lists were used for sequences and bullets for sets — the model learned the convention statistically and follows it as a strong prediction bias

C) A special numbered-list processing module separate from the main prediction engine

D) The context window giving priority to numbered items over bullet items

✔ Correct Answer: B) Training on enormous amounts of text where numbered lists were used for sequences and bullets for sets — the model learned the convention statistically and follows it as a strong prediction bias

Explanation: The model has no hard-coded list parser — it learned the convention from training data. Numbered lists appeared in sequences; bullets appeared in sets. This statistical pattern became a strong prediction bias: numbered = follow this order. The convention is learned, not programmed.

Q94. The course distinguishes between 'teaching the AI' (which is impossible at inference time) and 'putting information on the desk' (which is always possible). A user who consistently pastes a 2-page company style guide at the start of every chat is:

A) Wasting tokens on redundant information the model already knows

B) Teaching the model their company's style, which will persist after the chat

C) Correctly supplying context — putting information on the desk so the model has the specifics it cannot look up from frozen weights, compensating for both the knowledge cutoff and the model's lack of private-world knowledge

D) Triggering a context overflow that reduces output quality

✔ Correct Answer: C) Correctly supplying context — putting information on the desk so the model has the specifics it cannot look up from frozen weights, compensating for both the knowledge cutoff and the model's lack of private-world knowledge

Explanation: Idea 2 (frozen weights, can't know private world) + Idea 5 (context window is the only channel for specifics). The style guide is private-world knowledge not in the training data. Pasting it is the correct pattern — context-loading compensates for what the frozen weights cannot contain.

Q95. The course says 'Markdown In, HTML Out' is 'just the most common pair of answers' to the question 'who reads this last?' This phrasing implies:

A) Markdown and HTML are always the correct choices for professional work

B) The format decision is always driven by the who-reads-this-last question — Markdown and HTML are frequent answers but CSV, PDF, Word, plain text, and PPTX are all valid answers in the right context

C) Markdown and HTML should be used together in every project

D) Other formats (PDF, Word) are deprecated and should not be used in AI-native workflows

✔ Correct Answer: B) The format decision is always driven by the who-reads-this-last question — Markdown and HTML are frequent answers but CSV, PDF, Word, plain text, and PPTX are all valid answers in the right context

Explanation: The course explicitly says the question is 'who reads this last?' and Markdown/HTML is 'the most common pair of answers.' The format table lists 10+ different correct answers depending on who reads it and what they will do with it — the question, not the pair, is the principle.

Q96. The course says 'a SKILL.md written for Claude Code drops into .opencode/skills/ and runs unchanged.' This fact is MOST directly enabled by which principle from the broader Agent Factory ecosystem?

A) Both tools use the same underlying language model

B) The discipline is the constant, the tool is the variable — the spec-driven, Markdown-based method is designed to be portable across harnesses, so SKILL.md is a tool-agnostic format

C) OpenCode is built on Claude Code's open-source codebase

D) Anthropic and OpenCode share a common API standard for skill files

✔ Correct Answer: B) The discipline is the constant, the tool is the variable — the spec-driven, Markdown-based method is designed to be portable across harnesses, so SKILL.md is a tool-agnostic format

Explanation: The Agent Factory's core portability principle: 'the method is the constant, the tool is the variable.' SKILL.md is Markdown — the spec language of the method. Any tool that honors the format can read it. This is why the discipline outlives any specific vendor's product.

Q97. Idea 2 explains the knowledge cutoff. A user asks a model with an August 2025 cutoff about a January 2026 event. The model produces a plausible-sounding account of the event. The user should interpret this as:

A) Evidence that the model has been secretly updated beyond its stated cutoff

B) Idea 1 in action: the model predicts the most plausible continuation for a prompt about an event it has no real knowledge of — producing a confident, fluent, and potentially entirely fabricated account

C) Idea 7: this topic falls in the 'superhuman' zone of the jagged frontier

D) A cached response from a previous user who did ask about this event

✔ Correct Answer: B) Idea 1 in action: the model predicts the most plausible continuation for a prompt about an event it has no real knowledge of — producing a confident, fluent, and potentially entirely fabricated account

Explanation: Frozen weights (Idea 2) mean events after August 2025 are not in the model. But Idea 1 means the model doesn't say 'I don't know' — it predicts the most plausible continuation. For post-cutoff events, that continuation is a confident-sounding fabrication. Idea 3 confirms no internal flag marks it as a guess.

Q98. The course's 'one sentence to keep' from each section are: 'it predicts, not retrieves' (Idea 1) and 'write Markdown precise enough for a machine; demand HTML rich enough

for a human.' Both sentences share which underlying insight?

A) Both emphasize the importance of machine speed over human readability

B) Both center on the asymmetry between machine processing (prediction, Markdown) and human experience (truth-checking, readable output) — the human's role is to supply what the machine cannot: verification and readable output design

C) Both require users to learn technical languages to work with AI

D) Both recommend avoiding HTML in professional contexts

✔ Correct Answer: B) Both center on the asymmetry between machine processing (prediction, Markdown) and human experience (truth-checking, readable output) — the human's role is to supply what the machine cannot: verification and readable output design

Explanation: Idea 1: the machine predicts; humans verify (the missing second faculty). Markdown course: machines read Markdown; humans read HTML. Both sentences encode the same asymmetry — the machine handles the mechanical middle, and the human handles the two ends: precise input and readable, verified output.

Q99. A researcher asks the AI to 'summarize this 200-page report' and receives a fluent 3-page summary. She does NOT verify a single data point in the summary before publishing. Given Idea 3, the risk she is taking is:

A) Low — summarization is in the 'superhuman' zone of the jagged frontier

B) High — the same process that produces correct summaries produces incorrect ones with no internal flag; specific numbers, citations, and claims may be subtly wrong or invented, indistinguishable from correct ones by reading alone

C) Medium — the model is accurate on summaries but occasionally misses nuance

D) Zero — the original document was in the context window, so the model had a verified source to predict from

✔ Correct Answer: B) High — the same process that produces correct summaries produces incorrect ones with no internal flag; specific numbers, citations, and claims may be subtly wrong or invented, indistinguishable from correct ones by reading alone

Explanation: Idea 3: no truth-checker. Even with the document in context (Idea 5), the model predicts a continuation — and specific numbers, citations, or claims may be wrong or interpolated with no internal flag. Fluent, confident, and wrong arrives in the same voice as fluent, confident, and right.

Q100. The nine ideas of 'What AI Actually Is' close with: 'it is a prediction machine that learned by reading and has no organ for truth, so it is fluent everywhere, reliable only where the text was thick, and you are the part that checks.' The Markdown In, HTML Out course's single sentence closes with: 'write Markdown precise enough for a machine; demand HTML rich enough for a human.' Together, these two sentences define the USER'S complete role in agentic work as:

A) Programmer (Markdown) and Designer (HTML)

B) The precision-supplier on input (Markdown spec), the verification-supplier on output (checking), and the readability-demander on output (HTML) — the human is the quality gate at both ends of the machine

C) Approver (HTML review) and Trainer (Markdown teaching)

D) Publisher (HTML) and Archivist (Markdown storage)

✔ Correct Answer: B) The precision-supplier on input (Markdown spec), the verification-supplier on output (checking), and the readability-demander on output (HTML) — the human is the quality gate at both ends of the machine

Explanation: Both closing sentences together: human writes precise Markdown (input quality gate), machine predicts and executes (middle), human demands readable HTML (output usability gate) AND verifies truth (output quality gate). The machine is the fluent middle; the human is the precision-and-verification wrapper at both ends.

CodeVerse Soban

End of Exam Paper — Best of Luck!

AI Agent Factory | Level 1 | 100 Tricky MCQs | What AI Actually Is + Markdown In, HTML Out