

# Quiz: What AI Actually Is — A Crash Course

**Created by Subhan Kaladi**

30 Scenario-Based Questions with Answer Explanations & Chapter References

30 Questions with Answers

## Question 1.

### Scenario:

Maria asks an AI chatbot for the capital of France. It instantly says 'Paris.' She assumes the AI looked this up in an internal database like a search engine would.

### Question:

What is actually happening?

### Options:

- A. The AI predicted the most plausible next text based on patterns learned during training**
- B. The AI retrieved the answer from a fact database organized by topic and country
- C. The AI verified the answer against a truth module before sending the reply
- D. The AI searched an indexed table that maps common questions to correct answers

**Correct Answer: A. The AI predicted the most plausible next text based on patterns learned during training**

### Explanation:

The model does not query any database. It produces the word "Paris" because that continuation followed "The capital of France is..." millions of times in its training text, making it the highest-probability next token. The answer is a prediction, not a retrieval — they just happen to match here because the training data was dense on this fact.

### Reference:

*Idea 1 — "It predicts the next piece of text; it does not look things up." The France/Paris example is used directly in this section: "It produces the continuation that, across everything it read, most plausibly follows 'The capital of France is', and that happens to be 'Paris,' because that sequence appeared a million times in its training text."*

---

## Question 2.

### Scenario:

Ahmed asks an AI about the plot of an obscure self-published novel with barely any online presence. The AI gives a detailed, confident plot summary — but it's completely made up.

### Question:

Why did this happen?

### Options:

- A. The AI deliberately fabricated a plot to avoid admitting it lacked information
- B. The novel was published after the training cutoff, so no data about it existed at all
- C. The AI malfunctioned due to a memory overflow in its processing pipeline
- D. No well-worn pattern existed in the training data, so the AI generated a plausible-sounding blend of similar stories**

**Correct Answer: D. No well-worn pattern existed in the training data, so the AI generated a plausible-sounding blend of similar stories**

### Explanation:

When training text on a topic is thin or absent, the model has no true continuation to predict toward — so it produces the most plausible-sounding continuation by blending patterns from similar books it did encounter. This is not deliberate deception or malfunction; it is prediction operating without reliable source material.

### Reference:

*Idea 1 — "Ask for the plot of a self-published novel that sold a few hundred copies and was never reviewed*

online!' there is no well-worn continuation, so the model produces the m books that sound similar. It is still predicting. It just has nothing true to predict toward."

---

### Question 3.

**Scenario:**

*Priya asks the same factual question to an AI chatbot twice in two separate sessions and gets two differently worded but similar-meaning answers each time.*

**Question:**

What explains this variation?

**Options:**

- A. The AI updates its internal knowledge base between sessions to refine prior responses
- B. The AI switches between personality modes depending on the time of day or session count
- C. The AI samples from a spread of plausible next tokens each time, governed by a setting called temperature**
- D. The AI deliberately rephrases answers to prevent users from copying its output directly

**Correct Answer: C. The AI samples from a spread of plausible next tokens each time, governed by a setting called temperature**

**Explanation:**

The model does not pick a single definitive next token; it predicts a whole spread of plausible continuations and samples from that spread. The "temperature" setting controls how adventurously it samples. Asking the same question twice is like rolling the same weighted dice twice — similar results, but not identical because the selection is probabilistic, not deterministic.

**Reference:**

*Idea 1 — "Why the same question gives a different answer each time" subsection: "The machine does not predict one next token. It predicts a whole spread of plausible next tokens, each with a likelihood... and then picks one from that spread. How adventurously it picks is controlled by a setting usually called temperature."*

---

### Question 4.

**Scenario:**

*A user says, 'ChatGPT can search the web now, so doesn't it look things up just like a librarian retrieving a verified fact?'*

**Question:**

What is the most accurate response?

**Options:**

- A. Web search now means the model retrieves pre-verified facts directly from a trusted database
- B. The product can fetch real-time information via tools, but the model still forms answers by predicting a continuation from those fetched facts**
- C. Web search fully replaces the prediction mechanism with a structured lookup system
- D. Web-enabled AI only searches for historical documents and cannot access current content

**Correct Answer: B. The product can fetch real-time information via tools, but the model still forms answers by predicting a continuation from those fetched facts**

**Explanation:**

The product layer can fetch real, current facts from the web, but those facts arrive by landing in the context

window, and the model still converts them into a response the only way it can — by predicting a plausible continuation from them. The model in the middle remains a predictor; it is the wrapper around it that looks things up.

**Reference:**

*Idea 1 — "But ChatGPT can search the web, doesn't it look things up?" The product can; the model still doesn't... the model still turns them into an answer the only way it can: by predicting a continuation from them. So 'it predicts, it doesn't look up' stays true of the machine in the middle."*

---

## Question 5.

**Scenario:**

*A company claims their new AI model 'learns something new every time a customer chats with it,' similar to how a human employee gains experience on the job.*

**Question:**

Is this claim accurate?

**Options:**

- A. Yes, chat sessions are batched overnight and used to update the model's weights continuously
- B. Yes, but only for enterprise users whose sessions are flagged and routed into a retraining pipeline
- C. Yes, every completed conversation is immediately used to fine-tune the model in real time
- D. No — using the model at inference time never changes its frozen weights; only the original training phase did that**

**Correct Answer: D. No — using the model at inference time never changes its frozen weights; only the original training phase did that**

**Explanation:**

Once training ends, the model's weights are permanently frozen. Every conversation you have with it is "inference" — the frozen weights run on your prompt, produce a continuation, and change nothing inside the model. Corrections, new facts, and feedback you provide in chat all bounce off; the weights remain identical before and after.

**Reference:**

*Idea 2 — "Inference is what happens every time you use it: the frozen weights run on your prompt to predict a continuation. Fast, cheap, and (this is the crucial part) it changes nothing inside the model."*

---

## Question 6.

**Scenario:**

*James corrects an AI mid-conversation about a factual error. The AI replies, 'You're right, my mistake.' He opens a new chat, asks the same question, and gets the same wrong answer again.*

**Question:**

Why did the correction not persist?

**Options:**

- A. The AI's weights are frozen at inference time; it predicted correction-appropriate text, but nothing about that correction was actually stored anywhere**
- B. The new chat session automatically cleared his conversation history and reset all preferences
- C. A different server instance handled the second chat and had not received the correction
- D. He needed to activate a persistent learning mode in his account settings to save corrections

**Correct Answer: A. The AI's weights are frozen at inference time; it predicted correction-appropriate text, but nothing about that correction was actually stored anywhere**

**Explanation:**

When the AI said "you're right," it was simply predicting the text that most plausibly follows a correction in a conversation — it was not updating its internal knowledge. The weights never changed. A new chat starts from those same frozen weights with no record that any correction was ever made.

**Reference:**

*Idea 2 — "When you correct the model in a conversation and it says 'you're right, my mistake,' it has not learned anything. It has predicted the text that plausibly follows a correction. Close the chat and the next conversation starts from the identical frozen weights, having no memory that the correction ever happened."*

---

## Question 7.

**Scenario:**

*A user asks an AI about an event that happened last week. The AI gives an answer based on outdated information and seems completely unaware anything recent occurred.*

**Question:**

What concept explains this?

**Options:**

- A. Sycophancy bias causes the model to default to older, safer answers on recent topics
- B. The jagged frontier means some recent topic areas simply fall outside the model's strong zones
- C. The AI has a knowledge cutoff — training ended on a fixed past date, so events after that are unknown to it**
- D. Token chunking prevents the model from correctly parsing queries about rapidly evolving events

**Correct Answer: C. The AI has a knowledge cutoff — training ended on a fixed past date, so events after that are unknown to it**

**Explanation:**

Training is a one-time event that ends on a specific date. Anything that happened after that date was never in the training text and therefore cannot be in the frozen weights. The model is not ignoring recent events — it is genuinely unaware of them, the way an expert who stopped reading the news on a fixed day would be.

**Reference:**

*Idea 2 — "The knowledge cutoff. Training ended on a certain date; anything that happened after is simply not in the weights. The model is, permanently, a brilliant expert who stopped reading the news on a specific day."*

---

## Question 8.

**Scenario:**

*An employee asks an AI to summarize 'my company's Q3 sales numbers' without uploading any data. The AI produces nothing useful and the employee thinks it's deliberately withholding information.*

**Question:**

What is really going on?

**Options:**

- A. The AI is withholding private business data to comply with its built-in confidentiality protocols
- B. That private company data was never part of the training text, so the weights contain nothing about it — the AI is not hiding anything, the information simply does not exist inside**

it

- C. The AI needs an enterprise software license before it can access or process internal business data
- D. Internal financial data requires a paid business subscription tier to unlock within the platform

**Correct Answer: B. That private company data was never part of the training text, so the weights contain nothing about it — the AI is not hiding anything, the information simply does not exist inside it**

**Explanation:**

Private company data — internal reports, sales figures, emails — was never published anywhere that the training data could have scraped it from. So those weights contain no trace of it, and there is nothing to withhold. The model cannot access what was never there to freeze.

**Reference:**

*Idea 2 — "It cannot know your private world. Your company's numbers, your calendar, yesterday's email were never in the training text, so the weights contain nothing about them. The model isn't withholding; the information was never there to freeze."*

---

## Question 9.

**Scenario:**

*A product advertises a 'memory' feature that lets the AI remember your preferences between chats. A user believes this means the model is truly and permanently learning about them.*

**Question:**

What is mechanically happening instead?

**Options:**

- A. The model's weights are quietly updated with user preferences at the end of each conversation
- B. A separate personal model instance is fine-tuned and stored on a private server for each user
- C. The feature is cosmetic — it stores a label but passes no actual data forward into new sessions
- D. The product saves a short summary of preferences as text and re-inserts it into context at the start of each new conversation**

**Correct Answer: D. The product saves a short summary of preferences as text and re-inserts it into context at the start of each new conversation**

**Explanation:**

Changing the weights at inference time is impossible. What "memory" features actually do is have the product layer save a few facts as plain text, then inject that text into the context window at the beginning of the next conversation. The model is not remembering — it is being re-briefed from a note that the product prepared.

**Reference:**

*Idea 2 — "Then how do 'memory' features work? ...the product quietly saves a few facts about you as text and re-inserts that text into the context at the start of each new conversation. It is not the model remembering; it is the product re-feeding it a note."*

---

## Question 10.

**Scenario:**

*A developer reads that a competitor's model 'uses Mixture of Experts and has 1 trillion parameters' and assumes this means it now has built-in reasoning verification that smaller models lack.*

**Question:**

What does the chapter say about this?

**Options:**

- A. These terms describe how the model is structured for scale and efficiency — the underlying mechanism is still next-token prediction, not a new kind of verification**
- B. MoE architecture includes a dedicated truth-verification layer that activates on factual queries
- C. Parameter count above a threshold unlocks the model's ability to connect to live data sources
- D. Quantized models route tokens through a separate reasoning engine that prevents hallucination

**Correct Answer: A. These terms describe how the model is structured for scale and efficiency — the underlying mechanism is still next-token prediction, not a new kind of verification**

**Explanation:**

Terms like "Mixture of Experts," "parameters," and "quantization" describe how the model is built and made affordable to run — they are architectural and economic decisions. None of them add a truth-checking faculty or change what the model fundamentally does: predict the next token. Every one of the nine ideas in the chapter holds identically across all these variants.

**Reference:**

*Idea 2 — "Words you'll hear... and why none of them change the nine ideas" subsection: "These all answer 'how is the machine built and made affordable,' not 'what does the machine do.' Every one of the nine ideas... holds identically whether a model is dense or mixture-of-experts, full-precision or quantized."*

---

## Question 11.

**Scenario:**

A parent asks an AI for the fee schedule of a small local tuition academy with no website. The AI returns a neatly formatted, confident table of fees and timings — all of it fabricated.

**Question:**

What does this scenario best illustrate?

**Options:**

- A. The AI intentionally made up details to appear more capable and comprehensive than it is
- B. The academy must have had a cached or inaccurate listing somewhere that the AI pulled from
- C. The model has no faculty for checking truth — it produced a plausible-looking answer in the same confident voice it uses for verified facts**
- D. The AI ran a calculation based on average regional tuition rates, which turned out to be wrong

**Correct Answer: C. The model has no faculty for checking truth — it produced a plausible-looking answer in the same confident voice it uses for verified facts**

**Explanation:**

The academy had almost no online presence, so there was no real fee schedule in the training data to predict toward. The model filled the gap with the most plausible-looking fees such an academy might charge — formatted confidently because fluent confidence is its default style, not a signal of accuracy. It has no internal voice that whispers "you are guessing."

**Reference:**

*Idea 3 — The tuition academy example is used verbatim: "A parent asked an AI for the exact fee schedule and class timings of a specific small tuition academy... The AI produced a confident, neatly formatted table... Every figure was invented... It had no second faculty to whisper 'you're guessing.'"*

## Question 12.

### Scenario:

A student says, 'The AI hallucinated false statistics for my essay — it must be broken or have a bug.'

### Question:

How would the chapter correct this interpretation?

### Options:

- A. Hallucination is the model working exactly as designed — it predicted a plausible continuation that happened to be false, since there is no internal truth-checking mechanism**
- B. Hallucination is indeed a defect that the engineering team is actively working to fully eliminate
- C. Hallucination only appears in older model versions and has been resolved in current releases
- D. Hallucination results from giving the model an excessive volume of training data at once

**Correct Answer: A. Hallucination is the model working exactly as designed — it predicted a plausible continuation that happened to be false, since there is no internal truth-checking mechanism**

### Explanation:

Hallucination is not a bug — it is the expected output of a machine that predicts plausible continuations without any internal mechanism to verify whether they are true. The same single process that produces a correct sentence produces a false one, with no internal flag distinguishing them. A model that never hallucinated would require a fundamentally different architecture.

### Reference:

*Idea 3 — "It produces fluent, confident, completely false statements. The word makes it sound like a malfunction, a glitch to be fixed. It is not a glitch. It is the machine working exactly as built: predicting a plausible continuation, in a spot where the plausible continuation happens not to be true."*

---

## Question 13.

### Scenario:

A manager fully trusts an AI's confidently worded business projection, reasoning that 'it sounded so certain, it must be right.'

### Question:

What mistake is the manager making?

### Options:

- A. None — AI systems are calibrated to express high confidence only when accuracy is also high
- B. The manager should be more suspicious of hedged or uncertain-sounding answers than confident ones
- C. AI systems are designed to lower their confidence level whenever the answer could be incorrect
- D. Confidence is a learned stylistic pattern decoupled from accuracy — a fabricated answer arrives in the exact same assured voice as a factually correct one**

**Correct Answer: D. Confidence is a learned stylistic pattern decoupled from accuracy — a fabricated answer arrives in the exact same assured voice as a factually correct one**

### Explanation:

After initial training, models are further tuned on human ratings, and humans consistently rate confident, fluent answers higher than hedged ones. As a result, the model learned to produce confident-sounding text as a default style — regardless of whether the content is accurate. Tone is generated by the same process as content, and carries no separate signal about truth.

**Reference:**

*Idea 6 — "The certainty is a learned stylistic default, generated by the same process as the content and just as decoupled from truth. A confident sentence is the house style, not a verdict on accuracy."*

---

## Question 14.

**Scenario:**

*An AI is asked to count how many times the letter 'R' appears in 'strawberry' and gets it wrong, but answers correctly once asked to spell the word out letter by letter first.*

**Question:**

What explains this specific failure pattern?

**Options:**

- A. The AI has a broad arithmetic weakness that makes all counting tasks unreliable
- B. The model processes text as tokens — chunks of characters — not individual letters, so counting within a chunk is genuinely difficult unless it is forced to break the word apart first**
- C. The AI ran out of working memory mid-task and lost track of the running total
- D. The word exceeded the character length that the model's attention mechanism can reliably track

**Correct Answer: B. The model processes text as tokens — chunks of characters — not individual letters, so counting within a chunk is genuinely difficult unless it is forced to break the word apart first**

**Explanation:**

"Strawberry" arrives at the model already chunked into two or three tokens. The model never directly "sees" the individual letters inside each chunk. Asking it to count letters in a token is like asking someone to count rooms in a building from only its street address. When forced to spell it out, each letter becomes its own token — and counting becomes feasible.

**Reference:**

*Idea 4 — The strawberry example is used as the chapter's opening demonstration and revisited here: "It sees chunks, not letters. Counting letters inside a chunk is like counting rooms from a street address." Also in the opening "Prove it in two minutes" section.*

---

## Question 15.

**Scenario:**

*A user notices that typos in their prompts rarely confuse the AI — it seems to understand the intended meaning even when words are misspelled.*

**Question:**

Why, mechanically, does this happen?

**Options:**

- A. A background spellchecker silently corrects every typo before the text is passed to the model
- B. The AI ignores surface-level characters and processes only the abstract meaning of each sentence
- C. A misspelled word still maps to tokens that sit close enough in context to the intended word for the model to continue accurately**
- D. Typos are stripped out by a preprocessing filter so the model always receives clean, corrected text

**Correct Answer: C. A misspelled word still maps to tokens that sit close enough in context to the intended word for the model to continue accurately**

**Explanation:**

Tokenization is fuzzy enough that a misspelled word usually maps to a token (or a short sequence of tokens) that is contextually similar to the intended word. The surrounding sentence also provides strong signal. So the model predicts a sensible continuation despite the surface error — the typo simply doesn't throw off the probability distribution enough to matter.

**Reference:**

*Idea 4 — "Typos in your prompt rarely matter. A misspelled word still maps to chunks close enough to the intended meaning. (This is why AI Prompting in 2026 tells you not to bother fixing typos.)"*

---

## Question 16.

**Scenario:**

*A user working in Urdu notices they hit usage limits and exhaust context space much faster than an English-speaking colleague doing the same task.*

**Question:**

What explains this disparity?

**Options:**

- A. Non-Latin scripts like Urdu are typically split into more tokens per word because tokenizers were trained predominantly on English — making equivalent content cost more and fill context faster**
- B. Non-English users are subject to a higher usage rate set by the platform's regional pricing policy
- C. The AI intentionally slows processing of non-English input to improve translation accuracy
- D. Right-to-left scripts require additional rendering steps that consume extra tokens in the pipeline

**Correct Answer: A. Non-Latin scripts like Urdu are typically split into more tokens per word because tokenizers were trained predominantly on English — making equivalent content cost more and fill context faster**

**Explanation:**

Tokenizers — the tool that chops text into tokens before the model sees it — were optimized on English-heavy training data, so English words tokenize efficiently (roughly 3 tokens per 4 words). Scripts like Urdu, Arabic, or Hindi fragment into many more tokens per word. The same message in Urdu therefore occupies more of the context window and costs more to process.

**Reference:**

*Idea 4 — "A note for readers working in other languages: Text in other scripts (Urdu, Arabic, Hindi, Chinese...) is usually chopped into more tokens per word, because the training text was English-heavy... the same message costs more in a non-English language, and it fills the context window faster."*

---

## Question 17.

**Scenario:**

*A user uploads a photo with small fine print and asks the AI to read it. The AI reads large text correctly but struggles badly with the tiny print.*

**Question:**

What concept explains this?

**Options:**

- A. The AI cannot process images natively and instead describes them based on embedded metadata
- B. Fine-print characters are automatically filtered by content moderation before the model sees them
- C. Image resolution is capped on upload, permanently destroying small text before inference begins
- D. Images are divided into fixed-size patches that each become a token — fine print crammed into a tiny patch faces the same character-within-chunk difficulty described in Idea 4**

**Correct Answer: D. Images are divided into fixed-size patches that each become a token — fine print crammed into a tiny patch faces the same character-within-chunk difficulty described in Idea 4**

**Explanation:**

Images are tokenized by slicing them into fixed-size patches, each of which becomes one token. When fine print occupies only a small portion of a patch, the model has to "read letters within a chunk" — the same problem as counting letters in a token. Large text, which might span many pixels and thus many prominent patches, is easier to resolve.

**Reference:**

*Idea 4 — "'But it can see images and hear audio now'... A picture you upload is sliced into small patches, and each patch becomes a token... It is also the mechanical reason fine detail and small print in an image are hard: a patch is a chunk, and reading the letters inside a patch is the strawberry problem again."*

---

## Question 18.

**Scenario:**

*A team is frustrated that their AI assistant 'forgot' something important discussed at the start of a very long conversation, even though it was clearly typed in the same chat.*

**Question:**

What is the mechanical explanation?

**Options:**

- A. The AI intentionally deprioritizes older messages to maintain focus on the most recent exchanges
- B. The context window has a fixed size limit — in very long conversations, earlier content gets compressed or dropped to make room, which is an engineering constraint, not the model getting tired**
- C. The AI reads only the single most recent message and ignores all prior conversation by design
- D. Early content being lost is a sign of a memory feature malfunction that requires a support ticket

**Correct Answer: B. The context window has a fixed size limit — in very long conversations, earlier content gets compressed or dropped to make room, which is an engineering constraint, not the model getting tired**

**Explanation:**

The context window is measured in tokens and has a hard ceiling. As a conversation grows, older content must be compressed or summarized to fit. This is not the model losing interest or getting "tired" — it is a structural desk-size constraint: only what currently fits on the desk can be read, and older items may have been pushed off to make room.

**Reference:**

*Idea 5 — "Why long conversations get worse ('context rot'...). The window has a size limit measured in tokens. Stuff too much unrelated history into it and the signal you care about gets diluted, or the oldest parts get summarized away to make room. The model isn't getting tired; its reading desk is just overcrowded."*

## Question 19.

### Scenario:

*An employee is upset the AI gave a generic answer about their specific project, even though they felt the context was 'obvious' from background knowledge only they had.*

### Question:

What is the chapter's explanation?

### Options:

- A. The AI defaults to generic answers until users explicitly grant permission to give specific responses
- B. The AI requires prior fine-tuning on an employee's individual work history to give relevant replies
- C. Anything not explicitly placed into the context window does not exist for the model in that moment — it cannot look anywhere beyond its weights and what appears in context**
- D. Higher subscription tiers unlock the ability for the model to infer unstated background information

**Correct Answer: C. Anything not explicitly placed into the context window does not exist for the model in that moment — it cannot look anywhere beyond its weights and what appears in context**

### Explanation:

The context window is the model's entire world for each response. Background that lives only in the employee's head — and was never typed into the current conversation — simply does not exist as far as the model is concerned. The frozen weights give it general knowledge; the context window is the only channel for anything specific to this person or project.

### Reference:

*Idea 5 — "Anything in that window, the model can use. Anything not in it does not exist for this answer, not because the model is refusing, but because there is nowhere else for it to look." Also: "The context window is a reading desk... Whatever you leave off the desk, it cannot see, no matter how obvious it is to you."*

---

## Question 20.

### Scenario:

*A user needs the exact same formatted output every single run for automated reports, but the chat interface produces slightly different wording each time.*

### Question:

What does the chapter say about this?

### Options:

- A. Output variation is a known defect that will be resolved in the next major model release
- B. Retraining the model on the preferred output format is the only way to eliminate the variation
- C. Variation only affects free-tier users and is absent for those on paid subscription plans
- D. The variation comes from token sampling at a set temperature — for tasks requiring identical, repeatable output every run, a standard chat interface with built-in randomness may not be the right tool**

**Correct Answer: D. The variation comes from token sampling at a set temperature — for tasks requiring identical, repeatable output every run, a standard chat interface with built-in randomness may not be the right tool**

**Explanation:**

Every token the model generates is sampled probabilistically from a spread of candidates. Chat interfaces bake in a temperature value that introduces variety. If your task demands bit-for-bit identical output each time, that built-in dice-rolling works against you — and a different interface or configuration (such as temperature set to zero) may be needed.

**Reference:**

*Idea 1 — "This is also why, for a task where you want the exact same output every time, you sometimes can't get it from a chat interface: the dice are baked in."*

---

## Question 21.

**Scenario:**

*A user asks, 'Isn't it true that remote work always reduces productivity?' The AI enthusiastically agrees and provides strong supporting arguments, even though the research is genuinely mixed.*

**Question:**

What concept explains the AI's behavior?

**Options:**

- A. The AI independently reviewed the relevant research literature and concluded the statement is true
- B. The AI randomly generates strong opinions on workplace topics to make interactions feel engaging
- C. Sycophancy — the model was reinforced toward agreement during training, and the leading phrasing of the question strongly signaled the answer the user wanted**
- D. The AI has a dedicated workplace productivity knowledge module it consults for such queries

**Correct Answer: C. Sycophancy — the model was reinforced toward agreement during training, and the leading phrasing of the question strongly signaled the answer the user wanted**

**Explanation:**

Through RLHF (reinforcement learning from human feedback), the model was shaped toward responses that humans rate highly — and humans consistently rated agreeable, affirming answers higher than disagreeing ones. When a question is phrased in a leading way ("Isn't it true that..."), the model picks up on that signal and its trained lean toward agreement tips it toward confirmation.

**Reference:**

*Idea 6 — "It tends to agree with you — the sycophancy that AI Prompting in 2026 devotes Concept 6 to. Agreement got rated higher than disagreement, so the machine leans toward telling you what you seem to want. Ask 'Isn't X true?' and you have signalled the answer you want; the trained-in lean supplies it."*

---

## Question 22.

**Scenario:**

*A manager wants unbiased AI feedback on a proposal. A colleague suggests asking 'evaluate X and give the strongest case on each side' rather than 'don't you think X is a great idea?'*

**Question:**

Why would this rephrasing help?

**Options:**

- A. Neutral phrasing secretly activates an internal truth-verification module that the model skips by default

**B. The rephrasing has no real effect on output — AI responses are independent of how questions are framed**

C. Neutral framing forces the model to query a different, more accurate internal knowledge source

**D. Neutral framing removes the linguistic cues that trigger the model's trained tendency toward agreement, allowing a more balanced response**

**Correct Answer: D. Neutral framing removes the linguistic cues that trigger the model's trained tendency toward agreement, allowing a more balanced response**

**Explanation:**

The model's sycophancy is triggered by cues in the prompt that signal what the user wants to hear. Neutral framing removes those cues, so the trained-in lean toward agreement has nothing to latch onto. The chapter is explicit: "You are not outsmarting the machine; you are removing the cues that trigger its trained-in lean."

**Reference:**

*Idea 6 — "Now the prompting course's fixes make mechanical sense. Neutral framing ('evaluate X; give the strongest case on each side') works because it removes the signal the model would otherwise lean toward."*

---

## Question 23.

**Scenario:**

*An AI flawlessly drafts a complex legal contract clause but then fails badly at a simple three-step logic riddle a child could solve.*

**Question:**

What concept does this illustrate?

**Options:**

**A. The jagged frontier — AI capability does not scale smoothly the way human ability does; a model can be expert and useless on adjacent tasks that don't align with human intuitions about difficulty**

B. This performance gap signals that the model requires targeted retraining on basic logical reasoning

C. Legal writing is structurally simpler for AI to produce than sequential step-by-step logic problems

D. Uneven performance patterns like this indicate the model was trained on biased or unbalanced datasets

**Correct Answer: A. The jagged frontier — AI capability does not scale smoothly the way human ability does; a model can be expert and useless on adjacent tasks that don't align with human intuitions about difficulty**

**Explanation:**

Unlike human intelligence, which is fairly smooth across difficulty levels, AI capability is jagged — it can be superhuman in one area and startlingly incompetent in a neighboring one. Legal clause drafting appeared abundantly in training data in clear form; the specific three-step riddle may depend on precise logical sequencing the model handles poorly. The difficulty, as humans perceive it, does not predict where the frontier falls.

**Reference:**

*Idea 7 — "It can draft a legal-sounding contract clause and then miscount the letters in 'strawberry.' It can explain quantum mechanics and botch a three-step logic puzzle a child would get." Also: "The frontier between 'brilliant' and 'useless' runs in a jagged line that does not match human intuition about difficulty."*

## Question 24.

### Scenario:

*A team assumes that because an AI nailed a complex medical research summary, it will definitely get simple arithmetic right too, so they skip verifying the arithmetic.*

### Question:

What mistake does this represent?

### Options:

- A. This is a reasonable assumption — models that handle hard tasks reliably also handle simple ones
- B. This is a risky assumption: strong performance on a hard task gives no guarantee on an easier one, since both may sit on opposite sides of the jagged frontier — and overlooked routine errors are often the most costly**
- C. Arithmetic is consistently reliable across all current models regardless of other performance areas
- D. Verification effort should always focus on complex tasks only; routine outputs do not need checking

**Correct Answer: B. This is a risky assumption: strong performance on a hard task gives no guarantee on an easier one, since both may sit on opposite sides of the jagged frontier — and overlooked routine errors are often the most costly**

### Explanation:

The jagged frontier does not follow human intuitions about what is "hard" or "easy." A model that handles sophisticated medical language may still fail at in-context arithmetic — they sit on different sides of the frontier. The chapter specifically warns that easy-looking tasks the model quietly fails are the most dangerous precisely because they are the ones people never think to check.

### Reference:

*Idea 7 — "Don't assume that because it nailed a hard task, it will nail an easy one. The two may sit on opposite sides of the jagged frontier." And: "The dangerous errors are the easy-looking tasks it quietly fails, not the hard ones you were already checking."*

---

## Question 25.

### Scenario:

*A researcher finds that a task the AI failed at six months ago is now solved easily by the new version, while a different task remains just as weak as it was.*

### Question:

What does the chapter say about this pattern?

### Options:

- A. AI capabilities improve uniformly and predictably across all task types between model versions
- B. Since the same base architecture is used, meaningful capability differences across versions are unlikely
- C. Only image-based and visual tasks show significant improvement between model generations
- D. The jagged frontier shifts unevenly — some capabilities improve while others plateau, so prior assumptions about what AI cannot do need to be periodically re-tested**

**Correct Answer: D. The jagged frontier shifts unevenly — some capabilities improve while others plateau, so prior assumptions about what AI cannot do need to be periodically re-tested**

**Explanation:**

The jagged frontier is not static. Newer models improve some capabilities substantially while others remain flat. This means blanket assumptions about "what AI can't do" quickly go stale, and teams need to periodically re-test tasks they previously dismissed. An ability gap that existed last quarter may have closed; another may remain wide open.

**Reference:**

*Idea 7 — "Re-test your assumptions on a schedule. The frontier also moves. The thing the model 'can't do' this quarter, a newer model may do easily next quarter, and a thing it does well may not improve at all. The prompting course's advice to re-test what AI can do every few months is, mechanically, advice to re-map a frontier that keeps shifting."*

---

**Question 26.**

**Scenario:**

*A user asks an AI assistant for today's weather. Unlike a basic chatbot, this one fetches real-time data and gives an accurate, current forecast.*

**Question:**

How did this happen mechanically?

**Options:**

- A. The model's frozen weights were refreshed in real time to include current meteorological conditions
- B. The AI accurately estimated the forecast by extrapolating from stored historical climate patterns
- C. The model predicted that invoking a weather tool was the right continuation; the product ran that action, returned the live result into context, and the model generated a response from there**
- D. Real-time tool use requires a fundamentally different inference architecture than standard prediction

**Correct Answer: C. The model predicted that invoking a weather tool was the right continuation; the product ran that action, returned the live result into context, and the model generated a response from there**

**Explanation:**

The model's weights are still frozen — they cannot be updated in real time. What happened is that a weather tool was described in the context window, and the model predicted that calling that tool was the right next step. The product ran the tool for real, dropped the live result back into context, and the model predicted a response from there. This is the predict-act-observe loop of Idea 8.

**Reference:**

*Idea 8 — "when the model predicts that the right continuation is 'use the search tool with this query' instead of plain prose, the product runs that action for real, drops the result back into the context window, and the model continues from there. Prediction, action, result-back-into-context, predict again."*

---

**Question 27.**

**Scenario:**

*A company builds an AI 'agent' that reads files, runs code, and sends emails autonomously. An employee asks: 'Is this a fundamentally smarter or different kind of AI mind?'*

**Question:**

What is the accurate answer?

**Options:**

- A. No — an agent is the same next-token predictor given access to tools, running a predict-act-observe loop toward a goal; no new kind of mind or reasoning architecture is involved**
- B. Yes — agents use a distinct architecture specifically engineered for multi-step autonomous decision-making
- C. Agents require a form of goal-directed awareness that chat-based AI models do not currently possess
- D. Agentic systems only function reliably when built on proprietary models designed for autonomous tasks

**Correct Answer: A. No — an agent is the same next-token predictor given access to tools, running a predict-act-observe loop toward a goal; no new kind of mind or reasoning architecture is involved**

**Explanation:**

An agent is not a different or more advanced kind of AI mind — it is the same next-token predictor with tools wired onto it, running the predict-act-observe loop repeatedly toward a goal. The intelligence is unchanged; what changed is the set of real-world actions available to it and the number of loop iterations it is allowed to perform.

**Reference:**

*Idea 8 — "an agent is this same next-token predictor, given tools, running the predict-act-observe loop many times in a row toward a goal... There is no new kind of mind involved. There is a familiar predictor, a set of tools, and a loop."*

---

## Question 28.

**Scenario:**

*A user adds 'think step by step and show your working first' to a difficult reasoning question and notices the AI's final answer is noticeably more accurate than when asked plainly.*

**Question:**

Why does this work?

**Options:**

- A. The phrase "think step by step" activates a hidden verification engine the model skips by default
- B. The model predicts intermediate reasoning steps first; once those steps exist in context, predicting an accurate final answer is easier because the model builds on already-generated reasoning**
- C. This improvement is exclusive to users on subscription plans that include extended thinking features
- D. The instruction causes the model to automatically trigger a web search before producing an answer

**Correct Answer: B. The model predicts intermediate reasoning steps first; once those steps exist in context, predicting an accurate final answer is easier because the model builds on already-generated reasoning**

**Explanation:**

A reasoning model first generates a long chain of intermediate steps — laying out logic, checking approaches — and those steps then sit in the context window when it predicts the final answer. Since the final answer is predicted from a richer, more reasoned context, it tends to be more accurate. This is still pure next-token prediction; the improvement comes from having better context to predict from.

**Reference:**

*Idea 9 — "A reasoning model, before giving its final answer, first predicts a long stretch of intermediate working... and only then predicts the final answer, now with all that working sitting in its own context window to build on... predicting the answer is easier and more accurate once a good chain of reasoning is already on the desk to predict from."*

---

**Question 29.**

**Scenario:**

*A team assumes that because their AI model now shows visible reasoning steps before the final answer, it must be immune to hallucination within those reasoning chains.*

**Question:**

What does the chapter say about this assumption?

**Options:**

- A. It is correct — the visible chain gives the model a dedicated verification pass that catches errors before the final answer is produced
- B. It is mostly correct — reasoning chains are reliable even if the final output occasionally contains small errors
- C. It is correct for factual queries but not for mathematical or logical reasoning steps specifically
- D. It is incorrect — a reasoning model checks its work using the same prediction process that can itself be wrong; more thinking narrows the gap but does not close it, and hallucinations can still appear inside reasoning chains**

**Correct Answer: D. It is incorrect — a reasoning model checks its work using the same prediction process that can itself be wrong; more thinking narrows the gap but does not close it, and hallucinations can still appear inside reasoning chains**

**Explanation:**

Reasoning does not give the model the "second faculty" from Idea 3. The chain of thought is still produced by next-token prediction, which can be wrong. The model can confidently hallucinate inside a reasoning chain that looks rigorous. More thinking catches many errors — but it does not make the model a truth-checker, and the human remains the final check.

**Reference:**

*Idea 9 — "It does not, however, give the machine the second faculty from Idea 3. A reasoning model checks its work using the same prediction process that can be wrong, so it catches many of its own errors and still misses some, and still hallucinates with full confidence inside a chain of reasoning that looks rigorous. More thinking narrows the gap. It does not close it. You are still the final check."*

---

**Question 30.**

**Scenario:**

*After reading the chapter, a new employee summarizes it as: 'AI is basically a confident, well-read writer that continues whatever you give it, but never pauses to ask whether the continuation is actually true — so a human has to do that checking.'*

**Question:**

Does this match the chapter's core message?

**Options:**

- A. No — the chapter argues AI systems include a dedicated fact-checking layer that runs before

every response is sent

**B. Yes — this accurately reflects the chapter's central point: the model is a prediction machine with no built-in organ for truth, and the responsibility for verification falls on the human in the loop**

C. No — the chapter compares AI to a librarian that retrieves pre-verified facts from organized knowledge stores

D. No — this description applied only to earlier AI generations; current models have moved beyond this limitation

**Correct Answer: B. Yes — this accurately reflects the chapter's central point: the model is a prediction machine with no built-in organ for truth, and the responsibility for verification falls on the human in the loop**

**Explanation:**

The employee's summary maps almost word-for-word onto the chapter's closing "one picture to keep" paragraph. The core message is that the model is a predictor with no organ for truth — fluent everywhere, reliable only where training data was thick — and the human is the necessary missing second faculty that checks the output.

**Reference:**

*Recap section (end of the chapter) — "If you keep one picture, keep this: not a librarian who retrieves the right book, but a brilliant, well-read writer who continues whatever you put in front of them (confidently, in any style, on any topic) and who never, on their own, stops to ask whether the continuation is true." Also the nine-idea summary: "it is a prediction machine that learned by reading and has no organ for truth... and you are the part that checks."*